





Join a community dedicated to learning open source

The Red Hat® Learning Community is a collaborative platform for users to accelerate open source skill adoption while working with Red Hat products and experts.



Network with tens of thousands of community members



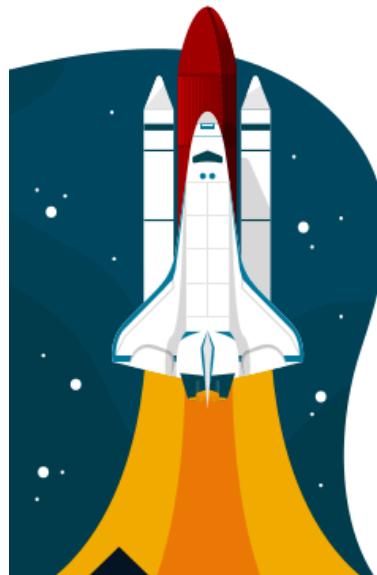
Engage in thousands of active conversations and posts



Join and interact with hundreds of certified training instructors



Unlock badges as you participate and accomplish new goals



This knowledge-sharing platform creates a space where learners can connect, ask questions, and collaborate with other open source practitioners.

Access free Red Hat training videos

Discover the latest Red Hat Training and Certification news

Connect with your instructor - and your classmates - before, after, and during your training course.

Join peers as you explore Red Hat products

Join the conversation learn.redhat.com



Copyright © 2020 Red Hat, Inc. Red Hat, Red Hat Enterprise Linux, the Red Hat logo, and Ansible are trademarks or registered trademarks of Red Hat, Inc. or its subsidiaries in the United States and other countries.



Red Hat OpenShift Administration

III: Scaling Deployments in the Enterprise

Red Hat OpenShift Container Platform 4.14 DO380
Red Hat OpenShift Administration III: Scaling Deployments in
the Enterprise
Edition 3 20250418
Publication date 20250418

Authors: Andrés Hernández, Álex Córcoles, Bernardo Gargallo,
Christopher Caillouet, François Andrieu,
Guillermo Badenes Gómez, Harpal Singh
Course Architect: Fernando Lozano
DevOps Engineer: Benjamin Marco
Editor: Julian Cable

© 2025 Red Hat, Inc.

The contents of this course and all its modules and related materials, including handouts to audience members, are © 2025 Red Hat, Inc.

No part of this publication may be stored in a retrieval system, transmitted or reproduced in any way, including, but not limited to, photocopy, photograph, magnetic, electronic or other record, without the prior written permission of Red Hat, Inc.

This instructional program, including all material provided herein, is supplied without any guarantees from Red Hat, Inc. Red Hat, Inc. assumes no liability for damages or legal action arising from the use or misuse of contents or details contained herein.

If you believe Red Hat training materials are being used, copied, or otherwise improperly distributed, please send email to training@redhat.com [mailto:training@redhat.com] or phone toll-free (USA) +1 (866) 626-2994 or +1 (919) 754-3700.

Red Hat, Red Hat Enterprise Linux, the Red Hat logo, JBoss, OpenShift, Fedora, Hibernate, Ansible, RHCA, RHCE, RHCSA, Ceph, and Gluster are trademarks or registered trademarks of Red Hat, Inc. or its subsidiaries in the United States and other countries.

Linux® is the registered trademark of Linus Torvalds in the United States and other countries.

Java® is a registered trademark of Oracle American, Inc. and/or its affiliates.

XFS® is a registered trademark of Hewlett Packard Enterprise Development LP or its subsidiaries in the United States and/or other countries.

MySQL® is a registered trademark of MySQL AB in the United States, the European Union and other countries.

Node.js® is a trademark of Joyent. Red Hat is not formally related to or endorsed by the official Joyent Node.js open source or commercial project.

The OpenStack word mark and the Square O Design, together or apart, are trademarks or registered trademarks of OpenStack Foundation, in the United States and other countries and are used with the OpenStack Foundation's permission. Red Hat, Inc. is not affiliated with, endorsed by, or sponsored by the OpenStack Foundation or the OpenStack community.

All other trademarks are the property of their respective owners.

Contributors: David Sacco, Ted Singdal

Document Conventions	ix
Admonitions	xi
Inclusive Language	x
Introduction	xi
Red Hat OpenShift Administration III: Scaling Deployments in the Enterprise	xi
Orientation to the Classroom Environment	xiii
Performing Lab Exercises	xix
1. Authentication and Identity Management	1
The OpenShift OAuth Server and Identity Providers	2
Quiz: The OpenShift OAuth Server and Identity Providers	10
LDAP Authentication and Group Synchronization	14
Guided Exercise: LDAP Authentication and Group Synchronization	27
Guided Exercise: Automate LDAP Group Synchronization	35
OIDC Authentication and Group Claims	41
Guided Exercise: OIDC Authentication and Group Claims	47
Guided Exercise: Solve User Sync Conflicts	59
Token and Client Certificate Authentication with <code>kubeconfig</code> Files	63
Guided Exercise: Token and Client Certificate Authentication with <code>kubeconfig</code> Files	72
Lab: Authentication and Identity Management	80
Summary	97
2. Backup, Restore, and Migration of Applications with OADP	99
Export and Import Application Data and Settings	100
Guided Exercise: Export and Import Application Data and Settings	117
OADP Operator Deployment and Features	133
Guided Exercise: OADP Operator Deployment and Features	146
Backup and Restore with OADP	157
Guided Exercise: Backup and Restore with OADP	179
Lab: Backup, Restore, and Migration of Applications with OADP	192
Summary	204
3. Cluster Partitioning	205
Node Pools	206
Quiz: Node Pools	212
Node Configuration with the Machine Configuration Operator	214
Guided Exercise: Node Configuration with the Machine Configuration Operator	225
Node Configuration with Special Purpose Operators	236
Guided Exercise: Node Configuration with Special Purpose Operators	239
Lab: Cluster Partitioning	245
Summary	255
4. Pod Scheduling	257
Pod Scheduling Concepts	258
Quiz: Pod Scheduling Concepts	264
Quiz: Pod Scheduling Scenarios	267
Node Selectors and Taints	271
Guided Exercise: Node Selectors and Taints	280
High Availability with Affinity Rules and Pod Disruption Budgets	292
Guided Exercise: High Availability with Affinity Rules and Pod Disruption Budgets	304
Lab: Pod Scheduling	316
Summary	327
5. OpenShift GitOps	329
GitOps for Kubernetes	330
Quiz: GitOps for Kubernetes	336
GitOps for Cluster Administration	340

Guided Exercise: GitOps for Cluster Administration	351
GitOps for Application Management	360
Guided Exercise: GitOps for Application Management	369
Lab: OpenShift GitOps	380
Summary	389
6. OpenShift Monitoring	391
Cluster Monitoring	392
Guided Exercise: Cluster Monitoring	402
Alerts and Notifications	408
Guided Exercise: Alerts and Notifications	419
Lab: OpenShift Monitoring	427
Summary	436
7. OpenShift Logging	437
Log Forwarding	438
Guided Exercise: Log Forwarding	450
Centralized Logging	461
Guided Exercise: Centralized Logging	470
Lab: OpenShift Logging	480
Summary	490
8. Comprehensive Review	491
Comprehensive Review	492
Lab: Cluster Administration	493

Document Conventions

This section describes various conventions and practices that are used throughout all Red Hat Training courses.

Admonitions

Red Hat Training courses use the following admonitions:



References

These describe where to find external documentation that is relevant to a subject.



Note

Notes are tips, shortcuts, or alternative approaches to the task at hand. Ignoring a note should have no negative consequences, but you might miss out on something that makes your life easier.



Important

Important sections provide details of information that is easily missed: configuration changes that apply only to the current session, or services that need restarting before an update applies. Ignoring these admonitions will not cause data loss, but might cause irritation and frustration.



Warning

Do not ignore warnings. Ignoring these admonitions will most likely cause data loss.

Inclusive Language

Red Hat Training is currently reviewing its use of language in various areas to help remove any potentially offensive terms. This is an ongoing process and requires alignment with the products and services that are covered in Red Hat Training courses. Red Hat appreciates your patience during this process.

Introduction

Red Hat OpenShift Administration III: Scaling Deployments in the Enterprise

This course teaches the required skills to manage OpenShift clusters at scale to productively support a growing number of stakeholders, applications, and users. You will learn how to configure and manage OpenShift clusters at scale to address increasing and special demands from applications and to ensure reliability, performance, and availability.

Course Objectives

- Configure pools of cluster nodes with special configurations, and ensure that only the intended workloads for those pools are scheduled on those nodes.
- Configure enterprise authentication and group management with legacy LDAP and cloud-native OIDC identity management systems.
- Deploy, manage, and query OpenShift Logging and configure log forwarding to external log aggregators and security information and event management (SIEM) systems.
- Automate cluster configuration and application deployment by using OpenShift GitOps.
- Troubleshoot application and cluster performance and availability issues by using OpenShift Monitoring.
- Configure and automate application-level backups by using OADP.

Audience

- Primary: Platform Engineers, System Administrators, Cloud Administrators, and other infrastructure-related IT roles who are responsible for implementing and managing infrastructure for applications.
- Secondary: Enterprise Architects, Site Reliability Engineers (SRE), DevOps Engineers, and other application-related IT roles who are responsible for designing infrastructure for applications.

Prerequisites

- Possess a current Red Hat Certified Specialist in OpenShift Administration certification, or equivalent experience and skills with managing OpenShift clusters and containerized applications, according to the topics of the DO180 and DO280 courses for version 4.12 and any later versions. Earlier 4.x releases of DO180 and DO280 are not sufficient.
- Experience with enterprise authentication, monitoring, and other IT infrastructure systems is helpful but not required.

Orientation to the Classroom Environment

In this course, the main computer system that is used for hands-on learning activities (exercises) is **workstation**.

The **workstation** machine has a standard user account, **student** with **student** as the password. Although no exercise in this course requires you to log in as **root**, if you must, then the **root** password on the **workstation** machine is **redhat**.

From the **workstation** machine, you type the **oc** commands to manage the OpenShift cluster, which comes preinstalled as part of your classroom environment.

Also from the **workstation** machine, you run the commands that are required to complete the exercises for this course.

If exercises require you to open a web browser to access any application or website, then you must use the graphical console of the **workstation** machine and use the Firefox web browser from there.



Note

During the initial start of your classroom environment, the OpenShift cluster takes more time to become fully available. The **lab** command at the beginning of each exercise checks and waits as required.

If you try to access your cluster by using either the **oc** command or the web console without first running a **lab** command, then your cluster might not yet be available. If the cluster is not available, then wait a few minutes and try again.

Log in to OpenShift from the Shell

To access your OpenShift cluster from the **workstation** machine, use the following API URL:

- <https://api.ocp4.example.com:6443>

```
[student@workstation ~]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
```

Besides the **admin** user, who has cluster administrator privileges, your OpenShift cluster also provides a **developer** user, with **developer** as the password, with no special privileges.

Accessing the OpenShift Web Console

If you prefer to use the OpenShift web console, then open a Firefox web browser on your **workstation** machine and access the following URL:

- <https://console-openshift-console.apps.ocp4.example.com>

Click **Red Hat Identity Management** and provide the login credentials for either the **admin** or the **developer** user.

The Classroom Environment

Every student gets a complete remote classroom environment. As part of that environment, every student gets a dedicated OpenShift cluster for administration tasks.

This environment contains all the necessary resources for the course. Because the course does not access resources outside the classroom environment, any failures of public resources, such as Git repositories or container image registries, should not affect the course.

The classroom environment runs entirely as virtual machines in a large Red Hat OpenStack Platform cluster, which is shared among many students.

Red Hat Training maintains many OpenStack clusters, in data centers across the globe, to provide lower latency to students from many countries.

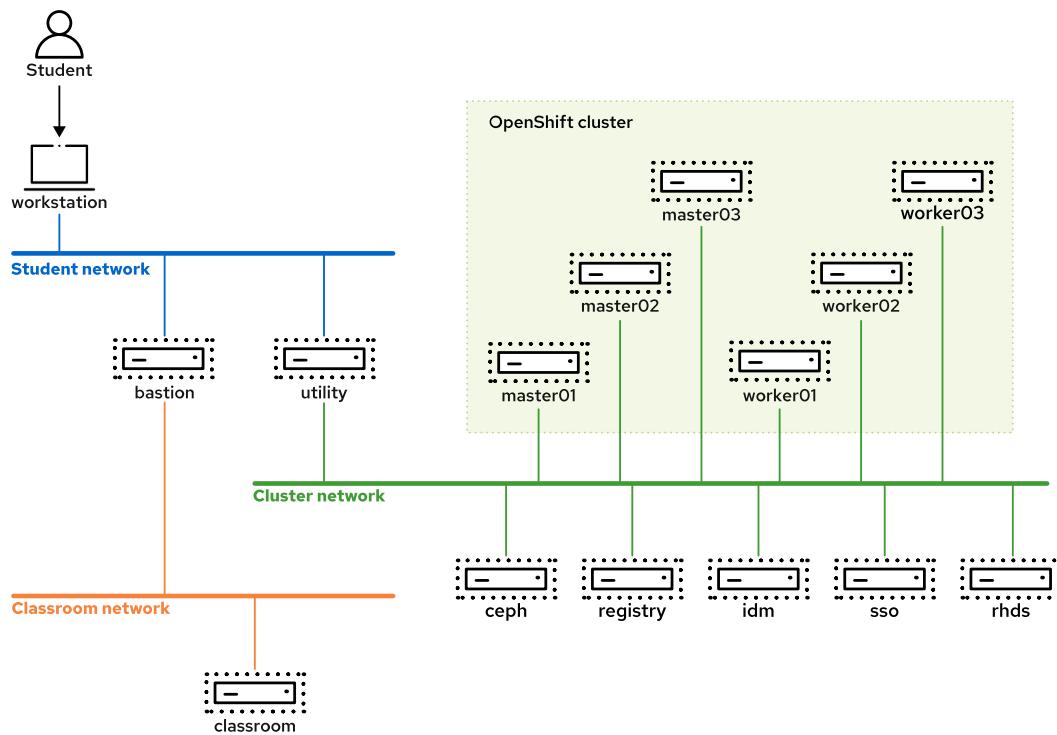


Figure 0.1: DO380 classroom architecture

All machines on the Student, Classroom, and Cluster networks run Red Hat Enterprise Linux 9 (RHEL 9), except those machines that are nodes of the OpenShift cluster. These cluster nodes run RHEL CoreOS.

The **bastion**, **utility**, **idm**, **registry**, and **classroom** systems must always be running. These systems provide infrastructure services that the classroom environment and its OpenShift cluster require. For most exercises, you are not expected to interact with any of these services directly.

Usually, the `lab` commands from the exercises access these machines to set up your environment for the exercise, and require no further action from you.

For the few exercises where you must access a system other than **workstation**, primarily the **utility** system, you receive explicit instructions and connection information as part of the exercise.

All systems in the *Student* network are in the `lab.example.com` DNS domain, and all systems in the *Classroom* network are in the `example.com` DNS domain.

The `masterXX` and `workerXX` systems are nodes of the OpenShift 4 cluster that is part of your classroom environment.

All systems in the *Cluster* network are in the `ocp4.example.com` DNS domain.

Classroom Machines

Machine name	IP addresses	Role
<code>workstation.lab.example.com</code>	172.25.250.9	Graphical workstation for system administration
<code>classroom.example.com</code>	172.25.254.254	Router to link the Classroom network to the internet
<code>bastion.lab.example.com</code>	172.25.250.254	Router to link the Student network to the Classroom network
<code>utility.lab.example.com</code>	172.25.250.253	Router to link the Student and Cluster networks
<code>ceph.ocp4.example.com</code>	192.168.50.30	Server with a Red Hat Storage Ceph preinstalled cluster
<code>idm.ocp4.example.com</code>	192.168.50.40	Red Hat Identity Management system
<code>registry.ocp4.example.com</code>	192.168.50.50	Server with Quay and GitLab
<code>sso.ocp4.example.com</code>	192.168.50.60	Red Hat Single Sign-On
<code>rhds.ocp4.example.com</code>	192.168.50.70	Red Hat Directory Server system
<code>master01.ocp4.example.com</code>	192.168.50.10	Control plane node
<code>master02.ocp4.example.com</code>	192.168.50.11	Control plane node
<code>master03.ocp4.example.com</code>	192.168.50.12	Control plane node
<code>worker01.ocp4.example.com</code>	192.168.50.13	Compute node
<code>worker02.ocp4.example.com</code>	192.168.50.14	Compute node
<code>worker03.ocp4.example.com</code>	192.168.50.15	Compute node

The Dedicated OpenShift Cluster

The Red Hat OpenShift Container Platform 4 cluster inside the classroom environment is preinstalled by using the Pre-existing Infrastructure installation method. All nodes are treated as bare metal servers, even though they are virtual machines in an OpenStack cluster.

OpenShift cloud provider integration capabilities are not enabled, and some features that depend on that integration, such as machine sets and autoscaling of cluster nodes, are not available.

Introduction

Your OpenShift cluster is in the state from running the OpenShift installer with default configurations, except for some day-2 customizations:

- The cluster provides a default storage class that is backed by a Network File System (NFS) storage provider.
- The cluster also provides storage classes that are backed by Ceph.
- The cluster uses an LDAP identity provider that is configured to use Red Hat Identity Management that runs on the `idm.ocp4.example.com` system.

The *Troubleshooting Access to your OpenShift Cluster* section provides information about how to access the **utility** machine.

Restoring Access to your OpenShift Cluster

If you suspect that you cannot log in to your OpenShift cluster as the `admin` user any more because you incorrectly changed your cluster authentication settings, then run the `lab finish` command from your current exercise and restart the exercise by running its `lab start` command.

If running a `Lab` command does not resolve the issue, then you can follow the instructions in the next section to use the **utility** machine to access your OpenShift cluster.

Troubleshooting Access to Your OpenShift Cluster

The **utility** machine ran the OpenShift installer inside your classroom environment, and it is a useful resource to troubleshoot cluster issues. You can view the installer manifests and logs in the `/home/lab/ocp4` directory of the **utility** machine.

Logging in to the **utility** server is rarely required to perform exercises. If your OpenShift cluster is taking too long to start, or is in a degraded state, then you can log in to the **utility** machine as the `lab` user to troubleshoot your classroom environment.

The `student` user on the **workstation** machine is already configured with SSH keys that enable logging in to the **utility** machine without a password.

```
[student@workstation ~]$ ssh lab@utility
```

In the **utility** machine, the `lab` user is preconfigured with a `.kube/config` file that grants access as `system:admin` without first requiring the `oc login` command.

You can then run troubleshooting commands, such as the `oc get node` command, if they fail from the **workstation** machine.

You should not require SSH access to your OpenShift cluster nodes for regular administration tasks, because OpenShift 4 provides the `oc debug` command. The `lab` user on the **utility** server is preconfigured with SSH keys to access all cluster nodes if necessary. For example:

```
[lab@utility ~]$ ssh -i ~/.ssh/lab_rsa core@master01.ocp4.example.com
```

In the preceding example, replace `master01` with the name of the intended cluster node.

Approving Node Certificates on Your OpenShift Cluster

Red Hat OpenShift Container Platform clusters are designed to run continuously, 24x7, until they are decommissioned. Unlike a production cluster, the classroom environment contains a cluster that was stopped after installation, and that stops and restarts several times before you finish this course. This scenario requires special handling unlike a production cluster.

The control plane and compute nodes in an OpenShift cluster communicate with each other. All communication between cluster nodes is protected by mutual authentication based on per-node TLS certificates.

The OpenShift installer handles creating and approving TLS certificate signing requests (CSRs) for the full stack automation installation method. The system administrator manually approves these CSRs for the Pre-existing Infrastructure installation method.

All per-node TLS certificates have a short expiration life of 24 hours (the first time) or 30 days (after renewal). When the affected cluster nodes are about to expire, they create CSRs, and the control plane automatically approves them. If the control plane is offline when the TLS certificate of a node expires, then a cluster administrator must approve the pending CSR.

The **utility** machine includes a system service that approves CSRs from the cluster when you start your classroom, to ensure that your cluster is ready when you begin the exercises. If you create or start your classroom and begin an exercise too quickly, then your cluster might not be ready. In this case, wait a few minutes while the **utility** machine handles CSRs, and then try again.

Sometimes, the **utility** machine fails to approve all the required CSRs, for example because the cluster took too long to generate all the required CSRs requests and the system service did not wait long enough. Also, some OpenShift cluster nodes might not have waited long enough for approval of their CSRs, and might issue new CSRs that superseded previous ones.

In these cases, your cluster takes too long to come up, and your `oc login` or `lab` commands fail. To resolve the problem, you can log in to the **utility** machine, as explained previously, and run the `sign.sh` script to approve any additional and pending CSRs.

```
[lab@utility ~]$ ./sign.sh
```

The `sign.sh` script loops a few times, in case that your cluster nodes issue new CSRs that supersede previous certificates.

After you approve, or the system service in the **utility** machine approves all CSRs, OpenShift must restart some cluster operators. It takes a few moments before your OpenShift cluster is ready to answer requests from clients. To help you handle this scenario, the **utility** machine provides the `wait.sh` script, which waits until your OpenShift cluster is ready to accept authentication and API requests from remote clients.

```
[lab@utility ~]$ ./wait.sh
```

In the unlikely event that neither the service on the **utility** machine nor running the `sign.sh` and `wait.sh` scripts makes your OpenShift cluster available to begin exercises, you can run troubleshooting commands from the **utility** machine, re-create your classroom, or open a customer support ticket.

Controlling Your Systems

In an instructor-led training classroom, you are assigned a physical computer (`foundationX.ilt.example.com`), which is used to access the virtual machines that run on that host. You are automatically logged in to the physical machine as the `kiosk` user, using `redhat` as the password. The classroom systems that you work with are virtual machines running on that host.

On `foundationX`, you use the `rht-vmctl` and `rht-vmview` commands to work with the virtual machines. Run the commands in the following table as the `kiosk` user on `foundationX`. These commands can be used with any of the virtual machines.

rht-vmctl Commands

Action	Command
Start the virtual machine named <code>server</code> .	<code>rht-vmctl start server</code>
Display the system console for the virtual machine named <code>server</code> , to log in and work on that system.	<code>rht-vmview view server</code>
Reset the virtual machine named <code>server</code> to its initial course state and restart it.	<code>rht-vmctl reset server</code>

At the start of a lab exercise, if the instruction "reset your servera" appears, then run the `rht-vmctl reset servera` command on the `foundationX` system as the `kiosk` user. Likewise, if the instruction "reset your workstation" appears, then run the `rht-vmctl reset workstation` command on `foundationX` as the `kiosk` user.

Performing Lab Exercises

You might see the following lab activity types in this course:

Guided exercise

A *guided exercise* is a hands-on practice exercise that follows a presentation section. It walks you through a procedure to perform, step by step.

Quiz

A *quiz* is typically used when checking knowledge-based learning, or when a hands-on activity is impractical for some other reason.

End of chapter lab

An *end-of-chapter lab* is a graded hands-on activity to help you to check your learning. You work through a set of high-level steps, based on the guided exercises in that chapter, but the steps do not walk you through every command. A solution is provided with a step-by-step walk-through.

Comprehensive review lab

A *comprehensive review lab* is used at the end of the course. It is also a graded hands-on activity, and might cover content from the entire course. You work through a specification of what to accomplish in the activity, without receiving the specific steps to do so. Again, a solution is provided with a step-by-step walk-through that meets the specification.

To prepare your lab environment at the start of each hands-on activity, run the `lab start` command with a specified activity name from the activity's instructions. Likewise, at the end of each hands-on activity, run the `lab finish` command with that same activity name to clean up after the activity. Each hands-on activity has a unique name within a course.

The syntax for running an exercise script is as follows:

```
[student@workstation ~]$ lab action exercise
```

The *action* is a choice of *start*, *grade*, or *finish*. All exercises support *start* and *finish*. Only end-of-chapter labs and comprehensive review labs support *grade*.

start

The *start* action verifies the required resources to begin an exercise. It might include configuring settings, creating resources, checking prerequisite services, and verifying necessary outcomes from previous exercises. You can perform an exercise at any time, even without performing preceding exercises.

grade

For graded activities, the *grade* action directs the `lab` command to evaluate your work, and shows a list of grading criteria with a **PASS** or **FAIL** status for each. To achieve a **PASS** status for all criteria, fix the failures and rerun the *grade* action.

finish

The *finish* action cleans up resources that were configured during the exercise. You can perform an exercise as many times as you want.

Introduction

The `lab` command supports tab completion. For example, to list all exercises that you can start, enter `lab start` and then press the Tab key twice.

Troubleshooting Lab Scripts

When you run the `lab` command with a valid exercise and action, it creates one log file in the `/tmp/log/labs` directory, and also creates the directory itself if it does not exist. The created exercise file captures any error messages from your terminal.

```
[student@workstation ~]$ ls -l /tmp/log/labs  
-rw-rw-r-- 1 student student 1024 Apr 25 12:00 comprehensive-review
```

Interpreting the Exercise Log Files

Exercise scripts send output to the log file when the scripts fail. Thus, the exercise log is useful for troubleshooting.

Although exercise scripts are always run from the `workstation` machine, they perform tasks on other systems in the course environment. Many course environments, including OpenStack and OpenShift, use a command-line tool from the `workstation` machine to communicate with server systems by using API calls.

Because script actions typically distribute tasks to many systems, additional troubleshooting is necessary to determine where a failed task occurred. Log in to those other systems and use Linux diagnostic skills to read local system log files and determine the root cause of a lab script failure.

Upgrading the Lab Package Version

The course lab package is installed on the `workstation` machine. If the instructor requests it, you can upgrade the version to fix any defect in the preinstalled lab package.

```
[student@workstation ~]$ lab upgrade DO380
```

Chapter 1

Authentication and Identity Management

Goal

Configure OpenShift clusters to authenticate by using LDAP and OIDC enterprise identity systems and to recognize groups that those systems define.

Sections

- The OpenShift OAuth Server and Identity Providers (and Quiz)
- LDAP Authentication and Group Synchronization (and Guided Exercise)
- Automate LDAP Group Synchronization (Guided Exercise)
- OIDC Authentication and Group Claims (and Guided Exercise)
- Solve User Sync Conflicts (Guided Exercise)
- Token and Client Certificate Authentication with kubeconfig Files (and Guided Exercise)
- Authentication and Identity Management

Lab

The OpenShift OAuth Server and Identity Providers

Objectives

- Define the concepts and custom resources of the OpenShift OAuth server, and explain how these resources augment Kubernetes authentication.

Authenticating API Requests

Authentication determines access to an OpenShift cluster, to ensure that only authenticated users can access the OpenShift cluster. To interact with an OpenShift cluster, users must authenticate to the Kubernetes API. After the user successfully authenticates to the cluster, the authorization layer either accepts or rejects the API request depending on the user's access permissions.

Authentication in Kubernetes

Authentication in Kubernetes verifies the signature validity for a token or a client certificate that the user provides. Kubernetes supports any authentication method, provided that the final result of the authentication method is a token or client certificate. Because Kubernetes cannot log in by any means by itself, you need extra tools to operationalize the external authentication methods. The core components for Kubernetes authentication are the API server and the authentication plug-ins. The API server is the central control plane component that is responsible for handling incoming requests. When a request reaches the API server, it forwards the request to the appropriate authentication plug-in. The authentication plug-in validates the provided credentials and determines whether the client is authenticated.

For example, Kubernetes considers to be authenticated any user that presents a valid certificate that is signed by a trusted certificate authority (CA). Kubernetes reads the `subject` field from the certificate, and assigns the username and the groups with the `common name (CN)` and the `organization (O)` fields, respectively. Kubernetes either accepts or rejects the API request for the user depending on the role-based access control (RBAC) rules that are configured for the username and the group.

Kubernetes rejects any request with an invalid access token or certificate by showing a `401 Unauthorized` error.

If the request does not present an access token or certificate, then Kubernetes assigns to the request the `system:anonymous` system user and the `system:unauthenticated` system group. After assigning the anonymous user to the request, Kubernetes uses RBAC policies to determine whether the anonymous user has appropriate access to perform the requested action. System groups are explained later in this section.

Authentication in OpenShift

OpenShift expands the Kubernetes authentication mechanisms to provide additional features and integrations with external systems. OpenShift includes its own built-in OAuth server, which is not present in Kubernetes. The OpenShift OAuth server handles user authentication, and it issues OAuth tokens that grant access to specific resources.

You can authenticate to OpenShift by using the following methods:

X.509 client certificates

Use X.509 client certificates to authenticate to the API server. The API server validates the client certificate against a trusted certificate authority. This authentication method requires an HTTPS connection to the API server. OpenShift authentication through client certificates is the same as in Kubernetes. Client certificates are explained later in this course.

OAuth access tokens

In OpenShift, the OpenShift control plane includes a built-in OAuth server that generates OAuth access tokens. OpenShift preconfigures Kubernetes to trust tokens that its own internal OAuth server issues. Thus, after you log in to the OpenShift OAuth server, Kubernetes authenticates the requests by trusting the tokens that the internal OAuth server issued.

The OpenShift OAuth Server

OpenShift provides by default the authentication operator, which runs an internal OAuth server. The OAuth server issues OAuth access tokens to users for authenticating through the API. To configure authentication, the OpenShift OAuth server provides various identity providers (IdPs) that enable integration with identity management systems such as OpenID Connect and LDAP servers. Some provided IdPs support the custom OAuth workflows from internet services, such as GitHub. The OpenShift OAuth server works only with certain supported IdPs, and you cannot add other IdPs.

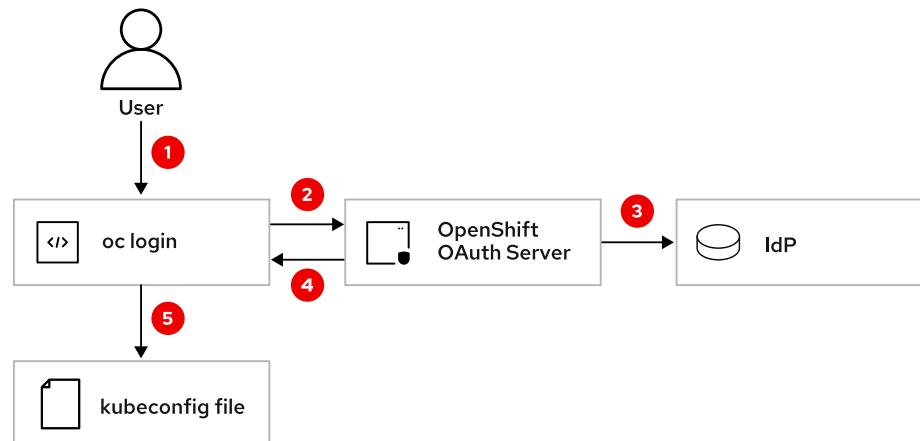


Note

A list of the supported IdPs is available at https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/authentication_and_authorization/index#supported-identity-providers

The most common way to authenticate to the OpenShift cluster is to use the `oc login` command in the command-line interface (CLI). The `oc login` command uses the OpenShift OAuth server to authenticate the user with an IdP. Then, the OpenShift OAuth server generates an OAuth access token for the user to authenticate their API requests.

The following diagram summarizes the authentication flow when you use the `oc login` command:



- ① A user runs the `oc login` command, and provides the username and password.
- ② OpenShift sends the credentials from the `oc login` command to the internal OAuth server.

Chapter 1 | Authentication and Identity Management

- ❸ The internal OAuth server validates the credentials with an IdP.
- ❹ The internal OAuth server creates an OAuth access token for the user.
- ❺ OpenShift stores the OAuth access token in the `kubeconfig` file, by default located at `~/.kube/config`.

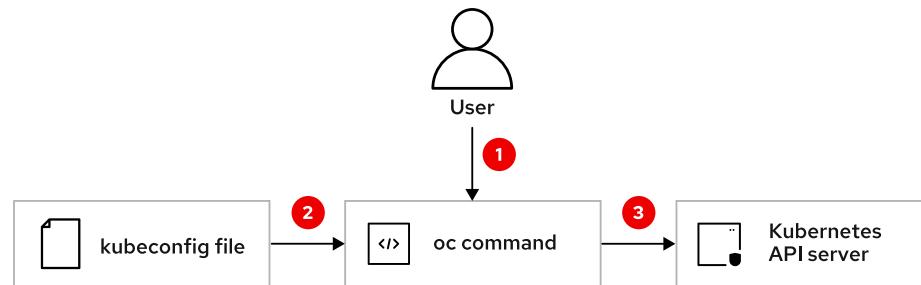
You can verify the two endpoints that the `oc login` command uses, which are the API server and the OAuth server routes, by increasing the logging level:

```
[user@host ~]$ oc login -u user -p password https://api.ocp4.example.com:6443 \
--loglevel 7
I0929 ...] HEAD https://api.ocp4.example.com:6443/
I0929 ...] Request Headers:
I0929 ...] Response Status: 403 Forbidden in 20 milliseconds
I0929 ...] GSSAPI Enabled
I0929 ...] GET https://api.ocp4.example.com:6443/.well-known/oauth-authorization-
server
I0929 ...] Request Headers:
I0929 ...]     X-Csrftoken: 1
I0929 ...] Response Status: 200 OK in 0 milliseconds
I0929 ...] using system roots as no error was encountered
I0929 ...] GET https://oauth.openshift.apps.ocp4.example.com/
oauth/authorize?client_id=openshift-challenging-
client&code_challenge=B8RP...aCtk&code_challenge_method=S256&redirect_uri=https
%3A%2F%2Foauth.openshift.apps.ocp4.example.com%2Foauth%2Ftoken
%2Fimplicit&response_type=code
...output omitted...
```

This authentication flow is similar when using the OpenShift web console instead of the CLI, except that the OpenShift web console stores the token in the browser memory instead of in a `kubeconfig` file. The two endpoints for login, which are the API server and the OAuth server, are the reason why you must accept two server certificates the first time that you log in to the OpenShift web console.

After getting the OAuth access token, every time that the user invokes the OpenShift CLI, OpenShift uses the stored credentials in the `kubeconfig` file to authenticate the user to the Kubernetes API server. The `kubeconfig` file details are explained later in this course.

The following diagram summarizes the flow when you use the OpenShift CLI after authenticating by using the `oc login` command:



- ❶ A user accesses the Kubernetes API by using the `oc` command.

Chapter 1 | Authentication and Identity Management

- ❷ OpenShift reads the credentials from the `kubeconfig` file.
- ❸ OpenShift uses the credentials from the `kubeconfig` file to authenticate to the Kubernetes API server.

You can also retrieve OAuth user access tokens from the OpenShift OAuth server by using the `namespace_route/oauth/authorize` and `namespace_route/oauth/token` endpoints.

**Note**

Requesting an OAuth access token through the OpenShift REST API is out of the scope of this course. For more information about this topic, refer to <https://access.redhat.com/solutions/6610781>

Configure the OpenShift OAuth Server

For an IdP to validate the identity of the requester, you must configure and enable at least one IdP in the OAuth server. If a user requests a new OAuth access token, then the OAuth server uses the configured IdPs to establish the requester's identity. After the user successfully logs in to the cluster, the OAuth server creates the OAuth access token for the user, and OpenShift creates the identity and user resources.

In production clusters, it is typical to configure multiple IdPs in the internal OAuth server. The following example shows the parameters for the custom resource (CR) to configure two IdPs in the OAuth server. The example uses `htpasswd` and `LDAP` as the IdPs.

**Note**

Configuring the `htpasswd` IdP for OpenShift authentication is explained in detail in the *DO280: Red Hat OpenShift Administration II: Operating a Production Kubernetes Cluster* course.

```
apiVersion: config.openshift.io/v1
kind: OAuth
metadata:
  name: cluster
spec:
  identityProviders: ❶
    - name: htpasswd-idp-name ❷
      mappingMethod: claim ❸
      type: HTPasswd
      htpasswd:
        fileData:
          name: htpasswd-secret ❹
    - name: ldap-idp-name
      mappingMethod: claim
      type: LDAP
      ldap:
        ...output omitted...
```

- ❶ The `identityProviders` array contains the information for all the IdPs that are configured in your cluster.

Chapter 1 | Authentication and Identity Management

- ❷ OpenShift prefixes the value of the identity claim with the provider name to form an identity name, and also uses the identity claim to build the redirect URL.
- ❸ The `mappingMethod` parameter controls how OpenShift establishes mappings between the IdP identities and the `User` resources.
- ❹ The OpenShift secret resource contains the users and passwords.

Users, Groups, and Identities

Users, groups, and identities play a crucial role in managing access and permissions within the cluster.

Users, Groups, and Identities in Kubernetes

Kubernetes does not have resources that represent user accounts, groups, or identities. Users and groups in Kubernetes are strings that Kubernetes reads from the client certificates or from the tokens.

For example, if you use a client certificate, then Kubernetes reads the `common name (CN)` and the `organization (O)` fields, and uses those strings for the username and the groups, respectively. Then, Kubernetes uses the username and groups to allow or deny the API request by using the RBAC policies.

For service accounts, which Kubernetes uses mainly to grant permissions to pods and to other resources to interact with the Kubernetes API server, Kubernetes adds only the `serviceaccount` string to the service account certificate or token to match RBAC policies. Service accounts are explained later in this course.

System Groups

System groups or virtual groups are predefined groups that are built into the system to grant specific permissions and access rights to certain categories of users. Kubernetes creates these groups during the setup to control access to critical components and resources within the cluster. These groups are a part of cluster roles and cluster role bindings. Thus, system groups are part of the RBAC system in Kubernetes.

system:authenticated

All authenticated users in the cluster, regardless of their role or namespace

system:authenticated:oauth

All authenticated users with an OAuth access token in the cluster

system:unauthenticated

Unauthenticated users, meaning those users who did not yet pass any authentication process

system:masters

Users with administrative access to the cluster. Members of this group have full control over all resources and actions within the cluster, including cluster-level operations and management.

The Kubernetes API server adds the strings from system groups to the user and group strings from the client certificate or from the token, whenever they apply. Then, Kubernetes uses RBAC rules to approve or deny the request depending on the user and groups.

For example, if a user does not provide an access token or certificate, then Kubernetes assigns the `system:unauthenticated` group string to the request, and RBAC policies allow or deny the request depending on the RBAC rule configuration.

Users, Groups, and Identities in OpenShift

OpenShift includes the `User` resource, which defines regular users in the cluster. Regular users typically represent human users who interact with the OpenShift cluster. The OpenShift internal OAuth server automatically creates user resources when users first log in to the cluster through an IdP. You can also manually add users to the cluster by using the Kubernetes API.

OpenShift includes the `Group` resource, which defines groups in the cluster. Groups are logical collections of users with common access requirements. OpenShift supports users as members of one or more groups, for flexible definition of access control policies. Use groups to simplify the permissions in your cluster, to grant permissions to multiple users at the same time instead of individually. The OpenShift internal OAuth server automatically creates group resources when users first log in to the cluster through an OpenID Connect (OIDC) IdP. However, you must manually manage group resources in OpenShift for most of the supported IdPs, such as from LDAP.

OpenShift includes the `Identity` resource, which defines identities in the cluster. The identity resource keeps a record of successful authentication attempts from a specific user and IdP. Identities in OpenShift refer to the external identities of users. OpenShift maps user identities from IdPs to user resources. With identities, users can log in to OpenShift with their existing credentials from IdPs.

OAuth Server Mapping Methods

When a user first logs in to the cluster through an IdP, OpenShift creates a user resource and an identity resource, and maps them. OpenShift creates the user resource by using the preferred username from the IdP.

When different IdPs are configured in the OpenShift OAuth server, more than one IdP can provide the same username for a user, which causes collisions. With the `mappingMethod` parameter in the OAuth server, you can control how OpenShift establishes mappings between the provider's identities and the user resource, and avoid those collisions.

You can use one of the following values:

claim

OpenShift provisions a user with the identity's preferred username. However, it fails if a user with that username is already mapped to another identity. This value is the default for the `mappingMethod` parameter.

add

OpenShift provisions a user with the identity's preferred username. If a user with that username exists in OpenShift, then it maps the identity to the existing user, and adds the new identity to any existing identity mappings for the user. Use this method when you configure multiple IdPs that identify the same set of users and map to the same usernames.

lookup

OpenShift looks up an existing identity, user identity mapping, and user, but does not automatically provision users or identities. Thus, with this method you must manually provision users. With this method, cluster administrators can set up identities and users manually, or by using an external process.

Authentication on a New OpenShift Cluster

OpenShift generates two authentication methods with administrative rights during the initial installation of the cluster: with the `kubeadmin` user, and with a client certificate for the `system:admin` user in a `kubeconfig` file.

**Note**

The `kubeadmin` user account is present only in self-managed OpenShift clusters. Cloud services do not use the `kubeadmin` user. Cloud services provide initial user credentials by using their cloud provider consoles and APIs in different ways.

The `kubeadmin` and `system:admin` users have the `cluster-admin` role, so OpenShift grants administrative access to these accounts by default. Use the `kubeadmin` user account during the cluster setup and initialization phase only. Red Hat recommends deleting the existing `kubeadmin` user after you create an administrator user with the `cluster-admin` role, to increase cluster security. Thus, configure an IdP in the OAuth server and create an administrator user with the `cluster-admin` role before you remove the `kubeadmin` user. Store the `system:admin` client certificate in a safe place so you can use it for recovery tasks.

The password for the `kubeadmin` user is available in the output that is generated when the OpenShift cluster finishes the installation, and in the corresponding log files. This log file also provides the path for the generated `kubeconfig` file that contains the client certificate for the `system:admin` user. The following example shows the OpenShift cluster output with the `kubeadmin` user details.

```
...output omitted...
INFO Install complete!
INFO To access the cluster as the system:admin user when using 'oc', run 'export
KUBECONFIG=/home/lab/ocp4/auth/kubeconfig' ①
INFO Access the OpenShift web-console here: https://console-openshift-
console.apps.ocp4.example.com
INFO Login to the console with user: "kubeadmin", and password: "iMcM5-eCy2Q-
PH66E-NkyPN" ②
INFO Time elapsed: 10m52s
...output omitted...
```

- ① Path for the generated `kubeconfig` file that contains the client certificate
- ② The `kubeadmin` user unique password

**Warning**

Removing the `kubeadmin` user cannot be undone. If you delete the `kubeadmin` user before creating another administrator user, then you lose administrator access to the OpenShift cluster and you must reinstall the cluster.

To remove the `kubeadmin` user from your cluster, log in as an administrator user and run the following command:

```
[user@host ~]$ oc delete secrets kubeadmin -n kube-system
```



References

For more information about authentication and authorization on OpenShift, refer to the Red Hat OpenShift Container Platform 4.14 *Authentication and Authorization* documentation at

https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/authentication_and_authorization/index

For more information about removing the `kubeadmin` user, refer to the *Removing the kubeadmin User* in the Red Hat OpenShift Container Platform 4.14 *Authentication and Authorization* documentation at

https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/authentication_and_authorization/index#removing-kubeadmin

OAuth Working Group Specifications

<https://oauth.net/specs/>

► Quiz

The OpenShift OAuth Server and Identity Providers

Choose the correct answers to the following questions:

- ▶ 1. Which authentication method handles user authentication in OpenShift, and is not present in Kubernetes?
 - a. Client certificates
 - b. OAuth server
 - c. Internal Certificate Authority

- ▶ 2. Which is the cluster administrator user that the OpenShift installation process creates?
 - a. root
 - b. admin
 - c. k8sadmin
 - d. kubeadmin

- ▶ 3. Which three API resources are API extension resources that OpenShift OAuth defines? (Choose three.)
 - a. Users
 - b. Identities
 - c. Roles
 - d. Groups
 - e. Role bindings

- ▶ 4. Your company configures two IdPs, htpasswd and LDAP, in the internal OpenShift OAuth server. Because both IdPs have some users in common, you must configure the mapping method parameter in the OAuth server to avoid any collision between user identities. Which mapping method maps the identities from both IdPs to the same user and avoids collisions?
 - a. claim
 - b. add
 - c. lookup

- ▶ 5. In OpenShift, how does the OAuth server handle user authentication?
 - a. The OAuth server relies on client certificates that the IdP issues.
 - b. The OAuth server delegates authentication to the IdP and generates a user OAuth token after the IdP validates the user credentials.
 - c. The OAuth server uses OAuth tokens that the IdP generates.

- **6. In Kubernetes and OpenShift, what is the default behavior for unauthenticated users when accessing the API server?**
- a. They are assigned to the system:unauthenticated group with limited access.
 - b. They are granted full cluster-admin privileges by default.
 - c. They are denied any access to the API resources.
 - d. They are assigned to the system:authenticated group with limited access.

► Solution

The OpenShift OAuth Server and Identity Providers

Choose the correct answers to the following questions:

- ▶ 1. Which authentication method handles user authentication in OpenShift, and is not present in Kubernetes?
 - a. Client certificates
 - b. OAuth server
 - c. Internal Certificate Authority

- ▶ 2. Which is the cluster administrator user that the OpenShift installation process creates?
 - a. root
 - b. admin
 - c. k8sadmin
 - d. kubeadmin

- ▶ 3. Which three API resources are API extension resources that OpenShift OAuth defines? (Choose three.)
 - a. Users
 - b. Identities
 - c. Roles
 - d. Groups
 - e. Role bindings

- ▶ 4. Your company configures two IdPs, htpasswd and LDAP, in the internal OpenShift OAuth server. Because both IdPs have some users in common, you must configure the mapping method parameter in the OAuth server to avoid any collision between user identities. Which mapping method maps the identities from both IdPs to the same user and avoids collisions?
 - a. claim
 - b. add
 - c. lookup

- ▶ 5. In OpenShift, how does the OAuth server handle user authentication?
 - a. The OAuth server relies on client certificates that the IdP issues.
 - b. The OAuth server delegates authentication to the IdP and generates a user OAuth token after the IdP validates the user credentials.
 - c. The OAuth server uses OAuth tokens that the IdP generates.

- **6. In Kubernetes and OpenShift, what is the default behavior for unauthenticated users when accessing the API server?**
- a. They are assigned to the system:unauthenticated group with limited access.
 - b. They are granted full cluster-admin privileges by default.
 - c. They are denied any access to the API resources.
 - d. They are assigned to the system:authenticated group with limited access.

LDAP Authentication and Group Synchronization

Objectives

- Configure an LDAP identity provider and automate group synchronization between OpenShift OAuth and an LDAP server.

The LDAP Service

After the installation of a new Red Hat OpenShift Container Platform (RHOCOP) cluster, only a `kubeadmin` user exists. A cluster administrator must next create users and groups, or more commonly, integrate a business system that both provides these entities and manages authentication. In doing so, an OpenShift cluster adopts a consistent corporate network authentication approach with other business systems, and inherits the existing security implementations in place, such as the company access restrictions that are available through a VPN-based connection.

Integrating an existing authentication method, and available users and groups in your organization, by using a Lightweight Directory Access Protocol (LDAP) server, is a common approach for many businesses. An LDAP directory server defines a standard application protocol for querying the system to authenticate users and groups. Additionally, LDAP integrations enable the use of Red Hat Identity Management (IdM) and Microsoft Active Directory (AD) services that many organizations use for managing authentication.

An LDAP server can be configured as an Identity Provider (IdP) for authentication bindings between the LDAP identity entries and cluster users through the OpenShift OAuth server. When users enter credentials for the cluster, the OAuth server initiates a connection and a corresponding query to the LDAP server, and creates bindings when matching unique entries are found.

Although this implementation is sufficient for authenticating the user, it is still necessary to define the access for the user, based on group memberships and roles. Furthermore, configuration within the OpenShift cluster is required to synchronize groups from the LDAP database, as well as to define the appropriate access and permissions for each group.



Note

This section does not cover LDAP administration in depth, but only the necessary information for configuring LDAP integrations in an RHOCOP cluster.

Authentication by Using LDAP

The LDAP protocol defines a query language to retrieve information from database entries of a remote LDAP server, which can contain entries for users and groups within the organization. When implemented as an OpenShift OAuth IdP, the cluster can rely on the LDAP server to validate user-provided login credentials for the cluster.

An LDAP URI exists and uses a standard syntax for connection parameters to an LDAP server, and is necessary when configuring client connections to an LDAP server. To perform queries, the URI includes the search criteria to match with entries in the LDAP database. Many LDAP servers also

require the use of an administrative bind DN, which is a set of privileged credentials with access to perform queries. This set of privileged credentials is included in the URI, when implemented.

A correctly formed search string describes the LDAP endpoint and the specific set of credentials, or a set of search criteria, such as user or group names, along with the bind Distinguished Name (bindDN) credentials when required for the interaction.

An LDAP URI is formed with the following standardized syntax:

```
ldap://host:port/basedn?attribute?scope?filter
```

ldap

The LDAP protocol designation. For LDAP over SSL, use the `ldaps` protocol. Red Hat recommends using StartTLS over regular `ldap`.

host:port

The LDAP server name and listening port. The default for `ldap` is `localhost:389`, and for `ldaps` is `localhost:636`.

basedn

The DN of the directory where a query begins. This directory is used as the root location for the search.

attributes

The target for the search, with the default as `uid`, which the LDAP lookup is intended to return. Multiple attributes are provided by a comma-separated list.

scope

The scope of the LDAP lookup. The scope is either `one` or `sub`. The default is `sub` if unspecified.

filter

An LDAP filter refines the results for the query. If omitted, the filter defaults to `(objectClass=*)`.

The following table provides an example set of information for an LDAP lookup for a specific username:

URL object	Value
Connection protocol	<code>ldap</code>
Server name and port	<code>ldap.example.com:389</code>
Base DN	<code>dc=example, dc=com</code>
Attributes	<code>givenName, sn, cn</code>
Query filter	<code>uid=payden.tomcheck</code>

The preceding information would result in the following URL for an LDAP lookup:

```
ldap://ldap.example.com:389/dc=example,dc=com?givenName,sn,cn?  
(uid=payden.tomcheck)
```

The result of this lookup returns one or more usernames from directory entries that match the given `uid` value in the LDAP server.

Chapter 1 | Authentication and Identity Management

In an OpenShift cluster, when a new user provides login credentials for the cluster, the OAuth LDAP IdP searches for the identity by using the configured LDAP search account. This LDAP search account is configured during the OAuth IdP configuration, and contains a Bind Distinguished Name (Bind DN) and a Bind Password. On a unique match from the search, the provided credentials are submitted to the LDAP server in a second interaction, and an authentication binding is created between an OpenShift user resource and the returned LDAP `id` and `preferredUsername` values.

If an authentication request has no matching result, then the query fails and no bind occurs.

Whereas the LDAP server provides the validation to create an authentication binding for the identity, a cluster administrator must also grant appropriate cluster roles for the identity through the OpenShift Role Based Access Controls (RBAC) settings. A new user configuration is complete when the binding and role are in place.

An administrator must synchronize groups from an LDAP server as an additional cluster-side configuration before a full LDAP IdP configuration is complete. It is necessary to create a cluster cron job to routinely query the LDAP server and to update the OpenShift groups to ensure consistency within the cluster to any organizational updates to the LDAP group entries.

LDAP CLI Queries

Although this course does not provide guidelines for LDAP administration, a tool is needed to inspect and validate available LDAP connectivity during cluster configuration.

The `ldapsearch` command is available from the `openldap-clients` RPM package, and provides a command-line query utility for LDAP interactions. The following example illustrates the format for an LDAP query that uses the `ldapsearch` command:

```
[user@host ~]$ ldapsearch -x -D bind_dn -H ldap_host_URI -b search_base \
--filter filter --requestedAttribute attributes -W
```

In this example, the following parameters are used:

- The `-x` option specifies using simple authentication.
- The `-W` option prompts you to provide the `bindPassword` secret for the specific `-D bind_dn` option.
- The `-b search_base` option denotes the root location for the lookup.
- The `--filter` and `--requestedAttribute` options provide further criteria to refine the search.

This approach provides a method for testing credentials for the LDAP server from the CLI, as well as retrieving listings for various LDAP entries, such as any defined groups. Although the previous example is too broad to provide meaningful results or detailed information, this approach establishes the validity of the credentials. Specifying additional attributes and filters returns more specific information, such as refining the query to search for an account with `uid=1001`, as shown in the following example:

```
[user@host ~]$ ldapsearch -x -D "cn=administrator,dc=example,dc=com" \
-H ldap://ldap.example.com -b "uid=1001" -W "objectclass=account" uid
Enter Password:
# extended LDIF
#
```

Chapter 1 | Authentication and Identity Management

```
# LDAPv3
# base <uid=1001> (default) with scope subtree
# filter: (objectclass=account)
# requesting: ALL

# example.com
dn: dc=example,dc=com

...output omitted...
```

The options for `ldapsearch` queries, and their resulting output, are consistent with the needed information for the configuration of an LDAP server as an OAuth IdP.

The organizational units that describe the path to the correct location for an information lookup within the LDAP system are similar to the structure of directories and files within a file system.

The following example DN shows the `uid`, `ou`, and `dc` values for information within the LDAP server:

```
dn: uid=ptomchek,ou=users,dc=ocp4,dc=redhat,dc=com
```

The following commands illustrate a corresponding organizational structure for the same information within a file system:

```
[user@host ~]$ mkdir -p ocp4.redhat.com/users
[user@host ~]$ touch ocp4.redhat.com/users/ptomchek
[user@host ~]$ tree ocp4.redhat.com/
ocp4.redhat.com/
    └── users
        └── ptomchek
```



Note

For more information about LDAP administration and the `ldapsearch` command, refer to the Red Hat Directory Server documentation in the references section.

Configuring an LDAP Server as an OpenShift IdP

Integrating an LDAP server as an IdP for an OpenShift cluster requires the following process:

- Create an OpenShift secret resource to store the password for LDAP queries. This OpenShift secret resource contains the base64-encoded secret in the `bindPassword` field.
- Create an OpenShift configuration map resource that contains the certificate authority bundle (if the LDAP server uses encryption).



Note

Create this configuration map resource only if your LDAP IdP has a CA certificate that is not present in the default system keystore.

This OpenShift ConfigMap resource belongs in the `openshift-config` namespace.

Chapter 1 | Authentication and Identity Management

- Update the CR to add the LDAP configuration to the existing OpenShift OAuth IdP entries.
- Redeploy the IdP CR to the cluster with the added LDAP IdP.

Create the bindPassword Secret Resource

The `bindPassword` for the IdP configuration is stored in an OpenShift secret resource. The following example is a YAML file for creating the `bindPassword` secret for an LDAP IdP:

```
apiVersion: v1
kind: Secret
metadata:
  name: ldap-secret
  namespace: openshift-config
type: Opaque
data:
  bindPassword: base64-encoded-bind-password
```


Note

Bear in mind that cluster administrators can read the password that is stored in the secret. Red Hat recommends using a service account with limited permissions as the bindDN within LDAP to avoid any security breaches.

Create the Certificate Configuration Map

LDAP configurations for OAuth IdPs use an OpenShift configuration map resource in the `openshift-config` namespace to define the certificate authority bundles.

Create an OpenShift configuration map resource that contains the certificate authority. You must store the certificate authority in the `ca.crt` key of the configuration map resource.

The following example is a YAML file that defines how to create a configuration map resource for the certificate bundle:

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: ca-config-map
  namespace: openshift-config
data:
  ca.crt: |
    CA_certificate_PEM
```

Update the OAuth CR

The following example shows the configuration for defining an LDAP IdP in the OAuth CR YAML file:

```
apiVersion: config.openshift.io/v1
kind: OAuth
metadata:
  name: cluster
```

```

spec:
  identityProviders:
    - name: ldapidp ①
      mappingMethod: claim ②
      type: LDAP
      ldap:
        attributes:
          id: ③
          - dn
          email: ④
          - mail
          name: ⑤
          - cn
        preferredUsername: ⑥
        - uid
      bindDN: 'cn=administrator,dc=example,dc=com' ⑦
      bindPassword: ⑧
        name: ldap-secret
      ca: ⑨
        name: ca-config-map
      insecure: false ⑩
      url: "ldaps://ldaps.example.com/ou=users,dc=acme,dc=com?uid" ⑪

```

- ① Provider names are shown on the web console login screen and in the list of configured IdPs.
- ② Controls how mappings are established between this provider's identities and user objects.
- ③ List of attributes where the first non-empty attribute is used. At least one attribute is required. If none of the listed attributes has a value, then authentication fails.
- ④ List of attributes to use as the email address. The first non-empty attribute is used.
- ⑤ List of attributes to use as the display name. The first non-empty attribute is used.
- ⑥ List of attributes to use as the preferred username when provisioning a user for this identity. The first non-empty attribute is used.
- ⑦ Required DN that is used during the search phase, unless anonymous searches are allowed. Must be set if the `bindPassword` parameter is defined.
- ⑧ Required OpenShift secret resource that contains the bind password, unless anonymous searches are allowed. Must be set if the `bindDN` parameter is defined.
- ⑨ Optional: Reference to an OpenShift configuration map resource that contains the privacy enhanced mail (PEM) encoded certificate authority bundle to validate server certificates for the configured URL. Used only when the `insecure` option is false.
- ⑩ When true, no TLS connection is made to the server. When false, `ldaps://` URLs connect by using TLS, and `ldap://` URLs are upgraded to TLS. The `insecure` option must be set to false when `ldaps://` URLs are in use, because these URLs always attempt to connect by using TLS.
- ⑪ An RFC 2255 URL, which specifies the LDAP host and maps the identity parameters in the LDAP database schema. This `url` value associates the user-supplied login credential with the specific entries in the LDAP server for identity lookups.

**Note**

The OAuth CR contains all configured IdPs. Any additional IdP that is configured is appended to the CR. Replacing or removing any entries can result in the loss of access to your cluster.

Log in by Using the Added LDAP IdP

After you update the OAuth CR, the new LDAP IdP is available to the cluster. The IdP Name is shown among any other configured IdPs on the web console login page, as well as in the list of configured OAuth IdPs.

Log in with this added LDAP IdP by using the `oc login` command or the added selection from the OpenShift web console, and by providing a username and password that are available through the LDAP server. During authentication, OpenShift generates a search filter by combining the attribute and filter in the configured OAuth CR `url` parameter, with the provided username. Then, OpenShift applies this filter to the LDAP directory to find a unique entry.

From the preceding configuration example, if the user enters `payden.tomcheck` as a username, then a query for that value is sent to the LDAP server. This query attempts to find a unique match in the `uid` LDAP field that the OAuth CR `url` value "`ldaps://ldaps.example.com/ou=users,dc=acme,dc=com?uid`" specifies.

If you configure the `mappingMethod` parameter for the LDAP IdP entry in the OAuth CR as `claim` or `add`, then OpenShift uses the resulting LDAP identity to create a cluster `User` resource on the first attempt to access the cluster.

Authentication fails if the LDAP lookup does not return a unique match. In this case, no binding is created, no user is added to the cluster, and access is denied.

OpenShift Group Synchronization with LDAP

Configuring an OAuth IdP for an LDAP server provides only a mechanism to validate and create users in the cluster, based on entries that are provided through the LDAP instance. Businesses also rely on LDAP to provide the groups that are used throughout the organization and to define appropriate RBAC configurations. By configuring LDAP group synchronization, a business can manage group memberships in a single location, and OpenShift can use these groups within the cluster.

OpenShift can synchronize groups from an LDAP server, which enables you to mirror the available LDAP groups within the cluster. LDAP group synchronization requires a sync configuration file with a client configuration and a query definition. Additionally, user-defined mappings can optionally provide granular details for establishing the custom relationships between group entries in the LDAP server and the corresponding OpenShift groups that you require.

Group synchronization requires a project where the pods run, a service account to perform the work, a cluster role where the right permissions are defined, and cluster role binding to associate the cluster role with the service account. When the synchronization configuration is in place, the synchronization of groups is automated by using a cron job to schedule the recurring execution of the synchronization task.

LDAP Sync Configuration File

To synchronize LDAP groups, create an `LDAPSsyncConfig` resource that contains the client LDAP configuration details and query definition for the LDAP group entities that are needed within OpenShift.

The LDAP client configuration part defines the connection parameters for the LDAP server, as well as the `bindDN` account and `bindPassword` secret for performing queries. This configuration uses the same connection and credentials that are used during the OAuth LDAP IdP implementation.

The following YAML code is an example of LDAP client configuration details:

```
url: ldap://1.2.3.4:389
bindDN: cn=administrator,dc=example,dc=com
bindPassword: password
insecure: false
ca: ca-bundle.crt
```

The LDAP query definition part describes the format of entries within the LDAP server that provide the groups to mirror to the OpenShift cluster.

The following YAML code is an example of an LDAP query definition:

```
baseDN: ou=groups,dc=example,dc=com
scope: sub
dereflAliases: never
timeout: 0
filter: (objectClass=group)
pageSize: 0
```

By assembling the client configuration details together with the query definition, you can create an `LDAPSsyncConfig` resource by using the parameters for your LDAP server:

```
kind: LDAPSsyncConfig
apiVersion: v1
url: ldap://example.com:389
insecure: false ①
rfc2307:
  groupsQuery:
    baseDN: "ou=groups,dc=example,dc=com"
    scope: sub
    dereflAliases: never
    pageSize: 0
    groupUIDAttribute: dn ②
    groupNameAttributes: [ cn ] ③
    groupMembershipAttributes: [ member ]
  usersQuery:
    baseDN: "ou=users,dc=example,dc=com"
    scope: sub
    dereflAliases: never
    pageSize: 0
    userUIDAttribute: dn ④
```

```
userNamesAttributes: [ mail ] ⑤
tolerateMemberNotFoundErrors: false
tolerateMemberOutOfScopeErrors: false
```

- ➊ When true, no TLS connection is made to the server. When false, `ldaps://` URLs connect by using TLS, and `ldap://` URLs are upgraded to TLS. The `insecure` option must be set to false when `ldaps://` URLs are in use, because these URLs always attempt to connect by using TLS.
- ➋ The field that corresponds to the unique identifier field on the LDAP server for a group.
- ➌ The attribute for the name of the group.
- ➍ The field that corresponds the unique identifier field of the LDAP server for a user.
- ➎ The attribute for the name of the user.

Example Active Directory Configuration

Active Directory (AD) is another source for providing group definitions for an OpenShift cluster. The AD configuration requires an LDAP query definition and attributes to represent users within the OpenShift group definitions.

The following YAML code is an example LDAP synchronization configuration for an AD schema:

```
kind: LDAPSsyncConfig
apiVersion: v1
url: ldap://ad.example.com:389
activeDirectory:
  usersQuery:
    baseDN: "ou=users,dc=example,dc=com"
    scope: sub
    derefAliases: never
    filter: (objectclass=person)
    pageSize: 0
    userNamesAttributes: [ mail ] ➊
    groupMembershipAttributes: [ memberOf ] ➋
```

- ➊ The attribute to use as the username in the OpenShift group record.
- ➋ The attribute on the user that stores the group membership information.

Note the schema that AD uses to map users to the corresponding group membership through the `memberOf` attribute.

Running an LDAP Synchronization

After the `LDAPSsyncConfig` file is created, an administrator can use the file to synchronize groups from the LDAP server to the OpenShift cluster. The following example command uses the `oc adm groups sync` command to synchronize all groups from the LDAP server to the cluster by using the `config.yaml` file that contains the `LDAPSsyncConfig` details:

```
[user@host ~]$ oc adm groups sync --sync-config=config.yaml --confirm
```

If the `--confirm` parameter is omitted, then the command performs a dry run of the operation. An initial dry run is advised when testing a new configuration, and to review the resulting groups that the cluster will inherit.

Automating LDAP Group Synchronizations

LDAP group synchronizations are automated through the use of a cron job. Configure this cron job from within a designated OpenShift project for this purpose. Execute the job by using a service account with administrative access to the cluster to manage groups.

Start by creating the project for the cron job to run. The following command creates a project named `ldap-group-sync` for this purpose:

```
[user@host ~]$ oc new-project ldap-group-sync
```

The client `bindPassword` secret resource and the configuration map resource that are used during the configuration of the LDAP IdP are also needed for this task. Additionally, create a service account with an appropriate cluster role and cluster role binding for the synchronization. The following YAML code shows an example definition for creating a service account named `ldap-group-sync-acct`:

```
kind: ServiceAccount
apiVersion: v1
metadata:
  name: ldap-group-sync-acct
  namespace: ldap-group-sync
```

Next, define a cluster role, as shown in the following example:

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: ldap-group-sync-acct
rules:
  - apiGroups:
    - ''
      ①
      - user.openshift.io
    resources:
      - groups
    verbs:
      - get
      - list
      - create
      - update
```

- ① The core API group is identified by an empty string.

Then, define a cluster role binding to bind the previous cluster role to the existing service account.

```
kind: ClusterRoleBinding
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: ldap-group-sync-acct
```

Chapter 1 | Authentication and Identity Management

```
subjects:
  - kind: ServiceAccount
    name: ldap-group-sync-acct
    namespace: ldap-group-sync
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: ldap-group-sync-acct
```

The cluster is configured to communicate with the LDAP IdP, and a service account is available to configure OpenShift groups from the LDAP group entities. The following YAML code is an example configuration map file to specify the sync configuration:

```
kind: ConfigMap
apiVersion: v1
metadata:
  name: ldap-group-sync-acct
  namespace: ldap-group-sync
data:
  sync.yaml: |
    kind: LDAPSyncConfig
    apiVersion: v1
    url: ldaps://1.2.3.4:389
    insecure: false
    bindDN: cn=administrator,dc=example,dc=com
    bindPassword:
      file: "/etc/secrets/bindPassword"
    ca: /etc/ldap-ca/ca.crt
    rfc2307:
      groupsQuery:
        baseDN: "ou=groups,dc=example,dc=com"
        scope: sub
        filter: "(objectClass=groupOfMembers)"
        derefAliases: never
        pageSize: 0
      groupUIDAttribute: dn
      groupNameAttributes: [ cn ]
      groupMembershipAttributes: [ member ]
      usersQuery:
        baseDN: "ou=users,dc=example,dc=com"
        scope: sub
        derefAliases: never
        pageSize: 0
      userUIDAttribute: dn
      userNameAttributes: [ uid ]
      tolerateMemberNotFoundErrors: false
      tolerateMemberOutOfScopeErrors: false
```

Finally, define and deploy a cron job for the LDAP group synchronization. The following YAML code is an example cron job definition that could be customized, such as by adapting the cron-specified schedule to meet business needs:

```

kind: CronJob
apiVersion: batch/v1
metadata:
  name: ldap-group-sync-acct
  namespace: ldap-group-sync
spec:
  schedule: "*/30 * * * *" ❶
  concurrencyPolicy: Forbid
  jobTemplate:
    spec:
      backoffLimit: 0
      ttlSecondsAfterFinished: 1800 ❷
      template:
        spec:
          containers:
            - name: ldap-group-sync
              image: "registry.redhat.io/openshift4/ose-cli:latest"
              command:
                - "/bin/bash"
                - "-c"
                - "oc adm groups sync --sync-config=/etc/config/sync.yaml --confirm" ❸
          volumeMounts:
            - mountPath: "/etc/config"
              name: "ldap-sync-volume"
            - mountPath: "/etc/secrets"
              name: "ldap-bind-password"
            - mountPath: "/etc/ldap-ca"
              name: "ldap-ca"
          volumes:
            - name: "ldap-sync-volume"
              configMap:
                name: "ldap-group-sync-acct"
            - name: "ldap-bind-password"
              secret:
                secretName: "ldap-secret" ❹
            - name: "ldap-ca"
              configMap:
                name: "ca-config-map"
          restartPolicy: "Never"
          terminationGracePeriodSeconds: 30
          activeDeadlineSeconds: 500
          dnsPolicy: "ClusterFirst"
          serviceAccountName: "ldap-group-sync-acct"

```

- ❶ The defined schedule for the job in cron format, to show every 30 minutes.
- ❷ The time, in seconds, to keep completed sync information, to correspond to the preceding cron schedule.
- ❸ The LDAP group sync command to execute by using the sync configuration file that is defined in the configuration map. **The oc adm groups sync command is defined on a single line**

- ④ The LDAP IdP bindDN and bindPassword.

After the file is created, the configured cron job executes the group synchronization command during the specified schedule to ensure that consistent group configurations that are available from the LDAP IdP are added to the OpenShift cluster. This job runs within the **project** namespace, by using the service account, at an interval that the cron time spec defines in the CronJob definition.



References

For more information about integrating LDAP IdP for a RHOCOP cluster, refer to the *Configuring LDAP Identity Providers* section in the *Authentication and Authorization* chapter in the Red Hat OpenShift Container Platform 4.14 *OpenShift Container Platform* documentation at

https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/authentication_and_authorization/index#configuring-ldap-identity-provider

For more information about LDAP administration and the `ldapsearch` command, refer to the Red Hat Directory Server documentation at

https://access.redhat.com/documentation/en-us/red_hat_directory_server/11/html-single/administration_guide/index#Examples-of-common-ldapsearches

► Guided Exercise

LDAP Authentication and Group Synchronization

Configure an LDAP identity provider and automate group synchronization between OpenShift OAuth and an LDAP server.

Outcomes

- Configure an LDAP identity provider (IdP) for RHOPC.

Before You Begin

As the `student` user on the `workstation` machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start auth-ldap
```

Instructions

Your company requires you to configure a Red Hat Directory Services (RHDS) Identity Provider (IdP), which is available on the `rhds` machine as an LDAP IdP for RHOPC.



Important

LDAP administration is outside the scope of this course. The RHDS server is already correctly configured for this exercise.

- 1. As the `admin` user, locate and then navigate to the Red Hat OpenShift web console.

- 1.1. Use the terminal to log in to the OpenShift cluster as the `admin` user with the `redhatocp` password.

```
[student@workstation ~]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
...output omitted...
```

- 1.2. Identify the URL for the OpenShift web console.

```
[student@workstation ~]$ oc whoami --show-console
https://console-openshift-console.apps.ocp4.example.com
```

- 1.3. Open a web browser and navigate to `https://console-openshift-console.apps.ocp4.example.com`. Either type the URL in a web browser, or right-click and select `Open Link` from the terminal.

- 2. Log in to the OpenShift web console as the `admin` user.

Chapter 1 | Authentication and Identity Management

- 2.1. Click **Red Hat Identity Management** and log in as the **admin** user with the **redhatocp** password.

- **3.** Navigate to the Configuration page for OAuth in the RHOC cluster web console.

 - 3.1. Click **Administration > Cluster Settings** from the sidebar.
 - 3.2. Click the **Configuration** tab to browse the list of configurable resources.

Cluster Settings

Details ClusterOperators Configuration

Edit the following resources to manage the configuration of your cluster.

Filter by name or description... /

- 3.3. Scroll through the alphabetical list of configuration resources and select **OAuth**.
- | Configuration resource | Description |
|--------------------------------|--|
| Authentication | Authentication specifies cluster-wide settings for authentication (like OAuth and webhook token authenticators). The canonical name of an instance is 'cluster'. Compatibility level 1: Stable within a major release for a minimum of 12 months or 3 minor releases (whichever is...) |
| OAuth | OAuth holds cluster-wide information about OAuth. The canonical name is 'cluster'. It is used to configure the integrated OAuth server. This configuration is only honored when the top level Authentication config has type set to IntegratedOAuth. Compatibility level 1: Stable within a... |

- **4.** View the existing YAML configuration for the Red Hat Identity Management IdP.

 - 4.1. Click the **YAML** tab and scroll down to the bottom of the configuration file.

Configuration > OAuth details

OA cluster

Details **YAML**

OAuth details

Name: cluster

Labels: No labels [Edit](#)

Annotations: 5 annotations

Created at: Jun 2, 2023, 6:53 AM

Owner: CV version

**Note**

You can view the cluster OAuth configuration from the CLI in YAML format, by using the `oc get oauth cluster --output yaml` command.

You can also view the cluster OAuth configuration details from the CLI, by using the `oc describe oauth` command.

- 4.2. The settings for the existing Red Hat Identity Management IdP are shown in the spec keys on lines 57-78.

```
spec:
  identityProviders:
    - ldap:
        attributes:
          email:
            - mail
          id:
            - dn
          name:
            - cn
        preferredUsername:
          - uid
      bindDN: 'uid=admin,cn=users,cn=accounts,dc=ocp4,dc=example,dc=com'
      bindPassword:
        name: ldap-secret
      ca:
        name: ca-config-map
      insecure: false
      url: >-
        ldap://idm.ocp4.example.com/
        cn=users,cn=accounts,dc=ocp4,dc=example,dc=com?uid
        mappingMethod: claim
        name: Red Hat Identity Management
        type: LDAP
```

- ▶ 5. From the CLI, test the RHDS credentials to ensure a valid LDAP connection, before configuring the additional IdP.

- 5.1. Open a terminal on the workstation, and create an `ldapsearch` query with the information in the following table:

Query option	Value
Bind DN (-D)	cn=Directory Manager
URI (-H)	<code>ldaps://rhds.ocp4.example.com</code>
Bind password (-w)	redhatocp

- 5.2. From the CLI, test the `ldapsearch` connection by using the authored query. With the `-w` option that is shown, you can supply the password, in plain text.

```
[student@workstation ~] ldapsearch -D "cn=Directory Manager" \
-w redhatocp -H ldaps://rhds.ocp4.example.com

# extended LDIF
#
# LDAPv3
# base <> (default) with scope subtree
# filter: (objectclass=*)
# requesting: ALL
#

# example.com
dn: dc=example,dc=com
objectClass: top
objectClass: domain
dc: example
description: dc=example,dc=com

# people, example.com
dn: ou=people,dc=example,dc=com
objectClass: top
objectClass: organizationalUnit
ou: people

# kristendelgado, people, example.com
dn: uid=kristendelgado,ou=people,dc=example,dc=com
objectClass: top
objectClass: account
objectClass: posixAccount
objectClass: shadowAccount
objectClass: nsMemberOf
cn: Kristen Delgado
uid: kristendelgado
uidNumber: 10001
gidNumber: 101
homeDirectory: /home/kristendelgado
loginShell: /bin/bash
gecos: kristendelgado
shadowLastChange: 0
shadowMax: 0
shadowWarning: 0
userPassword:: e1NTSEF9Wm1sMwd1WjJLajlRM1dKZGlFVnV6aTNEaCs5NzFPeFg=
memberOf: cn=administrators,ou=people,dc=example,dc=com

# administrators, people, example.com
dn: cn=administrators,ou=people,dc=example,dc=com
objectClass: top
objectClass: groupOfUniqueNames
cn: administrators
uniqueMember: uid=kristendelgado,ou=people,dc=example,dc=com

# search result
search: 2
```

```
result: 0 Success  
  
# numResponses: 5  
# numEntries: 4
```

- 6. In the web console, add a new IdP for RHDS by using the validated information.

- 6.1. Return to the previous page by clicking the **Details** tab.

**Note**

The existing Red Hat Identity Management IdP is shown in the list of available IdPs.

- 6.2. In the **Identity providers** section, select LDAP from the **Add** dropdown.

The screenshot shows a 'Identity providers' configuration interface. At the top, it says 'Identity providers determine how users log into the cluster.' Below that is a dropdown menu labeled 'Add' with a blue underline. A list of provider types is displayed: Basic Authentication, GitHub, GitLab, Google, HTPasswd, Keystone, LDAP, OpenID Connect, and Request Header. The 'LDAP' option is highlighted with a red rectangular box.

- 6.3. Complete the corresponding form fields by using the information in the following table:

Field	Value
Name	Red Hat Directory Server
URL and BaseDN	ldaps://rhds.ocp4.example.com/dc=example,dc=com?uid
Bind DN	cn=Directory Manager
Bind password	redhatocp
Email	mail

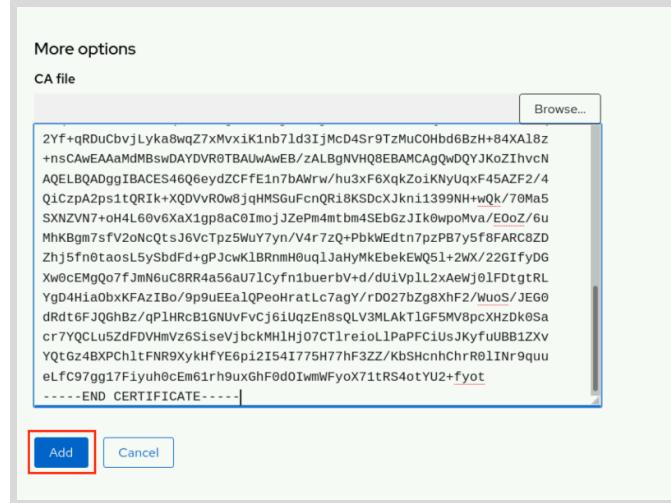
**Note**

Use the default values for all other fields.

- 6.4. In the CA File field, enter the certificate for the RHDS server from the rhds classroom machine, which is provided on the workstation machine in the file ~/D0380/labs/auth-ldap/rhds_ca.crt.

```
[student@workstation ~] cat ~/D0380/labs/auth-ldap/rhds_ca.crt
-----BEGIN CERTIFICATE-----
...output omitted...
```

- 6.5. Click Add.



- 7. Verify the addition of the new IdP.

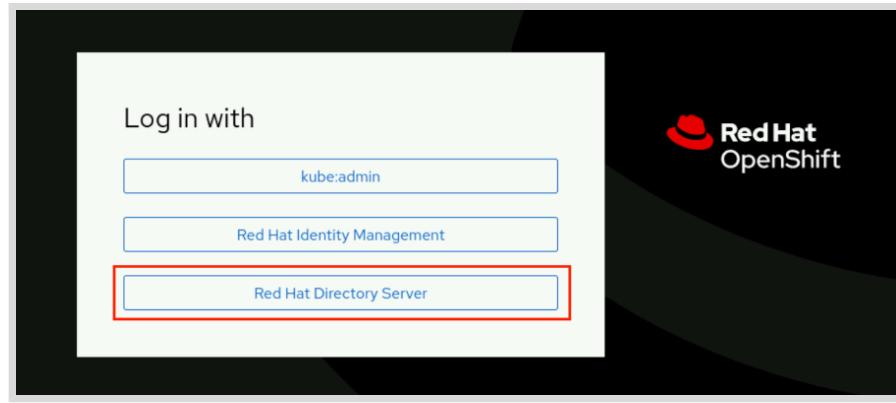
- 7.1. View the new entry in the IdP list on the OAuth details page.

OAuth details						
Name	cluster					
Labels	No labels Edit					
Annotations	5 annotations Edit					
Created at	Jun 2, 2023, 6:53 AM					
Owner	CV version					
Identity providers						
Identity providers determine how users log into the cluster.						
<div><p>Info New identity provider added.</p><p>Authentication is being reconfigured. The new identity provider will be available once reconfiguration is complete. View authentication conditions for reconfiguration status.</p></div>						
<div>Add ...</div>						
Name	Type	Mapping method				
Red Hat Identity Management	LDAP	claim				
Red Hat Directory Server	LDAP	claim				

7.2. Review the configuration additions on the **YAML** tab on lines 78-98.

```
spec:
  identityProviders:
    - ldap:
        attributes:
          email:
            - mail
        id:
          - dn
        name:
          - cn
      preferredUsername:
        - uid
    bindDN: 'cn=Directory Manager'
    bindPassword:
      name: ldap-bind-password-cjckt
    ca:
      name: ca-config-map
    insecure: false
    url: >-
      ldaps://rhds.ocp4.example.com/dc=example,dc=com?uid
  mappingMethod: claim
  name: Red Hat Directory Server
  type: LDAP
```

- 7.3. From the upper-right dropdown menu, click **Log out** to log out of the web console and return to the authentication page.
 - 7.4. The new IdP appears in the list as Red Hat Directory Server.



Note

It might take several minutes for the cluster to redeploy the pods and for the new IdP to appear.

- 8. Select the Red Hat Directory Server IdP, and verify the authentication function by using the following credentials to log in to the cluster.

Username	Password
kristendelgado	redhat123

Finish

On the **workstation** machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish auth-ldap
```

► Guided Exercise

Automate LDAP Group Synchronization

Automate group synchronization between OpenShift OAuth and an LDAP server.

Outcomes

- Configure an automated group synchronization for the Red Hat Directory Services (RHDS) LDAP identity provider (IdP).

Before You Begin

As the student user on the workstation machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start auth-sync
```

Instructions



Note

This exercise requires the completion of the previous Guided Exercise to configure the IdP. Ensure that you completed the previous section before proceeding.

Configure an automated group synchronization for the secondary RHDS IdP to maintain updated user and group information.

The implementation configures the cluster administrator privilege for the `administrators` group, which includes the `kristendelgado` user.

The RHDS LDAP IdP information is in the following table:

Parameter	Value
Bind DN (-D)	<code>cn=Directory Manager,dc=example,dc=com</code>
URI (-H)	<code>ldaps://rhds.ocp4.example.com</code>
Password (-w)	<code>redhatocp</code>

- 1. From the CLI, test the RHDS connection by using the `ldapsearch` command with the information in the table.

```
[student@workstation ~]$ ldapsearch -D "cn=Directory Manager" \
-w redhatocp -H ldaps://rhds.ocp4.example.com
# extended LDIF
#
# LDAPv3
# base <> (default) with scope subtree
```

```
# filter: (objectclass=*)
# requesting: ALL
#
# example.com
dn: dc=example,dc=com
objectClass: top
objectClass: domain
dc: example
description: dc=example,dc=com

# people, example.com
dn: ou=people,dc=example,dc=com
objectClass: top
objectClass: organizationalUnit
ou: people

# kristendelgado, people, example.com
dn: uid=kristendelgado,ou=people,dc=example,dc=com
objectClass: top
objectClass: account
objectClass: posixAccount
objectClass: shadowAccount
objectClass: nsMemberOf
cn: Kristen Delgado
uid: kristendelgado
uidNumber: 10001
gidNumber: 101
homeDirectory: /home/kristendelgado
loginShell: /bin/bash
gecos: kristendelgado
shadowLastChange: 0
shadowMax: 0
shadowWarning: 0
userPassword:::  
e1NTSEF9Wm1sMwd1WjJLajlRM1dKZG1FVnV6aTNEaCs5NzFPeFg=
memberOf: cn=administrators,ou=people,dc=example,dc=com

# administrators, people, example.com
dn: cn=administrators,ou=people,dc=example,dc=com
objectClass: top
objectClass: groupOfUniqueNames
cn: administrators
uniqueMember: uid=kristendelgado,ou=people,dc=example,dc=com

# search result
search: 2
result: 0 Success

# numResponses: 5
# numEntries: 4
```

Chapter1 | Authentication and Identity Management

- 2. Verify that you can log in to the cluster as the `kristendelgado` user with `redhat123` as the password, to ensure that the user from the RHDS IdP is available.

```
[student@workstation ~]$ oc login -u kristendelgado -p redhat123 \
https://api.ocp4.example.com:6443
Login successful.
...output omitted...
```

- 3. Change to the `~/D0380/labs/auth-sync-sync` directory and switch back to using the `admin` user and the `redhatocp` password.

- 3.1. Change to the `~/D0380/labs/auth-sync-sync` directory.

```
[student@workstation ~]$ cd ~/D0380/labs/auth-sync-sync
[student@workstation sync]$
```

- 3.2. Connect to the OpenShift cluster as the `admin` user with `redhatocp` as the password.

```
[student@workstation sync]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
Login successful.
...output omitted...
```

- 4. Create the `auth-rhds-sync` project.

```
[student@workstation sync]$ oc new-project auth-rhds-sync
Now using project "auth-rhds-sync" on server "https://api.ocp4.example.com:6443".
...output omitted...
```

- 5. Create a service account with permissions to synchronize groups from the RHDS server.

- 5.1. Create a service account called `rhds-group-syncer`.

```
[student@workstation sync]$ oc create sa rhds-group-syncer
serviceaccount/rhds-group-syncer created
```

- 5.2. Create the `rhds-group-syncer` cluster role with the `get`, `list`, `create`, `update` verbs.

```
[student@workstation sync]$ oc create clusterrole rhds-group-syncer \
--verb get,list,create,update --resource groups
clusterrole.rbac.authorization.k8s.io/rhds-group-syncer created
```

- 5.3. Create the cluster role binding for the `rhds-group-syncer` service account and cluster role.

```
[student@workstation sync]$ oc adm policy add-cluster-role-to-user \
rhds-group-syncer -z rhds-group-syncer
clusterrole.rbac.authorization.k8s.io/rhds-group-syncer added: "rhds-group-syncer"
```

Chapter 1 | Authentication and Identity Management

- 6. Create a secret that contains the provided RHDS bind password.

```
[student@workstation sync]$ oc create secret generic rhds-secret \
--from-literal bindPassword='redhatocp'
secret/rhds-secret created
```

- 7. Create the configuration files for the RHDS automated group synchronization.

- 7.1. Create the RHDS synchronization configuration from the example that is provided in the `rhds-sync.yaml` file in the working directory to supply the bind DN and password.

```
kind: LDAPSsyncConfig
apiVersion: v1
url: ldaps://rhds.ocp4.example.com:636
bindDN: 'cn=Directory Manager'
bindPassword:
  file: /etc/secrets/bindPassword
ca: /etc/config/ca.crt
augmentedActiveDirectory:
  groupsQuery:
    baseDN: "ou=people,dc=example,dc=com"
    scope: sub
    derefAliases: never
    pageSize: 0
  groupUIDAttribute: dn
  groupNameAttributes: [ cn ]
  usersQuery:
    baseDN: "ou=people,dc=example,dc=com"
    scope: sub
    derefAliases: never
    filter: (objectclass=account)
    pageSize: 0
  userNameAttributes: [ uid ]
  groupMembershipAttributes: [ memberOf ]
```

- 7.2. Create a configuration map that contains the `LDAPSsyncConfig` file and the trusted certificate.

```
[student@workstation sync]$ oc create configmap rhds-config \
--from-file rhds-sync.yaml=rhds-sync.yaml,ca.crt=rhds_ca.crt
configmap/rhds-config created
```

- 8. Create the cron job for the automated schedule to synchronize groups every minute.

- 8.1. Update the cron job configuration for the group synchronization example that is provided in the `rhds-groups-cronjob.yaml` file in the working directory to supply the secret and service account information.

```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: rhds-group-sync
```

```

namespace: auth-rhds-sync
spec:
  schedule: "*/1 * * * *"
  jobTemplate:
    spec:
      template:
        spec:
          restartPolicy: Never
          containers:
            - name: ldap-group-sync
              image: "registry.ocp4.example.com:8443/openshift4/ose-cli:v4.12"
              command:
                - "/bin/sh"
                - "-c"
                - "oc adm groups sync --sync-config=/etc/config/rhds-sync.yaml --confirm" ❶
              volumeMounts:
                - mountPath: "/etc/config"
                  name: "ldap-sync-volume"
                - mountPath: "/etc/secrets"
                  name: "ldap-bind-password"
              volumes:
                - name: "ldap-sync-volume"
              configMap:
                name: "rhds-config"
              - name: "ldap-bind-password"
                secret:
                  secretName: "rhds-secret"
            serviceAccountName: rhds-group-syncer
            serviceAccount: rhds-group-syncer

```

- ❶ Write this value in a single line.

8.2. Create the cron job from the `rhds-groups-cronjob.yaml` file.

```
[student@workstation sync]$ oc create -f rhds-groups-cronjob.yaml
...output omitted...
cronjob.batch/rhds-group-sync created
```

8.3. Wait for one minute for the cron job to trigger, and verify the synchronization of the `administrators` group from RHDS.

```
[student@workstation sync]$ oc get groups
NAME           USERS
Default SMB Group
administrators   kristendelgado
...output omitted...
```

- ▶ 9. Apply the `cluster-admin` role to the `administrators` group.

```
[student@workstation sync]$ oc adm policy add-cluster-role-to-group \
  cluster-admin administrators
clusterrole.rbac.authorization.k8s.io/cluster-admin added: "administrators"
```

- **10.** Log in to the cluster as the `kristendelgado` user and verify that the user has cluster administrator privileges from the `administrators` group membership.

10.1. Log in to the cluster.

```
[student@workstation sync]$ oc login -u kristendelgado -p redhat123
Login successful.
...output omitted...
```

10.2. Verify you can perform an administrative task.

```
[student@workstation sync]$ oc auth can-i create users -A
yes
```

- **11.** Change to the student HOME directory.

```
[student@workstation sync]$ cd
```

Finish

On the `workstation` machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish auth-sync
```

OIDC Authentication and Group Claims

Objectives

- Configure an OIDC identity provider and automate group synchronization between OpenShift OAuth and an OIDC server.

OpenID Connect

You can use OpenShift to configure OpenID Connect (OIDC) identity providers (IdPs) for synchronizing users and groups.

OIDC is a set of standards for delegating the authentication of a user who accesses a protected resource. OIDC provides a way for applications to verify the identity of users and to obtain user profile information.

OIDC is the standard for cloud authentication, social authentication, single sign-on (SSO), and two-factor authentication (2FA). OIDC removes the responsibility of setting, storing, and managing passwords locally, which is often associated with credential-based data breaches, and which needs administration resources.

Examples of OIDC providers that are tested and supported on OpenShift include Google, Microsoft identity platform, and Keycloak. Red Hat provides Red Hat Single Sign-On (SSO) to extend the capabilities of the OpenShift internal OAuth, and to serve as the solution for an OIDC identity infrastructure. Red Hat SSO can run on bare metal or virtualized environments, or as pods on OpenShift.



Note

For more information about Red Hat SSO, refer to the *Product Documentation for Red Hat Single Sign-On 7.6* at https://access.redhat.com/documentation/en-us/red_hat_single_sign-on/7.6

To see the full list of supported OIDC providers on OpenShift, you can refer to https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/authentication_and_authorization/index#identity-provider-oidc-supported_configuring-oidc-identity-provider

OIDC Tokens

OIDC relies on the JavaScript Object Signing and Encryption (JOSE) set of standards. The primary standard within JOSE is JSON Web Token (JWT), which serves as a standardized format for representing information in a readable piece of text data. The following list shows some JWT advantages:

- Can parse in many programming languages
- Can propagate across networks
- Validates message integrity without relying on external validators or resource-intensive checks

Chapter 1 | Authentication and Identity Management

Three parts, which are separated by a period, form the JWT token: the header, the payload, and the hash signature.

The header contains the type of the token and the signing algorithm.

The payload contains the user claims. User claims are attributes for the user with details about their identity, profile, privileges, and group membership. Examples of user claims are the `sub`, `iat`, `exp`, or `iss` parameters, which are the identity for the user, the issue date, the expiration date, and the token issuer, respectively.

Finally, the signature is composed of the concatenation of the encoded header, the encoded payload, and the result of applying a signing algorithm.

OIDC in OpenShift

An *identity broker* is a service that connects clients with IdPs. The identity broker delegates the authentication of the user to the external IdP. OpenShift includes a built-in OAuth server that you can configure to determine the user's identity from the configured IdPs. Thus, the OpenShift built-in OAuth server acts as an identity broker.

When a user tries to log in to OpenShift, the OpenShift built-in OAuth server redirects the user to a login screen to choose from the configured IdPs. For OIDC IdPs, you must previously configure the OIDC client in the IdP that authenticates the user to OpenShift. Then, use the connection parameters and credentials from the OIDC client to configure the OIDC IdP in the OpenShift built-in OAuth server. After authenticating the user in the configured IdP, the internal OAuth server creates an access token for the request and returns the token. The built-in OAuth server, as for any other IdP, also creates or updates the user resources. If you define OIDC group claims, then the OAuth server also creates or updates the groups.

OIDC Claims

OIDC claims are key/value pairs that contain information about the user. OpenShift reads user claims from the JWT token that the IdP issues, and uses these user claims to populate the user, identity, and group resources. You must configure one claim to use as the user's identity. The default identity claim is `sub`, which stands for *subject identifier*.

You can configure additional user parameters in other standard claims, such as the preferred username, display name, or email address. For any of those user parameters, you can specify multiple claims. OpenShift uses the first claim with a non-empty value.

The following list includes some standard claims that are defined in OIDC:

`sub`

The remote identity for the user at the IdP.

`preferred_username`

The preferred username when provisioning a user, which typically corresponds to the username in the authentication system for the user.

`email`

Email address.

`name`

Display name.

**Note**

To see the full list of standard claims that are defined in OIDC, refer to https://openid.net/specs/openid-connect-core-1_0.html#StandardClaims

OIDC Group Claims

You can also define additional claims for parameters that OIDC does not define as standard claims. One important additional claim for users is the `group` claim. OpenShift uses the `group` claim to map the group membership of a user in the IdP to a group object. Then, you can use role-based access control (RBAC) objects in OpenShift to assign permissions to that group, instead of assigning permissions to individual users.

**Note**

The Red Hat Communities of Practice Group Sync Operator GitHub repository provides an unsupported operator for synchronizing OIDC groups with external providers that cannot provide group claims as part of their tokens. You can find more information in <https://github.com/redhat-cop/group-sync-operator>

Configuring OIDC IdP

Follow the next steps to integrate an OIDC IdP into the OpenShift cluster:

1. Obtain the client ID and the client secret from the OIDC IdP client for the OpenShift integration.
2. Create an OpenShift secret object, which contains the client secret that is obtained from the OIDC client configuration, in the `openshift-config` namespace.
3. Create an OpenShift configuration map object, which contains the certificate authority bundle in the `ca.crt` file parameter, in the `openshift-config` namespace (required only if the CA certificate is not configured as a system-wide CA).
4. Create the OAuth CR YAML file to include the OIDC IdP.
5. Apply the configuration file to the OAuth CR.

Create the OAuth CR YAML File

After creating the OpenShift secret object that contains the client secret, and the configuration object that contains the certificate authority bundle (if necessary), you can create the OAuth CR YAML file that contains the information to configure the OIDC IdP. If you add an OIDC IdP to the OAuth CR, then you must include the OIDC IdP information in the `identityProviders` array. You can add multiple OIDC IdPs to the OAuth CR YAML file.

**Important**

The `identityProviders` array in the OAuth CR might not be empty. If you remove other IdPs when adding your OIDC IdP, then you cannot log in to the cluster through those IdPs.

Chapter 1 | Authentication and Identity Management

The following example shows a minimal configuration file for Red Hat SSO integration to OpenShift. The settings might differ for other OIDC providers, and you must work with your vendor or identity administrator to get all the necessary attributes for your specific setup.

```
apiVersion: config.openshift.io/v1
kind: OAuth
metadata:
  name: cluster
spec:
  identityProviders:
    - name: oidc_provider_name ①
      mappingMethod: claim
      type: OpenID
      openID:
        clientID: oidc_clientid ②
        clientSecret: ③
          name: secret_name
        ca: ④
          name: config_map_name
        claims: ⑤
          preferredUsername:
            - preferred_username
            - email
          name:
            - nickname
            - given_name
            - name
          email:
            - custom_email_claim
            - email
        groups:
          - groups
      issuer: https://external_idp_url.com ⑥
```

- ① OpenShift prefixes the value of the identity claim with the provider name to form an identity name, and uses the identity name to build the redirect URL.
- ② The client ID for the existing client in the IdP. You must enable the client to redirect to `https://oauth-openshift.apps.cluster_name.cluster_domain/oauth2callback/idp_provider_name`.
- ③ The name for the OpenShift secret object that contains the client secret.
- ④ (Optional) The name for the OpenShift configuration map object that contains the certificate authority bundle.
- ⑤ The list of claims to use as the identity, such as the preferred username, email address, or groups.
- ⑥ The URL for the IdP. OpenShift accepts only HTTPS URLs.

Log in to the Cluster Through the OIDC IdP

After adding the OIDC IdP to the cluster, you can log in through the IdP by using the OpenShift web console, and providing the username and the password.

If your OIDC IdP supports the grant flow for resource owner password credentials (ROPC), then you can also log in through the IdP by using the `oc login` command with the username and the password.

If your OIDC IdP does not support the ROPC grant flow, then you receive a `You must obtain an API token` login error when you use the `oc login` command with the username and the password. Then, you must first get an OAuth access token and use it to log in by using the `oc login --token=access_token` command. You can get the OAuth access token by logging in through the OpenShift web console and then clicking **Help > Command line tools > Copy login command**. You can also request the OAuth access token through the OpenShift REST API.



Note

Requesting an OAuth access token through the OpenShift REST API is out of the scope of this course. For more information about this topic, refer to <https://access.redhat.com/solutions/6610781>

IdP Mapping Methods

IdP mapping methods apply to any IdP. If you configure the `mappingMethod` parameter for the OIDC IdP as `claim` or `add`, then OpenShift establishes mappings between the provider's identity and the User object the first time that a user logs in to the cluster.

OAuth Access Tokens

The first time that a user logs in to the cluster, the OpenShift built-in OAuth server creates an OAuth token to authenticate to the API. OAuth access tokens are common to any IdP. OpenShift renews the token to authenticate to the API every time that the user logs in to the cluster. As a user, you can list all your user-owned OAuth access tokens by using the following command. This command lists all the user-owned OAuth access tokens from any IdP that is configured in the OpenShift built-in OAuth server.

```
[user@host ~]$ oc get useroauthaccesstokens
NAME      CLIENT NAME   CREATED     EXPIRES   ...
sha256~9BZ3...  openshift-...  69m        2023-06-14 13:24:42 +0000 UTC ...
sha256~lm10...  console       75m        2023-06-14 13:19:20 +0000 UTC ...
sha256~xmpc...  openshift-...  83m        2023-06-14 13:11:24 +0000 UTC ...
```

If you modify a user parameter in the OIDC IdP, then the OIDC IdP does not automatically synchronize to OpenShift. For example, if you remove a user account from the OIDC IdP and the user is logged in to OpenShift, then the user can still perform tasks in OpenShift until they log out, because they still have a valid token that the OpenShift built-in OAuth server emits.

For this reason, after you modify a parameter in the OIDC IdP, you must manually remove all the user access tokens from OpenShift to force a user logout. Then, after the user logs in again, OpenShift synchronizes the new user parameters.

You can use the following command to remove all the user access tokens from OpenShift:

```
[user@host ~]$ oc delete oauthaccesstoken $(oc get oauthaccesstoken -o \
  jsonpath='{.items[?(@.userName=="username")].metadata.name}')
```



References

For more information about OIDC, refer to the OpenID specifications web page at <https://openid.net/developers/specs/>

For more information about how to configure OIDC IdPs on OpenShift, refer to the *Configuring an OpenID Connect Identity Provider* section in the Red Hat OpenShift Container Platform 4.14 Authentication and Authorization documentation at https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/authentication_and_authorization/index#identity-provider-oidc-about_configuring-oidc-identity-provider

For more information about how to configure a GitHub or GitHub Enterprise IdP on OpenShift, refer to the *Configuring a GitHub or GitHub Enterprise Identity Provider* section in the Red Hat OpenShift Container Platform 4.14 Authentication and Authorization documentation at https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/authentication_and_authorization/index#identity-provider-overview_configuring-github-identity-provider

For more information about how to configure Azure Active Directory authentication on OpenShift, refer to the *Configure Azure Active Directory Authentication for an Azure Red Hat OpenShift 4 Cluster (CLI)* section in the Azure Red Hat OpenShift documentation at <https://learn.microsoft.com/en-us/azure/openshift/configure-azure-ad-cli>

► Guided Exercise

OIDC Authentication and Group Claims

Configure an OIDC identity provider and automate group synchronization between OpenShift OAuth and an OIDC server.

Outcomes

- Configure Red Hat Single Sign-On (SSO) as an OIDC identity provider (IdP) for OpenShift.
- Synchronize users and groups from Red Hat SSO to OpenShift.

Before You Begin

As the student user on the workstation machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start auth-oidc
```

Instructions

Your company requires you to configure Red Hat SSO, which is running in the `sso` machine as an OIDC IdP for OpenShift, to automate user and group synchronization between OpenShift OAuth and the Red Hat SSO OIDC server.

As a use case, the `lab` script configures Red Hat SSO to include the `external_providers` realm. The following table provides the details for three users and two groups that are available in this realm. The password for the three users is `redhat_sso`.

First name	Last name	Username	Group membership
Abby	Quincy	abbyquincy	contractors
Fricis	Ritcher	fricisritcher	contractors
Jaya	Lamont	jayalamont	partners

The company requests that you give `read` access to the `auth-oidc` OpenShift project for users in the `partners` group. Additionally, users in the `contractors` group need to be able to edit objects in the `auth-oidc` OpenShift project.

Finally, inspect the behavior in OpenShift after the deletion of a synchronized user or a group membership from Red Hat SSO.



Important

Red Hat SSO administration is outside the scope of this course. For more information about Red Hat SSO, refer to the *Red Hat Single Sign-On Administration (DO313)* course at <https://www.redhat.com/en/services/training/do313-red-hat-single-sign-on-administration>

- ▶ 1. Assign the `edit` cluster role in OpenShift to the `contractors` group, so users in that group can modify most of the objects in the `auth-oidc` project. Assign the `view` cluster role in OpenShift to the `partners` group, so users in that group can view most of the objects in the `auth-oidc` project but cannot make modifications.
- 1.1. Connect to the OpenShift cluster as the `admin` user with `redhatocp` as the password.

```
[student@workstation ~]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
Login successful.
```

```
You have access to 70 projects, the list has been suppressed.
You can list all projects with 'oc projects'

Using project "default".
```

- 1.2. Change to the `auth-oidc` project.

```
[student@workstation ~]$ oc project auth-oidc
Now using project "auth-oidc" on server "https://api.ocp4.example.com:6443".
```

- 1.3. Assign the `edit` cluster role in the `auth-oidc` project to the `contractors` group.



Note

Ignore the warning message, because OpenShift creates the `contractors` group after you synchronize the users from Red Hat SSO.

```
[student@workstation ~]$ oc adm policy add-role-to-group edit contractors
Warning: Group 'contractors' not found
clusterrole.rbac.authorization.k8s.io/edit added: "contractors"
```

- 1.4. Assign the `view` cluster role in the `auth-oidc` project to the `partners` group.



Note

Ignore the warning message, because OpenShift creates the `partners` group after you synchronize the users from Red Hat SSO.

```
[student@workstation ~]$ oc adm policy add-role-to-group view partners
Warning: Group 'partners' not found
clusterrole.rbac.authorization.k8s.io/view added: "partners"
```

- 2. List the users and groups in the `external_providers` realm in Red Hat SSO. List the group membership for the `abbyquincy` user.

- 2.1. Connect to the Red Hat SSO machine as the `rhsso` user.

```
[student@workstation ~]$ ssh rhsso@sso.ocp4.example.com
[rhsso@sso ~]$
```

- 2.2. Use the `kcadm` tool to connect to Red Hat SSO.

```
[rhsso@sso ~]$ /opt/rh-sso-7.6/bin/kcadm.sh config credentials \
--server https://sso.ocp4.example.com:8080/auth \
--user admin --password redhatocp --realm master
Logging into https://sso.ocp4.example.com:8080/auth as user admin of realm master
```

- 2.3. List the users in the `external_providers` realm. The user IDs would differ on your system.

```
[rhsso@sso ~]$ /opt/rh-sso-7.6/bin/kcadm.sh get users \
-r external_providers --fields 'id,username'
[ {
  "id" : "a175e1b7-6210-40f8-aeda-732021142e84",
  "username" : "abbyquincy"
}, {
  "id" : "958fde0c-360c-48f8-b5e5-942708fbb36e",
  "username" : "fricisritcher"
}, {
  "id" : "64898122-5875-4418-88ac-c9eeeaa0f409",
  "username" : "jayalamont"
} ]
```

- 2.4. List the groups in the `external_providers` realm. The group IDs would differ on your system.

```
[rhsso@sso ~]$ /opt/rh-sso-7.6/bin/kcadm.sh get groups \
-r external_providers --fields 'id,name'
[ {
  "id" : "e92319be-d5df-4a0c-833a-687fd25ca34c",
  "name" : "contractors"
}, {
  "id" : "3dcc3053-4ebf-4894-969a-f26d8e2bc22f",
  "name" : "partners"
} ]
```

- 2.5. List the groups for the `abbyquincy` user. Use the ID for the `abbyquincy` user from an earlier step.

```
[rhsso@sso ~]$ /opt/rh-sso-7.6/bin/kcadm.sh get \
  users/a175e1b7-6210-40f8-aeda-732021142e84/groups -r external_providers
[ {
  "id" : "e92319be-d5df-4a0c-833a-687fd25ca34c",
  "name" : "contractors",
  "path" : "/contractors"
} ]
```

► 3. Retrieve the Red Hat SSO client information and note the client secret.

3.1. List the information for the ocp_rhsso client from Red Hat SSO.

```
[rhsso@sso ~]$ /opt/rh-sso-7.6/bin/kcadm.sh get clients \
  -r external_providers -q clientId=ocp_rhss
[ {
  "id" : "f57e9ddc-8c60-4b40-8048-ec0120595be2",
  "clientId" : "ocp_rhss",
  "surrogateAuthRequired" : false,
  "enabled" : true,
  "alwaysDisplayInConsole" : false,
  "clientAuthenticatorType" : "client-secret",
  "redirectUris" : [ "https://oauth-openshift.apps.ocp4.example.com/*" ],
  "webOrigins" : [ "https://oauth-openshift.apps.ocp4.example.com" ],
  ...output omitted...
```

3.2. Generate a JSON file with the Red Hat SSO client information, which contains the client ID, the authentication server URL, and the client secret. Use the ocp_rhsso ID from the previous step. The client secret is necessary for configuring the OAuth Custom Resource (CR) in OpenShift, in a later step.

```
[rhsso@sso ~]$ /opt/rh-sso-7.6/bin/kcadm.sh get \
  clients/f57e9ddc-8c60-4b40-8048-ec0120595be2/installation\
/providers/keycloak-oidc-keycloak-json \ ①
  -r external_providers > rhsso.json
```

- ① The /providers/... text must come after the .../installation text in a single line without spaces.

3.3. View the content of the JSON file, which contains the Red Hat SSO client information. You would need some Red Hat SSO client information when configuring the OIDC IdP on OpenShift.

Note the secret, which you use in a later step. The client secret would differ on your system.

The JSON file also provides the Red Hat SSO client ID, ocp_rhsso, in the resource parameter. Use this value for the clientId parameter in the IdP configuration file on OpenShift.

The issuer parameter in the IdP configuration on OpenShift concatenates the value from the auth-server-url parameter, the /realms/ string, and the Red Hat SSO realm name, which in this case is external_providers.

```
[rhsso@sso ~]$ cat rhsso.json
{
  "realm" : "external_providers",
  "auth-server-url" : "https://sso.ocp4.example.com:8080/auth/",
  "ssl-required" : "external",
  "resource" : "ocp_rhsso",
  "credentials" : {
    "secret" : "X4ZTPfDr0b8loq0FArfidhaHq85bHyiy"
  },
  "confidential-port" : 0
```

3.4. Return to the workstation machine.

```
[rhsso@sso ~]$ exit
logout
Connection to sso.ocp4.example.com closed.
[student@workstation ~]$
```

- ▶ **4.** Configure the OpenShift OAuth CR to synchronize users from the Red Hat SSO OIDC client that was configured in the previous step.
- 4.1. Connect to the OpenShift cluster as the `admin` user with `redhatocp` as the password.

```
[student@workstation ~]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
Login successful.
...output omitted...
```

4.2. Create the `rhsso-oidc-client-secret` OpenShift secret for the Red Hat SSO client secret by using the client secret from a previous step.

```
[student@workstation ~]$ oc create secret generic rhsso-oidc-client-secret \
--from-literal clientSecret=X4ZTPfDr0b8loq0FArfidhaHq85bHyiy \
-n openshift-config
secret/rhsso-oidc-client-secret created
```

4.3. Create the OAuth CR YAML file. You can find an example for the CR in the `~/D0380/labs/auth-oidc/sso_config.yaml` file. The YAML file includes an LDAP IdP that you must preserve, because it provides the `admin` and `developer` users. Do not remove the LDAP IdP, and add the OIDC IdP for Red Hat SSO.

```
apiVersion: config.openshift.io/v1
kind: OAuth
metadata:
  name: cluster
spec:
  identityProviders:
    - ldap:
        ...output omitted...
    - openID:
```

Chapter 1 | Authentication and Identity Management

```

claims:
  email:
    - email
  name:
    - name
  preferredUsername:
    - preferred_username
  groups:
    - groups
clientID: ocp_rhsso
clientSecret:
  name: rhsso-oidc-client-secret
extraScopes: []
issuer: >-
  https://sso.ocp4.example.com:8080/auth/realms/external_providers
mappingMethod: claim
name: RHSSO_OIDC
type: OpenID

```

4.4. Apply the configuration to the OAuth CR.

```
[student@workstation ~]$ oc apply -f ~/D0380/labs/auth-oidc/sso_config.yml
oauth.config.openshift.io/cluster configured
```

4.5. Verify the status for the OAuth pods and wait for the OAuth pods to be redeployed. Press Ctrl+C when done.

```
[student@workstation ~]$ watch oc get pods -n openshift-authentication
Every 2.0s: oc get pods -n openshift-authentication  workstation: ...

NAME                               READY   STATUS    RESTARTS   AGE
oauth-openshift-79c7865785-6zxvp  1/1     Running   0          2m10s
oauth-openshift-79c7865785-bbp5w  1/1     Running   0          2m39s
oauth-openshift-79c7865785-jl6th  1/1     Running   0          102s
^C
```

- ▶ 5. Verify that you can log in to the cluster as the `abbyquincy` user with `redhat_sso` as the password, and create resources in the `auth-oidc` OpenShift project. The ability to create resources derives from the `edit` cluster role for the `contractors` group.

5.1. Log in to the cluster as the `abbyquincy` user.

```
[student@workstation ~]$ oc login -u abbyquincy -p redhat_sso
Login successful.
...output omitted...
Using project "auth-oidc".
```

5.2. Verify that the user can create a pod.

```
[student@workstation ~]$ oc run ubi9-date --restart 'Never' \
--image registry.ocp4.example.com:8443/ubi9/ubi -- date
pod/ubi9-date created
```

Chapter1 | Authentication and Identity Management

- 5.3. Verify that the user can view pod information in the project.

```
[student@workstation ~]$ oc get pods
NAME        READY   STATUS    RESTARTS   AGE
ubi9-date   0/1     Completed  0          85s
```

- 5.4. Verify that the user cannot review the user information, because the user is not a cluster administrator.

```
[student@workstation ~]$ oc get users
Error from server (Forbidden): users.user.openshift.io is forbidden:
User "abbyquincy" cannot list resource "users" in API group "user.openshift.io" at
the cluster scope
```

- 6. Verify that OpenShift creates the user and group the first time that they log in.

- 6.1. Verify that OpenShift synchronizes the user from Red Hat SSO the first time that they log in, and that the **abbyquincy** user is a member of the **contractors** group in OpenShift.

```
[student@workstation ~]$ oc login -u admin -p redhatocp
...output omitted...

[student@workstation ~]$ oc get users
NAME        UID           FULL NAME      IDENTITIES
abbyquincy  11bc3d49-...  Abby Quincy   RHSSO_OIDC:a175e1b7-...-732021142e84
admin       00fed1ea-...  Administrator  Red Hat Identity Management:dWlk...

[student@workstation ~]$ oc get groups
NAME        USERS
Default SMB Group
admins      admin
contractors  abbyquincy
developer
editors
ocpadmins   admin
ocpdevs     developer
```

- 7. Synchronize the **jayalamont** user to OpenShift, and verify that the user can view objects in the OpenShift project, but cannot edit them.

- 7.1. Log in as the **jayalamont** user with **redhat_sso** as the password.

```
[student@workstation ~]$ oc login -u jayalamont -p redhat_sso
Login successful.
...output omitted...
Using project "auth-oidc".
```

- 7.2. Verify that the user can view pod information in the project.

Chapter 1 | Authentication and Identity Management

```
[student@workstation ~]$ oc get pods
NAME        READY   STATUS    RESTARTS   AGE
ubi9-date   0/1     Completed   0          7m47s
```

7.3. Verify that the user cannot remove the pod in the project.

```
[student@workstation ~]$ oc delete pod ubi9-date
Error from server (Forbidden): pods "ubi9-date" is forbidden: User "jayalamont"
cannot delete resource "pods" in API group "" in the namespace "auth-oidc"
```

- 8. Log in to OpenShift as the **fricisritcher** user, which is another user in the **contractors** group, and which can edit resources in the **auth-oidc** OpenShift project.

- 8.1. Log in as the **fricisritcher** user in the OpenShift web console. To do so, open a web browser and navigate to <https://console-openshift-console.apps.ocp4.example.com>. Click **RHSSO_OIDC**.
- 8.2. Log in as the **fricisritcher** user with **redhat_sso** as the password.

- 9. Remove the user's membership of the **contractors** group in Red Hat SSO.

- 9.1. Change to the terminal window. Verify that OpenShift correctly synchronizes the user from Red Hat SSO, and that the user is a member of the **contractors** group.

```
[student@workstation ~]$ oc login -u admin -p redhatocp
...output omitted...

[student@workstation ~]$ oc get users
NAME        UID           FULL NAME      IDENTITIES
abbyquincy  11bc3d49-...  Abby Quincy    RHSSO_OIDC:a175e1b7-...-732021142e84
admin        00fed1ea-...  Administrator  Red Hat Identity Management:dwlk...
fricisritcher af752e5b-...  Fricis Ritcher  RHSSO_OIDC:958fde0c-...-942708fbb36e
jayalamont   e5f74993-...  Jaya Lamont   RHSSO_OIDC:64898122-...-c9eeeaa0f409

[student@workstation ~]$ oc get groups
NAME          USERS
Default SMB Group
admins        admin
contractors  abbyquincy, fricisritcher
developer
editors
ocpadmins    admin
ocpdevs       developer
partners      jayalamont
```

- 9.2. Connect to the Red Hat SSO machine as the **rhsso** user. Remove the **fricisritcher** user's membership of the **contractors** group in Red Hat SSO, by using the user and group ID from a previous step.

```
[student@workstation ~]$ ssh rhsso@sso.ocp4.example.com

[rhsso@sso ~]$ /opt/rh-sso-7.6/bin/kcadm.sh config credentials \
--server https://sso.ocp4.example.com:8080/auth \
--user admin --password redhatocp --realm master
Logging into https://sso.ocp4.example.com:8080/auth as user admin of realm master

[rhsso@sso ~]$ /opt/rh-sso-7.6/bin/kcadm.sh delete \
users/958fde0c-360c-48f8-b5e5-942708fbb36e \
/groups/e92319be-d5df-4a0c-833a-687fd25ca34c \ ❶
-r external_providers
```

❶ The /groups/... text must come after the user ID, without spaces.

- 9.3. Verify that, even though you remove the **fricisritcher** user's membership of the **contractors** group in Red Hat SSO, the user is still a member of the **contractors** group in OpenShift.

```
[rhsso@sso ~]$ exit
logout
Connection to sso.ocp4.example.com closed.

[student@workstation ~]$ oc get groups
NAME          USERS
Default SMB Group
admins        admin
contractors  abbyquincy, fricisritcher
developer
editors
ocpadmins    admin
ocpdevs      developer
partners     jayalamont
```

- ▶ 10. Verify that the **fricisritcher** user can still edit resources in the **auth-oidc** OpenShift project. Verify that the changes in the OpenShift user apply only when the user logs out and then logs back in.
 - 10.1. Try to remove the **ubi9-date** pod as the **fricisritcher** user. To do so, change to the web browser window, and from the perspective switcher select **Administrator**. Click **Workloads > Pods**. Select **auth-oidc** from the project drop-down menu. In the list with the **ubi9-date** pod, click the icon with three dots and then click **Delete Pod**.

The screenshot shows the Red Hat OpenShift web console interface. At the top, there's a navigation bar with the Red Hat OpenShift logo and a user dropdown for 'Fricis Ritcher'. Below the header, the project 'auth-oidc' is selected. The main area is titled 'Pods' with a 'Create Pod' button. There are filters for 'Name' and 'Search by name...', and columns for 'Name', 'Status', 'Ready', and 'Restarts'. A single pod named 'ubi9-date' is listed with a status of 'Completed' and '0/1 Ready'. To the right of the pod, a context menu is open with five options: 'Add PodDisruptionBudget', 'Edit labels', 'Edit annotations', 'Edit Pod', and 'Delete Pod'. The 'Delete Pod' option is highlighted with a red border.

Click **Delete**. The user can still edit objects in OpenShift, because it did not synchronize the group membership for the user.

- 10.2. Log out and log back in as the `fricisritcher` user in the OpenShift web console. Click the username in the upper right corner and click **Log out**. Click `RHSSO_OIDC` and log in again as the `fricisritcher` user with `redhat_sso` as the password.
- 10.3. Change to the terminal window. Verify that OpenShift synchronizes the group membership for the user, because the user is not a member of the `contractors` group.

```
[student@workstation ~]$ oc get groups
NAME          USERS
Default SMB Group
admins        admin
contractors   abbyquincy
developer
editors
ocpadmins     admin
ocpdevs       developer
partners      jayalamont
```

- ▶ 11. Remove the `fricisritcher` user from Red Hat SSO, and verify that the user session is still open in the OpenShift web console until the user logs out. Force the user to log out so they cannot reconnect. Verify that the user account is present as a leftover in OpenShift and requires manual removal.

Chapter1 | Authentication and Identity Management

- 11.1. Change to the web browser window and verify that the session for the **fricisritcher** user is not expired. If the session expires, then log in again as the **fricisritcher** user.
- 11.2. Change to the terminal window and remove the **fricisritcher** user from Red Hat SSO, by using the user ID from an earlier step.

```
[student@workstation ~]$ ssh rhsso@sso.ocp4.example.com

[rhsso@sso ~]$ /opt/rh-sso-7.6/bin/kcadm.sh config credentials \
--server https://sso.ocp4.example.com:8080/auth \
--user admin --password redhatocp --realm master
Logging into https://sso.ocp4.example.com:8080/auth as user admin of realm master

[rhsso@sso ~]$ /opt/rh-sso-7.6/bin/kcadm.sh delete \
users/958fde0c-360c-48f8-b5e5-942708fbb36e -r external_providers
```

- 11.3. Change to the web browser window and try to create a project called **fricis-project** as the **fricisritcher** user. To do so, click **Create a new project**. In the name field, enter **fricis-project**, and click **Create**. The user can create a project even when you remove them from Red Hat SSO.
- 11.4. Create a pod in the **fricis-project** project as the **fricisritcher** user. To do so, click **Workloads > Pods** and then **Create Pod**.
- 11.5. Change the default YAML definition file as follows. Then, click **Create**.

```
apiVersion: v1
kind: Pod
metadata:
  name: example
  labels:
    name: httpd
  namespace: fricis-project
spec:
...output omitted...
```

The user can create a pod even if you remove them from Red Hat SSO.

- 11.6. Change to the terminal window, and verify that the user has active access tokens in OpenShift.

```
[rhsso@sso ~]$ exit
logout
Connection to sso.ocp4.example.com closed.

[student@workstation ~]$ oc get oauthaccesstoken
NAME      USER NAME      CLIENT NAME      CREATED      EXPIRES      ...
sha256~4v47...  admin      openshift-...  159m        2023-05-31  ...
...output omitted...
sha256~Wb0Z...  fricisritcher  console      8m42s        2023-05-31  ...
sha256~Y52j...  admin      openshift-...  159m        2023-05-31  ...
...output omitted...
```

- 11.7. Remove all the access tokens for the **fricisritcher** user.

```
[student@workstation ~]$ oc delete oauthaccesstoken \
$(oc get oauthaccesstoken \
-o jsonpath='{.items[?(@.userName=="fricisritcher")].metadata.name}')
oauthaccesstoken.oauth.openshift.io "sha256~Wb0Z..." deleted
```

- 11.8. Change to the web browser window and wait until OpenShift automatically logs out the user. Verify that you can no longer log in as the **fricisritcher** user. Close the web browser window.

- 11.9. Change to the terminal window, and verify that the user and the user identity are still present in OpenShift, and require manual removal.

```
[student@workstation ~]$ oc get users
NAME          UID        FULL NAME      IDENTITIES
abbyquincy   11bc3d49-...  Abby Quincy   RHSSO_OIDC:a175e1b7-...-732021142e84
admin         00fed1ea-...  Administrator  Red Hat Identity Management:dwlk...
developer    514df291-...  Developer User  Red Hat Identity Management:dwlk...
fricisritcher af752e5b-...  Fricis Ritcher  RHSSO_OIDC:958fde0c-...-942708fbb36e
jayalamont   e5f74993-...  Jaya Lamont   RHSSO_OIDC:64898122-...-c9eeeaa0f409
```

- 11.10. Remove the **fricisritcher** user and identity from OpenShift.

```
[student@workstation ~]$ oc delete user fricisritcher
user.user.openshift.io "fricisritcher" deleted
[student@workstation ~]$ oc delete identity \
  RHSSO_OIDC:958fde0c-360c-48f8-b5e5-942708fbb36e
identity.user.openshift.io "RHSSO_OIDC:958fde0c-...-942708fbb36e" deleted
```

Finish

On the **workstation** machine, use the **lab** command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish auth-oidc
```

► Guided Exercise

Solve User Sync Conflicts

Solve conflicts when synchronizing users from more than one identity provider.

Outcomes

- Synchronize users and groups from Red Hat SSO to OpenShift.
- Solve conflicts when synchronizing a user from more than one IdP.

Before You Begin

As the student user on the workstation machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start auth-conflict
```

Instructions

In this exercise, you solve the conflicts that appear when two IdPs provide the same user.

As a use case, the lab script configures Red Hat SSO and htpasswd as OpenShift IdPs.

The htpasswd IdP provides the `abbyquincy` user with `redhat_htpasswd` as the password.

Red Hat SSO provides the `abbyquincy` user with `redhat_sso` as the password in the `external_providers` realm.

- 1. Log in as the `abbyquincy` user through the htpasswd IdP. The user also exists in Red Hat SSO. Try to synchronize the user from both IdPs.
 - 1.1. Log in as the `abbyquincy` user through the htpasswd IdP.



Note

If you get a 401 login error, then wait for a few moments for the `openshift-authentication` pods to re-create, or check them by logging in as an administrator and using the `oc get pods -n openshift-authentication` command.

```
[student@workstation ~]$ oc login -u abbyquincy -p redhat_htpasswd \
https://api.ocp4.example.com:6443
Login successful.
...output omitted...
```

- 1.2. Verify that OpenShift correctly synchronizes the user from the htpasswd IdP the first time that they log in.

Chapter 1 | Authentication and Identity Management

```
[student@workstation ~]$ oc login -u admin -p redhatocp
...output omitted...
[student@workstation ~]$ oc get users
NAME          UID        FULL NAME      IDENTITIES
abbyquincy   d43c5583-...   htpasswd_provider:abbyquincy
admin         00fed1ea-...   Administrator  Red Hat Identity Management:dwlk...
```

- 1.3. Try to log in as the abbyquincy user through Red Hat SSO.

```
[student@workstation ~]$ oc login -u abbyquincy -p redhat_sso
Error from server (InternalError): Internal error occurred: unexpected response:
500
```

- 1.4. Review the logs for the authentication pods to find the login error. OpenShift cannot synchronize the user from Red Hat SSO, because the identity was synchronized from the htpasswd IdP, and the `mappingMethod` parameter for both IdPs is set to `claim`. The pod names might differ on your system.

```
[student@workstation ~]$ oc get pods -n openshift-authentication
NAME           READY   STATUS    RESTARTS   AGE
oauth-openshift-69df8585dc-rhh7j   1/1     Running   0          4m46s
oauth-openshift-69df8585dc-rwvzb   1/1     Running   0          5m42s
oauth-openshift-69df8585dc-zlkzd   1/1     Running   0          5m14s

[student@workstation ~]$ oc logs oauth-openshift-69df8585dc-zlkzd \
-n openshift-authentication
...output omitted...
... Error authenticating login "abbyquincy" with provider "RHSSO_OIDC": user
"abbyquincy" cannot be claimed by identity "RHSSO_OIDC:a175e1b7-6210-40f8-
aeda-732021142e84" because it is already mapped to [htpasswd_provider:abbyquincy]
... AuthenticationError: user "abbyquincy" cannot be claimed by identity
"RHSSO_OIDC:a175e1b7-6210-40f8-aeda-732021142e84" because it is already mapped to
[htpasswd_provider:abbyquincy]
```

- 1.5. To log in by using the OIDC Red Hat SSO, delete the abbyquincy user and its identity.

```
[student@workstation ~]$ oc get user abbyquincy
NAME          UID        FULL NAME      IDENTITIES
abbyquincy   e9538f53-...   htpasswd_provider:abbyquincy
[student@workstation ~]$ oc get identity htpasswd_provider:abbyquincy
NAME          IDP NAME      IDP USER NAME  USER NAME
htpasswd_provider:abbyquincy   htpasswd_provider  abbyquincy  abbyquincy
[student@workstation ~]$ oc delete user abbyquincy
user.user.openshift.io "abbyquincy" deleted
[student@workstation ~]$ oc delete identity htpasswd_provider:abbyquincy
identity.user.openshift.io "htpasswd_provider:abbyquincy" deleted
```

- 1.6. Log in as the abbyquincy user through Red Hat SSO.

```
[student@workstation ~]$ oc login -u abbyquincy -p redhat_sso
Login successful.
...output omitted...
```

- 1.7. Verify that OpenShift correctly synchronizes the user from Red Hat SSO.

```
[student@workstation ~]$ oc login -u admin -p redhatocp
...output omitted...
[student@workstation ~]$ oc get user abbyquincy
NAME          UID        FULL NAME      IDENTITIES
abbyquincy   54f7...   Abby Quincy    RHSSO_OIDC:a175...
```

- ▶ 2. To avoid conflicts when two IdPs provide the same user, change the OIDC Red Hat SSO and htpasswd mappingMethod parameter from claim to add, so OpenShift adds the identity from Red Hat SSO and htpasswd to the user no matter which one is used first.

- 2.1. Change the OIDC Red Hat SSO mappingMethod parameter from claim to add.

```
apiVersion: config.openshift.io/v1
kind: OAuth
metadata:
  name: cluster
spec:
  identityProviders:
  ...output omitted...
  - htpasswd:
    fileData:
      name: htpasswd-secret
    mappingMethod: add
    name: htpasswd_provider
    type: HTPasswd
  - openID:
  ...output omitted...
    mappingMethod: add
    name: RHSSO_OIDC
    type: OpenID
```



Note

You can find the YAML file for the OAuth CR in the ~/D0380/labs/auth-conflict/oauth_config.yml file.

- 2.2. Apply the OAuth CR YAML file to the cluster.

```
[student@workstation ~]$ oc apply -f ~/D0380/labs/auth-conflict/oauth_config.yml
oauth.config.openshift.io/cluster configured
```

- 2.3. Wait for the OAuth pods to redeploy.

```
[student@workstation ~]$ watch oc get pods -n openshift-authentication
Every 2.0s: oc get pods -n openshift-authentication workstation: ...
NAME                  READY   STATUS    RESTARTS   AGE
oauth-openshift-f94c8d5fd-8w9nb   1/1     Running   0          92s
oauth-openshift-f94c8d5fd-jnt8r   1/1     Running   0          63s
oauth-openshift-f94c8d5fd-vk79l   1/1     Running   0          35s
^C
```

2.4. Try to log in again as the abbyquincy user through the htpasswd IdP.

```
[student@workstation ~]$ oc login -u abbyquincy -p redhat_htpasswd
Login successful.
...output omitted...
```

2.5. Verify that OpenShift correctly synchronizes the user from the htpasswd IdP. Verify that the abbyquincy user has two identities, one from each of the two IdPs.

```
[student@workstation ~]$ oc login -u admin -p redhatocp
...output omitted...
[student@workstation ~]$ oc get user abbyquincy
NAME      UID      FULL NAME      IDENTITIES
abbyquincy  54f7...  Abby Quincy  RHSSO_OIDC:a175..., 
                           htpasswd_provider:abbyquincy
```

Finish

On the **workstation** machine, use the **lab** command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish auth-conflict
```

Token and Client Certificate Authentication with kubeconfig Files

Objectives

- Generate a token and a client certificate and add them to a kubeconfig file.

External Client Authentication

In certain scenarios, you might authenticate external clients, such as CI/CD pipelines or monitoring tools, to the Kubernetes cluster. Kubernetes enables external clients to authenticate to the Kubernetes API by embedding either a client certificate or an authentication token into a `kubeconfig` configuration file. This feature ensures that only authorized external clients can access the Kubernetes cluster.

Service Accounts

A *service account* (SA) is a Kubernetes resource that provides an identity for processes that run in a pod. SA tokens authenticate the interactions between components within the OpenShift cluster, as well as with external resources. OpenShift uses SA tokens to grant permissions to pods and other resources to interact with the Kubernetes API server. External services use SA tokens to access other OpenShift resources, or to access the Kubernetes API. With SAs, you can control API access without the need to borrow a regular user's credentials. SAs are specific to a particular project and cannot be directly shared across projects.

For example, you can use SAs in the following scenarios:

- A replication controller makes API calls to create or delete pods.
- A pod that collects logs might need access to certain log storage resources.
- A backup and restore tool might require access to the cluster's configuration and data to back up and restore data.
- An external application makes monitoring or integration API calls.

Service Account Tokens

When you create an SA in OpenShift, the SA automatically contains two secrets: an API token and credentials for the OpenShift Container Registry. The automatically generated API token and credentials never expire. However, you can revoke the token and credentials by deleting the secrets, because OpenShift automatically generates a new secret.

Starting with Kubernetes 1.23, SAs do not automatically create the secrets that contain long-term credentials for accessing the Kubernetes API. Thus, you must obtain the credentials by using the `TokenRequest` API. The tokens that the `TokenRequest` API provides are more secure than the tokens that are stored in secrets, because the `TokenRequest` tokens have a bounded lifetime; other API clients cannot read the `TokenRequest` tokens; and OpenShift automatically invalidates the `TokenRequest` tokens when the pod that they are mounted into is deleted.

Although Red Hat OpenShift Container Platform (RHOC) 4.14 is based on Kubernetes 1.27, it still creates the SA token secrets to communicate with the Kubernetes API server, because some features and workloads need the SA secrets. However, this behavior will change in a future release.

Chapter 1 | Authentication and Identity Management

Thus, although you can use the automatically generated token secret to authenticate the SA to the Kubernetes API server, Red Hat recommends using the `TokenRequest` API to generate tokens that are bound to the SA, and not relying on the automatically generated token secret.



Note

You can still manually create long-lived API tokens for SAs. However, Red Hat recommends using long-lived API tokens only if you cannot use the `TokenRequest` API and if the security exposure of a non-expiring token is acceptable to you. To manually create long-lived API tokens for SAs, refer to https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/nodes/index#nodes-pods-secrets-creating-sa_nodes-pods-secrets

The following diagram summarizes the authentication flow by using the `TokenRequest` API:



- ➊ An application that is running in a pod needs to authenticate to the Kubernetes API. The `kubelet` agent requests a bound SA account token from the `TokenRequest` API.
- ➋ The `kubelet` agent receives the SA token from the `TokenRequest` API and mounts it into the pod.
- ➌ The application reads the SA token and uses it to authenticate to the Kubernetes API.

For external applications, you can manually generate a bound SA token through the `TokenRequest` API, and use it as the bearer token for authentication to the Kubernetes API. Generating a bound SA token in OpenShift by using the `TokenRequest` API does not remove or override the automatically generated token secret.

After creating the SA, use the following command to generate a bound SA token by using the `TokenRequest` API:

```
[user@host ~]# oc create token SA_name -n project
```

You can set the lifetime for the SA token by using the `--duration` option. The default lifetime for the SA token is one hour. You must manually refresh the manually generated SA token before it expires, so that your application can continue authenticating to the Kubernetes API.

The following example shows a decoded JWT SA token:

```

HEADER
...output omitted...
PAYLOAD:DATA
{
  "aud": [
    "https://kubernetes.default.svc"
  ],
  "exp": 1698059299, ➊
  "iat": 1698055699, ➋
}
  
```

```

"iss": "https://kubernetes.default.svc",
"kubernetes.io": {
  "namespace": "my-project", ③
  "serviceaccount": {
    "name": "my-sa", ④
    "uid": "3acb4630-babe-4e85-996d-cbc80f62146f"
  }
},
"nbf": 1698055699,
"sub": "system:serviceaccount:my-project:my-sa"
}
SIGNATURE
...output omitted...

```

- ①** SA token expiration date in Unix epoch format
- ②** SA token issued date in Unix epoch format
- ③** Project for the SA token
- ④** Name for the SA

OpenShift assigns an SA to every pod that you deploy. By default, OpenShift creates the following SAs for every project:

builder

OpenShift uses this SA to build pods. By default, this SA has the `system:image-builder` role, so the resource can push images to any image stream in the project by using the internal Docker registry.

deployer

OpenShift uses this SA in deployment pods. By default, this SA has the `system:deployer` role, so the resource can view and modify replication controllers and pods in the project. This SA exists only for applications that use OpenShift deployment configuration resources.

default

OpenShift assigns this default SA to pods if you do not specify a different SA when you create the pods.

Manage an SA

You can manage SAs by using the typical `oc` commands, such as `create`, `get`, or `describe`. You can also grant roles to SAs the same as for regular users.

If you grant roles to an SA, then you must use the name of the SA together with the name of the project and the `system:serviceaccount` string, according to the following syntax:

```
system:serviceaccount:project:name
```

You can also use the `-z` option to avoid using the long `system:serviceaccount:project:sa_name` name, and use instead the short `sa_name` name.

Whenever you create an SA, it is automatically a member of the following groups:

system:serviceaccounts

This group includes all the SAs in the cluster.

Chapter 1 | Authentication and Identity Management

`system:serviceaccounts:project`

This group includes all the SAs in the specified project.

Client Certificate Authentication

Client certificate authentication in Kubernetes clusters refers to the process of authenticating clients, such as users or services, which access the Kubernetes cluster by using TLS client certificates.

By default, OpenShift provides an internal certificate authority (CA). The OpenShift internal CA is a built-in component of the OpenShift cluster that manages and issues digital certificates in the cluster. The internal CA provides a trusted source for generating X.509 certificates for secure communication, authentication, and encryption in the OpenShift cluster.

The Kubernetes API server requires client authentication by using client certificates. OpenShift is preconfigured to trust the client certificates that the OpenShift internal CA signs.



Note

You can also configure additional client CAs for the Kubernetes API server. For more information about configuring additional client CAs for the Kubernetes API server, refer to <https://access.redhat.com/solutions/6054271>

This feature is mainly used to generate a client certificate for an administrator user, such as the predefined `system:admin` user, to use this client certificate as a backdoor for cluster administrators in the event of failure of the IdP that provides the administrator credentials.

Although Red Hat does not recommend doing so, you can also use client certificates in certain scenarios when running outside the cluster, such as CI/CD pipelines, automation playbooks, or monitoring tools. These clients need to present valid client certificates during the TLS handshake to establish a secure connection with the API server. This mechanism ensures that only authorized clients can access the API server from outside the cluster.



Warning

Red Hat recommends using SAs to run automation outside the cluster whenever you can, instead of using client certificates, because a certificate cannot be revoked in OpenShift. Revoking a client certificate would require invalidating all client certificates that the current CA ever signed, and creating replacement certificates for all client users and applications. For more information about this topic, refer to <https://github.com/kubernetes/kubernetes/issues/18982>

Create a Client Certificate

OpenShift assigns the username and the groups for the user by using the common name (CN) and the organization (O) fields from the certificate, respectively. You can use role-based access control (RBAC) rules to provide the minimal rights to the client account to perform the job. For example, you could provide view permissions for the entire cluster to a monitoring tool, or edit permissions on selected projects to a CI/CD pipeline.

To create client certificates by using the internal OpenShift CA, you must follow these steps:

1. Create a certificate signing request (CSR): The first step is to create a CSR for the client. You can use the OpenSSL tool to create the CSR. This request includes the client's information and the public key. You can use more than one group for the user if you require it.

```
[user@host ~]$ openssl req -nodes -newkey rsa:4096 -keyout key_filename \
-subj "/O=group1/O=group2/CN=username" -out csr_filename
```

2. Create the CSR YAML resource manifest. The following example shows the parameters for a CSR:

```
apiVersion: certificates.k8s.io/v1
kind: CertificateSigningRequest
metadata:
  name: csr_name 1
spec:
  signerName: kubernetes.io/kube-apiserver-client 2
  expirationSeconds: 604800 # one week 3
  request: $(base64 -w0 csr_filename) 4
  usages: 5
    - client auth
```

- 1** The CSR name.
 - 2** OpenShift provides the `kube-apiserver-client` built-in signer of certificates that the API server accepts. OpenShift provides different built-in signers that you can use to sign certificates. For more information, refer to <https://kubernetes.io/docs/reference/access-authn-authz/certificate-signing-requests/#kubernetes-signers>
 - 3** Expiration time for the certificate in seconds. If you do not specify a value, then the default value is 30 days.
 - 4** Contains the OpenSSL X.509 CSR, which is encoded in Base64.
 - 5** Specifies the use cases for the client certificate. It must include the `client auth` usage, which indicates that the certificate is intended for client authentication.
3. Submit the CSR to the Kubernetes API server. The API server interacts with the internal CA to process the CSR.

```
[user@host ~]$ oc apply -f csr.yaml
```

4. Review, approve, and sign the CSR: You can review the CSR details by using the `oc describe csr csr_name` command. After reviewing the CSR details, use the following command so the internal CA signs the CSR to generate the client certificate:

```
[user@host ~]$ oc adm certificate approve csr_name
```

5. Retrieve the client certificate: After you approve the CSR and the OpenShift internal CA generates the certificate, you can retrieve the signed certificate from the API server.

```
[user@host ~]$ oc get csr csr_name -o jsonpath='{.status.certificate}' \
| base64 -d > certificate_filename
```

6. Distribute the certificate: The signed certificate, together with the private key that generates the CSR, enable the client to authenticate to the cluster.

OpenShift CLI Configuration Files

You can use a `kubeconfig` file as a command-line interface (CLI) configuration file to set up profiles for use with the Kubernetes `kubectl` and OpenShift `oc` CLI tools. Moreover, most Kubernetes client libraries use `kubeconfig` files in the same way as the `kubectl` and `oc` CLI tools. Use `kubeconfig` files to authenticate external applications to the cluster by storing tokens and client certificates inside the `kubeconfig` files.

The `kubeconfig` file is defined as a YAML file that contains clusters, users, and contexts.

clusters

The `clusters` parameter in the `kubeconfig` file contains information about the OpenShift clusters, such as the IP or fully qualified domain name (FQDN), or the CA.

users

The `users` parameter contains the user credentials to interact with the Kubernetes API. This parameter contains information such as the username, the user password, or the user token.

contexts

The `contexts` parameter contains information about the combination of a cluster and a user to interact with the Kubernetes API. Whenever you run an `oc` command, you reference a context inside the `kubeconfig` file.

kubeconfig File Details

When you run any `oc` command, OpenShift reads a `kubeconfig` file in the following ways in turn:

1. The specified file in the `--kubeconfig` option, if you use it
2. The specified file in the `KUBECONFIG` environment variable, if it is set
3. The default `kubeconfig` file in `~/.kube/config`

When you log in to the OpenShift cluster through the `oc login` command for the first time, OpenShift creates a `kubeconfig` file in the default location at `~/.kube/config`, if the file does not exist. You can add more authentication and connection details to the `kubeconfig` file automatically by using the `oc login` and `oc project` commands, or by manually editing the `kubeconfig` file.

You can read the details from your `kubeconfig` file by using the `oc config view` command or by opening the file with a text editor. The following example shows the parameters for a `kubeconfig` file from an OpenShift cluster:

```
apiVersion: v1
clusters: ①
- cluster:
  server: https://api.ocp4.prod.com:6443
  name: production
- cluster:
  server: https://api.ocp4.stage.com:6443
  certificate-authority: ocp-apiserver-cert.crt
  name: stage
users: ②
- name: admin-production
  user:
    token: REDACTED
- name: admin-stage
```

```

user:
  client-certificate: admin-stage.crt
  client-key: tls.key
contexts: ③
- context:
  cluster: production
  namespace: prod-app
  user: admin-production
  name: prod-app/api-ocp4-prod-com:6443/admin-production
- context:
  cluster: stage
  namespace: demo-app
  user: admin-stage
  name: demo-app/api-ocp4-stage-com:6443/admin-stage
current-context: demo-app/api-ocp4-stage-com:6443/admin-stage
kind: Config
preferences: {}

```

- ① The list of all the clusters that you already connected to. The example defines two clusters, **production** and **stage**, with their FQDNs that are defined in the **server** parameter. The **stage** cluster definition also contains the public certificate from the API server in the **certificate-authority** parameter.
- ② The list of all the users that you already connected to the cluster. The example defines two users, **admin-production** and **admin-stage**. The **admin-production** user definition uses the token that is defined in the **token** parameter to authenticate to the API server. OpenShift truncates the token information, to prevent the configuration file from becoming too long. You can use the **--raw** option in the **oc config view** command to show the token information. The **admin-stage** user authenticates to the cluster by using a client certificate that the internal OpenShift CA previously signed, as defined in the **client-certificate** parameter, and the key that signed the client certificate, as defined in the **client-key** parameter.
- ③ The list of contexts that you can reference when using the **oc** command. Every context specifies a combination of a user and a cluster that are already defined in the **kubeconfig** file, through the **cluster** and **user** parameters, and a project through the **namespace** parameter.

Configure CLI Profiles

Although you can set most **kubeconfig** file parameters by using the **oc login** and **oc project** OpenShift commands, you can use the **oc config** command to manually configure these parameters if needed, instead of directly modifying the file. You can use the **oc config** command with the **--kubeconfig=file** option to create and store **kubeconfig** files with different parameters that you can use when required. The **oc config** command comes from the Kubernetes **kubectl** CLI tool, with no modifications from OpenShift.

Use the following **oc config** subcommands to manually configure your **kubeconfig** files:

set-cluster

Creates a cluster entry in the **kubeconfig** file. This subcommand accepts different cluster parameters as the server IP or the CA file.

set-credentials

Creates a user entry in the **kubeconfig** file. This subcommand accepts different user parameters as the username, the user password, or the token.

set-context

Creates a context entry in the `kubeconfig` file. Use this subcommand to define a context to specify a combination of a user and a cluster. You can also set the OpenShift project.

use-context

Sets the current context.

**Note**

For more information about manually configuring the `kubeconfig` files by using the `oc config` command, refer to https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/cli_tools/index#manual-configuration-of-cli-profiles_managing-cli-profiles

User Impersonation

User impersonation in OpenShift enables certain users or SAs to act on behalf of other users or SAs with different permissions and roles. This feature is useful when administrators or privileged users need to perform actions on behalf of regular users, or when SAs need to execute tasks on behalf of other SAs, or for regular users to perform actions with administrative permissions. For example, a system administrator can debug an RBAC rule by impersonating another user and verifying whether OpenShift accepts or denies a request. As another example, you want system administrators to escalate privileges when it is necessary instead of them logging in as cluster administrators.

System Administrator That Impersonates a Regular User

As a system administrator, use the `--as` and `--as-group` options when using the `oc` command to impersonate a user or group. Use the `oc auth can-i` command to test the user access to a particular resource in the cluster. Use the `-n` option to specify a project for the request, or use the `-A` option to verify the action in all the projects. Thus, use the following command to test the RBAC rules for a particular user or group by impersonating them:

```
[user@host ~]$ oc auth can-i command --as user_to_impersonate \
--as-group group_to_impersonate
```

You can also list all the permissions for a specific user or group by using the `oc auth can-i --list` command.

Regular User That Impersonates Another User

To grant a regular user permissions to impersonate another user, you must create a custom role with the appropriate permissions, and then create a role binding to assign the role to the user.

For example, to allow a regular user to run commands as an administrator, you can create the following cluster role, which enables anyone to impersonate the administrator user:

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: sudo-admin
rules:
- apiGroups: [""] ①
```

```
resources: ["users"]
verbs: ["impersonate"]
resourceNames: ["admin"]
```

- ➊ The core API group is identified with an empty string.

Then, bind the role to the user:

```
[user@host ~]$ oc create clusterrolebinding binding_name \
--clusterrole sudo-admin --user regularuser_
```

The user can impersonate the `admin` user by using the `--as=admin` option. The following example shows how the user cannot retrieve the node information when using their account, but can retrieve that information when impersonating the `admin` user:

```
[user@host ~]$ oc get nodes
Error from server (Forbidden): nodes is forbidden:
User "regularuser" cannot list resource "nodes" in API group "" at the cluster
scope

[user@host ~]$ oc get nodes --as admin
NAME      STATUS    ROLES          AGE     VERSION
master01   Ready     control-plane,master   62d    v1.25.7+eab9cc9
worker01   Ready     worker           57d    v1.25.7+eab9cc9
```



References

For more information about using SAs for applications on OpenShift, refer to the *Using Service Accounts in Applications* section in the Red Hat OpenShift Container Platform 4.14 *Authentication and authorization* documentation at https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/authentication_and_authorization/index#service-accounts-overview_using-service-accounts

For more information about how to configure CLI profiles on OpenShift, refer to the *Managing CLI Profiles* section in the Red Hat OpenShift Container Platform 4.14 *CLI Tools* documentation at https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/cli_tools/index#about-switches-between-cli-profiles_managing-cli-profiles

Kubernetes: Certificates and Certificate Signing Requests

<https://kubernetes.io/docs/reference/access-authn-authz/certificate-signing-requests/>

Kubernetes: Authenticating

<https://kubernetes.io/docs/reference/access-authn-authz/authentication/#user-impersonation>

► Guided Exercise

Token and Client Certificate Authentication with kubeconfig Files

Generate a token and a client certificate and add them to a kubeconfig file.

Outcomes

- Generate a service account (SA) token for authenticating a script that runs outside the OpenShift cluster.
- Generate a client certificate for a system administrator as a backdoor account for the OpenShift cluster.

Before You Begin

As the student user on the workstation machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start auth-tls
```

Instructions

Your company provides you a Bash script to verify the health of an OpenShift cluster. Because this script requires authentication to the cluster, you must generate an SA token to authenticate it. Use the `TokenRequest` API to generate the SA token. Do not use the automatically generated token secret. The username for the SA must be `health-robot`. The SA token must be valid for one week. The script verifies the health status for all the pods in the OpenShift cluster, and creates the `/tmp/cluster.log` file if a pod is in a pending, failed, or unknown state. Thus, because you need to read the pod information for the whole cluster, you must assign the `cluster-reader` cluster role to the SA. You can find the Bash script in the `~/D0380/labs/auth-tls/ocp_health.sh` file. Create a RHEL cron job to run the script every minute.

Your company also needs a backdoor administrator account to use if the IdP that provides the administrator account fails. As a solution, you decide to create a client certificate. The username for the administrator certificate must be `admin-backdoor`. The `admin-backdoor` account must be part of the `backdoor-administrators` group, and you must assign the `cluster-admin` cluster role to that group to have administrator permissions for the cluster. The expiration for the administrator certificate must be one week.

- 1. Create the `health-robot` SA in the `auth-tls` OpenShift project and assign it the `cluster-reader` role.
 - 1.1. Connect to the OpenShift cluster as the `admin` user with `redhatocp` as the password.

Chapter1 | Authentication and Identity Management

```
[student@workstation ~]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
Login successful.

You have access to 70 projects, the list has been suppressed. You can list all
projects with 'oc projects'

Using project "default".
Welcome! See 'oc help' to get started.
```

1.2. Change to the auth-tls project.

```
[student@workstation ~]$ oc project auth-tls
Now using project "auth-tls" on server "https://api.ocp4.example.com:6443".
```

1.3. Create the health-robot SA in the auth-tls project.

```
[student@workstation ~]$ oc create sa health-robot
serviceaccount/health-robot created
```

1.4. Assign the cluster-reader cluster role to the health-robot SA so the SA can retrieve information from most of the objects in the cluster.

```
[student@workstation ~]$ oc adm policy add-cluster-role-to-user \
cluster-reader system:serviceaccount:auth-tls:health-robot
clusterrole.rbac.authorization.k8s.io/cluster-reader added:
"system:serviceaccount:auth-tls:health-robot"
```

- ▶ 2. Create the health-robot SA token and store it in the HEALTHROBOT_TOKEN variable. Create a kubeconfig file with the health-robot user credentials. Store the kubeconfig file in the ~/D0380/labs/auth-tls/robot-cert/health-robot.config file.

2.1. Generate the health-robot SA account token and store it in the HEALTHROBOT_TOKEN variable. The SA token must be valid for one week, which is 604800 seconds.

```
[student@workstation ~]$ HEALTHROBOT_TOKEN=$(oc create token -n auth-tls \
health-robot --duration 604800s)
```

2.2. Change to the ~/D0380/labs/auth-tls directory.

```
[student@workstation ~]$ cd D0380/labs/auth-tls
```

2.3. Add the health-robot user credentials to the kubeconfig file.

```
[student@workstation auth-tls]$ oc config set-credentials health-robot \
--token $HEALTHROBOT_TOKEN --kubeconfig robot-cert/health-robot.config
User "health-robot" set.
```

Chapter 1 | Authentication and Identity Management

- 2.4. Set the cluster options in the kubeconfig file.

```
[student@workstation auth-tls]$ oc config set-cluster \
$(oc config view -o jsonpath='{.clusters[0].name}') \
--server https://api.ocp4.example.com:6443 \
--kubeconfig robot-cert/health-robot.config
Cluster "api-ocp4-example-com:6443" set.
```

- 2.5. Set the context for the health-robot user.

```
[student@workstation auth-tls]$ oc config set-context health-robot \
--cluster $(oc config view -o jsonpath='{.clusters[0].name}') \
--namespace auth-tls --user health-robot \
--kubeconfig robot-cert/health-robot.config
Context "health-robot" created.
```

- 2.6. Use the context for the health-robot user.

```
[student@workstation auth-tls]$ oc config use-context health-robot \
--kubeconfig robot-cert/health-robot.config
Switched to context "health-robot".
```

- 3. Verify the permissions for the health-robot SA. Use the health-robot SA token through the kubeconfig file and user impersonation.

- 3.1. Verify the identity for the health-robot user.

```
[student@workstation auth-tls]$ oc whoami \
--kubeconfig robot-cert/health-robot.config
system:serviceaccount:auth-tls:health-robot
```

- 3.2. Verify that the health-robot SA can retrieve information from the users in the OpenShift cluster.

```
[student@workstation auth-tls]$ oc auth can-i get users -A \
--as system:serviceaccount:auth-tls:health-robot
yes
```

- 3.3. Verify that the health-robot SA cannot create projects in the OpenShift cluster.

```
[student@workstation auth-tls]$ oc auth can-i create project -A \
--as system:serviceaccount:auth-tls:health-robot
no
```

- 3.4. Verify that the health-robot SA can retrieve information from the pods in the OpenShift cluster.

```
[student@workstation auth-tls]$ oc auth can-i get pods -A \
--as system:serviceaccount:auth-tls:health-robot
yes
```

Chapter1 | Authentication and Identity Management

- 3.5. Run the `cluster-health.sh` script to verify the status of the pods in the OpenShift cluster. The script uses the credentials that are stored in the `~/D0380/labs/auth-tls/robot-cert/health-robot.config` file.

```
[student@workstation auth-tls]$ sh cluster-health.sh
Connected to the cluster as the 'system:serviceaccount:auth-tls:health-robot' user
✓ OpenShift is reachable and up, at version: '4.14.12'
✓ All pods are either running or succeeded.
```

- 4. Create a local RHEL CRON job to execute the `cluster-health.sh` script every minute. Create a failing pod, and verify that the script creates a log file with the projects that contain pods in a pending, failed, or unknown state.

- 4.1. Open the `crontab` file for the student with the default text editor.

```
[student@workstation auth-tls]$ crontab -e
```

- 4.2. Insert the following line to define the CRON job.

```
* * * * * ~/D0380/labs/auth-tls/cluster-health.sh
```

- 4.3. Press Esc and type `:wq` to save the changes and exit the editor. When the editor exits, you should see the following output:

```
crontab: installing new crontab
[student@workstation auth-tls]$
```

- 4.4. Create a failing pod in the `auth-tls` project.

```
[student@workstation auth-tls]$ oc run failing-pod \
--image registry.ocp4.example.com:8443/failing-pod
...output omitted...
pod/failing-pod created
```

- 4.5. Wait for a few minutes and verify the log file. The log file contains the date and the projects with pods in a pending, failed, or unknown state.

```
[student@workstation auth-tls]$ cat /tmp/cluster.log
...output omitted...
Tue Jul 11 10:29:04 EDT 2023
x Namespaces with pending pods:
x auth-tls
```

- 5. Create a client certificate for a backdoor administrator account to use if the IdP that provides the administrator account fails. The username for the administrator certificate must be `admin-backdoor`, and the user must be part of the `backdoor-administrators` group with the `cluster-admin` cluster role. Set the expiration time for the certificate to one week.

- 5.1. Create the `backdoor-administrators` group in the OpenShift cluster.

Chapter 1 | Authentication and Identity Management

```
[student@workstation auth-tls]$ oc adm groups new backdoor-administrators
group.user.openshift.io/backdoor-administrators created
```

- 5.2. Assign the `cluster-admin` cluster role to the `backdoor-administrators` group so users in that group can act as cluster administrators.

```
[student@workstation auth-tls]$ oc adm policy add-cluster-role-to-group \
  cluster-admin backdoor-administrators
clusterrole.rbac.authorization.k8s.io/cluster-admin added: "backdoor-
administrators"
```

- 5.3. Create the `admin-cert` directory.

```
[student@workstation auth-tls]$ mkdir admin-cert
```

- 5.4. Create an OpenSSL certificate signing request (CSR) for the `admin-backdoor` username. The user must be part of the `backdoor-administrators` group.

```
[student@workstation auth-tls]$ openssl req -newkey rsa:4096 -nodes \
  -keyout tls.key -subj "/O=backdoor-administrators/CN=admin-backdoor" \
  -out admin-cert/admin-backdoor-auth.csr
Generating a RSA private key
.....+++++
.....+++++
writing new private key to 'tls.key'
-----
```

- 5.5. Define a signing request resource in a YAML file called `admin-backdoor-csr.yaml` inside the `admin-cert` directory. Set the expiration time for the certificate to one week. The name for the CSR must be `admin-backdoor-access`.

```
[student@workstation auth-tls]$ cat << EOF >> admin-cert/admin-backdoor-csr.yaml
apiVersion: certificates.k8s.io/v1
kind: CertificateSigningRequest
metadata:
  name: admin-backdoor-access
spec:
  signerName: kubernetes.io/kube-apiserver-client
  expirationSeconds: 604800 # one week
  request: $(base64 -w0 admin-cert/admin-backdoor-auth.csr)
  usages:
  - client auth
EOF
```

- 5.6. Create the CSR for the `admin-backdoor` user in the OpenShift cluster.

```
[student@workstation auth-tls]$ oc create -f admin-cert/admin-backdoor-csr.yaml
certificatesigningrequest.certificates.k8s.io/admin-backdoor-access created
```

- 5.7. Verify the CSR status for the `admin-backdoor` user in the OpenShift cluster.

```
[student@workstation auth-tls]$ oc get csr admin-backdoor-access
NAME          AGE   SIGNERNAME REQUESTOR REQUESTEDDURATION CONDITION
admin-backdoor-access 3m28s kuberne... admin      7d           Pending
```

- 5.8. Approve the CSR in the OpenShift cluster.

```
[student@workstation auth-tls]$ oc adm certificate approve admin-backdoor-access
certificatesigningrequest.certificates.k8s.io/admin-backdoor-access approved
```

- 5.9. Verify again the status for the CSR for the admin-backdoor user in the OpenShift cluster.

```
[student@workstation auth-tls]$ oc get csr admin-backdoor-access
NAME          AGE   SIGNERNAME REQUESTOR REQUESTEDDURATION CONDITION
admin-backdoor-access 4m23s kuberne... admin      7d           Approved,Issued
```

- ▶ 6. Create a kubeconfig file with the admin-backdoor user credentials. Store the kubeconfig file in the ~/D0380/labs/auth-tls/admin-cert/admin-backdoor.config file.

- 6.1. Extract the signed client certificate for the admin-backdoor user. Store the signed client certificate in the ~/D0380/labs/auth-tls/admin-cert/admin-backdoor-access.crt file.

```
[student@workstation auth-tls]$ oc get csr admin-backdoor-access \
-o jsonpath='{.status.certificate}' \
| base64 -d > admin-cert/admin-backdoor-access.crt
```

- 6.2. Add the admin-backdoor user credentials to the kubeconfig file.

```
[student@workstation auth-tls]$ oc config set-credentials admin-backdoor \
--client-certificate admin-cert/admin-backdoor-access.crt --client-key tls.key \
--embed-certs --kubeconfig admin-cert/admin-backdoor.config
User "admin-backdoor" set.
```

- 6.3. Get the public certificate from the API server.

```
[student@workstation auth-tls]$ openssl s_client -showcerts \
-connect api.ocp4.example.com:6443 </dev/null 2>/dev/null|openssl x509 \
-outform PEM > ocp-apiserver-cert.crt
```

- 6.4. Set the cluster options in the kubeconfig file.

```
[student@workstation auth-tls]$ oc config set-cluster \
$(oc config view -o jsonpath='{.clusters[0].name}') \
--certificate-authority ocp-apiserver-cert.crt --embed-certs=true \
--server https://api.ocp4.example.com:6443 \
--kubeconfig admin-cert/admin-backdoor.config
Cluster "api-ocp4-example-com:6443" set.
```

Chapter 1 | Authentication and Identity Management

6.5. Set the context for the admin-backdoor user.

```
[student@workstation auth-tls]$ oc config set-context admin-backdoor \
--cluster $(oc config view -o jsonpath='{.clusters[0].name}') \
--namespace default --user admin-backdoor \
--kubeconfig admin-cert/admin-backdoor.config
Context "admin-backdoor" created.
```

6.6. Use the context for the admin-backdoor user.

```
[student@workstation auth-tls]$ oc config use-context admin-backdoor \
--kubeconfig admin-cert/admin-backdoor.config
Switched to context "admin-backdoor".
```

- 7. Verify the permissions for the admin-backdoor user. Use the admin-backdoor user client certificate through the kubeconfig file and user impersonation.

7.1. Verify the identity for the admin-backdoor user.

```
[student@workstation auth-tls]$ oc whoami \
--kubeconfig admin-cert/admin-backdoor.config
admin-backdoor
```

7.2. Verify that the admin-backdoor user can retrieve information from the users in the OpenShift cluster.

```
[student@workstation auth-tls]$ oc auth can-i get users -A \
--as admin-backdoor --as-group backdoor-administrators
yes
```

7.3. Verify that the admin-backdoor user can create users in the OpenShift cluster.

```
[student@workstation auth-tls]$ oc auth can-i create users -A \
--as admin-backdoor --as-group backdoor-administrators
yes
```

7.4. Verify that the admin-backdoor user can create projects in the OpenShift cluster.

```
[student@workstation auth-tls]$ oc auth can-i create project -A \
--as admin-backdoor --as-group backdoor-administrators
yes
```

7.5. Change to the student HOME directory.

```
[student@workstation auth-tls]$ cd
```

Finish

On the workstation machine, use the lab command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish auth-tls
```

► Lab

Authentication and Identity Management

Define the concepts and custom resources of the OpenShift OAuth server, and explain how these resources augment Kubernetes authentication.

Configure an LDAP identity provider and automate group synchronization between OpenShift OAuth and an LDAP server.

Configure an OIDC identity provider and automate group synchronization between OpenShift OAuth and an OIDC server.

Generate a token and a client certificate and add them to a `kubeconfig` file.

Outcomes

- Use a client certificate for a system administrator account to recover access to the OpenShift cluster.
- Configure RHDS as an LDAP identity provider (IdP) for OpenShift.
- Configure Red Hat SSO as an OIDC IdP for OpenShift.
- Synchronize users and groups from the LDAP and Red Hat SSO IdPs to OpenShift.
- Generate a service account (SA) token for authenticating a script that runs from outside the OpenShift cluster.

Before You Begin

As the student user on the `workstation` machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start auth-review
```

Instructions

A cluster administrator in your company erroneously removed all the configured IdPs from the OAuth server. You must use a backdoor administrator certificate to re-create the IdPs. For this purpose, create a `kubeconfig` file that contains the client certificate, the certificate key, and the OpenShift CA server certificate. The certificate is already approved in the OpenShift CA server with administrator privileges.

After you create the administrator `kubeconfig` file, restore the LDAP IdP that contains one of the company managers, Payden Tomcheck, with the `paydентomcheck` username and the `redhat123` password. You must configure the LDAP automated group synchronization, because the user is part of the `administrators` group. This group must have administrative privileges to the cluster.

Then, configure an OIDC IdP. This IdP provides the consultant user, Lauren Chan, with the `laurenchan` username and the `redhat_sso` password. You must configure the OIDC IdP to

automatically synchronize the groups, because the user is part of the `consultants` group. This group must have read privileges to the `auth-review` project.

Finally, configure an external monitoring app, which runs outside the cluster and uses an SA account token that must have read access to the cluster. You must create the SA account in the `auth-review` project with the `audit-bot` name.

1. Create a `kubeconfig` file that contains the client certificate, the certificate key, and the OpenShift CA server certificate.
You must create the `kubeconfig` file in the `~/D0380/labs/auth-review/certificate` directory with the `admin.config` name.
You can find the `admin-access.crt` client certificate, the `tls.key` certificate key, and the `ocp-apiserver-cert.crt` CA server certificate in the same directory. Use `api-ocp4-example-com:6443` as the cluster name and `https://api.ocp4.example.com:6443` as the cluster IP in the `kubeconfig` file.
2. Use the administrator `kubeconfig` file to configure the LDAP IdP. Use the LDAP IdP information from the following table:

Parameter	Value
DN (-D)	<code>cn=Directory Manager</code>
URI (-H)	<code>ldaps://rhds.ocp4.example.com</code>
Password (-w)	<code>redhatocp</code>

Use the `~/D0380/labs/auth-review/rhds` directory to create any necessary files.

You can find the RHDS certificate in that directory with the `rhds_ca.crt` name. Use the `rhds-ldap-secret` and `rhds-ca-config-map` names for the RHDS secret and configuration map respectively. Use `Red Hat Directory Server` as the name for the RHDS LDAP IdP.

You can find an incomplete example for the OAuth CR in the `rhds-ldap-idp.yaml` file.

3. Configure the LDAP automated group synchronization in a project called `auth-ldapsync`.
Create the `auth-ldapsync` project, and create the `ldap-group-syncer` SA. Create the `ldap-group-syncer` cluster role with `get`, `list`, `create`, and `update` permissions to the `groups` resource, and assign the cluster role to the SA.
Use the `ldap-secret` and `ldap-config` names for the RHDS secret and configuration map, respectively. Mount the secret and the configuration map in the `/etc/secrets/bindPassword` and `/etc/config/ca.crt` files, respectively.
Use the `~/D0380/labs/auth-review/group-sync` directory to create any necessary files. You can find the RHDS certificate in that directory with the `rhds_ca.crt` name. You can find an incomplete example for the LDAP synchronization configuration in the `ldap-sync.yaml` file. For the cron job for the group synchronization, you can apply the ready-to-use `rhds-groups-cronjob.yaml` file.
Assign the `cluster-admin` role to the `administrators` group, and verify its permissions.
4. Configure the Red Hat SSO OIDC IdP.
You need the OIDC client secret to configure OpenShift. To get this information, access the Red Hat SSO machine by using the `ssh rhssouser@rhssouser.ocp4.example.com` command from

Chapter 1 | Authentication and Identity Management

the workstation machine. On this `sso.ocp4.example.com` machine, use the `/opt/rh-sso-7.6/bin/kcadm.sh` command with the following parameters to query Red Hat SSO:

Parameter	Value
Authentication URL	<code>https://sso.ocp4.example.com:8080/auth</code>
Authentication realm	<code>master</code>
Administrator username	<code>admin</code>
Administrator password	<code>redhatocp</code>
Dedicated realm for OpenShift	<code>external_providers</code>
Client ID	<code>ocp_rhsso</code>

In OpenShift, use the `rhsso-oidc-client-secret` name for the Red Hat SSO secret. Use `RHSSO_OIDC` as the name for the Red Hat SSO IdP.

You can find an incomplete example for the Red Hat SSO IdP configuration in the `~/D0380/labs/auth-review/sso/sso_config.yaml` file. You must include the IdP configuration for the LDAP server that you set up in a previous step.

- Configure an external monitoring app to run outside the cluster.

Create the `audit-bot` SA, generate an SA token, and store the SA token in the `AUDITBOT_TOKEN` variable. Use the `auth-review` project. Create a `kubeconfig` file with the `audit-bot` user credentials. Store the `kubeconfig` file in the `~/D0380/labs/auth-review/log_script/audit-bot.config` file.

Run the `cluster-health.sh` script to verify the status of the pods in the OpenShift cluster. You can find the script in the `~/D0380/labs/auth-review/log_script` directory. The script uses the credentials that are stored in the `~/D0380/labs/auth-review/log_script/audit-bot.config` file.

Evaluation

As the student user on the workstation machine, use the `lab` command to grade your work. Correct any reported failures and rerun the command until successful.

```
[student@workstation ~]$ lab grade auth-review
```

Finish

As the student user on the workstation machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish auth-review
```

► Solution

Authentication and Identity Management

Define the concepts and custom resources of the OpenShift OAuth server, and explain how these resources augment Kubernetes authentication.

Configure an LDAP identity provider and automate group synchronization between OpenShift OAuth and an LDAP server.

Configure an OIDC identity provider and automate group synchronization between OpenShift OAuth and an OIDC server.

Generate a token and a client certificate and add them to a kubeconfig file.

Outcomes

- Use a client certificate for a system administrator account to recover access to the OpenShift cluster.
- Configure RHDS as an LDAP identity provider (IdP) for OpenShift.
- Configure Red Hat SSO as an OIDC IdP for OpenShift.
- Synchronize users and groups from the LDAP and Red Hat SSO IdPs to OpenShift.
- Generate a service account (SA) token for authenticating a script that runs from outside the OpenShift cluster.

Before You Begin

As the student user on the workstation machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start auth-review
```

Instructions

A cluster administrator in your company erroneously removed all the configured IdPs from the OAuth server. You must use a backdoor administrator certificate to re-create the IdPs. For this purpose, create a `kubeconfig` file that contains the client certificate, the certificate key, and the OpenShift CA server certificate. The certificate is already approved in the OpenShift CA server with administrator privileges.

After you create the administrator `kubeconfig` file, restore the LDAP IdP that contains one of the company managers, Payden Tomcheck, with the `paydентomcheck` username and the `redhat123` password. You must configure the LDAP automated group synchronization, because the user is part of the `administrators` group. This group must have administrative privileges to the cluster.

Then, configure an OIDC IdP. This IdP provides the consultant user, Lauren Chan, with the `laurenchan` username and the `redhat_sso` password. You must configure the OIDC IdP to

Chapter 1 | Authentication and Identity Management

automatically synchronize the groups, because the user is part of the `consultants` group. This group must have read privileges to the `auth-review` project.

Finally, configure an external monitoring app, which runs outside the cluster and uses an SA account token that must have read access to the cluster. You must create the SA account in the `auth-review` project with the `audit-bot` name.

1. Create a `kubeconfig` file that contains the client certificate, the certificate key, and the OpenShift CA server certificate.

You must create the `kubeconfig` file in the `~/D0380/labs/auth-review/certificate` directory with the `admin.config` name.

You can find the `admin-access.crt` client certificate, the `tls.key` certificate key, and the `ocp-apiserver-cert.crt` CA server certificate in the same directory. Use `api-ocp4-example-com:6443` as the cluster name and `https://api.ocp4.example.com:6443` as the cluster IP in the `kubeconfig` file.

- 1.1. Change to the `~/D0380/labs/auth-review/certificate` directory.

```
[student@workstation ~]$ cd ~/D0380/labs/auth-review/certificate  
[student@workstation certificate]$
```

- 1.2. Add the `admin` user credentials to the `kubeconfig` file.

```
[student@workstation certificate]$ oc config set-credentials admin \  
--client-certificate admin-access.crt --client-key tls.key \  
--embed-certs --kubeconfig admin.config  
User "admin" set.
```

- 1.3. Set the cluster options in the `kubeconfig` file.

```
[student@workstation certificate]$ oc config set-cluster \  
api-ocp4-example-com:6443 --certificate-authority ocp-apiserver-cert.crt \  
--embed-certs --server https://api.ocp4.example.com:6443 \  
--kubeconfig admin.config  
Cluster "api-ocp4-example-com:6443" set.
```

- 1.4. Set the context for the `admin` user.

```
[student@workstation certificate]$ oc config set-context admin \  
--cluster api-ocp4-example-com:6443 --namespace auth-review --user admin \  
--kubeconfig admin.config  
Context "admin" created.
```

- 1.5. Use the context for the `admin` user.

```
[student@workstation certificate]$ oc config use-context admin \  
--kubeconfig admin.config  
Switched to context "admin".
```

- 1.6. Test the administrator certificate.

```
[student@workstation certificate]$ oc whoami --kubeconfig admin.config
system:admin
```

1.7. Change to the student HOME directory.

```
[student@workstation certificate]$ cd
```

2. Use the administrator kubeconfig file to configure the LDAP IdP. Use the LDAP IdP information from the following table:

Parameter	Value
DN (-D)	cn=Directory Manager
URI (-H)	ldaps://rhds.ocp4.example.com
Password (-w)	redhatocp

Use the ~/D0380/labs/auth-review/rhds directory to create any necessary files.

You can find the RHDS certificate in that directory with the rhds_ca.crt name. Use the rhds-ldap-secret and rhds-ca-config-map names for the RHDS secret and configuration map respectively. Use Red Hat Directory Server as the name for the RHDS LDAP IdP.

You can find an incomplete example for the OAuth CR in the rhds-ldap-idp.yaml file.

2.1. Change to the ~/D0380/labs/auth-review/rhds directory.

```
[student@workstation ~]$ cd ~/D0380/labs/auth-review/rhds
[student@workstation rhds]$
```

2.2. From the CLI, test the connection with the ldapsearch command by using the information in the table.

```
[student@workstation rhds]$ ldapsearch -D "cn=Directory Manager" \
-w redhatocp -H ldaps://rhds.ocp4.example.com
...output omitted...
# paydenthomcheck, people, example.com
dn: uid=paydenthomcheck,ou=people,dc=example,dc=com
objectClass: top
objectClass: account
objectClass: posixAccount
objectClass: shadowAccount
cn: Payden Tomcheck
uid: paydenthomcheck
uidNumber: 10001
gidNumber: 101
homeDirectory: /home/paydenthomcheck
loginShell: /bin/bash
gecos: paydenthomcheck
shadowLastChange: 0
shadowMax: 0
```

Chapter 1 | Authentication and Identity Management

```
shadowWarning: 0
userPassword::  
e1NTSEF9Wm1sMwd1WjjLajlRM1dKZG1lVnV6aTNEaCs5NzFPFg=  
  
# administrators, people, example.com
dn: cn=administrators,ou=people,dc=example,dc=com
objectClass: top
objectClass: groupOfUniqueNames
cn: administrators
uniqueMember: uid=paydentomcheck,ou=people,dc=example,dc=com
...output omitted...
```

2.3. Create the RHDS secret.

```
[student@workstation rhds]$ oc create secret generic rhds-ldap-secret \
--from-literal bindPassword=redhatocp -n openshift-config \
--kubeconfig ~/DO380/labs/auth-review/certificate/admin.config
secret/rhds-ldap-secret created
```

2.4. Create the configuration map that contains the certificate.

```
[student@workstation rhds]$ oc create configmap rhds-ca-config-map \
--from-file ca.crt=rhds_ca.crt -n openshift-config \
--kubeconfig ~/DO380/labs/auth-review/certificate/admin.config
configmap/rhds-ca-config-map created
```

2.5. Edit the RHDS LDAP IdP in the OAuth server. This configuration replaces the OAuth configuration by overriding any IdP configuration that was previously in the cluster. You can find an incomplete example for the OAuth CR in the `rhds-ldap-idp.yaml` file.

```
apiVersion: config.openshift.io/v1
kind: OAuth
metadata:
  name: cluster
spec:
  identityProviders:
    - ldap:
        attributes:
          id:
            - dn
          email:
            - mail
        name:
          - cn
        preferredUsername:
          - uid
      bindDN: 'cn=Directory Manager'
      bindPassword:
        name: rhds-ldap-secret
      ca:
        name: rhds-ca-config-map
      insecure: false
      url: >-
```

Chapter1 | Authentication and Identity Management

```
ldaps://rhds.ocp4.example.com/dc=example,dc=com?uid
mappingMethod: claim
name: Red Hat Directory Server
type: LDAP
```

- 2.6. Use the `oc apply` command to configure the RHDS LDAP IdP in the OAuth server.

```
[student@workstation rhds]$ oc apply -f rhds-ldap-idp.yaml \
--kubeconfig ~/D0380/labs/auth-review/certificate/admin.config
oauth.config.openshift.io/cluster configured
```

- 2.7. Verify the status for the OAuth pods and wait for the OAuth pods to be redeployed. Press `Ctrl+C` when all the new pods are running.

```
[student@workstation rhds]$ watch oc get pods -n openshift-authentication \
--kubeconfig ~/D0380/labs/auth-review/certificate/admin.config
Every 2.0s: oc get pods -n openshift-authentication    workstation: Tue Aug 22 ...
NAME                      READY   STATUS    RESTARTS   AGE
oauth-openshift-78cbdc45f7-2z887  1/1     Running   0          35s
oauth-openshift-78cbdc45f7-4jsdc  1/1     Running   0          62s
oauth-openshift-78cbdc45f7-rlv2q  1/1     Running   0          89s
^C
```

- 2.8. Verify that you can log in to the cluster as the `paydентомcheck` user with `redhat123` as the password.

```
[student@workstation group-sync]$ oc login -u paydентомcheck -p redhat123 \
https://api.ocp4.example.com:6443
Login successful.
...output omitted...
```

- 2.9. Change to the student HOME directory.

```
[student@workstation rhds]$ cd
```

3. Configure the LDAP automated group synchronization in a project called `auth-ldapsync`.

Create the `auth-ldapsync` project, and create the `ldap-group-syncer` SA. Create the `ldap-group-syncer` cluster role with `get`, `list`, `create`, and `update` permissions to the `groups` resource, and assign the cluster role to the SA.

Use the `ldap-secret` and `ldap-config` names for the RHDS secret and configuration map, respectively. Mount the secret and the configuration map in the `/etc/secrets/` `bindPassword` and `/etc/config/ca.crt` files, respectively.

Use the `~/D0380/labs/auth-review/group-sync` directory to create any necessary files. You can find the RHDS certificate in that directory with the `rhds_ca.crt` name. You can find an incomplete example for the LDAP synchronization configuration in the `ldap-sync.yaml` file. For the cron job for the group synchronization, you can apply the ready-to-use `rhds-groups-cronjob.yaml` file.

Assign the `cluster-admin` role to the `administrators` group, and verify its permissions.

- 3.1. Change to the `~/D0380/labs/auth-review/group-sync` directory.

Chapter 1 | Authentication and Identity Management

```
[student@workstation ~]$ cd ~/D0380/labs/auth-review/group-sync  
[student@workstation group-sync]$
```

- 3.2. Create a project called auth-ldapsync.

```
[student@workstation group-sync]$ oc new-project auth-ldapsync \  
--kubeconfig ~/D0380/labs/auth-review/certificate/admin.config  
Now using project "auth-ldapsync" on server "https://api.ocp4.example.com:6443".  
...output omitted...
```

- 3.3. Create a service account called ldap-group-syncer.

```
[student@workstation group-sync]$ oc create sa ldap-group-syncer \  
--kubeconfig ~/D0380/labs/auth-review/certificate/admin.config  
serviceaccount/ldap-group-syncer created
```

- 3.4. Create the ldap-group-syncer cluster role.

```
[student@workstation group-sync]$ oc create clusterrole ldap-group-syncer \  
--verb get,list,create,update --resource groups \  
--kubeconfig ~/D0380/labs/auth-review/certificate/admin.config  
clusterrole.rbac.authorization.k8s.io/ldap-group-syncer created
```

- 3.5. Create the cluster role binding for the ldap-group-syncer SA and cluster role.

```
[student@workstation ~]$ oc adm policy add-cluster-role-to-user \  
ldap-group-syncer -z ldap-group-syncer \  
--kubeconfig ~/D0380/labs/auth-review/certificate/admin.config  
clusterrole.rbac.authorization.k8s.io/ldap-group-syncer added: "ldap-group-syncer"
```

- 3.6. Create a secret that contains the LDAP bind password.

```
[student@workstation group-sync]$ oc create secret generic ldap-secret \  
--from-literal bindPassword='redhatocp' \  
--kubeconfig ~/D0380/labs/auth-review/certificate/admin.config  
secret/ldap-secret created
```

- 3.7. Create the LDAP synchronization configuration file. You can find an incomplete example for the LDAP synchronization configuration in the ldap-sync.yaml file.

```
kind: LDAPSNCConfig  
apiVersion: v1  
url: ldaps://rhds.ocp4.example.com:636  
bindDN: 'cn=Directory Manager'  
bindPassword:  
  file: /etc/secrets/bindPassword  
ca: /etc/config/ca.crt  
augmentedActiveDirectory:  
  groupsQuery:  
    baseDN: "ou=people,dc=example,dc=com"
```

```

scope: sub
derefAliases: never
pageSize: 0
groupUIDAttribute: dn
groupNameAttributes: [ cn ]
usersQuery:
  baseDN: "ou=people,dc=example,dc=com"
  scope: sub
  derefAliases: never
  filter: (objectclass=account)
  pageSize: 0
userNameAttributes: [ uid ]
groupMembershipAttributes: [ memberOf ]

```

- 3.8. Create a configuration map that contains the LDAPSsyncConfig file and the trusted certificate.

```
[student@workstation group-sync]$ oc create configmap ldap-config \
--from-file ldap-sync.yaml=ldap-sync.yaml,ca.crt=rhds_ca.crt \
--kubeconfig ~/DO380/labs/auth-review/certificate/admin.config
configmap/ldap-config created
```

- 3.9. Review the cron job resource manifest for the group synchronization in the rhds-groups-cronjob.yaml file.

```

apiVersion: batch/v1
kind: CronJob
metadata:
  name: group-sync
  namespace: auth-ldapsync
spec:
  schedule: "*/1 * * * *"
  jobTemplate:
    spec:
      template:
        spec:
          restartPolicy: Never
          containers:
            - name: ldap-group-sync
              image: "registry.ocp4.example.com:8443/openshift4/ose-cli:v4.12"
              command:
                - "/bin/sh"
                - "-c"
                - "oc adm groups sync --sync-config=/etc/config/ldap-sync.yaml --confirm" ❶
          volumeMounts:
            - mountPath: "/etc/config"
              name: "ldap-sync-volume"
            - mountPath: "/etc/secrets"
              name: "ldap-bind-password"
          volumes:
            - name: "ldap-sync-volume"
          configMap:
            name: "ldap-config"

```

Chapter 1 | Authentication and Identity Management

```
- name: "ldap-bind-password"
  secret:
    secretName: "ldap-secret"
  serviceAccountName: ldap-group-syncer
  serviceAccount: ldap-group-syncer
```

- ➊ This command should be written in a single line.

3.10. Create the cron job for the group synchronization.

```
[student@workstation group-sync]$ oc create -f rhds-groups-cronjob.yaml \
--kubeconfig ~/D0380/labs/auth-review/certificate/admin.config
...output omitted...
cronjob.batch/group-sync created
```

3.11. Wait for one minute for the cron job to start, and verify that OpenShift synchronizes the administrators group from RHDS.

```
[student@workstation group-sync]$ oc get groups \
--kubeconfig ~/D0380/labs/auth-review/certificate/admin.config
NAME          USERS
Default SMB Group
administrators  paydентомcheck
admins        admin
...output omitted...
```

3.12. Apply the cluster-admin role to the administrators group.

```
[student@workstation group-sync]$ oc adm policy add-cluster-role-to-group \
cluster-admin administrators \
--kubeconfig ~/D0380/labs/auth-review/certificate/admin.config
clusterrole.rbac.authorization.k8s.io/cluster-admin added: "administrators"
```

3.13. Log in to the cluster as the paydентомcheck user.

```
[student@workstation group-sync]$ oc login -u paydентомcheck -p redhat123
Login successful.
...output omitted...
```

3.14. Verify that the user has cluster administrator privileges.

```
[student@workstation group-sync]$ oc auth can-i create users -A
yes
```

3.15. Change to the student HOME directory.

```
[student@workstation group-sync]$ cd
```

4. Configure the Red Hat SSO OIDC IdP.

You need the OIDC client secret to configure OpenShift. To get this information, access the Red Hat SSO machine by using the ssh `rhssousername@sso.ocp4.example.com` command from

the workstation machine. On this `sso.ocp4.example.com` machine, use the `/opt/rh-sso-7.6/bin/kcadm.sh` command with the following parameters to query Red Hat SSO:

Parameter	Value
Authentication URL	<code>https://sso.ocp4.example.com:8080/auth</code>
Authentication realm	<code>master</code>
Administrator username	<code>admin</code>
Administrator password	<code>redhatocp</code>
Dedicated realm for OpenShift	<code>external_providers</code>
Client ID	<code>ocp_rhsso</code>

In OpenShift, use the `rhsso-oidc-client-secret` name for the Red Hat SSO secret. Use `RHSSO_OIDC` as the name for the Red Hat SSO IdP.

You can find an incomplete example for the Red Hat SSO IdP configuration in the `~/D0380/labs/auth-review/sso/sso_config.yaml` file. You must include the IdP configuration for the LDAP server that you set up in a previous step.

4.1. Connect to the Red Hat SSO machine as the `rhsso` user.

```
[student@workstation ~]$ ssh rhsso@sso.ocp4.example.com
[rhsso@sso ~]$
```

4.2. Use the `kcadm` tool to connect to Red Hat SSO.

```
[rhsso@sso ~]$ /opt/rh-sso-7.6/bin/kcadm.sh config credentials \
--server https://sso.ocp4.example.com:8080/auth \
--user admin --password redhatocp --realm master
Logging into https://sso.ocp4.example.com:8080/auth as user admin of realm master
```

4.3. List the information for the `ocp_rhsso` client from Red Hat SSO.

```
[rhsso@sso ~]$ /opt/rh-sso-7.6/bin/kcadm.sh get clients \
-r external_providers -q clientId=ocp_rhsso
[ {
"id" : "f57e9ddc-8c60-4b40-8048-ec0120595be2",
"clientId" : "ocp_rhsso",
"surrogateAuthRequired" : false,
"enabled" : true,
"alwaysDisplayInConsole" : false,
"clientAuthenticatorType" : "client-secret",
"redirectUris" : [ "https://oauth-openshift.apps.ocp4.example.com/*" ],
"webOrigins" : [ "https://oauth-openshift.apps.ocp4.example.com" ],
...output omitted...
```

4.4. Generate a JSON file, which contains the client ID, the authentication server URL, and the client secret. Use the `ocp_rhsso` ID from the previous step.

```
[rhsso@sso ~]$ /opt/rh-sso-7.6/bin/kcadm.sh get \
  clients/f57e9ddc-8c60-4b40-8048-ec0120595be2/installation\
/providers/keycloak-oidc-keycloak-json \ ①
  -r external_providers > rhsso.json
```

- ① The /providers/... text must come after the .../installation text in a single line without spaces.
- 4.5. View the content of the JSON file, which contains the Red Hat SSO client information. Note the secret, which you use in a later step. The client secret would differ on your system. Use the ocp_rhsso Red Hat SSO client ID for the clientId parameter in the IdP configuration file on OpenShift. The issuer parameter in the IdP configuration on OpenShift concatenates the value from the auth-server-url parameter, the /realms/ string, and the external_providers Red Hat SSO realm name.

```
[rhsso@sso ~]$ cat rhsso.json
{
  "realm" : "external_providers",
  "auth-server-url" : "https://sso.ocp4.example.com:8080/auth/",
  "ssl-required" : "external",
  "resource" : "ocp_rhsso",
  "credentials" : {
    "secret" : "X4ZTPfDr0b8loq0FArfidhaHq85bHyiy"
  },
  "confidential-port" : 0
}
```

- 4.6. Return to the workstation machine.

```
[rhsso@sso ~]$ exit
logout
Connection to sso.ocp4.example.com closed.
[student@workstation ~]$
```

- 4.7. Connect to the OpenShift cluster as the paydentitycheck user with redhat123 as the password.

```
[student@workstation ~]$ oc login -u paydentitycheck -p redhat123 \
  https://api.ocp4.example.com:6443
Login successful.
...output omitted...
```

- 4.8. Create the rhsso-oidc-client-secret OpenShift secret for the Red Hat SSO client secret by using the client secret from a previous step.

```
[student@workstation ~]$ oc create secret generic rhsso-oidc-client-secret \
  --from-literal clientSecret=X4ZTPfDr0b8loq0FArfidhaHq85bHyiy \
  -n openshift-config
secret/rhsso-oidc-client-secret created
```

- 4.9. Edit the Red Hat SSO IdP in the OAuth server. You can find an example for the CR in the `~/D0380/labs/auth-review/sso/sso_config.yaml` file. You must include the IdP configuration for the LDAP server that you set up in a previous step.

```
[student@workstation ~]$ oc edit oauth cluster
apiVersion: config.openshift.io/v1
kind: OAuth
metadata:
  name: cluster
spec:
  identityProviders:
    - ldap:
      ...output omitted...
    - openID:
      claims:
        email:
          - email
      name:
        - name
      preferredUsername:
        - preferred_username
      groups:
        - groups
      clientID: ocp_rhss0
      clientSecret:
        name: rhss0-oidc-client-secret
      extraScopes: []
      issuer: >-
        https://sso.ocp4.example.com:8080/auth/realms/external_providers
      mappingMethod: claim
      name: RHSSO_OIDC
      type: OpenID
```



Note

You can also use the solution in the `~/D0380/solutions/auth-review/sso/sso_config.yaml` file to configure the Red Hat SSO IdP in the OAuth server.

```
[student@workstation ~]$ oc apply -f \
~/D0380/solutions/auth-review/sso/sso_config.yaml
oauth.config.openshift.io/cluster configured
```

- 4.10. Verify the status for the OAuth pods and wait for the OAuth pods to be redeployed. Press **Ctrl+C** when all the new pods are running.

Chapter 1 | Authentication and Identity Management

```
[student@workstation ~]$ watch oc get pods -n openshift-authentication
Every 2.0s: oc get pods -n openshift-authentication workstation: Thu Aug 22 ...
NAME          READY   STATUS    RESTARTS   AGE
oauth-openshift-7d5f479864-6jktw  1/1     Running   0          3m25s
oauth-openshift-7d5f479864-86cfg  1/1     Running   0          3m42s
oauth-openshift-7d5f479864-dc6jw  1/1     Running   0          3m48s
^C
```

- 4.11. Verify that you can log in to the cluster as the `laurenchan` user with `redhat_sso` as the password.

```
[student@workstation ~]$ oc login -u laurenchan -p redhat_sso
Login successful.
...output omitted...
```

- 4.12. Log in to the cluster as the `paydентомcheck` user with `redhat123` as the password.

```
[student@workstation ~]$ oc login -u paydентомcheck -p redhat123
Login successful.
...output omitted...
```

- 4.13. Verify that OpenShift synchronizes the user from Red Hat SSO the first time that they log in.

```
[student@workstation ~]$ oc get users
NAME        UID           FULL NAME      IDENTITIES
laurenchan  dbe329a2-...  Lauren Chan   RHSSO_OIDC:a175e1b7-...
paydентомcheck 2554ba58-...  Payden Tomcheck  Red Hat Directory Server:dwlk...
```

- 4.14. Verify that the `laurenchan` user is a member of the `consultants` group.

```
[student@workstation ~]$ oc get groups
NAME        USERS
Default SMB Group
administrators  paydентомcheck
admins       admin
consultants   laurenchan
...output omitted...
```

- 4.15. Assign the `view` cluster role in the `auth-review` project to the `consultants` group.

```
[student@workstation ~]$ oc adm policy add-role-to-group view consultants \
-n auth-review
clusterrole.rbac.authorization.k8s.io/view added: "consultants"
```

5. Configure an external monitoring app to run outside the cluster.

Create the `audit-bot` SA, generate an SA token, and store the SA token in the `AUDITBOT_TOKEN` variable. Use the `auth-review` project. Create a `kubeconfig` file with

the audit-bot user credentials. Store the kubeconfig file in the ~/D0380/labs/auth-review/log_script/audit-bot.config file.

Run the cluster-health.sh script to verify the status of the pods in the OpenShift cluster. You can find the script in the ~/D0380/labs/auth-review/log_script directory. The script uses the credentials that are stored in the ~/D0380/labs/auth-review/log_script/audit-bot.config file.

5.1. Change to the ~/D0380/labs/auth-review/log_script directory.

```
[student@workstation ~]$ cd ~/D0380/labs/auth-review/log_script  
[student@workstation log_script]$
```

5.2. Change to the auth-review project.

```
[student@workstation log_script]$ oc project auth-review  
Now using project "auth-review" on server "https://api.ocp4.example.com:6443".
```

5.3. Create the audit-bot SA in the auth-review project.

```
[student@workstation log_script]$ oc create sa audit-bot  
serviceaccount/audit-bot created
```

5.4. Assign the cluster-reader cluster role to the audit-bot SA so that the SA can retrieve information from most of the objects in the cluster.

```
[student@workstation log_script]$ oc adm policy add-cluster-role-to-user \  
cluster-reader system:serviceaccount:auth-review:audit-bot  
clusterrole.rbac.authorization.k8s.io/cluster-reader added:  
"system:serviceaccount:auth-review:audit-bot"
```

5.5. Generate the audit-bot SA account token and store it in the AUDITBOT_TOKEN variable.

```
[student@workstation log_script]$ AUDITBOT_TOKEN=$(oc create token \  
-n auth-review audit-bot)
```

5.6. Add the audit-bot user credentials to the kubeconfig file.

```
[student@workstation log_script]$ oc config set-credentials audit-bot \  
--token $AUDITBOT_TOKEN --kubeconfig audit-bot.config  
User "audit-bot" set.
```

5.7. Set the cluster options in the kubeconfig file.

```
[student@workstation log_script]$ oc config set-cluster \  
api-ocp4-example-com:6443 --server https://api.ocp4.example.com:6443 \  
--kubeconfig audit-bot.config  
Cluster "api-ocp4-example-com:6443" set.
```

5.8. Set the context for the audit-bot user.

```
[student@workstation log_script]$ oc config set-context audit-bot \
--cluster api-ocp4-example-com:6443 --namespace auth-review --user audit-bot \
--kubeconfig audit-bot.config
Context "audit-bot" created.
```

5.9. Use the context for the audit-bot user.

```
[student@workstation log_script]$ oc config use-context audit-bot \
--kubeconfig audit-bot.config
Switched to context "audit-bot".
```

5.10. Verify the identity of the audit-bot user by using the kubeconfig file.

```
[student@workstation log_script]$ oc whoami --kubeconfig audit-bot.config
system:serviceaccount:auth-review:audit-bot
```

5.11. Run the `cluster-health.sh` script to verify the status of the pods in the OpenShift cluster. The script uses the credentials that are stored in the `~/D0380/labs/auth-review/log_script/audit-bot.config` file.

```
[student@workstation log_script]$ sh cluster-health.sh
Connected to the cluster as the 'system:serviceaccount:auth-review:audit-bot' user
✓ OpenShift is reachable and up, at version: '4.14.12'
✓ All pods are either running or succeeded.
```

5.12. Change to the student HOME directory.

```
[student@workstation log_script]$ cd
```

Evaluation

As the `student` user on the `workstation` machine, use the `lab` command to grade your work. Correct any reported failures and rerun the command until successful.

```
[student@workstation ~]$ lab grade auth-review
```

Finish

As the `student` user on the `workstation` machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish auth-review
```

Summary

- OpenShift expands the Kubernetes authentication mechanisms by providing a built-in OAuth server to configure external identity providers (IdPs).
- You can authenticate to OpenShift by using client certificates and OAuth access tokens.
- OpenShift includes users, groups, and identities.
- The mapping methods in OpenShift control how OpenShift establishes mappings between the identities of IdPs and the user resources.
- In the OpenShift built-in OAuth server, you can configure different IdPs, such as LDAP and OIDC.
- When you configure an LDAP IdP, OpenShift automatically synchronizes the users, but not the groups. Synchronizing LDAP groups requires manual action by an administrator and custom automation.
- Use OIDC to automatically synchronize users and groups from an OIDC IdP, such as Google, Microsoft Identity Platform, or Keycloak.
- OIDC claims enable OpenShift to read the user information from the OIDC token and to populate the user, identity, and group resources.
- OpenShift provides by default the `kubeadmin` user, which is a special system-level user account with administrative access.
- Use service accounts and client certificates to authenticate external resources to the Kubernetes API.

Chapter 2

Backup, Restore, and Migration of Applications with OADP

Goal

Back up and restore application settings and data with OpenShift API for Data Protection (OADP).

Sections

- Export and Import Application Data and Settings (and Guided Exercise)
- OADP Operator Deployment and Features (and Guided Exercise)
- Backup and Restore with OADP (and Guided Exercise)

Lab

- Backup, Restore, and Migration of Applications with OADP

Export and Import Application Data and Settings

Objectives

- Export and import application data and settings between projects.

Export and Import a Kubernetes Application

The ability to export and import a Kubernetes application is useful in many scenarios. The following list describes some of these scenarios:

- Partial or full restoration to a previous working state
- Disaster recovery and business continuity
- Migration from one cluster to another
- Duplication on multiple environments

Backup and restore procedures are critical for an organization to recover from data loss or corruption. The most common data loss scenarios are hardware failures, cyberattacks, software bugs, or human errors.

A Kubernetes application backup includes all the needed resources to restore that application to a previous working state. This backup must include the following artifacts:

- Kubernetes resources that define the application and its settings.
- Container images in the internal registry that containers of this application use.
- Data that is stored in persistent volumes or object storage for a stateful application.



Note

Application data or configuration that is hosted on external services such as Database-as-a-Service or object storage might be required but is outside the scope of this chapter.

Red Hat OpenShift and Red Hat partners provide data protection solutions for a faster recovery plan. The following list includes examples of data protection solutions:

- OpenShift API for Data Protection
- Veeam Kasten K10
- Storware Backup and Recovery
- IBM Spectrum Protect Plus
- Pure Storage Portworx Backup

Backing Up Application Resources

A Kubernetes application has many resources. The export and import include the following steps:

- List all required resources for the application.
- Export the listed resources.
- Clean the exported resource files.
- Deploy the cleaned resource files.

The first step to export an application is to list all the application resources.



Note

The `oc get all` command lists only a subset of resources in a project and does not show resources such as secrets, configuration maps, and so on.

You do not need to list all resources in a project. For example, a MySQL application might require the following resources:

- Deployment
- Service
- Secret

You can list these resources by using the `oc get` command.

```
[user@host ~]$ oc get deployment,svc,secret
NAME                   READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/mysql   1/1     1           1           112s

NAME          TYPE      CLUSTER-IP       EXTERNAL-IP   PORT(S)    AGE
service/mysql  ClusterIP  172.30.210.241  <none>        9001/TCP   115s

NAME          TYPE          DATA   AGE
secret/dockercfg-skqdg  kubernetes.io/dockercfg   1       2m16s
secret/token-cpmxf      kubernetes.io/service-account-token  4       2m16s
...output omitted...
secret/mysql-credentials Opaque          3       60s
```

A MySQL application might use persistent volume claim for data. The required resources for application export depend on the application.

Export Application Resources

You can export an application object to a YAML file by using the `oc get` command.

```
[user@host ~]$ oc -n prod get deployment/mysql -o yaml > backup_deployment.yaml
```

The `backup_deployment.yaml` file has all deployment details such as specifications, metadata, and status.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  annotations:
    deployment.kubernetes.io/revision: "2"
  labels:
    app: mysql
    name: mysql
    namespace: prod
  ...output omitted...
spec:
  progressDeadlineSeconds: 600
```

```

replicas: 1
selector:
  matchLabels:
    app: mysql
    ...output omitted...
status:
  availableReplicas: 1
...output omitted...

```

You can export all other application resources by using the same `oc get` command.

Import Application Resources

You can create the application resources in a new project by using the backup YAML files. You can remove the `metadata.namespace` field and use the `oc create -f` command to create each resource in the new project.

Apart from the `metadata.namespace` field, all backup resource files have runtime and status information such as `metadata.annotations`, `metadata.creationTimestamp`, `metadata.resourceVersion`, `metadata.generation`, and `status`. You can remove these fields and keep cleaned backup resource files to create resources in a new project.

The following example demonstrates importing some required resources for an application:

```

[user@host ~]$ oc create -f clean_backup_deployment.yaml -n prod-backup
deployment.apps/mysql created

[user@host ~]$ oc create -f service-mysql.yaml -n prod-backup
The Service "mysql" is invalid: spec.clusterIPs: Invalid value:
[]string{"172.30.210.241"}: failed to allocate IP 172.30.210.241: provided IP is
already allocated

```

The `mysql` deployment was created successfully in the `prod-backup` project by using the backup YAML file.

However, the `metadata` and `status` removal does not work for all resources. Some resources require additional modifications.

The service resource creation failed because the IP address is already allocated to the `mysql` service in the `prod` project. The service backup YAML file requires more modification for a successful creation.

The modification depends on the resource type. Service resource creation requires the removal of not only the `metadata` and `status` fields but also the `spec.clusterIP` field from the YAML file.

You can use a text editor or other tools to remove these fields. For example, you can use the `yq` tool to process the YAML file and remove a specific field.

```
[user@host ~]$ cat service-mysql.yaml \
| yq d - metadata.namespace \
| yq d - spec.clusterIP* > clean-service-mysql.yaml

[user@host ~]$ oc create -f clean-service-mysql.yaml -n prod-backup
service/mysql created
```



Note
The `yq` tool is a command-line YAML processor that is similar to the `jq` tool for JSON files. The `d` option removes the specified field from the output.

See the references section for more information about the `yq` tool.

Similar to the service resource, the route resource creation requires the removal of the `metadata.namespace` and `spec.host` fields from the YAML file.

```
[user@host ~]$ cat route-frontend.yaml \
| yq d - metadata.namespace \
| yq d - spec.host > clean-route-frontend.yaml

[user@host ~]$ oc create -f clean-route-frontend.yaml -n prod-backup
route.route.openshift.io/etherpad created
```

You do not need to create all the exported resources. The new project creates some resources, such as service accounts, secrets, and role bindings. An application also creates some resources. The `deployment` resource creates the `replicasets` resource, and the `deploymentconfig` resource creates the `replicationcontrollers` resource.

Backing Up Container Images

You can export container images by using container tools, such as `podman` or `skopeo`, to copy the images from one registry to another. For more details about Podman and Skopeo, refer to the *DO188: Red Hat OpenShift Development I: Introduction to Containers with Podman* training course.

Because the OpenShift internal registry is accessible only from within the cluster by default, you can use a Kubernetes job to export the images to a remote location. To access the registry from outside the cluster or to export images from your local machine, an OpenShift administrator must expose the registry externally.

Expose OpenShift Internal Registry

You can configure the OpenShift internal registry operator to expose the registry externally with the following command:

```
[user@host ~]$ oc patch \
configs.imageregistry.operator.openshift.io/cluster \
--patch '{"spec":{"defaultRoute":true}}' \
--type merge
```

**Note**

This action requires the `cluster-admin` role.

The modification to the image registry operator triggers a redeployment of the OpenShift API server. It can take up to 10 minutes for the cluster to stabilize.

The operator creates a `default-route` route to expose externally the registry that uses the following URL format:

```
default-route-openshift-image-registry.apps-domain
```

**Note**

You can use a custom hostname for the registry if needed. See the references section for more information about exposing the internal registry.

You can use the following command to retrieve the registry URL from the created route and save it in an environment variable for later use:

```
[user@host ~]$ REGISTRY=$(oc get \
  route default-route \
  -n openshift-image-registry \
  --template '{{.spec.host}}')
```

OpenShift users who do not have access to the `openshift-image-registry` namespace can retrieve the registry URL from any image stream:

```
[user@host ~]$ oc -n openshift get is/cli \
  -ojsonpath=".status.publicDockerImageRepository{`\n'}"

default-route-openshift-image-registry.apps.ocp4.example.com/openshift/cli
```

You can then log in to the internal registry by using your OpenShift username and authentication token with the following command:

```
[user@host ~]$ podman login \
  -u $(oc whoami) \
  -p $(oc whoami -t) \
  --tls-verify=false \
  $REGISTRY
```

**Note**

If the cluster's default ingress certificate is not trusted, then you must use the `--tls-verify=false` option to skip the certificate verification.

To configure a trusted certificate with the Ingress Operator, refer to the *Setting a custom default certificate* section in the *Ingress Operator in OpenShift Container Platform* chapter in the Red Hat OpenShift Container Platform 4.14 Networking documentation at https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/networking/index#nw-ingress-setting-a-custom-default-certificate_configuring-ingress

Alternatively, you can use the `oc registry login` command to log in to the internal registry without the need to specify your credential and registry URL. The `oc registry login` command automatically uses your authentication token and the internal registry URL from the OpenShift cluster.

The `oc` command stores the credentials in the `~/.docker/config.json` file in Base64 format. The Podman, Skopeo, and Docker clients can use the authentication details from that file to access an image registry.

```
[user@host ~]$ oc registry login
info: Using registry public hostname default-route-openshift-image-
registry.apps.ocp4.example.com
Warning: the default reading order of registry auth file will be changed from
"${HOME}/.docker/config.json" to podman registry config locations in the future
version of oc. "${HOME}/.docker/config.json" is deprecated, but can still be used
for storing credentials as a fallback. See https://github.com/containers/image/
blob/main/docs/containers-auth.json.5.md for the order of podman registry config
locations.
Saved credentials for default-route-openshift-image-registry.apps.ocp4.example.com
```

**Note**

The use of `~/.docker/config.json` file by the `oc` command is deprecated and will be changed to the `~/.config/containers/auth.json` file to store credentials in a future version. Although Podman and Skopeo can use the credentials from both files, the Docker client uses only the first file.

You can safely ignore the warning from the `oc registry login` command that mentions this deprecation, because the rest of the chapter uses only Podman and Skopeo.

Export Container Images

To export the image from outside the cluster, expose the internal registry and use the `skopeo` command to copy the image to a remote registry:

```
[user@host ~]$ skopeo copy \
  docker://${REGISTRY}/project_name/imagestream:tag \ ①
  docker://remote-registry.example.com/path/to/image:remotetag ②
```

Chapter 2 | Backup, Restore, and Migration of Applications with OADP

- ① Fully qualified source image in the OpenShift internal registry
- ② Destination registry URL with image and tag information

Skopeo can copy images to other locations such as a local directory or a .tar archive. See the references section for detailed use of the `skopeo copy` command.



Note

You can use the `skopeo sync` command to copy all the available tags in an image. See the references section for more information about the `skopeo sync` command.

If Skopeo is not available on your system, then you can use Podman or Docker to pull the image in the local container registry. You can then export the image as a .tar file with the `podman save` command or push the image to a remote registry.

```
[user@host ~]$ podman pull ${REGISTRY}/project_name/imagestream:tag  
[user@host ~]$ podman save ${REGISTRY}/project_name/imagestream:tag \  
| bzip2 > image_backup.tar.bz2
```

Alternatively, you can use the `oc image mirror` command to copy images to a local or remote location, similar to the `skopeo copy` command:

```
[user@host ~]$ oc image mirror ${REGISTRY}/project_name/imagestream:* \  
remote-registry.example.com/path/to/image ①
```

- ① You can use the wildcard character (*) to copy all the tags to the destination registry.



Note

The `oc` client that matches your OpenShift version is available in the `openshift/cli:latest` container image that is included in the OpenShift internal registry.

See the references section for additional examples of the `oc image mirror` command.

To export a container image from within the cluster, you can use any available container tools in a pod to copy the image to the location of your choice, such as a persistent volume on NFS storage, S3 storage, or a remote registry.

You can use the following YAML file as an example to create a Kubernetes job that exports the container image to a persistent volume by using the OpenShift client:

```
apiVersion: batch/v1  
kind: Job  
metadata:  
  name: backup-image  
  namespace: application  
  labels:  
    app: backup
```

```

spec:
  backoffLimit: 1
  template:
    metadata:
      labels:
        app: backup
    spec:
      containers:
        - name: backup
          image: image-registry.openshift-image-registry.svc:5000/openshift/
cli:latest
  env:
    - name: REGISTRY_AUTH_FILE
      value: /tmp/dockercfg.json
  command: ["/bin/bash", "-c"]
  args:
    - | ❶
      oc registry login
      oc image mirror \
        image-registry.openshift-image-registry.svc:5000/application/myapp:* \
        file://myapp --dir /backup
  volumeMounts: ❷
    - mountPath: /backup
      name: backup
  ...output omitted...
  volumes: ❸
    - name: backup
      persistentVolumeClaim:
        claimName: backup-volume

```

- ❶ Log in to the internal registry and export all tags of the `myapp` image stream to the `/backup` path.
- ❷ Volume mount definition for the backup location.
- ❸ Volume definition for the backup persistent volume claims (PVC).



Note

You need the `system:image-puller` role on the OpenShift project to pull images from any image streams in that project. Project users and administrators already have this permission, as well as the default service account.

Import Container Images

You can import an image to an image stream in any of your projects by using the same container tools as for the export. Similar to the export, you can use the following `skopeo` command to import an image from a remote repository:

```
[user@host ~]$ skopeo copy \
  docker://remote-registry.example.com/myimage:latest \
  docker://${REGISTRY}/project_name/mynewimagestream:latest
```

**Note**

If the image stream does not exist, then OpenShift creates it when the image is pushed.

You can also use the `oc image mirror` command to copy images in a similar way to the `skopeo copy` command:

```
[user@host ~]$ oc image mirror \
remote-registry.example.com/myimage:latest \
${REGISTRY}/project_name/mynewimagestream:latest
```

You can use Podman or Docker to import an image, although the process uses more steps, because you must import the image to the local container engine registry and rename the image to use the new registry name.

The following commands import an image in the Podman local registry from a compressed archive and push it to the OpenShift internal registry:

```
[user@host ~]$ podman load -i image_backup.tar.bz2
...output omitted...
Storing signatures
Loaded image(s): registry.apps.ocp4.example.com/application/myapp:1.2.3
```

If needed, rename the image with the `podman tag` command to match the new OpenShift project and image stream:

```
[user@host ~]$ podman tag \
registry.apps.ocp4.example.com/application/myapp:1.2.3 \
${REGISTRY}/newproject/myapp:1.2.3
```

Then, send the image to the OpenShift internal registry:

```
[user@host ~]$ podman push ${REGISTRY}/newproject/myapp:1.2.3
...output omitted...
Writing manifest to image destination
Storing signatures
```

**Note**

You need the `system:image-pusher` role on the OpenShift project to push images to any image streams in that project. Project users and administrators already have this permission, as well as the `builder` service account.

Backing Up Application Data

Several methods of data backup exist, depending on the application. Some applications provide dedicated tools or procedures to achieve the most reliable data protection and consistency. Different consistency levels can be achieved depending on the backup method.

Inconsistent backup

A backup is called inconsistent when the application alters data during the backup process. Traditional data copying when the application is running creates inconsistent backups.

Crash-consistent backup

A crash-consistent backup is created by suspending disk I/O during the backup, either by using snapshot technology or specialized tools, to ensure data consistency on disk. Application data in memory or pending I/O operations are not captured. The state of the application is kept as if the application was suddenly shut down due to power loss, or crashed.

Application-consistent backup

Application-consistent backup is the most reliable type of backup because it ensures that all in-memory data and pending I/O operations are written on disk before creating the backup.

Some applications provide a set of tools to flush memory data to disk and to pause file system operations on demand. These tools provide an application-consistent backup without any downtime by using snapshots, and is also known as *hot backup*.

For example, a MySQL database provides the `FLUSH TABLES WITH READ LOCK` statement to flush all operations to disk and to lock all tables before taking a snapshot. You can then unlock tables with the `UNLOCK TABLES` statement after the snapshot is created.

Database applications often come with specialized tools to create and restore backups without stopping the application or using volume snapshots. The following list includes examples of specialized tools:

- `mysqldump` for MySQL and MariaDB
- `pg_dump` for PostgreSQL



Note

Creating backups with specialized tools is out of the scope of this course.

A more universal way to create application-consistent backup is to stop the application, copy the data to another location, and then restart the application. This method is also called *cold backup*, because the application is down during the operation.

Depending on the amount of data, the application can be unavailable for a long time during the backup operation. By using volume snapshot, you can reduce this downtime to only a few minutes and use the cloned volume to back up while the application is back online.

Volume Snapshot

Volume snapshot capability is available only with Container Storage Interface (CSI) drivers. However, snapshot functions are not implemented for all CSI drivers. The following table includes examples of CSI drivers with the snapshot capability:

Storage provider	CSI driver
AWS Elastic Block Storage	<code>efs.csi.aws.com</code>
Azure Disk	<code>disk.csi.azure.com</code>
CephFS	<code>cephfs.csi.ceph.com</code>
Ceph RBD	<code>rbd.csi.ceph.com</code>

Storage provider	CSI driver
NetApp	csi.trident.netapp.io

Kubernetes provides similar API resources to `PersistentVolume` and `PersistentVolumeClaim` to create and manage volume snapshots.

VolumeSnapshotClass

Similar to the storage class for a persistent volume claim, a volume snapshot class describes the CSI driver and associated settings to create a volume snapshot.



Note

The `VolumeSnapshotClass` driver must match the `StorageClass` provisioner of the source PVC.

The following commands list all available storage and volume snapshot classes:

```
[user@host ~]$ oc get volumesnapshotclasses
NAME                                     DRIVER
ocs-storagecluster-cephfsplugin-snapclass  openshift-storage.cephfs.csi.ceph.com
ocs-storagecluster-rbdplugin-snapclass     openshift-storage.rbd.csi.ceph.com

[user@host ~]$ oc get storageclasses | egrep "^\w+|csi"
NAME                                     PROVISIONER
ocs-external-storagecluster-cephfs       openshift-storage.cephfs.csi.ceph.com
ocs-external-storagecluster-ceph-rbd      openshift-storage.rbd.csi.ceph.com
```

VolumeSnapshot

Similar to the `PersistentVolumeClaim` resource, a `VolumeSnapshot` resource requests the creation of a snapshot.

The following example is a YAML file for creating a volume snapshot:

```
apiVersion: snapshot.storage.k8s.io/v1
kind: VolumeSnapshot
metadata:
  name: my-snapshot ①
  namespace: application ②
spec:
  volumeSnapshotClassName: ocs-storagecluster-rbdplugin-snapclass ③
  source:
    persistentVolumeClaimName: application-data ④
```

- ① Name of the volume snapshot.
- ② Namespace of the volume snapshot. It must be the same as the source PVC.
- ③ Snapshot class name for the volume snapshot.
- ④ Name of the source PVC that is used for the snapshot.

The following command lists volume snapshots. A snapshot is successfully created when the `READYTOUSE` attribute is set to `true` and a `VolumeSnapshotContent` resource is created:

```
[user@host ~]$ oc get volumesnapshot
NAME          READYTOUSE   SOURCEPVC      ...  SNAPSHOTCONTENT
my-snapshot   true        application-data ...  snapcontent-798...cf6
```

**Note**

When creating an application-consistent backup, the application must be quiesced or scaled down before the snapshot creation. The application can safely be resumed or scaled up after the snapshot is created and ready to use.

VolumeSnapshotContent

Similar to the PersistentVolume resource, a VolumeSnapshotContent resource represents a snapshot that a VolumeSnapshot resource created.

After the snapshot is created, it can function as a source to create a PVC and for a pod to use it. The following example is a YAML file for creating a persistent volume from a snapshot:

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: my-snapshot-volume 1
  namespace: application 2
spec:
  storageClassName: ocs-external-storagecluster-ceph-rbd 3
  accessModes:
    - ReadWriteOnce
  dataSource:
    apiGroup: snapshot.storage.k8s.io
    kind: VolumeSnapshot
    name: my-snapshot 4
  resources:
    requests:
      storage: 1Gi 5
```

- 1** Name of the PVC.
- 2** Namespace of the PVC. It must be the same as the snapshot namespace.
- 3** Storage class name for the PVC.
- 4** Name of the snapshot.
- 5** Size of the new volume. Must be equal to or greater than the snapshot size.

Export Application Data

If your application provides specialized backup tools, you can use them to export the data to your chosen location. The following example is a cron job definition to back up a MariaDB database and to store the backup file on AWS S3 storage:

```
apiVersion: batch/v1
kind: CronJob
metadata:
```

```

name: backup-db
spec:
  schedule: "0 0 * * *"
  jobTemplate:
    spec:
      template:
        spec:
          initContainers:
            - envFrom:
              - secretRef:
                  name: mariadb
            image: quay.io/redhattraining/mariadb:10.5
            command: ["/bin/bash", "-c"]
            args: ①
              - >
                mariadb-dump -u "${MARIADB_USER}" -p"${MARIADB_PASSWORD}"
                -h mariadb "${MARIADB_DATABASE}"
                | bzip2 > /backup/backup-$(date '+%Y%m%d-%H%M').sql.bz2;
                ls -al /backup
          name: backup
          volumeMounts:
            - mountPath: /backup
              name: backup
        containers:
          - image: docker.io/amazon/aws-cli:latest
            command: ["/bin/bash", "-c"]
            args: ②
              - >
                aws s3 cp --no-progress /backup/backup* s3://backup/
          name: s3cli
          volumeMounts:
            - mountPath: /backup
              name: backup
            - mountPath: /root/.aws
              name: aws-creds
        volumes:
          - name: backup
            emptyDir: {}
          - name: aws-creds
            secret:
              secretName: s3config

```

- ① Use the `mariadb-dump` tool to export the database to a compressed file in an ephemeral volume.
- ② Use the AWS CLI tool to send the backup file to AWS S3-compatible storage.

If the application does not provide backup tools, then you can use volume snapshot to back up the application data. Depending on the storage provider that is available in the OpenShift cluster, you might rather export the snapshot content to an external storage location.

The following example is a job definition that uses an existing snapshot volume to archive the snapshot content, and exports the backup file to a remote S3 bucket:

```

apiVersion: batch/v1
kind: Job
metadata:
  name: backup
  namespace: application
  labels:
    app: backup
spec:
  backoffLimit: 1
  template:
    metadata:
      labels:
        app: backup
    spec:
      containers:
        - name: backup
          image: docker.io/d3fk/s3cmd:latest
          command: ["/bin/sh", "-c"]
          args:
            - | ①
              tar czf -C /snapshot /tmp/mybackup.tar.gz .
              s3cmd cp /tmp/mybackup.tar.gz s3://backup/
      volumeMounts: ②
        - mountPath: /snapshot
          name: snapshot
          ...output omitted...
      volumes: ③
        - name: snapshot
          persistentVolumeClaim:
            claimName: my-snapshot-volume

```

- ① Archive the snapshot content and copy the archive to a remote S3 bucket.
- ② Volume mount definition for the snapshot data in the container.
- ③ Volume definition for the snapshot PVC.



Important

If the volume storage class does not support volume snapshot, then you can mount the application volume instead to export the data. In this case, you must ensure that no other pods are using the volume during the backup, to avoid data inconsistencies.

Another option is to export the snapshot content locally from a pod by using the `oc cp` command. The following example is a pod definition, where the snapshot data is mounted, which you can use to export the data to your workstation:

```

apiVersion: v1
kind: Pod
metadata:
  name: export
spec:

```

```
containers:
- image: registry.access.redhat.com/ubi9:latest
  command: ["/bin/bash", "-c"]
  args:
    - sleep infinity ❶
    ...output omitted...
  volumeMounts:
    - mountPath: /snapshot
      name: snapshot
volumes: ❸
- name: snapshot
  persistentVolumeClaim:
    claimName: my-snapshot-volume
```

- ❶ The `sleep infinity` command ensures that the pod stays alive during the manual export.
- ❷ Volume mount definition for the snapshot data in the container.
- ❸ Volume definition for the snapshot PVC.

You can use the following `oc` command to copy the snapshot data from the `export` pod to your local machine:

```
[user@host ~]$ oc cp export:/snapshot /tmp/backup
```



Note

The `tar` binary must be installed in the remote container for the `oc cp` command to work.

After the snapshot content is exported to a remote location, you can safely remove the snapshot PVC and the volume snapshot, to free up space on the storage back end.



Important

If your application uses more than one persistent volume, then you must track the backup names and the volumes that they belong to. You need that information to restore each backup to the correct persistent volume.

Import Application Data

To import the data to another cluster or namespace, create a pod or a job where the application volume is mounted, and copy the exported data to the pod.

The following example is a YAML file for creating a job that fetches a remote backup from an S3 bucket and extracts the archive to the application volume:

```
apiVersion: batch/v1
kind: Job
metadata:
  name: restore
  namespace: application
  labels:
```

```
app: restore
spec:
  backoffLimit: 1
  template:
    metadata:
      labels:
        app: restore
    spec:
      containers:
        - name: restore
          image: docker.io/d3fk/s3cmd:latest
          command: ["/bin/bash", "-c"]
          args: ①
          - |
            s3cmd cp s3://backup/mybackup.tar.gz /tmp/
            tar xvzf /tmp/mybackup.tar.gz -C /data
      volumeMounts: ②
        - mountPath: /data
          name: application-data
      ...output omitted...
  volumes: ③
    - name: application-data
      persistentVolumeClaim:
        claimName: application-data
```

- ① Install the s3cmd tool, fetch the backup from the S3 bucket, and extract the archive inside the application data volume.
- ② Volume mount definition for the application data in the container.
- ③ Volume definition for the application PVC.

You can restore the data to an existing volume, or create a new one. If you use a new volume, then you must update the application to use that new volume.



Important

To avoid data corruption, ensure that the application that uses the volume is not running during the restoration procedure.



References

For accessing the internal registry, refer to the *Accessing the Registry* chapter in the Red Hat OpenShift Container Platform 4.14 *Registry* documentation at
https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/registry/index#accessing-the-registry

For exposing the internal registry, refer to the *Exposing the Registry* chapter in the Red Hat OpenShift Container Platform 4.14 *Registry* documentation at
https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/registry/index#securing-exposing-registry

For more information about Kubernetes Volume Snapshots, refer to the *Volume Snapshots* section in the Kubernetes documentation at
<https://kubernetes.io/docs/concepts/storage/volume-snapshots>

oc image mirror Usage Examples

https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/cli_tools/index#oc-image-mirror

skopeo-copy man page

<https://github.com/containers/skopeo/blob/main/docs/skopeo-copy.1.md>

skopeo-sync man page

<https://github.com/containers/skopeo/blob/main/docs/skopeo-sync.1.md>

yq documentation

<https://mikefarah.gitbook.io/yq/v/v3.x/>

► Guided Exercise

Export and Import Application Data and Settings

Export all application resources and data for a project, and import them to another project to create a functional copy of a live application.

Outcomes

- Export an OpenShift application, and include the settings, container images, and data.
- Restore an application to a different namespace.

Before You Begin

As the student user on the workstation machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

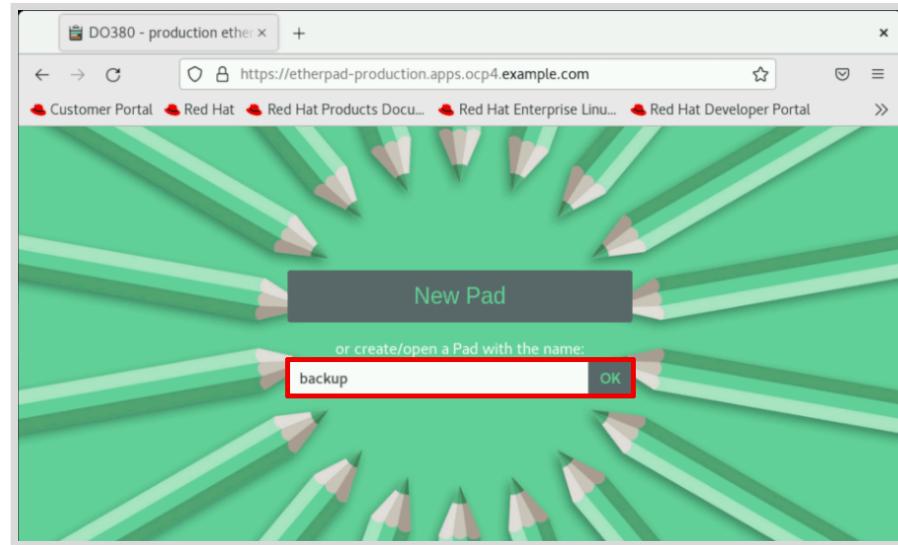
```
[student@workstation ~]$ lab start backup-export
```

Instructions

The development team published a new version of the etherpad application in the OpenShift internal registry and wants to deploy it in production. Before the deployment of the new version, your company requires you to create and restore a backup of the production application into a new stage project and to verify data integrity.

The application to back up is in the `production` project.

- 1. Connect to the production `etherpad` application and create a pad. You use this pad later in this exercise to validate the application restoration.
- 1.1. Open a web browser and navigate to the Etherpad URL:
 - `https://etherpad-production.apps.ocp4.example.com`
 - 1.2. Create a pad named `backup` and click `OK`.



- 1.3. Add a line to the pad with the current date and time, followed by **Production backup started**.
- 1.4. The application automatically saves changes to the pad. Close the browser tab.

► 2. Review the resources in the production project.

- 2.1. Open a terminal on the workstation and log in to the OpenShift cluster as the developer user with the developer password.

```
[student@workstation ~]$ oc login -u developer -p developer \
  https://api.ocp4.example.com:6443
Login successful.

You have one project on this server: "production"

Using project "production".
```

- 2.2. List the resources in the production project with the `oc get all` command.

```
[student@workstation ~]$ oc get all
NAME                                READY   STATUS    RESTARTS   AGE
pod/etherpad-6f9598bbb5-tmnt8     1/1     Running   0          109s

NAME           TYPE      CLUSTER-IP        EXTERNAL-IP      PORT(S)       AGE
service/etherpad ClusterIP  172.30.210.241  <none>        9001/TCP     111s

NAME           READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/etherpad  1/1     1           1          109s

NAME           DESIRED  CURRENT  READY   AGE
replicaset.apps/etherpad-6f9598bbb5  1       1       1      109s

NAME           ...      TAGS      UPDATED
imagestream.image.openshift.io/etherpad ...  1.8.18,1.9.1  109s
```

NAME	...	PORT	TERMINATION	WILDCARD
route.route.openshift.io/etherpad	...	http	edge/Redirect	None

**Note**

The `oc get all` command shows only a subset of all available resources in a namespace, and might omit some resources that the application requires.

2.3. The etherpad application deployment requires the following resource types:

- `PersistentVolumeClaim`
- `Service`
- `Route`
- `Deployment`

Use the `oc get` command to list these resource types in the `production` project, and compare the result with the previous step.

```
[student@workstation ~]$ oc get pvc,svc,route,deployment
NAME                      STATUS    ...    AGE
persistentvolumeclaim/etherpad   Bound    ...    43m

NAME          TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
service/etherpad   ClusterIP   172.30.96.26  <none>           9001/TCP   43m

NAME                      ...    PORT      TERMINATION      WILDCARD
route.route.openshift.io/etherpad  ...    http     edge/Redirect      None

NAME          READY    UP-TO-DATE      AVAILABLE      AGE
deployment.apps/etherpad  1/1       1            1            43m
```

▶ 3. Export the Kubernetes resources from the previous step and save them in the `~/D0380/labs/backup-export/production` directory.

3.1. Change to the `~/D0380/labs/backup-export` directory.

```
[student@workstation ~]$ cd ~/D0380/labs/backup-export
[student@workstation backup-export]$
```

3.2. Create a `production` directory.

```
[student@workstation backup-export]$ mkdir production
```

3.3. Export the persistent volume claim resource to a YAML file named `01-pvc.yaml`. Clean up the YAML file by removing the following fields:

- `metadata.annotations`
- `metadata.creationTimestamp`
- `metadata.namespace`
- `metadata.finalizers`
- `metadata.resourceVersion`
- `metadata.uid`

- spec.volumeName
- status

```
[student@workstation backup-export]$ oc get pvc etherpad -o yaml \
| yq d - metadata.annotations \
| yq d - metadata.creationTimestamp \
| yq d - metadata.namespace \
| yq d - metadata.finalizers \
| yq d - metadata.resourceVersion \
| yq d - metadata.uid \
| yq d - spec.volumeName \
| yq d - status \
> production/01-pvc.yml
```

- 3.4. Export the deployment resource to a YAML file named 02-deployment.yml. Clean up the YAML file by removing the following fields:

- metadata.annotations
- metadata.creationTimestamp
- metadata.namespace
- metadata.resourceVersion
- metadata.uid
- metadata.generation
- status

```
[student@workstation backup-export]$ oc get deployment etherpad -o yaml \
| yq d - metadata.annotations \
| yq d - metadata.creationTimestamp \
| yq d - metadata.namespace \
| yq d - metadata.resourceVersion \
| yq d - metadata.uid \
| yq d - metadata.generation \
| yq d - status \
> production/02-deployment.yml
```

- 3.5. Export the service resource to a YAML file named 03-service.yml. Clean up the YAML file by removing the following fields:

- metadata.annotations
- metadata.creationTimestamp
- metadata.namespace
- metadata.resourceVersion
- metadata.uid
- spec.clusterIP
- spec.clusterIPs
- status

```
[student@workstation backup-export]$ oc get svc etherpad -o yaml \
| yq d - metadata.annotations \
| yq d - metadata.creationTimestamp \
| yq d - metadata.namespace \
| yq d - metadata.resourceVersion \
| yq d - metadata.uid \
| yq d - spec.clusterIP* \
| yq d - status \
> production/03-service.yml
```

- 3.6. Export the route resource to a YAML file named `04-route.yml`. Clean up the YAML file by removing the following fields:

- `metadata.annotations`
- `metadata.creationTimestamp`
- `metadata.namespace`
- `metadata.resourceVersion`
- `metadata.uid`
- `spec.host`
- `status`

```
[student@workstation backup-export]$ oc get route etherpad -o yaml \
| yq d - metadata.annotations \
| yq d - metadata.creationTimestamp \
| yq d - metadata.namespace \
| yq d - metadata.resourceVersion \
| yq d - metadata.uid \
| yq d - spec.host \
| yq d - status \
> production/04-route.yml
```

- 3.7. Review the exported resource files.

```
[student@workstation backup-export]$ tree production
production/
├── 01-pvc.yml
├── 02-deployment.yml
├── 03-service.yml
└── 04-route.yml

0 directories, 4 files
```

- 4. As the `admin` user, expose the internal registry to enable users to export and import container images.

- 4.1. Log in to the OpenShift cluster as the `admin` user with the `redhatocp` password.

```
[student@workstation backup-export]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
Login successful.

...output omitted...
```

4.2. Expose the internal registry.

```
[student@workstation backup-export]$ oc patch \
configs.imageregistry.operator.openshift.io/cluster \
--patch '{"spec":{"defaultRoute":true}}' --type merge
config.imageregistry.operator.openshift.io/cluster patched
```

4.3. Wait until the openshift-apiserver operator is redeployed. It can take a couple of minutes. Press Ctrl+C to exit the watch command.

```
[student@workstation backup-export]$ watch -n10 oc get co openshift-apiserver
NAME          VERSION   AVAILABLE   PROGRESSING ... MESSAGE
openshift-apiserver  4.14.12   True        True      ...
APIServerDeploymentProgressing: deployment/apiserver.openshift-apiserver: 1/3
pods have been updated to the latest generation

openshift-apiserver  4.14.12   True        True      ...
APIServerDeploymentProgressing: deployment/apiserver.openshift-apiserver: 2/3
pods have been updated to the latest generation

openshift-apiserver  4.14.12   True        False     ...
^C
```

- 5. As the developer user, export all container images that are referenced in the etherpad image stream and save them in the ~/D0380/labs/backup-export/production directory.

5.1. Log in to the OpenShift cluster as the developer user with the developer password.

```
[student@workstation backup-export]$ oc login -u developer -p developer \
https://api.ocp4.example.com:6443
Login successful.

You have one project on this server: "production"

Using project "production".
```

5.2. Log in to the internal registry by using the `oc registry login` command.

```
[student@workstation backup-export]$ oc registry login
info: Using registry public hostname default-route-openshift-image-
registry.apps.ocp4.example.com
Saved credentials for default-route-openshift-image-registry.apps.ocp4.example.com
into /run/user/1000/containers/auth.json
```

- 5.3. Identify the fully qualified image name of the `etherpad` image and save it as the `IMAGE` environment variable.

```
[student@workstation backup-export]$ IMAGE=$(oc get is etherpad \
--template '{{.status.publicDockerImageRepository}}')

[student@workstation backup-export]$ echo ${IMAGE}
default-route-openshift-image-registry.apps.ocp4.example.com/production/etherpad
```

- 5.4. Export all container images from the `etherpad` image stream by using the `oc image mirror` command.

```
[student@workstation backup-export]$ oc image mirror \
${IMAGE} file://etherpad --dir=production
...output omitted...
sha256:057c...2513 file://etherpad:1.9.1
sha256:7265...8ec1 file://etherpad:1.8.18
info: Mirroring completed in 2.96s (116.3MB/s)
```

**Note**

The `oc image mirror` stores the exported images in the `v2` directory.

- 5.5. Review the exported images in the `production` directory.

```
[student@workstation backup-export]$ tree production
production/
...output omitted...
└── v2
    └── etherpad
        ├── blobs
        │   └── sha256:057c...a2513
        │   ...output omitted...
        └── manifests
            ├── 1.8.18 -> sha256:7265...28ec1
            ├── 1.9.1 -> sha256:057c...2513
            └── sha256:057c...2513
                └── sha256:7265...28ec1

4 directories, 38 files
```

- 6. Create a snapshot of the persistent volume to limit the downtime of the application during the backup operation.

- 6.1. List the persistent volumes and identify the associated storage class.

```
[student@workstation backup-export]$ oc get pvc
NAME      STATUS    ...      STORAGECLASS          AGE
etherpad  Bound     ...      ocs-external-storagecluster-ceph-rbd  25m
```

Chapter 2 | Backup, Restore, and Migration of Applications with OADP

- 6.2. Identify the driver that is used for the ocs-external-storagecluster-ceph-rbd storage class.

```
[student@workstation backup-export]$ oc get storageclass \
  ocs-external-storagecluster-ceph-rbd
NAME                                     PROVISIONER
ocs-external-storagecluster-ceph-rbd    openshift-storage.rbd.csi.ceph.com ...
```

- 6.3. Get the volume snapshot storage class name from the openshift-storage.rbd.csi.ceph.com driver.

```
[student@workstation backup-export]$ oc get volumesnapshotclasses \
  | egrep "^\w+|openshift-storage.rbd.csi.ceph.com"
NAME                                     ...
ocs-external-storagecluster-rbdplugin-snapclass ...
```

- 6.4. Scale down the application.

```
[student@workstation backup-export]$ oc scale deployment/etherpad --replicas 0
deployment.apps/etherpad scaled
```

- 6.5. Verify that no application pods are running.

```
[student@workstation backup-export]$ oc get pods
No resources found in production namespace.
```

- 6.6. Create a volume snapshot of the etherpad persistent volume by using the volume snapshot class from the previous steps. You can use the template in the ~/D0380/labs/backup-export/volumesnapshot.yaml path.

```
apiVersion: snapshot.storage.k8s.io/v1
kind: VolumeSnapshot
metadata:
  name: etherpad
spec:
  volumeSnapshotClassName: ocs-external-storagecluster-rbdplugin-snapclass
  source:
    persistentVolumeClaimName: etherpad
```

**Note**

You can use the solution in the ~/D0380/solutions/backup-export/volumesnapshot-etherpad.yaml file.

- 6.7. Create the volume snapshot resource.

```
[student@workstation backup-export]$ oc apply -f volumesnapshot.yaml
volumesnapshot.snapshot.storage.k8s.io/etherpad created
```

- 6.8. Verify that the volume snapshot is created and ready to use.

```
[student@workstation backup-export]$ oc get volumesnapshot
NAME      READYTOUSE   SOURCEPVC   ...   CREATIONTIME   AGE
etherpad   true        etherpad    ...   12s           13s
```

6.9. Scale up the application.

```
[student@workstation backup-export]$ oc scale deployment/etherpad --replicas 1
deployment.apps/etherpad scaled
```

- 7. Export the application data from the volume snapshot and save it in the ~/D0380/labs/backup-export/production directory.

- 7.1. Create a persistent volume claim by using the data from the volume snapshot. You can use the template in the ~/D0380/labs/backup-export/pvc-snapshot.yml path.

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: etherpad-snapshot
  labels:
    app: etherpad-snapshot
spec:
  storageClassName: ocs-external-storagecluster-ceph-rbd
  accessModes:
  - ReadWriteOnce
  dataSource:
    apiGroup: snapshot.storage.k8s.io
    kind: VolumeSnapshot
    name: etherpad
  resources:
    requests:
      storage: 1Gi
```



Note

You can use the solution in the ~/D0380/solutions/backup-export/pvc-snapshot.yml file.

- 7.2. Create the PVC snapshot resource.

```
[student@workstation backup-export]$ oc apply -f pvc-snapshot.yml
persistentvolumeclaim/etherpad-snapshot created
```

- 7.3. Verify that the persistent volume claim status is Bound.

```
[student@workstation backup-export]$ oc get pvc etherpad-snapshot
NAME      STATUS   ...   STORAGECLASS   AGE
etherpad-snapshot   Bound   ...   ocs-external-storagecluster-ceph-rbd   117s
```

- 7.4. Mount the snapshot volume on `/snapshot` in a new pod named `export-snapshot`. You can use the template in the `~/D0380/labs/backup-export/pod-snapshot.yml` file.

```
apiVersion: v1
kind: Pod
metadata:
  name: export-snapshot
  labels:
    app: export-snapshot
spec:
  containers:
    ...output omitted...
  volumeMounts:
    - mountPath: /snapshot
      name: snapshot
  terminationGracePeriodSeconds: 2
  volumes:
    - name: snapshot
      persistentVolumeClaim:
        claimName: etherpad-snapshot
```

**Note**

You can use the solution in the `~/D0380/solutions/backup-export/pod-snapshot.yml` file.

- 7.5. Create the pod snapshot resource.

```
[student@workstation backup-export]$ oc apply -f pod-snapshot.yml
pod/export-snapshot created
```

- 7.6. Verify that the `export-snapshot` pod is running.

```
[student@workstation backup-export]$ oc get pods export-snapshot
NAME          READY   STATUS    RESTARTS   AGE
export-snapshot 1/1     Running   0          31s
```

- 7.7. Copy the snapshot data from the pod to the `production/data` directory in the `workstation` machine.

```
[student@workstation backup-export]$ oc cp \
  export-snapshot:/snapshot production/data
tar: Removing leading '/' from member names
```

- 7.8. Review the exported data.

```
[student@workstation backup-export]$ tree production
production/
...output omitted...
├── data
│   └── dirty.db
```

```

|   └── lost+found
|   ├── minified_1f08...0f7e
|   ├── minified_1f08...0f7e.gz
|   ├── minified_5a47...57ca
|   ├── minified_5a47...57ca.gz
|   ├── minified_c24d...5b91
|   └── minified_c24d...5b91.gz
...output omitted...

```

6 directories, 45 files

7.9. Delete the export-snapshot pod.

```
[student@workstation backup-export]$ oc delete pod/export-snapshot
pod "export-snapshot" deleted
```

► 8. Import the application container to the stage project.

8.1. Create the stage project.

```
[student@workstation backup-export]$ oc new-project stage
Now using project "stage" on server "https://api.ocp4.example.com:6443".
...output omitted...
```

8.2. Import the etherpad container images to the stage project by using the oc image mirror command.

```
[student@workstation backup-export]$ oc image mirror \
--dir=production file://etherpad \
default-route-openshift-image-registry.apps.ocp4.example.com/stage/etherpad
...output omitted...
sha256:7265...8ec1 default-route.../stage/etherpad:1.8.18
sha256:057c...2513 default-route.../stage/etherpad:1.9.1
info: Mirroring completed in 4.98s (69.07MB/s)
```

8.3. Review the etherpad image stream.

```
[student@workstation backup-export]$ oc describe is/etherpad
...output omitted...
Image Repository: default-route.../stage/etherpad
Image Lookup: local=false
Unique Images: 2
Tags: 2

1.9.1
no spec tag

* image-registry...stage/etherpad@sha256:057c...2513
  24 seconds ago

1.8.18
no spec tag
```

```
* image-registry...stage/etherpad@sha256:7265...8ec1
  38 seconds ago
```

► 9. Import the production application data to the stage project.

9.1. Import the persistent volume claim definition.

```
[student@workstation backup-export]$ oc apply -n stage -f production/01-pvc.yml
persistentvolumeclaim/etherpad created
```

9.2. Verify that the persistent volume claim status is Bound.

```
[student@workstation backup-export]$ oc get pvc
NAME      STATUS    VOLUME          ...
etherpad   Bound     pvc-1d2139df-8969-4dc8-affd-c473e301480f ... 4s
...        AGE
```

9.3. Mount the volume on /data in a new pod named restore-snapshot. You can use the template in the ~/D0380/labs/backup-export/pod-restore.yml file.

```
apiVersion: v1
kind: Pod
metadata:
  name: restore-snapshot
  labels:
    app: restore-snapshot
spec:
  containers:
    ...output omitted...
    volumeMounts:
      - mountPath: /data
        name: data
  terminationGracePeriodSeconds: 2
  volumes:
    - name: data
      persistentVolumeClaim:
        claimName: etherpad
```



Note

You can use the solution in the ~/D0380/solutions/backup-export/pod-restore.yml file.

9.4. Create the restore-snapshot pod.

```
[student@workstation backup-export]$ oc apply -f pod-restore.yml
pod/restore-snapshot created
```

9.5. Verify that the restore-snapshot pod is running.

```
[student@workstation backup-export]$ oc get po restore-snapshot
NAME          READY   STATUS    RESTARTS   AGE
restore-snapshot   1/1     Running   0          65s
```

- 9.6. Copy the backup data from the workstation machine to the /data directory in the pod.

```
[student@workstation backup-export]$ oc rsync --no-perms \
./production/data/ restore-snapshot:/data/
sending incremental file list
dirty.db
minified_5dbe...03cb
minified_5dbe...03cb.gz
minified_8214...cacf
minified_8214...cacf.gz
minified_bd66...b58f
minified_bd66...b58f.gz

sent 776,429 bytes received 150 bytes 1,553,158.00 bytes/sec
total size is 775,441 speedup is 1.00
```

- 9.7. Review the imported data.

```
[student@workstation backup-export]$ oc exec restore-snapshot -- \
ls -1 /data
dirty.db
lost+found
minified_5dbe...03cb
minified_5dbe...03cb.gz
minified_8214...cacf
minified_8214...cacf.gz
minified_bd66...b58f
minified_bd66...b58f.gz
```

- 9.8. Delete the restore-snapshot pod.

```
[student@workstation backup-export]$ oc delete pod/restore-snapshot
pod "restore-snapshot" deleted
```

► 10. Import the application Kubernetes resources to the stage project.

- 10.1. Create a copy of the deployment YAML file and name it stage-deployment.yml.

```
[student@workstation backup-export]$ cp production/02-deployment.yml \
stage-deployment.yml
```

- 10.2. Modify the stage-deployment.yml file with the following parameters:

Parameter	Value
Container image	image-registry.openshift-image-registry.svc:5000/stage/etherpad:1.8.18
TITLE environment variable	D0380 - stage etherpad

```

apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app.kubernetes.io/name: etherpad
  name: etherpad
spec:
  ...
  spec:
    containers:
      - env:
          - name: TITLE
            value: D0380 - stage etherpad
          - name: DEFAULT_PAD_TEXT
            value: Add some content and write your ideas
          - name: SUPPRESS_ERRORS_IN_PAD_TEXT
            value: "true"
          - name: EXPOSE_VERSION
            value: "true"
        image: image-registry.openshift-image-registry.svc:5000/stage/
etherpad:1.8.18
  ...
  initContainers:
    - args:
        ...
      image: image-registry.openshift-image-registry.svc:5000/stage/
etherpad:1.8.18
      imagePullPolicy: IfNotPresent
      name: clean
  ...

```



Note

The TITLE environment variable defines the instance name for this application. It is also the name of the browser window when accessing the application.

10.3. Import the deployment definition to the stage project.

```
[student@workstation backup-export]$ oc apply -n stage -f stage-deployment.yml
deployment.apps/etherpad created
```

10.4. Import the service definition to the stage project.

```
[student@workstation backup-export]$ oc apply -n stage \
-f production/03-service.yml
service/etherpad created
```

10.5. Import the route definition to the `stage` project.

```
[student@workstation backup-export]$ oc apply -n stage \
-f production/04-route.yml
route.route.openshift.io/etherpad created
```

10.6. Review all imported resources.

```
[student@workstation backup-export]$ oc get pvc,is,svc,route,deployment
NAME                                     STATUS    ...
persistentvolumeclaim/etherpad   Bound     ...

NAME                               ...      TAGS        UPDATED
imagestream.image.openshift.io/etherpad ...  1.8.18,1.9.1  4 minutes ago

NAME          TYPE      CLUSTER-IP      EXTERNAL-IP    PORT(S)    AGE
service/etherpad ClusterIP  172.30.222.62 <none>       9001/TCP   29s

NAME                           HOST/PORT
route.route.openshift.io/etherpad  etherpad-stage.apps.ocp4.example.com  ...

NAME          READY  UP-TO-DATE  AVAILABLE  AGE
deployment.apps/etherpad  1/1    1          1          103s
```

► **11.** Connect to the stage `etherpad` application and open the backup pad.

11.1. Open a web browser and navigate to the stage Etherpad application URL. Confirm the instance name in the browser window.

- <https://etherpad-stage.apps.ocp4.example.com>

11.2. Open and review the pad named `backup`.

If the restoration is successful, then you should see the same content from the beginning of this section.

11.3. Close the browser tab.

► **12.** Clean up the resources.

12.1. Change to the student `HOME` directory.

```
[student@workstation backup-export]$ cd
[student@workstation ~]$
```

12.2. Log in to the OpenShift cluster as the `admin` user with the `redhatocp` password.

```
[student@workstation ~]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
Login successful.

...output omitted...
```

12.3. Remove the exposed route from the OpenShift internal registry.

```
[student@workstation ~]$ oc patch \
configs.imageregistry.operator.openshift.io/cluster \
--patch '{"spec":{"defaultRoute":false}}' --type merge
config.imageregistry.operator.openshift.io/cluster patched
```



Note

The cluster image registry patch triggers a kube-apiserver deployment rollout. The `lab finish` command might take up to 10 minutes for the cluster to stabilize.

12.4. Delete the `stage` project.

```
[student@workstation ~]$ oc delete project stage
project.project.openshift.io "stage" deleted
```

Finish

On the `workstation` machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish backup-export
```

OADP Operator Deployment and Features

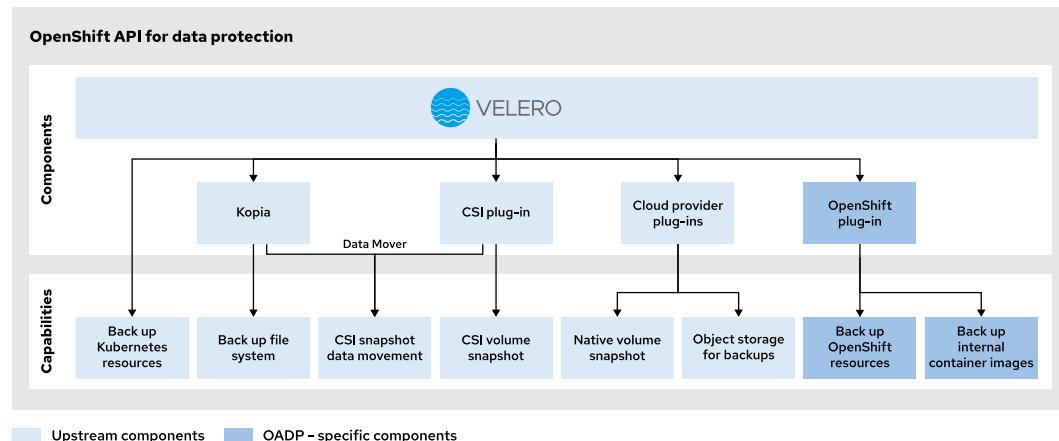
Objectives

- Deploy the OADP operator and describe its features and dependencies.

OpenShift API for Data Protection

The OpenShift API for Data Protection (OADP) product provides a native backup solution for applications that run on OpenShift.

OADP uses several components such as Velero and Kopia to back up all the Kubernetes resources, container images from the internal registry, and persistent volumes that are associated with an application.



The following list describes the main components that OADP uses:

Velero

Velero is the main upstream component of OADP. It provides the backup API to the users and uses plug-ins to add extended capabilities to OADP.

Data Mover

This feature enables exporting snapshot content to object storage by using Kopia and the CSI plug-in.

Kopia

Open source backup tool that Velero and Data Mover use to back up persistent volumes.

Those components are discussed in more detail later in this chapter. See the references section for more information about those components.

The OADP framework is initially designed for partners to provide the needed infrastructure for a complete backup solution. OpenShift administrators can use OADP as a basic backup solution, although it might not implement some required features, such as a graphical user interface or multi-tenancy support. The following examples are backup solutions that use OADP:

- Catalogic CloudCasa

Chapter 2 | Backup, Restore, and Migration of Applications with OADP

- IBM Spectrum Protect Plus
- Dell EMC PowerProtect Data Manager

Kubernetes Resources Backup

OADP uses Velero to back up all the Kubernetes resources to cloud object storage. You can filter objects to back up and restore by namespace, resource type, and label.

Velero not only exports and imports resources; it also performs additional processing on the resources. The following list describes some of this additional processing:

- Determine the correct restore order for each resource.
- Apply any user-provided filters.
- Modify the destination namespace, if it is different from the backup.
- Remove or alter resource attributes such as auto-assigned node ports, IP addresses, or hostnames.
- Skip managed resources such as a replica set that deployment owns.

You can add backup and export logic by using Velero plug-ins.

Data Volume Snapshot

OADP can back up volumes by using cloud native snapshot API or the Kubernetes Container Storage Interface (CSI) snapshot API.

Red Hat recommends the use of snapshots when possible, because this feature enables fast and consistent backup.

File System Backup

For volumes that are not supported by either a cloud-native snapshot API or the CSI snapshot API, OADP provides a file-system backup solution.

File System Backup (FSB), or Pod Volume Backup, enables backing up almost any kind of Kubernetes volume that does not support the snapshot capability, such as NFS and local storage.



Warning

FSB does not support hostPath volumes.

OADP uses the open source Kopia tool to back up the application volume file system.

Kopia accesses the volume file system from the volume mount point on the cluster node. For that reason, the application pod must be running during the backup.

During the file-system backup, the application can still alter the file system and corrupt the backup. To prevent such a scenario, the application must be quiesced to prevent any write operations while the backup is in progress.

Application-Consistent Backup

You can use backup and restore hooks to run commands before and after the backup and restore operation. By using a pre-hook, it is possible to quiesce an application, such as a transactional

database, to flush pending I/O operations to the storage back end and pause future application writing operations to avoid any data loss or corruption during the backup. When the backup is complete, a post-hook is used to resume the application operations.

If you use volume snapshots, then the hooks are executed just before and after the snapshot creation, which significantly reduces the time that the application is in read-only mode.

OADP Data Mover

Some CSI drivers might store the snapshot in the same storage location as the original volume and might be vulnerable to data loss in the event of a disaster. The storage solution that is used with Red Hat OpenShift Container Platform might not guarantee snapshot durability. In such scenarios, you can enable the OADP Data Mover to copy the snapshot content to a remote backup storage location and remove the snapshot from the storage back end.

Data Mover uses Kopia to upload and encrypt snapshot content to a unified backup repository on the object storage. The Kopia unified backup repository is explained later in this chapter.



Note

Earlier OADP versions than v1.3 use Restic instead of Kopia to back up persistent volumes. For OADP v1.3, Restic is still available at installation, but Data Mover requires you to enable Kopia.

Velero Plug-ins

You can extend the OADP capabilities and backup logic by using the Velero plug-in system. Support for various cloud providers, storage solutions, and additional CRDs are available with integrated plug-ins.

OADP includes the following plug-ins that you can enable during the installation:

`aws`

Stores and retrieves backups on AWS S3-compatible storage. Manages volume snapshots on AWS EBS volumes.

`gcp`

Stores and retrieves backups on Google Cloud Storage. Manages volume snapshots on Google Compute Engine Disks.

`azure`

Stores and retrieves backups on Azure Blob Storage. Manages volume snapshots on Azure Managed Disks.

`openshift`

Provides additional logic to back up and restore OpenShift resources, including container images from the internal registry.

`kubevirt`

Provides additional logic to back up and restore virtual machines and other OpenShift Virtualization resources.

`csi`

Provides volume snapshot support by using the Kubernetes CSI Snapshot API.

OADP can also use third-party plug-ins, such as the following examples, to extend support to additional storage solutions and custom Kubernetes resources.

Chapter 2 | Backup, Restore, and Migration of Applications with OADP

- OpenStack plug-in for object storage on Swift and volume snapshots on Cinder.
- HPE plug-in for volume snapshots on HPE Nimble Storage.
- DigitalOcean plug-in for volume snapshots on DigitalOcean Volumes Block Storage.



Important

Red Hat does not support third-party plug-ins.

OADP API Resources

The OADP operator provides the following API resources:

DataProtectionApplication

The `dataProtectionApplication` resource is the configuration for the OADP operator and its components.

BackupStorageLocation

The OADP operator creates and manages a `backupStorageLocation` resource for each storage location that is defined in the `dataProtectionApplication` configuration.

VolumeSnapshotLocation

The OADP operator creates and manages a `volumeSnapshotLocation` resource for each snapshot location that is defined in the `dataProtectionApplication` configuration.

Backup

The `backup` resource requests the Velero server to create a backup.

Restore

The `restore` resource requests the Velero server to restore a backup.

Schedule

The `schedule` resource creates a backup periodically on a given schedule by using the Cron format.

BackupRepository

The `BackupRepository` resource tracks the lifecycle of the backup repositories.

Requirements for OADP

Red Hat OpenShift API for Data Protection requires the following resources.

OpenShift User Permissions

OADP configuration requires a user with the `cluster-admin` role. Backup and restore operations require administrative privileges in the OADP namespace (`openshift-adp` by default).



Important

OADP deployment has administrative privileges on the cluster to back up and restore all resources in any namespaces.

Object Storage

OADP requires an object storage location to store the backups. The following list is an example of object storage providers that are available for storing backups:

- Amazon Web Services

- Microsoft Azure
- Google Cloud Platform
- Any AWS S3-compatible provider, such as MinIO or OpenShift Data Foundation

Object Storage with OpenShift Data Foundation

OpenShift Data Foundation (ODF) is the storage solution that is used in this classroom environment. ODF provides two object storage services: the MultiCloud Object Gateway (NooBaa MCG) and the Ceph RADOS Object Gateway (Ceph RGW). Both services are fully compatible with OADP by using the aws provider plug-in.



Note

For more details about OpenShift Data Foundation, refer to the *DO370: Enterprise Kubernetes Storage with Red Hat OpenShift Data Foundation* training course.

Similar to a persistent volume claim, you can use an object bucket claim to request an S3-compatible bucket. ODF creates a secret and a configuration map with the same name as the object bucket claim with the credentials and information to access the bucket.

The following is an example of an object bucket claim definition:

```
apiVersion: objectbucket.io/v1alpha1
kind: ObjectBucketClaim
metadata:
  name: my-bucket-claim
  namespace: my-namespace
spec:
  storageClassName: storage-class-name ①
  generateBucketName: my-bucket ②
```

- ① Storage class name for the object bucket. To use NooBaa MCG, set the value to `openshift-storage.noobaa.io`. To use the Ceph RGW gateway, set the value to `ocs-external-storagecluster-ceph-rgw`.
- ② Prefix to add to the generated bucket name.

A few minutes after the object bucket claim creation, the status is set to Bound when the bucket is created and ready to use.

```
[user@host ~]$ oc get obc
NAME           STORAGE-CLASS          PHASE   AGE
my-bucket-claim  openshift-storage.noobaa.io  Bound   5m
```

ODF creates a configuration map with the bucket information in the same namespace as the object bucket claim:

```
[user@host ~]$ oc get configmap/my-bucket-claim -o yaml
apiVersion: v1
data:
  BUCKET_HOST: s3.openshift-storage.svc ①
  BUCKET_NAME: my-bucket-9ce46e22-2fb8-4a46-af95-f6949d87c3fd ②
  BUCKET_PORT: "443" ③
```

```
BUCKET_REGION: "" ④
BUCKET_SUBREGION: ""
kind: ConfigMap
...output omitted...
```

- ①** S3 API URL, for internal use only.
- ②** The generated name of the bucket.
- ③** The TCP port for the S3 endpoint. The port 443 is for https protocol.
- ④** The region name, if used by the storage class.

If the BUCKET_HOST uses an internal service URL that ends with .svc, then the URL can be reached only from a workload that is running inside the cluster. The s3 route in the openshift-storage namespace provides the external endpoint URL.

```
[user@host ~]$ oc get route/s3 -n openshift-storage \
-o jsonpath='{.spec.host}{"\n"}'
s3-openshift-storage.apps.ocp4.example.com
```

The S3 bucket credentials are stored in a secret in the same namespace as the object bucket claim:

```
[user@host ~]$ oc extract secret/my-bucket-claim --to=-
# AWS_ACCESS_KEY_ID
YEAsbMJnG3o1bGANZprt
# AWS_SECRET_ACCESS_KEY
xjaeCDhskn7lfrdA7WqzoUxpiRYuyjc9uDaWlMw3
```

CSI Snapshot

To use volume snapshots with CSI drivers, a volume snapshot class must be created to register the CSI driver to use to create snapshots. The driver name must be the same as the provisioner attribute of the volume storage class to back up.

```
[user@host ~]$ oc get storageclasses
NAME                      PROVISIONER
nfs-storage (default)      k8s-sigs.io/nfs-subdir-external-provisioner
ocs-storagecluster-ceph-rbd  openshift-storage.rbd.csi.ceph.com
ocs-storagecluster-ceph-rgw   openshift-storage.ceph.rook.io/bucket
ocs-storagecluster-cephfs    openshift-storage.cephfs.csi.ceph.com
openshift-storage.noobaa.io  openshift-storage.noobaa.io/obc

[user@host ~]$ oc get volumesnapshotclasses
NAME                      DRIVER
ocs-storagecluster-cephfsplugin-snapclass  openshift-storage.cephfs.csi.ceph.com
ocs-storagecluster-rbdplugin-snapclass       openshift-storage.rbd.csi.ceph.com
```



Note

For storage classes that do not support snapshots, such as nfs-storage in the previous example, you can use the file-system backup.

OADP Installation and Configuration

You can install the OADP operator from the OperatorHub or by using the oc command. See the references section for instructions to install an operator by using the OperatorHub or the oc command.



Note

For more details about Kubernetes operators, refer to the *DO280: Red Hat OpenShift Administration II: Operating a Production Kubernetes Cluster* training course.

The `dataProtectionApplication` resource defines the configuration and components that the OADP operator manages. The following YAML file is an example of a `dataProtectionApplication` resource that stores backups on AWS S3 storage:

```
apiVersion: oadp.openshift.io/v1alpha1
kind: DataProtectionApplication
metadata:
  name: oadp-backup
  namespace: openshift-adp
spec:
  configuration:
    velero:
      defaultPlugins:
        - aws
        - openshift
        - csi
      defaultSnapshotMoveData: true
    nodeAgent:
      enable: true
      uploaderType: kopia
    backupLocations:
      - velero:
          config:
            profile: "default"
            region: "us-east-1"
          provider: aws
          default: true
          credential:
            key: cloud
            name: cloud-credentials
        objectStorage:
          bucket: my-bucket
          prefix: oadp
```

The OADP configuration is composed of several sections.

Components Configuration

The `configuration` section describes the configuration OADP components, such as Velero and Kopia.

```
configuration:  
  velero:  
    defaultPlugins: ❶  
      - aws  
      - openshift  
      - csi  
    defaultSnapshotMoveData: true ❷  
    nodeAgent: ❸  
      enable: true ❹  
    uploaderType: kopia ❺  
    podConfig: ❻  
      resourceAllocations:  
        limits:  
          cpu: "1"  
          memory: 8Gi  
        requests:  
          cpu: 500m  
          memory: 256Mi
```

- ❶ Velero plug-ins to enable. The `openshift` plug-in is mandatory. The `csi` plug-in is necessary for snapshots and Data Mover.
- ❷ Set this parameter to true to use the default Data Mover.
- ❸ The agent that manages volume backups.
- ❹ Enable the node agent if you use either FSB or Data Mover.
- ❺ You can use either Kopia or Restic as the uploader. If you use the Data Mover, then you must use Kopia.
- ❻ Optional pod configuration such as node selector, labels, and resource allocation for each component.

Backup Storage Location

The `backupLocation` section describes the object storage location for backups. You can define multiple backup locations. The following example uses the `aws` provider plug-in:

```
backupLocations:  
  - velero:  
    config: ❶  
      profile: "default" ❷  
      region: "us-east-1"  
      s3Url: https://s3.openshift-storage.svc  
      s3ForcePathStyle: "true"  
      insecureSkipTLSVerify: "true"  
    provider: aws ❸  
    default: true ❹  
    credential: ❺  
      key: cloud  
      name: cloud-credentials  
    objectStorage:
```

```
bucket: my-bucket 6
prefix: oadp 7
caCert: LLS0S0...LS0tLS0K 8
```

- 1** Storage provider-specific configuration. In this example, the configuration is for the aws provider.
- 2** The credentials profile name that is defined in the Velero secret.
- 3** The storage provider plug-in name.
- 4** Use this backup location as the default one if none is specified in the backup and restore resources.
- 5** Velero secret with the storage provider credentials.
- 6** Bucket name from the storage provider.
- 7** Optional subdirectory in the bucket to store backups. If not specified, then the backups are stored at the root level of the bucket.
- 8** Optional CA bundle in Base64 format to verify TLS connections to the storage provider.

With the aws provider plug-in, you can configure any S3-compatible object storage by setting the `s3Url` attribute with a custom S3 endpoint URL. OADP can access the S3 endpoint by using two addressing models:

- The path-style model, such as `https://s3.mydomain.com/bucket-name/myfile.txt`
- The DNS-style model (also known as virtual-hosted style), such as `\https://bucket-name.s3.mydomain.com/myfile.txt` This method is the default.

If the object storage solution does not support the DNS-style method, then you can configure OADP to use the path-style model instead with the `s3ForcePathStyle` option.

The `region` attribute must be set even if the storage solution does not use a region. You can then set this attribute to the standard AWS region, such as "us-east-1".



Note

ImageStream backup does not work with the `caCert` option. If the S3 endpoint certificate is not trusted, then the `insecureSkipTLSVerify` attribute is required to back up container images.

The Velero secret contains the credentials to access the object storage. If the secret name is not specified in the configuration, then Velero uses the `cloud-credentials` default secret name. You can create the Velero secret with a configuration file that is specific to the provider plug-in. The following Velero configuration file is for the aws provider plug-in:

```
[user@host ~]$ cat credentials-velero

[myProfile] 1
aws_access_key_id=AWS_ACCESS_KEY_ID
aws_secret_access_key=AWS_SECRET_ACCESS_KEY
```

- ① The profile name for the credentials. The name must match the profile name in the backup storage location configuration.

You can then create the Velero secret with the following `oc create secret` command:

```
[user@host ~]$ oc create secret generic cloud-credentials \
--from-file cloud=credentials-velero
```



Note

The Velero secret must be created before the `DataProtectionApplication` object, or the installation will fail.

Snapshot Storage Location

The `snapshotLocations` section describes the storage configuration to use for volume snapshots when using the cloud provider native snapshot capability. For CSI snapshots, this section is not needed, because volume snapshot classes are used instead. With Data Mover and FSB, backups are stored in the defined object storage location in the `backupLocations` section.

You can configure multiple snapshot locations per provider. The following example defines one snapshot location by using the `aws` provider plug-in:

```
snapshotLocations:
- name: default
  velero:
    provider: aws
    config:
      region: us-west-2 ①
      profile: "default" ②
      credential: ③
        key: cloud
        name: cloud-credentials
```

- ① AWS region to use to create the snapshots.
- ② The credentials profile name that is defined in the Velero secret.
- ③ **Optional:** Velero secret with the storage provider credentials.

If the secret name is not specified in the configuration, then Velero uses the `cloud-credentials` default secret name. The same secret can be shared with the snapshot location and the backup location.



Note

Depending on the storage provider, the region to store snapshots might be limited to the same region where the original volume is. Cross-provider snapshots are not supported.

See the references section for configuring OADP with other cloud providers such as Microsoft Azure and Google Cloud Platform.

CSI Snapshots

For OADP to use a volume snapshot class, you must add the `velero.io/csi-volumesnapshot-class: "true"` label. Only one volume snapshot class per driver needs this label.

You can use the following command to add the required label:

```
[user@host ~]$ oc label volumesnapshotclass my-snapshot-storageclass \
  velero.io/csi-volumesnapshot-class="true"
my-snapshot-storageclass labeled
```

Verification and Troubleshooting

When the configuration is complete, several OADP components are created depending on the enabled features.

```
[user@host ~]$ oc -n openshift-adp get deploy
NAME                      READY   UP-TO-DATE   AVAILABLE   AGE
openshift-adp-controller-manager   1/1     1           1          24h ①
velero                     1/1     1           1          25s ②
```

- ① OADP controller that is deployed on the OADP operator installation.
- ② Velero deployment that manages backup and restore operations.

If the Node Agent that manages volume backups is enabled, then the following daemon set is created:

```
[user@host ~]$ oc -n openshift-adp get daemonset
NAME      DESIRED  CURRENT  READY  UP-TO-DATE  AVAILABLE  NODE SELECTOR AGE
node-agent  3        3        3      3           3          <none>    23s ①
```

- ① The node-agent is deployed on all compute nodes, to back up volume file systems.

The OADP operator creates and validates each backup storage location that is defined in the `DataProtectionApplication` object. OADP tries to connect to the object storage with the provided credentials and validates the configuration. If the configuration is correct, then the backup storage location enters the `Available` phase.

```
[user@host ~]$ oc -n openshift-adp get backupstoragelocation
NAME      PHASE      LAST VALIDATED   AGE      DEFAULT
oadp-backup-1  Available  7s            2m44s  true
```

If the backup storage location remains in the `Unavailable` phase, then you can use the `oc describe` command to get more information and error messages from the `Status` field:

```
[user@host ~]$ oc -n openshift-adp describe \
  backupstoragelocation oadp-backup-1
...output omitted...
Status:
  Last Synced Time: 2023-09-22T10:27:27Z
  Last Validation Time: 2023-09-22T12:03:27Z
  Message: BackupStorageLocation "oadp-backup-1" is unavailable:
  rpc error: code = Unknown desc = InvalidAccessKeyId: The AWS access key Id you
  provided does not exist in our records.
  Phase: Unavailable
...output omitted...
```

In this example, the S3 credentials that are configured in the Velero secret are incorrect.

**Note**

The OADP operator checks the validity of the backup storage location configuration every minute. Review the `last validated` field of the backup storage location to determine when the last check ran.

You can use the AWS CLI or the s3cmd tool to validate the S3 configuration and to browse the content of the S3 bucket.

To configure the s3cmd tool, create a `.s3cfg` file in your home directory with your S3 credentials. You can use the following configuration template:

```
access_key = AWS_ACCESS_KEY_ID ①
secret_key = AWS_SECRET_ACCESS_KEY
host_base = s3.mydomain.com ②
host_bucket = s3.mydomain.com/%(bucket)s ③
signature_v2 = True ④
```

- ① S3 access key and secret key.
- ② S3 endpoint URL, if the AWS S3 endpoint is not used.
- ③ S3 bucket address model, if your storage provider does not support the DNS-style model.
- ④ If your storage provider does not support the new AWS v4 signature, then you can use the AWS v2 signature instead. OpenShift Data Foundation requires this setting.

To list the contents of a bucket and to validate the configuration, use the `s3cmd ls` command to list all available buckets and the `s3cmd la` command to list all objects in all buckets:

**Note**

If the S3 bucket is empty, then the `s3cmd la` command returns an empty line.

```
[user@host ~]$ s3cmd ls  
2023-09-22 12:33  s3://my-bucket  
  
[user@host ~]$ s3cmd la  
2023-09-22 12:31          143  s3://my-bucket/somefile.txt  
...output omitted...
```



References

For more information about installing operators by using the OperatorHub, refer to the *Installing from OperatorHub Using the Web Console* section in the *Administrator Tasks* chapter in the Red Hat OpenShift Container Platform 4.14 Operators documentation at

https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/operators/index#olm-installing-from-operatorhub-using-web-console_olm-adding-operators-to-a-cluster

For more information about installing operators by using the command line, refer to the *Installing from OperatorHub Using the CLI* section in the *Administrator Tasks* chapter in the Red Hat OpenShift Container Platform 4.14 Operators documentation at

https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/operators/index#olm-installing-operator-from-operatorhub-using-cli_olm-adding-operators-to-a-cluster

For more information about configuring OADP with other cloud object storage, refer to the *Installing and Configuring OADP* section in the *Application Backup and Restore* chapter in the Red Hat OpenShift Container Platform 4.14 Backup and Restore documentation at

https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/backup_and_restore/index#installing-and-configuring-oadp

For more information about the OADP v1.3 Data Mover, refer to the *OADP 1.3 Data Mover* section in the *Application Backup and Restore* chapter in the Red Hat OpenShift Container Platform 4.14 Backup and Restore documentation at

https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/backup_and_restore/index#oadp-1-3-data-mover

Velero

<https://velero.io/>

Kopia

<https://kopia.io/>

S3cmd: Command Line S3 Client

<https://s3tools.org/s3cmd>

► Guided Exercise

OADP Operator Deployment and Features

Deploy the OADP operator and perform a backup to validate that OADP is functional in a cluster.

Outcomes

- Install and configure the OpenShift API for Data Protection operator.
- Back up an application by using OADP to validate the configuration.

Before You Begin

As the student user on the `workstation` machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start backup-oadp
```

Instructions

Install and configure the OADP operator with the CSI snapshot and the OADP Data Mover feature.

Create an S3-compatible object storage bucket with OpenShift Data Foundation and configure it as a backup storage location.

To validate the configuration, back up the application in the `production` project by using OADP and verify that the backup is successfully created in the S3 bucket.



Important

All other exercises in this chapter depend on your cluster being correctly configured to use OpenShift API for Data Protection.

If you cannot complete this exercise, then delete and re-create your lab environment. After re-creating your lab environment, you can either attempt this exercise again, or the start script for another exercise can configure OADP for you.

► 1. As the `admin` user, locate and go to the OpenShift web console.

1.1. Log in to your OpenShift cluster as the `admin` user.

```
[student@workstation ~]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
Login successful.

...output omitted...
```

- 1.2. Identify the URL for the OpenShift web console.

```
[student@workstation ~]$ oc whoami --show-console
https://console-openshift-console.apps.ocp4.example.com
```

- 1.3. Open a web browser and go to <https://console-openshift-console.apps.ocp4.example.com>.
- 1.4. Click **Red Hat Identity Management** and log in as the **admin** user with **redhatocp** as the password.

► 2. Install the OADP operator.

- 2.1. Click **Operators > OperatorHub**. In the **Filter by keyword** field, type **OADP** to locate the OADP operator, and then click **OADP Operator**.

The screenshot shows the Red Hat OpenShift OperatorHub interface. On the left is a sidebar with navigation links: Home, Operators (selected), OperatorHub (highlighted), Installed Operators, Workloads, Networking, Storage, Builds, Observe, and Compute. The main area is titled 'OperatorHub' and displays a list of operators categorized under 'All Items'. The categories include Cloud Provider, Developer Tools, Logging & Tracing, Modernization & Migration, Monitoring, OpenShift Optional, Security, Storage, and Other. A search bar at the top right contains the text 'OADP', which is also highlighted with a red box. Below the search bar, a result card for the 'do380 Operator Catalog Cs' operator is shown, also highlighted with a red box. The result card includes a small icon of a person wearing a red hat, the name 'do380 Operator Catalog Cs', the provider 'provided by Red Hat', and a brief description: 'OADP (OpenShift API for Data Protection) operator sets up and installs Data Protection...'.

- 2.2. The web console displays information about the OADP operator. Select the **stable-1.3** channel and the **1.3.1** version. Click **Install** to proceed to the **Install Operator** page.



OADP Operator

1.3.1 provided by Red Hat

Install

Channel
stable-1.3 ▾

Version
1.3.1 ▾

Capability level

Basic Install
 Seamless Upgrades
 Full Lifecycle
 Deep Insights
 Auto Pilot

OpenShift API for Data Protection (OADP) operator sets up and installs Velero on the OpenShift platform, allowing users to backup and restore applications.

Backup and restore Kubernetes resources and internal images, at the granularity of a namespace, using a version of Velero appropriate for the installed version of OADP.

OADP backs up Kubernetes objects and internal images by saving them as an archive file on object storage. OADP backs up persistent volumes (PVs) by creating snapshots with the native cloud snapshot API or with the Container Storage Interface (CSI). For cloud providers that do not support snapshots, OADP backs up resources and PV data with Restic or Kopia.

- [Installing OADP for application backup and restore](#)
- [Installing OADP on a ROSA cluster and using STS, please follow the Getting Started Steps 1-3 in order to obtain the role ARN needed for using the standardized STS configuration flow via OLM](#)
- [Frequently Asked Questions](#)

Source

do380 Operator Catalog
Cs

- 2.3. Review the default configuration and click **Install** to install the operator.
 - 2.4. Wait until the installation is complete and the web console displays ready for use.

A screenshot of the OpenShift Operator Catalog interface. At the top, there's a search bar with the placeholder "Search Operators". Below it, a card for the "OADP Operator" is displayed. The card features a red icon of a person wearing a red hat, the text "OADP Operator", "oadp-operator.v1.3.1 provided by Red Hat", and a green circular button with a white checkmark. A large green banner at the bottom states "Installed operator: ready for use". Below the banner are two blue buttons: "View Operator" and "View installed Operators in Namespace openshift-adp".

- 3. Create an S3-compatible object storage bucket for OADP to store backups.

- 3.1. From the command line, change to the openshift-adp project.

```
[student@workstation ~]$ oc project openshift-adp
Now using project "openshift-adp" on server "https://api.ocp4.example.com:6443".
```

- ### 3.2. Change to the ~/D0380/labs/backup-oadp directory

```
[student@workstation ~]$ cd ~/D0380/labs/backup-oadp  
[student@workstation backup-oadp]$
```

- 3.3. Create an object bucket claim by using the following values

Parameter	Value
Name	backup
Namespace	openshift-adp
Storage class name	openshift-storage.noobaa.io
Generate bucket name	backup

You can use the resource definition in the ~/D0380/labs/backup-oadp/obc-backup.yml path.

```
apiVersion: objectbucket.io/v1alpha1
kind: ObjectBucketClaim
metadata:
  name: backup
  namespace: openshift-adp
spec:
  storageClassName: openshift-storage.noobaa.io
  generateBucketName: backup
```



Important

This environment uses ODF as a back end for both application data storage and backup storage for simplicity. In a disaster situation, both data and backups can be lost at the same time.

For production environment, Red Hat recommends using a different storage location for backups and for application data.

3.4. Create the backup OBC.

```
[student@workstation backup-oadp]$ oc apply -f obc-backup.yml
objectbucketclaim.objectbucket.io/backup created
```

3.5. Verify that the object bucket claim is created and in the Bound phase.

```
[student@workstation backup-oadp]$ oc get obc
NAME      STORAGE-CLASS          PHASE      ...
backup    openshift-storage.noobaa.io  Bound      ...
```

- ▶ 4. Retrieve the S3 bucket information and credentials. You can use a second terminal for this step to make it easier to copy and paste values in the remainder of this section.

When an object bucket claim is created, Red Hat OpenShift Data Foundation creates a matching secret and configuration map with the same name that contains the bucket information and credentials.

4.1. Retrieve the bucket name and bucket host from the generated configuration map.

```
[student@workstation backup-oadp]$ oc extract --to-- cm/backup
# BUCKET_HOST
s3.openshift-storage.svc
# BUCKET_NAME
backup-7d9...f4c
# BUCKET_PORT
443
# BUCKET_REGION

# BUCKET_SUBREGION
```

**Note**

The s3.openshift-storage.svc service uses a TLS certificate that is signed with the self-signed service CA that OpenShift manages. To prevent a certificate signed by unknown authority error, you must include the CA certificate in the OADP configuration.

- 4.2. Retrieve the service CA certificate for the s3.openshift-storage.svc endpoint. This certificate is available in the openshift-service-ca.crt configuration map in any namespace.

Encode the certificate in Base64 format and save the value for the next step.

```
[student@workstation backup-oadp]$ oc get cm/openshift-service-ca.crt \
-o jsonpath='{.data.service-ca\.crt}' | base64 -w0; echo
...output omitted...
```

- 4.3. Retrieve the bucket credentials from the generated secret.

```
[student@workstation backup-oadp]$ oc extract --to-- secret/backup
# AWS_ACCESS_KEY_ID
JNmbmhD0AQ3BFMABtXC4
# AWS_SECRET_ACCESS_KEY
xjkwp8bXeJazTgC4u/WJTbzgiD0tfWgt80tdADLz
```

**Note**

The values might be different in your environment.

- 4.4. Identify the public URL for the S3 endpoint from the s3 route in the openshift-storage namespace. You use this URL to connect to the S3 bucket from the workstation machine in the next step.

```
[student@workstation backup-oadp]$ oc get route s3 -n openshift-storage
NAME      HOST/PORT          ...
s3        s3-openshift-storage.apps.ocp4.example.com ...
```

- 5. Configure the s3cmd command to browse and validate the object storage configuration.

Chapter 2 | Backup, Restore, and Migration of Applications with OADP

- 5.1. Create a `~/.s3cfg` configuration file in the student home directory with the S3 credentials and S3 endpoint URL from the previous step.

You can use the sample configuration file in the `~/D0380/labs/backup-oadp/s3cfg` path.

```
[student@workstation backup-oadp]$ vim ~/.s3cfg
access_key = AWS_ACCESS_KEY_ID
secret_key = AWS_SECRET_ACCESS_KEY
host_base = s3-openshift-storage.apps.ocp4.example.com
host_bucket = s3-openshift-storage.apps.ocp4.example.com/%(bucket)s
signature_v2 = True
```

**Note**

Replace the `AWS_ACCESS_KEY_ID` and `AWS_SECRET_ACCESS_KEY` values with the values from the previous step.

- 5.2. List the content of the bucket by using the `s3cmd` command to validate the configuration.

```
[student@workstation backup-oadp]$ s3cmd ls
```

**Note**

Because the bucket is empty at this stage, the command returns an empty line. The command returns an error message if the configuration is incorrect.

- 6. Prepare the `DataProtectionApplication` configuration in a new `dpa-backup.yaml` YAML file.

Enable the `csi` plug-in and the Data Mover feature for the CSI snapshot.

Use the S3 bucket information from the previous steps to configure the default backup storage location.

**Note**

You can find an example of the resource definition in the `~/D0380/labs/backup-oadp/dpa-backup.yaml` file.

```
apiVersion: oadp.openshift.io/v1alpha1
kind: DataProtectionApplication
metadata:
  name: oadp-backup
  namespace: openshift-adp
spec:
  configuration:
    nodeAgent:
      enable: true
      uploaderType: kopia
    velero:
      defaultPlugins:
```

```

    - aws
    - openshift
    - csi
  defaultSnapshotMoveData: true
  backupLocations:
    - velero:
        config:
          profile: "default"
          region: "us-east-1"
          s3Url: https://s3.openshift-storage.svc
          s3ForcePathStyle: "true" ①
          insecureSkipTLSVerify: "true"
        provider: aws
        default: true
        credential:
          key: cloud
          name: cloud-credentials ②
      objectStorage:
        bucket: backup-9e2...20b ③
        prefix: oadp
        caCert: LLS0S0...LS0tLS0K ④

```

- ① The s3ForcePathStyle attribute must be set to true when using ODF.
 - ② Secret with the S3 credentials. You create the `cloud-credentials` secret in a following step.
 - ③ Use the bucket name from the previous step.
 - ④ Use the service CA certificate in Base64 format from the previous step.
- ▶ 7. Create the `cloud-credentials` secret in the `openshift-adp` namespace with the `backup` object bucket credentials.

- 7.1. Create a `cloud-credentials` file with the object bucket credentials. You can use the configuration example in the `~/D0380/labs/backup-oadp/cloud-credentials` file.

```
[default]
aws_access_key_id=AWS_ACCESS_KEY_ID
aws_secret_access_key=AWS_SECRET_ACCESS_KEY
```



Note

Replace the `AWS_ACCESS_KEY_ID` and `AWS_SECRET_ACCESS_KEY` values with the values from the previous step.

- 7.2. Create the `cloud-credentials` secret with the `cloud-credentials` file content.

Chapter 2 | Backup, Restore, and Migration of Applications with OADP

```
[student@workstation backup-oadp]$ oc create secret generic \
  cloud-credentials \
  -n openshift-adp \
  --from-file cloud=cloud-credentials
secret/cloud-credentials created
```

- 8. Apply the OADP configuration by using the `dpa-backup.yaml` YAML file from an earlier step.

```
[student@workstation backup-oadp]$ oc apply -f dpa-backup.yaml
dataprotectionapplication.openshift.io/oadp-backup created
```

- 9. Verify that both the `velero` deployment object and the `kopia` node-agent daemon set are created in the `openshift-adp` namespace and in the Ready state.

```
[student@workstation backup-oadp]$ oc get deploy,daemonset -n openshift-adp
NAME                                     READY   UP-TO-DATE ...
deployment.apps/openshift-adp-controller-manager  1/1     1      ...
deployment.apps/velero                      1/1     1      ...

NAME          DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE
daemonset.apps/node-agent    3         3         3        3           3
```

- 10. Verify that the backup storage location object is created and in the Available phase.

```
[student@workstation backup-oadp]$ oc get BackupStorageLocation
NAME      PHASE      LAST VALIDATED   AGE   DEFAULT
oadp-backup-1  Available   1m            10m   true
```

**Note**

It can take a minute for the backup storage location resource to enter the Available phase. If the resource is stuck in the Unavailable phase, then use the `oc describe BackupStorageLocation` command to get the status and error messages that relate to the configuration.

- 11. Configure all available volume snapshot classes for OADP.

- 11.1. List all available volume snapshot classes.

```
[student@workstation backup-oadp]$ oc get volumesnapshotclass
NAME
ocs-external-storagecluster-cephfsplugin-snapclass ...
ocs-external-storagecluster-rbdplugin-snapclass ...
```

- 11.2. For each volume snapshot class, set the `velero.io/csi-volumesnapshot-class` label to `true`.

```
[student@workstation backup-oadp]$ oc label volumesnapshotclass \
  velero.io/csi-volumesnapshot-class="true" --all
.../ocs-external-storagecluster-cephfsplugin-snapclass labeled
.../ocs-external-storagecluster-rbdplugin-snapclass labeled
```

► 12. Back up the **production** project to validate the OADP configuration.

- 12.1. Modify the resource definition in the `~/D0380/labs/backup-oadp/backup.yml` file to match the following specification.

```
apiVersion: velero.io/v1
kind: Backup
metadata:
  name: backup-production
  namespace: openshift-oadp
spec:
  includedNamespaces:
    - production
```

- 12.2. Create the backup resource.

```
[student@workstation backup-oadp]$ oc apply -f backup.yml
backup.velero.io/backup-production created
```

► 13. Wait a couple of minutes and check the backup completion status.

- 13.1. Verify that the backup object is in the **Completed** phase.

```
[student@workstation backup-oadp]$ oc describe backup backup-production
...output omitted...
Status:
  Backup Item Operations Attempted: 1
  Backup Item Operations Completed: 1
  Completion Timestamp:          2023-09-07T16:29:25Z
  Expiration:                   2023-10-07T16:28:10Z
  Format Version:               1.1.0
  Phase:                        Completed
  Progress:
    Items Backed Up: 42
    Total Items:     42
  Start Timestamp:              2023-09-07T16:28:10Z
  Version:                     1
  Events:                      <none>
```

- 13.2. Use the `s3cmd` command to review the content of the S3 storage.

```
[student@workstation backup-oadp]$ s3cmd la -r
2024-05-31 14:20      2597  s3://backup-78...1a/docker/registry/v2/.../data  ①
...output omitted...
2024-05-31 14:20      77529  s3://backup-78...1a/oadp/backups/...tar.gz  ②
...output omitted...
2024-05-31 14:20     286923  s3://backup-78...1a/oadp/kopia/...  ③
...output omitted...
```

- ① The /docker/registry path contains the container images.
- ② The /oadp/backups path contains the Kubernetes resources and the logs of the backup.
- ③ The /oadp/kopia path is the Kopia backup repository that contains backups of all data volumes in a project.

► 14. Clean up the resources.

14.1. Change to the student HOME directory.

```
[student@workstation backup-oadp]$ cd
[student@workstation ~]$
```

14.2. Use the `velero` command to remove the backup and all associated resources.

```
[student@workstation ~]$ oc exec deployment/velero -c velero -it -- \
./velero delete backup backup-production
Are you sure you want to continue (Y/N)? y
Request to delete backup "backup-production" submitted successfully.
The backup will be fully deleted after all associated data (disk snapshots, backup
files, restores) are removed.
```

14.3. Wait a couple of minutes and verify that the backup object is deleted.

```
[student@workstation ~]$ oc get backup
No resources found in openshift-adp namespace.
```

14.4. Use the `s3cmd` command to clean up the object storage bucket. Use the bucket name from previous steps.

```
[student@workstation ~]$ s3cmd rm -r --force \
s3://backup-7d9d5169-6c8d-4b40-bbba-931ddaebe6f4c/
...output omitted...
```

14.5. Delete the backuprepository resource in the openshift-adp namespace.

```
[student@workstation ~]$ oc get backuprepository -n openshift-adp
NAME                      AGE      REPOSITORY TYPE
production-oadp-backup-1-kopia-qfs64  7m59s    kopia

[student@workstation ~]$ oc delete backuprepository \
  production-oadp-backup-1-kopia-qfs64 -n openshift-adp
backuprepository.velero.io "production-oadp-backup-1-kopia-qfs64" deleted
```

Finish

On the **workstation** machine, use the **lab** command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish backup-oadp
```

Backup and Restore with OADP

Objectives

- Configure one-time and scheduled backups with OADP and restore from them.

OADP Custom Resources

Backup and restore are initiated by creating the corresponding Kubernetes resources in the `openshift-adp` namespace. Administrative access to the `openshift-adp` namespace, or to the `cluster-admin` role, is required to create those resources.

OADP provides the following custom resources for backup and restore:

Backup

The `backup` resource initiates a single backup attempt. This resource defines the namespaces and resources to include in the backup, and can also include a list of commands to run before or after the backup.

The backup resource definition and the backup information, such as backup logs and the list of included backup resources, are stored in the object storage with the backup.

OADP synchronizes backup definitions between the object storage and the OpenShift cluster to enable restoring backups to a different cluster with the same backup storage location.

If a backup resource exists in the OpenShift cluster but is deleted from the object storage, then OADP deletes the Kubernetes resource. Conversely, if a backup exists in the object storage, but not in OpenShift, then OADP creates the matching backup resource in the cluster.



Note

Only backups with a `Completed` state are synchronized. The object storage synchronization does not automatically create or remove backup resources with a `Failed` or `PartiallyFailed` state.

Restore

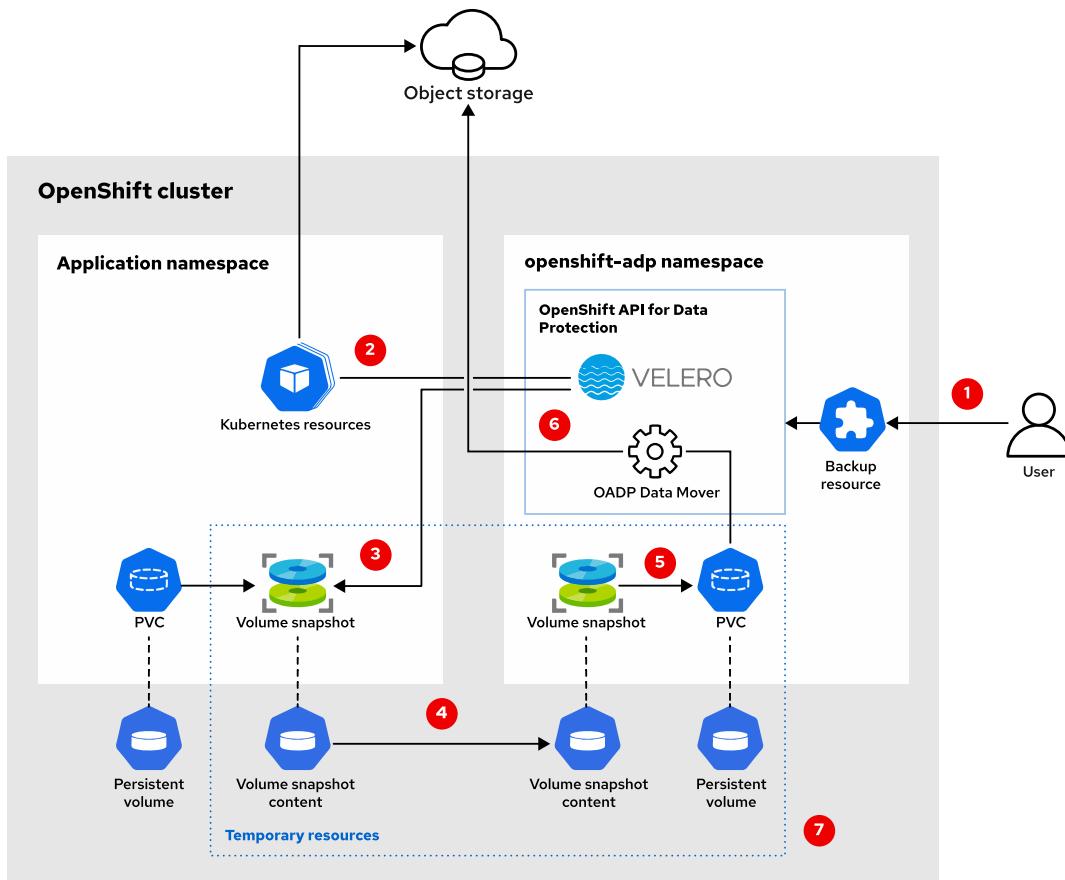
The `restore` resource starts restoring an existing backup resource. The restore result, the list of restored resources, and the restore logs are stored in the object storage.

Schedule

The `schedule` resource starts a backup on a given schedule that is written in Cron format. A `schedule` resource is similar to a cron job resource. A `schedule` defines a backup template to create a backup resource at a recurring interval.

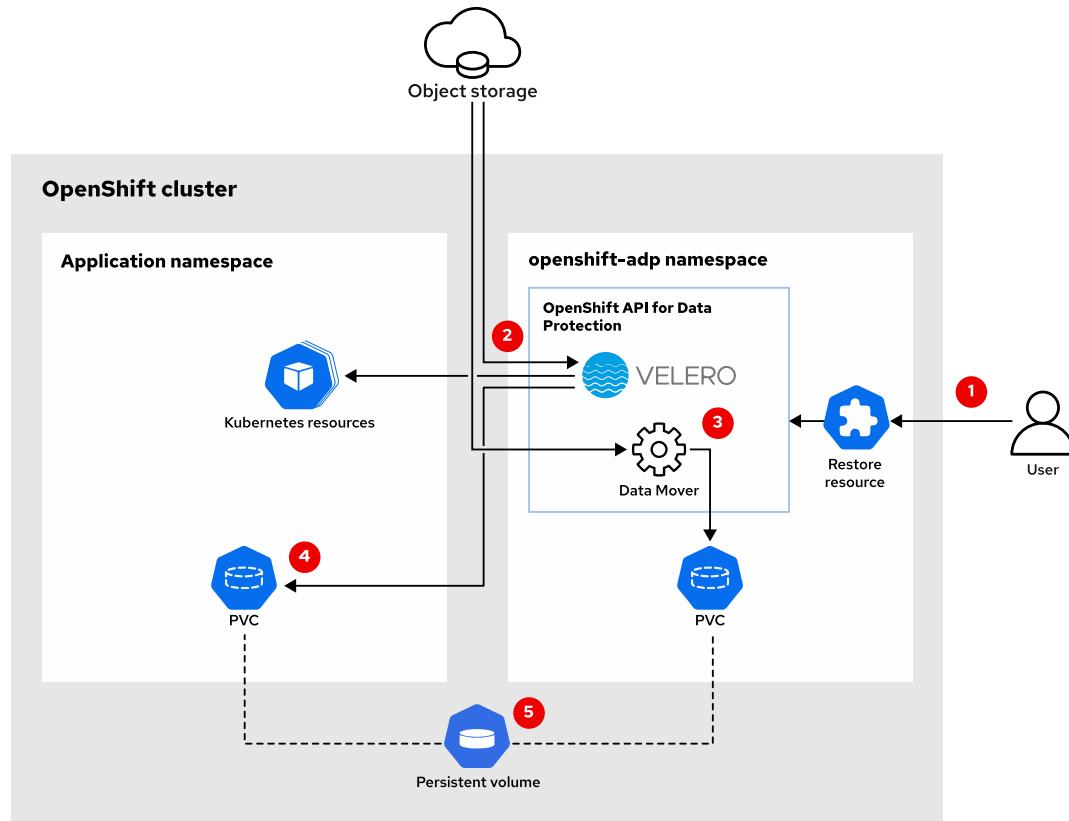
Backing up and Restoring an OpenShift Application

OADP backs up OpenShift applications by using the following process:



1. The administrator creates a backup resource in the **openshift-adp namespace**, which triggers the backup process.
2. Velero exports all Kubernetes resources from the application namespace to the backup storage location.
3. The Velero CSI plug-in creates a CSI snapshot of the application volume.
4. Data Mover clones the **volumeSnapshotContent** and **volumeSnapshot** resources to the **openshift-adp namespace**.
5. Data Mover creates a PVC from the volume snapshot.
6. Data Mover uses Kopia to transfer the volume data to the backup storage location.
7. After the backup completes, Data Mover deletes all volume snapshots and the PVCs that were created during the backup process.

For restoring backups, OADP uses the following process:



- ① The administrator creates a restore resource in the **openshift-adp** namespace, which triggers the restore process.
- ② Velero imports the Kubernetes resources from the backup storage location to the application namespace.
- ③ Data Mover creates a temporary PVC in the **openshift-adp** namespace, and transfers the application data from the backup storage location to the new volume.
- ④ Velero creates the PVC in the application namespace with the same PV as the one that Data Mover is restoring. The PVC stays in the pending state until Data Mover completes the restore.
- ⑤ After the transfer is complete, Data Mover releases the PV from the temporary PVC in the **openshift-adp** namespace, and binds that PV to the final PVC in the application namespace.

Backing up and Restoring a Stateless Application

You can start backing up a namespace by using the following backup definition:

```
apiVersion: velero.io/v1
kind: Backup
metadata:
  name: my-app-backup ①
  namespace: openshift-adp
spec:
```

```

includedNamespaces: ②
- my-app-project
ttl: 720h0m0s ③
labelSelector:
  matchLabels: ④
    app: my-app
includedResources: ⑤
- deployments
- configmaps
- secrets
- services
- routes

```

- ① Name of the backup
- ② List of namespaces to back up
- ③ Amount of time before the automatic deletion of the backup. If the `ttl` value is not specified, then the default is 30 days.
- ④ Labels that are required for the resources to be included in the backup
- ⑤ List of resource types to back up

The following example shows a stateless application, followed by the needed definition to back it up in the `website` namespace. The application is a static website that is built with Hugo and that uses a Source-to-Image (S2I) container image. This application uses the following resources that are created with the `app=hugo` label:

- Image Stream
- Build Configuration
- Deployment
- Service
- Route

You can back up this application with the following backup definition:

```

apiVersion: velero.io/v1
kind: Backup
metadata:
  name: website
  namespace: openshift-adp
spec:
  includedNamespaces:
  - website
  labelSelector:
    matchLabels:
      app: hugo
  includedResources:
  - imagestreams
  - buildconfigs
  - deployments
  - services
  - routes

```

You can review the configuration and the backup status with the `oc describe` command:

```
[user@host ~]$ oc -n openshift-adp describe backup/website
Name:          website
Namespace:     openshift-adp
...output omitted...
Spec:
  Csi Snapshot Timeout:      10m0s
  Default Volumes To Fs Backup:  false
  Included Namespaces:
    website
  Included Resources:
    imagestreams
    buildconfigs
    deployments
    services
    routes
  Item Operation Timeout:   1h0m0s
  Label Selector:
    Match Labels:
      App:           hugo
  Snapshot Move Data: true
  Storage Location: oadp-config-1
  Ttl:             720h0m0s
Status:
  Backup Item Operations Attempted: 1
  Backup Item Operations Completed: 1
  Completion Timestamp:            2023-12-11T14:01:42Z ①
  Expiration:                     2024-01-10T14:01:27Z ②
  Format Version:                 1.1.0
  Phase:                          Completed ③
  Progress:                      ④
    Items Backed Up:        7
    Total Items:           7
  Start Timestamp:            2023-12-11T14:01:27Z ⑤
  Version:                     1
```

- ① Backup completion time.
- ② Backup expiration date, which is set with the `ttl` setting. The backup is automatically removed from both the OpenShift cluster and the object storage after this date.
- ③ Status of the backup.
- ④ The backup progress, with the number of backed up items so far and the total items to back up.
- ⑤ Backup start time.

The backup resource can have the following statuses during its lifetime:

New

Initial status when a backup is created. If the backup definition is incorrect, then the backup is aborted and its status changes to `FailedValidation`. You can review the `validationErrors` status field for more information about the error.

If the backup definition is valid, then the status changes to **InProgress**.

InProgress

Status when the backup is in progress. During this phase, OADP backs up the resources that are specified in the backup definition and runs the backup hooks.

When OADP uses additional plug-ins, such as OADP Data Mover, the backup enters the **WaitingForPluginOperations** status. After the plug-in processes are complete, the backup enters the **Finalizing** status where OADP saves all the remaining backup items, such as backup logs and metadata, to the object storage.

If backing up some resources fails, then the status changes in turn to **WaitingForPluginOperationsPartiallyFailed**, **FinalizingPartiallyFailed**, and then **PartiallyFailed**.

PartiallyFailed and Failed

The final status of a backup, with some missing resources because of a backup failure. A partially failed backup can still restore a project, but incompletely. A backup with the **Failed** status cannot be restored.

You can review the **failureReason** status field for more information about the error.

Completed

Final status of a successfully completed backup. All the data is in the object storage and the backup is ready to use in a restore.

You can initiate a restore by using the following restore definition:

```
apiVersion: velero.io/v1
kind: Restore
metadata:
  name: restore-name ①
  namespace: openshift-adp
spec:
  backupName: backup-name ②
```

- ① Name of the restore resource
- ② Name of the backup resource to restore

OADP restores only the resources that are not already in the destination namespace. The destination namespace is the same as in the backup, unless you use the **namespaceMapping** field to specify a different namespace.

For example, you can use the following restore definition to restore the **website** project from the previous example to a new **website-stage** project:

```
apiVersion: velero.io/v1
kind: Restore
metadata:
  name: website-stage
  namespace: openshift-adp
spec:
  backupName: website
  namespaceMapping:
    website: website-stage
```

After creating the restore resource, you can review the configuration and monitor the progress of the restore with the `oc describe` command:

```
[user@host ~]$ oc -n openshift-adp describe restore/website-stage
Name:          website-stage
Namespace:     openshift-adp
...output omitted...
Spec:
  Backup Name:  website
  Excluded Resources: ①
    nodes
    events
    events.events.k8s.io
    backups.velero.io
    restores.velero.io
    resticrepositories.velero.io
    csinodes.storage.k8s.io
    volumeattachments.storage.k8s.io
    backuprepositories.velero.io
  Item Operation Timeout: 1h0m0s
  Namespace Mapping:
    Website:  website-stage
Status:
  Completion Timestamp:  2023-12-12T09:36:23Z ②
  Phase:                Completed ③
  Progress:              ④
    Items Restored:      7
    Total Items:         7
  Restore Item Operations Attempted: 1
  Restore Item Operations Completed: 1
  Start Timestamp:       2023-12-12T09:36:14Z ⑤
```

- ① OADP automatically excludes some resources from the restore, such as Kubernetes events and OADP custom resources. You can modify this list with the `excludedResources` field in the restore definition.
- ② Restore completion time.
- ③ Status of the restore.
- ④ The restore progress, with the number of restored items so far and the total items to restore.
- ⑤ Restore start time.

The possible statuses of the restore are the same as for the backup resource. A successful restore goes through the following statuses during its lifetime:

- New
- InProgress
- WaitingForPluginOperations
- Completed

During the restore, OADP adds two labels to the imported resources: A `velero.io/restore-name` label with the restore name, and a `velero.io/backup-name` with the backup name that the resource came from. You can use those labels to identify the resources that were restored from a specific backup or restore name.

Chapter 2 | Backup, Restore, and Migration of Applications with OADP

For example, you can list the resources with the `velero.io/restore-name=website-stage` label that are created by the restore from the previous example:

```
[user@host ~]$ oc -n website-stage get all \
  -l velero.io/restore-name=website-stage
NAME          TYPE      CLUSTER-IP      EXTERNAL-IP    PORT(S)        AGE
service/hugo   ClusterIP  172.30.61.176  <none>       8080/TCP,8443/TCP  65m

NAME          READY     UP-TO-DATE    AVAILABLE    AGE
deployment.apps/hugo  1/1       1           1           65m

NAME          TYPE      FROM        LATEST
buildconfig.build.openshift.io/hugo  Source      Git         0

NAME          HOST/PORT
route.route.openshift.io/hugo     hugo-website-stage.apps.ocp4.example.com  ...
...
```

OADP stores in the object storage detailed information about each backup and restore attempt, such as logs, the list of resources that are backed up or restored, and a summary of the errors or warnings that occur during the backup and restore process. This information is not stored in the OADP Kubernetes custom resources, and only limited information such as the runtime status of the backup and restore is available.

The following excerpt shows the object storage layout after the backup and restore from the previous examples:



- ① The `/docker/registry/v2` path contains a Docker registry with the container images that are included in all backups.
- ② The `/oadp/backups` path contains all the information about each backup, including the backed-up Kubernetes resources and the backup logs. Each subdirectory is unique to a single backup attempt and relates to the matching backup resource in the cluster.
- ③ The `/oadp/restores` path contains all the information about each restore attempt, including the backed-up Kubernetes resources and the restore logs. Each subdirectory is unique to a single restore attempt and relates to the matching restore resource in the cluster.

- ③ The /oadp/restores path contains all the information about each restore, including the restore logs. Each subdirectory is unique to a single restore attempt and relates to the matching restore resource in the cluster.

Introducing the Velero Tool

OADP provides the `velero` command-line tool that can retrieve backup and restore information from both the object storage and the OpenShift cluster.

The Velero CLI tool is available from the `velero` deployment in the `openshift-adp` namespace. You can define an alias to access the `velero` binary by using the following command:

```
[user@host ~]$ alias velero='\\
oc -n openshift-adp exec deployment/velero -c velero -it -- ./velero'
```



Note

The remainder of this section uses the `velero` alias on the command line to refer to the Velero CLI tool.

The `velero` command can use the same syntax as the `kubectl` or `oc` commands, but is limited to OADP Kubernetes custom resources such as the `backup`, `restore`, and `schedule` resources.

The `oc get` command displays limited information about OADP resources. However, the `velero get` command provides the same runtime information from the Kubernetes resources as the `oc describe` command:

```
[user@host ~]$ oc -n openshift-adp get backup,restore
NAME                           AGE
backup.velero.io/website        24h

NAME                           AGE
restore.velero.io/website-stage 4h27m

[user@host ~]$ velero get backup
NAME      STATUS     ERRORS  WARNINGS  CREATED   EXPIRES  STORAGE LOCATION  SELECTOR
website  Completed  0        0          ...       29d      oadp-config-1  app=hugo

[user@host ~]$ velero get restore
NAME      BACKUP  STATUS     STARTED  COMPLETED  ERRORS  WARNINGS  ...
website-stage website  Completed  ...      ...      0        0          ...
```

The Velero tool provides a `describe` command that is similar to the `oc describe` command, and adds a `--details` option that retrieves additional information about the resource from the object storage:

```
[user@host ~]$ velero describe backup website --details
Name:      website
Namespace:  openshift-adp
...output omitted...

Phase:  Completed ❶
```

Chapter 2 | Backup, Restore, and Migration of Applications with OADP

```
Namespaces: ②
  Included: website
  Excluded: <none>

Resources:
  Included:      imagestreams, buildconfigs, deployments, services, routes
  Excluded:      <none>
  Cluster-scoped: auto

Label selector: app=hugo

...output omitted...

Started: 2023-12-11 14:01:27 +0000 UTC
Completed: 2023-12-11 14:01:42 +0000 UTC

Expiration: 2024-01-10 14:01:27 +0000 UTC

Total items to be backed up: 7
Items backed up: 7

Resource List: ③
  apps/v1/Deployment:
    - website/hugo
  build.openshift.io/v1/BuildConfig:
    - website/hugo
  image.openshift.io/v1/ImageStream:
    - website/hugo
    - website/nginx-122
  route.openshift.io/v1/Route:
    - website/hugo
  v1/Service:
    - website/hugo

Velero-Native Snapshots: <none included>
```

- ① Status of the backup.
- ② Definition of the namespaces and resources to back up.
- ③ Resources that are included in the backup. The `--details` option adds this information, which comes from the object storage.

If the backup or restore resource has errors or warnings, then the `--details` option adds them to the command output.

In the following example, the `website` project is backed up without filtering by resource type. OADP backs up all resources with the `app=hugo` label in the `website` namespace.

```
apiVersion: velero.io/v1
kind: Backup
metadata:
  name: website-label
  namespace: openshift-adp
spec:
```

```

includedNamespaces:
- website
labelSelector:
  matchLabels:
    app: hugo

```

You can use the `velero create restore` command to restore the previous backup in a new `website-dev` namespace, as follows:

```

[user@host ~]$ velero create restore website-dev \
--from-backup=website-label --namespace-mappings=website:website-dev
Restore request "website-dev" submitted successfully.
Run 'velero restore describe website-dev' or 'velero restore logs website-dev' for
more details.

```

The restore completes with one warning about an existing resource in the target namespace:

```

[user@host ~]$ velero describe restore website-dev --details
Name:          website-dev
Namespace:     openshift-adp
...output omitted...

Phase:           Completed
Total items to be restored: 11
Items restored:   11

Started:        2023-12-12 10:21:21 +0000 UTC
Completed:      2023-12-12 10:21:27 +0000 UTC

Warnings:
  Velero:      <none>
  Cluster:     <none>
  Namespaces:
    website-dev: could not restore, Endpoints "hugo" already exists. Warning: the
                  in-cluster version is different than the backed-up version.

Backup:         website-label
...output omitted...

Resource List:
  apps/v1/Deployment:
    - website-dev/hugo(created)
  build.openshift.io/v1/Build:
    - website-dev/hugo-1(skipped)
  build.openshift.io/v1/BuildConfig:
    - website-dev/hugo(created)
  discovery.k8s.io/v1/EndpointSlice:
    - website-dev/hugo-5hq2x(created)
  image.openshift.io/v1/ImageStream:
    - website-dev/hugo(skipped)
  image.openshift.io/v1/ImageStreamTag:
    - website-dev/hugo:latest(skipped)

```

```
image.openshift.io/v1/ImageTag:  
  - website-dev/hugo:latest(skipped)  
route.openshift.io/v1/Route:  
  - website-dev/hugo(created)  
v1/Endpoints:  
  - website-dev/hugo(failed)  
v1/Namespace:  
  - website-dev(created)  
v1/Service:  
  - website-dev/hugo(created)
```

In this example, OADP tries unsuccessfully to restore the hugo endpoint. The hugo service in the backup automatically creates this resource.

Avoid including resources that other resources manage, such as builds, endpoints, or replica sets. It is unnecessary and can cause issues during the restore. You must filter to include in the backup only those resources that your application requires for a successful deployment.

Backing up a Stateful Application with Backup Hooks

OADP creates crash-consistent backups of your application by using volume snapshots. To back up persistent volumes, you must include the `persistentvolumeclaims` and `persistentvolumes` resource types in the backup definition.

OpenShift assigns to each namespace a unique UID and GID that the application pod uses to write data to persistent volumes. If you restore an application to a new namespace, then OpenShift assigns a new set of UIDs and GIDs that prevent the application from accessing its data.



Note

For more details about user and group ID assignments, refer to the *DO180: Red Hat OpenShift Administration I: Operating a Production Cluster* training course.

So that OADP can restore the UID and GID, include the `namespace` resource type in the backup definition. If you use a label selector in the backup definition, then you must add the corresponding label to the namespace. Alternatively, you can use the `kubernetes.io/metadata.name` label that Kubernetes automatically sets on all namespaces.

To improve the consistency of a backup, you can use backup hooks to specify a list of commands to execute in the application pod before and after the backup is created. You can then use those hooks to quiesce the application and perform an application-consistent backup.

Some applications require additional steps when using volume snapshots to create a usable backup. For more details about the backup procedure, refer to the application documentation. To configure backup hooks, you must specify the target pod (by using labels), the container name, and the commands to run on that container.

You can configure the following hook types according to your needs:

Pre backup hooks

A pre backup hook is executed before any other backup action on the pod. If the command fails, then the backup stops immediately with the `Failed` status.

Chapter 2 | Backup, Restore, and Migration of Applications with OADP

As an example, you can use this type of hook to quiesce and prepare the application for backup.

Post backup hooks

A post backup hook is executed after the backup of the pod and its attached volumes. If the command fails, then the backup stops immediately with the `PartiallyFailed` status.

As an example, you can use this type of hook to resume or unlock the application after the backup is complete.

Init restore hooks

An init restore hook is executed after the pod and its attached volumes are restored, but before any container on that pod starts. The `init` restore hook defines one or more init containers that follow the same specification as the init container in a pod definition.

OADP does not monitor the status of the init container. Therefore, if the command fails, then the restore continues without any error or warning, but the application pod is in error with the `Init:Error` status.

As an example, you can use this type of hook to restore a database that the application requires and that is external to the OpenShift cluster.

Post restore hooks

A post restore hook is executed when the application pod is restored and running. If the command fails, then the error is logged and the restore continues.

As an example, you can use this type of hook to run an integrity check on the restored database.

The following example is a backup definition for a MongoDB database. It uses backup hooks to flush all pending writes to the disk and to lock the database to prevent any writes during the backup.

```
apiVersion: velero.io/v1
kind: Backup
metadata:
  name: mongodb
  namespace: openshift-adp
spec:
  includedNamespaces:
    - mongodb
  orLabelSelectors: ①
    - matchLabels:
        app: mongodb
    - matchLabels:
        kubernetes.io/metadata.name: mongodb
  includedResources: ②
    - deployments
    - services
    - secret
    - pvc
    - pv
    - pods
    - namespace
  hooks:
    resources:
      - name: mongodb-lock
```

```

labelSelector: ③
  matchLabels:
    app: mongodb
pre: ④
  - exec:
    container: mongodb
    command:
      - /usr/bin/mongosh
      - --eval
      - db.fsyncLock();
post: ⑤
  - exec:
    container: mongodb
    command:
      - /usr/bin/mongosh
      - --eval
      - db.fsyncUnlock();

```

- ① Resources with the `app: mongodb` or `kubernetes.io/metadata.name: mongodb` label are included in the backup.
- ② The `pvc` and `pv` resource types must be specified in the `includedResources` key to back up the application volume. The `pods` resource type must also be specified for the backup hooks to be executed. The `namespace` resource type must be specified to preserve the UID and GID that are used in the application volume.
- ③ The hook runs on all pods that match the label selector.
- ④ The `pre` backup hook executes the `db.fsyncLock()` MongoDB command in the `mongodb` container to lock the database before the volume snapshot.
- ⑤ The `post` backup hook executes the `db.fsyncUnlock()` MongoDB command to unlock the database after the backup is completed.



Important

Because backup and restore hooks are executed only on pods that are included in the backup, you must include the `pod` resource type in the backup.

If the hooks are configured to run on a pod that is not included in the backup, then the hooks are ignored without any error or warning.

The following restore resource definition restores the MongoDB backup from the previous example:

```

apiVersion: velero.io/v1
kind: Restore
metadata:
  name: mongodb
  namespace: openshift-adp
spec:
  backupName: mongodb
  hooks:
    resources:

```

```

- name: mongodb-unlock
  labelSelector:
    matchLabels:
      app: mongodb
  postHooks:
    - init: ①
      initContainers:
        - name: remove-lock
          image: mongo/mongodb-community-server:7.0-ubi8
          volumeMounts:
            - name: mongodb-data
              mountPath: /data/db
          command:
            - /usr/bin/rm
            - /data/db/mongod.lock

```

- ① The `init` restore hook removes the database lock from the backup before the database starts.

OADP Data Mover stores the volume backup in the object storage inside a Kopia unified repository.

A *Kopia unified repository* is a backup repository that Kopia manages for storing volume backups. Kopia encrypts, deduplicates, and compresses the data in the backup repository. OADP uses a dedicated backup repository for each namespace to store all volume backups for that namespace. OADP uses this unified repository to store backups from both volume snapshots and file-system backups.

When the first backup of a namespace is created, OADP initializes a Kopia repository in the object storage, and creates a matching `BackupRepository` resource in the `openshift-adp` namespace.

OADP uses a `BackupRepository` custom resource to store information about the backup repository such as the repository encryption keys and object storage information.

The following excerpt shows the object storage layout after the previous backup and restore examples:

```

s3
└── oadp
    ├── backups
    │   └── mongodb
    │       ├── mongodb-logs.gz
    │       └── ...
    ├── kopia ①
    │   └── mongodb
    │       ├── kopia.repository
    │       └── ...
    └── restores
        └── mongodb
            ├── restore-mongodb-logs.gz
            └── ...

```

- ❶ The /oadp/kopia path contains the backup repositories with data from volume backups. Each subdirectory is a unique Kopia repository for a single namespace in the cluster and contains all backups from all volumes in that namespace.

Scheduling a Recurring Backup

You can back up an application at a recurring interval by using the schedule resource. A schedule resource is similar to a cron job resource. A schedule resource requires a schedule in Cron format and a definition template of the backup resource to create, at the specified time.

The following example is a schedule definition of a daily backup for the previous static website example:

```
apiVersion: velero.io/v1
kind: Schedule
metadata:
  name: website-daily
  namespace: openshift-adp
  labels: ❶
    app: hugo
spec:
  schedule: "0 7 * * *" ❷
  paused: false ❸
  template: ❹
    includedNamespaces:
    - website
    labelSelector:
      matchLabels:
        app: hugo
    includedResources:
    - imagestreams
    - buildconfigs
    - deployments
    - services
    - routes
    ttl: 720h0m0s
```

- ❶ The labels that you set on the schedule resource are automatically copied to the backup resources that the schedule creates.
- ❷ Specifies the schedule for the job in Cron format.
- ❸ You can disable the schedule by setting the paused parameter to true.
- ❹ Sets the backup definition template by using the same settings as in a backup resource.

See the references section for more information about the schedule API definition.

To get detailed information about a schedule, use the `velero get schedule` command:

```
[user@host ~]$ velero get schedule
NAME      STATUS  CREATED  SCHEDULE  BACKUP TTL  LAST BACKUP  SELECTOR PAUSED
website-daily  Enabled ...  0 7 * * *  720h0m0s  ...  app=hugo  false
```

The status of an activated schedule is Enabled. If the schedule is disabled with the paused parameter, then its status is New. The LAST BACKUP column shows the time of the latest backup that the schedule created.

With the Velero tool, you can create a one-time backup by using the same definition as in an existing schedule. To back up an application on demand, you can create a disabled schedule as a backup template. You can then start a new backup with this schedule when you need it.

In this example, you are creating a pre-upgrade-1.1 backup from the website-daily schedule:

```
[user@host ~]$ velero create backup pre-upgrade-1.1 \
--from-schedule website-daily
Creating backup from schedule, all other filters are ignored.
Backup request "pre-upgrade-1.1" submitted successfully.
...output omitted...
```

Backups that are created from a schedule inherit the schedule's labels. In addition, the velero.io/schedule-name label is set on the backup resources with the schedule name. You can use those labels to identify the schedules and backups for your application on the openshift-adp namespace:

```
[user@host ~]$ velero get backup -l app=hugo
NAME           STATUS    ERRORS   WARNINGS ...
post-upgrade-1.1   Finalizing   0        0        ...
pre-upgrade-1.1    Completed    0        0        ...
website-daily-20231215100856  Completed    0        0        ...
website-daily-20231215091728  Completed    0        0        ...
```

Because the openshift-adp namespace contains the backup and restore resources for all applications that are running on the cluster, it is important to use labels to identify OADP resources that relate to your application.

Cleaning Backups

Because OADP synchronizes backup definitions with the object storage, OADP automatically recreates any backup that you delete with the oc command. To permanently delete a backup, you must delete it from the object storage.

Use the velero command to delete backup and restore information from the object storage, and all the associated resources from the cluster:

```
[user@host ~]$ velero delete backup backup-name
Are you sure you want to continue (Y/N)? y
Request to delete backup "backup-name" submitted successfully.
The backup will be fully deleted after all associated data (disk snapshots, backup
files, restores) are removed.
```

The backup status changes to Deleting, and OADP removes all restore resources that are attached to this backup from both the OpenShift cluster and the object storage. Then the backup itself is removed from the cluster and the object storage.

Chapter 2 | Backup, Restore, and Migration of Applications with OADP

You can instruct OADP to delete a backup automatically after a specified elapsed time by using the TTL (Time To Live) setting on the backup and schedule definition. By default, OADP deletes backups after 30 days. The minimum lifetime of a backup is 1 hour.



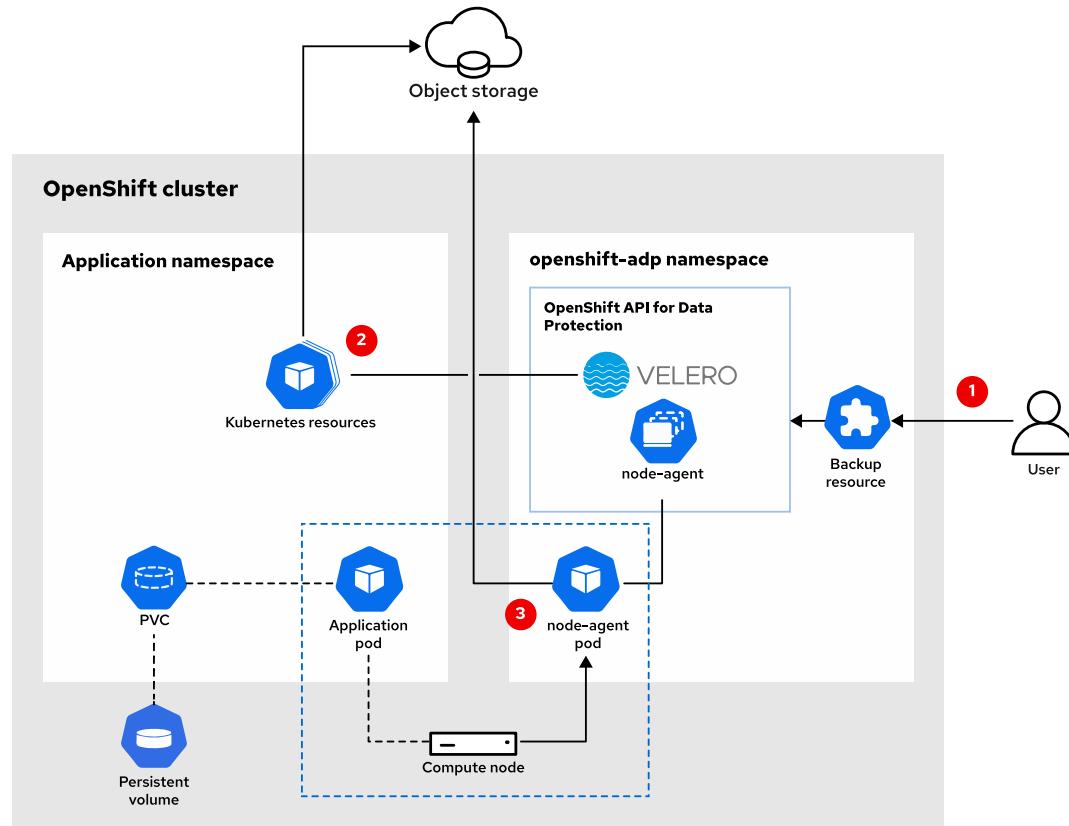
Important

OADP does not immediately delete data from the backup repository on the object storage. OADP occasionally runs repository maintenance, such as to delete data that backups no longer need. OADP might start deleting data from the object storage up to 24 hours after a backup is deleted.

Backing up Volumes with File System Backup

With the File System Backup (FSB) feature, you can back up volumes that are not compatible with snapshots.

OADP uses the following process to back up OpenShift applications with FSB:

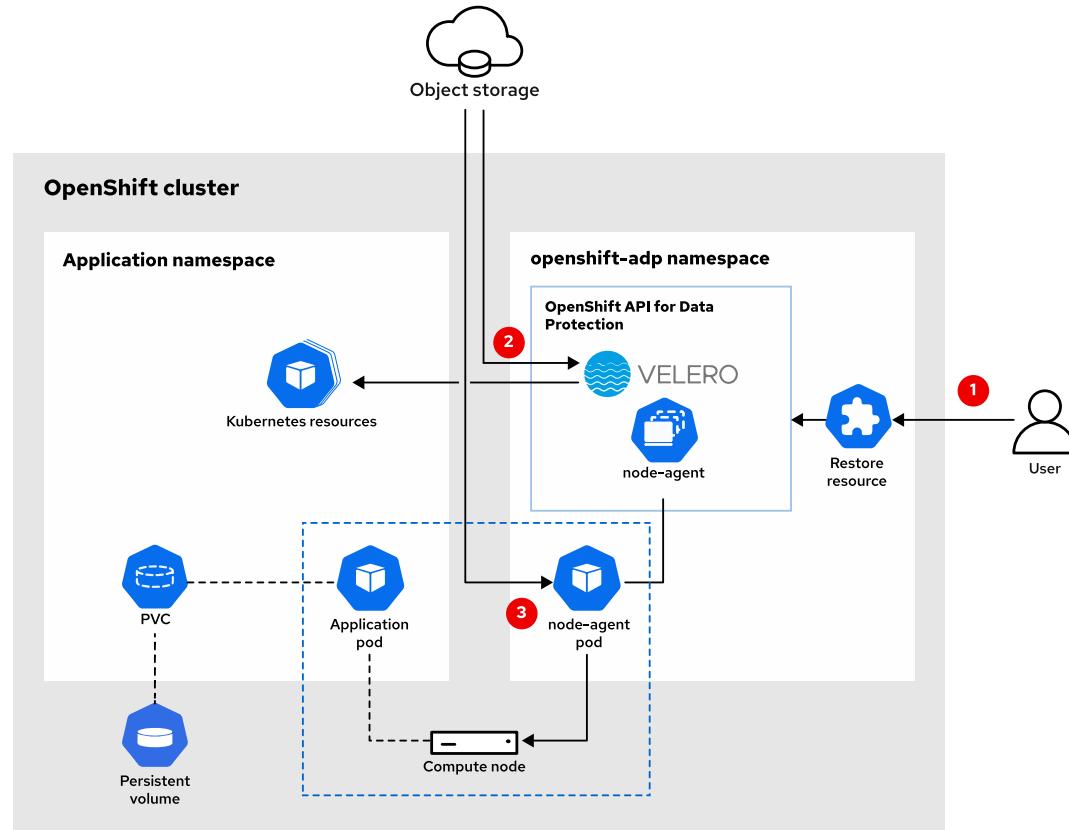


- ① The administrator creates a backup resource in the `openshift-adp` namespace that triggers the backup process.
- ② Velero exports all Kubernetes resources from the application namespace to the backup storage location.
- ③ The node-agent daemon set, which runs on the same node as the application pod, exports the volume data from the volume mount point on the cluster node to the backup storage location.

**Important**

OADP copies the data from the application file system while the application is still running. You must use backup hooks to ensure that the application data is not altered during the backup process.

The FSB restore process is as follows:



- ① The administrator creates a restore resource in the `openshift-adp` namespace, which triggers the restore process.
- ② Velero imports the Kubernetes resources from the backup storage location to the application namespace.
- ③ The `node-agent` daemon set, which runs on the same node as the application pod, imports the application data from the backup storage location to the mount point for the application volume on the cluster node.

You can instruct OADP to back up all volumes in a backup definition with FSB by using the `defaultVolumesToFsBackup` option:

```
apiVersion: velero.io/v1
kind: Backup
metadata:
  name: <backup-name>
  namespace: openshift-adp
spec:
  defaultVolumesToFsBackup: true
```

You can also annotate the application pod to specify which volumes to back up with FSB by using the `backup.velero.io/backup-volumes` annotation. If your application uses multiple volumes from different storage classes, then OADP backs up the volumes in the annotation with FSB, and all other volumes with snapshots.

In the following example, the application to back up is an Nginx web server that serves a website from an Amazon Elastic File System (AWS EFS). Because the AWS EFS CSI driver does not support volume snapshots, you must back up the volume with FSB.

To enable FSB for the `wwwdata` volume, you must set the `backup.velero.io/backup-volumes` annotation in the deployment, as follows:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: website-nginx
  namespace: org-website
spec:
  template:
    metadata:
      annotations:
        backup.velero.io/backup-volumes: wwwdata ①
    spec:
      containers:
        name: web
        image: registry.access.redhat.com/ubi9/nginx-120
        ...output omitted...
        volumeMounts:
          - mountPath: /opt/app-root/src
            name: wwwdata
            ...output omitted...
        volumes:
          - name: wwwdata
            persistentVolumeClaim:
              claimName: nginx-wwwdata
...output omitted...
```

- ① List of volumes to back up with FSB. You must use the same volume name that is defined in the pod definition.



Note

You can specify multiple volumes with the `backup.velero.io/backup-volumes` annotation, with a comma to separate each volume. For example, `backup.velero.io/backup-volumes: volume1,volume2,volume3`

**Important**

Only volumes that are compatible with volume snapshots are included in the backup by default. You must enable FSB to include volumes in your backup that are not compatible with volume snapshots.

OADP stores file-system backups in the same unified backup repository on the object storage that Data Mover uses for volume snapshots.

Troubleshooting Backups and Restores

If a backup or restore fails, you can get more information about the failure with the `velero` command:

```
[user@host ~]$ velero get backup mybackup
NAME      STATUS      ERRORS   WARNINGS   ...
mybackup  PartiallyFailed  1        0          ...
```

Use the `velero describe` command for more details about the errors and warnings:

```
[user@host ~]$ velero describe backup mybackup
Name:           mybackup
Namespace:      openshift-adp
Labels:         velero.io/storage-location=oadp-backup-1
Annotations:    velero.io/source-cluster-k8s-gitversion=v1.25.7+eab9cc9
                velero.io/source-cluster-k8s-major-version=1
                velero.io/source-cluster-k8s-minor-version=25

Phase:  PartiallyFailed (run `velero backup logs mybackup` for more information)

Errors:
  Velero:      <none>
  Cluster:     <none>
  Namespaces:
    database:  resource: /pods name: /mariadb-757c5bdc88-mrwhb error: /command
    terminated with exit code 1
    ...output omitted...
```

You can view the backup logs with the `velero backup logs` command:

```
[user@host ~]$ velero backup logs mybackup
...output omitted...
```

Because the logs are stored in the object storage, you can also download them with the `s3cmd` command, and view them with any text editor.

The logs are stored in the `s3://bucket-name/oadp/backups/backup-name/backup-name-logs.gz` path for backups and in the `s3://bucket-name/oadp/restores/restore-name/restore-name-logs.gz` path for restores.

You can use the following command to download the log file to the current directory with `s3cmd`:

```
[user@host ~]$ s3cmd get \
  s3://backup-bucket/oadp/backups/mybackup/mybackup-logs.gz
download: 's3://backup-bucket/oadp/backups/mybackup/mybackup-logs.gz' ...
  7454 of 7454   100% in    0s    96.14 KB/s  done
```

For a backup that uses hooks, you can search the log with the `hookPhase` keyword to review the status of the hooks. The standard and error outputs of the hook command are recorded in the `stdout` and `stderr` lines, respectively.

```
[user@host ~]$ velero backup logs mybackup | grep hookPhase
...
... msg="running exec hook" hookCommand="/bin/bash -c mariadb -u root -e \"set
global read_only=1;flush tables\"]" ... hookPhase=pre

...
... msg="stdout: " ... hookPhase=pre

...
... msg="stderr: ERROR 1045 (28000): Access denied for user
'root'@'localhost' (using password: NO)\n" ... hookPhase=pre

...
... msg="Error executing hook" error="command terminated with exit code 1" ...
hookPhase=pre
```

In the previous example, the backup hook was misconfigured and is missing the password to connect to the database.



References

For more information about backing up applications with OADP, refer to the *OADP Backing Up* section in the *OADP Application Backup and Restore* chapter in the Red Hat OpenShift Container Platform 4.14 *Backup and Restore* documentation at https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/backup_and_restore/index#oadp-backing-up

For more information about troubleshooting OADP, refer to the *Troubleshooting* section in the *OADP Application Backup and Restore* chapter in the Red Hat OpenShift Container Platform 4.14 *Backup and Restore* documentation at https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/backup_and_restore/index#troubleshooting

Velero API Definitions

<https://velero.io/docs/v1.12/api-types/>

► Guided Exercise

Backup and Restore with OADP

Perform a crash-consistent and an application-consistent backup of a database, and restore each backup to a different project.

Outcomes

- Back up and restore an OpenShift project.
- Schedule a backup of an OpenShift project.
- Remove previous backups.

Before You Begin

As the student user on the **workstation** machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start backup-restore
```

Instructions

Create a crash-consistent backup of the **database** project by using OADP and CSI snapshots, and restore it to a new **database-crash** project.

Schedule a weekly backup of the **database** project. Use hooks to quiesce the database and produce an application-consistent backup.

Trigger a backup from the schedule and restore the application-consistent backup to a new **database-backup** project.

Remove the backups and all related resources.

► 1. As the **admin** user, review the content of the **database** project.

1.1. Log in to your OpenShift cluster as the **admin** user.

```
[student@workstation ~]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
Login successful.

...output omitted...
```

1.2. Change to the **database** project.

```
[student@workstation ~]$ oc project database
Now using project "database" on server "https://api.ocp4.example.com:6443".
```

1.3. List the resources with the `app=mariadb` label in the **database** project.

```
[student@workstation ~]$ oc get \
  pvc,svc,deployment,secret,configmap -l app=mariadb
NAME          ...   STORAGECLASS
persistentvolumeclaim/mariadb  ...   ocs-external-storagecluster-ceph-rbd

NAME      TYPE      ...   PORT(S)
service/mariadb  LoadBalancer  ...   3306:31402/TCP

NAME          READY  UP-TO-DATE  AVAILABLE  AGE
deployment.apps/mariadb  1/1    1          1          5m51s

NAME      TYPE      DATA  AGE
secret/mariadb Opaque  4     5m54s

NAME      DATA  AGE
configmap/database  1     5m53s
```

- 1.4. Review the UID and GID assignment for the database namespace. The values might be different in your environment.

```
[student@workstation ~]$ oc get namespace database -o yaml
...output omitted...
metadata:
  annotations:
    ...
  openshift.io(sa.scc.supplemental-groups: 1000810000/10000
  openshift.io(sa.scc.uid-range: 1000810000/10000
labels:
  kubernetes.io/metadata.name: database
  ...output omitted...
```

- 1.5. Use the DO380/labs/backup-restore/view-db.sh script to connect to the MariaDB database in the mariadb deployment, and list the rows in the application_logs table from the application database.

```
[student@workstation ~]$ DO380/labs/backup-restore/view-db.sh
+ oc exec -c mariadb deploy/mariadb -- bash -c 'mariadb -u ${MARIADB_USER} \
-p"${MARIADB_PASSWORD}" \
${MARIADB_DATABASE} -te \
"SELECT * FROM application_logs ORDER BY id DESC LIMIT 5;"'
+-----+-----+-----+
| id | time           | app       | message        |
+-----+-----+-----+
| 25 | 2023-12-05 10:48:12 | db-loader | ad42...785d |
| 24 | 2023-12-05 10:48:02 | db-loader | 26b3...be3e |
| 23 | 2023-12-05 10:47:52 | db-loader | ed74...6ed8 |
| 22 | 2023-12-05 10:47:42 | db-loader | f9f5...3270 |
| 21 | 2023-12-05 10:47:32 | db-loader | 2c65...b85c |
+-----+-----+-----+
```

**Note**

The db-loader application is deployed in the application project. This application creates a record in the application_logs table every 10 seconds to simulate database activity during the backup operations.

- 2. Create a crash-consistent backup of the database project.

- 2.1. Change to the ~/D0380/labs/backup-restore directory.

```
[student@workstation ~]$ cd ~/D0380/labs/backup-restore
[student@workstation backup-restore]$
```

- 2.2. Create a db-manual backup resource in the openshift-adp namespace to back up the database project. Use resource filtering to back up only the listed MariaDB resources in the previous step.

Modify the partial resource definition in the ~/D0380/labs/backup-restore/backup-db-manual.yml file as follows:

```
apiVersion: velero.io/v1
kind: Backup
metadata:
  name: db-manual
  namespace: openshift-adp
spec:
  includedNamespaces:
    - database
  orLabelSelectors:
    - matchLabels:
        app: mariadb
    - matchLabels:
        kubernetes.io/metadata.name: database
  includedResources:
    - namespace
    - deployments
    - configmaps
    - secrets
    - pvc
    - pv
    - services
```

**Important**

So that OADP can restore the UID and GID assigned to the database namespace, include the namespace resource type and the kubernetes.io/metadata.name: database label in the backup definition.

- 2.3. Apply the configuration for the backup resource.

```
[student@workstation backup-restore]$ oc apply -f backup-db-manual.yml
backup.velero.io/db-manual created
```

► 3. Review the status of the backup.

- 3.1. Create an alias to access the `velero` binary from the Velero deployment in the `openshift-adp` namespace.

```
[student@workstation backup-restore]$ alias velero='\
  oc -n openshift-adp exec deployment/velero -c velero -it -- ./velero'
```

- 3.2. Use the `velero` command to get the status of the `db-manual` backup. Monitor the output to verify that the backup status is `Completed`. The backup process takes several minutes. Note the backup creation time for the next step.

```
[student@workstation backup-restore]$ velero get backup db-manual
NAME      STATUS      ERRORS   WARNINGS   CREATED
db-manual  Completed    0         0          2023-12-04 12:16:40  ...
```

- 3.3. Use the `view-db.sh` script to review the content of the `application_logs` table from the `mariadb` database in the `mariadb` deployment. Use the backup creation time from the previous step as a script argument to list the rows that were created before the backup time. You use the result to compare with the database that you restore in a later step.

```
[student@workstation backup-restore]$ ./view-db.sh "2023-12-04 12:16:40"
+ oc exec -c mariadb deploy/mariadb -- bash -c 'mariadb -u ${MARIADB_USER}
-p"${MARIADB_PASSWORD}"      ${MARIADB_DATABASE} -te      "SELECT * FROM
application_logs WHERE time <= \"2023-12-04 12:16:40\"      ORDER BY id DESC LIMIT
5;"'
+-----+-----+-----+
| id | time           | app       | message
+-----+-----+-----+
| 36 | 2023-12-05 10:50:03 | db-loader | 2232...4341 |
| 35 | 2023-12-05 10:49:53 | db-loader | b8a2...6ddf |
| 34 | 2023-12-05 10:49:43 | db-loader | 9cb1...a6d1 |
| 33 | 2023-12-05 10:49:33 | db-loader | c0e5...5e8b |
| 32 | 2023-12-05 10:49:23 | db-loader | e3c9...1929 |
+-----+-----+-----+
```

► 4. Restore the `db-manual` backup to a new `database-crash` project.

- 4.1. Create a `db-crash` restore resource in the `openshift-adp` namespace to restore the `db-manual` backup. Restore the backup to the `database-crash` project. Modify the partial resource definition in the `~/DO380/labs/backup-restore/restore-db-crash.yaml` file as follows:

```
apiVersion: velero.io/v1
kind: Restore
metadata:
  name: db-crash
  namespace: openshift-adp
spec:
  backupName: db-manual
```

```
restorePVs: true
namespaceMapping:
  database: database-crash
```

- 4.2. Apply the configuration for the restore resource.

```
[student@workstation backup-restore]$ oc apply -f restore-db-crash.yml
restore.velero.io/db-crash created
```

- 4.3. Use the `velero` command to get the status of the `db-crash` restore resource. Monitor the output to verify that the backup status is `Completed`. The restoration process takes several minutes.

```
[student@workstation backup-restore]$ velero get restore db-crash
NAME      BACKUP      STATUS      ...
db-crash  db-manual   Completed   ...
```

5. Review the content of the `database-crash` project and verify the database integrity.

- 5.1. Change to the `database-crash` project.

```
[student@workstation backup-restore]$ oc project database-crash
Now using project "database-crash" on server "https://api.ocp4.example.com:6443".
```

- 5.2. List the resources in the `database-crash` project and compare them with the resources in the `database` project from the previous step.

```
[student@workstation backup-restore]$ oc get \
  pvc,svc,deployment,secret,configmap -l app=mariadb
NAME          ...      STORAGECLASS
persistentvolumeclaim/mariadb ...  ocs-external-storagecluster-ceph-rbd

NAME          TYPE      ...      PORT(S)
service/mariadb LoadBalancer ...  3306:31402/TCP

NAME          READY    UP-TO-DATE  AVAILABLE  AGE
deployment.apps/mariadb 1/1      1           1          5m51s

NAME          TYPE      DATA     AGE
secret/mariadb Opaque   4        5m54s

NAME          DATA     AGE
configmap/database 1       5m53s
```

- 5.3. Review the UID and GID assignment for the `database-crash` namespace. Verify that the numbers are the same as for the `database` namespace from the previous step.

```
[student@workstation ~]$ oc get namespace database-crash -o yaml
...output omitted...
metadata:
```

```
annotations:  
  ...output omitted...  
  openshift.io(sa.scc.supplemental-groups: 1000810000/10000  
  openshift.io(sa.scc.uid-range: 1000810000/10000  
labels:  
  kubernetes.io/metadata.name: database-crash  
...output omitted...
```

- 5.4. Use the `mariadb-check` command to check the integrity and to repair the database tables.

```
[student@workstation backup-restore]$ oc exec -c mariadb deploy/mariadb -- \  
bash -c \  
'mariadb-check -u root -p"${MARIADB_ROOT_PASSWORD}" -A --auto-repair'  
  
application.application_logs  
warning : 1 client is using or hasn't closed the table properly  
warning : Size of datafile is: 13296      Should be: 13120  
error   : Wrong bytesec: 0-0-0 at linkstart: 13120  
error   : Corrupt  
mysql.column_stats                                OK  
...output omitted...  
  
Repairing tables  
application.application_logs  
info    : Wrong bytesec: 0- 0- 0 at 13120; Skipped  
status  : OK
```



Note

In this example, the `application_logs` table is corrupted and some data might be lost after the repair. The result might be different in your lab environment.

- 5.5. Use the `view-db.sh` script to list the records in the `application_logs` table from the `application` database in the `database-crash` project. Compare the result with the content of the database in the `database` project from the previous step.

```
[student@workstation backup-restore]$ ./view-db.sh  
+ oc exec -c mariadb deploy/mariadb -- bash -c 'mariadb -u ${MARIADB_USER} \  
-p"${MARIADB_PASSWORD}" \  
${MARIADB_DATABASE} -te \  
"SELECT * FROM application_logs ORDER BY id DESC LIMIT 5;"  
+-----+-----+-----+  
| id | time           | app       | message     |  
+-----+-----+-----+  
| 35 | 2023-12-05 10:49:53 | db-loader | b8a2...6ddf |  
| 34 | 2023-12-05 10:49:43 | db-loader | 9cb1...a6d1 |  
| 33 | 2023-12-05 10:49:33 | db-loader | c0e5...5e8b |  
| 32 | 2023-12-05 10:49:23 | db-loader | e3c9...1929 |  
| 31 | 2023-12-05 10:49:13 | db-loader | dc45...b7d8 |  
+-----+-----+-----+
```

**Note**

Compare the latest row ID with the ID from the previous step. Rows can be missing. This outcome can result from the database corruption that was detected in the previous step, or because the records were not written to the disk before the backup started.

- ▶ 6. Delete the crash-consistent backup and all related resources.

- 6.1. Review the content of the S3 bucket by using the `s3cmd` command.

```
[student@workstation backup-restore]$ s3cmd ls -r
...output omitted...
... s3://backup-.../oadp/backups/db-manual/velero-backup.json ①
... s3://backup-.../oadp/backups/db-manual/....gz
...output omitted...
... s3://backup-.../oadp/kopia/database/kopia.repository ②
... s3://backup-.../oadp/kopia/database/551...c2b
...output omitted...
... s3://backup-.../oadp/restores/db-crash/restore-db-crash-logs.gz ③
... s3://backup-.../oadp/restores/db-crash/....gz
...output omitted...
```

- ① The `/oadp/backups/backup_name` path contains the backup of the Kubernetes resources and the logs of the backup.
- ③ The `/oadp/kopia/project_name` path is the Kopia backup repository that contains backups of all data volumes in a project.
- ② The `/oadp/restores/restore_name` path contains the restore logs.

**Note**

The database uses a container image that is stored on an external registry and is therefore not included in the backup.

- 6.2. Delete the db-manual backup by using the `velero` command.

```
[student@workstation backup-restore]$ velero delete backup db-manual
Are you sure you want to continue (Y/N)? y
Request to delete backup "db-manual" submitted successfully.
The backup will be fully deleted after all associated data (disk snapshots, backup files, restores) are removed.
```

**Note**

Velero deletes all backup and restore resources that are associated with the specified backup in the OpenShift cluster. Velero also deletes the backup and restore files from the object storage.

The Kopia backup repository is never deleted, even if all backups from the project are removed. Kopia removes expired and unused volume backups from the object storage during the repository maintenance which happens occasionally.

- 6.3. Monitor the status of the db-manual backup with the `velero` command and wait until OADP deletes the backup.

```
[student@workstation backup-restore]$ velero get backup db-manual
NAME      STATUS    ERRORS   WARNINGS   ...
db-manual  Deleting  0        0          ...
```

**Note**

If the command returns the following error, then the backup is removed from both OpenShift and the object storage:

```
An error occurred: backups.velero.io "db-manual" not found
command terminated with exit code 1
```

- 6.4. Verify that the backup and restore resources are removed from the `openshift-adp` namespace. It can take a few minutes for the OADP operator to remove the resources.

```
[student@workstation backup-restore]$ oc -n openshift-adp get backup,restore
No resources found in openshift-adp namespace.
```

- 6.5. Review the content of the S3 bucket. The db-manual backup and restore are automatically deleted and only the Kopia backup repository for the database project remains.

```
[student@workstation backup-restore]$ s3cmd la -r
...
... s3://backup-..../oadp/kopia/database/_log_2024...5f5
... s3://backup-..../oadp/kopia/database/kopia.repository
... s3://backup-..../oadp/kopia/database/551...c2b
... output omitted...
```

- 7. Schedule a recurring application-consistent backup of the database project.

Use backup hooks to ensure that all database operations are flushed to the file system and that the database is locked in read-only mode during the backup.

Backups must be automatically deleted after 15 days.

- 7.1. Create a db-backup schedule resource in the `openshift-adp` namespace to back up the database project every week.

Set the schedule to 7 AM every Sunday. For this exercise, disable the schedule with the paused field to prevent unexpected backups from starting during the hands-on activity. You manually trigger a backup from this schedule in a later step.

Modify the partial resource definition in the `~/DO380/labs/backup-restore/schedule-db-backup.yml` file as follows:

```
apiVersion: velero.io/v1
kind: Schedule
metadata:
  name: db-backup
  namespace: openshift-adp
spec:
  schedule: "0 7 * * 0"
  paused: true
  template:
    ttl: 360h0m0s
    includedNamespaces:
      - database
    orLabelSelectors:
      - matchLabels:
          app: mariadb
      - matchLabels:
          kubernetes.io/metadata.name: database
    includedResources:
      - namespace
      - deployments
      - configmaps
      - secrets
      - pvc
      - pv
      - services
      - pods
    hooks:
      resources:
        - name: mariadb-readonly
          pre:
            - exec:
                container: mariadb
                command:
                  - /bin/bash
                  - -c
                  - |
                    mariadb -u "root" \
                      -p"${MARIADB_ROOT_PASSWORD}" \
                      ${MARIADB_DATABASE} -e \
                      "set global read_only=1; \ ①
                      BACKUP STAGE START; \
                      BACKUP STAGE BLOCK_COMMIT;"; ②
                    sync; ③
          post:
            - exec:
                container: mariadb
                command:
```

```

- /bin/bash
- -c
- |
mariadb -u "root" \
-p"${MARIADB_ROOT_PASSWORD}" \
${MARIADB_DATABASE} -e \
"set global read_only=0;"; ④

```

- ① The `set global read_only=1` command locks the database in read-only mode to prevent any writes during the backup.
- ② The `BACKUP STAGE` commands ensure that all database tables are flushed to the disk.
- ③ The `sync` command ensures that all cached writes are flushed to the persistent storage.
- ④ The `set global read_only=0` command removes the database lock after the backup is completed.



Warning

Because backup hooks are executed only on pods that are included in the backup, you must include the pod resource in the `includedResources` field.

- 7.2. Apply the configuration for the schedule resource.

```
[student@workstation backup-restore]$ oc apply -f schedule-db-backup.yml
schedule.velero.io/db-backup created
```

- 7.3. Verify the status of the db-backup schedule.

```
[student@workstation backup-restore]$ velero get schedule
NAME      STATUS ... SCHEDULE    BACKUP TTL   LAST BACKUP  SELECTOR  PAUSED
db-backup New     ... 0 7 * * 0 360h0m0s  n/a        <none>   true
```

- 7.4. Use the `velero` command to trigger a new backup from the db-backup schedule. Save the backup name for a later step.

```
[student@workstation backup-restore]$ velero create backup \
--from-schedule=db-backup
...output omitted...
Creating backup from schedule, all other filters are ignored.
Backup request "db-backup-20231205114600" submitted successfully.
...output omitted...
```

- 7.5. Use the `velero` command to monitor the status of the backup and wait until the backup is in the `Completed` state. The backup process takes several minutes. Note the creation time of the backup from the `CREATED` field.

Press `Ctrl+C` to exit the `watch` command when done.

```
[student@workstation backup-restore]$ watch -n 10 oc -n openshift-adp \
exec deployment/velero -c velero -it -- ./velero get backup \
-l velero.io/schedule-name=db-backup
NAME          STATUS    ERRORS  WARNINGS  CREATED
db-backup-2023...00  Completed   0        0      2023-12-05 11:46:00 ...
```

- 7.6. Use the `view-db.sh` script to identify the rows immediately before the backup time in the `application_logs` table from the application database in the `mariadb` deployment. Use the date and time from the previous step as argument of the script.

```
[student@workstation backup-restore]$ ./view-db.sh -n database "2023-12-05
11:46:00"
+ oc exec -n database -c mariadb deploy/mariadb -- bash -c 'mariadb -u
${MARIADB_USER} -p"${MARIADB_PASSWORD}" ${MARIADB_DATABASE} -te      "SELECT *
FROM application_logs WHERE time <= \"2023-12-05 11:46:00\"      ORDER BY id DESC
LIMIT 5;"'
+-----+-----+-----+
| id | time           | app       | message    |
+-----+-----+-----+
| 371 | 2023-12-05 11:45:58 | db-loader | 325ef...f337 |
| 370 | 2023-12-05 11:45:48 | db-loader | 2505f...3215 |
| 369 | 2023-12-05 11:45:38 | db-loader | 77ba0...561a |
| 368 | 2023-12-05 11:45:27 | db-loader | 6bffb...8f5f |
| 367 | 2023-12-05 11:45:17 | db-loader | a895d...3dfd |
+-----+-----+-----+
```

► 8. Restore the application-consistent backup to the `database-backup` project.

- 8.1. Create a `db-backup` restore resource in the `openshift-adp` namespace. Restore the backup from the previous step to the `database-backup` project.

Modify the partial resource definition in the `~/D0380/labs/backup-restore/restore-db-backup.yml` file as follows:

```
apiVersion: velero.io/v1
kind: Restore
metadata:
  name: db-backup
  namespace: openshift-adp
spec:
  backupName: db-backup-20231205114600
  restorePVs: true
  namespaceMapping:
    database: database-backup
```

- 8.2. Apply the configuration for the restore resource.

```
[student@workstation backup-restore]$ oc apply -f restore-db-backup.yml
```

- 8.3. Use the `velero` command to get the status of the `db-backup` restore resource. Monitor the output to verify that the status is `Completed`. The restoration process takes several minutes.

```
[student@workstation backup-restore]$ velero get restore db-backup
NAME      BACKUP          STATUS    ...  ERRORS  WARNINGS
db-backup  db-backup-20231205114600  Completed  ...  0        0
```

- 9. Review the content of the database-backup project.

- 9.1. Change to the database-backup project.

```
[student@workstation backup-restore]$ oc project database-backup
Now using project "database-backup" on server "https://api.ocp4.example.com:6443".
```

- 9.2. Use the mariadb-check command to verify the integrity of the database tables.

```
[student@workstation backup-restore]$ oc exec -c mariadb deploy/mariadb -- \
bash -c 'mariadb-check -u root -p"${MARIADB_ROOT_PASSWORD}" -A'

application.application_logs           OK
mysql.column_stats                     OK
...output omitted...
mysql.transaction_registry             OK
```



Note

All tables are OK and no corruption is detected.

- 9.3. Use the view-db.sh script to list the rows in the application_logs table from the application database in the database-backup project. Compare the result with the content of the database in the database project from the previous step.

```
[student@workstation backup-restore]$ ./view-db.sh
+ oc exec -c mariadb deploy/mariadb -- bash -c 'mariadb -u ${MARIADB_USER} \
-p"${MARIADB_PASSWORD}" \
${MARIADB_DATABASE} -te \
"SELECT * FROM application_logs ORDER BY id DESC LIMIT 5;"'
+-----+-----+-----+
| id   | time            | app       | message   |
+-----+-----+-----+
| 371  | 2023-12-05 11:45:58 | db-loader | 325e...f337 |
| 370  | 2023-12-05 11:45:48 | db-loader | 2505...3215 |
| 369  | 2023-12-05 11:45:38 | db-loader | 77ba...561a |
| 368  | 2023-12-05 11:45:27 | db-loader | 6bff...8f5f |
| 367  | 2023-12-05 11:45:17 | db-loader | a895...3dfd |
+-----+-----+-----+
```



Note

Compare the latest row ID with the ID from the previous step. All rows until the start of the backup are present.

- 10. Clean up the resources.

10.1. Change to the student HOME directory.

```
[student@workstation backup-restore]$ cd  
[student@workstation ~]$
```

10.2. Delete the database-crash and database-backup projects.

```
[student@workstation ~]$ oc delete project database-crash database-backup  
project.project.openshift.io "database-crash" deleted  
project.project.openshift.io "database-backup" deleted
```

10.3. Delete the backup that you created in the previous step by using the `velero` command. Use the `velero.io/schedule-name=db-backup` label to delete all backups that were created from the `db-backup` schedule.

```
[student@workstation backup-restore]$ velero delete backup \  
-l velero.io/schedule-name=db-backup  
Are you sure you want to continue (Y/N)? y  
Request to delete backup "db-backup-20231204141509" submitted successfully.  
The backup will be fully deleted after all associated data (disk snapshots, backup  
files, restores) are removed.
```

Finish

On the workstation machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish backup-restore
```

► Lab

Backup, Restore, and Migration of Applications with OADP

Deploy the OADP operator and describe its features and dependencies.

Configure one-time and scheduled backups with OADP and restore from them.

Outcomes

- Configure the OpenShift API for Data Protection operator to back up OpenShift applications.
- Schedule a complete backup of an application.
- Create and restore an application-consistent backup by using backup hooks.
- Clean up previous backups.

Before You Begin

As the student user on the workstation machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start backup-review
```

Instructions

Your team deployed a wiki application on OpenShift to host its operational documentation. You are asked to configure a scheduled daily backup by using OADP.

The application is deployed in the `wiki` project and is accessible at the following URL: <https://mediawiki-wiki.apps.ocp4.example.com>. The application uses the MediaWiki software and a PostgreSQL database with two volumes: uploaded images and documents are stored on a CephFS volume, and the database uses a Ceph RBD volume.

A colleague started installing OADP and performed the following tasks:

- Install the required operators for OADP.
- Create the S3 bucket and configure the `s3cmd` command with the S3 credentials.
- Create partial OADP configuration files in the `~/D0380/labs/backup-review` path with the S3 information and credentials.

Review and complete the configuration of OADP. Ensure that all storage classes that the application is using can be backed up with CSI snapshots. Configure OADP to store volume backups on the S3 object storage.

Configure a scheduled daily backup of the `wiki` project. The backup must start at 11 PM every day. To ensure backup consistency, you must lock the application to prevent any writes during the backup, and prepare the database for a snapshot before the backup by using backup hooks.

All Kubernetes resources required for the application have the `app.kubernetes.io/part-of=mediawiki` label. Resources specific to the MediaWiki component have the `app.kubernetes.io/name: mediawiki` label. Resources specific to the PostgreSQL database have the `app.kubernetes.io/name: postgresql` label. The backup must include only the resources with those labels.

To lock the MediaWiki application in read-only mode, create a `/data/images/lock_yBgMBwiR` file in the `mediawiki` container. This lock file must contain a single line with the "backup in progress" string to display the lock to the users in the application. To unlock the application after the backup is complete, remove the lock file.



Note

The `/data/images` path is in the `mediawiki` volume. The lock file is therefore included in the backup.

To prepare the PostgreSQL database for a snapshot and to flush all in-memory transactions to the disk, use the `psql -c "CHECKPOINT;"` command in the `postgresql` container.

To prepare for an upcoming update of the software stack, you are asked to restore an exact copy of the `wiki` application to a new `wiki-staging` project, so your team can test the update in a staging environment first.

Trigger an immediate backup from the scheduled backup and restore it to the `wiki-staging` project. Ensure that the application is unlocked after the restore by deleting the MediaWiki lock file with a restore hook.

To distinguish the production environment from the staging environment, change the MediaWiki site name to `D0380 Team Wiki Staging`. This setting is controlled by the `MEDIAWIKI_SITE_NAME` environment variable on the `mediawiki` deployment.

Configure the MediaWiki URL on the `mediawiki` deployment with the `MEDIAWIKI_SITE_SERVER` environment variable to match the route URL in the `wiki-staging` project.

Finally, remove all manual backup and restore resources from OpenShift and from the object storage.



Note

The S3 object storage in the lab environment uses a custom certificate that is signed with the OpenShift service CA. If you use the `velero describe --details` or `velero logs` commands, you must specify the CA certificate with the `--cacert` parameter.

The CA certificate for the OpenShift service CA is available from the `velero` pod in the `/run/secrets/kubernetes.io/serviceaccount/service-ca.crt` path.

1. Complete the configuration files in the `~/D0380/labs/backup-review` path and configure OADP with the following requirements:
 - Ensure that OADP can back up volumes with CSI snapshots.
 - Volume backups must be stored in the S3 bucket.
 - S3 credentials are in the `credentials-velero` configuration file.

Chapter 2 | Backup, Restore, and Migration of Applications with OADP

- Partial OADP configuration with S3 information is in the `oadp-config.yaml` file.
2. Identify the storage classes that the application uses. Ensure that the application volumes can be backed up by using CSI snapshots. Configure the matching volume snapshot classes for OADP.
3. Create a scheduled daily backup of the `wiki` project. The backup must start at 11 PM every day. Use backup hooks to prepare the application for backup by using the provided commands on both the MediaWiki and PostgreSQL pods. Ensure that both CephFS and RBD volumes are backed up.

For this exercise, disable the schedule with the `paused` field to prevent unexpected backups from starting during the hands-on activity. You manually trigger a backup from this schedule in a later step.

Only the resources that the application uses or required by OADP should be in the backup. Use the application labels to filter the resources to back up. The application uses the following resource types:

- Deployments
- Stateful Sets
- Secrets
- Services
- Routes
- Persistent Volumes

You can use the partial resource definition files in the `~/D0380/labs/backup-review` path.

4. Trigger an immediate backup from the scheduled backup, and restore it to the `wiki-staging` project. Use restore hooks to unlock the application.
- The backup and restore should be completed without any errors or warnings.
5. Change the MediaWiki site name and site URL in the `wiki-staging` project.
6. Remove the manual backup and restore resources from the previous steps.

Evaluation

As the `student` user on the `workstation` machine, use the `lab` command to grade your work. Correct any reported failures and rerun the command until successful.

```
[student@workstation ~]$ lab grade backup-review
```

Finish

As the `student` user on the `workstation` machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish backup-review
```

► Solution

Backup, Restore, and Migration of Applications with OADP

Deploy the OADP operator and describe its features and dependencies.

Configure one-time and scheduled backups with OADP and restore from them.

Outcomes

- Configure the OpenShift API for Data Protection operator to back up OpenShift applications.
- Schedule a complete backup of an application.
- Create and restore an application-consistent backup by using backup hooks.
- Clean up previous backups.

Before You Begin

As the student user on the workstation machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start backup-review
```

Instructions

Your team deployed a wiki application on OpenShift to host its operational documentation. You are asked to configure a scheduled daily backup by using OADP.

The application is deployed in the `wiki` project and is accessible at the following URL: <https://mediawiki-wiki.apps.ocp4.example.com>. The application uses the MediaWiki software and a PostgreSQL database with two volumes: uploaded images and documents are stored on a CephFS volume, and the database uses a Ceph RBD volume.

A colleague started installing OADP and performed the following tasks:

- Install the required operators for OADP.
- Create the S3 bucket and configure the `s3cmd` command with the S3 credentials.
- Create partial OADP configuration files in the `~/D0380/labs/backup-review` path with the S3 information and credentials.

Review and complete the configuration of OADP. Ensure that all storage classes that the application is using can be backed up with CSI snapshots. Configure OADP to store volume backups on the S3 object storage.

Configure a scheduled daily backup of the `wiki` project. The backup must start at 11 PM every day. To ensure backup consistency, you must lock the application to prevent any writes during the backup, and prepare the database for a snapshot before the backup by using backup hooks.

Chapter 2 | Backup, Restore, and Migration of Applications with OADP

All Kubernetes resources required for the application have the `app.kubernetes.io/part-of=mediawiki` label. Resources specific to the MediaWiki component have the `app.kubernetes.io/name: mediawiki` label. Resources specific to the PostgreSQL database have the `app.kubernetes.io/name: postgresql` label. The backup must include only the resources with those labels.

To lock the MediaWiki application in read-only mode, create a `/data/images/lock_yBgMBwiR` file in the `mediawiki` container. This lock file must contain a single line with the "backup in progress" string to display the lock to the users in the application. To unlock the application after the backup is complete, remove the lock file.



Note

The `/data/images` path is in the `mediawiki` volume. The lock file is therefore included in the backup.

To prepare the PostgreSQL database for a snapshot and to flush all in-memory transactions to the disk, use the `psql -c "CHECKPOINT;"` command in the `postgresql` container.

To prepare for an upcoming update of the software stack, you are asked to restore an exact copy of the wiki application to a new `wiki-staging` project, so your team can test the update in a staging environment first.

Trigger an immediate backup from the scheduled backup and restore it to the `wiki-staging` project. Ensure that the application is unlocked after the restore by deleting the MediaWiki lock file with a restore hook.

To distinguish the production environment from the staging environment, change the MediaWiki site name to `D0380 Team Wiki Staging`. This setting is controlled by the `MEDIAWIKI_SITE_NAME` environment variable on the `mediawiki` deployment.

Configure the MediaWiki URL on the `mediawiki` deployment with the `MEDIAWIKI_SITE_SERVER` environment variable to match the route URL in the `wiki-staging` project.

Finally, remove all manual backup and restore resources from OpenShift and from the object storage.



Note

The S3 object storage in the lab environment uses a custom certificate that is signed with the OpenShift service CA. If you use the `velero describe --details` or `velero logs` commands, you must specify the CA certificate with the `--cacert` parameter.

The CA certificate for the OpenShift service CA is available from the `velero` pod in the `/run/secrets/kubernetes.io/serviceaccount/service-ca.crt` path.

1. Complete the configuration files in the `~/D0380/labs/backup-review` path and configure OADP with the following requirements:
 - Ensure that OADP can back up volumes with CSI snapshots.
 - Volume backups must be stored in the S3 bucket.
 - S3 credentials are in the `credentials-velero` configuration file.

- Partial OADP configuration with S3 information is in the `oadp-config.yml` file.

1.1. Connect to the OpenShift cluster as the `admin` user with `redhatocp` as the password.

```
[student@workstation ~]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
Login successful.
...output omitted...
```

1.2. Change to the `~/D0380/labs/backup-review` directory.

```
[student@workstation ~]$ cd ~/D0380/labs/backup-review
```

1.3. Create the `cloud-credentials` secret in the `openshift-adp` namespace with the `credentials-velero` file content.

```
[student@workstation backup-review]$ oc create secret generic \
cloud-credentials -n openshift-adp --from-file cloud=credentials-velero
secret/cloud-credentials created
```

1.4. Complete the OADP configuration in the `oadp-config.yml` file to enable CSI snapshots and Data Mover.

```
apiVersion: oadp.openshift.io/v1alpha1
kind: DataProtectionApplication
metadata:
  name: oadp-config
  namespace: openshift-adp
spec:
  ...output omitted...
  configuration:
    nodeAgent:
      enable: true
      uploaderType: kopia
    velero:
      defaultPlugins:
        - openshift
        - aws
        - csi
      defaultSnapshotMoveData: true
```

1.5. Apply the OADP configuration by using the `oadp-config.yml` file.

```
[student@workstation backup-review]$ oc apply -f oadp-config.yml
dataprotectionapplication.oadp.openshift.io/oadp-config created
```

1.6. Verify that the `BackupStorageLocation` object is created and in the Available phase.

Chapter 2 | Backup, Restore, and Migration of Applications with OADP

```
[student@workstation backup-review]$ oc get -n openshift-adp \
  BackupStorageLocations
NAME          PHASE      LAST VALIDATED   AGE      DEFAULT
oadp-config-1  Available   18s            2m51s   true
```

2. Identify the storage classes that the application uses. Ensure that the application volumes can be backed up by using CSI snapshots. Configure the matching volume snapshot classes for OADP.
 - 2.1. List the persistent volume claims in the `wiki` namespace and identify the storage classes.

```
[student@workstation backup-review]$ oc -n wiki get pvc
NAME           ...   STORAGECLASS           AGE
mediawiki       ...   ocs-external-storagecluster-cephfs   5m45s
postgresql-postgresql-0 ...   ocs-external-storagecluster-ceph-rbd   5m43s
```

- 2.2. Identify the provisioner for each storage class from the previous step.

```
[student@workstation backup-review]$ oc get storageclasses \
  ocs-external-storagecluster-ceph-rbd ocs-external-storagecluster-cephfs
NAME           PROVISIONER           ...
ocs-external-storagecluster-ceph-rbd  openshift-storage.rbd.csi.ceph.com ...
ocs-external-storagecluster-cephfs    openshift-storage.cephfs.csi.ceph.com ...
```

- 2.3. List the available volume snapshot classes. Ensure that each storage class from the previous step has a matching volume snapshot class with the same driver.

```
[student@workstation backup-review]$ oc get volumesnapshotclasses
NAME           DRIVER
ocs-...-cephfsplugin-snapclass  openshift-storage.cephfs.csi.ceph.com
ocs-...-rbdplugin-snapclass     openshift-storage.rbd.csi.ceph.com
```

- 2.4. For each volume snapshot class, set the `velero.io/csi-volumesnapshot-class` label to `true`.

```
[student@workstation backup-review]$ oc label volumesnapshotclass \
  velero.io/csi-volumesnapshot-class="true" --all
.../ocs-external-storagecluster-cephfsplugin-snapclass labeled
.../ocs-external-storagecluster-rbdplugin-snapclass labeled
```

3. Create a scheduled daily backup of the `wiki` project. The backup must start at 11 PM every day. Use backup hooks to prepare the application for backup by using the provided commands on both the MediaWiki and PostgreSQL pods. Ensure that both CephFS and RBD volumes are backed up.

For this exercise, disable the schedule with the `paused` field to prevent unexpected backups from starting during the hands-on activity. You manually trigger a backup from this schedule in a later step.

Only the resources that the application uses or required by OADP should be in the backup. Use the application labels to filter the resources to back up. The application uses the following resource types:

- Deployments
- Stateful Sets
- Secrets
- Services
- Routes
- Persistent Volumes

You can use the partial resource definition files in the `~/DO380/labs/backup-review` path.

- 3.1. Modify the partial resource definition in the `~/DO380/labs/backup-review/schedule.yaml` file as follows:

```
apiVersion: velero.io/v1
kind: Schedule
metadata:
  name: wiki-backup
  namespace: openshift-adp
spec:
  schedule: "0 23 * * *"
  paused: true
  template:
    includedNamespaces:
      - wiki
    orLabelSelectors:
      - matchLabels:
          app.kubernetes.io/part-of: mediawiki
      - matchLabels:
          kubernetes.io/metadata.name: wiki
    includedResources:
      - deployments
      - statefulsets
      - secrets
      - pvc
      - pv
      - services
      - routes
      - pods
      - namespace
    hooks:
      resources:
        - name: mediawiki-lock
        labelSelector:
          matchLabels:
            app.kubernetes.io/name: mediawiki
      pre:
        - exec:
            container: mediawiki
            command:
              - /bin/bash
              - -c
              - echo "backup in progress" > /data/images/lock_yBgMBwiR;
      post:
        - exec:
```

```
container: mediawiki
command:
- rm
- ./data/images/lock_yBgMBwiR
- name: postgresql-checkpoint
labelSelector:
matchLabels:
app.kubernetes.io/name: postgresql
pre:
- exec:
  container: postgresql
  command:
- psql
- -c
- "CHECKPOINT;"
```

- 3.2. Apply the configuration for the schedule resource.

```
[student@workstation backup-review]$ oc apply -f schedule.yml
schedule.velero.io/wiki-backup created
```

- 3.3. Create an alias to access the `velero` binary from the Velero deployment in the `openshift-adp` namespace.

```
[student@workstation backup-review]$ alias velero='\
oc -n openshift-adp exec deployment/velero -c velero -it -- ./velero'
```

- 3.4. Verify the status of the schedule with the `velero` command.

```
[student@workstation backup-restore]$ velero get schedule
NAME      STATUS     CREATED          SCHEDULE   ...
wiki-backup  New       2023-11-17 14:37:39 +0000 UTC  0 23 * * *
```

4. Trigger an immediate backup from the scheduled backup, and restore it to the `wiki-staging` project. Use restore hooks to unlock the application.

The backup and restore should be completed without any errors or warnings.

- 4.1. Use the `velero` command to start a backup by using the schedule definition from the previous step. Note the name of the backup that the command creates, to use in the next step.

```
[student@workstation backup-review]$ velero backup create \
--from-schedule wiki-backup
...output omitted...
Creating backup from schedule, all other filters are ignored.
Backup request "wiki-backup-20231115113447" submitted successfully.
Run velero backup describe wiki-backup-20231115113447 or velero backup logs wiki-backup-20231115113447 for more details.
```

**Note**

The S3 object storage that is configured in the lab environment uses a custom certificate signed with the OpenShift service CA. You must add the CA certificate to the `velero backup describe --details` and `velero backup logs` commands as follow:

```
[user@host]$ velero backup logs \
--cacert=/run/secrets/kubernetes.io/serviceaccount/service-ca.crt \
wiki-backup-20231115113447
```

- 4.2. Monitor the status of the backup and verify that the backup ends with the Completed status. The backup process takes several minutes.

```
[student@workstation backup-review]$ velero get backup
NAME                  STATUS        ERRORS   WARNINGS   ...
wiki-backup-20231115113447  Completed      0         0          ...
```

- 4.3. Create the restore resource in the `openshift-adp` namespace to restore the `wiki` project to a new `wiki-staging` project. Use the backup name from the previous step. Configure a post-restore hook to remove the MediaWiki lock file from the `mediawiki` volume.

You can use the incomplete `restore.yml` file as the starting point.

```
apiVersion: velero.io/v1
kind: Restore
metadata:
  name: wiki-staging
  namespace: openshift-adp
spec:
  backupName: wiki-backup-20231115113447
  namespaceMapping:
    wiki: wiki-staging
  hooks:
    resources:
      - name: mediawiki-unlock
        labelSelector:
          matchLabels:
            app.kubernetes.io/name: mediawiki
    postHooks:
      - exec:
          container: mediawiki
          command:
            - rm
            - /data/images/lock_yBgMBwiR
```

- 4.4. Create the resource by using the YAML manifest.

```
[student@workstation backup-review]$ oc apply -f restore.yml
restore.velero.io/wiki-staging created
```

Chapter 2 | Backup, Restore, and Migration of Applications with OADP

- 4.5. Use the `velero` command to get the status of the restore. Monitor the output to verify that the restore status is **Completed**. The restore process takes several minutes.

```
[student@workstation backup-review]$ velero get restore
NAME      BACKUP          STATUS    ...  ERRORS  WARNINGS
wiki-staging  wiki-backup-20231115113447  Completed  ...  0        0
```

- 4.6. Review the restored resources in the `wiki-staging` project.

```
[student@workstation backup-review]$ oc -n wiki-staging get all
NAME           READY   STATUS    RESTARTS   AGE
pod/mediawiki-77dfffc8bd-nlc74  1/1     Running   0          67s
pod/postgresql-0                1/1     Running   0          66s

NAME            TYPE        CLUSTER-IP      EXTERNAL-IP    PORT(S)    AGE
service/mediawiki   ClusterIP   172.30.86.62  <none>        8080/TCP   67s
service/postgresql  ClusterIP   172.30.216.157 <none>        5432/TCP   67s

NAME           READY   UP-TO-DATE  AVAILABLE  AGE
deployment.apps/mediawiki  1/1       1           1          66s

NAME           DESIRED  CURRENT  READY   AGE
replicaset.apps/mediawiki-77dfffc8bd  1         1         1          66s

NAME           READY   AGE
statefulset.apps/postgresql  1/1     66s

NAME           HOST/PORT
route.route.openshift.io/mediawiki mediawiki-wiki-staging.apps.ocp4.example.com
```

5. Change the MediaWiki site name and site URL in the `wiki-staging` project.

- 5.1. Review the `mediawiki` route URL in the `wiki-staging` project.

```
[student@workstation backup-review]$ oc -n wiki-staging get route
NAME      HOST/PORT          PATH  SERVICES
mediawiki  mediawiki-wiki-staging.apps.ocp4.example.com          mediawiki
```

- 5.2. Set the `MEDIAWIKI_SITE_SERVER` environment variable on the `mediawiki` deployment to match the route URL from the previous step. Set the `MEDIAWIKI_SITE_NAME` environment variable to `D0380 Team Wiki Staging`.

```
[student@workstation backup-review]$ oc -n wiki-staging \
  set env deployment/mediawiki \
  MEDIAWIKI_SITE_SERVER=https://mediawiki-wiki-staging.apps.ocp4.example.com \
  MEDIAWIKI_SITE_NAME="D0380 Team Wiki Staging"
deployment.apps/mediawiki updated
```

- 5.3. Verify that the `mediawiki` deployment is ready.

```
[student@workstation backup-review]$ oc get deployment -n wiki-staging
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
mediawiki  1/1     1           1           14m
```

- 5.4. Open a web browser and navigate to <https://mediawiki-wiki-staging.apps.ocp4.example.com>. Verify the MediaWiki site name in the browser window and that you can edit any page.
6. Remove the manual backup and restore resources from the previous steps.
- 6.1. Use the `velero` command to delete the backup. Use the backup name from the previous step.

```
[student@workstation backup-review]$ velero delete \
  backup wiki-backup-20231115113447
Are you sure you want to continue (Y/N)? y
Request to delete backup "wiki-backup-20231115113447" submitted successfully.
The backup will be fully deleted after all associated data (disk snapshots, backup files, restores) are removed.
```

- 6.2. Change to the student HOME directory.

```
[student@workstation backup-review]$ cd
[student@workstation ~]$
```

Evaluation

As the `student` user on the `workstation` machine, use the `lab` command to grade your work. Correct any reported failures and rerun the command until successful.

```
[student@workstation ~]$ lab grade backup-review
```

Finish

As the `student` user on the `workstation` machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish backup-review
```

Summary

- A Kubernetes application backup must contain the required Kubernetes resources, container images, and persistent volumes to restore the application to a previous working state.
- OpenShift API for Data Protection provides a native backup solution for applications that run on OpenShift.
- OADP stores backups on object storage.
- OADP uses CSI snapshots to back up volumes and provides a data movement feature to move the volume snapshot content to object storage.
- Use the File System Backup feature with Restic to back up volumes that do not support volume snapshots.
- Prepare the application for backup with backup hooks to ensure data consistency.
- You can back up an application periodically by using a backup schedule.
- You must have administrative access to the `openshift-adp` namespace to back up and restore with OADP.

Chapter 3

Cluster Partitioning

Goal

Configure a subset of cluster nodes to be dedicated to a type of workload.

Sections

- Node Pools (and Quiz)
- Node Configuration with the Machine Configuration Operator (and Guided Exercise)
- Node Configuration with Special Purpose Operators (and Guided Exercise)

Lab

- Cluster Partitioning

Node Pools

Objectives

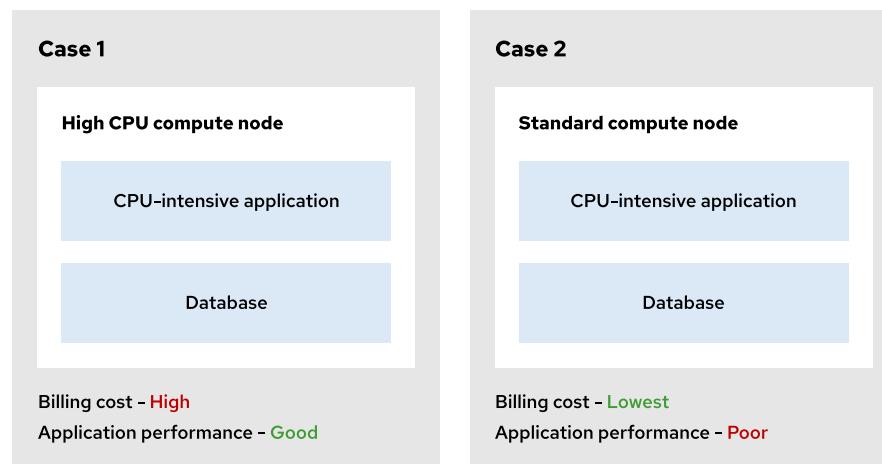
- Illustrate methods of adding and configuring OpenShift cluster nodes, by using node labels to configure nodes and to schedule pods to those nodes.

Introduction

The OpenShift cluster has control plane nodes and compute nodes. Compute nodes are also known as *worker nodes*. The control plane nodes manage the Red Hat OpenShift Container Platform cluster and orchestrate the distribution of the workloads on the compute nodes.

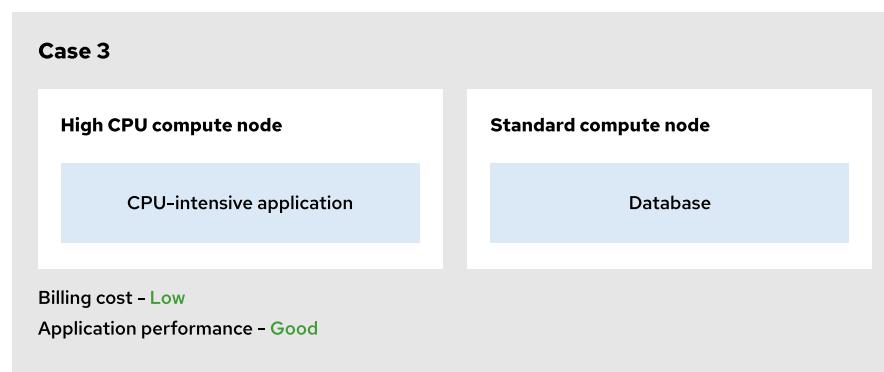
The applications that run in the OpenShift cluster have different specifications and requirements. Some applications might require compute nodes with graphics processing units (GPU) for graphical applications, artificial intelligence (AI) computations, or machine learning (ML) workloads. Some applications, such as simple websites, might not require specific hardware or specifications, and work smoothly on a small compute node.

For example, you run an application that is composed of many microservices that require high CPU capacity, and a database that does not have such a requirement.



You can run both the CPU-intensive application and the database instances on compute nodes with high CPU capacity. In this case, the CPU-intensive application runs smoothly, but the database does not use the high CPU node to full capacity, and the billing cost increases.

In OpenShift, you can deploy both the CPU-intensive application and the database instances on compute nodes with low CPU capacity. In this case, the billing cost reduces but the application does not work correctly because of the low CPU compute node.



The best solution is to deploy the CPU-intensive application instances on high CPU compute nodes and to deploy the database application on the low CPU compute nodes.

You need compute nodes with the different capacity and hardware to accommodate these applications, and a partition of these nodes to deploy the application. You can deploy these applications on specific nodes by using *node pools*.

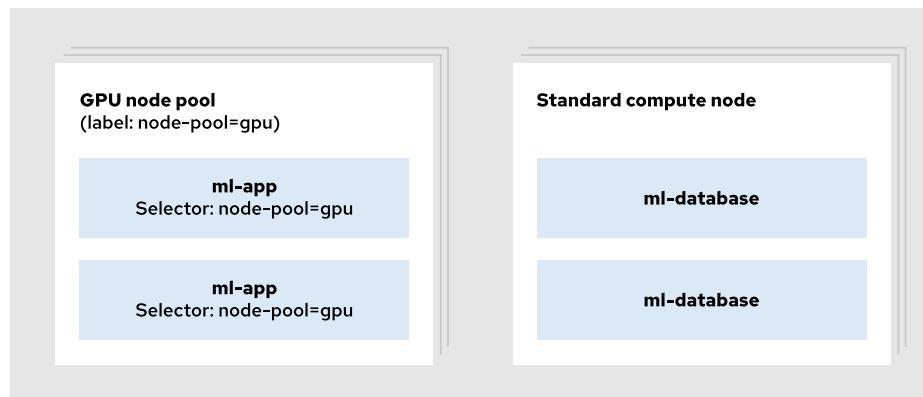
In some cases, you might face a "noisy neighbor" issue where a application gives poor performance and high latency due to other applications. You can configure compute nodes to be part of node pools for critical applications, to avoid the noisy neighbor issue.

Node Pools

A node pool is a logical group of OpenShift compute nodes. Compute nodes with similar hardware configurations can be organized into a node pool. Collecting these similar nodes into node pools provides a method of targeting placement for workload deployments.

You can configure multiple node pools for your workloads to ensure that your applications run on the required hardware.

An instance often requires a specific configuration for a compute node to run smoothly.



In this scenario, you have an application that uses ML and that also uses another database application. The ML application requires nodes with a GPU to run smoothly, whereas a database can run on a standard compute node.

Chapter 3 | Cluster Partitioning

In this case, provision compute nodes with GPU hardware and configure GPU compute nodes as part of a gpu node pool by using the `node-pool=gpu` label. You can schedule the ML application pods on the gpu node pool, whereas database pods deploy on other compute nodes.

You might have critical applications without specific hardware requirements and that can run on a standard compute node. These critical applications might face the "noisy neighbor" issue.

For example, consider a messaging application. The application does not require specific hardware and can run on any standard node, but cannot afford high network latency. Other non-critical applications might run on the same cluster node and increase network traffic and latency for the messaging application.

Although the application would typically run adequately on that compute node, other non-critical applications can affect the network performance of the compute node. The messaging application faces another "noisy neighbor" issue. To avoid this issue, you can use a node pool to organize your compute nodes to run critical applications.

Defining node pools for your cluster hardware is specific to the approach for application deployments in your organization. In a scenario with multi-tiered infrastructure deployments, such as a development, staging, or production tier of hardware, administrators can organize node pools to correspond to these tiers. Earlier hardware can populate a "development" node pool, whereas later infrastructure can be designated for production workloads through a "production" node pool. With this approach, developers can target these node pools to separate development, staging, and production applications that run in the cluster.

Node Provisioning

Compute node provisioning depends on the method of cluster installation.

In a user-provisioned infrastructure (UPI), you must manually provision new instances. The strategy for provisioning nodes is specific to your data center and your IT processes. However, the base steps remain the same, as follows:

1. Update the compute node Ignition file with an updated TLS certificate.
2. Install Red Hat Enterprise Linux CoreOS (RHCOS) from an ISO image or by using a Preboot eXecution Environment (PXE) boot approach.
3. Add the new instance to the ingress load balancer.
4. Approve the Certificate Signing Requests (CSRs).

In an installer-provisioned infrastructure (IPI), the machine API automatically performs scaling operations for supported cloud providers. Thus, you can modify the specified number of replicas in a machine set resource, and OpenShift communicates with the cloud provider to provision or remove instances. You can scale up and scale down by using autoscaling based on the workload requirements.

For more details about Red Hat OpenShift installation and node provisioning, refer to the DO322: *Red Hat OpenShift Installation Lab* training course.

Some managed clusters have protected machine sets for direct modification and support node provisioning by using node pools. For example, Red Hat OpenShift on AWS (ROSA) clusters use machine pools. You create a machine pool and select an instance type that is specific to the workload. The machine set is protected against direct modification. You can scale a machine pool manually or you can use the *cluster autoscaler*.

Chapter 3 | Cluster Partitioning

For more details about ROSA and autoscaling, refer to the *CS220: Creating and Configuring Production Red Hat OpenShift on AWS (ROSA) Clusters* training course.

Node Labels

Use node labels to configure compute nodes to be a member of a specific node pool. Node labels are key-value pairs that are attached to the node.

You can add a label to multiple compute nodes to designate each as a member of one node pool. In UPI clusters, specify node labels for the members of a node pool within a `MachineSet` object since the node labels can be overridden by the machine or cluster API.



Note

When applying node labels to existing members of a `MachineSet`, it is necessary to redeploy the nodes before the node label is available. For more information on defining and applying node labels within a `MachineSet`, see the following Red Hat article: <https://access.redhat.com/solutions/6955894>.

You can list the node and labels by using the `oc get nodes` command:

```
[user@host ~]$ oc get nodes worker01 --show-labels
NAME      STATUS    ROLES   ...   LABELS
worker01  Ready     worker   ...   beta.kubernetes.io/
arch=amd64,beta.kubernetes.io/os=linux,cluster.ocs.openshift.io/openshift-
storage=,kubernetes.io/arch=amd64,kubernetes.io/hostname=worker01,kubernetes.io/
os=linux,node-role.kubernetes.io/worker=,node.openshift.io/os_id=rhcos
```

Compute nodes can have many node labels. Add as many node labels to a compute node as needed for workload placement onto the intended hardware. You can add a custom node label to the compute nodes by using the `oc label node` command.

```
[user@host ~]$ oc label node/worker01 node-pool=gpu
node/worker01 labeled ①
[user@host ~]$ oc get nodes worker01 --show-labels
NAME      STATUS    ROLES   ...   LABELS
worker01  Ready     worker   ...   beta.kubernetes.io/
arch=amd64,beta.kubernetes.io/os=linux,cluster.ocs.openshift.io/openshift-
storage=,kubernetes.io/arch=amd64,kubernetes.io/hostname=worker01,kubernetes.io/
os=linux,node-role.kubernetes.io/worker=,node.openshift.io/os_id=rhcos, node-
pool=gpu ②
```

- ① The `oc label` command adds the `node-pool=gpu` label to the `worker01` compute node.
- ② The `oc get nodes` command displays the `node-pool=gpu` label for the `worker01` compute node.

You can configure multiple GPU compute nodes to be part of a gpu node pool by adding the same node label to all GPU-enabled compute nodes.

```
[user@host ~]$ oc label node/worker02 node-pool=gpu
node/worker02 labeled
[user@host ~]$ oc get nodes --selector node-pool=gpu
NAME      STATUS    ROLES   AGE     VERSION
worker01   Ready     worker  13d    v1.25.7+eab9cc9
worker02   Ready     worker  13d    v1.25.7+eab9cc9
```

The `oc get nodes` command with the selector filter shows all nodes with the `node-pool=gpu` label. You can schedule the application to target a compute node in the gpu node pool by using node labels. The next chapter discusses pod scheduling in detail.

Node Configuration

The OpenShift cluster manages upgrades to the nodes by using the machine configuration operator (MCO). You can manage nodes by defining groups of these nodes, which are called *machine configuration pools (MCP)*. The MCO uses the `master` and `worker` MCPs, by default. You can create custom MCPs for node pools. MCPs use labels to match one or more MCs to one or more nodes.

For example, you can create the gpu MCP. The gpu MCP selects nodes based on its assigned gpu node role.

Similarly, cloud providers also provide different compute and storage instances to meet customer requirements. For example etcd storage and I/O-intensive workloads require a low-latency storage solution for optimal performance. AWS has an `io2` EBS volume type to support these workloads. You can use the `gp2` EBS volume type for other workloads where intensive I/O is not required.



References

For more information about the control plane architecture, refer to the *Control Plane Architecture* chapter in the Red Hat OpenShift Container Platform 4.14 Architecture documentation at

https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/architecture/index#control-plane

For more information about adding compute nodes to a cluster, refer to the *About Adding RHEL Compute Nodes to a Cluster* section in the *Adding RHEL Compute Machines to an OpenShift Container Platform Cluster* chapter in the Red Hat OpenShift Container Platform 4.14 Machine Management documentation at

https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/machine_management/index#adding-rhel-compute

For more information about the installation prerequisites, refer to the *Prerequisites* section in the *Deploying Installer-provisioned Clusters on Bare Metal* chapter in the Red Hat OpenShift Container Platform 4.14 Installing documentation at

https://docs.redhat.com/en/documentation/openshift_container_platform/4.14/html-single/deploying_installer-provisioned_clusters_on_bare_metal/index#ipi-install-prerequisites

For more details about ROSA and Microsoft Azure Red Hat OpenShift, refer to the *DO120: Introduction to Red Hat OpenShift Service on AWS (ROSA)* and *DO121: Introduction to Microsoft Azure Red Hat OpenShift* training courses respectively.

Creating and Adding Additional Worker Nodes to the Cluster

<https://access.redhat.com/solutions/4270972>

Infrastructure Nodes in OpenShift 4

<https://access.redhat.com/solutions/5034771>

Kubernetes: Assign Pods to Nodes

<https://kubernetes.io/docs/tasks/configure-pod-container/assign-pods-nodes>

► Quiz

Node Pools

Choose the correct answers to the following questions:

► 1. Which is not a default node role for a Red Hat OpenShift Container Platform 4.14 cluster?

- a. worker
- b. master
- c. control-plane
- d. storage

► 2. What is the purpose of using node pools within your OpenShift cluster?

- a. Organizing cluster hardware to aid in targeting the most appropriate infrastructure hardware for each component of your deployed applications.
- b. Providing an automated monitoring solution for the regions in your data center.
- c. Ensuring uptime for workloads by always using the fastest hardware.
- d. Alleviating storage constraints for large data sets.

► 3. Which command adds a compute node to a node pool?

- a. `oc set label node/NODE_NAME pool=NODE_POOL_NAME`
- b. `oc label node/NODE_NAME node-pool=NODE_POOL_NAME`
- c. `oc get node NODE_NAME --selector node-pool=NODE_POOL_NAME`
- d. `oc create node-pool=NODE_POOL_NAME NODE_NAME`

► 4. Which is not an advantage of using node pools when running applications?

- a. Allocate resources efficiently for network-intensive applications to mitigate "noisy neighbor" issues.
- b. Automate security of applications in the cluster.
- c. Manage costs effectively by labeling expensive hardware.
- d. Separate infrastructure during maintenance and management.

► Solution

Node Pools

Choose the correct answers to the following questions:

- ▶ 1. Which is not a default node role for a Red Hat OpenShift Container Platform 4.14 cluster?
 - a. worker
 - b. master
 - c. control-plane
 - d. storage

- ▶ 2. What is the purpose of using node pools within your OpenShift cluster?
 - a. Organizing cluster hardware to aid in targeting the most appropriate infrastructure hardware for each component of your deployed applications.
 - b. Providing an automated monitoring solution for the regions in your data center.
 - c. Ensuring uptime for workloads by always using the fastest hardware.
 - d. Alleviating storage constraints for large data sets.

- ▶ 3. Which command adds a compute node to a node pool?
 - a. `oc set label node/NODE_NAME pool=NODE_POOL_NAME`
 - b. `oc label node/NODE_NAME node-pool=NODE_POOL_NAME`
 - c. `oc get node NODE_NAME --selector node-pool=NODE_POOL_NAME`
 - d. `oc create node-pool=NODE_POOL_NAME NODE_NAME`

- ▶ 4. Which is not an advantage of using node pools when running applications?
 - a. Allocate resources efficiently for network-intensive applications to mitigate "noisy neighbor" issues.
 - b. Automate security of applications in the cluster.
 - c. Manage costs effectively by labeling expensive hardware.
 - d. Separate infrastructure during maintenance and management.

Node Configuration with the Machine Configuration Operator

Objectives

- Apply operating system settings to cluster nodes with the machine configuration operator.

The Machine Configuration Operator

OpenShift uses Red Hat Enterprise Linux CoreOS (RHCOS) as the underlying operating system in the hosts. The entire operating system is updated as a single image, instead of on a package-by-package basis. The machine configuration operator (MCO) manages the RHCOS operating system upgrades and configuration changes.

Do not use traditional Red Hat Enterprise Linux management approaches, such as manually editing files or using system commands such as the `systemctl` command, for RHCOS operating system upgrades or configuration changes. Those changes can conflict with the MCO and the MCO can override them.

The MCO comprises the following pods:

`machine-config-operator`

These pods form the main operator workload that manages the rest of the MCO components.

`machine-config-controller`

The machine configuration controller (MCC) manages the synchronization of machine upgrades according to specified configurations through a machine configuration object. The MCC offers options to upgrade individual sets of machines.

`machine-config-server`

The machine configuration server (MCS) provides instance customizations to machines that join the cluster. The MCS uses Ignition to provide the instance customizations to cluster nodes that join the cluster. The RHCOS operating system downloads and processes Ignition files at boot time.



Note

Installing an OpenShift cluster and using Ignition files to configure the cluster nodes is explained in detail in the *DO322: Red Hat OpenShift Installation Lab* course.

The MCO also requires the `machine-config-daemon` daemonset that RHCOS provides. The machine configuration daemon (MCD) implements updates to machine configurations and validates each machine's state in accordance with the requested configuration.

The MCO can manage the following resources:

- `systemd` services
- SSH keys
- Kernel arguments
- Files in the `/var` or `/etc` directories.

Chapter 3 | Cluster Partitioning

The MCO can also manage directories, such as the /opt and /usr/local directories, which can be writeable if symbolically linked to the /var or /etc directories.

The MCO defines two custom resources (CRs), which are part of the machineconfiguration.openshift.io/v1 API group.

MachineConfig

A machine configuration (MC) CR declares instance customizations by using the Ignition configuration format.

MachineConfigPool

A machine configuration pool (MCP) CR uses labels to match one or more MCs to one or more nodes by means of the machineConfigSelector and nodeSelector parameters, respectively. This resource creates a pool of nodes with the same configuration. The MCO uses the MCP to track status when the MCO applies MCs to the nodes.

OpenShift administrators set custom node configurations by declaring the MachineConfig and MachineConfigPool CRs.

The following diagram shows the relationship between the node, MC, and MCP labels:



Figure 3.4: Relationship between the node, MC, and MCP labels

In the previous diagram, the MCP CR uses the nodeSelector parameter to select all the nodes with the worker role, and applies to them the MCs with the machineconfiguration.openshift.io/role: worker label that is selected by using the machineConfigSelector parameter.

The MCO also manages two custom resources (CRs) for modifying CRI-O container runtime settings and the Kubelet service: the ContainerRuntimeConfig and KubeletConfig CRs, respectively.

Machine Configurations

MCs declare instance configurations by using the Ignition configuration format.

Ignition files encode file contents by using the Base64 encoding scheme. Other items, such as systemd units or SSH keys, do not use the Base64 encoding scheme. In a terminal, use the base64 and base64 -d commands to encode and decode files or standard input.

You can also use Butane to create Ignition files and MCs. Butane simplifies handling Base64 encoding.

**Note**

For more information about how to use Butane to create MCs, refer to https://docs.redhat.com/en/documentation/openshift_container_platform/4.14/html-single/installation_configuration/index#installation-special-config-butane_installing-customizing

The following example shows an MC to modify the configuration file for the `journald` service by using the Base64 encoding scheme:

```
apiVersion: machineconfiguration.openshift.io/v1
kind: MachineConfig
metadata:
  labels:
    machineconfiguration.openshift.io/role: worker ①
  name: 60-journald ②
spec:
  config:
    ignition:
      version: 3.2.0
    storage:
      files:
        - contents:
            source: data:text/plain;charset=utf-8;base64,VGVzdG1...E8gKITAo= ③
          filesystem: root
          mode: 0644
          path: /etc/systemd/journald.conf
```

- ① The label for the MC. The MC uses the `machineconfiguration.openshift.io/role` MC role label. By default, OpenShift comes with preinstalled MCs for control plane and compute nodes.
- ② Prefix the name with a two-digit number that specifies when to apply the configuration, relative to MCs that belong to the same MCP. Higher numbers have precedence.
- ③ Use the data URL format to embed escaped file content. Base64 encoding is common for Ignition files.

The following example shows an MC to change the default kernel to a real-time kernel. This configuration does not use the Base64 encoding scheme:

```
apiVersion: machineconfiguration.openshift.io/v1
kind: MachineConfig
metadata:
  labels:
    machineconfiguration.openshift.io/role: worker
  name: 99-kernel-realtime
spec:
  kernelType: realtime
```

You can use the `oc explain mc` command to read the parameters that you can change by using an MC.

Chapter 3 | Cluster Partitioning

You can list the MCs on your cluster by using the `oc get machineconfig` or `oc get mc` commands. To list the MCs for a specific node label, use the `--selector` argument. The following example lists the MCs for the `worker` role in the cluster:

```
[user@host ~]$ oc get machineconfig \
--selector machineconfiguration.openshift.io/role=worker
NAME                GENERATEDBYCONTROLLER  IGNITIONVERSION AGE
00-worker            52fe...bf97          3.2.0          108d
01-worker-container-runtime 52fe...bf97          3.2.0          108d
01-worker-kubelet    52fe...bf97          3.2.0          108d
99-worker-chrony-conf-override 52fe...bf97          3.2.0          108d
99-worker-generated-registries 52fe...bf97          3.2.0          108d
99-worker-ssh        52fe...bf97          3.2.0          108d
```

The MCO reads MCs alphanumerically by the name, from `00-*` to `99-*`. OpenShift stores the resulting compilation of MCs in a rendered MC resource. Thus, if two or more MCs apply changes to the same file, then the MCO applies only the changes from the MC with a higher number. For two MCs with the same number, the MCO uses the last one alphabetically. For example, the `99-worker-ssh-new` MC has precedence over the `99-worker-ssh-last` MC.

Machine Configuration Pools

MCPs use labels to match one or more MCs to one or more nodes. By using multiple MCs, you can split the node configuration and focus every MC on one aspect of server configuration. For example, you can configure one MC for DNS resolution and another MC to synchronize time.

You can also use the same MCs in multiple MCPs. For example, although a time synchronization MC would apply to multiple MCPs, only the MCPs for a specific node pool would include the configuration that a special hardware accelerator card requires.

MCPs specify a `machineConfigSelector` MC label to select one or more MCs, and a `nodeSelector` node label to select one or more nodes.

By default, OpenShift includes MCPs for the `master` and `worker` roles. However, you can add custom MCPs to the cluster. Red Hat recommends creating custom MCPs as a composition of `worker` and custom MCs that are applied to the nodes. Assigning the `worker` role to the custom MCP is critical so that OpenShift applies operating system updates that are labeled as `worker` to the machines in the pool. Thus, the `machineConfigSelector` match expression selects both the `worker` role and the custom label. The nodes that are part of the custom pool use the MCs from the `worker` role with the additions from the custom label.

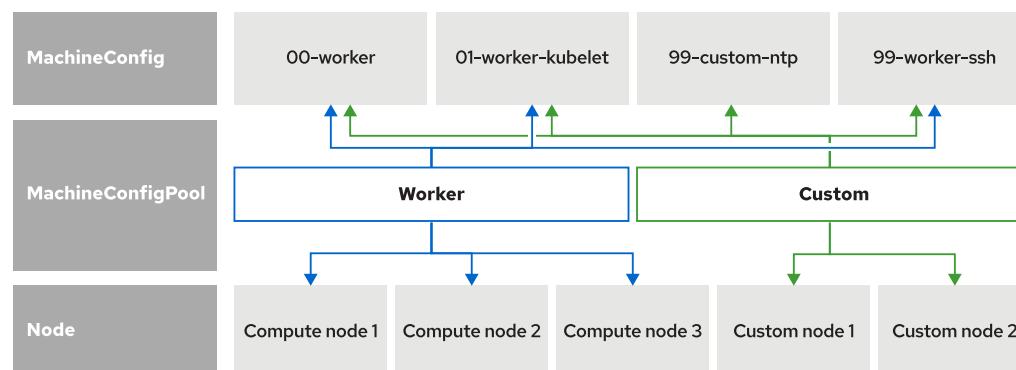


Figure 3.5: Custom pool with the worker MCs and an additional MC

Chapter 3 | Cluster Partitioning

In the previous diagram, the `custom` MCP targets two nodes with the `custom` label, and the `worker` MCP targets three nodes with the `worker` label. The `worker` pool applies to the `worker` nodes the `00-worker`, `01-worker-kubelet`, and `99-worker-ssh` MCs. The `custom` pool uses the MCs from the `worker` role, and adds the `99-custom-ntp` MC to the configuration of the `custom` nodes.



Important

Nodes with the `worker` role can be part of only one custom pool, and nodes with the `master` role cannot be part of a custom pool. In these cases, the MCO does not apply any changes that are specific to the custom pools, and shows an error in the MCC pod logs.

The following MCP specification demonstrates creating a separate pool for `custom` nodes:

```
apiVersion: machineconfiguration.openshift.io/v1
kind: MachineConfigPool
metadata:
  name: custom ①
spec:
  machineConfigSelector:
    matchExpressions:
      - key: machineconfiguration.openshift.io/role
        operator: In
        values: [worker, custom] ②
  nodeSelector:
    matchLabels:
      node-role.kubernetes.io/custom: "" ③
```

- ① The name for the custom MCP.
- ② The MC selector includes both the `worker` and `custom` MCs.
- ③ The node selector includes the nodes in the cluster with the `custom` role.



Note

You can find more information about labels and selectors in <https://kubernetes.io/docs/concepts/overview/working-with-objects/labels/>

You can list the MCPs in your cluster by using the `oc get machineconfigpool` or `oc get mcp` commands.

```
[user@host ~]$ oc get machineconfigpool
NAME      CONFIG          UPDATED   UPDATING   DEGRADED
MACHINECOUNT  READYMACHINECOUNT  UPDATEDMACHINECOUNT  DEGRADEDMACHINECOUNT  AGE
master    rendered-master-716...267  True     False      False
3          3                3          3          3          108d
worker    rendered-worker-573...3db  True     False      False
3          3                3          3          3          108d
```

The previous command shows the following information for the MCPs:

Chapter 3 | Cluster Partitioning**CONFIG**

The name for the rendered MC that is applied to the nodes in the MCP. During the first update, this field remains blank.

UPDATED

The **True** status indicates that the MCO applied the current MC to the nodes in that MCP. The **False** status indicates that nodes in the MCP are updating.

UPDATING

The **True** status indicates that the MCO is applying the intended MC. Nodes in the **UPDATING** state might not be available for scheduling. The **False** status indicates that all nodes in the MCP are updated.

DEGRADED

A **True** status indicates that the MCO is blocked from applying the current or intended MC to at least one node in that MCP. A possible reason is a detection of configuration drift. Configuration drift is explained later in this section. Nodes that are degraded might not be available for scheduling. A **False** status indicates that all nodes in the MCP are ready.

MACHINECOUNT

Indicates the total number of machines in the MCP.

READYMACHINECOUNT

Indicates the total number of machines in the MCP that are ready for scheduling.

UPDATEDMACHINECOUNT

Indicates the total number of machines in the MCP with the current MC.

DEGRADEDMACHINECOUNT

Indicates the total number of machines in that MCP that are degraded or irreconcilable.

Label Nodes

You can add labels to your node by using the `oc label` command.

The following example adds the `custom` role to the `worker03` node in the cluster.

```
[user@host ~]$ oc label node/worker03 node-role.kubernetes.io/custom=
```

```
node/worker03 labeled
```

You can verify the node roles in the cluster by using the `oc get nodes` command:

```
[user@host ~]$ oc get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
master01	Ready	control-plane, master	114d	v1.25.7+eab9cc9
master02	Ready	control-plane, master	114d	v1.25.7+eab9cc9
master03	Ready	control-plane, master	114d	v1.25.7+eab9cc9
worker01	Ready	worker	12d	v1.25.7+eab9cc9
worker02	Ready	worker	12d	v1.25.7+eab9cc9
worker03	Ready	custom , worker	12d	v1.25.7+eab9cc9

Infrastructure Nodes

One important node label is the `infra` role. Use the `infra` role for nodes that host only infrastructure components, such as cluster logging, cluster monitoring, or the integrated container

Chapter 3 | Cluster Partitioning

image registry. Adding the `infra` role for nodes is recommended for larger clusters to ensure the performance and stability of OpenShift cluster services, such as the router or OAuth services, or to prevent the impact of heavy infrastructure components, such as metrics and logging, on user workloads.

Infrastructure nodes do not count towards the total number of required OpenShift subscriptions to run the environment. You can create a custom MCP to apply MCs to the infrastructure nodes.



Note

For more information about creating infrastructure nodes, refer to https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/nodes/index#nodes-nodes-creating-infrastructure-nodes

For more information about using infrastructure nodes to separate maintenance and management, and to prevent incurring billing costs against subscription counts, refer to <https://access.redhat.com/solutions/5034771>

Machine Configuration Operator Updates

The MCD reports the state of the node updates by using node annotations. You can use these annotations to assess the state of the update.

You can list the node annotations by using the `oc describe` command:

```
[user@host ~]$ oc describe node worker01
Name:           worker01
Roles:          worker
Labels:         beta.kubernetes.io/arch=amd64
                ...output omitted...
                node-role.kubernetes.io/worker=
                node.openshift.io/os_id=rhcos
Annotations:   machineconfiguration.openshift.io/currentConfig: rendered-
                worker-370...bfd ①
                machineconfiguration.openshift.io/desiredConfig: rendered-
                worker-370...bfd ②
                machineconfiguration.openshift.io/reason:
                machineconfiguration.openshift.io/state: Done ③
                volumes.kubernetes.io/controller-managed-attach-detach: true
...output omitted...
```

- ① The current rendered MC that is applied to the node.
- ② The intended rendered MC to be applied to the node.
- ③ The current node state regarding MCO.

When the intended configuration does not match the current configuration, the MCD applies the intended rendered MC, drains all the pods from the node, and reboots the node.

Configuration Drift

A *configuration drift* is the state where the configuration on a node does not fully match what the currently applied rendered MC specifies.

Chapter 3 | Cluster Partitioning

The MCD checks for configuration drifts when a node boots, or when any of the specified files in the MC are modified outside the MC, or before a new MC is applied.

When the MCD detects a configuration drift, the MCD performs the following tasks:

- Logs an error message to the console.
- Generates a Kubernetes event.
- Stops additional drift detection on the affected node.
- Marks both the node and the MCP with the degraded state.

The MCO marks the node in the degraded state until an administrator corrects the node configuration. Although a degraded node is online and operational, you cannot update it.

You can correct configuration drift and return the node to the Ready state with one of the following remediations:

Generate

Generate a force file on the degraded node to bypass the configuration drift detection and reapply the current MC. To generate the force file, create a debug pod on the node with the degraded state and create the /run/machine-config-daemon-force file. Then, OpenShift skips the MC validation, restarts the node, and applies the current MC to the node. The force file does not force the node upgrade; it instead skips validation of configurations on the system and attempts an update regardless of the difference. Depending on the issue on your node, skipping the validation process might not help you move past your node error.

Rewrite

Rewrite the file contents or change the file permissions of the files on the node to match the MC configuration. This manual procedure requires you to review the logs and manually fix the conflicting file. This remediation does not require rebooting the node in a degraded state, and thus avoids possible downtime in your applications.

For information about configuration drift, refer to the Status field for the pool with the degraded node:

```
[user@host ~]$ oc describe mcp worker
...output omitted...
Status:
  Conditions:
    ...output omitted...
    Last Transition Time: 2023-10-02T10:11:37Z
    Message:           Node worker01 is reporting: "content mismatch for file '/etc/containers/registries.conf'"
    Reason:            1 nodes are reporting degraded status on sync
    Status:            True
    Type:              NodeDegraded
  Last Transition Time: 2023-10-02T10:11:37Z
  Message:
  Reason:
  Status:            True
  Type:              Degraded
...output omitted...
```

In the previous example, the MCO detects a configuration drift for the /etc/containers/registries.conf file.

You can get the MCD pod name for a specific node by using a filter.

```
[user@host ~]$ oc get pod -n openshift-machine-config-operator \
--field-selector spec.nodeName=worker01
NAME                  READY   STATUS    RESTARTS   AGE
machine-config-daemon-jsrzm   2/2     Running   2          19d
```

You can also review the logs for the MCD that gives you more information about the configuration drift.

```
[user@host ~]$ oc logs machine-config-daemon-jsrzm \
-n openshift-machine-config-operator
...output omitted...
E1002 10:11:36.163676    2667 daemon.go:589] Preflight config drift check failed:
content mismatch for file "/etc/containers/registries.conf"
E1002 10:11:36.163692    2667 writer.go:200] Marking Degraded due to: content
mismatch for file "/etc/containers/registries.conf"
W1002 10:11:40.193224    2667 daemon.go:1763] current+desiredConfig is rendered-
worker-d4f45006b2d83725d98af944c8296774 but state is Degraded
I1002 10:11:40.510208    2667 rpm-ostree.go:394] Running captured: rpm-ostree
kargs
E1002 10:11:40.568388    2667 on_disk_validation.go:207] content mismatch for file
"/etc/containers/registries.conf" (-want +got):
[ ]uint8(
"""
-
- unqualified-search-registries = ["registry.access.redhat.com", "docker.io"]
+ unqualified-search-registries = ["docker.io"]
short-name-mode = ""

...
// 107 identical lines
"""
)
E1002 10:11:40.568446    2667 daemon.go:589] Preflight config drift check failed:
content mismatch for file "/etc/containers/registries.conf"
E1002 10:11:40.568467    2667 writer.go:200] Marking Degraded due to: content
mismatch for file "/etc/containers/registries.conf"
E1002 10:11:46.987190    2667 daemon.go:1077] mode mismatch for file: "/
etc/containers/registries.conf"; expected: -rw-r--r--/420/0644; received: -
rwxrwxrwx/511/0777
E1002 10:11:46.987222    2667 writer.go:200] Marking Degraded due to: mode
mismatch for file: "/etc/containers/registries.conf"; expected: -rw-r--r--/420/0644;
received: -rwxrwxrwx/511/0777
...output omitted...
```

In the previous example, the MCD marks the `worker01` node in a degraded state due to the mismatches of content and permissions for the file.

Configure MCO-related Custom Resources

The MCO provides the `ContainerRuntimeConfig` CR to modify CRI-O container runtime settings, and the `KubeletConfig` CR to manage the Kubelet service. You can use these CRs to configure a subset of CRI-O and Kubelet configuration parameters. Always use valid values for the configuration parameter, because invalid values might render cluster nodes unusable.

Chapter 3 | Cluster Partitioning

Although you can modify the CRI-O container runtime settings and the Kubelet service by using the MachineConfig CR, using either of these two special CRs simplifies node deployment and configuration management, provides API checking, and prevents misconfigurations. Moreover, because OpenShift does not support changing all the settings of the Kubelet service and the container runtime, these CRs provide only the configuration changes that OpenShift supports.

Create a Custom Resource for Kubelet Configuration

OpenShift provides a kubelet configuration controller to the MCC. You can use the KubeletConfig CR to edit the kubelet parameters.

The MCO can write the `kubelet.conf` configuration file and the `kubelet.systemd` unit file to Ignition, so Ignition writes these two files to configure the `kubelet` agent when it starts on a node. If you create an MC to change the `kubelet` parameters, then the MCD reboots the nodes to write the new configuration. With this approach, OpenShift can restore the default `kubelet` configuration if you delete the KubeletConfig instance.



Note

For a list of all the parameters that you can modify by using the KubeletConfig CR, you can refer to the `KubeletConfiguration` API object in Kubernetes that uses the same parameters. Refer to <https://kubernetes.io/docs/reference/config-api/kubelet-config.v1beta1/#kubelet-config-k8s-io-v1beta1-KubeletConfiguration>

Create a Custom Resource for Container Runtime Configuration

You can change the settings for the OpenShift CRI-O runtime by using the ContainerRuntimeConfig CR. The MCO can write the `crio.conf` and `storage.conf` configuration files on the associated nodes with the updated values.

You can modify the following parameters by using a ContainerRuntimeConfig CR:

Logging level

The `LogLevel` parameter sets the level of verbosity for logging messages. The default level is `info`. Other options include the `fatal`, `panic`, `error`, `warn`, `debug`, and `trace` options.

Overlay size

The `overlaySize` parameter sets the maximum size of a container image.

Container runtime

The `defaultRuntime` parameter sets the container runtime to either `runc` (the default) or `crun`.



Important

Support for the `crun` container runtime is a Technology Preview feature only. Technology Preview features are not supported with Red Hat production service level agreements and might not be functionally complete. Red Hat does not recommend using Technology Preview features in production.

You can also use the ContainerRuntimeConfig CR to change the limit of PIDs or the maximum logging size. However, Red Hat recommends using the KubeletConfig CR to change these parameters, because they will likely be deprecated in a future version.



References

For more information about the Machine Configuration Operator, refer to the *Post-installation Machine Configuration Tasks* section in the Red Hat OpenShift Container Platform 4.14 *Post-installation Configuration* documentation at
https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/postinstallation_configuration/post-install-machine-configuration-tasks

How does Machine Config Pool work?

<https://www.redhat.com/en/blog/openshift-container-platform-4-how-does-machine-config-pool-work>

► Guided Exercise

Node Configuration with the Machine Configuration Operator

Label a cluster node and apply a machine configuration to it, without affecting other cluster nodes.

Outcomes

- Use the machine configuration operator (MCO) to modify the Network Time Protocol (NTP) client configuration.
- Create a machine configuration pool (MCP) and a machine configuration (MC).
- Observe the status of MC updates.

Before You Begin

As the student user on the `workstation` machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start nodes-mco
```

Instructions

Your company requires you to configure an NTP server for the nodes in the cluster to synchronize services among them. The new NTP server is `0.rhel.pool.ntp.org`.

In production environments, it is common to change the default NTP configuration to refer to NTP servers in the corporate LAN, instead of to public NTP servers from the internet, to prevent clock skew. However, in this exercise, you apply a new NTP configuration only to the `worker01` node, to eliminate waiting time for node reboots when applying MCs, and to show how to apply MCs only to a subset of nodes by using labels.

To configure the NTP server in the `worker01` node, create the custom MCP. The custom MCP must apply the MCs with the `worker` and `custom` labels to the nodes with the `custom` role. Then, create an MC for the NTP configuration with the `custom` label.

- ▶ 1. Connect to the OpenShift cluster and verify the roles for the nodes in the cluster. List the MCs that apply to the `worker` and `custom` roles.
 - 1.1. Connect to the OpenShift cluster as the `admin` user with `redhatocp` as the password.

```
[student@workstation ~]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
Login successful.
...output omitted...
```

- 1.2. Change to the `nodes-mco` project.

Chapter 3 | Cluster Partitioning

```
[student@workstation ~]$ oc project nodes-mco
Now using project "nodes-mco" on server "https://api.ocp4.example.com:6443".
```

- 1.3. Verify the nodes in the cluster and their roles.

```
[student@workstation ~]$ oc get nodes
NAME      STATUS    ROLES          AGE     VERSION
master01   Ready     control-plane, master   109d    v1.27.10+28ed2d7
master02   Ready     control-plane, master   109d    v1.27.10+28ed2d7
master03   Ready     control-plane, master   109d    v1.27.10+28ed2d7
worker01   Ready     worker           21d    v1.27.10+28ed2d7
worker02   Ready     worker           21d    v1.27.10+28ed2d7
worker03   Ready     worker           21d    v1.27.10+28ed2d7
```

- 1.4. List the MCs that apply to the `worker` roles.

```
[student@workstation ~]$ oc get machineconfig \
-l machineconfiguration.openshift.io/role=worker
NAME                  GENERATEDBYCONTROLLER  IGNITIONVERSION  AGE
00-worker              52fe2...ebf97        3.4.0          12d
01-worker-container-runtime 52fe2...ebf97        3.4.0          12d
01-worker-kubelet       52fe2...ebf97        3.4.0          12d
97-worker-generated-kubelet 52fe2...ebf97        3.4.0          12d
98-worker-generated-kubelet 52fe2...ebf97        3.4.0          12d
99-worker-chrony-conf-override 52fe2...ebf97        3.2.0          12d
99-worker-generated-registries 52fe2...ebf97        3.4.0          12d
99-worker-ssh            52fe2...ebf97        3.2.0          12d
```

- 1.5. List the MCs that apply to the `custom` roles.

```
[student@workstation ~]$ oc get machineconfig \
-l machineconfiguration.openshift.io/role=custom
No resources found
```

- 2. Label the `worker01` node with the `node-role.kubernetes.io/custom` label. Verify that the node has the correct label.

- 2.1. Apply the `node-role.kubernetes.io/custom` label to the `worker01` node.

```
[student@workstation ~]$ oc label node/worker01 node-role.kubernetes.io/custom=
node/worker01 labeled
```

- 2.2. Verify that the node has the correct label.

Chapter 3 | Cluster Partitioning

```
[student@workstation ~]$ oc get nodes
NAME      STATUS    ROLES          AGE     VERSION
master01   Ready     control-plane,master  33d    v1.27.10+28ed2d7
master02   Ready     control-plane,master  33d    v1.27.10+28ed2d7
master03   Ready     control-plane,master  33d    v1.27.10+28ed2d7
worker01   Ready     custom,worker       12d    v1.27.10+28ed2d7
worker02   Ready     worker           12d    v1.27.10+28ed2d7
worker03   Ready     worker           12d    v1.27.10+28ed2d7
```

- 3. Create an MCP for the node with the `custom` label. The MCP must include the MCs for the `worker` and `custom` labels.

- 3.1. Change to the `~/D0380/labs/nodes-mco` directory.

```
[student@workstation ~]$ cd ~/D0380/labs/nodes-mco
```

- 3.2. Create the MCP custom resource (CR) YAML file. You can find an example for the CR in the `~/D0380/labs/nodes-mco/custom-mcp.yaml` file.

```
apiVersion: machineconfiguration.openshift.io/v1
kind: MachineConfigPool
metadata:
  name: custom
spec:
  machineConfigSelector:
    matchExpressions:
      - key: machineconfiguration.openshift.io/role
        operator: In
        values: [worker,custom]
  nodeSelector:
    matchLabels:
      node-role.kubernetes.io/custom: ""
```

- 3.3. Apply the configuration for the MCP CR.

```
[student@workstation nodes-mco]$ oc create -f custom-mcp.yaml
machineconfigpool.machineconfiguration.openshift.io/custom created
```

- 3.4. Verify that the `custom` MCP is correctly created. If the `UPDATING` field is marked as true for the `custom` MCP, then wait until the MCO updates the node.

```
[student@workstation nodes-mco]$ oc get mcp
NAME      CONFIG          UPDATED  UPDATING  DEGRADED  MACHINECOUNT ...
custom    rendered-custom... True     False     False     1          ...
master    rendered-master... True     False     False     3          ...
worker    rendered-worker... True     False     False     2          ...
```

- 3.5. Verify that the `worker01` node is assigned for the `custom` MCP, and that the MCO applies the rendered MC to the node.

```
[student@workstation nodes-mco]$ oc get events -n \
  openshift-machine-config-operator \
  --sort-by='{"lastTimestamp"}' \
  --field-selector involvedObject.name=custom
...  REASON          OBJECT            MESSAGE
...  RenderedConfigGenerated  machineconfigpool/custom  rendered-
custom-2f56...81d0 successfully generated ...
...  SetDesiredConfig   machineconfigpool/custom  Targeted node worker01 to
MachineConfig rendered-custom-2f56...81d0
```

- ▶ 4. Because an OpenShift preconfigured MC configures the NTP client, verify the name for the MCs that modify the `chrony` configuration for the nodes with the `worker` label. Create an MC for the `chrony` configuration that applies to the `custom` label, with higher precedence than the preconfigured MC. The name for the MC must be `99-zcustom-chrony` to have higher precedence.
- 4.1. Verify the names for the MCs that modify the `chrony` configuration for the nodes with the `worker` label.

```
[student@workstation nodes-mco]$ oc get mc --selector \
  machineconfiguration.openshift.io/role=worker \
  -o jsonpath="{range .items[*]}{.metadata.name} \
  {'': '{range .spec.config.storage.files[*]}{.path} \
  {'': '\n'}{end}}{\n}{end}" | grep chrony
99-worker-chrony-conf-override: /etc/chrony.conf
```

- 4.2. Create a `chrony` configuration file. You can find an example for the CR in the `~/D0380/labs/nodes-mco/chrony-mod.conf` file.

```
pool 0.rhel.pool.ntp.org iburst
driftfile /var/lib/chrony/drift
makestep 1.0 3
rtcsync
logdir /var/log/chrony
```

- 4.3. Encode the `chrony` configuration file in Base64 format. You use the encoded configuration file in a later step. The encoded configuration file would differ on your system.

```
[student@workstation nodes-mco]$ base64 -w0 chrony-mod.conf; echo
cG9vbCAwLnJoZWhUCG9...C92YXIVbG9nL2Nocm9ueQo=
```

- 4.4. Create the `99-zcustom-chrony` MC CR YAML file. Use the encoded `chrony` configuration file from a previous step. You can find an example for the CR in the `~/D0380/labs/nodes-mco/99-zcustom-chrony.yaml` file.

```
apiVersion: machineconfiguration.openshift.io/v1
kind: MachineConfig
metadata:
  labels:
    machineconfiguration.openshift.io/role: custom
```

Chapter 3 | Cluster Partitioning

```
name: 99-zcustom-chrony
spec:
  config:
    ignition:
      version: 3.2.0
    storage:
      files:
        - contents:
            compression: ""
            source: data:;base64,cG9vbCAwLnJ0ZWwucG9...C92YXIVbG9nL2Nocm9ueQo=
            mode: 420
            overwrite: true
            path: /etc/chrony.conf
```

4.5. Apply the configuration for the MC CR.

```
[student@workstation nodes-mco]$ oc create -f 99-zcustom-chrony.yml
machineconfig.machineconfiguration.openshift.io/99-zcustom-chrony created
```

4.6. Verify that the 99-zcustom-chrony MC is correctly created.

```
[student@workstation nodes-mco]$ oc get mc
NAME                      GENERATEDBYCONTROLLER IGNITIONVERSION AGE
...output omitted...
99-master-generated-registries 52fe26136643a946ff... 3.2.0       109d
99-master-ssh                3.2.0       109d
99-worker-chrony-conf-override 3.2.0       109d
99-worker-generated-registries 52fe26136643a946ff... 3.2.0       109d
99-worker-ssh                3.2.0       109d
99-zcustom-chrony           3.2.0       2m5s
...output omitted...
```

4.7. Verify that the MCP starts updating the resources.

```
[student@workstation nodes-mco]$ oc get mcp
NAME      CONFIG          UPDATED  UPDATING  DEGRADED  MACHINECOUNT ...
custom   rendered-custom-... False    True     False     1         ...
master    rendered-master-... True     False     False     3         ...
worker    rendered-worker-... True     False     False     2         ...
```

4.8. Verify that the MCO applies the rendered MC to the worker01 node.

Chapter 3 | Cluster Partitioning

```
[student@workstation nodes-mco]$ oc get events \
-n openshift-machine-config-operator \
--sort-by='{"lastTimestamp"}' \
--field-selector involvedObject.name=custom
... REASON OBJECT MESSAGE
... RenderedConfigGenerated machineconfigpool/custom rendered-
custom-2f56...81d0 successfully generated ...
... SetDesiredConfig machineconfigpool/custom Targeted node worker01 to
MachineConfig rendered-custom-2f56...81d0
... RenderedConfigGenerated machineconfigpool/custom rendered-
custom-009f...3e5d successfully generated ...
... SetDesiredConfig machineconfigpool/custom Targeted node worker01 to
MachineConfig rendered-custom-009f...3e5d
```

4.9. Wait a few minutes until the MCO finishes the update.

[student@workstation nodes-mco]\$ oc get mcp						
NAME	CONFIG	UPDATED	UPDATING	DEGRADED	MACHINECOUNT	...
custom	rendered-custom-...	True	False	False	1	...
master	rendered-master-...	True	False	False	3	...
worker	rendered-worker-...	True	False	False	2	...

- ▶ 5. Verify that the MCO applies the chrony configuration to the worker01 node, and that the worker02 and worker03 nodes remain unaltered.

5.1. Create a debug pod for the worker01 node.

```
[student@workstation nodes-mco]$ oc debug node/worker01
Temporary namespace openshift-debug-vlrjg is created for debugging node...
Starting pod/worker01-debug ...
To use host binaries, run chroot /host
Pod IP: 192.168.50.13
If you don't see a command prompt, try pressing enter.
sh-4.4#
```

5.2. Run the chroot /host command to use host binaries.

```
sh-4.4# chroot /host
```

5.3. Verify the content for the chrony configuration file. The MC correctly applies the changes to the /etc/chrony.conf file.

```
sh-4.4# cat /etc/chrony.conf
pool 0.rhel.pool.ntp.org iburst
driftfile /var/lib/chrony/drift
makestep 1.0 3
rtcsync
logdir /var/log/chrony
```

5.4. Remove the debug pod for the worker01 node.

Chapter 3 | Cluster Partitioning

```
sh-4.4# exit ①
exit
sh-4.4# exit ②
exit

Removing debug pod ...
Temporary namespace openshift-debug-vlrjg was removed.
```

- ① This command exit the chroot environment.
- ② This command exit the debug pod.

5.5. Create a debug pod for the worker02 node.

```
[student@workstation nodes-mco]$ oc debug node/worker02
...output omitted...
```

5.6. Run the chroot /host command to use host binaries.

```
sh-4.4# chroot /host
```

5.7. Verify the content for the chrony configuration file. The MC does not change the content in the /etc/chrony.conf file.

```
sh-4.4# cat /etc/chrony.conf
# The Machine Config Operator manages this file.
server classroom.example.com iburst

stratumweight 0
driftfile /var/lib/chrony/drift
rtcsync
makestep 10 3
bindcmdaddress 127.0.0.1
bindcmdaddress ::1
...output omitted...
```

5.8. Remove the debug pod for the worker02 node.

```
sh-4.4# exit ①
exit
sh-4.4# exit ②
exit
...output omitted...
```

- ① This command exit the chroot environment.
- ② This command exit the debug pod.

► 6. Manually modify the /etc/chrony.conf file on the worker01 node, and verify that the MCO detects a mismatch in the chrony configuration file.

6.1. Create a debug pod for the worker01 node.

Chapter 3 | Cluster Partitioning

```
[student@workstation nodes-mco]$ oc debug node/worker01
...output omitted...
```

- 6.2. Run the `chroot /host` command to use host binaries.

```
sh-4.4# chroot /host
```

- 6.3. Modify the content for the `/etc/chrony.conf` configuration file as follows. If the vi editor does not show the first line, then resize the window to show it.

```
pool 1.rhel.pool.ntp.org iburst
driftfile /var/lib/chrony/drift
makestep 1.0 3
rtcsync
logdir /var/log/chrony
```

- 6.4. Remove the debug pod for the `worker01` node.

```
sh-4.4# exit ①
exit
sh-4.4# exit ②
exit
...output omitted...
```

- ① This command exit the `chroot` environment.
- ② This command exit the debug pod.

- 6.5. Verify that the MCO marks the `worker01` node with the degraded state.

```
[student@workstation nodes-mco]$ oc get mcp
NAME      CONFIG          UPDATED   UPDATING  DEGRADED  MACHINECOUNT ...
custom    rendered-custom-...  False     True      True       1
master    rendered-master-...  True      False     False      3
worker    rendered-worker-...  True      False     False      2
```

- 6.6. Verify the information for the `custom` MCP, and that the MCO detects a mismatch in the `chrony` configuration file.

```
[student@workstation nodes-mco]$ oc describe mcp custom
...output omitted...
Status:
  Conditions:
    ...output omitted...
    Last Transition Time: 2023-09-20T11:03:53Z
    Message:           Node worker01 is reporting: "unexpected on-disk state
validating against rendered-custom-009f7b72211eb0d37a8ba1d03a343e5d: content
mismatch for file '/etc/chrony.conf'"
    Reason:            1 nodes are reporting degraded status on sync
    Status:            True
    Type:              NodeDegraded
```

Chapter 3 | Cluster Partitioning

```
Last Transition Time: 2023-09-20T11:03:53Z
Message:
Reason:
Status: True
Type: Degraded
...output omitted...
```

- 6.7. Get the name of the MCO daemon pod. The name for the daemon pod might be different on your system.

```
[student@workstation nodes-mco]$ oc get pod -n openshift-machine-config-operator \
--field-selector spec.nodeName=worker01
NAME                   READY   STATUS    RESTARTS   AGE
machine-config-daemon-jsrzm   2/2     Running   4          16d
```

- 6.8. Review the logs for the MCO daemon for more information about the configuration drift.

```
[student@workstation nodes-mco]$ oc logs machine-config-daemon-jsrzm \
-n openshift-machine-config-operator
...output omitted...
W0929 09:23:20.539413    2520 daemon.go:1763] current+desiredConfig is rendered-
custom-11677dc3a3e84870c3a359ab1be4eafc but state is Degraded
I0929 09:23:20.858668    2520 rpm-ostree.go:394] Running captured: rpm-ostree
kargs
E0929 09:23:20.905408    2520 on_disk_validation.go:207] content mismatch for file
"/etc/chrony.conf" (-want +got):
bytes.Join({
  "pool ",
- "0",
+ "1",
  ".rhel.pool.ntp.org iburst\n driftfile /var/lib/chrony/drift\n makest",
  "ep 1.0 3\n rtsync\n logdir /var/log/chrony\n",
}, "")
```

E0929 09:23:20.905438 2520 daemon.go:589] Preflight config drift check failed:
content mismatch for file "/etc/chrony.conf"

E0929 09:23:20.905457 2520 writer.go:200] Marking Degraded due to: content
mismatch for file "/etc/chrony.conf"

- 6.9. Create a debug pod for the worker01 node.

```
[student@workstation nodes-mco]$ oc debug node/worker01
...output omitted...
sh-4.4# chroot /host
```

- 6.10. Fix the /etc/chrony.conf file.

```
pool 0.rhel.pool.ntp.org iburst
driftfile /var/lib/chrony/drift
makestep 1.0 3
rtcsync
logdir /var/log/chrony
```

Chapter 3 | Cluster Partitioning

- 6.11. Exit the debug pod.

```
sh-4.4# exit  
exit  
sh-4.4# exit  
exit  
...output omitted...
```

- 6.12. Verify that the MCO marks the node with the `ready` state, which might take a few minutes.

```
[student@workstation nodes-mco]$ oc get mcp  
NAME      CONFIG          UPDATED  UPDATING  DEGRADED  MACHINECOUNT ...  
custom    rendered-custom-...  True     False     False     1         ...  
master    rendered-master-...  True     False     False     3         ...  
worker    rendered-worker-...  True     False     False     2         ...
```

- 7. Remove the `node-role.kubernetes.io/custom` label for the `worker01` node. After the MCO updates the node with the new configuration, delete the `99-zcustom-chrony` MC and the `custom-ntp` MCP.

- 7.1. Remove the `node-role.kubernetes.io/custom` label from the `worker01` node.

```
[student@workstation nodes-mco]$ oc label node/worker01 \  
  node-role.kubernetes.io/custom-  
node/worker01 unlabeled
```

- 7.2. Verify that no node uses the `node-role.kubernetes.io/custom` label.

```
[student@workstation nodes-mco]$ oc get nodes -l node-role.kubernetes.io/custom=  
No resources found
```

- 7.3. Verify that the MCO applies the changes to the `worker01` node. Wait until the MCO updates the node. The MCO finishes updating the node when the `desiredConfig` and `currentConfig` annotations point to the same rendered MC, and the `state` annotation is `Done`.

```
[student@workstation nodes-mco]$ oc describe node worker01 \  
| grep machineconfiguration  
machineconfiguration.openshift.io/controlPlaneTopology: HighlyAvailable  
machineconfiguration.openshift.io/currentConfig: rendered-custom-1167...eafc  
machineconfiguration.openshift.io/desiredConfig: rendered-worker-d4f4...6774  
machineconfiguration.openshift.io/desiredDrain: drain-rendered-worker-...  
machineconfiguration.openshift.io/lastAppliedDrain: drain-rendered-worker-...  
machineconfiguration.openshift.io/reason:  
machineconfiguration.openshift.io/state: Working
```

- 7.4. Verify that the MCP updated status is set to `True`.

Chapter 3 | Cluster Partitioning

```
[student@workstation nodes-mco]$ oc get mcp
NAME        CONFIG          UPDATED  UPDATING  DEGRADED  MACHINECOUNT ...
custom      rendered-custom...  True     False     False      0       ...
master      rendered-master...  True     False     False      3       ...
worker      rendered-worker...  True     False     False      3       ...
```

7.5. Delete the 99-zcustom-chrony MC.

```
[student@workstation nodes-mco]$ oc delete mc 99-zcustom-chrony
machineconfig.machineconfiguration.openshift.io "99-zcustom-chrony" deleted
```

7.6. Delete the custom MCP.

**Warning**

Remove the custom MCP only after the MCO applies the changes to the worker01 node.

If you remove the custom MCP before the MCO applies the changes to the node, then the MCP gets stuck in a degraded state, because the currentConfig MC of the node does not exist. Thus, none of the guided exercises from this point on will work.

To solve this issue, either recreate your entire cluster or apply the solution in <https://access.redhat.com/solutions/4970731>

```
[student@workstation nodes-mco]$ oc delete mcp custom
machineconfigpool.machineconfiguration.openshift.io "custom" deleted
```

7.7. Change to the student HOME directory.

```
[student@workstation nodes-mco]$ cd
```

Finish

On the workstation machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish nodes-mco
```

Node Configuration with Special Purpose Operators

Objectives

- Apply operating system settings to cluster nodes with higher-level operators instead of the low-level machine configuration operator.

Special Purpose Operators

The Red Hat Enterprise Linux CoreOS (RHCOS) operating system that runs on each node in an OpenShift cluster is not administered with the traditional methods for Red Hat Enterprise Linux. To adapt to the paradigm shift of using RHCOS as an infrastructure operating system, learning the methods for managing RHCOS is a large part of gaining the required knowledge to manage OpenShift clusters.

Although the Machine Configuration Operator (MCO) can handle many of the preferred customizations for the operating system that runs on your nodes, other special purpose operators are available for higher-level configurations, such as to focus on specialized hardware tuning, cluster organization, and advanced optimizations.

OperatorHub is a library of shared operators that provides a curated catalog for identifying and installing available operators in a cluster. The CLI and web console approaches are available for installing these special purpose operators, including those operators that are not available in OperatorHub.

Each operator runs in a designated namespace, and uses CRDs and CRs to define the implementation. An operator's function might also extend to actions across other namespaces. Red Hat recommends running each special purpose operator within its own namespace and using a service account with the appropriate permissions. Although operators typically configure the namespace and service account as part of the installation process, some operators might require manual setup.

The Node Tuning Operator

The Node Tuning Operator is a cluster operator that manages node-level optimizations by using the TuneD service and the tuned daemon. The operator manages the containerized tuned daemon for RHOCP as a Kubernetes daemon set. Creating custom TuneD profiles can enable granular customizations to cluster nodes. This approach helps the cluster nodes provide ideal environments for the running applications.

The Node Tuning Operator is installed by default in RHOCP clusters, and runs in the `openshift-cluster-node-tuning-operator` namespace. You can view the default TuneD profile for the cluster by running the following command:

```
[user@host ~]$ oc get Tuned/default -o yaml \
-n openshift-cluster-node-tuning-operator
```

You can view the applied TuneD profiles for each node with the following command:

```
[user@host ~]$ oc get profile -n openshift-cluster-node-tuning-operator
```

When customizing a cluster node, create a TuneD profile and specify the optimized parameters from a TuneD plug-in. Each of the TuneD plug-ins, such as `cpu`, `sysctl`, or `vm`, provides various customizations for the cluster nodes. Consult the TuneD documentation from Red Hat to explore the full list of TuneD profiles.

**Note**

For more information about TuneD plug-ins, refer to https://access.redhat.com/documentation/en-us/red_hat_enterprise_linux/8/html-single/monitoring_and_managing_system_status_and_performance/index#available-tuned-plug-ins_customizing-tuned-profiles

The Node Feature Discovery Operator

Node Feature Discovery (NFD) is a Kubernetes add-on in OperatorHub for installation, to expose node-level information. The NFD operator detects hardware details, such as processor architecture, specialized PCI hardware such as GPUs or networking cards, and other system configurations on your cluster nodes. The NFD operator runs inside the `openshift-nfd` namespace for these operations.

**Note**

Due to the minimal environment within a classroom cluster, the special purpose operators might not be available in OperatorHub in this course.

The operator profiles each of the cluster nodes and collects information about the hardware and system attributes. This operator collects information about connected PCI devices, CPU specifications, kernel information, the operating system version, and other hardware information. The operator adds labels to the cluster nodes with the gathered information, to aid deployment placement onto nodes that provide the required specialized hardware and environments.

Consider a cluster that added high-throughput networking cards to a subset of its nodes. After the NFD operator profiles and labels these nodes, applications that require the higher network throughput can target the labeled nodes with this hardware.

The Kernel Module Management Operator

The Kernel Module Management (KMM) Operator, which is available for installation in OperatorHub, manages, builds, signs, and deploys out-of-tree kernel modules and device plug-ins on OpenShift Container Platform clusters. This operator is useful for adding extra, alternative, or custom modules to the nodes in your cluster that run workloads that use those features.

Many applications, especially those applications that adopt later technologies before a full integration into modern operating systems is available, can require customizing the kernel modules for the cluster nodes. Developers who use Artificial Intelligence (AI) and Machine Learning (ML) applications often author custom kernel modules to optimize the computational operations.

**Note**

After deployment, kernel modules can run any type of operation. For this reason, carefully consider security implications and permissions when adding kernel modules.

You can install the KMM by searching in OperatorHub to add the operator to a cluster. The KMM uses the `openshift-kmm` namespace to manage kernel module deployments. A Module CRD specifies the kernel that can be added to one or more cluster nodes. A corresponding Module CR applies this CRD to install the required kernel version for each cluster node.

Vendor-provided Operators

Many hardware vendors author special purpose operators for OpenShift to provide unique functions that correspond to a particular product that the company sells. Companies such as IBM, Nvidia, and Intel author and distribute these operators to help cluster administrators in using the specialized vendor hardware that is installed on cluster nodes. Each of these special purpose operators can vary in availability, and might or might not be available in OperatorHub. When adding hardware to cluster nodes, consider vendors who provide OpenShift operators that help with the management and function of cluster nodes.



References

For more information about the Node Tuning Operator, refer to the *Node Tuning Operator* section in the Red Hat OpenShift Container Platform 4.14 *Working with Nodes* documentation at
https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/nodes/index#nodes-node-tuning-operator

For more information about the Node Feature Discovery Operator, refer to the *Node Feature Discovery Operator* section in the Red Hat OpenShift Container Platform 4.14 *Specialized Hardware and Driver Enablement* documentation at
https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/specialized_hardware_and_driver_enablement/index#node-feature-discovery-operator

For more information about the Kernel Module Management Operator, refer to the *Kernel Module Management Operator* section in the Red Hat OpenShift Container Platform 4.14 *Specialized Hardware and Driver Enablement* documentation at
https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/specialized_hardware_and_driver_enablement/index#kernel-module-management-operator

For more information about TuneD, refer to the *Getting Started with TuneD* section in the Red Hat Enterprise Linux 8 *Monitoring and Managing System Status and Performance* documentation at
https://access.redhat.com/documentation/en-us/red_hat_enterprise_linux/8/html-single/monitoring_and_managing_system_status_and_performance/index

► Guided Exercise

Node Configuration with Special Purpose Operators

Label a cluster node and apply a TuneD configuration to it, without affecting other cluster nodes.

Outcomes

Use the Node Tuning Operator to create a TuneD profile that configures a compute node setting for optimizing a database workload.

Before You Begin

As the student user on the workstation machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start nodes-operators
```

Instructions

The software vendor for the database that is used in a company application recommends adjusting a specific setting to improve performance. The suggestion is to disable the `transparent_hugepages` setting, which is enabled by default, on the compute node where the database runs, to improve memory management.

Your company requests that you configure an RHOCP cluster node for the database that requires this memory management adjustment. Create a TuneD profile through the Node Tuning Operator, which configures the `worker01` node to run the database with the customization for the node memory management, which the software vendor recommends.

Verify the initial setting and updated value on the node by inspecting the value in the `/sys/kernel/mm/transparent_hugepage/enabled` file.

- 1. Connect to the OpenShift cluster and verify the roles for the nodes in the cluster.

- 1.1. Connect to the OpenShift cluster as the `admin` user with `redhatocp` as the password.

```
[student@workstation ~]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
Login successful.
...output omitted...
```

- 2. Verify that the `worker01` node has the `stabledb` label.

Chapter 3 | Cluster Partitioning

```
[student@workstation ~]$ oc get nodes
NAME      STATUS    ROLES          AGE     VERSION
master01   Ready     control-plane,master  47d    v1.25.7+eab9cc9
master02   Ready     control-plane,master  47d    v1.25.7+eab9cc9
master03   Ready     control-plane,master  47d    v1.25.7+eab9cc9
worker01   Ready     stabledb,worker   16h    v1.25.7+eab9cc9
worker02   Ready     worker           16h    v1.25.7+eab9cc9
worker03   Ready     worker           16h    v1.25.7+eab9cc9
```

- 3. View the pods that are running in the `openshift-cluster-node-tuning-operator` namespace. View the existing TuneD profiles and configuration for the cluster.

3.1. View the running pods in the namespace.

```
[student@workstation ~]$ oc get pods -n \
  openshift-cluster-node-tuning-operator
NAME                           READY   STATUS    RESTARTS   AGE
cluster-node-tuning-operator-... 1/1     Running   2          47d
tuned-45hlf                     1/1     Running   1          18h
tuned-glr54                      1/1     Running   1          18h
...output omitted...
```

3.2. View the current tuned profiles in use for each cluster node.

```
[student@workstation ~]$ oc get profile -n openshift-cluster-node-tuning-operator
NAME      TUNED          APPLIED   DEGRADED   AGE
master01  openshift-control-plane  True      False      47d
master02  openshift-control-plane  True      False      47d
master03  openshift-control-plane  True      False      47d
worker01  openshift-node          True      False      20h
worker02  openshift-node          True      False      20h
worker03  openshift-node          True      False      20h
```

3.3. View the TuneD configuration in the cluster.

```
[student@workstation ~]$ oc get tuneds.tuned.openshift.io \
  -n openshift-cluster-node-tuning-operator -o yaml
apiVersion: v1
items:
- apiVersion: tuned.openshift.io/v1
  kind: Tuned
  metadata:
    creationTimestamp: "2023-10-04T10:12:26Z"
    generation: 1
    name: default
    namespace: openshift-cluster-node-tuning-operator
    resourceVersion: "8714"
    uid: a4b818f2-dab5-4091-9beb-c763d216a75a
  spec:
    profile:
      - data: |
          [main]
```

```

summary=Optimize systems running OpenShift (provider specific parent
profile)
include=-provider-${f:exec:cat:/var/lib/tuned/provider},openshift
name: openshift
recommend:
- match:
  - label: node-role.kubernetes.io/master
  - label: node-role.kubernetes.io/infra
operand:
tunedConfig: {}
priority: 30
...output omitted...

```

- ▶ 4. Create a debug pod to connect to the `worker01` node and inspect the current setting for the `transparent_hugepages` setting.
- 4.1. Open a new terminal window and create a debug pod to connect to the `worker01` node.

```
[student@workstation ~]$ oc debug node/worker01
...output omitted...
sh-4.4#
```

- 4.2. View the current `transparent_hugepages` setting by displaying the value in the `setting enabled` file. The `[always]` indicator denotes that the setting is always enabled.

```
sh-4.4# cat /sys/kernel/mm/transparent_hugepage/enabled
[always] madvise never
```

- 4.3. Log out and remove the debug pod for the `worker01` node.

```
sh-4.4# exit
...output omitted...
```

- ▶ 5. Navigate to the exercise directory and create a Tuned profile CR file named `stabledbCR.yaml` that defines a Tuned profile `stabledb-tuning` to disable the `transparent_hugepage` setting for nodes with the `stabledb` label. You can find an incomplete example for the CR in the `~/D0380/labs/nodes-operators/stabledbCR.yaml` file.

- 5.1. Change to the working directory for the exercise.

```
[student@workstation ~]$ cd ~/D0380/labs/nodes-operators
```

- 5.2. Create the `stabledbCR.yaml` file and define the resource kind as `Tuned` for the metadata name of `stabledb-tuning`.

Chapter 3 | Cluster Partitioning

```
apiVersion: tuned.openshift.io/v1
kind: Tuned
metadata:
  name: stabledb-tuning
  namespace: openshift-cluster-node-tuning-operator
```

- 5.3. Specify the profile name as **stabledb-tuning**. You can add an optional description in the summary key for the profile.

```
spec:
  profile:
    - name: stabledb-tuning
      data: |
        [main]
        summary=Worker optimization disabling transparent_hugepages.
```

- 5.4. Supply the arguments for the customization that disables the `transparent_hugepages` value.

```
[vm]
transparent_hugepage=never
```

- 5.5. Add the `recommend` section to the CR to target the node with the `stabledb` label.

```
recommend:
- match:
  - label: node-role.kubernetes.io/stabledb
```

- 5.6. Finally, provide the priority that is equal to or less than the currently applied profile. Profiles with lower priority values take precedence over ones with a higher value.

```
priority: 30
```

- 5.7. Verify that the file resembles the following CR:

```
apiVersion: tuned.openshift.io/v1
kind: Tuned
metadata:
  name: stabledb-tuning
  namespace: openshift-cluster-node-tuning-operator
spec:
  profile:
    - name: stabledb-tuning
      data: |
        [main]
        summary=Worker optimization disabling transparent_hugepages.
        [vm]
        transparent_hugepage=never

  recommend:
```

Chapter 3 | Cluster Partitioning

```
- match:  
  - label: node-role.kubernetes.io/stabledb  
    priority: 30  
    profile: stabledb-tuning
```

- 6. Apply the TuneD CR and verify that the profile is applied to the `worker01` node.

- 6.1. Apply the TuneD CR.

```
[student@workstation nodes-operators]$ oc apply -f stabledbCR.yml  
tuned.tuned.openshift.io/stabledb-tuning created
```

- 6.2. Verify the `stabled-tuning` profile for the `worker01` node.

```
[student@workstation nodes-operators]$ oc get profile \  
-n openshift-cluster-node-tuning-operator  
NAME      TUNED          APPLIED DEGRADED AGE  
master01  openshift-control-plane  True   False   47d  
master02  openshift-control-plane  True   False   47d  
master03  openshift-control-plane  True   False   47d  
worker01  stabledb-tuning        True   False   20h  
worker02  openshift-node         True   False   20h  
worker03  openshift-node         True   False   20h
```

- 7. From the second terminal window, re-create the debug pod to connect to the `worker01` node and verify the updated `transparent_hugepages` setting. This setting should reflect the changes from the new TuneD profile and no longer be enabled.

- 7.1. Open a new terminal window and create a debug pod to connect to the `worker01` node.

```
[student@workstation nodes-operators]$ oc debug node/worker01  
...output omitted...  
sh-4.4#
```

- 7.2. View the current setting for `transparent_hugepages` by displaying the value in the setting `enabled` file. The `[never]` value denotes that the setting is disabled.

```
sh-4.4# cat /sys/kernel/mm/transparent_hugepage/enabled  
always madvise [never]
```

- 7.3. Log out and remove the debug pod for the `worker01` node.

```
sh-4.4# exit  
...output omitted...
```

- 8. Change to the student HOME directory.

```
[student@workstation nodes-review]$ cd
```

Finish

On the `workstation` machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish nodes-operators
```

▶ Lab

Cluster Partitioning

Apply operating system settings to cluster nodes with the machine configuration operator, and with higher-level operators instead of the low-level machine configuration operator.

Outcomes

- Create machine configuration pools (MCPs) to apply custom configurations to some nodes in the cluster.
- Use the node tuning operator (NTO) to enable non-uniform memory access (NUMA) balancing.
- Use the machine configuration operator (MCO) to enable a real-time kernel.

Before You Begin

As the `student` user on the `workstation` machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start nodes-review
```

Instructions

Your company has a workload that requires many CPUs and much memory. Thus, the bandwidth between the CPUs and the memory is limiting the performance of the workload. In your cluster, the nodes with the `numa` role have NUMA hardware systems. These nodes perform best when the threads of their processes are accessing memory on the same NUMA node as where the threads are scheduled. Thus, your company requires you to enable the NUMA balancing in the kernel through the NTO for the nodes with that role. In this classroom environment, only one node has the `numa` role due to the classroom size. However, no nodes in the cluster have real NUMA hardware systems.

Your company also needs to deploy a workload for a 5G core user-plane function. This workload requires a real-time kernel to ensure low latency. Although using a `PerformanceProfile` custom resource (CR) is the recommended approach for configuring low-latency workloads which require real-time kernels, in this exercise you must use the MCO instead to practice by using its syntax. The NTO creates, among other resources, a similar machine configuration (MC) to the one that you create manually in the exercise. Although in production environments, this application type is deployed in multiple nodes and the kernel configuration must be applied to all of them, in this exercise you apply the configuration only to one node due to the classroom size.

Use the `admin` user with `redhatocp` as the password. For the resources that you must create, you can use the incomplete CR YAML files in the `~/D0380/labs/nodes-review` directory.

1. List the labels for the nodes in the cluster. Verify that the `worker01` node has the `numa` role. Create the `numa-bal` MCP to include in the pool the nodes with the `numa` role. This MCP must select MCs with the `numa-bal` and `worker` labels. You can find an incomplete example for the MCP CR in the `~/D0380/labs/nodes-review/numa-mcp.yaml` file.

Chapter 3 | Cluster Partitioning

2. Create the `kernel-numabal` TuneD profile CR to enable NUMA balancing. The TuneD profile CR applies to all the nodes in the `numa-bal` MCP, and includes the default `openshift-node` TuneD profile. You can find an incomplete example for the TuneD profile CR in the `~/D0380/labs/nodes-review/numa-tuned.yml` file. Verify that the NUMA balancing kernel parameter is enabled in the node. You can verify the NUMA balancing by running the `sysctl` command on the node and checking the `kernel.numa_balancing` parameter. Verify that the NUMA balancing feature is disabled on other nodes in the cluster.
3. Label the `worker02` node with the `kernel-rt` role. Create the `kernel-rt` MCP to include the nodes with the `kernel-rt` role in the pool. You can find an incomplete example for the MCP CR in the `~/D0380/labs/nodes-review/kernel-mcp.yml` file.
4. Use the `oc explain` command to find the MC parameter that changes the default kernel to a real-time kernel. Create the `99-kernel-realtime` MC to change the default kernel to a real-time kernel for the nodes in the `kernel-rt` pool. You can find an incomplete example for the MC CR in the `~/D0380/labs/nodes-review/kernel-mc.yml` file. Verify that the kernel for the `worker02` node is a real-time kernel. You can verify that the node uses a real-time kernel by using the `oc get nodes -o wide` command and verifying that the kernel version shows the `rt` string.

Evaluation

As the `student` user on the `workstation` machine, use the `lab` command to grade your work. Correct any reported failures and rerun the command until successful.

```
[student@workstation ~]$ lab grade nodes-review
```

Finish

As the `student` user on the `workstation` machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish nodes-review
```

► Solution

Cluster Partitioning

Apply operating system settings to cluster nodes with the machine configuration operator, and with higher-level operators instead of the low-level machine configuration operator.

Outcomes

- Create machine configuration pools (MCPs) to apply custom configurations to some nodes in the cluster.
- Use the node tuning operator (NTO) to enable non-uniform memory access (NUMA) balancing.
- Use the machine configuration operator (MCO) to enable a real-time kernel.

Before You Begin

As the `student` user on the `workstation` machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start nodes-review
```

Instructions

Your company has a workload that requires many CPUs and much memory. Thus, the bandwidth between the CPUs and the memory is limiting the performance of the workload. In your cluster, the nodes with the `numa` role have NUMA hardware systems. These nodes perform best when the threads of their processes are accessing memory on the same NUMA node as where the threads are scheduled. Thus, your company requires you to enable the NUMA balancing in the kernel through the NTO for the nodes with that role. In this classroom environment, only one node has the `numa` role due to the classroom size. However, no nodes in the cluster have real NUMA hardware systems.

Your company also needs to deploy a workload for a 5G core user-plane function. This workload requires a real-time kernel to ensure low latency. Although using a `PerformanceProfile` custom resource (CR) is the recommended approach for configuring low-latency workloads which require real-time kernels, in this exercise you must use the MCO instead to practice by using its syntax. The NTO creates, among other resources, a similar machine configuration (MC) to the one that you create manually in the exercise. Although in production environments, this application type is deployed in multiple nodes and the kernel configuration must be applied to all of them, in this exercise you apply the configuration only to one node due to the classroom size.

Use the `admin` user with `redhatocp` as the password. For the resources that you must create, you can use the incomplete CR YAML files in the `~/D0380/labs/nodes-review` directory.

1. List the labels for the nodes in the cluster. Verify that the `worker01` node has the `numa` role. Create the `numa-bal` MCP to include in the pool the nodes with the `numa` role. This MCP must select MCs with the `numa-bal` and `worker` labels. You can find an incomplete example for the MCP CR in the `~/D0380/labs/nodes-review/numa-mcp.yaml` file.
 - 1.1. Connect to the OpenShift cluster as the `admin` user with `redhatocp` as the password.

Chapter 3 | Cluster Partitioning

```
[student@workstation ~]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
Login successful.
...output omitted...
```

- 1.2. List the labels for the nodes. Verify that the **worker01** node has the **numa** role.

```
[student@workstation ~]$ oc get nodes
NAME      STATUS   ROLES          AGE     VERSION
master01  Ready    control-plane, master  49d    v1.25.7+eab9cc9
master02  Ready    control-plane, master  49d    v1.25.7+eab9cc9
master03  Ready    control-plane, master  49d    v1.25.7+eab9cc9
worker01 Ready    numa,worker    2d4h   v1.25.7+eab9cc9
worker02  Ready    worker           2d4h   v1.25.7+eab9cc9
worker03  Ready    worker           2d4h   v1.25.7+eab9cc9
```

- 1.3. Change to the `~/D0380/labs/nodes-review` directory.

```
[student@workstation ~]$ cd ~/D0380/labs/nodes-review
```

- 1.4. Create the MCP CR YAML file. You can find an incomplete example for the CR in the `~/D0380/labs/nodes-review/numa-mcp.yaml` file.

```
apiVersion: machineconfiguration.openshift.io/v1
kind: MachineConfigPool
metadata:
  name: numa-bal
spec:
  machineConfigSelector:
    matchExpressions:
      - key: machineconfiguration.openshift.io/role
        operator: In
        values: [worker, numa-bal]
  nodeSelector:
    matchLabels:
      node-role.kubernetes.io/numa: ""
```

- 1.5. Apply the configuration for the MCP CR.

```
[student@workstation nodes-review]$ oc create -f numa-mcp.yaml
machineconfigpool.machineconfiguration.openshift.io/numa-bal created
```

- 1.6. Verify that the **numa-bal** MCP is correctly created. Wait until the MCO updates the node and marks the **UPDATED** field as **True**. You might have to repeat this command many times.

```
[student@workstation nodes-review]$ oc get mcp
NAME          CONFIG              UPDATED   UPDATING   DEGRADED   MACHINECOUNT   ...
master        rendered-master-...  True      False      False      3           ...
numa-bal    rendered-numa-ba...  True     False      False      1           ...
worker        rendered-worker-...  True      False      False      2           ...
```

- 1.7. Verify that the worker01 node is assigned for the numa-bal MCP, and that the MCO applies the rendered MC to the node.

```
[student@workstation nodes-review]$ oc get events \
-n openshift-machine-config-operator \
--sort-by='{.lastTimestamp}' \
--field-selector involvedObject.name=numa-bal
... REASON          OBJECT          MESSAGE
... RenderedConfigGenerated machineconfigpool/numa-bal  rendered-numa-
bal-2f56...81d0 successfully generated ...
... SetDesiredConfig  machineconfigpool/numa-bal  Targeted node worker01
to MachineConfig rendered-numa-bal-2f56...81d0
... RenderedConfigGenerated machineconfigpool/numa-bal  rendered-numa-
bal-2f56...81d0 successfully generated ...
... SetDesiredConfig  machineconfigpool/numa-bal  Targeted node worker01
to MachineConfig rendered-numa-bal-2f56...81d0
```

2. Create the kernel-numabal Tuned profile CR to enable NUMA balancing. The Tuned profile CR applies to all the nodes in the numa-bal MCP, and includes the default openshift-node Tuned profile. You can find an incomplete example for the Tuned profile CR in the ~/D0380/labs/nodes-review/numa-tuned.yaml file. Verify that the NUMA balancing kernel parameter is enabled in the node. You can verify the NUMA balancing by running the sysctl command on the node and checking the kernel.numa_balancing parameter. Verify that the NUMA balancing feature is disabled on other nodes in the cluster.
- 2.1. Create the Tuned profile CR YAML file. You can find an incomplete example for the CR in the ~/D0380/labs/nodes-review/numa-tuned.yaml file.

```
apiVersion: tuned.openshift.io/v1
kind: Tuned
metadata:
  name: kernel-numabal
  namespace: openshift-cluster-node-tuning-operator
spec:
  profile:
    - data: |
        [main]
        summary=Custom NUMA balancing profile
        include=openshift-node
    ...output omitted...
  recommend:
    - machineConfigLabels:
        machineconfiguration.openshift.io/role: "numa-bal"
    priority: 30
  profile: kernel-numabal
```

Chapter 3 | Cluster Partitioning

2.2. Apply the configuration for the TuneD CR.

```
[student@workstation nodes-review]$ oc create -f numa-tuned.yml
tuned.tuned.openshift.io/kernel-numabal created
```

2.3. View the current TuneD profiles in use for each cluster node.

```
[student@workstation nodes-review]$ oc get profile \
-n openshift-cluster-node-tuning-operator
NAME      TUNED          APPLIED   DEGRADED AGE
master01  openshift-control-plane  True      False    49d
master02  openshift-control-plane  True      False    49d
master03  openshift-control-plane  True      False    49d
worker01 kernel-numabal        True      False    13m
worker02  openshift-node         True      False    2d4h
worker03  openshift-node         True      False    2d4h
```

2.4. Create a debug pod for the worker01 node.

```
[student@workstation nodes-review]$ oc debug node/worker01
...output omitted...
sh-4.4#
```

2.5. Run the sysctl command and verify that the kernel.numabalance parameter is enabled.

```
sh-4.4# sysctl -n kernel.numabalance
1
```

2.6. Remove the debug pod for the worker01 node.

```
sh-4.4# exit
...output omitted...
```

2.7. Create a debug pod for the worker02 node.

```
[student@workstation nodes-review]$ oc debug node/worker02
...output omitted...
sh-4.4#
```

2.8. Run the sysctl command and verify that the kernel.numabalance parameter is disabled.

```
sh-4.4# sysctl -n kernel.numabalance
0
```

2.9. Remove the debug pod for the worker02 node.

```
sh-4.4# exit
...output omitted...
```

Chapter 3 | Cluster Partitioning

3. Label the `worker02` node with the `kernel-rt` role. Create the `kernel-rt` MCP to include the nodes with the `kernel-rt` role in the pool. You can find an incomplete example for the MCP CR in the `~/D0380/labs/nodes-review/kernel-mcp.yml` file.

- 3.1. Add the `kernel-rt` role to the `worker02` node.

```
[student@workstation nodes-review]$ oc label node/worker02 \
node-role.kubernetes.io/kernel-rt=
```

- 3.2. List the labels for the nodes. Verify that the `worker02` node has the `kernel-rt` role.

```
[student@workstation ~]$ oc get nodes
NAME      STATUS   ROLES          AGE     VERSION
master01  Ready    control-plane,master  49d    v1.25.7+eab9cc9
master02  Ready    control-plane,master  49d    v1.25.7+eab9cc9
master03  Ready    control-plane,master  49d    v1.25.7+eab9cc9
worker01  Ready    numa,worker        2d4h   v1.25.7+eab9cc9
worker02 Ready    kernel-rt,worker  2d4h   v1.25.7+eab9cc9
worker03  Ready    worker          2d4h   v1.25.7+eab9cc9
```

- 3.3. Create the MCP CR YAML file. You can find an incomplete example for the CR in the `~/D0380/labs/nodes-review/kernel-mcp.yml` file.

```
apiVersion: machineconfiguration.openshift.io/v1
kind: MachineConfigPool
metadata:
  name: kernel-rt
spec:
  machineConfigSelector:
    matchExpressions:
      - key: machineconfiguration.openshift.io/role
        operator: In
        values: [worker,kernel-rt]
  nodeSelector:
    matchLabels:
      node-role.kubernetes.io/kernel-rt: ""
```

- 3.4. Apply the configuration for the MCP CR.

```
[student@workstation nodes-review]$ oc create -f kernel-mcp.yml
machineconfigpool.machineconfiguration.openshift.io/kernel-rt created
```

- 3.5. Verify that the `kernel-rt` MCP is correctly created. If the `UPDATING` field is marked as true for the `kernel-rt` MCP, then wait until the MCO finishes updating the node.

```
[student@workstation nodes-review]$ oc get mcp
NAME      CONFIG          UPDATED  UPDATING  DEGRADED  MACHINECOUNT ...
kernel-rt rendered-kernel-...  True    False     False     1         ...
master    rendered-master-... True    False     False     3         ...
numa-bal  rendered-numa-ba... True    False     False     1         ...
worker    rendered-worker-... True    False     False     1         ...
```

Chapter 3 | Cluster Partitioning

- 3.6. Verify that the `worker02` node is assigned for the `kernel-rt` MCP, and that the MCO applies the rendered MC to the node.

```
[student@workstation nodes-review]$ oc get events \
-n openshift-machine-config-operator \
--sort-by='{"lastTimestamp"}' \
--field-selector involvedObject.name=kernel-rt
... REASON OBJECT MESSAGE
...output omitted...
... RenderedConfigGenerated machineconfigpool/kernel-rt rendered-kernel-
rt-2f56...81d0 successfully generated ...
... SetDesiredConfig machineconfigpool/kernel-rt Targeted node worker02
to MachineConfig rendered-kernel-rt-2f56...81d0
```

- 4.** Use the `oc explain` command to find the MC parameter that changes the default kernel to a real-time kernel. Create the `99-kernel-realtime` MC to change the default kernel to a real-time kernel for the nodes in the `kernel-rt` pool. You can find an incomplete example for the MC CR in the `~/D0380/labs/nodes-review/kernel-mc.yml` file. Verify that the kernel for the `worker02` node is a real-time kernel. You can verify that the node uses a real-time kernel by using the `oc get nodes -o wide` command and verifying that the kernel version shows the `rt` string.

- 4.1. Find the MC parameter that changes the default kernel to a real-time kernel.

```
[student@workstation nodes-review]$ oc explain mc.spec
KIND:     MachineConfig
VERSION:  machineconfiguration.openshift.io/v1
...output omitted...
FIELDS:
...output omitted...

kernelType <string>
  Contains which kernel we want to be running like default (traditional),
  realtime

osImageURL <string>
  OSImageURL specifies the remote location that will be used to fetch the OS
```

- 4.2. Create the YAML file for the MC CR. You can find an incomplete example for the CR in the `~/D0380/labs/nodes-review/kernel-mc.yml` file.

```
apiVersion: machineconfiguration.openshift.io/v1
kind: MachineConfig
metadata:
  name: 99-kernel-realtime
  labels:
    machineconfiguration.openshift.io/role: "kernel-rt"
spec:
  kernelType: realtime
```

- 4.3. Apply the configuration for the MC CR.

Chapter 3 | Cluster Partitioning

```
[student@workstation nodes-review]$ oc create -f kernel-mc.yml
machineconfig.machineconfiguration.openshift.io/99-kernel-realtime created
```

- 4.4. Verify that the 99-kernel-realtime MC is correctly created.

```
[student@workstation nodes-review]$ oc get mc
NAME                GENERATEDBYCONTROLLER  IGNITIONVERSION  AGE
...output omitted...
01-worker-kubelet      52fe26136643a946ff...  3.2.0          49d
99-infra-chrony-conf-override   3.2.0          49d
99-kernel-realtime
99-master-chrony-conf-override   3.2.0          49d
99-master-generated-registries  52fe26136643a946ff...  3.2.0          49d
...output omitted...
```

- 4.5. Verify that the MCP starts updating the resources.

```
[student@workstation nodes-review]$ oc get mcp
NAME    CONFIG           UPDATED  UPDATING  DEGRADED  MACHINECOUNT ...
kernel-rt  rendered-kernel-rt...  False    True     False     1        ...
master   rendered-master-...  True     False     False     3        ...
numa-bal  rendered-numa-bal...  True     False     False     1        ...
worker   rendered-worker-...  True     False     False     1        ...
```

- 4.6. Verify that OpenShift applies the rendered MC to the worker02 node.

```
[student@workstation nodes-review]$ oc get events \
-n openshift-machine-config-operator \
--sort-by='lastTimestamp' \
--field-selector involvedObject.name=kernel-rt
... REASON           OBJECT           MESSAGE
...output omitted...
... RenderedConfigGenerated  machineconfigpool/kernel-rt  rendered-kernel-
rt-2f56...81d0 successfully generated ...
... SetDesiredConfig       machineconfigpool/kernel-rt  Targeted node
worker02 to MachineConfig rendered-kernel-rt-2f56...81d0
... RenderedConfigGenerated  machineconfigpool/kernel-rt  rendered-kernel-rt-
f6ab...234c successfully generated ...
... SetDesiredConfig       machineconfigpool/kernel-rt  Targeted node
worker02 to MachineConfig rendered-kernel-rt-f6ab...234c
```

- 4.7. Verify that the MCO updates the node. If the UPDATING field is marked as true for the kernel-rt MCP, then wait until the MCO finishes updating the node.

```
[student@workstation nodes-review]$ oc get mcp
NAME    CONFIG           UPDATED  UPDATING  DEGRADED  MACHINECOUNT ...
kernel-rt  rendered-kernel-rt...  True    False     False     1        ...
master   rendered-master-...  True     False     False     3        ...
numa-bal  rendered-numa-bal...  True     False     False     1        ...
worker   rendered-worker-...  True     False     False     1        ...
```

- 4.8. Verify that the `worker02` node is using a real-time kernel. You can do so by verifying that the kernel version contains the `rt` string.

```
[student@workstation nodes-review]$ oc get nodes -o wide
NAME      STATUS    ROLES          ...   KERNEL-VERSION
master01  Ready     control-plane,master ...  5.14.0-284.41.1.el9_2.x86_64
master02  Ready     control-plane,master ...  5.14.0-284.41.1.el9_2.x86_64
master03  Ready     control-plane,master ...  5.14.0-284.41.1.el9_2.x86_64
worker01  Ready     numa,worker       ...  5.14.0-284.41.1.el9_2.x86_64
worker02  Ready     kernel-rt,worker  ...  5.14.0-284.41.1.rt14.326.el9_2.x86_64
worker03  Ready     worker          ...  5.14.0-284.41.1.el9_2.x86_64
```

- 4.9. Change to the student HOME directory.

```
[student@workstation nodes-review]$ cd
```

Evaluation

As the `student` user on the `workstation` machine, use the `lab` command to grade your work. Correct any reported failures and rerun the command until successful.

```
[student@workstation ~]$ lab grade nodes-review
```

Finish

As the `student` user on the `workstation` machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish nodes-review
```

Summary

- The OpenShift cluster has control plane nodes and compute nodes.
- Control plane nodes manage the cluster and orchestrate the distribution of the workloads on the compute nodes.
- A node pool is a logical group of OpenShift compute nodes. Use node pools to organize the nodes on your cluster according to your cluster requirements.
- Use node labels to configure compute nodes to be a member of a specific node pool.
- The MCO manages the RHCOS operating system upgrades and configuration changes.
- Use the MCO to manage files in the /var and /etc directories, systemd services, SSH keys, and kernel arguments.
- The MCO defines two CRs: the MC CR, which declares instance customizations, and the MCP CR, which uses labels to match one or more MCs to one or more nodes.
- The MCO can manage two special CRs for modifying CRI-O container runtime settings and the Kubelet service.
- Use special-purpose operators, which are high level, to apply operating system settings to cluster nodes instead of applying them by using low-level operators.
- Use the NTO to tune the kernel for high-performance applications.

Chapter 4

Pod Scheduling

Goal

Configure workloads to run on a dedicated set of cluster nodes and prevent other workloads from using those cluster nodes.

Sections

- Pod Scheduling Concepts (and Matching Quiz)
- Pod Scheduling Scenarios (Quiz)
- Node Selectors and Taints (and Guided Exercise)
- High Availability with Affinity Rules and Pod Disruption Budgets (and Guided Exercise)

Lab

- Pod Scheduling

Pod Scheduling Concepts

Objectives

- Describe the main settings that affect pod placement on Kubernetes cluster nodes.

Pod Scheduling

Pod scheduling is the process of determining in which node to place a pod in the OpenShift cluster.

The OpenShift built-in scheduler identifies the most suitable node for the pods when you create them. The default scheduler meets the needs of most OpenShift users.

However, in some situations, OpenShift administrators might want more control over which node is used for pod placement. You can use the OpenShift advanced scheduling features to configure pods to run on a particular node or alongside a specific pod.

OpenShift includes the following advanced scheduling features:

- Scheduler profiles, to control how OpenShift schedules pods on nodes.
- Pod affinity rules, to keep sets of pods close to each other, on the same nodes. For example, you can run a REST service and its database on the same node to minimize network latency.
- Pod anti-affinity rules, to keep sets of pods far away from each other, on different nodes. For example, you can run replica pods of the same deployment in different nodes, so that if a node fails, then you do not lose all the pods for the workload.
- Node affinity, to keep sets of pods running on the same group of nodes. For example, you can configure a pod to run on nodes with a specific CPU.
- Node selectors, to schedule pods to a specific set of nodes. For example, you can schedule a pod to a node that provides special hardware that the pod needs.
- Taints and tolerations, to avoid scheduling pods to a specific set of nodes. For example, you can block a pod to run on a node that is reserved for OpenShift cluster services or control plane services.

OpenShift also includes the pod disruption budget resource, which controls how many instances can be down at the same time during voluntary disruptions, such as when scaling down, updating applications, or draining a node for maintenance.

Chapter 4 | Pod Scheduling

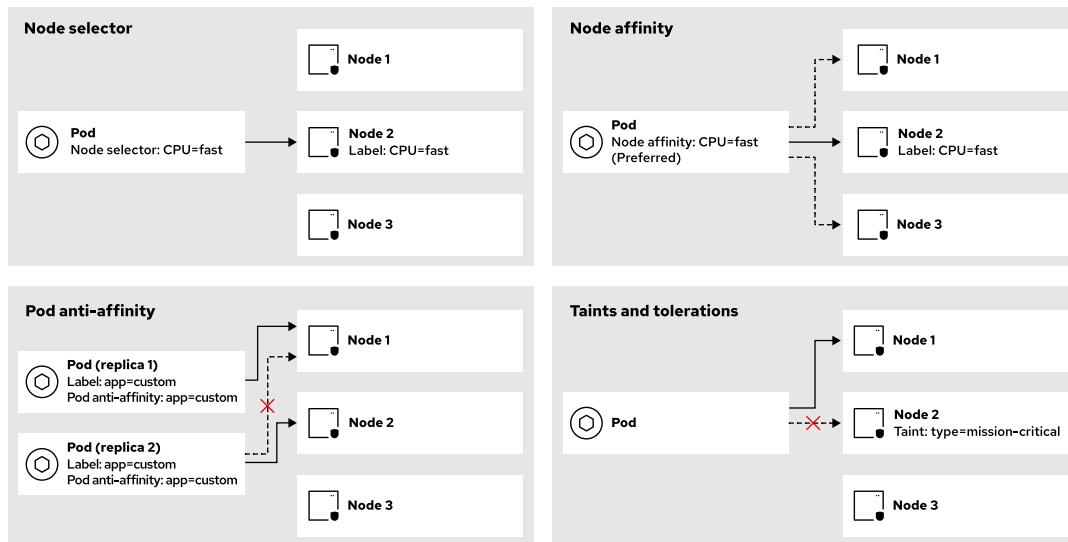


Figure 4.1: OpenShift advanced scheduling features

The previous diagram shows some of the advanced scheduling features. For example, the pod with a node selector for the `CPU=fast` label can be placed only on the second node, which contains the `CPU=fast` label.

For a pod that has a preferred node affinity rule for the `CPU=fast` label, OpenShift first tries to schedule it on the node with the `CPU=fast` label. If OpenShift cannot place the pod on that node, then OpenShift schedules the pod on another node.

If you create a deployment with the `app=custom` label and a pod anti-affinity rule for the same label, then OpenShift places the replicas for the deployment on different nodes.

If you taint a node with the `type=mission-critical` label, then OpenShift can schedule a pod on that node only if the pod has the correct toleration that matches the taint.

It is common in production environments to use more than one advanced scheduling feature at a time. For example, your cluster includes a workload that has a node selector to get GPU-enabled nodes, and pod anti-affinity rules to avoid downtime in the event of missing a node. The same workload also has a pod disruption budget to specify accepted degradation during a cluster upgrade, so the workload could run slower but still be available to its users.

Default Pod Scheduler

The default OpenShift pod scheduler determines the placement of pods onto nodes within the cluster. The scheduler reads data from the pod and identifies a suitable node based on configured profiles. After identifying the most suitable node, the scheduler creates a binding that associates the pod with a specific node, without modifying the pod.

The OpenShift default pod scheduler determines the placement of a pod onto a particular node according to the following procedure:

1. **Node filtering:** The OpenShift scheduler filters the nodes based on the configured constraints or requirements by means of functions called *predicates*. These predicates include available CPU capacity on the node to satisfy a pod's CPU resource request, free ports, or volume availability, among others.
2. **Prioritize the filtered list of nodes:** In this step, the scheduler assesses each node by using a set of priority or scoring functions, and assigns each node a score from 0 to 10. A score

Chapter 4 | Pod Scheduling

of 0 indicates a poor fit, and a score of 10 indicates an excellent fit for hosting the pod. Additionally, OpenShift administrators can assign a numeric weight to each scoring function in the scheduler's configuration. With this weight attribute, administrators can prioritize certain scoring functions.

3. Select the best node: OpenShift sorts the nodes based on their scores and selects the node with the highest score. If multiple nodes receive the same score, then OpenShift selects one node randomly.

If a pod does not specify its resource requests, then the scheduler could place it on a node that is already full, which could lead to poor performance or even killing the pod if the node is out of memory.

Scheduler Profiles

The OpenShift scheduler profile controls how OpenShift schedules pods onto nodes. You can set the scheduler profile by using the `oc edit scheduler cluster` command and modifying the `spec.profile` parameter.

The following scheduler profiles are available:

LowNodeUtilization

By using this profile, OpenShift distributes pods evenly across nodes to ensure a low resource usage per node. This profile provides the default scheduler behavior.

HighNodeUtilization

With this profile, OpenShift places as many pods as possible onto as few nodes as possible. This profile minimizes the node count but increases the resource usage per node.

NoScoring

This profile aims for quick scheduling cycles by disabling all the score plug-ins. Thus, by using this profile you sacrifice better scheduling decisions for faster ones.

Advanced Pod Scheduling

In some situations, you might want more control over which node is used for pod placement. You can use the following OpenShift advanced scheduling features to configure pods to run on a particular node or alongside a specific pod.

Node Selectors

Because a node selector specifies target node labels on your pod, OpenShift schedules the pod only on nodes that match the target labels. Thus, with a node selector, you must define the target labels in the pod specification file and label the nodes with those labels.

You can use node selectors to place pods on specific nodes. You can also define cluster-wide node selectors to place pods on specific nodes anywhere in the cluster, and define project node selectors to place pods in a project on specific nodes.

For example, you can label nodes in your cluster by using their geographical location. Then, application developers can use node selectors to deploy pods on nodes in a geographical region for availability and latency.

Affinity and Anti-affinity Rules

Affinity

Affinity refers to a pod property that controls the pod preference for a node.

Anti-affinity

Anti-affinity is a pod property that restricts the placement of the pod on a node.

You can use affinity and anti-affinity rules to control the following scenarios:

- The node to place a pod in, which is called *node affinity*.
- The node to place a pod in, relative to other pods, which is called *pod affinity*.

Node affinity

With node affinity, you can specify in a pod the affinity in relation to a group of nodes.

For example, you can configure a pod to run on nodes with a specific CPU or in a specific availability zone.

Node affinity, which is conceptually similar to node selectors, provides a more flexible way to specify constraints on node selection. Whereas node selectors specify the target nodes by using a set of label key-value pairs, node affinity enables a conditional approach with logical operators in the matching process.

With node affinity, you can define required and preferred rules. Because required rules work similarly to node selectors, the OpenShift scheduler schedules the pod only if the conditions are met. By using preferred rules, the OpenShift scheduler tries to schedule the pod in a node that meets the rules. If the scheduler does not find a node, then the scheduler schedules the pod on a node that does not meet the rules.

Pod affinity and anti-affinity

With pod affinity and pod anti-affinity, you can control the nodes for your pod based on the labels of other pods.

By using pod affinity rules, you can locate a pod on the same group of nodes as other pods based on its label selector. As an example, you can configure the scheduler to place pods of two services onto the same group of nodes, because the services communicate with each other regularly.

By using pod anti-affinity rules, you can prevent the scheduler from locating a pod on the same group of nodes as other pods based on its label selector. For example, you can configure the scheduler to prevent a pod of a particular service from being on the same nodes as pods of another service, because the two services interfere and reduce the performance of the services. Another typical scenario is to use node topology labels, such as availability zones in a cloud provider or racks in a data center, and to use anti-affinity rules to avoid putting pods of the same workload on the same failure domain.

Pod affinity, pod anti-affinity rules, and node affinity rules each have two types: required and preferred.

Required rules must be met before the scheduler places a pod on a node. If you use required rules and those rules are not met, then the scheduler cannot find an appropriate node to place the pod.

Preferred rules specify conditions that the scheduler tries to enforce, but the scheduler still schedules the pod if it cannot find a matching node.

You can configure pod and node affinity and anti-affinity rules in the pod specification YAML files by setting the `spec.config.affinity` parameter. For that parameter, you can set the `nodeAffinity`, `podAffinity`, and `podAntiAffinity` properties, and specify a required rule, or a preferred rule, or both. If you specify both required and preferred rules, then the node must first meet the required rules and after the scheduler attempts to meet the preferred rules.

Taints and Tolerations

You can configure taints on nodes so they schedule a pod only if the pod has a matching toleration. The node with a taint repels all the pods except those pods that have a toleration for that taint. For example, OpenShift automatically taints the control plane nodes, so that the pods that manage the control plane can be scheduled on them. However, the control plane nodes reject any other data plane pods that users deploy.

You can apply taints to a node in the node specification YAML file by setting the `spec.taints` parameter, and apply tolerations to a pod in the pod specification YAML file by setting the `spec.tolerations` parameter.

Taints and tolerations consist of a key, value, and effect.

- The key is any string, up to 253 characters.
- The value is any string, up to 63 characters.
- The effect is one of the following options:

NoSchedule

If the pod does not match the taint, then the node prevents scheduling it. Existing pods on the node continue running on the node.

PreferNoSchedule

If the pod does not match the taint, then the scheduler tries to avoid scheduling it in the node. Existing pods on the node continue running on the node.

NoExecute

If the pod does not match the taint, then the node prevents scheduling it. The scheduler removes any existing pods on the node that do not have a matching toleration.

Tolerations also include the `operator` parameter. The following values for the `operator` parameter are possible:

Equal

The toleration matches the taint if the key, value, and effect parameters match. This behavior is the default.

Exists

The toleration matches the taint if the key and effect parameters match. You must leave a blank value parameter.

Pod Disruption Budget

A *pod disruption budget* is a policy resource to control the disruption of pods during voluntary disruptions, such as scaling down, updating applications, or draining a node for maintenance. Pod disruption budgets do not apply on node failures.

By setting up pod disruption budgets, you can manage the availability and stability of your applications during disruptions, to reduce the risk of service degradation or downtime when maintaining or updating your cluster.

The configuration for a pod disruption budget consists of the following parts:

- A label selector to target a specific set of pods. By using the label selector, you define the pods that the pod disruption budget protects. Only pods that match the label selector are subject to the budget constraints. The pod selector in a pod disruption budget typically selects only pods from the same deployment.

- The availability level, which specifies the minimum number of pods that must be available simultaneously. You can define the following availability levels:

minAvailable

Defines the minimum number of pods that must always be available, even during a disruption. A pod is available when it has the `Ready` condition with the `True` value.

maxUnavailable

Defines the maximum number of pods that can be unavailable during a disruption.



References

For more information about pod scheduling on OpenShift, refer to the *Controlling Pod Placement onto Nodes (Scheduling)* section in the Red Hat OpenShift Container Platform 4.14 Nodes documentation at

https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/nodes/index#controlling-pod-placement-onto-nodes-scheduling

For more information about pod disruption budgets, refer to the *Understanding How to Use Pod Disruption Budgets to Specify the Number of Pods That Must Be Up* section in the Red Hat OpenShift Container Platform 4.14 Nodes documentation at

https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/nodes/index#nodes-pods-pod-disruption-about_nodes-pods-configuring

► Quiz

Pod Scheduling Concepts

Match the following items to their counterparts in the table.

Node affinity rules (preferred)

Node selectors

Pod affinity rules

Pod anti-affinity rules

Pod disruption budget

Taints and tolerations

Scenario	Advanced scheduling feature
You use labels for availability zones in a cloud provider, so you use this advanced scheduling feature in your application to avoid bandwidth fees between availability zones. Avoiding bandwidth fees between availability zones is a requirement.	
One of your cluster nodes must undergo maintenance activities, so you use this advanced scheduling feature to ensure that existing pods on the node are evicted before maintenance activities begin.	
You use labels for availability zones in a cloud provider, so you use this advanced scheduling feature in your application to avoid bandwidth fees between availability zones. Avoiding bandwidth fees between availability zones is preferred, but not required.	
You need to run two services that must communicate with each other with low latency, so you use this advanced scheduling feature to locate the pods from both services on the same group of nodes without constraining the number of nodes.	
You use labels for availability zones in a cloud provider, so you use this advanced scheduling feature to avoid putting pods of the same workload on the same failure domain.	
One of your applications needs at least two pods to be running on nodes, even in the event of a disruption, so you use this advanced scheduling feature to ensure this requirement.	

► Solution

Pod Scheduling Concepts

Match the following items to their counterparts in the table.

Scenario	Advanced scheduling feature
You use labels for availability zones in a cloud provider, so you use this advanced scheduling feature in your application to avoid bandwidth fees between availability zones. Avoiding bandwidth fees between availability zones is a requirement.	Node selectors
One of your cluster nodes must undergo maintenance activities, so you use this advanced scheduling feature to ensure that existing pods on the node are evicted before maintenance activities begin.	Taints and tolerations
You use labels for availability zones in a cloud provider, so you use this advanced scheduling feature in your application to avoid bandwidth fees between availability zones. Avoiding bandwidth fees between availability zones is preferred, but not required.	Node affinity rules (preferred)
You need to run two services that must communicate with each other with low latency, so you use this advanced scheduling feature to locate the pods from both services on the same group of nodes without constraining the number of nodes.	Pod affinity rules
You use labels for availability zones in a cloud provider, so you use this advanced scheduling feature to avoid putting pods of the same workload on the same failure domain.	Pod anti-affinity rules
One of your applications needs at least two pods to be running on nodes, even in the event of a disruption, so you use this advanced scheduling feature to ensure this requirement.	Pod disruption budget

► Quiz

Pod Scheduling Scenarios

Choose the correct answers to the following questions:

- 1. **One of your applications needs at least two pods to be running on nodes, even in the event of a disruption. Which option would ensure that at least two pods are always running even in the event of a disruption?**
- Set the `minAvailable` parameter to 100%, to ensure that the application is always available.
 - Set the `maxUnavailable` parameter to 0, to ensure that the application is always available.
 - Set the `minAvailable` parameter to 2, to ensure that at least two pods are always running.
 - Set the `maxUnavailable` parameter to 2, to ensure that at least two pods are always running.
- 2. **Your company needs you to ensure that a workload runs on a group of GPU-enabled nodes. The workload must avoid downtime in the event of missing a node, and must also accept degradation during a cluster update. Which set of advanced scheduling features would ensure these requirements?**
- Taints and tolerations, node affinity rules, and pod affinity rules
 - Node selectors, taints and tolerations, and node affinity rules
 - Node selectors, a pod disruption budget, and pod anti-affinity rules
 - Node selectors, a pod disruption budget, and pod affinity rules
- 3. **You have a multi-tier application that consists of web, application, and database components. You must distribute the web and application components across multiple nodes for high availability. You must place the database component on specific nodes with high I/O capacity. Which set of advanced scheduling features would ensure these requirements?**
- Pod anti-affinity rules for the web and application pods, and node selectors for the database pods.
 - Pod affinity rules for the web and application pods, and pod anti-affinity rules for the database pods.
 - Node affinity rules for the web pods, pod affinity rules for the application pods, and node selectors for the database pods.
 - Node selectors for the web pods, node affinity rules for the application pods, and pod affinity rules for the database pods.

- 4. You have GPU-enabled nodes in your cluster. Your company needs you to optimize scheduling the training workloads of a machine learning model, which requires GPU-enabled nodes. Moreover, as these workloads require resource-intensive calculations, you must distribute the pods across multiple nodes. Which set of advanced scheduling features would ensure these requirements?
- a. Node selectors, and taints and tolerations
 - b. Node selectors and pod affinity rules
 - c. Node affinity rules, and taints and tolerations
 - d. Node affinity rules and pod anti-affinity rules

► Solution

Pod Scheduling Scenarios

Choose the correct answers to the following questions:

- 1. **One of your applications needs at least two pods to be running on nodes, even in the event of a disruption. Which option would ensure that at least two pods are always running even in the event of a disruption?**
- a. Set the `minAvailable` parameter to 100%, to ensure that the application is always available.
 - b. Set the `maxUnavailable` parameter to 0, to ensure that the application is always available.
 - c. Set the `minAvailable` parameter to 2, to ensure that at least two pods are always running.
 - d. Set the `maxUnavailable` parameter to 2, to ensure that at least two pods are always running.
- 2. **Your company needs you to ensure that a workload runs on a group of GPU-enabled nodes. The workload must avoid downtime in the event of missing a node, and must also accept degradation during a cluster update. Which set of advanced scheduling features would ensure these requirements?**
- a. Taints and tolerations, node affinity rules, and pod affinity rules
 - b. Node selectors, taints and tolerations, and node affinity rules
 - c. Node selectors, a pod disruption budget, and pod anti-affinity rules
 - d. Node selectors, a pod disruption budget, and pod affinity rules
- 3. **You have a multi-tier application that consists of web, application, and database components. You must distribute the web and application components across multiple nodes for high availability. You must place the database component on specific nodes with high I/O capacity. Which set of advanced scheduling features would ensure these requirements?**
- a. Pod anti-affinity rules for the web and application pods, and node selectors for the database pods.
 - b. Pod affinity rules for the web and application pods, and pod anti-affinity rules for the database pods.
 - c. Node affinity rules for the web pods, pod affinity rules for the application pods, and node selectors for the database pods.
 - d. Node selectors for the web pods, node affinity rules for the application pods, and pod affinity rules for the database pods.

- 4. You have GPU-enabled nodes in your cluster. Your company needs you to optimize scheduling the training workloads of a machine learning model, which requires GPU-enabled nodes. Moreover, as these workloads require resource-intensive calculations, you must distribute the pods across multiple nodes. Which set of advanced scheduling features would ensure these requirements?
- a. Node selectors, and taints and tolerations
 - b. Node selectors and pod affinity rules
 - c. Node affinity rules, and taints and tolerations
 - d. Node affinity rules and pod anti-affinity rules

Node Selectors and Taints

Objectives

- Configure a workload to run on dedicated nodes, and prevent other workloads from running on those nodes.

Ensure and Prevent Workloads on Nodes

OpenShift advanced scheduling features include *node selectors*, which ensure that a pod is placed in certain nodes, and *taints and tolerations*, which prevent certain nodes from running some workloads.

Use node selectors to schedule pods to a dedicated set of nodes. For example, schedule a pod to a node that provides special hardware that the pod needs.

Use taints and tolerations to avoid scheduling pods to a dedicated set of nodes. For example, block a pod to run on a node that is reserved for OpenShift cluster services or for control plane services.

Node Selectors

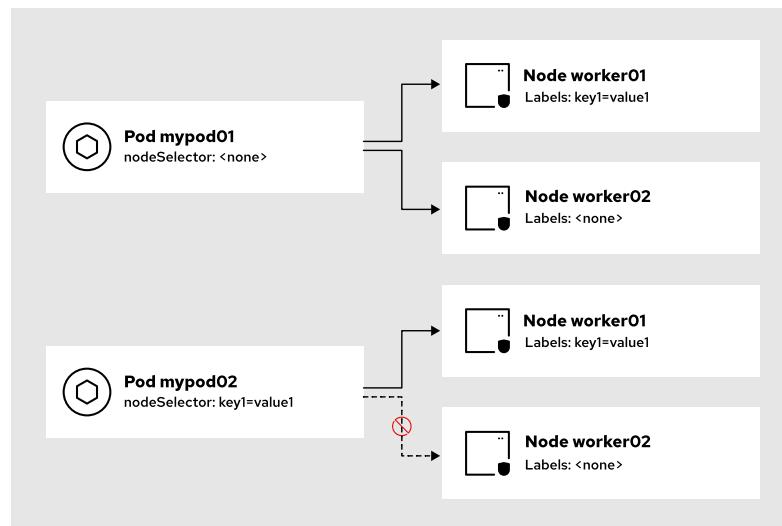
By using node selectors, you can specify target node labels on your pod, so the OpenShift scheduler tries to place the pod only on those nodes that match the target labels. If the OpenShift scheduler does not find a node that matches the target labels, then your pod will be unavailable.

You can use pod node selectors to place pods on dedicated nodes. You define pod node selectors as attributes on the pod template.

Use project node selectors to place pods in a project on dedicated nodes. You define project node selectors as annotations in the project CR.

Use cluster-wide node selectors to place pods on dedicated nodes anywhere in the cluster. You define cluster-wide node selectors as attributes on the OpenShift scheduler CR.

Specific syntax details for node selectors are explained later in this section.

**Figure 4.2: Node selectors**

The previous diagram shows how pod node selectors work. In the figure, the `worker01` node has the `key1=value1` label, but the `worker02` node does not have that label. For the first pod, OpenShift can schedule it in either node, because it does not have a node selector.

The second pod has a node selector that is defined as an attribute in the pod template. Then, OpenShift can schedule the pod only in nodes with a label that matches the node selector. Thus, the OpenShift scheduler can place the pod only in the `worker01` node.

Node Labels

To ensure that the OpenShift scheduler finds a node that matches the target labels, first label your nodes. You can add labels directly to a node resource or indirectly by using a machine set resource.

If your cluster is installed in a way that supports machine set resources, then you must use machine sets to set labels on nodes. In that case, do not label nodes directly. If your cluster does not support machine set resources, because for example it is installed by using the user-provisioned infrastructure method, then you must label nodes directly.



Note

If your cluster supports machine set resources, then OpenShift can delete and re-create the nodes at any time. Thus, if you directly label the node, then the node loses the label outside the machine set when OpenShift re-creates it.

You can add one or more labels to a node by using the following command:

```
[user@host ~]$ oc label node node1 key1=value1 ... keyN=valueN
node/node1 labeled
```

You can remove labels from a node by using the following command:

```
[user@host ~]$ oc label node node1 key1- ... keyN-
node/node1 unlabeled
```

Pod Node Selectors

Use node selectors on a pod to specify target node labels where the OpenShift scheduler tries to place the pod.

You can add one or more node selectors to a pod by using the following syntax:

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
  ...output omitted...
spec:
  nodeSelector:
    key1: value1
    key2: value2
  ...output omitted...
```

In the previous example, the OpenShift scheduler places the pod only in nodes with both the key1=value1 and key2=value2 labels.

You can read the node selectors for a pod by reviewing the `Node-Selectors` parameter as follows:

```
[user@host ~]$ oc describe pod myapp-95664666d-5l22c
...output omitted...
Node-Selectors:      key1:value1
                    key2:value2
...output omitted...
```

Cluster-wide Node Selectors

By default, the OpenShift scheduler tries to place the pods onto any available nodes in your cluster. However, you can modify the OpenShift scheduler default behavior so it includes default node selectors to any new pods that are created in the cluster. By using cluster-wide node selectors, you specify the default node selectors that OpenShift adds to the pods. You require administrative privileges to specify cluster-wide node selectors.

To define cluster-wide node selectors, you can edit the OpenShift scheduler CR as follows:

```
[user@host ~]$ oc edit scheduler cluster
```

Then specify the `defaultNodeSelector` parameter.

```
apiVersion: config.openshift.io/v1
kind: Scheduler
metadata:
  name: cluster
  ...output omitted...
spec:
  defaultNodeSelector: key1=value1,key2=value2
  ...output omitted...
```

Chapter 4 | Pod Scheduling

After you edit the OpenShift scheduler CR, wait for the pods in the `openshift-kube-apiserver` project to redeploy.

In the previous example, any new pods that you create in the cluster have the `key1=value1` and `key2=value2` node selectors, by default.

Project Node Selectors

You can specify node selectors for a project, so OpenShift includes those node selectors to any new pods that are created in the project. You require administrative privileges to specify project node selectors.

To define project node selectors, you can create a YAML file for an OpenShift project CR, and define the node selectors in the `metadata.annotations` section, as follows:

```
apiVersion: v1
kind: Namespace
metadata:
  name: myproject
  annotations:
    openshift.io/node-selector: key3=value3, key4=value4
...output omitted...
```

In the previous example, any new pods that you create in the project have by default the `key3=value3` and `key4=value4` node selectors.

Levels for Node Selectors

You can define node selectors at the following levels: pod, project, and cluster-wide.

Project node selectors take precedence over cluster-wide node selectors. Thus, if you define both project and cluster-wide node selectors, then OpenShift applies only the project node selectors to the new pods but not the cluster-wide node selectors.

Pod node selectors are additive to cluster-wide and project node selectors. Thus, if you define both pod and project node selectors, then OpenShift creates the pod with both the pod node selectors and the project node selectors. If you define both pod and cluster-wide node selectors, then OpenShift creates the pod with both the pod node selectors and the cluster-wide node selectors.



Important

Creating a pod node selector with the same key string but with a different value string from a project node selector or from a cluster-wide node selector causes a conflict, and OpenShift fails to create your pod.

Taints and Tolerations

You can use taints on nodes to prevent certain nodes from running some workloads. The OpenShift scheduler schedules pods in those nodes only if the pods have a matching toleration.

Chapter 4 | Pod Scheduling

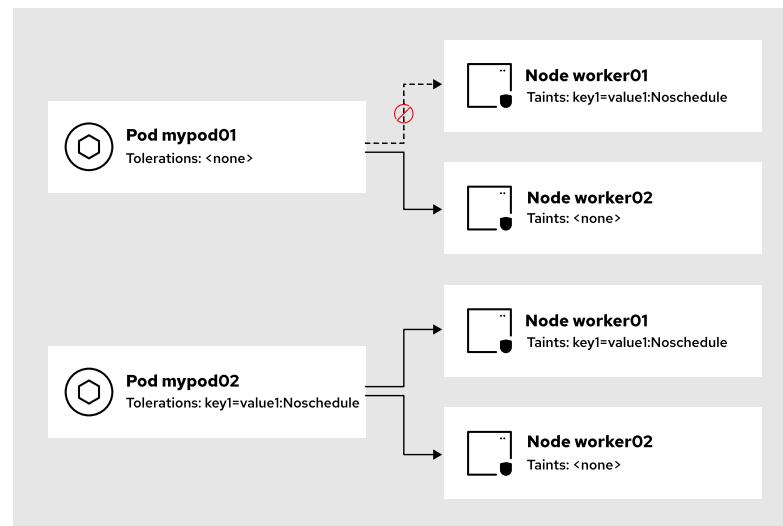


Figure 4.3: Node taints and pod tolerations

In the previous figure, the `worker01` node has a taint, but the `worker02` node does not have a taint. The OpenShift scheduler cannot place the first pod in the `worker01` node; the node rejects the pod because the toleration does not match the taint. Thus, OpenShift can schedule the pod only in the `worker02` node.

OpenShift can schedule the second pod in both nodes, because the toleration matches the taint.

Node Taints

Taints consist of a key, a value, and an effect.

- The key is any string, up to 253 characters.
- The value is any string, up to 63 characters.
- The effect is one of the following options:

NoSchedule

If the pod does not match the taint, then the node prevents scheduling it. Existing pods on the node continue running on the node.

PreferNoSchedule

If the pod does not match the taint, then the scheduler tries to avoid scheduling it in the node. Existing pods on the node continue running on the node.

NoExecute

If the pod does not match the taint, then the node prevents scheduling it. The scheduler removes any existing pods on the node without a matching toleration.

You can apply a taint to a node in the node specification YAML file by setting the `spec.taints` parameter or by using the following command:

```
[user@host ~]$ oc adm taint nodes node1 key1=value1:NoSchedule
node/node1 tainted
```

After you apply a taint to a node, you can verify the taint by using the following command:

Chapter 4 | Pod Scheduling

```
[user@host ~]$ oc describe nodes node1
Name:           worker01
Roles:          worker
...output omitted...
Taints:         key1=value1:NoSchedule
...output omitted...
```

To remove a taint from a node, use the following command:

```
[user@host ~]$ oc adm taint nodes node01 key1=value1:NoSchedule-
node/node01 untainted
```

Pod Tolerations

Tolerations consist of a key, a value, and an effect, similar to the node taints. Tolerations also include the `operator` parameter. The following values for the `operator` parameter are possible:

Equal

The toleration matches the taint if the key, the value, and the effect parameters match. This behavior is the default.

Exists

The toleration matches the taint if the key and the effect parameters match. You must leave the value parameter blank.

When you define a toleration with the `NoExecute` effect, you can also define the `tolerationSeconds` parameter. The `tolerationSeconds` parameter defines the time that a pod stays within a node before the node evicts it.

You can apply a toleration to a pod in the pod specification YAML file by setting the `spec.tolerations` parameter as follows:

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
  ...output omitted...
spec:
  tolerations:
  - key: "key1"
    operator: "Equal"
    value: "value1"
    effect: "NoExecute"
    tolerationSeconds: 3600
  ...output omitted...
```

In the previous example, OpenShift can schedule the pod in nodes with the `key1=value1:NoExecute` taint. If OpenShift schedules the pod in a node with the taint, then the pod runs in the node for 3600 seconds before the node evicts it and OpenShift reschedules the pod to another node.

Pod Scheduling and Node Conditions

By default, OpenShift taints and evicts nodes by condition. For example, OpenShift automatically taints nodes under some conditions, such as memory or disk pressure.

The following taints are built into OpenShift:

`node.kubernetes.io/not-ready`

The node is not ready. This taint corresponds to the `Ready=False` node condition.

`node.kubernetes.io/unreachable`

The node is unreachable from the node controller. This taint corresponds to the `Ready=Unknown` node condition.

`node.kubernetes.io/memory-pressure`

The node has memory pressure issues. This taint corresponds to the `MemoryPressure=True` node condition.

`node.kubernetes.io/disk-pressure`

The node has disk pressure issues. This taint corresponds to the `DiskPressure=True` node condition.

`node.kubernetes.io/network-unavailable`

The node network is unavailable.

`node.kubernetes.io/unschedulable`

The node is unschedulable.

`node.cloudprovider.kubernetes.io/uninitialized`

This taint sets the node as unusable when you start the node controller in an external cloud provider. After the cloud controller manager initializes the node, the kubelet removes the taint.

`node.kubernetes.io/pid-pressure`

The node has process identifier (PID) pressure. This taint corresponds to the `PIDPressure=True` node condition.

If a node reports one of the conditions from the previous table, then OpenShift applies the taint to the node until the condition clears.

For example, you can verify how OpenShift automatically applies the `node.kubernetes.io/unschedulable` taint to an unschedulable pod by following the next example. First, verify the taints for a node by using the following command:

```
[user@host ~]$ oc describe nodes node1 | grep Taints
Taints:      <none>
```

Then, mark the node as unschedulable.

```
[user@host ~]$ oc adm cordon node1
node/node1 cordoned
```

Finally, verify that OpenShift automatically applies the `node.kubernetes.io/unschedulable` taint with the `NoSchedule` effect to the node.

```
[user@host ~]$ oc describe nodes node1 | grep Taints
Taints:      node.kubernetes.io/unschedulable:NoSchedule
```

Chapter 4 | Pod Scheduling

OpenShift applies the taints to the nodes from the previous table with the `NoSchedule` effect, except for the `node.kubernetes.io/not-ready` and `node.kubernetes.io/unreachable` taints where OpenShift uses the `NoExecute` effect.

For the taints with the `NoSchedule` effect, the OpenShift scheduler schedules the pod to the node if the pod has a matching toleration, or if the node returns to a normal state, but running pods remain unaltered.

For the taints with the `NoExecute` effect, the OpenShift scheduler schedules the pod to the node if the pod has a matching toleration, or if the node returns to a normal state, and starts evicting and rescheduling running pods to different nodes. For pods that do not tolerate the taint and that are running on the node, OpenShift evicts them immediately. For pods that match the toleration, OpenShift evicts them only if the `tolerationSeconds` parameter is defined on the pod. The `tolerationSeconds` parameter defines the time in seconds that a pod remains in a node until the node evicts it.

The `tolerationSeconds` parameter is useful in certain scenarios, such as when an application must create local files when running for the first time. In this case, waiting for a few seconds for the pod to recover might be preferable to directly re-creating the pod on another node. If you define the `tolerationSeconds` parameter for an application, then the node waits the specified time before evicting the pods.

Project Tolerations

You can specify tolerations for a project, so OpenShift includes those tolerations to any new pods that are created in the project. You require administrative privileges to specify project tolerations. If you define project and pod tolerations, then OpenShift creates the pod with both of these tolerations.

To define project tolerations, you can create a YAML file for an OpenShift project CR and define the toleration in the `metadata.annotations` section as follows:

```
kind: Project
apiVersion: project.openshift.io/v1
metadata:
  name: myproject
  annotations:
    scheduler.alpha.kubernetes.io/defaultTolerations: >-
      [{"operator": "Equal", "effect": "NoSchedule", "key": "key1", "value": "value1"}]
```

In the previous example, any new pods that you create in the project have by default the `key1=value1:NoSchedule` toleration.

Tolerating All Taints

You can configure a pod to tolerate all the node taints by using the operator: "Exists" toleration with no key or value parameters.

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
  ...output omitted...
spec:
  tolerations:
    - operator: "Exists"
```



Note

OpenShift does not remove the pods with the `operator="Exists"` toleration from a node that has taints.



References

For more information about node selectors, refer to the *Placing Pods on Specific Nodes Using Node Selectors* section in the Red Hat OpenShift Container Platform 4.14 Nodes documentation at
https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/nodes/index#nodes-scheduler-node-selectors

For more information about taints and tolerations, refer to the *Controlling Pod Placement Using Node Taints* section in the Red Hat OpenShift Container Platform 4.14 Nodes documentation at
https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/nodes/index#nodes-scheduler-taints-tolerations

► Guided Exercise

Node Selectors and Taints

Configure a workload to run on dedicated nodes, and prevent other workloads from running on those nodes.

Outcomes

- Create deployments with node selectors to specify the node pool for a workload.
- Use taints and tolerations to prevent a node from running workloads.

Before You Begin

As the student user on the workstation machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start scheduling-selector
```

Instructions

Your company has a cluster with three compute nodes in three different racks, as in the following table:

Node name	Labels
worker01	<ul style="list-style-type: none">• rack=1• cpu=standard
worker02	<ul style="list-style-type: none">• rack=2• cpu=standard
worker03	<ul style="list-style-type: none">• rack=3• cpu=fast

The nodes include a label with the name of the rack that they are part of. The `worker01` and `worker02` nodes have a standard CPU and include the `cpu=standard` label. The `worker03` node has a fast CPU and includes the `cpu=fast` label.

This guided exercise includes the following scenarios:

In the first scenario, you create a deployment with two replicas to verify the default OpenShift scheduler behavior, which distributes the deployment pods between the three available compute nodes. Then, your company requires you to create another deployment, with two replicas, which needs a fast CPU. Thus, you must use a node selector to ensure that the OpenShift scheduler places the deployment in the fast CPU compute node. Finally, you drain the `worker03` node from the cluster to simulate a failure in the rack, and verify the status for the deployments.

In the second scenario, you create a project node selector for the standard CPU nodes. Then, you create a deployment with two replicas and verify that OpenShift applies the project node selector to the deployment.

Chapter 4 | Pod Scheduling

In this scenario, you also create another deployment with two replicas and with a node selector for the compute node in the first rack. This deployment explores how project and pod node selectors interact with each other. Thus, you verify that pods that the deployment creates have both the project and pod node selectors.

In the third scenario, the student taints the `worker01` and `worker03` nodes to reserve some workload capacity, because your cluster has limited capacity. The taint is `type=mission-critical` with the `NoSchedule` option, which ensures that previous node workloads continue working but that new workloads need a toleration to be placed in those nodes. In this scenario, you first try to create a deployment with a node selector for the fast CPU label, to meet your application's requirements. Because your deployment does not include a toleration for the taint, the deployment is not available. Then, you create another deployment, to include the node selector and the toleration, to verify that the deployment is available.

Create all the deployments as the developer user.

- ▶ 1. As the `admin` user, connect to the OpenShift cluster and verify the labels for the nodes in the cluster.
 - 1.1. Connect to the OpenShift cluster as the `admin` user with `redhatocp` as the password.

```
[student@workstation ~]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
Login successful.
...output omitted...
```

- 1.2. List the labels for the compute nodes. Verify the following node labels:
 - The `rack=1` and `cpu=standard` labels for the `worker01` node
 - The `rack=2` and `cpu=standard` labels for the `worker02` node
 - The `rack=3` and `cpu=fast` labels for the `worker03` node

NAME	STATUS	ROLES	AGE	VERSION	RACK	CPU
master01	Ready	control-plane, master	140d	v1.25.7+eab9cc9		
master02	Ready	control-plane, master	140d	v1.25.7+eab9cc9		
master03	Ready	control-plane, master	140d	v1.25.7+eab9cc9		
worker01	Ready	worker	37d	v1.25.7+eab9cc9	1	standard
worker02	Ready	worker	37d	v1.25.7+eab9cc9	2	standard
worker03	Ready	worker	37d	v1.25.7+eab9cc9	3	fast

- ▶ 2. Create a deployment called `myapp` with two replicas. Verify the default OpenShift scheduler behavior, which distributes the deployment pods between the three available compute nodes.
 - 2.1. Connect to the OpenShift cluster as the `developer` user with `developer` as the password, and verify that OpenShift uses the `scheduling-selector` project.

```
[student@workstation ~]$ oc login -u developer -p developer
...output omitted...
Using project "scheduling-selector".
```

- 2.2. Change to the `~/D0380/labs/scheduling-selector` directory.

Chapter 4 | Pod Scheduling

```
[student@workstation ~]$ cd ~/D0380/labs/scheduling-selector
```

- 2.3. Create a deployment CR YAML file with two replicas. You can find an incomplete example for the deployment CR in the `~/D0380/labs/scheduling-selector/deployment.yml` file.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp
spec:
  replicas: 2
  selector:
    matchLabels:
      app: myapp
...output omitted...
```

- 2.4. Apply the configuration for the deployment CR.

```
[student@workstation scheduling-selector]$ oc create -f deployment.yml
deployment.apps/myapp created
```

- 2.5. Verify that the deployment is correctly created with two replicas. Wait until all the pods are marked as ready and available. You might have to repeat this command many times.

```
[student@workstation scheduling-selector]$ oc get deployment
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
myapp     2/2     2           2           43s
```

- 2.6. Verify that the OpenShift scheduler distributes the two pods between the available compute nodes. The pod names and the nodes that run on your pods might differ in your system.

```
[student@workstation scheduling-selector]$ oc get pods -o wide
NAME                  READY   STATUS    ...   NODE    ...
myapp-847cf74d9-84tqf  1/1     Running   ...   worker02 ...
myapp-847cf74d9-npjrn  1/1     Running   ...   worker01 ...
```

- ▶ 3. Your company requires you to create a deployment, with two replicas, which needs a fast CPU. Create a deployment called `myapp-ns-fastcpu` with two replicas. The deployment must have a node selector for the `fastCPU` label.

- 3.1. Create a deployment CR YAML file with two replicas and a node selector for the `fastCPU` label. You can find an incomplete example for the deployment in the `~/D0380/labs/scheduling-selector/deployment-ns-fastcpu.yml` file.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp-ns-fastcpu
```

Chapter 4 | Pod Scheduling

```
spec:
  replicas: 2
  selector:
    matchLabels:
      app: myapp-ns-fastcpu
  template:
    metadata:
      labels:
        app: myapp-ns-fastcpu
    spec:
      ...output omitted...
      nodeSelector:
        cpu: fast
```

- 3.2. Apply the configuration for the deployment CR.

```
[student@workstation scheduling-selector]$ oc create \
-f deployment-ns-fastcpu.yml
deployment.apps/myapp-ns-fastcpu created
```

- 3.3. Verify that the deployment is correctly created with two replicas. Wait until all the pods are marked as ready and available. You might have to repeat this command many times.

```
[student@workstation scheduling-selector]$ oc get deployment
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
myapp         2/2     2           2           2m57s
myapp-ns-fastcpu   2/2     2           2           27s
```

- 3.4. Verify that the OpenShift scheduler creates two pods on the `worker03` node, because that node has the `cpu=fast` label. Pod names might differ in your system.

```
[student@workstation scheduling-selector]$ oc get pods -o wide \
-l app=myapp-ns-fastcpu
NAME                                     READY   STATUS    ...   NODE   ...
myapp-ns-fastcpu-5577dd75d8-m6s9f   1/1     Running   ...   worker03   ...
myapp-ns-fastcpu-5577dd75d8-pcw7j   1/1     Running   ...   worker03   ...
```

- ▶ 4. Drain the `worker03` node from the cluster to simulate a failure in the third rack. Verify the status for the deployments. The `myapp` deployment continues to be available, because the OpenShift scheduler can place its pods in the `worker01` and `worker02` nodes. However, the node selector in the `myapp-ns-fastcpu` deployment does not enable the scheduler to place the pods in the available compute nodes. Thus, the deployment is not available.

- 4.1. Connect to the OpenShift cluster as the `admin` user with `redhatocp` as the password.

```
[student@workstation scheduling-selector]$ oc login -u admin -p redhatocp
...output omitted...
```

- 4.2. Drain all the pods from the `worker03` node. This action might take a few minutes.

Chapter 4 | Pod Scheduling

```
[student@workstation scheduling-selector]$ oc adm drain worker03 \
--ignore-daemonsets --delete-emptydir-data
...output omitted...
node/worker03 drained
```

- 4.3. Verify the status for the deployments. Although the myapp deployment is available, the myapp-ns-fastcpu deployment is not available due to the node selector.

```
[student@workstation scheduling-selector]$ oc get deployment
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
myapp         2/2     2           2           10m
myapp-ns-fastcpu 0/2     2           0           6m4s
```

- 4.4. Verify that the scheduler cannot use the worker01 and worker02 nodes for the myapp-ns-fastcpu deployment due to the node selector.

```
[student@workstation scheduling-selector]$ oc get pods -o wide
NAME                           READY   STATUS    ...   NODE   ...
myapp-847cf74d9-84tqf          1/1     Running   ...   worker02   ...
myapp-847cf74d9-npjrn          1/1     Running   ...   worker01   ...
myapp-ns-fastcpu-5577dd75d8-b422c 0/1     Pending   ...   <none>   ...
myapp-ns-fastcpu-5577dd75d8-vdtsx 0/1     Pending   ...   <none>   ...
```

- 4.5. Review the project events. The OpenShift scheduler does not find a node for the myapp-ns-fastcpu deployment due to the node selector.

```
[student@workstation scheduling-selector]$ oc get events \
--sort-by='.lastTimestamp'
...output omitted...
4m11s      Normal    SuccessfulCreate  replicaset/myapp-ns-fastcpu-5577dd75d8
Created pod: myapp-ns-fastcpu-5577dd75d8-vdtsx
4m11s      Normal    SuccessfulCreate  replicaset/myapp-ns-fastcpu-5577dd75d8
Created pod: myapp-ns-fastcpu-5577dd75d8-b422c
4m11s      Warning   FailedScheduling  pod/myapp-ns-fastcpu-5577dd75d8-b422c
  0/6 nodes are available: 1 node(s) were unschedulable, 2 node(s) didn't
  match Pod's node affinity/selector, 3 node(s) had untolerated taint {node-
  role.kubernetes.io/master: }. preemption: 0/6 nodes are available: 6 Preemption is
  not helpful for scheduling.
```

- 4.6. Review the pod information. Pod names might differ in your system.

Chapter 4 | Pod Scheduling

```
[student@workstation scheduling-selector]$ oc describe pod \
myapp-ns-fastcpu-5577dd75d8-b422c
...output omitted...
Events:
  Type      Reason          Age           From            Message
  ----      ----          ----          ----            -----
  Warning   FailedScheduling 2m10s (x2 over 7m12s) default-scheduler 0/6 nodes
are available: 1 node(s) were unschedulable, 2 node(s) didn't match Pod's node
affinity/selector, 3 node(s) had untolerated taint {node-role.kubernetes.io/
master: }. preemption: 0/6 nodes are available: 6 Preemption is not helpful for
scheduling.
```

- 4.7. Mark the `worker03` node as schedulable.

```
[student@workstation scheduling-selector]$ oc adm uncordon worker03
node/worker03 uncordoned
```

- ▶ 5. Create the `scheduling-ns` project with a project node selector for the standard CPU nodes.

- 5.1. As the `admin` user, create the project CR YAML file with a node selector for the standard CPU nodes. You can find an incomplete example for the project in the `~/DO380/labs/scheduling-selector/project-scheduling.yml` file.

```
apiVersion: v1
kind: Namespace
metadata:
  name: scheduling-ns
  annotations:
    openshift.io/node-selector: cpu=standard
```

- 5.2. Apply the configuration for the project CR.

```
[student@workstation scheduling-selector]$ oc create \
-f project-scheduling.yml
namespace/scheduling-ns created
```

- 5.3. Give edit permission to the `developer` user in the `scheduling-ns` project.

```
[student@workstation scheduling-selector]$ oc adm policy add-role-to-user \
edit developer -n scheduling-ns
clusterrole.rbac.authorization.k8s.io/edit added: "developer"
```

- 5.4. Connect to the OpenShift cluster as the `developer` user with `developer` as the password.

```
[student@workstation scheduling-selector]$ oc login -u developer -p developer
...output omitted...
```

- 5.5. Change to the `scheduling-ns` project.

Chapter 4 | Pod Scheduling

```
[student@workstation scheduling-selector]$ oc project scheduling-ns
...output omitted...
```

- 5.6. Verify that the project includes the node selector.

```
[student@workstation scheduling-selector]$ oc describe project scheduling-ns
Name:   scheduling-ns
Created: 4 minutes ago
...output omitted...
Annotations:    openshift.io/node-selector=cpu=standard
                  openshift.io/sa.scc.mcs=s0:c26,c25
                  openshift.io/sa.scc.supplemental-groups=1000700000/10000
                  openshift.io/sa.scc.uid-range=1000700000/10000
Display Name:  <none>
Description:   <none>
Status:        Active
Node Selector:  cpu=standard
Quota:         <none>
Resource limits: <none>
...output omitted...
```

- 5.7. Create a deployment CR YAML file with two replicas. You can find an incomplete example for the deployment in the ~/D0380/labs/scheduling-selector/deployment-project-ns.yml file.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: project-ns
spec:
  replicas: 2
  selector:
    matchLabels:
      app: project-ns
...output omitted...
```

- 5.8. Apply the configuration for the deployment CR.

```
[student@workstation scheduling-selector]$ oc create -f deployment-project-ns.yml
deployment.apps/project-ns created
```

- 5.9. Verify the status for the deployment. Wait until all the pods are marked as ready and available. You might have to repeat this command many times.

```
[student@workstation scheduling-selector]$ oc get deployment
NAME        READY   UP-TO-DATE   AVAILABLE   AGE
project-ns  2/2     2           2           131s
```

- 5.10. Verify that the scheduler uses the worker01 and worker02 nodes due to the project node selector.

```
[student@workstation scheduling-selector]$ oc get pods -o wide
NAME           READY   STATUS    ...   NODE      ...
project-ns-79c4798c49-k8c92  1/1     Running   ...   worker01   ...
project-ns-79c4798c49-r2dn7  1/1     Running   ...   worker02   ...
```

- 5.11. Verify that the pods in the deployment include the project node selector. Pod names might differ in your system.

```
[student@workstation scheduling-selector]$ oc describe pod \
project-ns-79c4798c49-k8c92
Name:           project-ns-79c4798c49-k8c92
Namespace:      scheduling-ns
...output omitted...
Node-Selectors:  cpu=standard
Tolerations:    node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
                 node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
...output omitted...
```

- ▶ 6. Create a deployment with two replicas and a node selector for the first rack. Verify that the deployment includes both the project and the pod node selectors.
- 6.1. Create a deployment CR YAML file with two replicas and a node selector for the fast CPU node. You can find an incomplete example for the deployment in the ~/DO380/labs/scheduling-selector/deployment-project-podsel.yml file.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: project-podsel
spec:
  replicas: 2
  selector:
    matchLabels:
      app: project-podsel
  template:
    metadata:
      labels:
        app: project-podsel
    spec:
  ...output omitted...
  nodeSelector:
    rack: "1"
```

- 6.2. Apply the configuration for the deployment CR.

```
[student@workstation scheduling-selector]$ oc create \
-f deployment-project-podsel.yml
deployment.apps/project-podsel created
```

- 6.3. Verify the status for the deployment. Wait until all the pods are marked as ready and available. You might have to repeat this command many times.

Chapter 4 | Pod Scheduling

```
[student@workstation scheduling-selector]$ oc get deployment
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
project-ns    2/2     2           2           12m
project-podsel 2/2     2           2           72s
```

- 6.4. Verify that the OpenShift scheduler places all the pods for the `project-podsel` deployment in the `worker01` node, because the pods have both labels.

```
[student@workstation scheduling-selector]$ oc get pods -o wide \
-l app=project-podsel
NAME                           READY   STATUS    ...   NODE   ...
project-podsel-649f68f65d-44xxk 1/1     Running  ...   worker01 ...
project-podsel-649f68f65d-qvvjz 1/1     Running  ...   worker01 ...
```

- 6.5. Verify that the pods in the deployment include both the project and the pod node selectors. Pod names might differ in your system.

```
[student@workstation scheduling-selector]$ oc describe pod \
project-podsel-649f68f65d-44xxk
Name:           project-podsel-649f68f65d-44xxk
Namespace:      scheduling-ns
...output omitted...
Node-Selectors: cpu=standard
                rack=1
Tolerations:   node.kubernetes.io/not-ready:NoExecute op=Exists for
300s
                  node.kubernetes.io/unreachable:NoExecute op=Exists
for 300s
...output omitted...
```

- 7. As an OpenShift administrator, apply the `type=mission-critical` taint to the `worker01` and `worker03` nodes to reserve workload capacity. Use the `NoSchedule` option to ensure that previous node workloads continue working.

- 7.1. Connect to the OpenShift cluster as the `admin` user with `redhatocp` as the password.

```
[student@workstation scheduling-selector]$ oc login -u admin -p redhatocp
...output omitted...
```

- 7.2. Apply the `type=mission-critical:NoSchedule` taint to the `worker01` and `worker03` nodes.

```
[student@workstation scheduling-selector]$ oc adm taint nodes worker01 \
type=mission-critical:NoSchedule
node/worker01 tainted

[student@workstation scheduling-selector]$ oc adm taint nodes worker03 \
type=mission-critical:NoSchedule
node/worker03 tainted
```

Chapter 4 | Pod Scheduling

- 7.3. Verify the taints for the `worker01` and `worker03` nodes.

```
[student@workstation scheduling-selector]$ oc describe node worker01 | grep Taints
Taints:    type=mission-critical:NoSchedule
```

```
[student@workstation scheduling-selector]$ oc describe node worker03 | grep Taints
Taints:    type=mission-critical:NoSchedule
```

- 8. Create a deployment with two replicas with a node selector for the `cpu=fast` label. Verify that the deployment is not available because the two nodes with the `cpu=fast` label have a taint, and the deployment does not have a toleration for the taint.

- 8.1. Connect to the OpenShift cluster as the `developer` user with `developer` as the password.

```
[student@workstation scheduling-selector]$ oc login -u developer -p developer
...output omitted...
```

- 8.2. Create the `scheduling-taint` project.

```
[student@workstation scheduling-selector]$ oc new-project scheduling-taint
...output omitted...
```

- 8.3. Create a deployment CR YAML file with two replicas and a node selector for the `cpu=fast` label. You can find an incomplete example for the deployment in the `/home/student/D0380/labs/scheduling-selector/deployment-taint-fastcpu.yaml` file.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp-taint-fastcpu
spec:
  replicas: 2
  selector:
    matchLabels:
      app: myapp-taint-fastcpu
  template:
    metadata:
      labels:
        app: myapp-taint-fastcpu
    spec:
      ...output omitted...
      nodeSelector:
        cpu: fast
```

- 8.4. Apply the configuration for the deployment CR.

```
[student@workstation scheduling-selector]$ oc create \
-f deployment-taint-fastcpu.yaml
deployment.apps/myapp-taint-fastcpu created
```

Chapter 4 | Pod Scheduling

- 8.5. Verify the status for the deployment. The `myapp-taint-fastcpu` deployment is not available due to node taint without a toleration.

```
[student@workstation scheduling-selector]$ oc get deployment
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
myapp-taint-fastcpu   0/2       2           0           104s
```

- 8.6. Verify that the OpenShift scheduler cannot create the `myapp-taint-fastcpu` pods in a compute node.

```
[student@workstation scheduling-selector]$ oc get pods -o wide
NAME                           READY   STATUS    ...   NODE   ...
myapp-taint-fastcpu-5c646fb6b-k5r9g   0/1     Pending   ...   <none> ...
myapp-taint-fastcpu-5c646fb6b-tz2nr   0/1     Pending   ...   <none> ...
```

- 8.7. Review the project events. The OpenShift scheduler does not find a node for the `myapp-taint-fastcpu` deployment due to the node selector and the node taint.

```
[student@workstation scheduling-selector]$ oc get events \
--sort-by='-.lastTimestamp'
...output omitted...
2m11s      Warning  FailedScheduling  pod/myapp-taint-fastcpu-5c646fb6b-
tz2nr  0/6 nodes are available: 1 node(s) didn't match Pod's node affinity/
selector, 2 node(s) had untolerated taint {type: mission-critical}, 3 node(s) had
untolerated taint {node-role.kubernetes.io/master: }. preemption: 0/6 nodes are
available: 6 Preemption is not helpful for scheduling.
2m11s      Normal   SuccessfulCreate  replicaset/myapp-taint-
fastcpu-5c646fb6b  Created pod: myapp-taint-fastcpu-5c646fb6b-tz2nr
2m11s      Normal   SuccessfulCreate  replicaset/myapp-taint-
fastcpu-5c646fb6b  Created pod: myapp-taint-fastcpu-5c646fb6b-k5r9g
2m11s      Normal   ScalingReplicaSet  deployment/myapp-taint-fastcpu
Scaled up replica set myapp-taint-fastcpu-5c646fb6b to 2
```

- 8.8. Remove the `myapp-taint-fastcpu` deployment.

```
[student@workstation scheduling-selector]$ oc delete deployment \
myapp-taint-fastcpu
deployment.apps "myapp-taint-fastcpu" deleted
```

- ▶ 9. Create a deployment with the node selector for the `cpu=fast` label and the toleration for the taint. Verify that the deployment is available.

- 9.1. Modify the YAML file for the `myapp-taint-fastcpu` deployment CR by adding a toleration for the `type:mission-critical:NoSchedule` taint.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp-taint-fastcpu
spec:
  replicas: 2
  selector:
```

Chapter 4 | Pod Scheduling

```

matchLabels:
  app: myapp-taint-fastcpu
template:
  metadata:
    labels:
      app: myapp-taint-fastcpu
spec:
  ...output omitted...
  nodeSelector:
    cpu: fast
  tolerations:
  - key: "type"
    value: "mission-critical"
    operator: "Equal"
    effect: "NoSchedule"

```

- 9.2. Apply the configuration for the deployment CR.

```
[student@workstation scheduling-selector]$ oc create -f \
deployment-taint-fastcpu.yml
deployment.apps/myapp-taint-fastcpu created
```

- 9.3. Verify the status for the deployment. The `myapp-taint-fastcpu` deployment is available because it includes a toleration for the node taint.

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
myapp-taint-fastcpu	2/2	2	2	62s

- 9.4. Verify that the OpenShift scheduler places the two pods for the `myapp-taint-fastcpu` deployment in the `worker03` node, because the deployment has a toleration for the taint, and also because this node includes the specified label in the node selector.

```
[student@workstation scheduling-selector]$ oc get pods -o wide
NAME                               READY   STATUS    ...   NODE   ...
myapp-taint-fastcpu-86cc9d87cc-7x4n8  1/1    Running   ...   worker03   ...
myapp-taint-fastcpu-86cc9d87cc-8p4c4  1/1    Running   ...   worker03   ...
```

- 9.5. Change to the student HOME directory.

```
[student@workstation scheduling-selector]$ cd
```

Finish

On the `workstation` machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish scheduling-selector
```

High Availability with Affinity Rules and Pod Disruption Budgets

Objectives

- Configure workloads for resilience against node failures.

Scheduling Affinity

You can instruct the Kubernetes scheduler how to distribute the pods across the compute nodes in the cluster. The affinity rules enable the Kubernetes scheduler to set a preference to run the pods in a given cluster node, or to avoid running pods on a node where another set of pods is running.

Cluster administrators can set topology labels in the nodes to indicate the failure domain that they belong to. Application developers can configure the deployment resources to use the pod anti-affinity settings with custom topology labels to distribute the workloads across the failure domains.

Sometimes, the applications require multiple replicas to be available if a cluster disruption occurs. This type of requirement comes from load testing, which might indicate a need for a given number of replicas to give a suitable response time for users. The pod disruption budget resources set the minimum availability constraints to require that a given number or percentage of pods remain running even when a node is drained.

Pod Affinity and Anti-affinity

Pod affinity is a group of rules that are set in the pod specification. The scheduler uses the rules to place pods in the same cluster node where another workload is present, or to avoid placing pods in the same node where the pods from other workloads are running.

The pod affinity terms define a set of selectors to match the labels of other pods that are running in the cluster nodes. The scheduler can evaluate the pod affinity terms on a best-effort basis if the affinity setting is preferred, or require the conditions to be present to schedule the pod.

Kubernetes implements two types of pod affinity rules: *pod affinity* and *pod anti-affinity*.

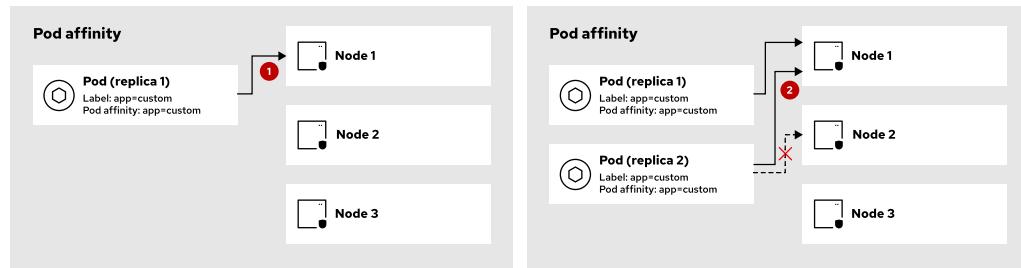
Pod affinity

The pod affinity setting enables scheduling of pods in the same node where the pods from another workload are running. The pod specification contains the pod affinity setting, which uses a selector to match the labels from other pods by using an operator.

Pod affinity helps reduce the network latency between specific pods. Pod affinity can help improve the performance of workloads that require significant communication between pods.

In the following image, the pods have a pod affinity setting, to direct the scheduler to place them in the same node. The deployment is configured to run two replicas.

Chapter 4 | Pod Scheduling



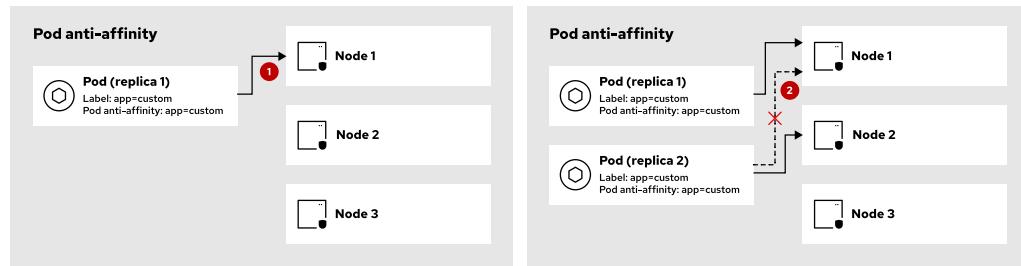
- ➊ The scheduler places the first pod in the first node. No preference exists, because no pods match the `app=custom` label selector.
- ➋ The scheduler looks for pods that match the `app=custom` label selector and sets a preference to run the second pod in the same node as the first pod. The second pod is then scheduled in the first node.

Pod anti-affinity

The pod anti-affinity setting prevents pods from being scheduled in the node where the pods from another workload are running. You can use a selector to match the labels from other pods. The pod specification contains the pod affinity setting that uses a selector to match the labels from other pods.

Pod anti-affinity prevents some workloads from running on the same node. Pod anti-affinity can help distribute a workload across failure domains to improve reliability, or prevent interference between workloads.

In the following image, the pods have a pod anti-affinity setting and direct the scheduler not to place them in the same node. The deployment is configured to run two replicas.



- ➊ The scheduler places the first pod in the first node. No preference exists, because no pods match the `app=custom` label selector.
- ➋ The scheduler looks for pods that match the `app=custom` label selector and sets a preference to avoid running the second pod in the same node as the first pod. The second pod is then scheduled on the second node.

Similar to node affinity, you can apply the required or preferred affinity terms to the pod affinity or pod anti-affinity settings.



Note

The scheduler evaluates the affinity terms only when the pod is created.

Required

The pod affinity terms are enforced when the pod is created, and the scheduler does not move the pod to another node if the conditions change.

The pod affinity rules include a list of selectors to match the labels of pods that are running in a cluster node.

Preferred

The scheduler evaluates the pod affinity terms on a best-effort basis. Thus, the scheduler might choose a node that does not satisfy all the requirements if no node matches all the conditions. The scheduler does not move the pod if the conditions change.

The preferred scheduling term has a list of weighted pod selector terms. Each term can match a label from running pods. Each rule has a weight value that is defined in the range 1-100 inclusive.

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
  labels:
    app: custom
spec:
  affinity:
    podAntiAffinity:
      preferredDuringSchedulingIgnoredDuringExecution: ①
      - weight: 100 ②
        podAffinityTerm:
          labelSelector: ③
          matchExpressions:
            - key: app
              operator: In ④
              values:
                - custom
          topologyKey: rack ⑤
...content omitted...
```

- ① The expressions are evaluated on a best-effort basis and the pod is scheduled in a node even if the constraints are not met.
- ② The weight value is between 1-100 inclusive.
- ③ The expression that matches a node label.
- ④ The operator that matches the labels.
- ⑤ The topology key that the scheduler uses to distribute pods across the failure domains.

Topology Keys

The Kubernetes scheduler can use node labels to indicate the failure domain that each node belongs to. The failure domain can be a rack that is connected to a different power source or networking equipment, or an indicator that the node is in another building. Kubernetes uses the following node labels as default topology keys to distribute the workloads across the cluster nodes:

kubernetes.io/hostname

This label is set to match the hostname value of each node in the cluster.

topology.kubernetes.io/zone

Represents the availability zone of the node. This label is set by the cloud provider and might not be present in on-premise deployments.

topology.kubernetes.io/region

Represents the region of the node. This label is set by the cloud provider and might not be present in on-premise deployments.

The topology key helps the scheduler to spread the pods across the failure domains. The cluster administrator must label all the nodes to indicate their corresponding failure domain. Application developers can use the node label as a custom topology key in the pod affinity terms, to direct the scheduler how to place the pods in the cluster nodes.

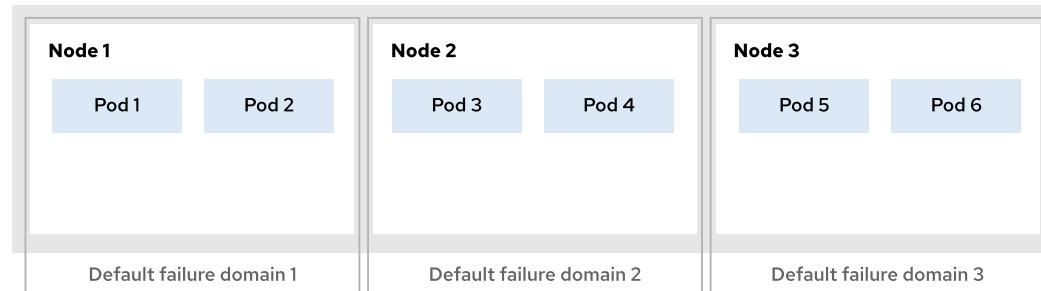
The following scenario describes a cluster distributed between two racks. Each rack is connected to a different power source. The nodes are distributed by rack according to the following table:

Topology Key	Control plane nodes	Compute nodes
rack=rack-1	master02 , master03	node01
rack=rack-2	master01	node02 , node03

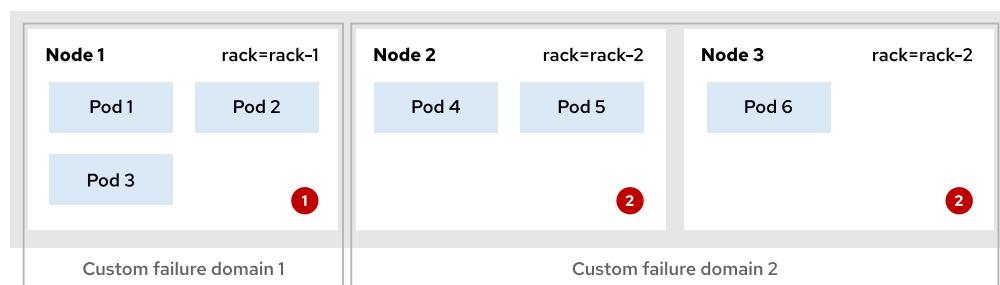
The administrator added the `rack` label to indicate the location and failure domain of each node. This label is intended as a custom topology key, so that the scheduler can spread the pods evenly across the compute nodes in different racks.

Without a custom topology key

The value of the `kubernetes.io/hostname` topology key is different on each node, and the scheduler identifies each node as a separate failure domain. The scheduler uses the default topology keys, and distributes the pods evenly across all the compute nodes.

**With a custom topology key**

The `rack` topology key specifies the failure domain that each node belongs to. The scheduler spreads pods evenly across the failure domains.



- ➊ Three pods are placed in the `node01` compute node in the `rack-1` failure domain.
- ➋ The other three pods are placed in the nodes in the `rack-2` failure domain.

Pod Disruption Budgets

The *pod disruption budget* is a policy object where you can specify the percentage or number of pods that must remain running when a voluntary disruption occurs in the cluster. The pod disruption budget resource is part of the `policy/v1` API group.

Kubernetes clusters have two types of disruptions: involuntary and voluntary.

Involuntary disruptions

Involuntary disruptions occur when a node fails and is disconnected from the cluster. The involuntary disruptions could be related to a hardware problem, a network issue within a zone, or a power issue in a rack. The Kubernetes cluster detects that a node is not responsive, creates replacement pods in another node, and eventually deletes all the pods that are scheduled on the failed node.

Voluntary disruptions

Voluntary disruptions occur when nodes are cordoned or taken offline for maintenance. When a voluntary disruption occurs, all the pods on the affected node are evicted at the same time, and the scheduler creates replacement pods in another node. If all the pods of a workload are running in the same node and that node is drained, then the minimum availability constraints for the application are not fulfilled.

You can define the pod disruption budget resource with either the `minAvailable` or the `maxUnavailable` attribute. The values for these settings are often discovered after load testing on the application. Sometimes, a given number of pods must run all the time to prevent the application from queuing the user requests. Another example is that the application delivers an acceptable response time for the users even if only 60% of the pods are running.

`minAvailable`

The minimum number of pods to be available, even during a voluntary disruption. You can set this attribute to a percentage or integer value to specify the minimum number of replica pods to be available.

```
apiVersion: policy/v1
kind: PodDisruptionBudget
metadata:
  name: nginx-minavailable
spec:
  minAvailable: 80% ➊
```

Chapter 4 | Pod Scheduling

```
selector: ②
  matchLabels:
    app: nginx
```

- ① The minimum percentage or number of pods to be available.
- ② The label selector for the affected pods.

maxUnavailable

The maximum number of pods that can be unavailable during a voluntary disruption. You can set this attribute to a percentage or integer value to specify the maximum number of replica pods that can be unavailable.

```
apiVersion: policy/v1
kind: PodDisruptionBudget
metadata:
  name: nginx-maxunavailable
spec:
  maxUnavailable: 33% ①
  selector: ②
  matchLabels:
    app: nginx
```

- ① The maximum percentage or number of pods that can be unavailable.
- ② The label selector for the affected pods.

You can list or describe the pod disruption budget resources to inspect their properties.

NAME	MIN AVAILABLE	MAX UNAVAILABLE	ALLOWED DISRUPTIONS	AGE
nginx-minavailable	80%	N/A	2	45s ①

- ① The minimum percentage or number of pods to be available.

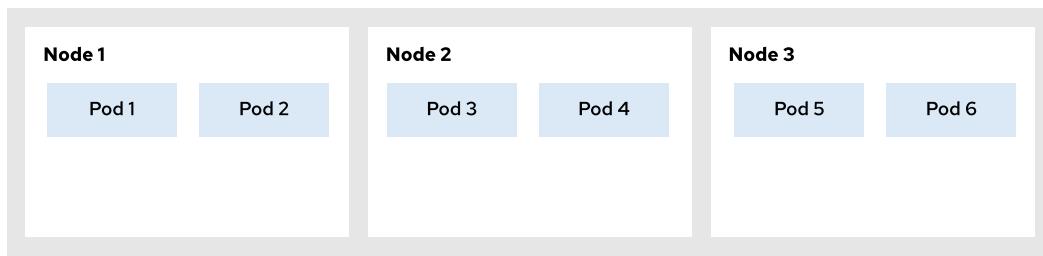
You can also describe the resource to view the label selector for the affected pods.

```
[user@host ~]$ oc describe pdb nginx-maxunavailable
Name:          nginx-maxunavailable
Namespace:     scheduling-pdb
Max unavailable: 33% ①
Selector:      app=nginx ②
Status:
  Allowed disruptions: 4
  Current:           10
  Desired:            6
  Total:              10
Events:        <none>
```

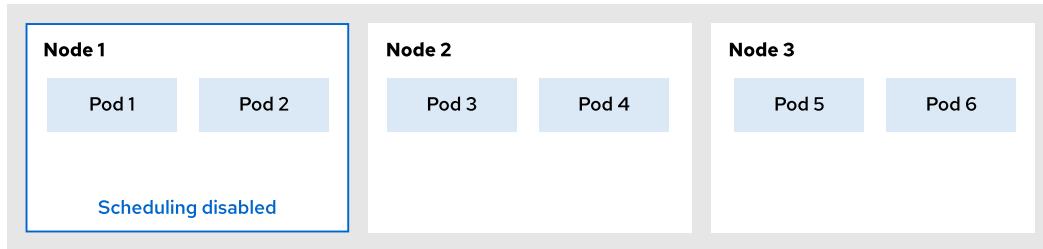
- ① The maximum percentage or number of pods that can be unavailable.
- ② The label selector for the affected pods.

Node Drain Without Pod Disruption Budget

The deployment is configured to run six replica pods that are distributed between the compute nodes. The application has an appropriate response time if five or more pods are running.



The first node is drained by using the `oc adm drain` command. The node is cordoned and marked as unschedulable.

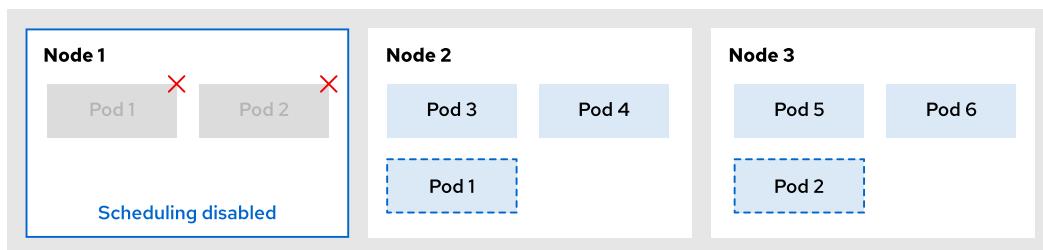


```
[user@host ~]$ oc adm drain node/node01 ...
node/node01 cordoned
Warning: ignoring DaemonSet-managed Pods: ...output omitted...
... request.go:682] Waited for ... due to client-side throttling ...output omitted...
...output omitted...
```

**Note**

You can safely ignore the warnings about managed pods and client-side throttling.

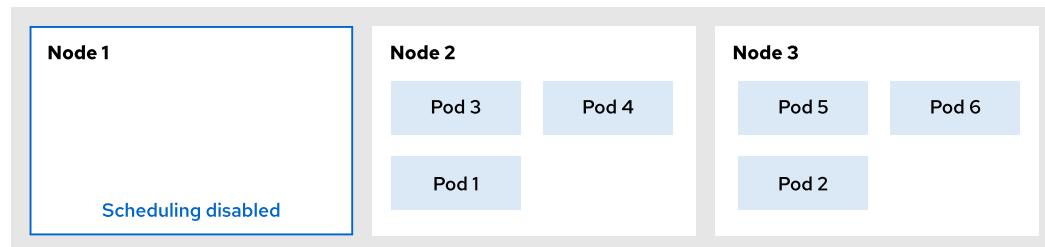
All the pods that are running in the drained node are marked for eviction. The scheduler creates replacement pods in another node. Only four pods are ready and the application starts responding slowly.



Chapter 4 | Pod Scheduling

```
...output omitted...
evicting pod scheduling-pdb/nginx-56bdf7d8c6-kf98k
evicting pod scheduling-pdb/nginx-56bdf7d8c6-8kk2b
pod/nginx-56bdf7d8c6-kf98k evicted
pod/nginx-56bdf7d8c6-8kk2b evicted
...output omitted...
```

The compute node is finally marked as drained. The replacement pods are marked as ready. The application now runs all six pods and the response time is within user expectations.



```
...output omitted...
node/node01 drained
```

The full output of the `oc adm drain` command is as follows:

```
[user@host ~]$ oc adm drain node/node01 --ignore-daemonsets --delete-emptydir-data
node/node01 cordoned
Warning: ignoring DaemonSet-managed Pods: ...output omitted...
... request.go:682] Waited for ... due to client-side throttling ...output omitted...
...output omitted...
evicting pod scheduling-pdb/nginx-56bdf7d8c6-kf98k ①
evicting pod scheduling-pdb/nginx-56bdf7d8c6-8kk2b
...output omitted...
pod/nginx-56bdf7d8c6-kf98k evicted ②
pod/nginx-56bdf7d8c6-8kk2b evicted
...output omitted...
node/node01 drained
```

- ① All the pods are marked for eviction when the node is drained.
- ② All the pods are evicted from the node at the same time and the application availability constraint is not met.

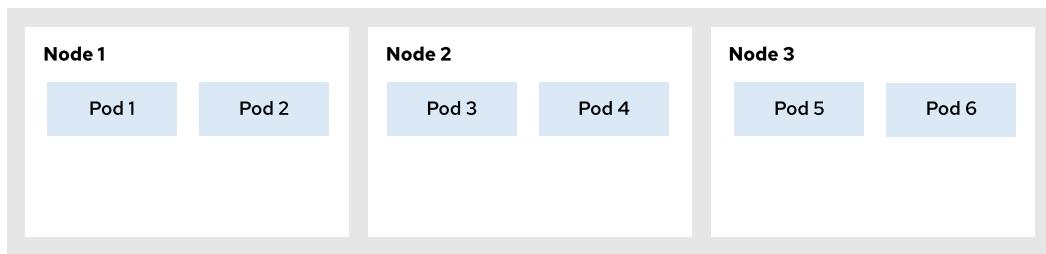
Node Drain with Pod Disruption Budget

With a pod disruption budget resource, you can specify the minimum availability constraints for the application. The node drain is blocked until the replacement pods for the application are scheduled in another node and become ready. The pod eviction is blocked until the pod disruption budget availability constraints are met. The pod eviction operation retries after five seconds.

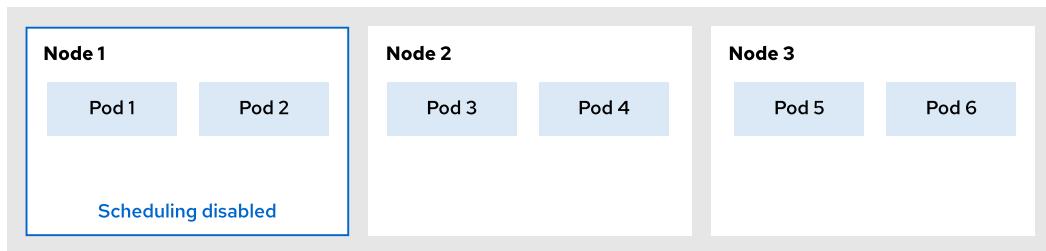
The deployment is configured to run six replica pods that are distributed between the compute nodes. The application has an appropriate response time if five or more pods are running. The

Chapter 4 | Pod Scheduling

developer creates a pod disruption budget resource where the `spec.minAvailable` parameter is set to five, and specifies the selector to match the labels of the deployment pods.



The first node is drained by using the `oc adm drain` command. The node is cordoned and marked as unschedulable.



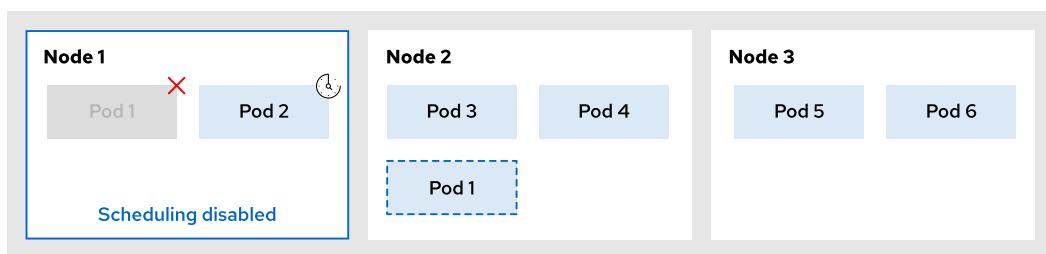
```
[user@host ~]$ oc adm drain node/node01 ...
node/node01 cordoned
Warning: ignoring DaemonSet-managed Pods: ...output omitted...
... request.go:682] Waited for ... due to client-side throttling ...output omitted...
...output omitted...
```



Note

You can safely ignore the warnings about managed pods and client-side throttling.

All the pods that are running in the drained node are marked for eviction. The scheduler creates a replacement pod in another node and waits for the new pod to be ready. The second pod cannot be evicted, because the availability constraint states that five out of six pods must be running. The eviction operation for the second pod retries after five seconds.



Chapter 4 | Pod Scheduling

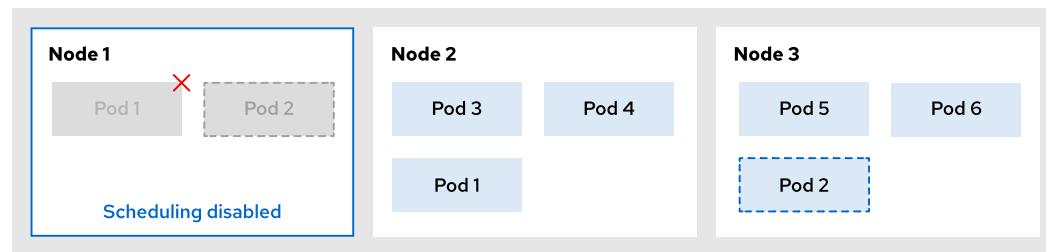
```
...output omitted...
evicting pod scheduling-pdb/nginx-647b4c98f5-nqvs6
evicting pod scheduling-pdb/nginx-647b4c98f5-98ptm
error when evicting pods/"nginx-647b4c98f5-98ptm" -n "scheduling-pdb" (will retry
after 5s): Cannot evict pod as it would violate the pod's disruption budget.
...output omitted...
```



Important

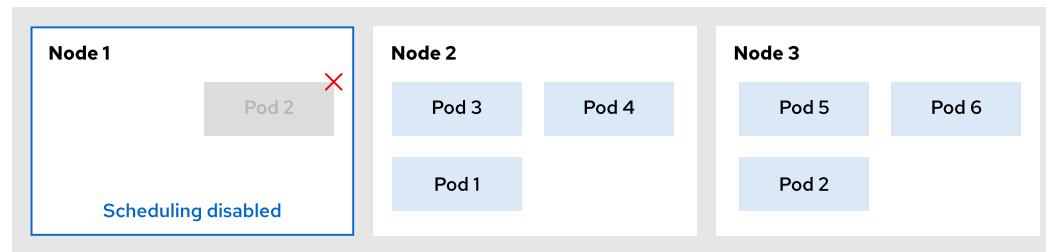
The drain process is blocked until all the pods are evicted from the node.

The first replacement pod is marked as ready, and the first pod is evicted from the drained node. The scheduler creates a second replacement pod in another node and waits for it to be ready.



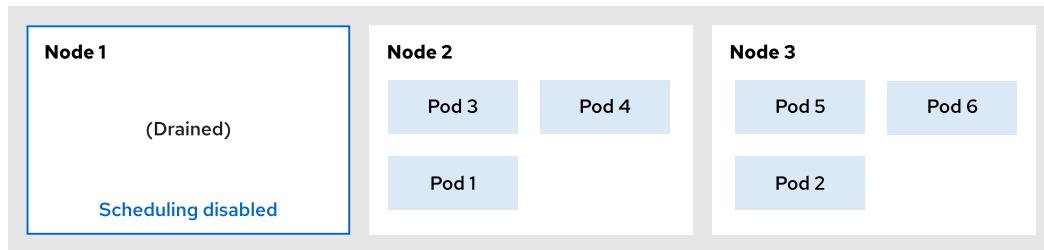
```
...output omitted...
pod/nginx-647b4c98f5-nqvs6 evicted
evicting pod scheduling-pdb/nginx-647b4c98f5-98ptm
...output omitted...
```

The second replacement pod is marked as ready, and the second pod is evicted from the drained node.



```
...output omitted...
pod/nginx-647b4c98f5-98ptm evicted
...output omitted...
```

The compute node is finally marked as drained.



```
...output omitted...
node/node01 drained
```

The full output of the `oc adm drain` command is as follows:

```
[user@host ~]$ oc adm drain node/node01 --ignore-daemonsets --delete-emptydir-data
node/node01 cordoned
Warning: ignoring DaemonSet-managed Pods: ...output omitted...
... request.go:682] Waited for ... due to client-side throttling ...output omitted...
...output omitted...
evicting pod scheduling-pdb/nginx-647b4c98f5-nqvs6
evicting pod scheduling-pdb/nginx-647b4c98f5-98ptm ①
error when evicting pods/"nginx-647b4c98f5-98ptm" -n "scheduling-pdb" (will retry
after 5s): Cannot evict pod as it would violate the pod's disruption budget. ②
...output omitted...
error when evicting pods/"nginx-647b4c98f5-98ptm" -n "scheduling-pdb" (will retry
after 5s): Cannot evict pod as it would violate the pod's disruption budget.
pod/nginx-647b4c98f5-nqvs6 evicted
evicting pod scheduling-pdb/nginx-647b4c98f5-98ptm
...output omitted...
pod/nginx-647b4c98f5-98ptm evicted ③
...output omitted...
node/node01 drained
```

- ① The pod is marked for eviction.
- ② The pod eviction is blocked until the availability constraints from the pod disruption budget are met.
- ③ The pod is finally evicted from the drained node.



Note

You can safely ignore the warnings about managed pods and client-side throttling.



References

For more information about pod scheduling on OpenShift, refer to the *Controlling Pod Placement onto Nodes (Scheduling)* section in the Red Hat OpenShift Container Platform 4.14 Nodes documentation at https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/nodes/index#controlling-pod-placement-onto-nodes-scheduling

For more information about pod disruption budgets, refer to the *Understanding How to Use Pod Disruption Budgets to Specify the Number of Pods That Must Be Up* section in the Red Hat OpenShift Container Platform 4.14 Nodes documentation at https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/nodes/index#nodes-pods-pod-disruption-about_nodes-pods-configuring

For more information about pod disruption budgets, refer to the *Pod Disruption Budgets* section in the Red Hat OpenShift Container Platform 4.14 Post-installation Configuration documentation at https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/postinstallation_configuration/index#post-install-pod-disruption-budgets

Kubernetes - Assigning Pods to Nodes

<https://kubernetes.io/docs/tasks/configure-pod-container/assign-pods-nodes-using-node-affinity>

Kubernetes - Specifying a Disruption Budget for Your Application

<https://kubernetes.io/docs/tasks/run-application/configure-pdb>

Kubernetes Well-known Labels, Annotations and Taints

<https://kubernetes.io/docs/reference/labels-annotations-taints>

► Guided Exercise

High Availability with Affinity Rules and Pod Disruption Budgets

Configure a workload to spread its pods between nodes from different failure domains and set minimum availability requirements. Then, drain a node to simulate a cluster update, and prove that the application keeps minimum capacity and availability.

Outcomes

- Add pod anti-affinity settings to a deployment resource manifest to spread the pods across the failure domains.
- Create a pod disruption budget to set a minimum availability constraint that is fulfilled when the cluster has a voluntary disruption.
- Drain a compute node to simulate a voluntary disruption.

Before You Begin

As the student user on the workstation machine, use the `lab` command to prepare your environment for this exercise.

```
[student@workstation ~]$ lab start scheduling-pdb
```

Instructions

The company has a local OpenShift cluster that is distributed between two racks. Each rack is connected to a different power source. The nodes are distributed by rack according to the following table:

Rack	Control plane nodes	Compute nodes
rack-a	master01, master02	worker03
rack-b	master03	worker01, worker02

The administrator needs to drain the compute nodes for maintenance. The administrator added the `rack` label to indicate the location and failure domain of each node. This label is intended as a custom topology key, so that the scheduler can spread the pods evenly across the compute nodes in different racks.

The application runs six pods and requires five of them to be available if a voluntary cluster disruption occurs, to achieve the intended response time.

The developer modifies the deployment resource to add the pod anti-affinity settings that use the custom topology, and also creates a pod disruption budget to indicate the minimum availability constraint of the application.

- 1. Verify that the nodes are labeled according to their rack location.

Chapter 4 | Pod Scheduling

- 1.1. Log in to the cluster as the `admin` user.

```
[student@workstation ~]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
Login successful.

...output omitted...
```

- 1.2. Verify that all the nodes have the rack labels according to the previous table.

NAME	STATUS	ROLES	AGE	VERSION	RACK
master01	Ready	control-plane,master	8d	v1.27.6+...	rack-a
master02	Ready	control-plane,master	8d	v1.27.6+...	rack-a
master03	Ready	control-plane,master	8d	v1.27.6+...	rack-b
worker01	Ready	worker	7d	v1.27.6+...	rack-b
worker02	Ready	worker	7d	v1.27.6+...	rack-b
worker03	Ready	worker	7d	v1.27.6+...	rack-a

- ➊➋ The worker01 and worker02 compute nodes are placed in the rack-b rack.
- ➌ The worker03 compute node is placed in the rack-a rack.

- ▶ 2. Create the deployment without pod affinity or a pod disruption budget.

- 2.1. Log in as the `developer` user and verify that you are using the `scheduling-pdb` project.

```
[student@workstation ~]$ oc login -u developer -p developer
Login successful.

You have one project on this server: "scheduling-pdb"

Using project "scheduling-pdb".
```

- 2.2. Change to the `~/D0380/labs/scheduling-pdb` directory.

```
[student@workstation ~]$ cd ~/D0380/labs/scheduling-pdb
```

- 2.3. Create the deployment by using the YAML resource manifest.

```
[student@workstation scheduling-pdb]$ oc apply -f deployment.yaml
deployment.apps/nginx created
```

- 2.4. Open a new terminal window, and then execute the following command to see the status of the pod disruption budget, deployment, and pods.

Wait until all pods are running, and verify that all the pods from the `nginx` deployment are marked as ready and available.

This process might take a few minutes.

```
[student@workstation scheduling-pdb]$ watch oc get pdb,deployments,pods -o wide
Every 2.0s: oc get pdb,deployments,pods ... workstation: Wed Jan  3 15:59:55 2024

NAME          READY   UP-TO-DATE   AVAILABLE   AGE   ...
deployment.apps/nginx   6/6      6           6          60s   ...

NAME          READY   STATUS    RESTARTS   AGE   IP     NODE   ...
pod/nginx-5676948d76-75l7z  1/1    Running   0          60s   ...   worker01 ...
pod/nginx-5676948d76-pkdr7  1/1    Running   0          60s   ...   worker01 ...
pod/nginx-5676948d76-gcst5  1/1    Running   0          60s   ...   worker01 ...
pod/nginx-5676948d76-njzv2  1/1    Running   0          60s   ...   worker02 ...
pod/nginx-5676948d76-94zmk  1/1    Running   0          60s   ...   worker03 ...
pod/nginx-5676948d76-mcdbj  1/1    Running   0          60s   ...   worker03 ...
```



Note

Keep this terminal window open to view the status of the resources from this exercise.

- 2.5. Return to the first terminal window and count the pods that are running on each compute node by using the `count-pods.sh` shell script.

The replica pods are distributed across the cluster nodes but are not distributed evenly between the `rack-a` and `rack-b` failure domains.

```
[student@workstation scheduling-pdb]$ ./count-pods.sh
NODE      PODS
worker01   3  ①
worker02   1  ②
worker03   2  ③
```

① ② The `worker01` and `worker02` compute nodes are placed in the `rack-b` rack.

③ The `worker03` compute node is placed in the `rack-a` rack.



Note

Although the exact number of pods that are running on each node might be different, the total replica count is six pods.

- 3. Simulate a voluntary disruption where the cluster administrator takes the `worker01` node offline for maintenance.



Important

The selected node for draining must have at least two pods running.

- 3.1. Log in as the `admin` user.

Chapter 4 | Pod Scheduling

```
[student@workstation scheduling-pdb]$ oc login -u admin -p redhatocp
Login successful.

...output omitted...
```

- 3.2. Drain the `worker01` node to simulate taking it offline for maintenance.

This command might take a few minutes to complete. Leave it running and continue with the next step. You review the output of this command in a later step.

```
[student@workstation scheduling-pdb]$ oc adm drain node/worker01 \
--ignore-daemonsets --delete-emptydir-data

...output omitted...
```

- 3.3. Switch to the second terminal window to view the eviction of the pods from the drained node. Wait until all pods are running in another node and are marked as ready. This process might take a few minutes.

All the application pods in the drained node are evicted at the same time and the minimum availability constraints are not met. Use the values in the age column to determine which pods were evicted from the drained node and were scheduled in a different node.

This situation happens because no pod disruption budget is associated with the deployment pods, and the deployment resource also does not have an affinity setting that uses the `rack` label as a custom topology key.

```
Every 2.0s: oc get pdb,deployments,pods ... workstation: Wed Jan 3 16:02:33 2024

NAME          READY  UP-TO-DATE  AVAILABLE  AGE   ...
deployment.apps/nginx  3/6    6           3          14m  ...
... ①

NAME          READY  STATUS     ...  AGE  IP      NODE   ...
pod/nginx-5676948d76-njzv2  1/1   Running   ...  3m   ...  worker02 ...
pod/nginx-5676948d76-94zmk  1/1   Running   ...  3m   ...  worker03 ...
pod/nginx-5676948d76-mcdbj  1/1   Running   ...  3m   ...  worker03 ...
pod/nginx-5676948d76-hfbv9  0/1   Init:0/1  ...  1s   ...  worker02 ...
pod/nginx-5676948d76-zxxlg  0/1   Init:0/1  ...  1s   ...  worker02 ...
pod/nginx-5676948d76-dh6dh  0/1   Init:0/1  ...  1s   ...  worker03 ...
... ②
```

- ① Only three replica pods are available.
- ② Three replica pods are evicted from the drained compute node.

**Note**

Although the exact number of pods that are running on each node might be different, the total replica count is six pods.

- 3.4. Return to the first terminal window and inspect the output of the `oc adm drain` command.

Chapter 4 | Pod Scheduling

Observe the pod eviction messages of the nginx pods. All the application pods in the drained node are evicted at the same time and the minimum availability constraint is not met.

```
[student@workstation scheduling-pdb]$ oc adm drain node/worker01 \
--ignore-daemonsets --delete-emptydir-data
node/worker01 cordoned
Warning: ignoring DaemonSet-managed Pods: ...output omitted...
...output omitted...
I1221 21:29:52.102938 111157 request.go:696] ...output omitted...
...output omitted...
evicting pod scheduling-pdb/nginx-5676948d76-pkdr7 ①
evicting pod scheduling-pdb/nginx-5676948d76-75l7z
evicting pod scheduling-pdb/nginx-5676948d76-gcst5
...output omitted...
pod/nginx-5676948d76-gcst5 evicted ②
pod/nginx-5676948d76-75l7z evicted
pod/nginx-5676948d76-pkdr7 evicted
...output omitted...
node/worker01 drained
```

- ① All the pods are marked for eviction when the node is drained.
- ② All the pods are evicted from the node at the same time and the application availability constraint is not met.

**Note**

You can safely ignore the warnings about managed pods and client-side throttling.

- 3.5. Get the state of the nodes to verify that the drained node is marked as not schedulable.

```
[student@workstation scheduling-pdb]$ oc get nodes
NAME      STATUS            ROLES
master01  Ready
master02  Ready
master03  Ready
worker01 Ready, SchedulingDisabled worker ①
worker02  Ready
worker03  Ready
```

- ① The compute node is drained for maintenance.
- 3.6. Count the pods that are running on each compute node. The scheduler placed replacement pods for the evicted pods in the worker02 and worker03 compute nodes.

Chapter 4 | Pod Scheduling

```
[student@workstation scheduling-pdb]$ ./count-pods.sh
NODE          PODS
worker01      0  ①
worker02        3
worker03        3
```

- ① No pods are on this node, because it was just drained.

3.7. Delete the nginx deployment.

```
[student@workstation scheduling-pdb]$ oc delete deployment/nginx
deployment.apps "nginx" deleted
```

3.8. Uncordon the worker01 node that you drained previously to remove the SchedulingDisabled status.

```
[student@workstation scheduling-pdb]$ oc adm uncordon node/worker01
node/worker01 uncordoned
```

3.9. List the cluster nodes and verify that all the compute nodes are marked as ready.

```
[student@workstation ~]$ oc get nodes -L rack
NAME     STATUS   ROLES            AGE    VERSION   RACK
master01 Ready    control-plane,master 8d    v1.27.6+... rack-a
master02 Ready    control-plane,master 8d    v1.27.6+... rack-a
master03 Ready    control-plane,master 8d    v1.27.6+... rack-b
worker01 Ready    worker           7d    v1.27.6+... rack-b
worker02 Ready    worker           7d    v1.27.6+... rack-b
worker03 Ready    worker           7d    v1.27.6+... rack-a
```

► 4. Create the nginx deployment with pod anti-affinity to spread the pods evenly across the compute nodes.

4.1. Log in as the developer user.

```
[student@workstation scheduling-pdb]$ oc login -u developer -p developer
Login successful.

...output omitted...
```

4.2. Edit the deployment-affinity.yaml file and set the affinity properties according to the following specification. Then, save and close the file.

```
...output omitted...
spec:
  ...output omitted...
  template:
    ...output omitted...
    spec:
      ...output omitted...
```

Chapter 4 | Pod Scheduling

```

containers:
  ...output omitted...
affinity:
  podAntiAffinity:
    preferredDuringSchedulingIgnoredDuringExecution: ①
      - weight: 100
    podAffinityTerm:
      topologyKey: rack ②
      labelSelector: ③
        matchExpressions:
          - key: app
            operator: In
            values:
              - nginx

```

- ① The weighted pod affinity term is evaluated only during pod scheduling, on a best-effort basis.
- ② The node label that indicates the failure domain for the nodes.
- ③ The label to select the pods that this affinity setting affects.

**Note**

The ~/DO380/solutions/scheduling-pdb/deployment-affinity.yaml file contains the correct configuration, and you can use it for comparison.

4.3. Create the application deployment resource by using the YAML manifest.

```
[student@workstation scheduling-pdb]$ oc apply -f deployment-affinity.yaml
deployment.apps/nginx created
```

4.4. Switch to the second terminal window. Wait until all pods are running and verify that all the pods from the nginx deployment are marked as ready and available.

This process might take a few minutes.

```
Every 2.0s: oc get pdb,deployments,pods ... workstation: Wed Jan  3 16:24:11 2024

NAME           READY   UP-TO-DATE   AVAILABLE   AGE   ...
deployment.apps/nginx   6/6       6           6           120s  ...

NAME           READY   STATUS    RESTARTS   AGE   IP     NODE   ...
pod/nginx-d5b9c7498-5hbkw  1/1    Running   0          99s   ...   worker01 ...
pod/nginx-d5b9c7498-nkx9j  1/1    Running   0          99s   ...   worker01 ...
pod/nginx-d5b9c7498-g6ztb  1/1    Running   0          99s   ...   worker02 ...
pod/nginx-d5b9c7498-bk7g6  1/1    Running   0          99s   ...   worker03 ...
pod/nginx-d5b9c7498-djn8p  1/1    Running   0          99s   ...   worker03 ...
pod/nginx-d5b9c7498-pz8g5  1/1    Running   0          99s   ...   worker03 ...
```

4.5. Return to the first terminal window and count the pods that are running on each compute node.

The pods are evenly distributed across the racks, because of the pod anti-affinity settings.

Chapter 4 | Pod Scheduling

- Three pods are running in the **rack-b** rack nodes.
- Three pods are running in the **rack-a** rack nodes.

```
[student@workstation scheduling-pdb]$ ./count-pods.sh
NODE          PODS
worker01      2  ①
worker02      1  ②
worker03      3  ③
```

① ② The worker01 and worker02 compute nodes are in the **rack-b** rack.

③ The worker03 compute node is in the **rack-a** rack.

► **5.** Create the pod disruption budget with the intended constraints.

- 5.1. Edit the `pod-disruption-budget.yaml` file and set the minimum available percentage and the label selector according to the following specification. Then, save and close the file.

```
apiVersion: policy/v1
kind: PodDisruptionBudget
metadata:
  name: nginx
  labels:
    app: nginx
spec:
  minAvailable: 80%
  selector:
    matchLabels:
      app: nginx
```

**Note**

The `~/DO380/solutions/scheduling-pdb/pod-disruption-budget.yaml` file contains the correct configuration, and you can use it for comparison.

5.2. Create the pod disruption budget by using the YAML manifest.

```
[student@workstation scheduling-pdb]$ oc apply -f pod-disruption-budget.yaml
poddisruptionbudget.policy/nginx created
```

- 5.3. Verify that the `nginx` pod disruption budget was created, and that it has the intended minimum available attribute.

```
[student@workstation scheduling-pdb]$ oc describe pdb nginx
Name:          nginx
Namespace:     scheduling-pdb
Min available: 80%
Selector:      app=nginx
Status:
  Allowed disruptions: 1  ①
```

Chapter 4 | Pod Scheduling

Current:	6
Desired:	5
Total:	6
Events:	<none>

- ❶ Only one pod can be evicted at a time from a drained node.

- ▶ 6. Drain a compute node to simulate a voluntary disruption.

**Important**

The selected node for draining must have at least two pods running.

- 6.1. Log in again as the `admin` user.

```
[student@workstation scheduling-pdb]$ oc login -u admin -p redhatocp
Login successful.

...output omitted...
```

- 6.2. Drain the `worker03` node to simulate taking it offline for maintenance.

This command might take a few minutes to complete. Leave it running and continue with the next step. You review the output of this command in a later step.

```
[student@workstation scheduling-pdb]$ oc adm drain node/worker03 \
--ignore-daemonsets --delete-emptydir-data

...output omitted...
```

- 6.3. Switch to the second terminal window to view the eviction of the pods from the drained node. Wait until all pods are running in another node and are marked as ready. This process might take a few minutes.

One pod is evicted at a time from the drained node and the availability constraints are met. Use the values in the age column to determine which pods were evicted from the drained node and were scheduled in a different node.

```
Every 2.0s: oc get pdb,deployments,pods ... workstation: Wed Jan  3 16:28:12 2024

NAME                                     MIN AVAIL... MAX UNAVAIL... ALLOWED DISRUPTIONS
poddisruptionbudget.policy/nginx      80%          N/A           1

NAME          READY  UP-TO-DATE   AVAILABLE   AGE   ...
deployment.apps/nginx      5/6       6            5          30m   ...
                                                 ①

NAME          READY  STATUS    RESTARTS  AGE   IP     NODE   ...
pod/nginx-d5b9c7498-5hbkw  1/1    Running    0         5m   ...  worker01 ...
pod/nginx-d5b9c7498-nkx9j  1/1    Running    0         5m   ...  worker01 ...
pod/nginx-d5b9c7498-g6ztb  1/1    Running    0         5m   ...  worker02 ...
pod/nginx-d5b9c7498-pz8g5  1/1    Running    0         5m   ...  worker03 ...
pod/nginx-d5b9c7498-q6rcf  1/1    Running    0         50s  ...  worker01 ...
                                                 ②
                                                 ③
```

Chapter 4 | Pod Scheduling

```
pod/nginx-d5b9c7498-pxx86 0/1    Init:0/1 0          10s ... worker02 ... ④
```

```
^C
```

- ① The pod eviction follows the pod disruption budget.
- ② The pods in the drained node continue to run until the scheduler evicts them.
- ③ The replacement pods are scheduled in another compute node.
- ④ Only one pod is evicted at a time from the drained node.

Press **Ctrl+C** and close the second terminal window when done.

- 6.4. Return to the first terminal window and inspect the output of the **oc adm drain** command.

From the pod eviction messages of the **nginx** pods, observe that one pod is evicted at a time from the drained node and the availability constraints are met.

The pod eviction is blocked until the PDB availability constraints are met. The pod eviction operation is retried after five seconds.

```
[student@workstation scheduling-pdb]$ oc adm drain node/worker03 \
--ignore-daemonsets --delete-emptydir-data
node/worker03 cordoned
Warning: ignoring DaemonSet-managed Pods: ...output omitted...
...output omitted...
I0103 16:27:16.659505 29741 request.go:696] Waited for ... due to client-side
throttling, not priority and fairness, request: ...output omitted...
...output omitted...
evicting pod scheduling-pdb/nginx-d5b9c7498-bk7g6
evicting pod scheduling-pdb/nginx-d5b9c7498-djn8p ①
evicting pod scheduling-pdb/nginx-d5b9c7498-pz8g5
...output omitted...
pod/nginx-d5b9c7498-bk7g6 evicted
...output omitted...
error when evicting pods/"nginx-d5b9c7498-djn8p" -n "scheduling-pdb" (will retry
after 5s): Cannot evict pod as it would violate the pod's disruption budget. ②
error when evicting pods/"nginx-d5b9c7498-pz8g5" -n "scheduling-pdb" (will retry
after 5s): Cannot evict pod as it would violate the pod's disruption budget.
...output omitted...
pod/nginx-d5b9c7498-djn8p evicted ③
evicting pod scheduling-pdb/nginx-d5b9c7498-pz8g5
error when evicting pods/"nginx-d5b9c7498-pz8g5" -n "scheduling-pdb" (will retry
after 5s): Cannot evict pod as it would violate the pod's disruption budget.
...output omitted...
pod/nginx-d5b9c7498-pz8g5 evicted
node/worker03 drained
```

- ① The pod is marked for eviction.
- ② The pod eviction is blocked until the PDB availability constraints are met.
- ③ The pod is finally evicted from the drained node.

**Note**

You can safely ignore the warnings about managed pods and client-side throttling.

- 6.5. List the cluster nodes and verify that the `worker01` node status is `SchedulingDisabled`.

```
[student@workstation scheduling-pdb]$ oc get nodes
NAME      STATUS        ROLES          AGE   VERSION
master01   Ready        control-plane, master   28d   v1.27.6+...
master02   Ready        control-plane, master   28d   v1.27.6+...
master03   Ready        control-plane, master   28d   v1.27.6+...
worker01   Ready        worker           8d    v1.27.6+...
worker02   Ready        worker           8d    v1.27.6+...
worker03   Ready, SchedulingDisabled   worker           8d    v1.27.6+... 1
```

- 1 The compute node is marked as not schedulable.

- 6.6. Count the pods that are running on each compute node.

```
[student@workstation scheduling-pdb]$ ./count-pods.sh
NODE      PODS
worker01  3 1
worker02  3 2
worker03  0 3
```

- 1 2 The pods are evenly distributed on the remaining nodes.

- 3 No pods are on this node, because it was just drained.

- 6.7. Switch to the student HOME directory.

```
[student@workstation scheduling-pdb]$ cd
[student@workstation ~]$
```

► 7. **Optional:** Clean up the resources that were used in this exercise.

- 7.1. Delete the `scheduling-pdb` project.

```
[student@workstation ~]$ oc delete project scheduling-pdb
project.project.openshift.io "scheduling-pdb" deleted
```

- 7.2. Uncordon all the compute nodes.

```
[student@workstation ~]$ oc adm uncordon -l node-role.kubernetes.io/worker
...output omitted...
```

- 7.3. Remove the `rack` label from all nodes.

```
[student@workstation ~]$ oc label node --all rack-
...output omitted...
```

Finish

On the **workstation** machine, use the **lab** command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish scheduling-pdb
```

▶ Lab

Pod Scheduling

Configure a workload to ensure that its pods run on dedicated nodes, and configure a node to prevent a pod's workloads from running on them.

Configure applications for resilience against node failures.

Outcomes

- Create deployments with node selectors to specify node pools for a workload.
- Configure pods with tolerations for nodes that use taints.
- Use pod disruption budgets to minimize downtime during cluster updates and other node maintenance processes.
- Use pod affinity rules to ensure that pods from the different workloads run on the same nodes.

Before You Begin

As the student user on the workstation machine, use the `lab start scheduling-review` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start scheduling-review
```

Instructions

Your company has a cluster with three compute nodes, as in the following diagram:

Node name	Labels	Taints
worker01	disk=ssd	application=ml
worker02	disk=ssd	(none)
worker03	disk=nvme	(none)

Servers might have different hardware. The `worker01` and `worker02` nodes have SATA SSD disks and include the `disk=ssd` label. The `worker03` node has NVMe disks and includes the `disk=nvme` label. The NVMe disks are faster than the SATA SSD disks. Because the cluster has limited resources, the `worker01` node includes the `application=ml:NoSchedule` taint to reserve some workload capacity for a machine learning application that is critical for your company.

Your company requires you to create a machine learning deployment that can run on the node with the taint. The name for the deployment must be `review-toleration`. This deployment uses the `registry.example.com:8443/ubi9/ubi:9.0.0-1468` container with eight replicas. Moreover, this deployment requires high availability in the cluster, with no more than 25% of the pods being unavailable. This requirement means that at least six pods must be always running on the cluster. Thus, if you try to drain a node that contains more than 25% of the pods,

then OpenShift cannot drain it until the pods that exceed the 25% threshold are placed onto other nodes. The pod disruption budget name must be `review-pdb`.

Your company also requires you to create a deployment that needs a fast disk. Thus, you must create the deployment with a node selector for the node with the `disk=nvme` label. The deployment uses the `registry.ocp4.example.com:8443/redhattraining/hello-world-nginx:v1.0` container with four replicas and with the `review-ns` name.

Finally, your company requires you to create a deployment that must communicate with the `review-ns` deployment frequently. Thus, you must create this deployment with a required pod affinity rule for the pods in the `review-ns` deployment. Use the `kubernetes.io/hostname` label as the `topologyKey` parameter. The deployment uses the `registry.ocp4.example.com:8443/rhel9/mysql-80:1-237` container with four replicas and with the `review-affinity` name.

Create all the deployments as the `developer` user with `developer` as the password in the `scheduling-review` project. For the deployments, you can use the incomplete CR YAML files in the `~/DO380/labs/scheduling-review` directory. If you need administrator permissions, then use the `admin` user with `redhatocp` as the password. You might use the `count-pods.sh` script, which counts the pods of a deployment in each node.

1. As the `admin` user, verify the labels and taints of your cluster nodes.
2. Create the `review-toleration` deployment with a toleration for the `worker01` node taint.
3. Create the `review-pdb` pod disruption budget for the `review-toleration` deployment. This pod disruption budget ensures that no more than 25% of the pods can become unavailable.
4. Verify that the pod disruption budget works as expected, and that you get error messages when you try to drain a node that contains more than 25% of the `review-toleration` deployment pods. After verifying that the pod disruption budget works as expected, uncordon all nodes to mark them as schedulable.
5. Create the `review-ns` deployment with a node selector for the node with the `disk=nvme` label, and verify that the pods are scheduled in the node with the `disk=nvme` label.
6. Create the `review-affinity` deployment with a required pod affinity rule for the pods in the `review-ns` deployment. Use the `kubernetes.io/hostname` node label as the topology key.

Evaluation

As the `student` user on the `workstation` machine, use the `lab` command to grade your work. Correct any reported failures and rerun the command until successful.

```
[student@workstation ~]$ lab grade scheduling-review
```

Finish

As the `student` user on the `workstation` machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish scheduling-review
```

► Solution

Pod Scheduling

Configure a workload to ensure that its pods run on dedicated nodes, and configure a node to prevent a pod's workloads from running on them.

Configure applications for resilience against node failures.

Outcomes

- Create deployments with node selectors to specify node pools for a workload.
- Configure pods with tolerations for nodes that use taints.
- Use pod disruption budgets to minimize downtime during cluster updates and other node maintenance processes.
- Use pod affinity rules to ensure that pods from the different workloads run on the same nodes.

Before You Begin

As the student user on the `workstation` machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start scheduling-review
```

Instructions

Your company has a cluster with three compute nodes, as in the following diagram:

Node name	Labels	Taints
worker01	disk=ssd	application=ml
worker02	disk=ssd	(none)
worker03	disk=nvme	(none)

Servers might have different hardware. The `worker01` and `worker02` nodes have SATA SSD disks and include the `disk=ssd` label. The `worker03` node has NVMe disks and includes the `disk=nvme` label. The NVMe disks are faster than the SATA SSD disks. Because the cluster has limited resources, the `worker01` node includes the `application=ml:NoSchedule` taint to reserve some workload capacity for a machine learning application that is critical for your company.

Your company requires you to create a machine learning deployment that can run on the node with the taint. The name for the deployment must be `review-toleration`. This deployment uses the `registry.ocp4.example.com:8443/ubi9/ubi:9.0.0-1468` container with eight replicas. Moreover, this deployment requires high availability in the cluster, with no more than 25% of the pods being unavailable. This requirement means that at least six pods must be always

Chapter 4 | Pod Scheduling

running on the cluster. Thus, if you try to drain a node that contains more than 25% of the pods, then OpenShift cannot drain it until the pods that exceed the 25% threshold are placed onto other nodes. The pod disruption budget name must be `review-pdb`.

Your company also requires you to create a deployment that needs a fast disk. Thus, you must create the deployment with a node selector for the node with the `disk=nvme` label. The deployment uses the `registry.ocp4.example.com:8443/redhattraining/hello-world-nginx:v1.0` container with four replicas and with the `review-ns` name.

Finally, your company requires you to create a deployment that must communicate with the `review-ns` deployment frequently. Thus, you must create this deployment with a required pod affinity rule for the pods in the `review-ns` deployment. Use the `kubernetes.io/hostname` label as the `topologyKey` parameter. The deployment uses the `registry.ocp4.example.com:8443/rhel9/mysql-80:1-237` container with four replicas and with the `review-affinity` name.

Create all the deployments as the `developer` user with `developer` as the password in the `scheduling-review` project. For the deployments, you can use the incomplete CR YAML files in the `~/DO380/labs/scheduling-review` directory. If you need administrator permissions, then use the `admin` user with `redhatocp` as the password. You might use the `count-pods.sh` script, which counts the pods of a deployment in each node.

- As the `admin` user, verify the labels and taints of your cluster nodes.

- Connect to the OpenShift cluster as the `admin` user with `redhatocp` as the password.

```
[student@workstation ~]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
Login successful.
...output omitted...
```

- List the labels for the compute nodes. Verify that the `worker01` and `worker02` nodes include the `disk=ssd` label, and that the `worker03` node includes the `disk=nvme` label.

NAME	STATUS	ROLES	AGE	VERSION	DISK
master01	Ready	control-plane,master	140d	v1.25.7+eab9cc9	
master02	Ready	control-plane,master	140d	v1.25.7+eab9cc9	
master03	Ready	control-plane,master	140d	v1.25.7+eab9cc9	
worker01	Ready	worker	37d	v1.25.7+eab9cc9	ssd
worker02	Ready	worker	37d	v1.25.7+eab9cc9	ssd
worker03	Ready	worker	37d	v1.25.7+eab9cc9	nvme

- Verify the taint for the `worker01` node.

```
[student@workstation ~]$ oc describe node worker01 | grep Taints
Taints:           application=ml:NoSchedule
```

- Create the `review-toleration` deployment with a toleration for the `worker01` node taint.

- Connect to the OpenShift cluster as the `developer` user with `developer` as the password, and verify that OpenShift uses the `scheduling-review` project.

Chapter 4 | Pod Scheduling

```
[student@workstation ~]$ oc login -u developer -p developer
Login successful.
...output omitted...
Using project "scheduling-review".
```

- 2.2. Change to the ~/D0380/labs/scheduling-review directory.

```
[student@workstation ~]$ cd ~/D0380/labs/scheduling-review
```

- 2.3. Edit the ~/D0380/labs/scheduling-review/review-toleration.yml file for the review-toleration deployment, and add the toleration for the application=ml:NoSchedule taint.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: review-toleration
spec:
  ...output omitted...
  spec:
    ...output omitted...
    tolerations:
      - key: "application"
        value: "ml"
        operator: "Equal"
        effect: "NoSchedule"
```

- 2.4. Create the review-toleration deployment with the toleration.

```
[student@workstation scheduling-review]$ oc create -f review-toleration.yml
deployment.apps/review-toleration created
```

- 2.5. Verify that the deployment is correctly created. Wait until all the pods are marked as ready and available. You might have to repeat this command many times.

```
[student@workstation scheduling-review]$ oc get deployment
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
review-toleration   8/8     8           8           53s
```

- 2.6. Verify that the OpenShift scheduler distributes the pods between the available compute nodes. OpenShift places some pods in the worker01 node, because the deployment includes a toleration for the taint. The pod names and the nodes that run on your pods might differ on your system.

```
[student@workstation scheduling-review]$ oc get pods -o wide
NAME                  READY   STATUS    ...   NODE   ...
review-toleration-76f8c74d7-87jjr  1/1    Running   ...   worker02   ...
review-toleration-76f8c74d7-gtwhx  1/1    Running   ...   worker03   ...
review-toleration-76f8c74d7-hfssj  1/1    Running   ...   worker01 ...
review-toleration-76f8c74d7-lmdcn  1/1    Running   ...   worker03   ...
```

Chapter 4 | Pod Scheduling

```
review-toleration-76f8c74d7-ph724 1/1   Running ... worker02 ...
review-toleration-76f8c74d7-tntwz 1/1   Running ... worker01 ...
review-toleration-76f8c74d7-w4tfj  1/1   Running ... worker03 ...
review-toleration-76f8c74d7-zr7d5  1/1   Running ... worker02 ...
```

- 2.7. Use the `count-pods.sh` script to count the pods for the `review-toleration` deployment in each node. The output might differ on your system, because the OpenShift scheduler distributes the pods across the three compute nodes depending on their workload.

```
[student@workstation scheduling-review]$ ./count-pods.sh review-toleration
NODE      PODS
worker01   2
worker02   3
worker03   3
```

3. Create the `review-pdb` pod disruption budget for the `review-toleration` deployment. This pod disruption budget ensures that no more than 25% of the pods can become unavailable.
- 3.1. Edit the `~/DO380/labs/scheduling-review/review-pdb.yml` file for the `review-pdb` pod disruption budget.

```
apiVersion: policy/v1
kind: PodDisruptionBudget
metadata:
  name: review-pdb
spec:
  maxUnavailable: 25%
  selector:
    matchLabels:
      app: review-toleration
```

- 3.2. Create the `review-pdb` pod disruption budget.

```
[student@workstation scheduling-review]$ oc create -f review-pdb.yml
poddisruptionbudget.policy/review-pdb created
```

- 3.3. Verify that the pod disruption budget is correctly created, with a maximum of two unavailable pods.

```
[student@workstation scheduling-review]$ oc describe pdb review-pdb
Name:           review-pdb
Namespace:      scheduling-review
Max unavailable: 25%
Selector:       app=review-toleration
Status:
  Allowed disruptions: 2
  Current:            8
  Desired:            6
  Total:              8
Events:          <none>
```

Chapter 4 | Pod Scheduling

4. Verify that the pod disruption budget works as expected, and that you get error messages when you try to drain a node that contains more than 25% of the **review-toleration** deployment pods. After verifying that the pod disruption budget works as expected, uncordon all nodes to mark them as schedulable.

- 4.1. Connect to the OpenShift cluster as the **admin** user with **redhatocp** as the password.

```
[student@workstation ~]$ oc login -u admin -p redhatocp
...output omitted...
```

- 4.2. To test that the pod disruption budget works as expected, first mark the **worker02** and the **worker03** nodes as unschedulable.

```
[student@workstation scheduling-review]$ oc adm cordon worker02 worker03
node/worker02 cordoned
node/worker03 cordoned
```

- 4.3. Perform a rollout restart of the **review-toleration** deployment to ensure that OpenShift creates all the pods in the **worker01** node.

```
[student@workstation scheduling-review]$ oc rollout restart \
deployment/review-toleration
deployment.apps/review-toleration restarted
```

- 4.4. Verify that all the pods are scheduled to execute in the **worker01** node. For any pods in a terminating status, wait until OpenShift removes those pods. You might have to repeat this command many times.

```
[student@workstation scheduling-review]$ oc get pods -o wide
NAME                  READY   STATUS    ...   NODE   ...
review-toleration-76f8c74d7-4494g  1/1    Running   ...   worker01 ...
review-toleration-76f8c74d7-5jw2q  1/1    Running   ...   worker01 ...
review-toleration-76f8c74d7-86mf8  1/1    Running   ...   worker01 ...
review-toleration-76f8c74d7-kp6wp  1/1    Running   ...   worker01 ...
review-toleration-76f8c74d7-ljxpd  1/1    Running   ...   worker01 ...
review-toleration-76f8c74d7-nbjz6  1/1    Running   ...   worker01 ...
review-toleration-76f8c74d7-t78bg  1/1    Running   ...   worker01 ...
review-toleration-76f8c74d7-tl29v  1/1    Running   ...   worker01 ...
```

- 4.5. Use the **count-pods.sh** script to count the pods for the **review-toleration** deployment in each node.

```
[student@workstation scheduling-review]$ ./count-pods.sh review-toleration
NODE      PODS
worker01   8
worker02   0
worker03   0
```

- 4.6. Mark the **worker02** and **worker03** nodes as schedulable.

Chapter 4 | Pod Scheduling

```
[student@workstation scheduling-review]$ oc adm uncordon worker02 worker03
node/worker02 uncordoned
node/worker03 uncordoned
```

- 4.7. Drain all the pods from the `worker01` node. You receive some error messages, because OpenShift can drain the node only after the pods from the `review-toleration` deployment that exceed the 25% threshold are placed onto other nodes. This action might take a few minutes.

```
[student@workstation scheduling-review]$ oc adm drain worker01 \
--ignore-daemonsets --delete-emptydir-data
...output omitted...
evicting pod scheduling-review/review-toleration-76f8c74d7-kp6wp
evicting pod scheduling-review/review-toleration-76f8c74d7-ljxpd
evicting pod scheduling-review/review-toleration-76f8c74d7-5jw2q
error when evicting pods/"review-toleration-76f8c74d7-5jw2q" -n "scheduling-review" (will retry after 5s): Cannot evict pod as it would violate the pod's disruption budget.
error when evicting pods/"review-toleration-76f8c74d7-ljxpd" -n "scheduling-review" (will retry after 5s): Cannot evict pod as it would violate the pod's disruption budget.
evicting pod scheduling-review/review-toleration-76f8c74d7-nbzj6
error when evicting pods/"review-toleration-76f8c74d7-nbzj6" -n "scheduling-review" (will retry after 5s): Cannot evict pod as it would violate the pod's disruption budget.
...output omitted...
node/worker01 drained
```

- 4.8. Mark the `worker01` node as schedulable.

```
[student@workstation scheduling-review]$ oc adm uncordon worker01
node/worker01 uncordoned
```

5. Create the `review-ns` deployment with a node selector for the node with the `disk=nvme` label, and verify that the pods are scheduled in the node with the `disk=nvme` label.

- 5.1. Connect to the OpenShift cluster as the `developer` user with `developer` as the password.

```
[student@workstation scheduling-review]$ oc login -u developer -p developer
...output omitted...
```

- 5.2. Edit the `~/DO380/labs/scheduling-review/review-ns.yaml` file for the `review-ns` deployment, and add a node selector for the `disk=nvme` label.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: review-ns
spec:
...output omitted...
```

Chapter 4 | Pod Scheduling

```
spec:
  ...output omitted...
  nodeSelector:
    disk: nvme
```

- 5.3. Create the `review-ns` deployment.

```
[student@workstation scheduling-review]$ oc create -f review-ns.yml
deployment.apps/review-ns created
```

- 5.4. Verify that the deployment is correctly created. Wait until all the pods are marked as ready and available. You might have to repeat this command many times.

```
[student@workstation scheduling-review]$ oc get deployment
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
review-ns      4/4     4           4           61s
review-toleration 8/8     8           8           46m
```

- 5.5. Verify that the OpenShift scheduler creates all the pods in the `worker03` node, which has the `disk=nvme` label. The pod names might differ on your system.

```
[student@workstation scheduling-review]$ oc get pods -o wide -l app=review-ns
NAME          READY   STATUS    ...   NODE   ...
review-ns-7475cbc8ff-d7dbq 1/1     Running   ...   worker03 ...
review-ns-7475cbc8ff-hpl5b 1/1     Running   ...   worker03 ...
review-ns-7475cbc8ff-p7lhb 1/1     Running   ...   worker03 ...
review-ns-7475cbc8ff-wjxjz 1/1     Running   ...   worker03 ...
```

6. Create the `review-affinity` deployment with a required pod affinity rule for the pods in the `review-ns` deployment. Use the `kubernetes.io/hostname` node label as the topology key.

- 6.1. Edit the `~/D0380/labs/scheduling-review/review-affinity.yml` file for the `review-affinity` deployment, and add a required pod affinity rule for the pods in the `review-ns` deployment. Use the `kubernetes.io/hostname` node label as the topology key.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: review-affinity
spec:
  ...output omitted...
  spec:
    ...output omitted...
    affinity:
      podAffinity:
        requiredDuringSchedulingIgnoredDuringExecution:
        - labelSelector:
            matchExpressions:
            - key: app
              operator: In
```

Chapter 4 | Pod Scheduling

```
values:
- review-ns
topologyKey: kubernetes.io/hostname
```

- 6.2. Create the review-affinity deployment.

```
[student@workstation scheduling-review]$ oc create -f review-affinity.yml
deployment.apps/review-affinity created
```

- 6.3. Verify that the deployment is correctly created. Wait until all the pods are marked as ready and available. You might have to repeat this command many times.

```
[student@workstation scheduling-review]$ oc get deployment
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
review-affinity  4/4      4           4           27s
review-ns       4/4      4           4           17m
review-toleration  8/8      8           8           64m
```

- 6.4. Verify that all the pods for the review-affinity deployment are placed in the worker03 node. The pod names might differ on your system.

```
[student@workstation scheduling-review]$ oc get pods -o wide \
-l app=review-affinity
NAME                           READY   STATUS    ...   NODE   ...
review-affinity-79d6b5bcfd-8qmh8  1/1     Running   ...   worker03 ...
review-affinity-79d6b5bcfd-cpkhz  1/1     Running   ...   worker03 ...
review-affinity-79d6b5bcfd-h8sq6  1/1     Running   ...   worker03 ...
review-affinity-79d6b5bcfd-mcv7k  1/1     Running   ...   worker03 ...
```

- 6.5. Verify that the OpenShift scheduler creates all the pods for the review-affinity deployment in the same compute node as the review-ns deployment. The pod names might differ on your system.

```
[student@workstation scheduling-review]$ oc get pods -o wide -l app=review-ns
NAME                           READY   STATUS    ...   NODE   ...
review-ns-7475cbc8ff-d7dbq    1/1     Running   ...   worker03 ...
review-ns-7475cbc8ff-hpl5b    1/1     Running   ...   worker03 ...
review-ns-7475cbc8ff-p7lhb    1/1     Running   ...   worker03 ...
review-ns-7475cbc8ff-wjxjz    1/1     Running   ...   worker03 ...
```

- 6.6. Change to the student HOME directory.

```
[student@workstation scheduling-review]$ cd
```

Evaluation

As the student user on the workstation machine, use the lab command to grade your work. Correct any reported failures and rerun the command until successful.

```
[student@workstation ~]$ lab grade scheduling-review
```

Finish

As the student user on the workstation machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish scheduling-review
```

Summary

- Pod scheduling is the process of determining in which node to place a pod in the OpenShift cluster.
- The OpenShift built-in scheduler identifies the most suitable node for the pods, and its default behavior meets the needs of most OpenShift users.
- OpenShift advanced scheduling features enable you to control which node is used for pod placement.
- The scheduler profiles control how OpenShift schedules pods on nodes.
- Use pod affinity rules to keep sets of pods close to each other, on the same nodes. Pod affinity rules can be preferred or required.
- Use pod anti-affinity rules to keep sets of pods far away from each other, on different nodes. Pod anti-affinity rules can be preferred or required.
- Use node affinity rules to keep sets of pods running on the same group of nodes. Node affinity rules can be preferred or required.
- Use node selectors to schedule pods to a specific set of nodes. You can define node selectors at the pod, project, and cluster-wide levels.
- Use taints and tolerations to avoid scheduling pods to a specific set of nodes.
- Pod disruption budgets enable you to control the disruption of pods during voluntary disruptions, by specifying the minimum available number of pods simultaneously.

Chapter 5

OpenShift GitOps

Goal

Deploy OpenShift GitOps for managing clusters and applications.

Sections

- GitOps for Kubernetes (and Quiz)
- GitOps for Cluster Administration (and Guided Exercise)
- GitOps for Application Management (and Guided Exercise)

Lab

- OpenShift GitOps

GitOps for Kubernetes

Objectives

- Define the fundamentals of GitOps and its use with Kubernetes clusters and applications.

Introduction to GitOps

Red Hat OpenShift GitOps is an operator that helps implement GitOps practices to manage OpenShift clusters.

Kubernetes can import and export resource definitions as text files. By working with resource definitions in text files, administrators describe their workloads instead of using a sequence of operations to create them. This approach is called *declarative resource management*.

By storing resource definitions in text files, Kubernetes administrators can use any tool or process for text files. For example, by keeping cluster resource definitions in a version control system, administrators can track changes to resources.

GitOps or *infrastructure as code* are terms to describe practices that relate to these concepts.

By enabling cluster users to manage Kubernetes resources by using a Git repository, administrators can further restrict direct access to the cluster. If users can modify Kubernetes resources only through a Git repository, then administrators can use Git features to enforce controls such as reviews and approvals. For example, many organizations use pull request workflows to update repositories.

Additionally, cluster management presents significant challenges when multiple processes can update cluster resources. Cluster administrators can face problems if they cannot clearly visualize the intended state of the cluster. When multiple users and processes update the state of the cluster, the intended state for the cluster can become unclear. The term *drift* describes the difference between the intended and actual states of a cluster.

With GitOps, a single process changes the cluster based on a definition of the intended cluster state. Administrators can then reference a single cluster definition to track what is happening with the cluster. Also, administrators who invest resources in documenting their changes can have reliable information about each change to the cluster.

Additionally, with a single cluster definition, administrators can reproduce the deployments of a cluster on separate clusters. For example, a cluster administrator can create a test cluster based on their production cluster definition. Then, the administrator can test cluster updates and other changes, some of which might be risky or hard to revert in a production environment, with reasonable guarantees that results from the test cluster apply to the production cluster. Production clusters often have different loads and interact with external services, so test clusters cannot reproduce the exact conditions. However, tests can provide the level of confidence that your organization requires for such changes.

GitOps and Kubernetes

Although the Kubernetes resource model is suited to GitOps, some Kubernetes idiosyncrasies require extra care.

Kubernetes resources contain information about both the definition and the status of the resource.

Consider the following excerpt from a Kubernetes deployment:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: "2023-12-14T17:02:18Z"
  generation: 1
  name: hello-node
  namespace: test
  ...output omitted...
spec:
  replicas: 1
  selector:
    matchLabels:
      app: hello-node
  strategy:
    rollingUpdate:
      maxSurge: 25%
      maxUnavailable: 25%
    type: RollingUpdate
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: hello-node
    spec:
      containers:
        - command:
          - /agnhost
          - serve-hostname
          image: registry.k8s.io/e2e-test-images/agnhost:2.43
          name: agnhost
          restartPolicy: Always
          schedulerName: default-scheduler
  status:
    availableReplicas: 1
    conditions:
      - lastTransitionTime: "2023-12-14T17:02:22Z"
        lastUpdateTime: "2023-12-14T17:02:22Z"
        message: Deployment has minimum availability.
        reason: MinimumReplicasAvailable
        status: "True"
        type: Available
      ...output omitted...
    observedGeneration: 1
    readyReplicas: 1
    replicas: 1
    updatedReplicas: 1
```

The `status` key contains information about the number of available replicas, which can change during the lifetime of the deployment. Although the content of the `metadata` and `spec` keys is mostly the definition of the deployment, Kubernetes adds some fields, such as the

Chapter 5 | OpenShift GitOps

creationTimestamp field, which are not necessary to re-create the deployment. Additionally, you might choose not to include in your resource definition some fields in the spec key with default values, such as the strategy key.

The `kubectl` command has subcommands to manipulate resource definitions, such as the `apply`, `edit`, and `patch` subcommands. These commands handle some complexities with updating resources.

Other Kubernetes features, such as mutating admission webhooks, can modify resources. These features can increase the complexity of updating resources.

Finally, although you can usually create resources in any order, with resources ultimately reaching the intended state, in some cases creating resources can bring more complexity. For example, you must create a namespace before creating the resources in the namespace. You can create custom resources only after the installation of the operator deploys the custom resource definitions.

Due to these factors, implementing GitOps involves additional sophistication. Simple GitOps implementations that are based on built-in features, such as the `kubectl apply` command, can present problems when resources increase in complexity.

System administrators can use other tools, such as Terraform and Ansible, to implement GitOps practices. Tools might differ in their approach and scope. Argo CD focuses on synchronizing Kubernetes resources from Git repositories to the cluster.

Introduction to OpenShift GitOps

OpenShift GitOps is an operator that manages Argo CD in OpenShift clusters. Argo CD is a tool that can synchronize Kubernetes resources to definitions that are stored in Git repositories.

Cluster administrators can use OpenShift GitOps to create Argo CD instances with specific access rights. Administrators can create an Argo CD instance that can manage specific namespaces and grant specific users access to the instance. With these instances, teams can be autonomous in deploying applications to specific namespaces.

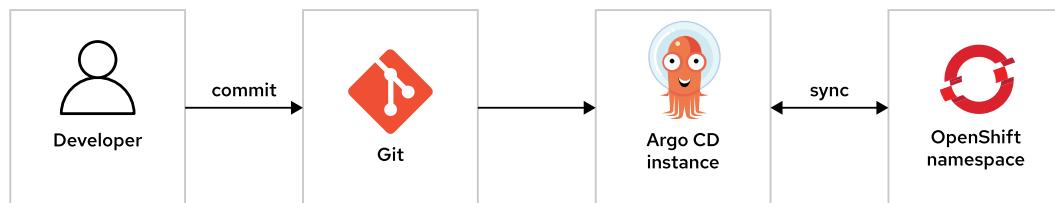


Figure 5.1: OpenShift GitOps architecture for developers

OpenShift GitOps also includes a default administrative Argo CD instance that can manage Kubernetes resources that require administrative access, such as non-namespaced resources or authentication resources. OpenShift administrators can use Argo CD instances with these configurations to manage the cluster itself.

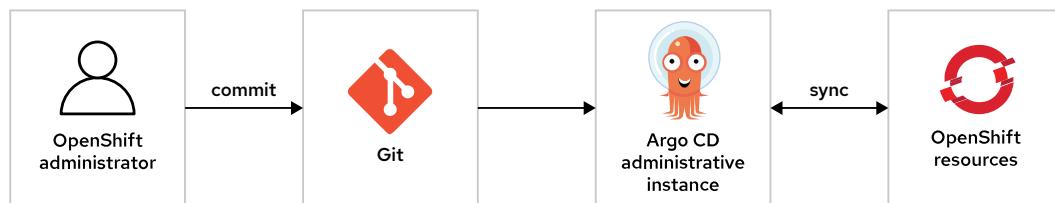


Figure 5.2: OpenShift GitOps architecture for cluster administrators

Chapter 5 | OpenShift GitOps

Argo CD uses the concept of *applications*. An application is an Argo CD resource that references a Git repository with Kubernetes resource definitions. Users can create and manage Argo CD applications from the web console or with the Kubernetes API.

Applications define a synchronization policy. Argo CD can detect and synchronize changes automatically, or users can trigger synchronizations on demand.

The Argo CD web console displays the status of application resources graphically. This representation includes both the resources that are defined in the application, and resources that are related to the application. For example, for an application that contains an operator subscription, the Argo CD web console also shows the install plans, operator workloads, and other related resources.

In this representation, you can view whether the resources are synchronized with the application definition, and inspect their status.

With OpenShift GitOps, cluster administrators can restrict permissions so that only Argo CD instances can update cluster resources, and Git repositories fully describe the intended cluster state. Even if other processes can modify the cluster state, Argo CD can detect and address issues with drift. Additionally, cluster administrators might use Argo CD to manage multiple clusters from the same cluster definition. Besides testing, having a cluster definition can help with other tasks, such as sharing resource definitions between clusters.

Continuous Integration, Continuous Delivery, and Continuous Deployment

The Argo CD name refers to the *continuous delivery* and *continuous deployment* practices. The CI/CD abbreviation refers to these practices and to *continuous integration*.

Continuous integration

Originally, continuous integration referred to integrating changes often to reduce conflicts from developing changes in isolation for long periods. Nowadays, continuous integration often refers to running tests automatically, and to practices such as requiring tests to pass before merging changes.

Continuous delivery

Continuous delivery is the practice of making each change a potential release of the project. With continuous delivery, the team can deploy the project at any time.

Continuous deployment

Continuous deployment is a variant of continuous delivery where every change triggers an automated deployment. Continuous deployment maximizes the release frequency of the project and shortens feedback loops.

CI/CD practices aim to reduce the lead time to deliver a new product or feature. When production issues are found, you can roll back small increments to known stable states.

OpenShift GitOps helps implement continuous delivery and deployment, but does not cover continuous integration. Red Hat provides other products that can help with implementing CI/CD:

Red Hat OpenShift Pipelines

The OpenShift Pipelines product is based on the open source Tekton project, and can run processes when changes occur to a Git repository. OpenShift Pipelines can implement the modern interpretation of CI, by validating changes. OpenShift Pipelines can also update container images with changes to source repositories.

Chapter 5 | OpenShift GitOps

With third-party applications, the application provider often provides updated images that you can deploy with Argo CD. However, with in-house applications, you can create updated images with OpenShift Pipelines.

Jenkins

Jenkins is a generic automation tool, to implement any type of CI/CD process.

For OpenShift automation, Red Hat recommends OpenShift Pipelines and OpenShift GitOps instead of Jenkins. The functions of the latter products require extra effort to implement with Jenkins. The OpenShift documentation covers migrating from Jenkins to OpenShift Pipelines.

By combining these products to implement CI/CD practices, organizations can create highly automated workflows that speed up development and that follow the necessary quality and security guidelines.

GitOps Workflows

The following steps describe an example of such a workflow:

- A developer updates the source code of an application.
- The developer pushes the change to a Git repository branch that contains the source code of the application, and creates a pull request.
- OpenShift Pipelines detects the new pull request and runs automated checks:
 - The application source code is checked for mechanical errors and incorrect formatting. OpenShift Pipelines posts an update to the Git repository, to indicate whether the pull request passed the checks or conversely contains errors.
 - OpenShift Pipelines triggers an Argo CD job that deploys the application with the specified change to a testing environment. Any team members can access the deployment to validate the change. Argo CD can mark the change as failed if the application fails to deploy.
- After the changes are validated, other team members can review the pull request. If validation failed, then a team member can fix the issues and resubmit. With this process, reviewers review only changes that passed automated validation.
- For any issues that reviewers find in the change, they can request changes from the original developer to address issues.
- Approved changes are merged.
- OpenShift Pipelines detects the merged change and builds an updated container image from the source code of the application, and triggers an Argo CD synchronization of the production environment.
- Argo CD deploys the updated image.

To ensure adequate velocity, the organization should ensure that the automated processes can complete quickly, and limit manual approvals to what is strictly required. For example, minor changes might not require a review; typical changes might require a single approver; and significant changes might require approvals from multiple specific stakeholders. These processes aim to automate and speed up the delivery of changes, and introduce only necessary validation. Organizations can choose their tradeoffs between speed and quality, according to their needs.



Warning

Following GitOps procedures can reduce security risks by introducing validation steps. However, version control systems such as Git can also introduce security risks.

For example, if you push sensitive data to a repository, then although your process might prevent pushing the data to production, unintended people might access the sensitive data. Furthermore, permanently removing information from a Git repository can require significant effort. Invalidating the data (for example, by creating an authentication token and disabling the leaked token) can be more efficient than permanently removing the sensitive data from the repository.

You can use elements, such as Git hooks, to improve the security of version control systems. Git hooks can prevent users from pushing sensitive information from their private workstation environment to a more accessible environment.

The following list describes some improvements to GitOps workflows that Argo CD can implement:

- Propagate changes through multiple static environments, such as testing, pre-production, and production environments.
- Limit synchronizations to specific days of the week or periods during the day.
- Add hooks for further processes, such as for sending notifications, pausing monitoring, or triggering backups before an Argo CD synchronization.



References

For more information, refer to the *Understanding OpenShift GitOps* chapter in the Red Hat OpenShift GitOps 1.10 *Understanding OpenShift GitOps* documentation at https://access.redhat.com/documentation/en-us/red_hat_openshift_gitops/1.10/html-single/understanding_openshift_gitops/index#what-is-gitops

For more information about Red Hat OpenShift Pipelines, refer to the *About Red Hat OpenShift Pipelines* chapter in the Red Hat OpenShift Pipelines 1.13 *About OpenShift Pipelines* documentation at https://access.redhat.com/documentation/en-us/red_hat_openshift_pipelines/1.13/html-single/about_openshift_pipelines/index#understanding-openshift-pipelines

► Quiz

GitOps for Kubernetes

Choose the correct answers to the following questions:

► 1. **What is the purpose of GitOps in the context of Kubernetes?**

- a. Deploying changes to production automatically after reviewing and testing them
- b. Storing Kubernetes resource definitions in Git repositories
- c. Improving the tools and diagnostics that engineers use to resolve issues
- d. Using Kubernetes and cloud providers

► 2. **What are two direct benefits of GitOps? (Choose two.)**

- a. Applying techniques such as version control to Kubernetes cluster administration
- b. Automated testing for infrastructure changes
- c. Aligning objectives between the development and operations teams
- d. Improving control over cluster changes, helping to audit changes, and using tools for approvals and reviews
- e. Improving visibility of cluster application behavior

► 3. **How is Red Hat OpenShift GitOps delivered?**

- a. A Helm chart
- b. A Kustomization
- c. An operator
- d. A base OpenShift component

► 4. **In which two ways does OpenShift GitOps help with controlling changes to a cluster? (Choose two.)**

- a. OpenShift GitOps manages permissions of Git repositories, to limit who can push changes.
- b. OpenShift GitOps helps administrators limit access to the Kubernetes API, because changes can be submitted through Git.
- c. With OpenShift GitOps, administrators can create multiple Argo CD instances to improve control over permissions.
- d. OpenShift GitOps validates changes to the cluster, so administrators can implement policies to prevent certain changes.

► **5. Which practice does OpenShift GitOps relate to?**

- a. DevOps
- b. Agile
- c. Continuous delivery and deployment
- d. Public cloud

► Solution

GitOps for Kubernetes

Choose the correct answers to the following questions:

► 1. **What is the purpose of GitOps in the context of Kubernetes?**

- a. Deploying changes to production automatically after reviewing and testing them
- b. Storing Kubernetes resource definitions in Git repositories
- c. Improving the tools and diagnostics that engineers use to resolve issues
- d. Using Kubernetes and cloud providers

► 2. **What are two direct benefits of GitOps? (Choose two.)**

- a. Applying techniques such as version control to Kubernetes cluster administration
- b. Automated testing for infrastructure changes
- c. Aligning objectives between the development and operations teams
- d. Improving control over cluster changes, helping to audit changes, and using tools for approvals and reviews
- e. Improving visibility of cluster application behavior

► 3. **How is Red Hat OpenShift GitOps delivered?**

- a. A Helm chart
- b. A Kustomization
- c. An operator
- d. A base OpenShift component

► 4. **In which two ways does OpenShift GitOps help with controlling changes to a cluster? (Choose two.)**

- a. OpenShift GitOps manages permissions of Git repositories, to limit who can push changes.
- b. OpenShift GitOps helps administrators limit access to the Kubernetes API, because changes can be submitted through Git.
- c. With OpenShift GitOps, administrators can create multiple Argo CD instances to improve control over permissions.
- d. OpenShift GitOps validates changes to the cluster, so administrators can implement policies to prevent certain changes.

► **5. Which practice does OpenShift GitOps relate to?**

- a. DevOps
- b. Agile
- c. Continuous delivery and deployment
- d. Public cloud

GitOps for Cluster Administration

Objectives

- Configure the default Argo CD instance from OpenShift GitOps for cluster administration.

The Red Hat OpenShift GitOps Operator

Red Hat OpenShift GitOps is an operator in OperatorHub, from the Red Hat operator catalog. OperatorHub provides a `latest` channel, with the most recent stable version of the operator. OperatorHub also provides `gitops-version` channels, such as the `gitops-1.10` channel, to install specific versions of the operator.

You can install the operator in a similar way to any other operator, for example by using the installation wizard from the web console. You can use the default operator installation options, although most deployments require customization after the operator installation.

The operator creates its workload in the `openshift-gitops-operator` namespace. This workload includes a deployment that monitors Argo CD custom resources and manages their workloads.

Argo CD Custom Resources

The OpenShift GitOps operator installs the ArgoCD custom resource definition. Cluster administrators can use this custom resource definition to declaratively create multiple Argo CD instances in the cluster. The operator manages the workloads to run the Argo CD instances.

With this capability, administrators can introduce instances with different permissions. Because Argo CD applications can create and modify any type of OpenShift resource, limiting the permissions of Argo CD instances can reduce risks. With the OpenShift GitOps operator, you can create Argo CD instances by creating a custom resource that describes an instance (including RBAC configuration), and the operator manages the Argo CD instance. This process helps administrators apply the principle of least privilege, by creating Argo CD instances with granular permissions.



Warning

Any user that can push resources to an application Git repository can perform operations with the permissions of the Argo CD instance, even without direct access to the Kubernetes API or Argo CD. For example, these operations might delete data or override access control, whether by accident or malicious intent. Although Argo CD has role-based access control for its resources, Argo CD trusts the content of the application Git repositories. Follow the same considerations for access control for the Git repositories as for granting cluster permissions.

You can also set repositories as private, and configure Argo CD to use credentials to access repositories.

OpenShift GitOps includes further resources. For example, Argo CD uses the Application custom resource to manage applications. When you create an application with the Argo CD web

interface, Argo CD creates an Application custom resource. This course does not cover working with those resources.

The Default Argo CD Instance

By default, the OpenShift GitOps operator creates a default Argo CD instance with the `openshift-gitops` name for both the instance and the namespace. This default instance has a set of permissions to support cluster administration. Red Hat recommends using this instance only for administering the cluster or for creating other instances for regular application deployment with minimal privileges.

The default instance has a limited configuration that might require further configuration to meet organizational requirements.

With the `openshift-gitops-server` route in the `openshift-gitops` namespace, the Argo CD web console is accessible outside the OpenShift cluster. By default, the web console is at the `openshift-gitops-server.openshift-gitops.apps.cluster_domain` URL.

Setting the Default Argo CD Instance as Unmanaged

The operator manages the default Argo CD instance. If you modify the instance, then the operator can revert your changes.

```
apiVersion: argoproj.io/v1beta1
kind: ArgoCD
metadata:
  creationTimestamp: "2024-01-02T11:27:35Z"
  finalizers:
  - argoproj.io/finalizer
  generation: 1
  name: openshift-gitops
  namespace: openshift-gitops
  ownerReferences:
  - apiVersion: pipelines.openshift.io/v1alpha1
    blockOwnerDeletion: true
    controller: true
    kind: GitopsService
    name: cluster
    uid: a7a0a073-0328-449e-8ffc-038854feb1c9
  resourceVersion: "262080"
  uid: 1d530e04-8068-4f6f-a066-c3ae36ba34f4
spec:
  applicationSet:
  ...output omitted...
```

For example, if you remove the `ownerReferences` key from the instance metadata, then the operator does not restore the instance.

However, setting the default instance as unmanaged might prevent operator updates from applying updates to the default instance. You can also consider ignoring the default instance (or setting it as unmanaged and deleting it, to save resources), and always creating your own instance. Depending on the criticality of your environment, consider extra testing of operator updates in a non-production environment.

Configuring Argo CD RBAC

Users can interact with Argo CD through the Kubernetes API, or through the Argo CD API, web interface, or command-line tool. Each Argo CD instance has a role-based access control system to manage permissions in the Argo CD interfaces.

Argo CD can use both local users that are defined in Argo CD and users from external authentication systems. The OpenShift GitOps operator includes features to automate integration with the OpenShift cluster authentication.

The default Argo CD instance includes a local `admin` user. This user has full permissions in the Argo CD interfaces. Argo CD updates the cluster by using a Kubernetes service account. Even unrestricted users in Argo CD are limited in how they can update the cluster, according to the permissions of the Argo CD service account.

The `openshift-gitops-cluster` secret in the `openshift-gitops` namespace contains the password for this user. Although this user can be suitable for small environments or experiments in test clusters, administrators usually integrate Argo CD with other authentication sources, such as the OpenShift cluster authentication.

The following excerpt from the default Argo CD instance shows the default authentication configuration:

```
apiVersion: argoproj.io/v1beta1
kind: ArgoCD
metadata:
  name: openshift-gitops
  namespace: openshift-gitops
  ...output omitted...
spec:
  ...output omitted...
  rbac:
    defaultPolicy: ""
    policy: |
      g, system:cluster-admins, role:admin
      g, cluster-admins, role:admin
    scopes: '[groups]'
    ...output omitted...
  sso:
    dex:
      openShiftOAuth: true
...output omitted...
```

rbac

The `rbac` key contains the Argo CD role-based access control configuration. This configuration grants the `admin` role to the `cluster-admins` group.

You can create a `cluster-admins` group in the cluster and add users to have administrator privileges for the Argo CD instance. Alternatively, you can edit the policy to grant administrator privileges to existing groups.

**Warning**

Argo CD resources, such as applications, are Kubernetes resources. Argo CD RBAC controls the use of resources through Argo CD, not through the Kubernetes API.

Use Kubernetes role-based access control to prevent access to Argo CD Kubernetes resources through the Kubernetes API.

sso

The sso key contains the authentication configuration. In the default Argo CD instance, the openShiftOAuth key is set to the true value. With this configuration, OpenShift users can log in to the Argo CD web console by clicking **LOG IN VIA OPENSHIFT**. By default, OpenShift users have limited privileges.

Configuring Trusted Certificates for Git Repository Access

Argo CD validates HTTPS certificates when accessing Git repositories, and can access only trusted repositories.

The *repo server* component of Argo CD handles communicating with Git repositories. To change the trusted certificate authorities, modify the certificate bundle in the `/etc/pki/ca-trust/extracted/pem/tls-ca-bundle.pem` path of the repo server.

The Argo CD custom resource contains fields to customize the repo server, similar to the `volumes` and `volumeMounts` keys of a pod template. The following example shows how to replace the certificate bundle with a bundle that is stored in a configuration map:

```
apiVersion: argoproj.io/v1beta1
kind: Argocd
metadata:
  name: openshift-gitops
  namespace: openshift-gitops
  ...output omitted...
spec:
  ...output omitted...
  repo:
    volumeMounts:
      - mountPath: /etc/pki/ca-trust/extracted/pem/tls-ca-bundle.pem
        name: volume_name
        subPath: bundle_name
    volumes:
      - configMap:
          name: configuration_map_name
        name: volume_name
  ...output omitted...
```

Configuring the Cluster Certificate Authority

If your repositories use certificates that the cluster certificate authority signed (for example, you expose them by using OpenShift routes), then you can use certificate injection to help configure Argo CD to trust the repositories.

Create an empty configuration map in the `openshift-gitops` namespace. Add the `config.openshift.io/inject-trusted-cabundle` label to the configuration map with the `true` value. OpenShift adds the bundle with the cluster certificate authority to the configuration map, in the `ca-bundle.crt` key. Then, you can follow the previous example to configure the repo server to trust the bundle.

Customizing the Argo CD Web Interface Certificate

The default Argo CD instance has its own self-signed HTTPS certificate. The Argo CD resource has a `termination` key to change the route termination. Use the `reencrypt` termination so that Argo CD uses the default router certificate.

```
apiVersion: argoproj.io/v1alpha1
kind: ArgoCD
metadata:
  ...output omitted...
spec:
  ...output omitted...
  server:
    ...output omitted...
    route:
      enabled: true
      tls:
        termination: reencrypt
...output omitted...
```

The Argo CD Web Interface

By default, the web console for an Argo CD instance is at the `argo_cd_instance-server-namespace.apps.cluster_domain` URL. Log in to Argo CD by accessing this URL.

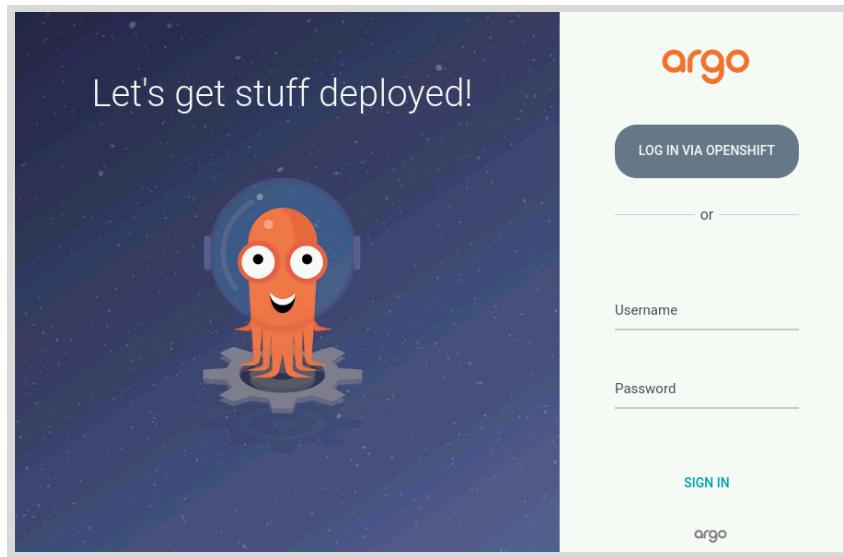


Figure 5.3: The Argo CD login page

If the Argo CD instance is configured with the `openShiftOAuth` key, then click `LOG IN VIA OPENSHIFT` to log in by using the cluster authentication.

Alternatively, you can use the **Username** and **Password** fields, and the **SIGN IN** button to use Argo CD local users.

When you log in, the web console displays the **Applications** page. The web console contains a navigation panel:

You can view and create applications in the **Applications** page. You can use the **User Info** page to view details about the logged-in user. These details can be useful to troubleshoot authentication issues.

The screenshot shows the Argo CD web interface. On the left is a sidebar with the Argo logo and version (v2.8.4+ebab8ec). It has links for **Applications**, **Settings**, **User Info**, and **Documentation**. Below these are filters for **Favorites Only** and **SYNC STATUS** (Unknown: 0, Synced: 1, OutOfSync: 0). The main area is titled "Applications" and shows a card for the "example" application. The card details are as follows:

- Project:** default
- Status:** Healthy Synced
- Repo:** https://github.com/argoproj/...
- Target:** HEAD
- Path:** .
- Destination:** in-cluster
- Name:** example
- Created:** 01/02/2024 08:01:07 (7 min...)

At the bottom of the card are three buttons: **SYNC**, **C**, and **X**.

Figure 5.4: The Argo CD applications page

The screenshot shows the "APPLICATION DETAIL" view for the "example" application. At the top, there are tabs for **APP DETAILS**, **APP DIFF**, **SYNC**, **SYNC STATUS**, **HISTORY AND ROLLBACK**, **DELETE**, and **REFRESH**. Below the tabs, there are two sections: "Sync Status" and "History".

Sync Status:

- Auto sync is not enabled.
- Author: Anand Francis Joseph <anandfrancis.josep...
- Comment: Template variable for container port (#251)
- Last sync failed a few seconds ago (Tue Jan 02 2024 08:12:43 GMT-0500).
- Author: Anand Francis Joseph <anandfrancis.josep...
- Comment: Template variable for container port (#251)

History:

The history shows four recent sync events:

- helm-hooks (ns) - a few seconds ago
- kustomize-guestbook (ns) - a few seconds ago
- sync-waves (ns) - a few seconds ago
- helm-guestbook (application) - 12 minutes ago

Figure 5.5: Application details

Chapter 5 | OpenShift GitOps

The application details page displays the Kubernetes resources that are associated with an Argo CD application. For each resource, you can view the synchronization status and resource details.



Important

When Argo CD displays a large amount of resources, the web interface might not be easy to work with.

Because Argo CD displays both the resources in your application and related resources, more resources might be displayed than expected.

To reduce the number of displayed resources, you can collapse resources to hide dependent resources.

With large applications, you might need to use other means of examining resources, such as the `kubectl` and `oc` commands, or the Kubernetes API.

Creating Applications

You can create Argo CD applications by using the Kubernetes API, the `argocd` command-line application, or the web console.

By default, the Argo CD web console shows a form to customize the application. The console also provides a YAML resource editor, which is similar to the OpenShift web console.

Figure 5.6: Application creation form

Use the following fields to create your application:

Application Name

The application resource name.

Project Name

Argo CD uses *projects* to group applications and to apply further access controls. Projects are not covered in this course. You can create applications in the `default` project that Argo CD provides.

Sync Policy

The `Automatic` policy synchronizes periodically. With the `Manual` policy, administrators must trigger synchronizations.

Retry

As described later in this lecture, you can use these fields to add retries to the synchronization process. You might need to configure retries when working with complex applications.

Repository URL

The URL of the repository with the application definition. If you use the HTTPS protocol to access the repository, then the repo server must trust the HTTPS certificate.

Path

The path in the Git repository that contains the resource definitions. You can define multiple applications from different paths in the same repository. A repository might also contain other content.

Cluster URL

The URL of the Kubernetes API for the cluster to deploy the application to. Use the `https://kubernetes.default.svc` URL to deploy to the same cluster where Argo CD is deployed.

When you create an application, Argo CD retrieves the contents of the repository. In the exercises of this course, applications are YAML files that describe application resources. Argo CD supports other ways to define applications, including the following technologies:

- Kustomize
- Helm charts
- Jsonnet
- Other configuration management tools that additional plug-ins support

Deploying Complex Resources by Using Sync Waves and Retries

Although you can usually create resources in any order and ultimately resources reach the intended state, in some cases creating resources can require further work. For example, you must create a namespace before creating the resources in the namespace. With custom resources, you usually need to wait until their operator installs the custom resource definitions.

You can use the `argocd.argoproj.io/sync-wave` annotation on resources to help with dependencies. The annotation value must be a positive or negative integer or zero. Resources without the annotation default to 0.

A group of resources with the same annotation is a *sync wave*. Argo CD synchronizes sync waves in order, starting with sync waves with lower numbers. Argo CD synchronizes a sync wave only when all lower sync waves are healthy.

Argo CD performs a dry run before synchronizations. When deploying custom resources, the synchronization continues only if the corresponding custom resource definitions exist when Argo CD performs the dry run. To prevent this problem, add the `argocd.argoproj.io/sync-options` annotation to these resources by adding the `SkipDryRunOnMissingResource=true` option.

For example, the following manifest defines an operator subscription and a custom resource:

```
---  
apiVersion: operators.coreos.com/v1alpha1  
kind: Subscription  
metadata:  
  name: example-operator  
  ...output omitted...
```

```
---  
kind: ExampleCustomResource  
metadata:  
  name: example  
  annotations:  
    argocd.argoproj.io/sync-options: SkipDryRunOnMissingResource=true  
    argocd.argoproj.io/sync-wave: "1"  
  ...output omitted...
```

Because the subscription does not have an explicit sync wave, Argo CD considers the subscription to be part of the `0` sync wave. Because the custom resource belongs to the `1` sync wave, Argo CD creates the custom resource only when the subscription is healthy. The sync options also exclude the custom resource from the dry run.

However, the previous example might still fail. The next section describes additional issues with deploying complex applications.

Eventual Consistency in Kubernetes

Kubernetes is a resilient distributed system, to fulfill its requirement to be a framework to run distributed systems resiliently. Building resilient distributed systems adds significant complexity.

Kubernetes uses the *eventual consistency* model to help achieve reliability and also preserving scalability. In distributed systems, changes must propagate to all nodes. Because systems that use eventual consistency do not apply updates immediately, updates can take a perceivable amount of time to be effective across the cluster or even in a single node.

Eventual consistency avoids some scalability and reliability issues. However, non-immediate updates add complexity to many processes, because some operations might need to wait until other operations are effective. For changes for non-immediate updates, one option is to retry operations until they succeed after dependent operations take effect. Software that works with Kubernetes often uses retries to work correctly.

Sync waves are often insufficient when working with complex applications such as operators. If you face issues with deploying applications, then consider configuring retries in your Argo CD applications for more reliable synchronization.

However, retries can clutter logs, slow down feedback loops, and hide problems. To limit problems that derive from excessive retries, consider investing effort in finding an adequate retry configuration.

Administrating Clusters with OpenShift GitOps

Although organizations often use Argo CD to deploy applications, cluster administrators can also use the default Argo CD instance to administer a cluster. By using Argo CD, administrators can customize authentication or deploy operators with all the benefits of GitOps.

When administering a cluster with Argo CD, you might need to patch existing resources, which is an uncommon scenario when deploying applications.

For example, cluster authentication uses the `cluster` resource of the `OAuth` type. You must edit this resource instead of creating other resources.

Argo CD can use the `server-side apply` process to edit existing resources. The server-side apply process is an alternative to the client-side apply process that `kubectl apply` and other processes use by default.

With client-side apply, commands such as `kubectl apply` start by retrieving the resource definition from the API server. If the resource exists, then the client compares this retrieved resource definition with the definition that you provide. Then, the client sends a PATCH request to the API server with the updated parts.

To edit an existing resource with server-side apply, your resource definition can include only your changes and exclude parts of the resource that you do not want to update. The client submits this partial resource definition, and the API server updates only the submitted parts.

For patching, use the `ServerSideApply=true` sync option. Also, add the `Validate=false` sync option to disable validation, because the partial resource definition might not be a valid complete resource definition.

The following manifest shows how to patch a resource:

```
apiVersion: operator.openshift.io/v1
kind: Console
metadata:
  name: cluster
  annotations:
    argocd.argoproj.io/sync-options: ServerSideApply=true,Validate=false
spec:
  customization:
    customProductName: Production
```

This specification does not contain all the fields that a full console specification requires. The `Validate=false` sync option ensures that even if the partial definition lacks mandatory fields, Argo CD submits the change.

This change is less likely to interfere with other changes to the cluster such as cluster updates.



References

For more information about the installation process, refer to the *Installing Red Hat OpenShift GitOps* chapter in the Red Hat OpenShift GitOps 1.10 *Installing GitOps* documentation at

https://access.redhat.com/documentation/en-us/red_hat_openshift_gitops/1.10/html-single/installing_gitops/index#installing-openshift-gitops

For more information about certificate injection, refer to the *Certificate Injection Using Operators* section in the *Configuring a Custom PKI* chapter in the Red Hat OpenShift Container Platform 4.14 *Networking* documentation at

https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/networking/index#certificate-injection-using-operators_configuring-a-custom-pki

For more information about Argo CD role-based access control, refer to the *Access Control and User Management* chapter in the Red Hat OpenShift GitOps 1.10 *Access Control and User Management* documentation at

https://access.redhat.com/documentation/en-us/red_hat_openshift_gitops/1.10/html-single/access_control_and_user_management/index#configuring-argo-cd-rbac

For more information about cluster administration with OpenShift GitOps, refer to the *Declarative Cluster Configuration* chapter in the Red Hat OpenShift GitOps 1.10 *Declarative Cluster Configuration* documentation at

https://access.redhat.com/documentation/en-us/red_hat_openshift_gitops/1.10/html-single/declarative_cluster_configuration/index#configuring-an-openshift-cluster-by-deploying-an-application-with-cluster-configurations

Kubernetes: Server-Side Apply

<https://kubernetes.io/docs/reference/using-api/server-side-apply/>

► Guided Exercise

GitOps for Cluster Administration

Deploy an add-on operator and custom resources by using GitOps.

Outcomes

- Install Red Hat OpenShift GitOps.
- Configure the default Argo CD instance.
- Create an Argo CD application that deploys an operator and custom resources from the operator.
- Patch a cluster resource by using server-side apply.

Before You Begin

As the student user on the workstation machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start gitops-admin
```

Instructions

Your cluster had downtime. During incident analysis, system administrators discovered changes to the cluster without proper documentation.

Your organization wants to implement an audit track for changes to the cluster. You decide to experiment with Red Hat OpenShift GitOps to achieve this objective.

Your experiment is to create a test cluster, install OpenShift GitOps, and perform some administrative changes to the test cluster. You test installing and using an operator, and patching existing resources.

► 1. As the `admin` user, locate and then navigate to the Red Hat OpenShift web console.

- 1.1. Use the terminal to log in to the OpenShift cluster as the `admin` user with `redhatocp` as the password.

```
[student@workstation ~]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
...output omitted...
```

- 1.2. Identify the URL for the OpenShift web console.

```
[student@workstation ~]$ oc whoami --show-console
https://console-openshift-console.apps.ocp4.example.com
```

Chapter 5 | OpenShift GitOps

- 1.3. Open a web browser and navigate to `https://console-openshift-console.apps.ocp4.example.com`. Either type the URL in a web browser, or right-click and select **Open Link** from the terminal.
 - 1.4. Click **Red Hat Identity Management** and log in as the **admin** user with `redhatocp` as the password.
- ▶ 2. Install the OpenShift GitOps operator from OperatorHub.
- 2.1. Navigate to **Operators > OperatorHub**.
 - 2.2. Click **Red Hat OpenShift GitOps**, and then click **Install**.
 - 2.3. Review the default configuration and click **Install**. The Operator Lifecycle Manager (OLM) can take a few minutes to install the operator. Click **View Operator** to navigate to the operator details.
 - 2.4. Open a separate tab and open the default Argo CD instance. You can use the application menu, which is the grid icon on the top navigation bar, by clicking **Cluster Argo CD**. You can also use the `https://openshift-gitops-server-openshift-gitops.apps.ocp4.example.com` URL.
The browser displays a warning because Argo CD uses a self-signed certificate. Argo CD might take a few minutes before starting to handle requests.
- ▶ 3. Disconnect the default instance from the operator so the operator does not revert changes to the instance.
- If you use the web console for this operation, then click the **Argo CD** tab to display the default `openshift-gitops` Argo CD instance in the `openshift-gitops` namespace. Click the name of the instance, and then click the **YAML** tab to display the resource editor. Perform the modifications that are described later, and then click **Save**.
- If you use the terminal for this operation, then run the following command:
- ```
[student@workstation ~]$ oc edit -n openshift-gitops argocd openshift-gitops
```
- Remove the `ownerReferences` key from the `metadata` key in the resource, and save your changes. The resulting resource definition should resemble this extract:
- ```
apiVersion: argoproj.io/v1alpha1
kind: ArgoCD
metadata:
  creationTimestamp: "2023-11-29T18:47:29Z"
  finalizers:
  - argoproj.io/finalizer
  generation: 1
  name: openshift-gitops
  namespace: openshift-gitops
  resourceVersion: "398331"
  uid: 00778335-f7d8-457d-92c2-121cd13f5d26
spec:
  applicationSet:
  ...output omitted...
```
- ▶ 4. For the route of the default instance, change the termination type to the `reencrypt` type. Edit the Argo CD resource to match the following example:

Chapter 5 | OpenShift GitOps

```
apiVersion: argoproj.io/v1alpha1
kind: ArgoCD
metadata:
  ...output omitted...
spec:
  ...output omitted...
  server:
    ...output omitted...
    route:
      enabled: true
      tls:
        termination: reencrypt
  ...output omitted...
```

Reload the Argo CD browser tab. Instead of the certificate warning, the Argo CD login page is shown.

▶ **5.** Grant administrator rights to the `ocpadmins` group.

5.1. Edit the Argo CD resource to match the following example:

```
apiVersion: argoproj.io/v1alpha1
kind: ArgoCD
metadata:
  ...output omitted...
spec:
  ...output omitted...
  prometheus:
    enabled: false
    ingress:
      enabled: false
    route:
      enabled: false
  rbac:
    policy: |
      g, ocpadmins, role:admin
      scopes: '[groups]'
    redis:
  ...output omitted...
```

▶ **6.** Configure the default instance to trust the cluster certificate authority. Argo CD accesses only trusted repositories.

6.1. Create a `cluster-root-ca-bundle` configuration map in the `openshift-gitops` namespace.

```
[student@workstation ~]$ oc create configmap -n openshift-gitops \
cluster-root-ca-bundle
```

6.2. Add the `config.openshift.io/inject-trusted-cabundle` label to the configuration map with the `true` value. OpenShift injects the bundle with the cluster certificate authority into the configuration maps with this label. This bundle contains the signing certificate for the classroom GitLab instance.

Chapter 5 | OpenShift GitOps

```
[student@workstation ~]$ oc label configmap -n openshift-gitops \
  cluster-root-ca-bundle config.openshift.io/inject-trusted-cabundle=true
configmap/cluster-root-ca-bundle labeled
```

- 6.3. Edit the Argo CD default instance to inject the bundle.

You can use the following command to edit the resource:

```
[student@workstation ~]$ oc edit argocd -n openshift-gitops openshift-gitops
```

Edit the resource to mount the ca-bundle.crt file in the configuration map to the /etc/pki/ca-trust/extracted/pem/tls-ca-bundle.pem path of the repository server container.

```
...output omitted...
spec:
  ...output omitted...
  repo:
    resources:
      limits:
        cpu: "1"
        memory: 1Gi
      requests:
        cpu: 250m
        memory: 256Mi
    volumeMounts:
      - mountPath: /etc/pki/ca-trust/extracted/pem/tls-ca-bundle.pem
        name: cluster-root-ca-bundle
        subPath: ca-bundle.crt
    volumes:
      - configMap:
          name: cluster-root-ca-bundle
          name: cluster-root-ca-bundle
  resourceExclusions: |
...output omitted...
```

► 7. Create a public repository in the classroom GitLab.

- 7.1. Open a web browser and navigate to <https://git.ocp4.example.com>. Log in as the developer user with d3v3lop3r as the password.
- 7.2. Click **New project**, and then click **Create blank project**. Use **gitops-admin** as the project slug (repository name), select the **Public** visibility level, and use the default values for all other fields. Click **Create project**.

► 8. Populate the repository.

- 8.1. Click **Clone**, and then copy the <https://git.ocp4.example.com/developer/gitops-admin.git> HTTPS URL.
- 8.2. Change to the ~/D0380/labs/gitops-admin directory.

```
[student@workstation ~]$ cd ~/D0380/labs/gitops-admin
```

Chapter 5 | OpenShift GitOps

8.3. In a terminal, run the following command to clone the new repository.

```
[student@workstation gitops-admin]$ git clone \
https://git.ocp4.example.com/developer/gitops-admin.git
Cloning into 'gitops-admin'...
...output omitted...
```

8.4. Change to the cloned repository directory.

```
[student@workstation gitops-admin]$ cd gitops-admin
```

The default configuration for new repositories adds a `README.md` initial file.

8.5. Copy the provided `operator.yaml` file to the repository.

```
[student@workstation gitops-admin]$ cp ./operator.yaml .
```

8.6. Examine the file and edit the file to match the following text:

```
apiVersion: v1
kind: Namespace
metadata:
  name: openshift-compliance
---
apiVersion: operators.coreos.com/v1
kind: OperatorGroup
...output omitted...
---
apiVersion: operators.coreos.com/v1alpha1
kind: Subscription
metadata:
  name: compliance-operator
...output omitted...
---
apiVersion: compliance.openshift.io/v1alpha1
profiles:
  - apiGroup: compliance.openshift.io/v1alpha1
    name: rhcos4-moderate
    kind: Profile
settingsRef:
  apiGroup: compliance.openshift.io/v1alpha1
  name: default
  kind: ScanSetting
kind: ScanSettingBinding
metadata:
  name: nist-moderate
  namespace: openshift-compliance
  annotations:
    argocd.argoproj.io/sync-options: SkipDryRunOnMissingResource=true
    argocd.argoproj.io/sync-wave: "1"
```

The file contains manifests to install the compliance operator, with the namespace, operator group, and subscription.

The file also contains a scan setting binding. The compliance operator examines scan setting bindings and scans the cluster after they are defined.

Because scan setting bindings are custom resources that the compliance operator installs, you must use extra configuration so that Argo CD can create the resource correctly:

- The dry run must be skipped, because when the Argo CD application is created, the custom resource definition does not exist, so the Argo CD validation would fail.
- The scan setting binding must be created after the custom resource definition exists. To delay the creation, you specify that the scan setting binding is created in a sync wave after the default sync wave. However, due to the Kubernetes use of eventual consistency, later you also configure retries for the application deployment to prevent further issues.

8.7. Add the `operator.yaml` file to the Git index.

```
[student@workstation gitops-admin]$ git add operator.yaml
```

8.8. Commit the changes.

```
[student@workstation gitops-admin]$ git commit -m "Add compliance manifests"
[main 6785970] Add compliance manifests
...output omitted...
```

8.9. Push the changes to the repository. Use the `developer` user with `d3v3lop3r` as the password.

```
[student@workstation gitops-admin]$ git push
...output omitted...
```

► 9. Log in to the Argo CD web console as the `admin` user.

9.1. Go back to the Argo CD browser tab.

9.2. Click **LOG IN VIA OPENSHIFT**, and then click **Red Hat Identity Management**. Log in as the `admin` user with `redhatocp` as the password, and then allow the `user:info` permission.

► 10. Create an application with the repository and observe the results.

10.1. Click **CREATE APPLICATION**.

10.2. Create an application with the information in the following table:

Chapter 5 | OpenShift GitOps

Field	Value
Application Name	gitops-admin
Project Name	default
Retry	Checked
Repository URL	https://git.ocp4.example.com/developer/gitops-admin.git
Path	.
Cluster URL	https://kubernetes.default.svc

Then, click CREATE.

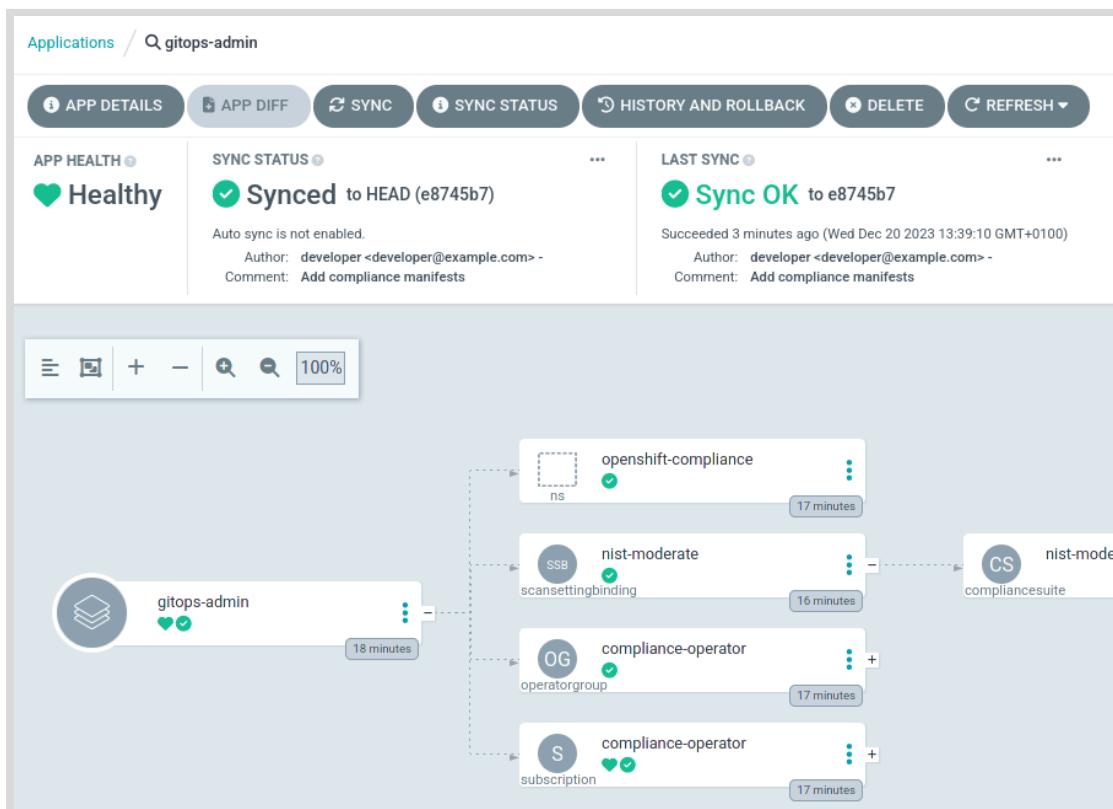
► 11. Synchronize the application.

11.1. Click **gitops-admin** to view the application.

11.2. Click **SYNC** to display the synchronization panel, and then click **SYNCHRONIZE**.

Argo CD starts synchronizing the application. After about one minute, the console shows the compliance operator as synchronized and healthy.

If you wait about four minutes, then the compliance operator scans and creates a compliance suite custom resource.



You can display the results of the scan from Argo CD by clicking the **nist-moderate** compliance suite. The compliance suite is in the done phase, with the non-compliant result.

You can also monitor the progress of the scan by running the `watch oc get compliancesuite -A` command.



Note

The many resources that the application tracks can cause the Argo CD web console to become unstable. If the web console does not respond, then you can close the browser tab and reopen the Argo CD web console.

Collapsing the resources in the Argo CD web console can help with exploring the resources.

► 12. Customize the OpenShift console.

12.1. Copy the provided `console.yaml` file to the repository.

```
[student@workstation gitops-admin]$ cp ./console.yaml .
```

12.2. Edit the file to match the following text:

```
apiVersion: operator.openshift.io/v1
kind: Console
metadata:
  name: cluster
  annotations:
    argocd.argoproj.io/sync-options: ServerSideApply=true,Validate=false
spec:
  customization:
    customProductName: Production
```

Because the `cluster` console resource exists, this manifest defines only a patch to update the product name. For Argo CD to patch the resource, you must use an annotation. Because some patches can be valid resources, but other patches can be invalid resources, disable validation to ensure that patches are not validated as full resources.

12.3. Add the `console.yaml` file to the Git index.

```
[student@workstation gitops-admin]$ git add console.yaml
```

12.4. Commit the changes.

```
[student@workstation gitops-admin]$ git commit -m "Customize console"
...output omitted...
```

12.5. Push the changes to the repository.

```
[student@workstation gitops-admin]$ git push
...output omitted...
```

- 12.6. Click **SYNC** to display the synchronization panel, and then click **SYNCHRONIZE**.
 - 12.7. Reload the OpenShift web console to observe the changes. After you reload the console, the browser tab title ends with the **Production** text.
- **13.** Change to the student HOME directory.

```
[student@workstation gitops-admin]$ cd
```

Finish

On the **workstation** machine, use the **lab** command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish gitops-admin
```

GitOps for Application Management

Objectives

- Deploy an Argo CD instance, from OpenShift GitOps, for application developers and application administrators.

Argo CD Instances

By default, the Red Hat OpenShift GitOps operator creates an Argo CD instance in the `openshift-gitops` namespace. This Argo CD instance has permissions to manage cluster-wide resources. Thus, Red Hat recommends using this instance only for administration purposes, and to create other, less privileged instances for regular application deployment.

To deploy regular applications, you can deploy new Argo CD instances. By default, the new Argo CD instances have permissions to manage resources only in the namespace where they are deployed.

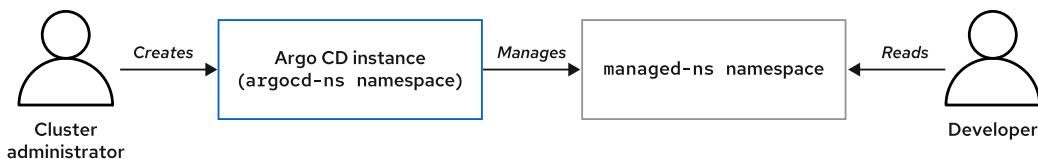
To create an Argo CD instance, you can create a YAML file as follows:

```
apiVersion: argoproj.io/v1beta1
kind: ArgoCD
metadata:
  name: argocd-name
  namespace: argocd-ns
spec:
  ...output omitted...
```

You can allow an Argo CD instance to manage resources in other namespaces than where it is deployed by adding the `argocd.argoproj.io/managed-by` label to the namespace. The following example allows the Argo CD instance in the `argocd-ns` namespace to manage the resources in the `managed-ns` namespace.

```
[user@host ~]$ oc label namespace managed-ns \
  argocd.argoproj.io/managed-by=argocd-ns
```

The following diagram shows an example of using an Argo CD instance to manage resources in other namespaces.



In this diagram, a cluster administrator creates an Argo CD instance in the `argocd-ns` namespace. Thus, the Argo CD instance has permissions to manage resources only in that namespace.

The cluster administrator includes the `argocd.argoproj.io/managed-by` label in the `managed-ns` namespace. Thus, the Argo CD instance can manage resources in that namespace.

Developer users need only read permission in the `managed-ns` namespace for troubleshooting, according to the principle of least privilege, because they would be managing application resources in their projects by using Git.

Argo CD Authentication

Argo CD can both use local users and integrate with other identity providers.

The OpenShift GitOps operator configures the default Argo CD instance to use Dex. Dex is an identity service that uses the OpenID connect protocol to connect Argo CD with the identity provider.

The default Argo CD instance in the `openshift-gitops` namespace has permission to manage cluster-wide resources. You can log in with write permission to this instance only by using the local `admin` user with the credentials in the `openshift-gitops-cluster` secret in the `openshift-gitops` namespace, or by using an OpenShift user from the `cluster-admins` group in OpenShift.

You can also configure the Argo CD instances to use Keycloak instead of Dex. For example, use the groups in Keycloak to determine the privileges in Argo CD. Keycloak acts as an identity broker between the Argo CD instance and OpenShift.



Note

Configuring Keycloak as the SSO authentication provider for Argo CD is outside the scope of this course.

For more information, refer to https://access.redhat.com/documentation/en-us/red_hat_openshift_gitops/1.10/html-single/access_control_and_user_management/index#configuring-sso-for-argo-cd-using-keycloak

Dex

The OpenShift GitOps operator configures Dex to delegate the authentication to the built-in OAuth server in OpenShift. Thus, Dex provides the users and groups that are defined in OpenShift.

You can manually configure Dex as the SSO authentication provider and use the OpenShift OAuth server by setting the `.spec.sso` parameter as in the following example:

```
...output omitted...
spec:
  sso:
    provider: dex
    dex:
      openShiftOAuth: true
...output omitted...
```

Argo CD Permissions

By default, only the local Argo CD `admin` user and the OpenShift users in the `cluster-admins` group can log in to the default Argo CD instance with write permission. Other OpenShift

Chapter 5 | OpenShift GitOps

users have only read permission. However, you can change the user and group level access by configuring the RBAC section in the Argo CD custom resource.



Important

Argo CD RBAC controls the use of resources through Argo CD, not through the Kubernetes API. Users with access to the Kubernetes API can modify those resources and modify Argo CD RBAC. You must set up Kubernetes role-based access control to prevent access to Argo CD Kubernetes resources.

Argo CD comes with the following predefined roles:

- `role:readonly` for read access to all the resources.
- `role:admin` for read and write access to all the resources.

For example, the default Argo CD instance includes the following RBAC configuration:

```
...output omitted...
spec:
  ...output omitted...
  rbac:
    defaultPolicy: ""
    policy: |
      g, system:cluster-admins, role:admin
      g, cluster-admins, role:admin
      scopes: '[groups]'
...output omitted...
```

Thus, for OpenShift users in the `system:cluster-admins` and `cluster-admins` groups, the predefined `role:admin` role is assigned by default.



Note

System groups, such as the `system:cluster-admins` group, are predefined groups that are built into the system to grant specific permissions to certain categories of users. Kubernetes creates these groups during the setup to control access to critical components and resources within the cluster. These groups are a part of cluster roles and cluster role bindings.

For more information, refer to <https://kubernetes.io/docs/reference/access-authn-authz/rbac/#referring-to-subjects>

To configure the Argo CD instance to assign the predefined `role:readonly` role to the users in the `cluster-readers` group, you can use the following RBAC configuration:

```
...output omitted...
spec:
  ...output omitted...
  rbac:
    defaultPolicy: ""
    policy: |
      g, system:cluster-admins, role:admin
      g, cluster-admins, role:admin
      g, cluster-readers, role:readonly
...output omitted...
```

```

g, cluster-readers, role:readonly
scopes: '[groups]'
...output omitted...

```

You can use the `defaultPolicy` field to define a policy for other users that do not conform to an existing policy.

Fine-grained RBAC

If you want a finer-grained RBAC for your Argo CD instance, then you can break down the permissions for all the resources in Argo CD. The resources in Argo CD are accounts, applications, applicationsets, certificates, clusters, exec, extensions, gpgkeys, logs, projects, and repositories.

For all the resources, the Argo CD actions include `create`, `delete`, `get`, and `update`. For the `applications` resource, Argo CD also includes the `action/group/kind/action-name`, `override`, and `sync` actions.

You can define specific fine-grained RBAC permissions with the following syntax in the `spec.rbac.policy` parameter:

```
p, role/user/group, resource, action, target
```

The `target` field differs between the application resources and the other resources:

- For the `applications`, `applicationsets`, `logs`, and `exec` resources, which belong to a project, the `target` field is `project/object`.
- For other resources, the `target` field is `object`.

As an example, the following excerpt creates the `project-devs` and `project-admins` roles, and assigns them to the users in the OpenShift groups with the same names.

```

...output omitted...
spec:
...output omitted...
rbac:
  defaultPolicy: ''
  policy: [
    g, project-devs, role:project-devs
    p, role:project-devs, applications, get, /*, allow
    p, role:project-devs, projects, get, *, allow
    p, role:project-devs, clusters, get, *, allow
    g, project-admins, role:project-admins
    p, role:project-admins, applications, *, /*, allow
    p, role:project-admins, projects, get, *, allow
    p, role:project-admins, clusters, get, *, allow
  scopes: '[groups]'
...output omitted...

```

Users in both the `project-devs` and `project-admins` groups can read the `clusters` and `projects` resources in Argo CD.

Users in the `project-devs` group can also read all the `applications` resources in any of the projects.

Chapter 5 | OpenShift GitOps

Users in the project-admins group can perform any action for the applications resources in any of the projects.



Note

Argo CD includes other RBAC features that are outside the scope of this course.

For more information, refer to <https://argo-cd.readthedocs.io/en/stable/operator-manual/rbac/>

Monorepo and Polyrepo Environments

If you give a user access to a Git repository, then the user has access to all the files in the repository. You cannot allow users to read or write only in certain paths within a repository. However, some products that build on top of Git, such as GitHub or GitLab, add features to improve protecting specific paths within a repository.

When setting up Git workflows for GitOps, you can either create only one repository for all your files, which is called a *monorepo environment*, or create repositories for different concerns, which is called a *polyrepo environment*.

In a monorepo environment, all the application sources, container files, Kubernetes manifests, and policy manifests are in a single Git repository.

A monorepo environment offers a centralized location for all the configuration changes. Thus, tracking, testing, and releasing processes are limited to a single Git repository. Moreover, having all your code in a single repository simplifies managing dependencies between your applications. However, users with access to the Git repository have access to all the code in the projects, which can lead to security issues.

From an Argo CD perspective, you can either treat the entire monorepo as a single application, to deploy changes across various services or components concurrently, or use paths to create several applications from a single repository. With the second option, you can use different paths for the application files and for the application configurations, such as network policies or quotas. This separation prevents application configuration changes from triggering unnecessary application rebuilds.

Another option when setting up Git workflows for GitOps is to separate the concerns by repository. A polyrepo environment enables different designs depending on your company needs.

For example, a polyrepo environment has several Git repositories for applications and configurations. Red Hat recommends using different repositories for application files and for application configurations, to prevent unnecessary application rebuilds. Independent lifecycles for application code and application configuration, along with specific approval processes, ensure smoother continuous integration. However, coordinating dependencies across repositories requires extra effort.

From an Argo CD perspective, you can create applications from the various repositories in a polyrepo environment. However, a polyrepo environment requires more complex configuration than a monorepo environment to manage the dependencies between repositories.

Organizing Applications by Environment

One GitOps practice separates the test and production environments by directory inside a Git repository, instead of creating different branches for them. This approach avoids the

complexity of managing branches for various environments, where merging changes might not be straightforward because of configuration differences.

The goal is to avoid resource duplication across environments. Given that organizations often maintain many environments from development to production, replicating the entire set of deployment assets across these environments can lead to inconsistencies.

Tools such as Kustomize or Helm create reusable and modular configurations, to separate base manifests and environment-specific adjustments. With Kustomize declarative and scalable configuration management, you can create a set of base resource files to deploy the application, and then create a set of overlay files to adapt the application for the environment. With Helm, you can store in Git a Helm chart that contains all the necessary resources to deploy your application, and then use different Git repositories to store the application customization for each environment. Then, Argo CD uses the Helm chart and you can specify in Argo CD the values file to use for each environment.



Note

Kustomize is outside the scope of this course.

Using Kustomize in an OpenShift cluster is explained in the *DO280: Red Hat OpenShift Administration II: Operating a Production Kubernetes Cluster* course.

For more information, refer to <https://kubernetes.io/docs/tasks/manage-kubernetes-objects/kustomization/>

Managing Sensitive Data in GitOps

To prevent security issues, you must ensure encryption or other protection of sensitive data or secrets. Sensitive data and secrets include passwords, APIs, encryption keys, SSH keys, tokens, or other digital credentials that allow users and applications access to other sensitive data, systems, and services.

Because Git users with access to the repository have access to all the files in the repository, they have access to any sensitive data that is stored in the repository. Thus, storing your sensitive data in Git repositories can introduce security concerns.

You can manually store your sensitive data in secret objects, outside a Git repository. However, if you do not correctly configure RBAC rules on your cluster, then anyone with API or etcd access can retrieve or modify a secret. Moreover, a user with permission to create a pod in a namespace can read any secret in that namespace. By default, cluster administrators can see the secrets of all users.

A mitigation is to use third-party secret management tools. Red Hat recommends using these secret management tools to mitigate security issues with features such as credential encryption, credential rotation, and short-lived tokens. Both commercial and open source solutions exist for this purpose.

For example, you can use the Secrets Store CSI Driver operator to manage multiple secrets, keys, and certificates. This driver enables Kubernetes pods to mount secrets that are stored in external secret management systems such as Azure Key Vault, AWS Secrets Manager, or HashiCorp Vault as volumes. This driver enhances security by centralizing secret management outside the cluster and allowing applications to access these secrets seamlessly through the Kubernetes API. The operator simplifies integrating and managing external secrets within Kubernetes deployments, to handle sensitive information more securely and efficiently.

Application Validation and Rollback

When you deploy an application to production, even with thorough testing you cannot predict or prevent all potential issues. Also, deploying an application requires actions before, during, and after application synchronization, and if that process fails.

For any production issues, you can roll back to known stable states the small increments of your application that CI/CD practices release. Increasing the release size also results in larger rollbacks, and implies more risks.

Argo CD provides some of these features for deploying your application.

Argo CD Resource Hooks

Argo CD provides *resource hooks*, which are customizable triggers in deployment lifecycles. With Argo CD resource hooks, you can define and execute actions during the lifecycle of an application synchronization process.

Argo CD can run resource hooks before, during, and after an application synchronization process, and also if the synchronization process fails. For example, you can run a resource hook before the synchronization process to migrate a database before deploying a new application version, or you can run a resource hook after the synchronization process for a smoke test of your application by verifying the application health.

Argo CD defines the following resource hooks:

PreSync

Argo CD executes these hooks before applying the manifests.

Sync

Argo CD executes these hooks only after it successfully completes all the PreSync hooks. These hooks are applied at the same time as applying the manifests.

Skip

This hook indicates to Argo CD to skip applying the manifests. Use this hook for resources that Argo CD must create, but that other resources modify and that must be out of Argo CD synchronization.

PostSync

Argo CD executes these hooks only after it successfully completes all the Sync hooks, and the application and its resources are in a healthy state.

SyncFail

Argo CD executes these hooks whenever the synchronization operation fails.

Argo CD resource hooks are Kubernetes resources that include the `argocd.argoproj.io/hook` annotation. Typically, Argo CD resource hooks are implemented through Kubernetes jobs.

The following example shows the annotation for a PreSync hook that migrates the database before deploying the new application version:

```
apiVersion: batch/v1
kind: Job
metadata:
  generateName: schema-migrate-database
  annotations:
    argocd.argoproj.io/hook: PreSync
...output omitted...
```

You can add the `argocd.argoproj.io/hook-delete-policy` annotation for your hooks, so Argo CD deletes them automatically according to one of the following policies:

HookSucceeded

Delete the hook resource after the hook succeeds.

HookFailed

Delete the hook resource after the hook fails.

BeforeHookCreation

Delete any existing hook resource before a new one is created.

The following excerpt shows the annotation for a `PostSync` hook that tests the application after the synchronization process and removes the hook if it succeeds:

```
apiVersion: batch/v1
kind: Job
metadata:
  generateName: integration-test-app
  annotations:
    argocd.argoproj.io/hook: PostSync
    argocd.argoproj.io/hook-delete-policy: HookSucceeded
...output omitted...
```

Argo CD Rollback

Argo CD enables you to roll back your application to a known stable state if the new version has errors.

You can manually roll back your application by clicking the **HISTORY AND ROLLBACK** button. You can then access previous application deployments and restore the application to an earlier version. For example, with this feature you can restore the application to a working version until you push a fix in your application repository.

Argo CD can also automate rollbacks in certain conditions, such as for failing resource hooks.



Note

The Argo CD rollback feature is disabled if you enable automatic synchronization for your application, because Argo CD automatically synchronizes the application to the latest Git commit version.

To roll back your application, you can temporarily disable automatic synchronization in Argo CD, or revert the last commit in your Git repository.

Argo Rollouts

Argo Rollouts is a Kubernetes controller and a set of custom resource definitions that provide advanced deployment capabilities. These capabilities include update strategies, such as progressive rollouts, and automated rollbacks and promotions.

You can integrate Argo Rollouts with ingress controllers and service meshes, to gradually shift the traffic from your previous application version to a new one.

For example, you can use Argo Rollouts to deploy a new application version for only a small percentage of the production traffic, and then analyze if everything works as expected. Argo Rollouts can automatically roll back to the previous stable version if the deployment fails.



Note

Argo Rollouts is a Technology Preview feature only, and is outside the scope of this course.



References

For more information about configuring SSO on Argo CD by using Dex or Keycloak, refer to the Red Hat OpenShift GitOps 1.10 *Access Control and User Management* documentation at

https://access.redhat.com/documentation/en-us/red_hat_openshift_gitops/1.10/html-single/access_control_and_user_management/index

For more information about Argo Rollouts, refer to the Red Hat OpenShift GitOps 1.10 *Argo Rollouts* documentation at

https://access.redhat.com/documentation/en-us/red_hat_openshift_gitops/1.10/html-single/argo_rollouts/index

For more information about best practices for structuring Git workflows, refer to the *The Path to GitOps* book at

<https://developers.redhat.com/e-books/path-gitops>

For more information about implementing GitOps with Argo CD in OpenShift, refer to the *Getting GitOps: A Practical Platform with OpenShift, Argo CD, and Tekton* book at

<https://developers.redhat.com/e-books/getting-gitops-practical-platform-openshift-argo-cd-and-tekton>

► Guided Exercise

GitOps for Application Management

Deploy an application and its dependencies by using OpenShift GitOps.

Outcomes

- Create an Argo CD instance for developers with the appropriate permissions.
- Use Argo CD to deploy an application in the cluster.

Before You Begin

As the student user on the workstation machine, use the `lab start gitops-app` command to prepare your environment for this exercise, and to ensure that all required resources are available.

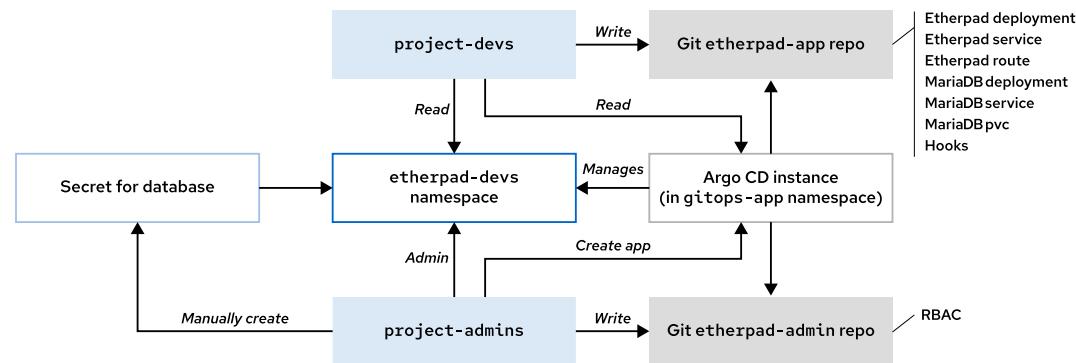
```
[student@workstation ~]$ lab start gitops-app
```

Instructions

Your company requires you to create a separate instance of Argo CD in the `gitops-app` namespace, for a project administration team and a project developer team that must deploy an application.

For the Argo CD instance, the `project-admins` group has full permissions for Argo CD applications, but only read permission for projects and clusters. For the same instance, the `project-devs` group has only read permission for Argo CD applications, projects, and clusters. The `project-admin` user is part of the `project-admins` group, and the `developer` user is part of the `project-devs` group.

The following diagram summarizes the relationship between the components in the exercise:



In this exercise, GitLab is configured with two repositories, one for project administration and one for project developer teams. The repository for the `developer` user is populated with the necessary files, including a MariaDB database and the Etherpad application. The `project-admin` user has one empty `etherpad-admin` repository.

Chapter 5 | OpenShift GitOps

Create the `etherpad-devs` namespace by using the OpenShift `admin` user. You must create the application in the `etherpad-devs` namespace. The Argo CD instance must manage the `etherpad-devs` namespace. The OpenShift `admin` user gives administrative access to the `etherpad-devs` namespace for users in the `project-admins` group.

The `project-admin` admin user creates a secret with the credentials for the MariaDB database to be deployed. This setup ensures that the database credentials are not stored in Git; storing in Git could lead to security issues.

At the beginning, the `developer` user has no permissions on the `etherpad-devs` namespace. Thus, the `project-admin` user must upload an RBAC file to the `etherpad-admin` repository and use Argo CD to give the `developer` user read permission in the project. Then, the `developer` user can view logs for troubleshooting if necessary.

Next, the `project-admin` user creates the MariaDB and Etherpad applications. The `project-admin` user also creates a PVC for the database backup.

The `developer` user creates two Argo CD hooks: one pre-synchronization hook that makes a database backup in a separate PVC, and one post-synchronization hook that verifies that the Etherpad application is running.

Finally, the `developer` user modifies the title for the Etherpad application in the repository and sees how the application is updated in Argo CD.

- ▶ 1. Create the Argo CD instance in the `gitops-app` namespace. The Argo CD must trust the classroom certificate.
 - 1.1. Connect to the OpenShift cluster as the `admin` user with `redhatocp` as the password.

```
[student@workstation ~]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
Login successful.

You have access to 70 projects, the list has been suppressed. You can list all
projects with 'oc projects'

Using project "default".
```

- 1.2. Change to the `gitops-app` project.

```
[student@workstation ~]$ oc project gitops-app
Now using project "gitops-app" on server "https://api.ocp4.example.com:6443".
```

- 1.3. Create a `cluster-root-ca-bundle` configuration map in the `gitops-app` namespace.

```
[student@workstation ~]$ oc create configmap cluster-root-ca-bundle
configmap/cluster-root-ca-bundle created
```

- 1.4. Add the `config.openshift.io/inject-trusted-cabundle` label to the configuration map with the `true` value. OpenShift injects the bundle with the cluster certificate authority to the configuration maps with this label. This bundle contains the signing certificate for the classroom GitLab instance.

```
[student@workstation ~]$ oc label configmap cluster-root-ca-bundle \
  config.openshift.io/inject-trusted-cabundle=true
configmap/cluster-root-ca-bundle labeled
```

15. Create the Argo CD instance YAML file. Set the termination type of the route to the reencrypt type. Grant read permission to projects and clusters to the `project-admins` and `project-devs` groups. For applications, grant full permissions to the `project-admins` group, and read permission to the `project-devs` groups. The Argo CD instance must mount the `ca-bundle.crt` certificate in the configuration map to the `/etc/pki/ca-trust/extracted/pem/tls-ca-bundle.pem` path of the repository server container. You can find an incomplete example for the Argo CD CR in the `~/DO380/labs/gitops-app/argocd-instance.yaml` file.

```
apiVersion: argoproj.io/v1beta1
kind: ArgoCD
metadata:
  name: argocd
  namespace: gitops-app
spec:
  server:
    ...output omitted...
  route:
    enabled: true
    tls:
      termination: reencrypt
  ...output omitted...
  rbac:
    defaultPolicy: ''
    policy: |
      g, system:cluster-admins, role:admin
      g, project-devs, role:project-devs
      p, role:project-devs, applications, get, /*, allow
      p, role:project-devs, projects, get, *, allow
      p, role:project-devs, clusters, get, *, allow
      g, project-admins, role:project-admins
      p, role:project-admins, applications, *, /*, allow
      p, role:project-admins, projects, get, *, allow
      p, role:project-admins, clusters, get, *, allow
    scopes: '[groups]'
  repo:
    resources:
      limits:
        cpu: 1000m
        memory: 1024Mi
      requests:
        cpu: 250m
        memory: 256Mi
    volumeMounts:
      - mountPath: /etc/pki/ca-trust/extracted/pem/tls-ca-bundle.pem
        name: cluster-root-ca-bundle
        subPath: ca-bundle.crt
    volumes:
      - configMap:
```

Chapter 5 | OpenShift GitOps

```
name: cluster-root-ca-bundle  
name: cluster-root-ca-bundle  
...output omitted...
```

1.6. Create the Argo CD instance.

```
[student@workstation ~]$ oc create -f \  
~/DO380/labs/gitops-app/argocd-instance.yaml  
argocd.argoproj.io/argocd created
```

- 2. Create the etherpad-devs project and label it as managed by the Argo CD instance. Assign administrative permission to users in the `project-admins` group.

2.1. Create the etherpad-devs project.

```
[student@workstation ~]$ oc new-project etherpad-devs  
Now using project "etherpad-devs" on server "https://api.ocp4.example.com:6443".  
...output omitted...
```

2.2. Label the etherpad-devs project as managed by the Argo CD instance in the gitops-app project.

```
[student@workstation ~]$ oc label namespace etherpad-devs \  
argocd.argoproj.io/managed-by=gitops-app  
namespace/etherpad-devs labeled
```

2.3. Assign administrative permission to the etherpad-devs project for users in the `project-admins` group.

```
[student@workstation ~]$ oc adm policy add-role-to-group admin \  
project-admins -n etherpad-devs  
clusterrole.rbac.authorization.k8s.io/admin added: "project-admins"
```

2.4. Log in as the developer user.

```
[student@workstation ~]$ oc login -u developer -p developer  
...output omitted...
```

2.5. Verify that the developer user has no access to the etherpad-devs project.

```
[student@workstation ~]$ oc projects  
You are not a member of any projects. You can request a project to be created with  
the 'new-project' command.
```

2.6. Log in as the `project-admin` user.

```
[student@workstation ~]$ oc login -u project-admin -p redhat  
...output omitted...
```

- 3. As the `project-admin` user in OpenShift, create the `mariadb` secret with the MariaDB database credentials. Create the `database-backup` persistent volume claim (PVC), for a later step to back up the MariaDB database.
- 3.1. Create the `mariadb` secret YAML file that contains the MariaDB credentials. You can find an incomplete example for the secret in the `~/D0380/labs/gitops-app/secret.yaml` file.

```
apiVersion: v1
kind: Secret
metadata:
  name: mariadb
  namespace: etherpad-devs
  labels:
    app: mariadb
stringData:
  MARIADB_DATABASE: etherpad_lite_db
  MARIADB_USER: appuser
  MARIADB_PASSWORD: securepassword
  MARIADB_ROOT_PASSWORD: supersecurepassword
```

- 3.2. Create the `mariadb` secret.

```
[student@workstation ~]$ oc create -f \
~/D0380/labs/gitops-app/secret.yaml
secret/mariadb created
```

- 4. As the `project-admin` user, create an RBAC file in the `etherpad-admin` repository in GitLab to give the `developer` user read permission in the `etherpad-devs` project. The `project-admin` user uses this repository in Argo CD in a later step to give read permission to the `developer` user.
- 4.1. Create the RBAC YAML file to give the `developer` user read permission in the `etherpad-devs` project. You can find an incomplete YAML example in the `~/D0380/labs/gitops-app/etherpad-admin/rbac.yaml` file.

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: developer-view
  namespace: etherpad-devs
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: view
subjects:
- kind: Group
  apiGroup: rbac.authorization.k8s.io
  name: project-devs
```

- 4.2. Open a web browser and navigate to `https://git.ocp4.example.com`. Log in as the `project-admin` user with `r3dh4tgit` as the password.
- 4.3. Click the `etherpad-admin` project.

Chapter 5 | OpenShift GitOps

- 4.4. Click the **Upload File** button and then click **upload**. Select the `~/DO380/labs/gitops-app/etherpad-admin/rbac.yaml` file, and then click the **Open** button. Click **Upload file**.
- 5. Open the Argo CD instance as the **project-admin** user and create the RBAC rule for the **developer** user.
- 5.1. Change to the terminal window and log in as the **admin** user.

```
[student@workstation ~]$ oc login -u admin -p redhatocp
...output omitted...
```

- 5.2. Verify the route for the Argo CD instance.

```
[student@workstation ~]$ oc get route -n gitops-app
NAME          HOST/PORT           PATH      ...
argocd-server  argocd-server-gitops-app.apps.ocp4.example.com ...
```

- 5.3. Open a web browser tab and navigate to `https://argocd-server-gitops-app.apps.ocp4.example.com`.
- 5.4. Click **LOG IN VIA OPENSHIFT** and log in as the **project-admin** user with **redhat** as the password, by using the **Red Hat Identity Management** identity provider, and allowing the **user:info** permission.
- 5.5. Click **CREATE APPLICATION**.
- 5.6. Create an Argo CD application with the information in the following table. Then, click **CREATE**.

Field	Value
Application Name	<code>rbac-rule</code>
Project Name	<code>default</code>
Sync Policy	<code>Automatic</code>
Repository URL	<code>https://git.ocp4.example.com/project-admin/etherpad-admin.git</code>
Path	<code>.</code>
Cluster URL	<code>https://kubernetes.default.svc</code>
Namespace	<code>etherpad-devs</code>

- 5.7. Change to the terminal window and log in as the **developer** user. The **developer** user has access to the `etherpad-devs` project.

```
[student@workstation ~]$ oc login -u developer -p developer
...output omitted...
Using project "etherpad-devs".
```

- 6. Create the Etherpad application by using the Argo CD instance and verify that it works. The `project-admin` user creates the application by using the `developer` user repository. Thus, although the `developer` user cannot modify the application in Argo CD or in OpenShift, the `developer` user can change the files in the repository to update the application.
- 6.1. Go back to the Argo CD browser tab.
 - 6.2. Click **NEW APP**.
 - 6.3. Create an application with the information in the following table:

Field	Value
Application Name	<code>etherpad-app</code>
Project Name	<code>default</code>
Sync Policy	<code>Automatic</code>
Repository URL	https://git.ocp4.example.com/developer/etherpad-app.git
Path	<code>.</code>
Cluster URL	https://kubernetes.default.svc
Namespace	<code>etherpad-devs</code>

Then, click **CREATE**.

- 6.4. Click `etherpad-app` to view the application. Argo CD starts synchronizing the application. After about one minute, the console shows the Etherpad application as synchronized and healthy.
- 6.5. Change to the terminal window and get the URL for the Etherpad application. The `developer` user has read access to the `etherpad-devs` project in OpenShift.

```
[student@workstation ~]$ oc get route
NAME      HOST/PORT          PATH      ...
etherpad   etherpad-etherpad-devs.apps.ocp4.example.com ...
```

- 6.6. Open a web browser tab and navigate to <https://etherpad-etherpad-devs.apps.ocp4.example.com> to verify that the application is up and running. Notice the D0380 - etherpad title in the tab.

► 7. As the `developer` user in Git, create two Argo CD hooks in the `etherpad-app` Git repository. The `developer` user Git credentials are already configured in the classroom environment. The pre-synchronization hook must back up the database in the `database-backup` PVC. The post-synchronization hook verifies that the Etherpad application is running. Modify the Etherpad application title.

 - 7.1. Change to the terminal window, and change to the `~/D0380/labs/gitops-app/etherpad-app` directory. The Git repository is synchronized to that directory.

Chapter 5 | OpenShift GitOps

```
[student@workstation ~]$ cd ~/DO380/labs/gitops-app/etherpad-app
```

- 7.2. Copy the provided hook files to the repository.

```
[student@workstation etherpad-app]$ cp ./hooks/{presync,postsync}.yaml .
```

- 7.3. Edit the pre-synchronization hook file to match the following text. The pre-synchronization hook uses the `mysqldump` command to back up the database in the `database-backup` PVC.

```
apiVersion: batch/v1
kind: Job
metadata:
  generateName: backup-mariadb
  annotations:
    argocd.argoproj.io/hook: PreSync
...output omitted...
```

- 7.4. Edit the post-synchronization hook file to match the following text. The post-synchronization hook uses the `curl` command to ensure that the Etherpad application is up and running.

```
apiVersion: batch/v1
kind: Job
metadata:
  generateName: test-etherpad
  annotations:
    argocd.argoproj.io/hook: PostSync
spec:
  template:
    metadata:
      name: test-etherpad
    spec:
      containers:
        - name: test-etherpad
          image: registry.ocp4.example.com:8443/redhattraining/mariadb:10.5
          command: ["curl", "-k", "-s", "https://etherpad-etherpad-
devs.apps.ocp4.example.com"] ①
          restartPolicy: Never
          backoffLimit: 2
```

① This command should be written in a single line.

- 7.5. Open the `etherpad-deployment.yaml` file and modify the Etherpad application title.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: etherpad
  labels:
```

Chapter 5 | OpenShift GitOps

```
app.kubernetes.io/name: etherpad
spec:
...output omitted...
spec:
  securityContext:
    runAsNonRoot: true
    seccompProfile:
      type: RuntimeDefault
  containers:
    - env:
      - name: TITLE
        value: My test Etherpad app
...output omitted...
```

- 7.6. Add the hook and deployment files to the Git index.

```
[student@workstation etherpad-app]$ git add presync.yaml postsync.yaml \
etherpad-deployment.yaml
```

- 7.7. Commit the changes.

```
[student@workstation etherpad-app]$ git commit -m \
"Add synchronization hooks and change title"
...output omitted...
```

- 7.8. Push the changes to the repository.

```
[student@workstation etherpad-app]$ git push
...output omitted...
```

- 8. As the developer user in Argo CD, verify that the application is updated to the last repository version with the two hooks.
- 8.1. Change to the Argo CD browser tab and click **Log out**.
 - 8.2. Click **LOG IN VIA OPENSHIFT** and log in as the developer user with developer as the password, by using the **Red Hat Identity Management** identity provider, and allowing the **user:info** permission.
 - 8.3. Click **etherpad-app** to view the application.
 - 8.4. Argo CD automatically detects the changes in the repository and starts synchronizing the application.

**Note**

The Argo CD automatic synchronization interval is set to 3 minutes by default. Thus, if the last synchronization does not match the last commit, then you can click **REFRESH** so Argo CD detects the changes.

If the Argo CD application is in the **Syncing** stage, then wait until Argo CD finishes synchronizing it.

- 8.5. Verify that Argo CD creates the `backup-mariadb` and `test-etherpad` jobs according to the hooks. To inspect the logs for the hooks, click the job name and then click `LOGS`. The `backup-mariadb` pre-synchronization hook connects to and backs up the MariaDB database in the `database-backup` PVC. The `test-etherpad` post-synchronization hook uses the `curl` command to verify that the Etherpad application is up and running.

- 8.6. Change to the terminal window and log in as the `project-admin` user.

```
[student@workstation etherpad-app]$ oc login -u project-admin -p redhat  
...output omitted...
```

- 8.7. Create a pod to inspect the contents of the `database-backup` PVC. You can use the `~/DO380/labs/gitops-app/backup-inspector.yaml` file for this purpose.

```
[student@workstation etherpad-app]$ oc create -f ./backup-inspector.yaml  
pod/backup-inspector created
```

- 8.8. Wait until the `backup-inspector` pod is running and open a remote shell on it.

```
[student@workstation etherpad-app]$ oc rsh backup-inspector
```

- 8.9. Review the contents of the `external-data` directory. The pre-synchronization hook creates a database backup every time that the application is updated to a new version.

```
sh-5.1$ ls /external-data/  
database-backup-files-202312201021 lost+found
```

- 8.10. Close the remote shell.

```
sh-5.1$ exit
```

- 8.11. Remove the `backup-inspector` pod.

```
[student@workstation etherpad-app]$ oc delete pod backup-inspector  
pod "backup-inspector" deleted
```

- 8.12. Change to the Etherpad web browser tab and reload it. The tab title must change to `My test Etherpad app`.

- ▶ 9. Close the web browser and change to the student HOME directory in the terminal window.

```
[student@workstation etherpad-app]$ cd
```

Finish

On the `workstation` machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish gitops-app
```

► Lab

OpenShift GitOps

Define the fundamentals of GitOps and its use with Kubernetes clusters and applications.

Deploy a GitOps instance for cluster administration.

Deploy another GitOps instance for application developers and administrators.

Outcomes

- Create an Argo CD instance for developers with the appropriate permissions.
- Use GitOps practices to create a cron job.

Before You Begin

As the student user on the workstation machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start gitops-review
```

Instructions

Some teams in your organization execute periodic tasks manually. Your organization wants to evaluate automating the tasks with OpenShift and GitOps practices.

In this exercise, you install the Red Hat OpenShift GitOps operator for this purpose. You create a custom Argo CD instance, a Git repository with a sample periodic task, and an Argo CD application.

Developers can create cron jobs only by pushing them to the Git repository. Developers must be able to troubleshoot the cron jobs.

1. Install the OpenShift GitOps operator from OperatorHub. Use the `admin` user with `redhatocp` as the password.
2. Create the Argo CD instance in the `gitops-review` namespace. The Argo CD must trust the classroom certificate. Grant read permission to applications, projects, and clusters to the users in the `project-devs` group.
You can find an incomplete example for the Argo CD CR in the `~/DO380/labs/gitops-review/argocd-instance.yaml` file.
3. Create the `project-devs` group in OpenShift, and add the `developer` user to it. Grant the `project-devs` group the `view` role to the `gitops-review` project.
4. Log in to the Argo CD web console as the `admin` user by using their OpenShift credentials. Create an Argo CD application by using the `cron-job` repository from the `developer` user. Use the information in the following table:

Field	Value
Application Name	cron-job
Project Name	default
SYNC POLICY	Automatic
Retry	Checked
Repository URL	https://git.ocp4.example.com/developer/cron-job.git
Path	.
Cluster URL	https://kubernetes.default.svc

5. As the `developer` user, create a `periodic-process` cron job in the Git repository. The URL for the Git instance is <https://git.ocp4.example.com>. The password for the `developer` user in Git is `d3v3lop3r`. The `developer` user Git credentials are already configured in the classroom environment.

The cron job runs the `echo hello` command every minute with the `registry.ocp4.example.com:8443/ubi9/ubi` image.

You can use the following command to generate a template for the cron job:

```
[student@workstation ~]$ oc create cronjob periodic-process \
-n gitops-review --schedule "* * * * *" \
--image registry.ocp4.example.com:8443/ubi9/ubi \
--dry-run=client -o yaml -- echo hello
```

6. As the `developer` user in Argo CD, verify that the application is updated to the last repository version with the cron job. The password for the `developer` user in OpenShift is `developer`. Verify that the user can review the application logs in Argo CD.
7. As the `developer` user in OpenShift, verify that the developer user can retrieve the logs for the cron job pods.

Evaluation

As the `student` user on the `workstation` machine, use the `lab` command to grade your work. Correct any reported failures and rerun the command until successful.

```
[student@workstation ~]$ lab grade gitops-review
```

Finish

As the `student` user on the `workstation` machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish gitops-review
```

► Solution

OpenShift GitOps

Define the fundamentals of GitOps and its use with Kubernetes clusters and applications.

Deploy a GitOps instance for cluster administration.

Deploy another GitOps instance for application developers and administrators.

Outcomes

- Create an Argo CD instance for developers with the appropriate permissions.
- Use GitOps practices to create a cron job.

Before You Begin

As the student user on the workstation machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start gitops-review
```

Instructions

Some teams in your organization execute periodic tasks manually. Your organization wants to evaluate automating the tasks with OpenShift and GitOps practices.

In this exercise, you install the Red Hat OpenShift GitOps operator for this purpose. You create a custom Argo CD instance, a Git repository with a sample periodic task, and an Argo CD application.

Developers can create cron jobs only by pushing them to the Git repository. Developers must be able to troubleshoot the cron jobs.

1. Install the OpenShift GitOps operator from OperatorHub. Use the `admin` user with `redhatocp` as the password.
 - 1.1. Use the command line to log in to the OpenShift cluster as the `admin` user with `redhatocp` as the password.

```
[student@workstation ~]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
...output omitted...
```

- 1.2. Identify the URL for the OpenShift web console.

```
[student@workstation ~]$ oc whoami --show-console
https://console-openshift-console.apps.ocp4.example.com
```

- 1.3. Open a web browser and navigate to `https://console-openshift-console.apps.ocp4.example.com`. Either type the URL in a web browser, or right-click and select **Open Link** from the command line.
- 1.4. Click **Red Hat Identity Management** and log in as the `admin` user with `redhatocp` as the password.
- 1.5. Navigate to **Operators > OperatorHub**.
- 1.6. Click **Red Hat OpenShift GitOps**, and then click **Install**.
- 1.7. Review the default configuration and click **Install**. The Operator Lifecycle Manager can take a few minutes to install the operator. Click **View Operator** to navigate to the operator details.
2. Create the Argo CD instance in the `gitops-review` namespace. The Argo CD must trust the classroom certificate. Grant read permission to applications, projects, and clusters to the users in the `project-devs` group.

You can find an incomplete example for the Argo CD CR in the `~/D0380/labs/gitops-review/argocd-instance.yaml` file.

- 2.1. Change to the command line, and change to the `gitops-review` project.

```
[student@workstation ~]$ oc project gitops-review
Now using project "gitops-review" on server "https://api.ocp4.example.com:6443".
```

- 2.2. Create a `cluster-root-ca-bundle` configuration map in the `gitops-review` namespace.

```
[student@workstation ~]$ oc create configmap cluster-root-ca-bundle
configmap/cluster-root-ca-bundle created
```

- 2.3. Add the `config.openshift.io/inject-trusted-cabundle` label to the configuration map with the `true` value. OpenShift injects the bundle with the cluster certificate authority to the configuration maps with this label. This bundle contains the signing certificate for the classroom GitLab instance.

```
[student@workstation ~]$ oc label configmap cluster-root-ca-bundle \
  config.openshift.io/inject-trusted-cabundle=true
configmap/cluster-root-ca-bundle labeled
```

- 2.4. Create the Argo CD instance YAML file. Set the termination type of the route to the `reencrypt` type. Grant read permission to applications, projects, and clusters to the `project-devs` groups. The Argo CD instance must mount the `ca-bundle.crt` certificate in the configuration map to the `/etc/pki/ca-trust/extracted/pem/tls-ca-bundle.pem` path of the repository server container. You can find an incomplete example for the Argo CD CR in the `~/D0380/labs/gitops-review/argocd-instance.yaml` file.

```
apiVersion: argoproj.io/v1beta1
kind: ArgocD
metadata:
```

```

name: argocd
namespace: gitops-review
spec:
  server:
    ...output omitted...
  route:
    enabled: true
  tls:
    termination: reencrypt
...output omitted...
rbac:
  defaultPolicy: ''
  policy: |
    g, system:cluster-admins, role:admin
    g, project-devs, role:project-devs
    p, role:project-devs, applications, get, /*, allow
    p, role:project-devs, projects, get, *, allow
    p, role:project-devs, clusters, get, *, allow
  scopes: '[groups]'
repo:
  resources:
    limits:
      cpu: 1000m
      memory: 1024Mi
    requests:
      cpu: 250m
      memory: 256Mi
  volumeMounts:
    - mountPath: /etc/pki/ca-trust/extracted/pem/tls-ca-bundle.pem
      name: cluster-root-ca-bundle
      subPath: ca-bundle.crt
  volumes:
    - configMap:
        name: cluster-root-ca-bundle
        name: cluster-root-ca-bundle
...output omitted...

```

2.5. Create the Argo CD instance.

```
[student@workstation ~]$ oc create -f \
~/DO380/labs/gitops-review/argocd-instance.yaml
argocd.argoproj.io/argocd created
```

3. Create the `project-devs` group in OpenShift, and add the `developer` user to it. Grant the `project-devs` group the `view` role to the `gitops-review` project.

3.1. Create the `project-devs` group.

```
[student@workstation ~]$ oc adm groups new project-devs
```

3.2. Add the `developer` user to the `project-devs` group.

```
[student@workstation ~]$ oc adm groups add-users project-devs developer
```

Chapter 5 | OpenShift GitOps

- 3.3. Add the `view` role to the `gitops-review` project to users in the `project-devs` group.

```
[student@workstation ~]$ oc adm policy add-role-to-group view project-devs \
-n gitops-review
```

4. Log in to the Argo CD web console as the `admin` user by using their OpenShift credentials. Create an Argo CD application by using the `cron-job` repository from the `developer` user. Use the information in the following table:

Field	Value
Application Name	<code>cron-job</code>
Project Name	<code>default</code>
SYNC POLICY	<code>Automatic</code>
Retry	Checked
Repository URL	https://git.ocp4.example.com/developer/cron-job.git
Path	.
Cluster URL	https://kubernetes.default.svc

- 4.1. Verify the route for the Argo CD instance.

```
[student@workstation ~]$ oc get route -n gitops-review
NAME           HOST/PORT          PATH      ...
argocd-server  argocd-server-gitops-review.apps.ocp4.example.com ...
```

- 4.2. Open a web browser tab and navigate to <https://argocd-server-gitops-review.apps.ocp4.example.com>
- 4.3. Click **LOG IN VIA OPENSHIFT** and log in as the `admin` user with `redhatocp` as the password, by using the **Red Hat Identity Management** identity provider, and allowing the `user:info` permission.
- 4.4. Click **CREATE APPLICATION**.
- 4.5. Fill the application information with the details in the table. Then, click **CREATE**.
- 4.6. Wait until Argo CD shows the `cron-job` application as synchronized and healthy. Then, click **Log out**.
5. As the `developer` user, create a `periodic-process` cron job in the Git repository. The URL for the Git instance is <https://git.ocp4.example.com>. The password for the `developer` user in Git is `d3v3lop3r`. The `developer` user Git credentials are already configured in the classroom environment.
The cron job runs the `echo hello` command every minute with the `registry.ocp4.example.com:8443/ubi9/ubi` image.
You can use the following command to generate a template for the cron job:

Chapter 5 | OpenShift GitOps

```
[student@workstation ~]$ oc create cronjob periodic-process \
-n gitops-review --schedule "* * * * *" \
--image registry.ocp4.example.com:8443/ubi9/ubi \
--dry-run=client -o yaml -- echo hello
```

- 5.1. In the web browser, open a tab navigate to the Git URL at <https://git.ocp4.example.com>. Log as the developer user with d3v3lop3r as the password. Click the cron-job project. Click Clone, and then copy the <https://git.ocp4.example.com/admin/gitops-review.git> HTTPS URL.
- 5.2. Change to the command line. Change to the ~/D0380/labs/gitops-review directory.

```
[student@workstation ~]$ cd ~/D0380/labs/gitops-review
```

- 5.3. Run the following command to clone the cron-job repository.

```
[student@workstation gitops-review]$ git clone \
https://git.ocp4.example.com/developer/cron-job.git
Cloning into 'cron-job'...
...output omitted...
```

- 5.4. Change to the cloned repository directory.

```
[student@workstation gitops-review]$ cd cron-job
```

The default configuration for new repositories adds a README.md initial file.

- 5.5. Create a periodic-process.yaml file with the following content:

```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: periodic-process
  namespace: gitops-review
spec:
  jobTemplate:
    metadata:
      name: periodic-process
    spec:
      template:
        spec:
          containers:
            - command:
              - echo
              - hello
            image: registry.ocp4.example.com:8443/ubi9/ubi
            name: periodic-process
            restartPolicy: OnFailure
  schedule: '* * * * *'
```

You can create the file with the output of the previous command, and removing the parts that are not present in the previous content.

- 5.6. Add the `periodic-process.yaml` file to the repository.

```
[student@workstation cron-job]$ git add periodic-process.yaml
```

- 5.7. Commit the changes.

```
[student@workstation cron-job]$ git commit -m "add job"  
...output omitted...
```

- 5.8. Push the changes.

```
[student@workstation cron-job]$ git push  
...output omitted...
```

6. As the `developer` user in Argo CD, verify that the application is updated to the last repository version with the cron job. The password for the `developer` user in OpenShift is `developer`. Verify that the user can review the application logs in Argo CD.

- 6.1. Change to the Argo CD browser tab.
- 6.2. Click **LOG IN VIA OPENSHIFT** and log in as the `developer` user with `developer` as the password, by using the **Red Hat Identity Management** identity provider, and allowing the `user:info` permission.
- 6.3. Click `cron-job` to view the application.
- 6.4. Argo CD automatically detects the changes in the repository and starts synchronizing the application.



Note

The Argo CD automatic synchronization interval is set to 3 minutes by default. Thus, if the last synchronization does not match the last commit, then you can click **REFRESH** so Argo CD detects the changes.

If the Argo CD application is in the **Syncing** stage, then wait until Argo CD finishes synchronizing it.

- 6.5. Verify that Argo CD creates the `periodic-process` cron job. The cron job runs the `echo hello` command every minute. Wait until the cron job creates a pod at least one time. To inspect the logs for the jobs, click the job name and then click **LOGS**. The logs show the word `hello` as the output.
7. As the `developer` user in OpenShift, verify that the developer user can retrieve the logs for the cron job pods.
 - 7.1. Change to the command line and log in as the `developer` user.

```
[student@workstation cron-job]$ oc login -u developer -p developer  
...output omitted...  
Using project "gitops-review".
```

7.2. List the pods for the gitops-review project.

```
[student@workstation cron-job]$ oc get pods  
NAME                      READY   STATUS    RESTARTS   AGE  
argocd-application-controller-0   1/1     Running   0          111m  
argocd-dex-server-85d586c89b-ndxtj   1/1     Running   0          111m  
argocd-redis-7545d85b5-mt98r      1/1     Running   0          111m  
argocd-repo-server-5c8bc5758d-9mtv9   1/1     Running   0          111m  
argocd-server-68cbc7589c-md52w      1/1     Running   0          111m  
periodic-process-28597732-wcrxq     0/1     Completed  0          2m41s  
periodic-process-28597733-z99pk     0/1     Completed  0          101s  
periodic-process-28597734-p9wc4      0/1     Completed  0          41s  
...output omitted...
```

7.3. Check the logs for one of the periodic-process pods.

```
[student@workstation cron-job]$ oc logs periodic-process-28597734-p9wc4  
hello
```

7.4. Close the web browser and change to the /home/student directory in the command line.

```
[student@workstation cron-job]$ cd
```

Evaluation

As the student user on the workstation machine, use the lab command to grade your work. Correct any reported failures and rerun the command until successful.

```
[student@workstation ~]$ lab grade gitops-review
```

Finish

As the student user on the workstation machine, use the lab command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish gitops-review
```

Summary

- With GitOps, files in version control describe Kubernetes resources.
- Red Hat OpenShift GitOps is an operator that helps implement GitOps practices in OpenShift clusters.
- The operator helps you deploy Argo CD instances with declarative configuration.
- Argo CD works with applications, which are Kubernetes resources that reference the repositories with the resource definitions to deploy.
- Argo CD implements retries and sync waves to help deploy resources that pose deployment challenges due to the use of eventual consistency in Kubernetes.
- Argo CD can use the server-side apply process to patch existing resources.
- Managing sensitive data when using GitOps requires additional measures to protect this data.
- Argo CD provides resource hooks to execute actions at different application lifecycle stages.
- Argo CD can roll back applications either on demand, or automatically after events such as a failed hook.

Chapter 6

OpenShift Monitoring

Goal

Troubleshoot performance and availability issues with applications and clusters.

Sections

- Cluster Monitoring (and Guided Exercise)
- Alerts and Notifications (and Guided Exercise)

Lab

- OpenShift Monitoring

Cluster Monitoring

Objectives

- Describe the architecture of OpenShift Monitoring and query the information in its dashboards.

OpenShift Observability

In the cloud computing era, systems are becoming ever more complex. All organizations aim for reliable, efficient, and secure applications and infrastructure. Observability plays a key role in achieving this goal. Observability is the ability to understand a system's or an application's state by collecting and analyzing its output and logs.

Red Hat OpenShift Observability collects logs, traces, events, and system metrics to provide real-time monitoring. Real-time monitoring helps to identify and troubleshoot issues for OpenShift applications and clusters.

Some key features of the Red Hat OpenShift Observability portfolio are as follows:

OpenShift Logging

Red Hat OpenShift Logging aggregates all the logs from the pods and nodes of an OpenShift cluster to a centralized location. Centralized logging improves searching, visualizing, and reporting of data.

OpenShift Monitoring

Red Hat OpenShift Monitoring provides monitoring for core platform components.

Network Observability

Network observability monitors and analyzes network traffic and helps to resolve connectivity issues.

Distributed Tracing

Distributed tracing collects observability data in distributed systems. Distributed tracing is based on the OpenTelemetry project.

For more details about network observability and distributed tracing, refer to the References section.

You can use these features with any stand-alone OpenShift cluster. Red Hat Advanced Cluster Management for Kubernetes (RHACM) provides multicluster observability features.

OpenShift Monitoring

Red Hat OpenShift Container Platform comes with a monitoring stack. The cluster monitoring operator manages the monitoring components and ensures that they are always available and updated. The default monitoring stack collects metrics and generates alerts for the cluster, which includes core platform components and all projects. However, the default monitoring stack does not support custom metrics for user-defined projects.

You can enable monitoring for user-defined projects. You can collect custom metrics and generate alerts by using the monitoring stack for user-defined projects.

Chapter 6 | OpenShift Monitoring

You can configure persistent storage for OpenShift monitoring. With this configuration, you can keep a record of the past cluster status, to investigate and correlate current and past issues within the cluster. The monitoring dashboard provides visuals for cluster metrics.

OpenShift Monitoring Stack

The OpenShift monitoring stack is based on the Prometheus open source project. The stack includes the components that are shown in the following figure and are then explained in the following section:

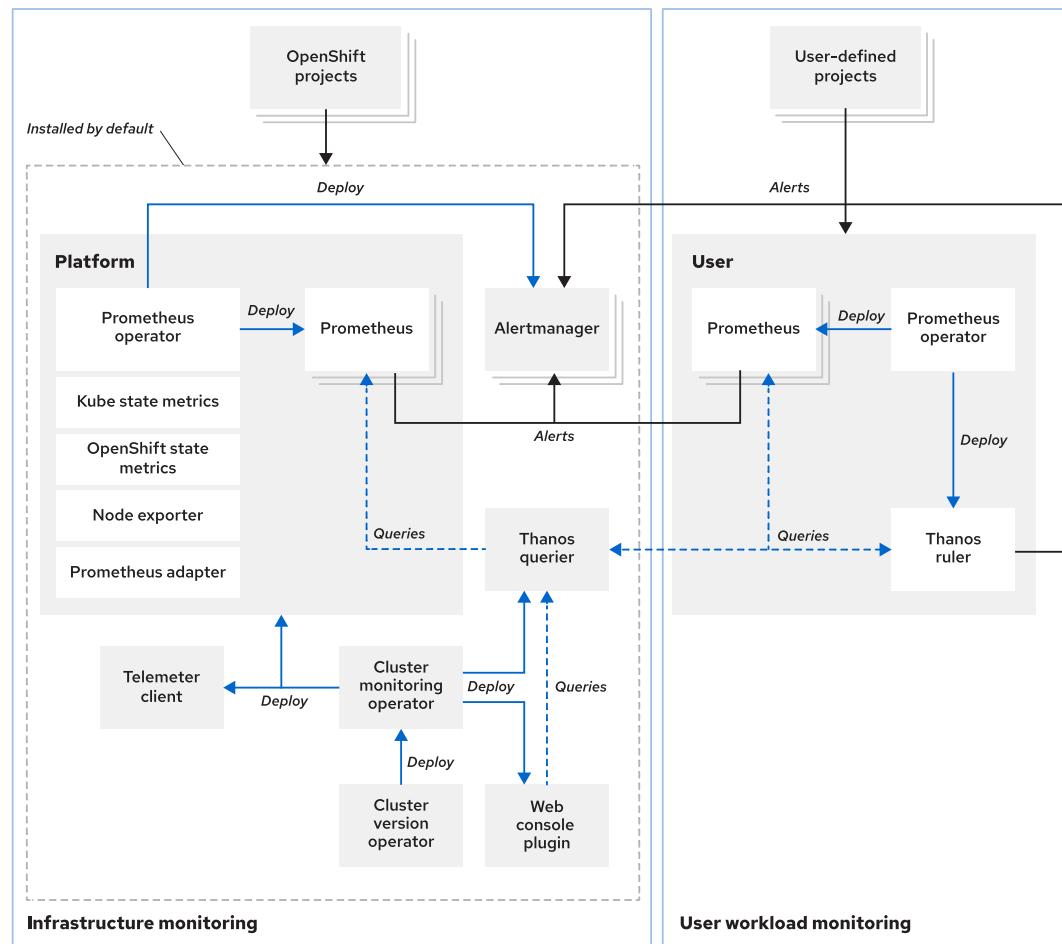


Figure 6.1: OpenShift monitoring architecture

Components for Default Monitoring Stack

The default monitoring stack is included with OpenShift Container Platform. All components of the default monitoring stack are installed in the `openshift-monitoring` project and provide monitoring features for core platform components.

Modifying any existing resource and creating additional `ServiceMonitor`, `PodMonitor`, or `PrometheusRule` resources in the `openshift-monitoring` project is not supported. The monitoring stack resets modified resources to ensure that its resources always remain in the expected state.

The monitoring stack deploys the following components in the environment for monitoring the infrastructure, receiving alerts, and consulting performance graphs.

Chapter 6 | OpenShift Monitoring

Cluster Monitoring Operator

The cluster monitoring operator is the central component of the monitoring stack. The cluster monitoring operator controls the deployed monitoring components and ensures that they are always in sync with the latest version of the cluster monitoring operator.

Prometheus Operator

The Prometheus operator deploys and configures both Prometheus and Alertmanager. The operator also manages the generation of configuration targets (service monitors and pod monitors).

Prometheus

Prometheus is the monitoring server.

Prometheus Adapter

The Prometheus adapter exposes cluster resources for Horizontal Pod Autoscaling (HPA).

Prometheus Alertmanager

Alertmanager handles alerts from the Prometheus server. An alert is a rule that evaluates to true or false and is often based on cluster observations, such as cluster CPU utilisation. An alert fires when the alert rule meets the true condition. You can configure Prometheus Alertmanager to group and route the alerts to the receiver.

Kube state metrics

The kube-state-metrics converter agent exports Kubernetes objects to metrics that Prometheus can parse.

OpenShift state metrics

The openshift-state-metrics agent is based on the kube-state-metrics agent and adds monitoring for OpenShift-specific resources (such as image registry metrics).

Node exporter

The node-exporter agent exports low-level metrics for compute nodes.

Thanos Querier

Thanos Querier is a single, multitenant interface that enables aggregating and deduplicating cluster and user workload metrics.

Telemeter Client

Telemeter Client sends a data portion from Prometheus instances to Red Hat for remote health monitoring.

The monitoring web console

The OpenShift Container Platform web console provides the **Observe** section to access and manage monitoring features. In the Observe section, you can access monitoring dashboards, metrics, alerts, and metrics targets.

Components for Monitoring User-defined Projects

You can enable monitoring for user-defined projects. You can collect application-specific custom metrics and also create alerts by using custom metrics with this monitoring stack. The monitoring components for user-defined projects are installed in the openshift-user-workload-monitoring project.

The monitoring stack for user-defined projects deploys the following components:

- Prometheus Operator
- Prometheus

- Thanos Ruler
- Alertmanager

For more details about the OpenShift monitoring stack for user-defined projects, see the References section.

Prometheus

Prometheus is an open source project for system monitoring and alerting.

Both Red Hat OpenShift Container Platform and Kubernetes integrate Prometheus to enable cluster metrics, monitoring, and alerting capabilities.

Prometheus gathers and stores streams of data from the cluster as *time-series data*. Time-series data consists of a sequence of samples, where each sample contains the following elements:

- A timestamp
- A numeric value (such as an integer, float, or Boolean)
- A set of labels in the form of key/value pairs

The key/value pairs isolate groups of related values for filtering.

For example, the `machine_cpu_cores` metric in Prometheus contains a sequence of measurement samples of the number of CPU cores for each machine.

OpenShift Monitoring Web Console

You can access and manage the monitoring features by using the web console. The OpenShift Container Platform web console provides the **Observe** section, with the following subsections:

- Alerting
- Metrics
- Dashboards
- Targets

Cluster Monitoring Alerting

You access cluster alerts from the OpenShift web console at **Observe > Alerting**. For each alert, the **Alerting** page displays a brief description, the state, and the severity. You can view alert details by clicking the name of the alert. The **Alert details** page also displays a time-series graphic.

The screenshot shows the 'Alerting' page in the OpenShift web interface. At the top, there are tabs for 'Alerts', 'Silences', and 'Alerting rules'. Below the tabs is a search bar with filters for 'Source' (set to 'Platform'), 'Alert State' (set to 'Firing'), and a 'Clear all filters' button. The main area displays a table of alerts:

Name	Severity	State	Source	...
AL_AlertmanagerReceiversNotConfigured	⚠ Warning	⚠ Firing Since Feb 8, 2024, 1:16 AM	Platform	...
AL_InsightsDisabled	ⓘ Info	⚠ Firing Since Feb 8, 2024, 7:45 AM	Platform	...
AL_KubeDaemonSetMisScheduled	⚠ Warning	⚠ Firing Since Feb 8, 2024, 1:17 AM	Platform	...
AL_KubeDaemonSetRolloutStuck	⚠ Warning	⚠ Firing Since Feb 8, 2024, 2:54 AM	Platform	...
AL_Watchdog	ⓘ None	⚠ Firing Since Feb 8, 2024, 7:45 AM	Platform	...

Figure 6.2: Cluster monitoring alerting

For more details about forwarding alerts to other systems, see the next section.

Cluster Monitoring Metrics

OpenShift integrates Prometheus metrics at **Observe > Metrics**.

From the **Metrics** page, enter an expression, such as a metric name, and then click **Run Queries** to retrieve the most recent sample for the metric.

The following example displays the `instance:node_cpu_utilisation:rate1m` metric over time.

Chapter 6 | OpenShift Monitoring

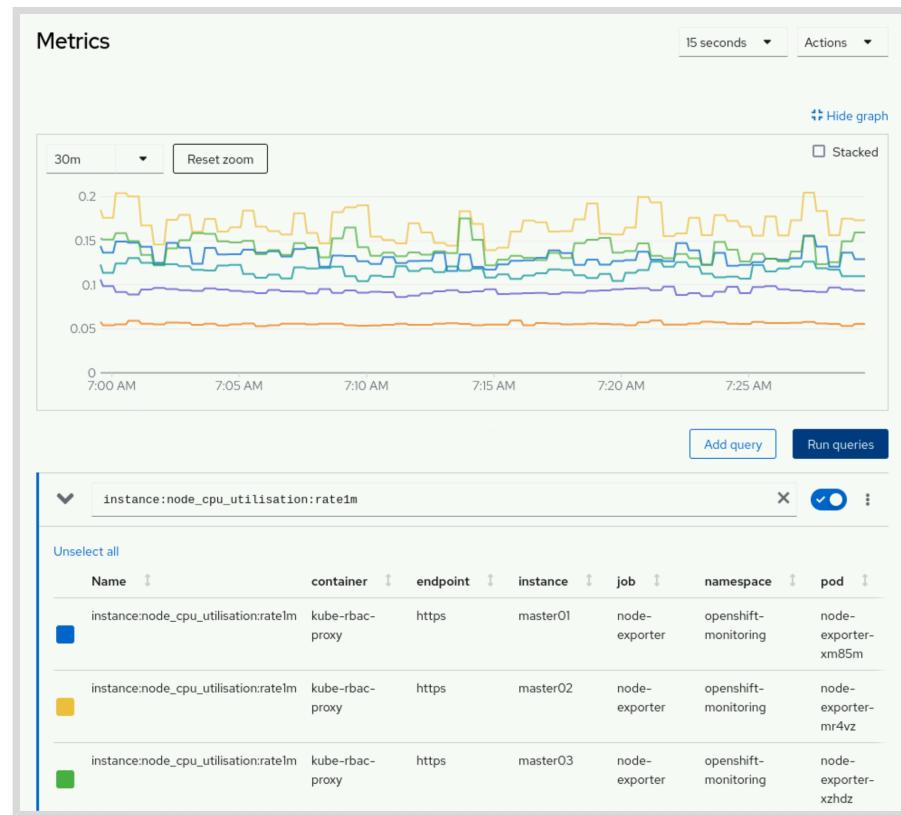


Figure 6.3: Monitoring metrics

The metric contains data for each node instance in the cluster.

The OpenShift has three monitoring stack components to gather the metrics from the Kubernetes API: the `kube-state-metrics`, `openshift-state-metrics`, and `node-exporter` agents.

The dashboards in OpenShift cluster monitoring combine metrics from the three agents.

See the References section to learn about the complete list of exposed metrics.

Prometheus provides a query language, PromQL, to select and aggregate time-series data.

You can filter a metric to include only certain key/value pairs. For example, you can modify the previous query to show only metrics for the `worker02` node by using the following expression:

```
instance:node_cpu_utilisation:rate1m{instance="worker02"}
```

Prometheus Query Language provides several operators to compute new time-series metrics. PromQL contains arithmetic operators, including addition, subtraction, multiplication, and division operators. PromQL contains comparison operators, including equality, greater-than, and less-than operators.

PromQL contains built-in functions, including the following ones, that you can include in PromQL expressions:

`sum()`

Adds the value of all sample entries at a given time.

`rate()`

Computes the per-second average of a time series for a given time range.

Chapter 6 | OpenShift Monitoring

`count()`
Counts the number of sample entries at a given time.

`max()`
Selects the maximum value out of the sample entries.

The following examples of Prometheus Query Language expressions use one metric from the `node-exporter` agent, and another metric from the `kube-state-metrics` agent:

`node_memory_MemAvailable_bytes/node_memory_MemTotal_bytes*100<50`
Shows nodes with less than 50% of available memory.

`kube_persistentvolumeclaim_status_phase{phase="Pending"} == 1`
Shows persistent volume claims in the pending state.

The Red Hat add-on operators can define extra metrics and alerts. For example, the `compliance` operator exposes additional metrics to Prometheus. You can get a list of exposed metrics by using the following query:

```
{name=~"compliance.*"}
```

Cluster Monitoring Dashboards

The OpenShift console integrates dashboards based on the gathered metrics at **Observe > Dashboards**. These dashboards refresh periodically to display current summary metrics and graphs.

In the graphs, which are interactive, you can further explore data features and characteristics that you observe.

The cluster monitoring dashboards serve as a good starting point for near real-time observability of cluster metrics and health.

After receiving an alert, an administrator might use the dashboards to investigate the problem. This investigation might include determining whether a specific node or project has a problem. Additionally, cluster monitoring dashboards can help identify whether a problem was temporary or appears to be persistent.

OpenShift cluster monitoring includes several default dashboards.

Some of the default monitoring default dashboards are as follows:

Kubernetes / Compute Resources / Cluster

This dashboard displays a high-level view of cluster resources. The `Kubernetes / Compute Resources / Cluster` dashboard page shows percentage values for CPU such as `CPU Utilisation`, `CPU Requests Commitment`, and `CPU Limits Commitment`. Similar values are also available for memory.

Chapter 6 | OpenShift Monitoring

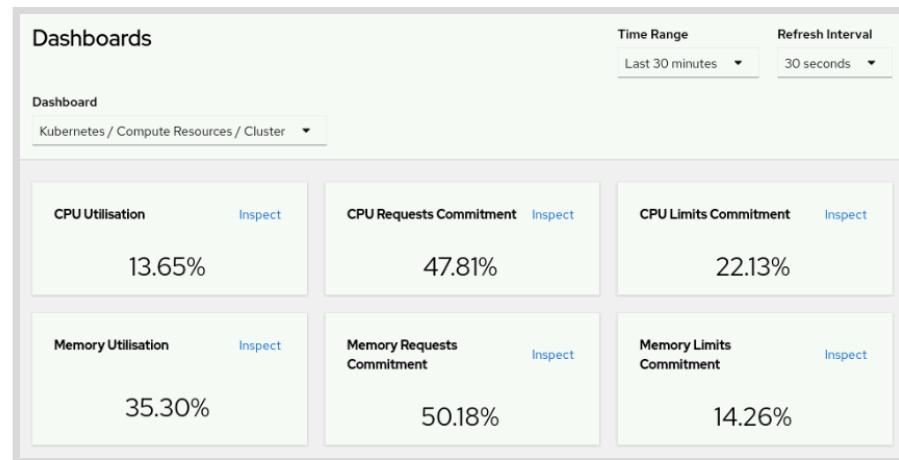


Figure 6.4: Kubernetes / Compute Resources / Cluster dashboard

You can see metrics by clicking **Inspect** for each parameter. Clicking **Inspect** shows the **Metrics** page where you can see metrics and a related graph.

For example, clicking **Inspect** for **CPU Utilisation** shows a graph and values for the following metrics:

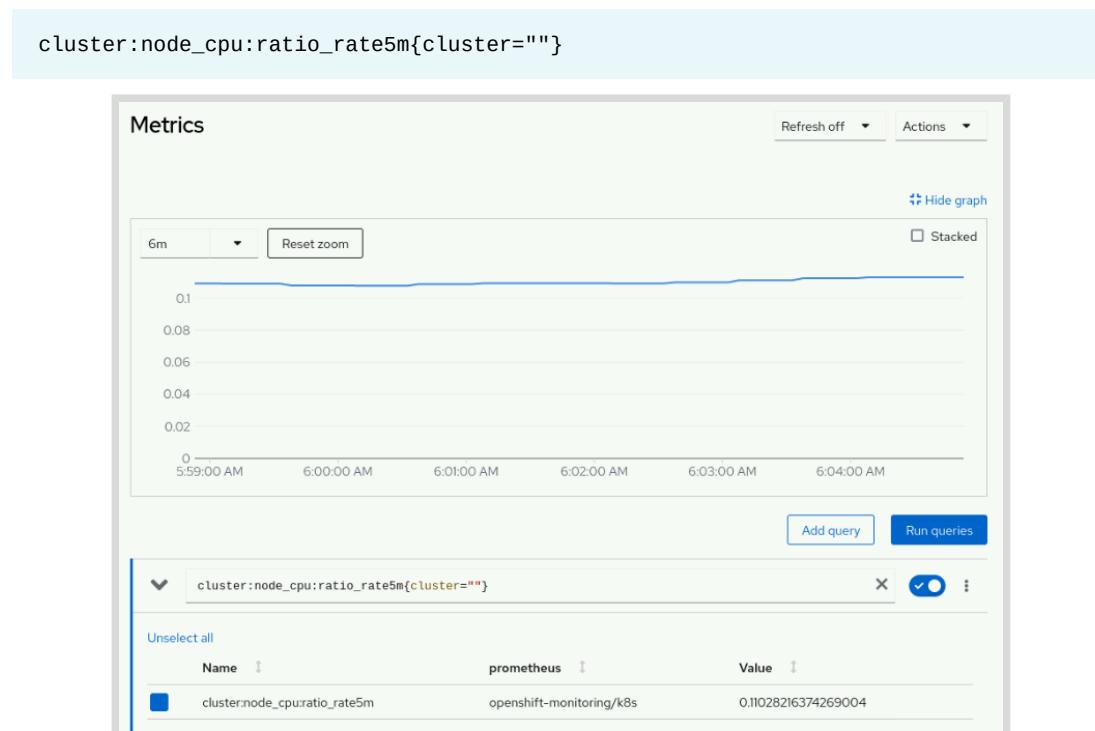


Figure 6.5: Metrics CPU utilisation

The **Kubernetes / Compute Resources / Cluster** dashboard page also shows graphs for CPU, memory, and network, such as **CPU Usage**, **CPU Quota**, **Memory Usage**, and **Memory Quota**.

These graphs are common to some dashboard pages. The only difference is data filtration. For example, the **Kubernetes / Compute Resources / Namespace (Workloads)** dashboard filters resource usage, first by namespace and then by workload type, such as by deployment, daemon set, and stateful set.

USE Method / Cluster

USE stands for *Utilisation Saturation and Errors*. This dashboard displays several graphics to identify whether the cluster is overutilised, oversaturated, or experiencing many errors. Because the dashboard displays all nodes in the cluster, you might be able to identify a node that is not behaving in the same way as the other nodes in the cluster.

The following graphic indicates that the `worker03` node is experiencing higher memory saturation than other nodes in the cluster.

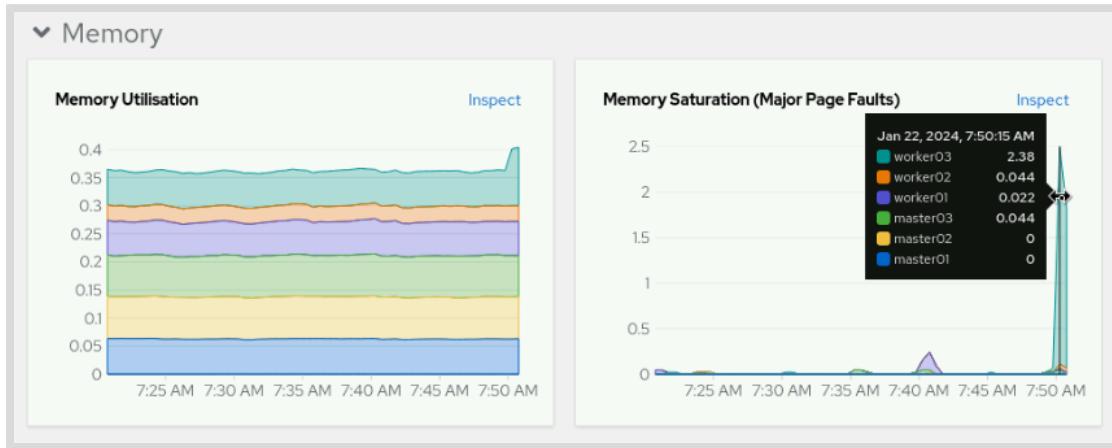


Figure 6.6: USE Method / Cluster - Memory Saturation



References

For more information about OpenShift monitoring, refer to the *Monitoring* chapter in the Red Hat OpenShift Container Platform 4.14 *Observability* documentation at https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/monitoring/index#monitoring-overview

For more information about enabling monitoring for user-defined projects, refer to the *Enabling Monitoring for User-defined Projects* chapter in the Red Hat OpenShift Container Platform 4.14 *Observability* documentation at https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/monitoring/index#enabling-monitoring-for-user-defined-projects

Red Hat OpenShift Observability

<https://www.redhat.com/en/technologies/cloud-computing/openshift/observability>

Check Out the new Network Observability Support in OpenShift 4.12

<https://www.redhat.com/en/blog/check-out-the-new-network-observability-support-in-openshift-4.12>

Observability Across OpenShift Cluster Boundaries with Distributed Data Collection

<https://www.redhat.com/en/blog/observability-across-openshift-cluster-boundaries-with-distributed-data-collection>

Prometheus Overview

<https://prometheus.io/docs/introduction/overview>

Prometheus Data Model

https://prometheus.io/docs/concepts/data_model

Node Exporter

https://github.com/prometheus/node_exporter

Exposed Metrics - kube-state-metrics

<https://github.com/kubernetes/kube-state-metrics/tree/master/docs#exposed-metrics>

Exposed Metrics - openshift-state-metrics

<https://github.com/openshift/openshift-state-metrics/blob/master/docs/README.md#exposed-metrics>

Querying Prometheus

<https://prometheus.io/docs/prometheus/latest/querying/basics/#querying-prometheus>

► Guided Exercise

Cluster Monitoring

Extract specific insights from monitoring dashboards and metrics queries to troubleshoot performance and availability issues with cluster nodes.

Outcomes

- Use the OpenShift monitoring default stack to identify a deployment that consumes excessive CPU and memory.
- Use the OpenShift monitoring default stack to find the cause of an availability issue.

Before You Begin

As the student user on the workstation machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start monitoring-cluster
```

Instructions

Your company has an OpenShift cluster for development and testing environments with three compute nodes. The cluster has two node pools.

The first node pool is for the development environment and uses the `env=dev` label. In this classroom environment, the node pool for the development environment has one `worker03` compute node. The development environment node pool has the `dev-monitor` and `dev-finance` namespaces. The development applications are deployed in the `dev-monitor` and `dev-finance` namespaces. The latest performance testing for each application suggests that with the given load, the applications use minimal resources and do not exceed 100% of the CPU and memory requests. Use the OpenShift monitoring default stack to identify whether a user workload consumes excessive CPU and memory. The load generator `traffic-simulator.sh` script is available at the `~/DO380/labs/monitoring-cluster` location.

The second node pool is for the testing environment and uses the `env=test` label. In this classroom environment, the node pool for the testing environment has one `worker02` compute node. The testing environment node pool has a `test-monitor` namespace. The `frontend-test` and `budget-test` test applications are deployed in the `test-monitor` namespace. The URLs for both applications is as follows:

```
frontend-test
  • http://frontend-test-monitor.apps.ocp4.example.com

budget-test
  • http://budget-test-monitor.apps.ocp4.example.com
```

Both applications are giving an `Application Not Available` message. Although the developers are trying to deploy more applications for testing, the pods are stuck at the `Pending` state. Use the OpenShift monitoring default stack to find the cause of this availability issue.

Chapter 6 | OpenShift Monitoring

- ▶ 1. Run the load generator for applications that are deployed in the development environment. The load generator `traffic-simulator.sh` script is available at the `~/D0380/labs/monitoring-cluster` location.

- 1.1. Open a terminal window and navigate to the `~/D0380/labs/monitoring-cluster` directory.

```
[student@workstation ~]$ cd ~/D0380/labs/monitoring-cluster
```

- 1.2. Execute the `traffic-simulator.sh` script to generate the load for applications that are deployed in the development environment.

```
[student@workstation monitoring-cluster]$ ./traffic-simulator.sh
...output omitted...
```

- ▶ 2. Use the cluster monitoring dashboards to display cluster resource usage information within the OpenShift web console.

- 2.1. Log in as the `admin` user to the OpenShift web console. To do so, open a web browser window and navigate to `https://console-openshift-console.apps.ocp4.example.com`.

Then click **Red Hat Identity Management**.

- 2.2. Log in as the `admin` user with `redhatocp` as the password.

- 2.3. Navigate to **Observe > Dashboards**.

- 2.4. Select the **Kubernetes / Compute Resources / Cluster** dashboard.

- 2.5. Scroll to the **CPU Quota** section. Click the **CPU Usage** column header, if necessary more than once, until it turns blue with a downward arrow. Rows are sorted by namespace in descending CPU order.

CPU Quota						
Namespace	Pods	Workloads	CPU Usage	CPU Requ...	CPU Requ...	
dev-monitor	13	4	3.02	0.65	464.49%	
openshift-kube-apiserver	19	-	0.62	0.9	68.84%	
dev-finance	11	2	0.589	0.6	98.22%	
openshift-monitoring	22	11	0.551	0.307	179.46%	
openshift-etcd	9	-	0.295	1.11	26.56%	

Figure 6.7: Namespaces with the highest CPU usage

The `dev-monitor` and `dev-finance` namespaces are likely listed among the five namespaces that use the most CPU resources. Although the `dev-monitor` namespace requests only 0.65 CPU resources, it uses considerably more. The result is that the **CPU Requests %** column reports more than 100%.

If you do not see the `dev-monitor` namespace with high resource consumption, then reload the page in your browser.

Chapter 6 | OpenShift Monitoring

- 2.6. Scroll to the **Memory Requests** section. Click the **Memory Usage** column header, if necessary more than once, until it turns blue with a downward arrow. Rows are sorted by namespaces in descending memory order.

Requests by Namespace						
Namespace	Pods	Workloads	Memory U...	↓	Memory R...	Memory R...
openshift-kube-apiserver	19	-	5.57 GiB	3.6 GiB	154.79%	
openshift-monitoring	22	11	5.31 GiB	3.49 GiB	151.96%	
dev-monitor	13	4	1.86 GiB	350 MiB	543.01%	
openshift-storage	22	14	1.61 GiB	12.02 GiB	13.43%	
openshift-etcd	9	-	1.14 GiB	2.68 GiB	42.67%	

Figure 6.8: Namespaces with the highest memory usage

The **dev-monitor** namespace is likely listed among the five namespaces that use the most memory resources. As with CPU usage, the **dev-monitor** namespace uses considerably more memory than the 350 MiB that the pods in the namespace request. The result is that the **Memory Requests %** column reports more than 100%.

- 3. Use the cluster monitoring dashboards to identify the workload and pod that uses the most cluster resources in the **dev-monitor** namespace.

- 3.1. Select the **Kubernetes / Compute Resources / Namespace (Workloads)** dashboard. Select the **dev-monitor** namespace and leave deployment as the workload type.
- 3.2. Scroll to the **CPU Quota** section. The **CPU Quota** section shows that the **frontend** deployment requests a total of 0.5 CPU resources and currently uses 0.001 CPU resources. The **exoplanets-app** and **exoplanets-db** deployments request a total of 0.05 CPU resources, and currently use 0.007 and 0.003 CPU resources respectively. The **python-app** deployment requests a total of 0.1 CPU resources, but the one pod in the deployment consumes more than 2 CPU resources.

CPU Quota						
Workload	Workload Ty...	Running Pods	CPU Usage	CPU Requests	CPU Reques...	
exoplanets-app	deployment	1	0.007	0.025	29.28%	
exoplanets-db	deployment	1	0.003	0.025	10.19%	
frontend	deployment	10	0.001	0.5	0.20%	
python-app	deployment	1	2.02	0.1	2,022.83%	

Figure 6.9: Deployment CPU usage

Although the actual CPU usage in your cluster might differ, it is clear that the **python-app** deployment uses the most CPU resources.

- 3.3. Scroll to the **Memory Usage** section. The **Memory Usage** graph shows that the **python-app** deployment has notable increases and decreases in memory usage. However, the ten pods in the **frontend** deployment consistently use a total of about 63 MiB of memory. One pod in the **exoplanets-app** deployment uses a total of

Chapter 6 | OpenShift Monitoring

about 15.02 MiB of memory, and one pod in the `exoplanets-db` deployment uses a total of about 129.2 MiB of memory.

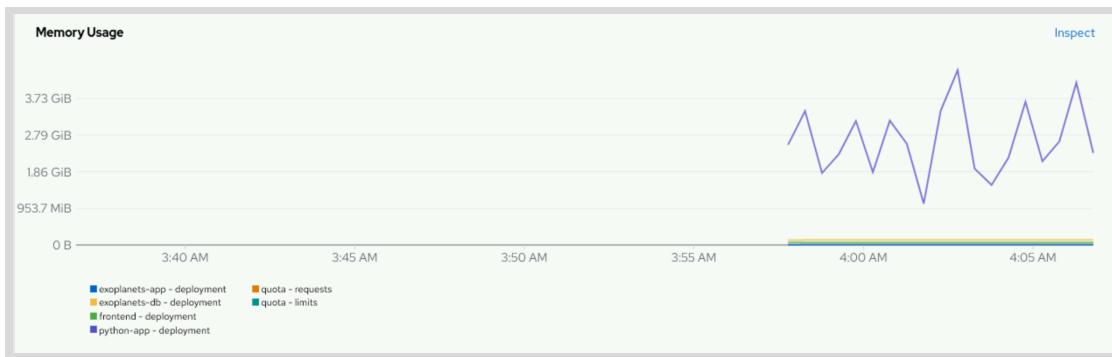


Figure 6.10: Deployment memory usage

- ▶ 4. You identify a deployment that appears to be behaving erratically and that uses more resources than intended. The next step might be to identify the owner of the `python-app` deployment to verify the application.

At minimum, the owner should add resource limits for both CPU and memory to the `python-app` deployment. Although a cluster administrator can add limits, setting incorrect limits might cause the pod to enter a `CrashLoopBackOff` state.

Similarly, a cluster administrator can implement a quota for the `dev-monitor` namespace to limit the total CPU and memory resources that the project uses. Because the `python-app` deployment does not specify limits for either CPU or memory, restricting CPU and memory resources with a quota would prevent the `python-app` pod from running.

- ▶ 5. Return to the terminal where the script is running. Press `Ctrl+C` to stop the script, and then close the terminal window.
- ▶ 6. Verify that the `frontend-test` and `budget-test` test applications are giving an `Application Not Available` message.
 - 6.1. Open a web browser window and navigate to the URL for the `budget-test` application <http://budget-test-monitor.apps.ocp4.example.com>. The web page shows the `Application Not Available` message.

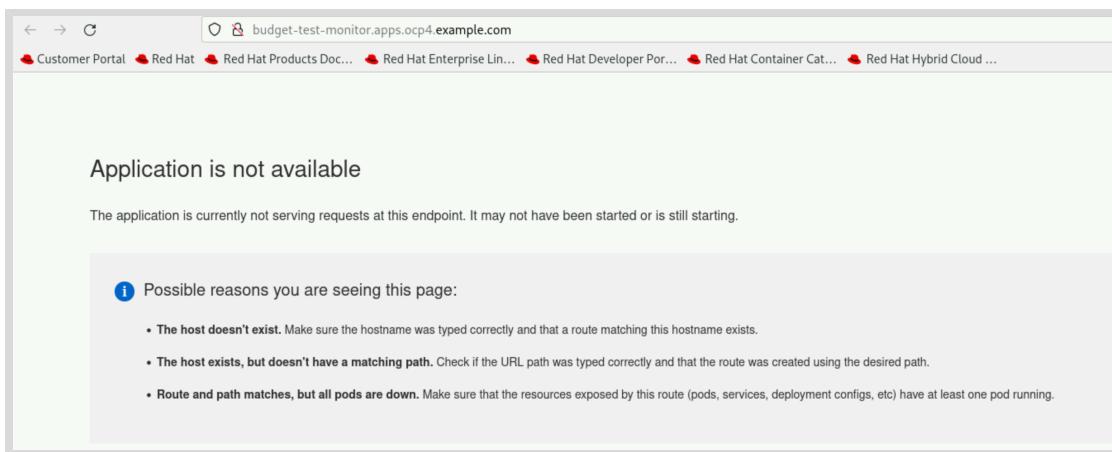


Figure 6.11: Application not available

Chapter 6 | OpenShift Monitoring

- 6.2. Open a web browser tab and navigate to the URL for the `frontend-test` application `http://frontend-test-monitor.apps.ocp4.example.com`. The web page shows the `Application Not Available` message.
- ▶ 7. Use the cluster monitoring dashboards to monitor the workloads and pods in the `test-monitor` namespace.
 - 7.1. Switch to the OpenShift web console web browser tab and select the `Kubernetes / Compute Resources / Namespace (Workloads)` dashboard. Select the `test-monitor` namespace and leave deployment as the workload type.
 - 7.2. Scroll to the `CPU Quota` section. The `CPU Quota` section shows that the `frontend-test` deployment requests a total of 1 CPU resource and currently uses almost no CPU resources. The `budget-test` deployment requests a total of 1 CPU and currently uses almost no CPU resources.

CPU Quota				
Workload	Workload Type	Running Pods	CPU Usage	CPU Requests
<code>budget-test</code>	deployment	2	-	0.05
<code>frontend-test</code>	deployment	20	-	1

Figure 6.12: Deployment CPU usage for the test environment

- 7.3. Scroll to the `Memory Quota` section. The `Memory Quota` section shows that the `budget-test` and `frontend-test` deployments currently use no memory.

Memory Quota				
Workload	Workload Type	Running Pods	Memory Usage	Memory Requests
<code>budget-test</code>	deployment	2	-	40 MiB
<code>frontend-test</code>	deployment	20	-	400 MiB

Figure 6.13: Deployment memory usage for the test environment

Similarly, neither deployment currently shows any network usage.

The applications are showing no CPU, memory, or network usage, and developers cannot deploy applications in the test environment. The testing environment has one `worker02` compute node.

- ▶ 8. Use the cluster monitoring dashboards to monitor the state of compute node usage for the test environment.
 - 8.1. Select the `Node Cluster` dashboard. The `NotReadyNodesCount` section shows that one node is in the not ready state.
 - 8.2. Click `Inspect` in the `NotReadyNodesCount` section. The metric shows that the sum of nodes with an unknown or false status is 1. This metric value implies that one node is not ready.
 - 8.3. Use the following metric query to the filter to identify the node with an unknown or false status.

Chapter 6 | OpenShift Monitoring

```
kube_node_status_condition{condition="Ready", status=~"unknown|false"} == 1
```

The metric lists all entries, and the == 1 condition filters to show the entries with an unknown or false status.

Click **Run queries**.

The screenshot shows the OpenShift Metrics UI with a single query defined:

```
kube_node_status_condition{condition="Ready", status=~"unknown|false"} == 1
```

The results table has the following columns: Name, condition, container, endpoint, job, namespace, node, prometheus, service, and status. There is one row of data:

Name	condition	container	endpoint	job	namespace	node	prometheus	service	status
kube_node_status_condition	Ready	kube-rbac-proxy-main	https-main	kube-state-metrics	openshift-monitoring	worker02	openshift-monitoring/k8s	kube-state-metrics	unknown

Figure 6.14: Node status metric

The output shows that the `worker02` node has unknown status.

- ▶ 9. You identify that the `worker02` compute node is not ready. The test applications are giving an `Application Not Available` message, because the test environment node pool has only compute node, which is not ready.

The cluster administrator can add a healthy node to the test node pool for application availability and troubleshoot the `worker02` compute node.

- ▶ 10. Close the web browser.

Finish

On the `workstation` machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish monitoring-cluster
```

Alerts and Notifications

Objectives

- View alerting rules.
- Configure and silence alerts.

OpenShift Alerts

In OpenShift Container Platform, you can use the Alerts UI to manage alerts, silences, and alerting rules.

Alerting rules

Alerting rules contain a set of conditions that outline a particular state within a cluster. Alerts are triggered when those conditions are true. An alerting rule can be assigned a severity that defines how the alerts are routed.

Alerts

An alert is fired when the conditions that are defined in an alerting rule are true. Alerts notify that a set of conditions apply in an OpenShift Container Platform cluster.

Alert receivers

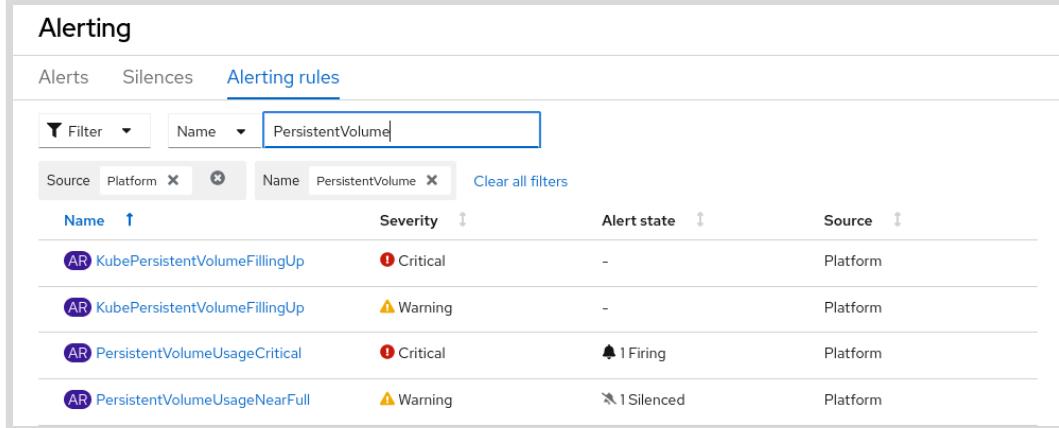
You can configure the alerting system to route alerts to a receiver and send notifications via email, send pager notifications, or forward them to another system by using a webhook.

Silences

A silence can be applied to an alert to prevent sending notifications when the conditions for an alert are true. You can mute an alert after the initial notification, while you work on resolving the underlying issue.

Viewing Alerting Rules

Click **Observe > Alerting** and then click **Alerting rules** to list the alerting rules in the OpenShift cluster. You can apply custom filters to search for a specific alert rule. The alert state column indicates whether the alert is currently firing or is silenced.



The screenshot shows the 'Alerting' interface with the 'Alerting rules' tab selected. A search bar at the top contains the filter 'PersistentVolume'. Below the search bar, there are two sets of filters: 'Source' (Platform) and 'Name' (PersistentVolume). A 'Clear all filters' link is also present. The main table lists four alerting rules:

Name	Severity	Alert state	Source
KubePersistentVolumeFillingUp	Critical	-	Platform
KubePersistentVolumeFillingUp	Warning	-	Platform
PersistentVolumeUsageCritical	Critical	1 Firing	Platform
PersistentVolumeUsageNearFull	Warning	1 Silenced	Platform

Click the alerting rule name to view its details such as its severity, description, and log message. You can also view the PromQL expression, which checks for the alert.

Name: PersistentVolumeUsageNearFull

Severity: Warning

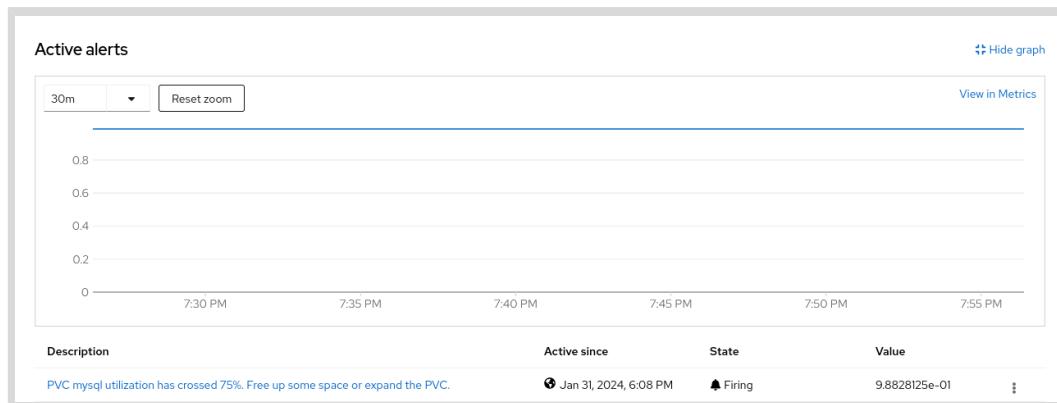
Description: PVC `{{ $labels.persistentvolumeclaim }}` utilization has crossed 75%. Free up some space or expand the PVC.

Message: PVC `{{ $labels.persistentvolumeclaim }}` is nearing full. Data deletion or PVC expansion is required.

Labels: prometheus=openshift-monitoring/k8s, severity=warning

```
(kubelet_volume_stats_used_bytes * on (namespace, persistentvolumeclaim) group_left (storageclass, provisioner) (kube_persistentvolumeclaim_info * on (storageclass) group_left (provisioner) kube_storageclass_info(provisioner=".*rbd.csi.ceph.com") (.cephfs.csi.ceph.com")) / (kubelet_volume_stats_capacity_bytes * on (namespace, persistentvolumeclaim) group_left (storageclass, provisioner) (kube_persistentvolumeclaim_info * on (storageclass) group_left (provisioner) kube_storageclass_info(provisioner=".*rbd.csi.ceph.com") (.cephfs.csi.ceph.com")) > 0.75
```

Scroll down to view the graph that displays the time when the alert was detected.



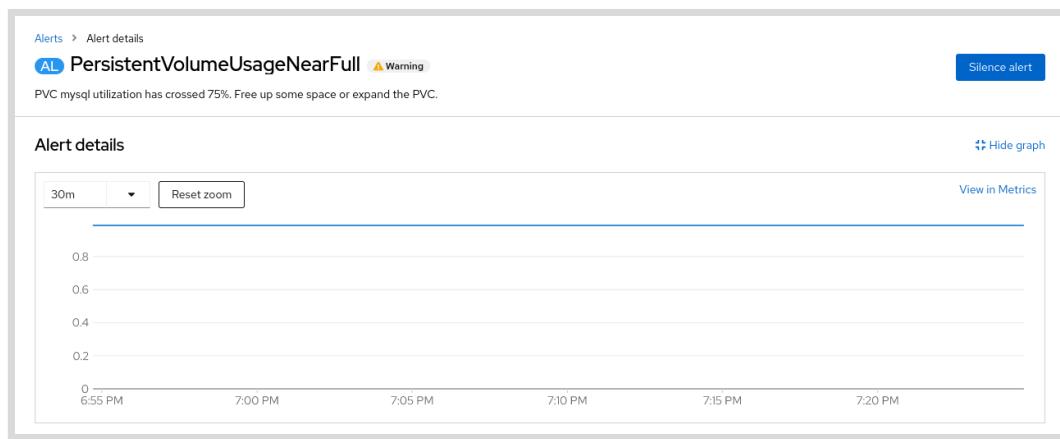
Viewing Firing Alerts

Click **Observe > Alerting** to view the alerts that are currently firing in the cluster. You can apply custom filters to find a specific alert.

Name	Severity	State	Source
AL PersistentVolumeUsageCritical	Critical	Firing	Platform
AL PersistentVolumeUsageNearFull	Warning	Firing	Platform
AL Watchdog	None	Firing	Platform

Chapter 6 | OpenShift Monitoring

Click the alert name to view its properties. The graph displays the time when the alert was detected.

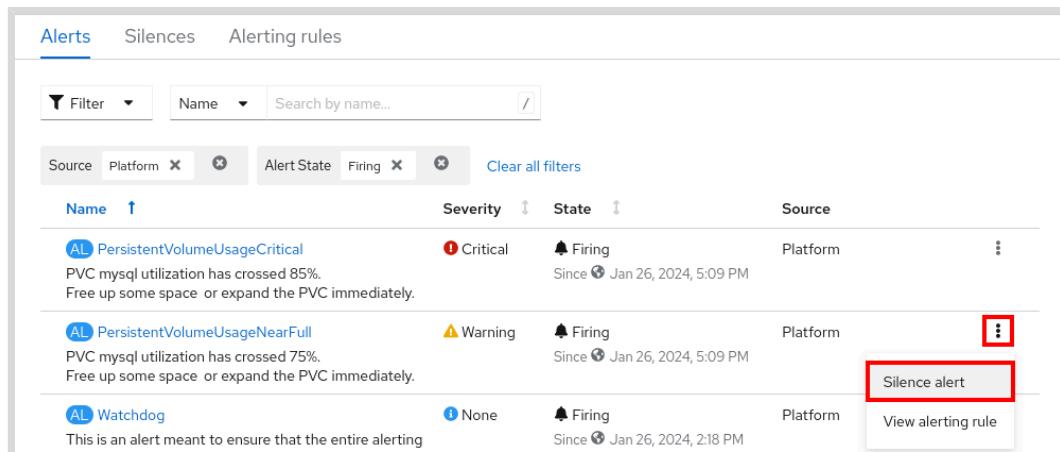


Scroll down to view the alert details and labels. You can create a receiver to match a label and forward the alert.



Silence Alerts

You can silence an alert for a period of time to stop sending new notifications to a receiver. For example, you can silence an alert while you are troubleshooting a problem, to stop sending new alert emails, or silence certain alerts that are expected to be fired during a scheduled maintenance window. Click **Observe > Alerting** to display a list of the alerts that are currently firing. Then click the three dots icon : at the right of the alert, and click **Silence alert**.



The alert silence is configured to start immediately after it is created. You can clear the **Start immediately** field to set a specific start time.

You can select a predefined duration for the alert silence. Select - in the **For** field to set when the silence begins and finishes.

Silence alert

Silences temporarily mute alerts based on a set of label selectors that you define. Notifications will not be sent for alerts that match all the listed values or regular expressions.

Duration

Silence alert from... Now For... 30m Until... 30m from now

Start immediately

Alert labels

Alerts with labels that match these selectors will be silenced instead of firing. Label values can be matched exactly or with a [regular expression](#).

Label name	Label value
alertname	PersistentVolumeUsageNearFull

RegEx Negative matcher

A comment is required when you create an alert silence. This comment is saved in the alert silence details.

Add label

Info

Creator
admin

Comment*
Silenced during troubleshooting

Silence **Cancel**

Viewing Alert Silences

Click **Observe > Alerting** and then click **Silences** to view the alert silences in the cluster.

Alerts	Silences	Alerting rules	
Filter Search by name... / Create silence Expire 0 silences			
Silence State Active Pending Clear all filters			
Name	Firing alerts	State	Creator
SL_PersistentVolumeUsageNearFull alertname=PersistentVolumeUsageNearFull endpoint=https-metrics instance=192.168.50.13:10250 job=kubelet metrics_path=/metrics severity=warning namespace=monitoring-alerts node=worker01	⚠ 1	Active Ends Feb 2, 2024, 5:19 PM	admin

You can see the silence name, start time, end time, matching labels, state, and associated comment.

Chapter 6 | OpenShift Monitoring

The screenshot shows the 'Silence details' page for a silence named 'PersistentVolumeUsageNearFull'. The 'Matchers' section lists several filter criteria: alertname=PersistentVolumeUsageNearFull, endpoint=https-metrics, severity=warning, instance=192.168.50.13:10250, job=kubelet, metrics_path=/metrics, service=kubelet, namespace=monitoring-alerts, node=worker01, persistentvolumeclaim=mysql, provisioner=openshift-storage.cephfs.csi.ceph.com, and storageclass=ocs-external-storagecluster-cephfs. The 'State' section shows it is 'Active'. The 'Starts at' and 'Ends at' sections show the silence began at 4:49 PM on Feb 2, 2024, and ends at 5:19 PM on Feb 2, 2024. The 'Created by' section shows it was created by 'admin'. The 'Comment' section notes it was silenced during troubleshooting. The 'Firing alerts' section shows one warning alert.

You can also view the alerts that match the silence.

The screenshot shows the 'Firing alerts' section. It displays a single alert named 'PersistentVolumeUsageNearFull' with a severity of 'Warning'. The alert message states: 'PVC mysql utilization has crossed 75%. Free up some space or expand the PVC.' A 'View alerting rule' button is available.

Viewing Silenced Alerts

Click **Observe > Alerting** to return to the alerting main section. The **PersistentVolumeUsageNearFull** alert is not displayed in the list because it is currently silenced. Clear the **Alert State** filter by clicking **x** to display all the alerts.

The screenshot shows the 'Alerts' tab of the alerting rules page. The 'Alert State' filter is cleared, showing two alerts: 'PersistentVolumeUsageCritical' (severity Critical, state Firing) and 'Watchdog' (severity None, state Firing). The 'PersistentVolumeUsageNearFull' alert is not visible because it is marked as silenced.

Observe that the **PersistentVolumeUsageNearFull** alert is listed and marked as silenced.

The screenshot shows the 'Alerts' tab of the alerting rules page. The 'Alert State' filter is cleared, showing three alerts: 'PersistentVolumeUsageCritical' (severity Critical, state Firing), 'PersistentVolumeUsageNearFull' (severity Warning, state Silenced), and 'Watchdog' (severity None, state Firing). The 'PersistentVolumeUsageNearFull' alert is highlighted with a red border and shows its silencing information: it was silenced on Jan 26, 2024, at 8:16 PM.

Sending Alert Notifications

You can view the firing alerts in the Alerting UI. Alerts are not configured by default to be sent to any notification systems. You can configure OpenShift Container Platform to send alerts to the following receiver types:

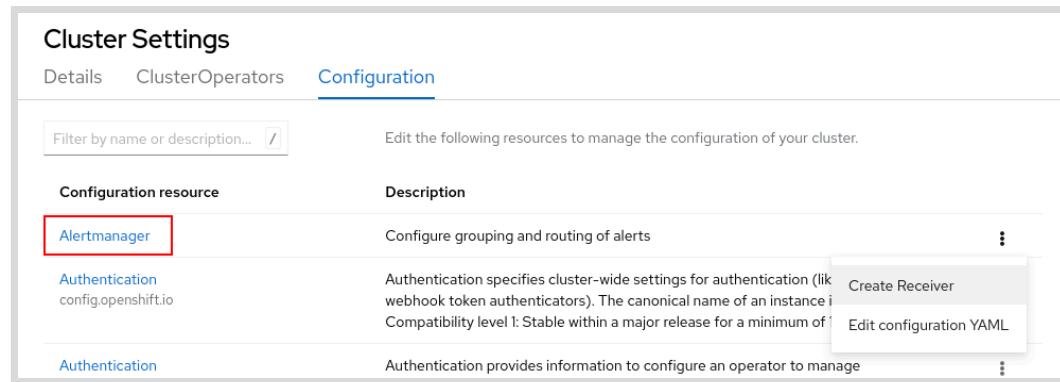
- Email
- PagerDuty
- Slack
- Webhook

By routing alerts to receivers, you can send timely notifications to the appropriate teams when failures occur. For example, critical alerts require immediate attention and are typically paged to an individual or to a critical response team. Alerts that provide non-critical warning notifications might instead be routed to a ticketing system for non-immediate review.

The alerting system uses Alertmanager, which is a component from the Prometheus monitoring stack. The Alertmanager configuration is saved in the `alertmanager-main` secret in the `openshift-monitoring` namespace.

Configuring Alertmanager with the Web Console

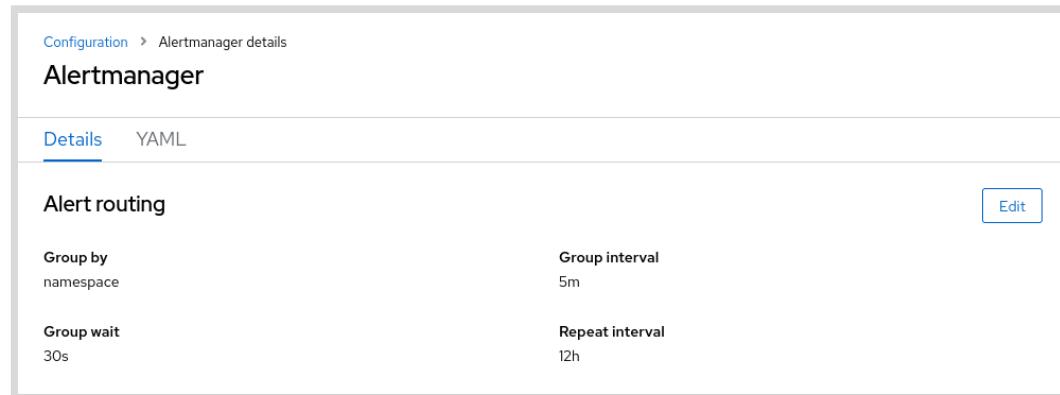
Click Administration > Cluster Settings, and then click Configuration and Alertmanager to open the Alertmanager configuration page. You can click the vertical ellipsis icon  at the right to create an alert receiver or to edit the Alertmanager configuration file.



The screenshot shows the 'Cluster Settings' interface with the 'Configuration' tab selected. A search bar at the top allows filtering by name or description. Below it, a table lists configuration resources. The 'Alertmanager' row is highlighted with a red box and has a vertical ellipsis icon on its right. A context menu is open over this icon, containing options: 'Create Receiver' and 'Edit configuration YAML'. Other rows in the table include 'Authentication config.openshift.io' and 'Authentication'.

Configuration resource	Description	Actions
Alertmanager	Configure grouping and routing of alerts	 Create Receiver Edit configuration YAML
Authentication config.openshift.io	Authentication specifies cluster-wide settings for authentication (like webhook token authenticators). The canonical name of an instance is	
Authentication	Authentication provides information to configure an operator to manage	

The Alertmanager configuration page lists the current alert routing parameters.



The screenshot shows the 'Alertmanager details' page under the 'Configuration' section. It has tabs for 'Details' and 'YAML'. The 'Alert routing' section is visible, containing two groups of parameters: 'Group by namespace' (Group interval: 5m) and 'Group wait' (30s, Repeat interval: 12h). An 'Edit' button is located in the top right corner of this section.

Scroll down to view the configured alert receivers. Click **Create receiver** to add a receiver. You can click the three dots icon  at the right of each receiver to edit its settings or to delete it.

Receivers

Filter Receivers by name... /

Create Receiver

Name	Integration type	Routing labels
Critical	Configure	
Default	Configure	All (default receiver)
null		
Watchdog		

Edit Receiver
Delete Receiver

You can view and edit the current Alertmanager configuration by clicking **YAML**.

Configuration > Alertmanager details

Alertmanager

Details [YAML](#)

Update this YAML to configure Routes, Receivers, Groupings and other Alertmanager settings.

Alt + F1 Accessibility help Show tooltips

```

1 global:
2   resolve_timeout: 5m
3   inhibit_rules:
4     - equal:

```

✓ alertmanager-main has been updated to version 123456789

[Save](#) [Cancel](#) [Download](#)

Configuring the Alerting Email Receiver

Click **Create Receiver** to create a receiver to send the alerts by email. Complete the values according to the following table, and then scroll down and click **Create**.

Field	Value
Receiver name	email
Receiver type	Email
To address	ocp-admins@example.com
SMTP configuration	
Save as default SMTP configuration	<input checked="" type="checkbox"/>
From address	alerts@ocp4.example.com
SMTP smarthost	192.168.50.254:25
SMTP hello	localhost
Auth username	smtp_training
Auth password	Red_H4T@!
Auth identity	(empty)

Field	Value
Auth secret	(empty)
Require TLS	<i>Unset</i>
Routing labels	
Value	alertname=PersistentVolumeUsageNearFull

The Alertmanager configuration changes are applied and the alerting system reloads in a few minutes.

Configuring Alertmanager with the Command Line

The Alertmanager configuration is saved in the `alertmanager-main` secret in the `openshift-monitoring` namespace.

```
[user@host ~]$ oc get secret/alertmanager-main -n openshift-monitoring
NAME          TYPE      DATA   AGE
alertmanager-main   Opaque    1     7d
```

You can extract the `alertmanager-main` secret to view the `alertmanager.yaml` configuration file.

```
[user@host ~]$ oc extract secret/alertmanager-main -n openshift-monitoring \
--to ./ --confirm
alertmanager.yaml
```

The default `alertmanager.yaml` configuration file contains many unnecessary quotation marks. Remove the quotation marks by using the `sed` command to improve readability.

```
[user@host ~]$ sed -f script.sed alertmanager.yaml
```



Important

Although removing the extraneous quotation mark characters is not required, it improves readability. The quotation mark characters are not required in a YAML file, except to represent `null` as a string.

The previous `sed` command uses this script file to remove all the quotation marks and converts `null` to a string.

```
#!/usr/bin/sed -f
s://" ①
s/>\<\(null\)\>/'\\"1'/"g ②
```

- ① Remove all the quotation marks from the file.
- ② Enclose `null` between quotation marks.

Chapter 6 | OpenShift Monitoring

The monitoring stack can send alerts by email through an SMTP server. The following example sends the `PersistentVolumeUsageNearFull` alerts to the `ocp-admins@example.com` email address.

```
global:
  resolve_timeout: 5m
  smtp_from: alerts@ocp4.example.com ①
  smtp_smarthost: '192.168.50.254:25' ②
  smtp_hello: localhost ③
  smtp_auth_username: smtp_training ④
  smtp_auth_password: Red_H4T@! ⑤
  smtp_require_tls: false ⑥
...output omitted...
inhibit_rules:
  ...output omitted...
receivers:
  ...output omitted...
  - name: email ⑦
    email_configs: ⑧
      - to: ocp-admins@example.com ⑨
route:
  group_by:
    - namespace
  group_interval: 2m ⑩
  group_wait: 30s
  receiver: Default
  repeat_interval: 1m ⑪
  routes:
    ...output omitted...
    - receiver: email ⑫
      match:
        alertname: PersistentVolumeUsageNearFull ⑬
        continue: true ⑭
...output omitted...
```

- ① The global SMTP host. If you do not define `smarthost` in the `email_configs` field for a receiver, then this field is the default host in use.
- ② The global email sender address. If you do not define `from` in the `email_configs` field for a receiver, then this field is the default address in use.
- ③ The hello parameter for the SMTP connection.
- ④ The global SMTP username for optional authentication. If you do not define `auth_username` in the `email_configs` field for a receiver, then this field is the default username in use.
- ⑤ The global SMTP password for optional authentication. This password is used if `auth_password` is not defined in the `email_configs` field for a receiver. If you do not define `auth_password` in the `email_configs` field for a receiver, then this field is the default password in use.
- ⑥ A global setting to specify whether TLS is required for SMTP. You can override this setting by using `require_tls` in the `email_configs` field for a receiver.

- ⑦ An arbitrary name for the receiver. A route specifies this receiver name for a match.
- ⑧ This setting indicates that the receiver sends alerts by email.
- ⑨ The `to` setting must be specified in the `email_configs` field, and does not have an equivalent global SMTP setting.
- ⑩⑪ Configure the `group_interval` and `repeat_interval` fields so the alert email notifications are sent more frequently.
- ⑫ The receiver to use if the `match` evaluates as `true` for the alert.
- ⑬ The expression to match a specific alert name.
- ⑭ By default, every alert that enters the routing tree stops after the first matching child. The `continue: true` parameter permits the alert to continue through the routing tree matching subsequent routes.

You can update the Alertmanager configuration by setting the data of the `alertmanager-main` secret in the `openshift-monitoring` namespace with the content of the `alertmanager.yaml` file.

```
[user@host ~]$ oc set data secret/alertmanager-main -n openshift-monitoring \
--from-file alertmanager.yaml
secret/alertmanager-main data updated
```

You can view the progression in the Alertmanager stateful set logs. A successful update generates the log message: `Completed loading of configuration file`.

```
[user@host ~]$ oc logs -f statefulset.apps/alertmanager-main -c alertmanager \
-n openshift-monitoring
Found 2 pods, using pod/alertmanager-main-0
...output omitted...
ts=2024-01-31T01:02:03.064Z caller=coordinator.go:113 level=info
component=configuration msg="Loading configuration file" file=/etc/alertmanager/
config_out/alertmanager.env.yaml
ts=2024-01-31T01:02:03.128Z caller=coordinator.go:126 level=info
component=configuration msg="Completed loading of configuration file" file=/etc/
alertmanager/config_out/alertmanager.env.yaml
```

An incorrect configuration generates a failed log message. If you see configuration errors in the logs, then modify the `alertmanager.yaml` file and reapply your changes to the `alertmanager-main` secret in the `openshift-monitoring` namespace.

```
...output omitted...
ts=2024-01-31T02:03:04.256Z caller=coordinator.go:113 level=info
component=configuration msg="Loading configuration file" file=/etc/alertmanager/
config_out/alertmanager.env.yaml
ts=2024-01-31T02:03:04.512Z caller=coordinator.go:118 level=error
component=configuration msg="Loading configuration file failed" file=/etc/
alertmanager/config_out/alertmanager.env.yaml err="no global SMTP smarthost set"
```



References

For more information about managing alerts in OpenShift, refer to the *Managing Alerts* chapter in the Red Hat OpenShift Container Platform 4.14 *Monitoring* documentation at

https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/monitoring/index#managing-alerts

Querying Prometheus

<https://prometheus.io/docs/prometheus/latest/querying/basics/#querying-prometheus>

AlertManager Configuration

<https://prometheus.io/docs/alerting/0.25/configuration/>

► Guided Exercise

Alerts and Notifications

Configure and silence alerts.

Outcomes

- Configure OpenShift to send email notifications.
- Review an alert notification email.
- Silence a firing alert.

Before You Begin

As the student user on the workstation machine, use the `lab` command to prepare your environment for this exercise.

```
[student@workstation ~]$ lab start monitoring-alerts
```

Instructions

In this guided exercise, you extract the Alertmanager secret and configure the values to send notifications about the `PersistentVolumeUsageNearFull` alerts to the `ocp-admins@example.com` email address.

Then, you deploy an application that triggers the `PersistentVolumeUsageNearFull` alert and wait for the notification email to be delivered to the `ocp-admins@example.com` address.

Finally, you create a silence for the alert for a period, to stop sending new notifications to the email address while you are troubleshooting the problem.

► 1. List and extract the secret that contains the Alertmanager configuration.

- 1.1. Log in to the cluster as the `admin` user.

```
[student@workstation ~]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
Login successful.

...output omitted...
```

- 1.2. Change to the `~/D0380/labs/monitoring-alerts` directory.

```
[student@workstation ~]$ cd ~/D0380/labs/monitoring-alerts
```

- 1.3. List the `alertmanager-main` secret in the `openshift-monitoring` namespace.

Chapter 6 | OpenShift Monitoring

```
[student@workstation monitoring-alerts]$ oc get secret/alertmanager-main \
-n openshift-monitoring
NAME          TYPE    DATA   AGE
alertmanager-main  Opaque  1      7d
```

- 1.4. Extract the existing alertmanager-main secret from the openshift-monitoring namespace to the current directory.

```
[student@workstation monitoring-alerts]$ oc extract secret/alertmanager-main \
-n openshift-monitoring --to ./ --confirm
alertmanager.yaml
```

- ▶ 2. Modify the Alertmanager configuration to send notifications about the persistent volume alerts to the ocp-admins@example.com email address.
 - 2.1. The default alertmanager-main secret contains many unnecessary quotation marks. Remove the quotation marks by using the sed command to improve readability.

```
[student@workstation monitoring-alerts]$ sed -i -f script.sed alertmanager.yaml
```



Important

Although removing the extraneous quotation mark characters is not required, it improves readability. The quotation mark characters are not required in a YAML file, except to represent null as a string.

The previous sed command removes all quotation marks, including possible required quotation marks.

- 2.2. Edit the alertmanager.yaml file and add the global SMTP settings.

```
global:
  resolve_timeout: 5m
  smtp_from: alerts@ocp4.example.com
  smtp_smarthost: '192.168.50.254:25'
  smtp_hello: localhost
  smtp_auth_username: smtp_training
  smtp_auth_password: Red_H4T@!
  smtp_require_tls: false
inhibit_rules:
  ...output omitted...
```

- 2.3. Add the email receiver to the alertmanager.yaml file.

```
...output omitted...
receivers:
- name: Default
- name: Watchdog
- name: Critical
- name: 'null'
```

```
- name: email
  email_configs:
    - to: ocp-admins@example.com
route:
  ...output omitted...
```

- 2.4. Configure the `group_interval` and `repeat_interval` settings so the alert email notifications are sent more frequently.

```
...output omitted...
route:
  group_by:
    - namespace
  group_interval: 2m
  group_wait: 30s
  receiver: Default
  repeat_interval: 1m
  ...output omitted...
```

- 2.5. Finally, add the match rule at the bottom of the file. Then, save and close the file.

```
...output omitted...
route:
  ...output omitted...
  routes:
    ...output omitted...
    - matchers:
      - severity = critical
      receiver: Critical
    - receiver: email
    match:
      alertname: PersistentVolumeUsageNearFull
```



Note

The `~/D0380/solutions/monitoring-alerts/alertmanager.yaml` file contains the correct configuration, and you can use it for comparison.

- 3. Apply the Alertmanager configuration changes.

- 3.1. Update the existing `alertmanager-main` secret in the `openshift-monitoring` namespace with the content of the `alertmanager.yaml` file.

```
[student@workstation monitoring-alerts]$ oc set data secret/alertmanager-main \
-n openshift-monitoring --from-file alertmanager.yaml
secret/alertmanager-main data updated
```

- 3.2. Follow the `alertmanager` container logs and look for the message: "Completed loading of configuration file" file=`/etc/alertmanager/config_out/alertmanager.env.yaml`

New log messages can take up to one minute to be displayed. Press `Ctrl+C` to exit the `oc logs` command.

Chapter 6 | OpenShift Monitoring

```
[student@workstation monitoring-alerts]$ oc logs -f \
  statefulset.apps/alertmanager-main -c alertmanager -n openshift-monitoring
Found 2 pods, using pod/alertmanager-main-0
...output omitted...
ts=2024-01-26T23:51:34.123Z caller=coordinator.go:256 level=info
  component=configuration msg="Loading configuration file" file=/etc/alertmanager/
config_out/alertmanager.env.yaml
ts=2024-01-26T23:51:34.456Z caller=coordinator.go:512 level=info
  component=configuration msg="Completed loading of configuration file" file=/etc/
alertmanager/config_out/alertmanager.env.yaml
...output omitted...
^C
```

**Note**

If you see configuration errors in the logs, then modify the `alertmanager.yaml` file and reapply your changes.

- 4. Create a workload that triggers persistent volume alerts and inspect the alerts.

- 4.1. Switch to the `monitoring-alerts` project.

```
[student@workstation monitoring-alerts]$ oc project monitoring-alerts
Now using project "monitoring-alerts" on server ...
```

- 4.2. Create the workload by using the YAML resource manifest.

```
[student@workstation monitoring-alerts]$ oc apply -f workload.yaml
configmap/mysql created
secret/mysql created
persistentvolumeclaim/mysql created
deployment.apps/mysql created
service/mysql created
```

- 4.3. Wait until the `mysql` pod is running, and verify that the deployment is marked as ready and available.

```
[student@workstation monitoring-alerts]$ watch oc get deployments,pods
Every 2.0s: oc get deployments,pods      workstation: Fri Jan 26 17:07:10 2024

NAME          READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/mysql   1/1     1           1           90s

NAME          READY   STATUS    RESTARTS   AGE
pod/mysql-69f4d555b4-hthp4   1/1     Running   0           90s
```

- 4.4. Get the web console URL.

```
[student@workstation monitoring-alerts]$ oc whoami --show-console
https://console.openshift-console.apps.ocp4.example.com
```

- 4.5. Open the web console URL in the browser. Click **Red Hat Identity Management** and log in with the following credentials:
- Username: admin
 - Password: redhatocp
- 4.6. Click **Observe > Alerting**, and then wait for the following alerts to fire:
- **PersistentVolumeUsageCritical**
 - **PersistentVolumeUsageNearFull**

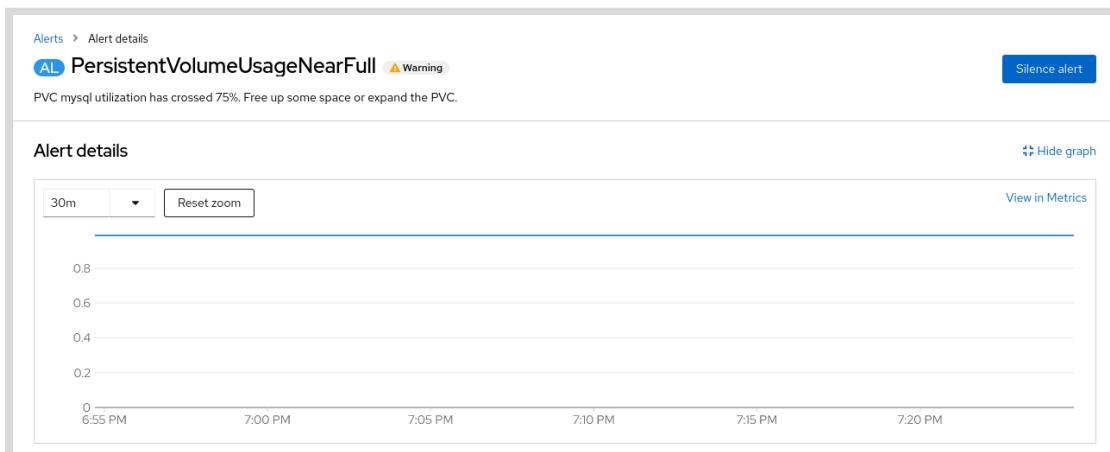
The screenshot shows the 'Alerts' tab in the OpenShift Alerting interface. There are two alerts listed:

- PersistentVolumeUsageCritical**: Critical severity, Firing since Jan 26, 2024, 5:09 PM. Description: PVC mysql utilization has crossed 85%. Action: Free up some space or expand the PVC immediately.
- PersistentVolumeUsageNearFull**: Warning severity, Firing since Jan 26, 2024, 5:09 PM. Description: PVC mysql utilization has crossed 75%. Action: Free up some space or expand the PVC immediately.



The alerts might take a few minutes to be triggered.

- 4.7. Click the **PersistentVolumeUsageNearFull** alert to view its properties. The graph displays the time when the alert was detected.



Scroll down to view the alert details and labels.

Chapter 6 | OpenShift Monitoring

Name: PersistentVolumeUsageNearFull

Severity: Warning

Description: PVC mysql utilization has crossed 75%. Free up some space or expand the PVC.

Labels:

- alertname=PersistentVolumeUsageNearFull
- endpoint=https-metrics
- instance=192.168.50.13:10250
- job=kubelet
- metrics_path=/metrics
- namespace=monitoring-alerts
- node=worker01
- persistentvolumeclaim=mysql
- provisioner=openshift-storage.cephfs.csi.ceph.com
- service=kubelet
- severity=warning
- storageclass=ocs-external-storagecluster-cephfs

Alerting rule: AR PersistentVolumeUsageNearFull

- 5. Verify that the Alertmanager sent an email notification for each firing alert.

The lab user on the utility.lab.example.com host receives emails that are sent to the ocp-admins@example.com email address.

- 5.1. Connect to the utility.lab.example.com host as the lab user.

```
[student@workstation monitoring-alerts]$ ssh lab@utility.lab.example.com
...output omitted...
```

- 5.2. Use the mutt command to access the mail messages.

```
[lab@utility ~]$ mutt
```

- 5.3. The existing emails that are sent from alerts@ocp4.example.com demonstrate that Alertmanager sent email notifications for persistent volume alerts.

```
q:Quit    d:Del    u:Undel    s:Save    m:Mail    r:Reply    g:Group    ?:Help

1 N  Jan 26 alerts@ocp4.exa... ( 251) [FIRING:1] monitoring-alerts (Persistent
...output omitted...

---Mutt: /var/spool/mail/lab [Msgs:1 New:1 123K]-----(date/date)-----(all)---
```

- 5.4. Select the first email message and press Enter to open the email.

```
i:Exit  -:PrevPg <Space>:NextPg  v:View Attach...  d:Del  r:Reply  j:Next  ?:Help

Date: Fri, 26 Jan 2024 23:36:49 +0000
From: alerts@ocp4.example.com
To: ocp-admins@example.com
Subject: [FIRING:1] monitoring-alerts (PersistentVolumeUsageNearFull
https-metrics 192.168.50.13:10250 kubelet /metrics worker01 platform
mysql openshift-monitoring/k8s openshift-storage.cephfs.csi.ceph.com
kubelet warning ocs-external-storagecluster-cephfs)

<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Transitional//EN"
...output omitted...
-N - 1/1: alerts@ocp4.example. [FIRING:1] monitoring-alerts (Persistent -- (1%)
```

**Note**

Alertmanager sends notification emails in HTML format.

5.5. Press **q** to close the message.

► **6.** Silence the **PersistentVolumeUsageNearFull** alert.

6.1. Return to the web browser window and click **Observe > Alerting** to display a list of the alerts that are currently firing.

6.2. Click the three dots icon : at the right of the **PersistentVolumeUsageNearFull** alert, and then click **Silence alert**.

Name	Severity	State	Source
AL PersistentVolumeUsageCritical	Critical	Firing	Platform
AL PersistentVolumeUsageNearFull	Warning	Firing	Platform
AL Watchdog	None	Firing	Platform

6.3. Set the alert silence time to 30 minutes. Complete the values according to the following table, and then scroll down and click **Silence**.

Field	Value
For	30m
Start immediately	Checked
Alert labels	
<i>Leave all the label names and values as they are presented</i>	
Info	
Creator	admin
Comment	Silenced during troubleshooting

6.4. Click **Observe > Alerting** to return to the alerting main section. The **PersistentVolumeUsageNearFull** alert is not displayed in the list, because it is currently silenced.

Clear the **Alert State** filter by clicking **x** to display all the alerts.

The screenshot shows the 'Alerts' tab in the OpenShift Monitoring interface. There are two alerts listed:

Name	Severity	State	Source
AL PersistentVolumeUsageCritical	Critical	Firing	Platform
AL Watchdog	None	Firing	Platform

A red box highlights the 'Clear all filters' button at the top right of the alert list.

- 6.5. Observe that the `PersistentVolumeUsageNearFull` alert is listed and marked as silenced.

The screenshot shows the 'Alerts' tab in the OpenShift Monitoring interface. Three alerts are listed, with the third one highlighted by a red box:

Name	Severity	State	Source
AL PersistentVolumeUsageCritical	Critical	Firing	Platform
AL PersistentVolumeUsageNearFull	Warning	Silenced	Platform
AL Watchdog	None	Firing	Platform

The alert `AL PersistentVolumeUsageNearFull` has a red box around it, indicating it is the target of the exercise step.

- 6.6. Return to the terminal window and observe that no new emails are received. Press `q` to exit `mutt`.

- 6.7. Exit the SSH session.

```
[lab@utility ~]$ exit
```

- 6.8. Change to the student `HOME` directory.

```
[student@workstation monitoring-alerts]$ cd  
[student@workstation ~]$
```

Finish

On the `workstation` machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish monitoring-alerts
```

▶ Lab

OpenShift Monitoring

Configure alerts.

Extract specific insights from monitoring dashboards to troubleshoot performance issues with cluster nodes.

Outcomes

- Configure the OpenShift alert manager to send email notifications for the NodeCPUOvercommit alert rule.
- Review the alert email.
- Use the OpenShift monitoring dashboards to identify issues with cluster nodes.
- Fix the problematic deployment, by reducing its number of pods to zero.

Before You Begin

As the student user on the workstation machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start monitoring-review
```

Instructions

Your company has a node pool for the stage environment in the OpenShift cluster. This node pool uses the `env=stage` label and includes the `worker03` node.

The lab script deploys an alert rule for the stage environment that triggers the `NodeCPUOvercommit` alert. This alert fires when applications request more than the available CPU.

The lab script deploys three applications, which are called `budget-app`, `frontend`, and `python-app`, for the stage environment. These applications are deployed in the `monitor-review` namespace. One of the applications is requesting more vCPU than is available, so find the application and scale it to zero pods. Thus, logs and events for the application are preserved for developers to fix it.

Use the `admin` user with `redhatocp` as the password. For the resources that you must create, you can use the incomplete YAML files in the `~/D0380/labs/monitoring-review` directory.

- As the `admin` user, connect to the OpenShift cluster and configure the OpenShift alert manager to send all notifications for alerts with the `NodeCPUOvercommit` name to the `ocp-admins@example.com` email address. Use the email configuration information from the following table:

Chapter 6 | OpenShift Monitoring

Field	Value
Host	192.168.50.254:25
Username	smtp_training
Password	Red_H4T@!
Sender	alerts@ocp4.example.com
TLS	False

Use one minute for the `group_interval`, `group_wait`, and `repeat_interval` parameters. Use the preconfigured `cpu-overcommit` receiver. You can find an incomplete example for the alert manager configuration in the `~/DO380/labs/monitoring-review/alertmanager.yaml` file.

2. Verify that you receive email alerts from the OpenShift alert manager. The `lab` user on the `utility.lab.example.com` host receives emails that are sent to the `ocp-admins@example.com` email address.

Use the `mutt` command to access the mail messages. The `mutt` command updates in real-time, so it automatically displays any new mail messages.

It can take a few minutes for OpenShift to fire the alert and to start sending emails.



Note

The alert manager sends alerts as emails in HTML format. The `lab` script configures `mutt` to display HTML messages.

3. Use the cluster alerting dashboards to verify the firing alert.
4. Use the cluster monitoring dashboards to display cluster CPU and memory resource usage information within the OpenShift web console. The compute nodes have four vCPUs. Identify the namespace that requests more than the available vCPU.
5. Use the cluster monitoring dashboard with the namespace that was identified in the previous step to investigate further and find the application that is responsible for requesting more than the available vCPU.
6. Use the information from the previous step to fix the problematic deployment by reducing its number of pods to zero. You scale the number of pods to zero to preserve application logs and events for troubleshooting. Verify that after fixing the deployment, the `NodeCPUOvercommit` alert is not fired.

Evaluation

As the `student` user on the `workstation` machine, use the `lab` command to grade your work. Correct any reported failures and rerun the command until successful.

```
[student@workstation ~]$ lab grade monitoring-review
```

Finish

As the student user on the workstation machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish monitoring-review
```

► Solution

OpenShift Monitoring

Configure alerts.

Extract specific insights from monitoring dashboards to troubleshoot performance issues with cluster nodes.

Outcomes

- Configure the OpenShift alert manager to send email notifications for the NodeCPUOvercommit alert rule.
- Review the alert email.
- Use the OpenShift monitoring dashboards to identify issues with cluster nodes.
- Fix the problematic deployment, by reducing its number of pods to zero.

Before You Begin

As the student user on the workstation machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start monitoring-review
```

Instructions

Your company has a node pool for the stage environment in the OpenShift cluster. This node pool uses the `env=stage` label and includes the `worker03` node.

The lab script deploys an alert rule for the stage environment that triggers the `NodeCPUOvercommit` alert. This alert fires when applications request more than the available CPU.

The lab script deploys three applications, which are called `budget-app`, `frontend`, and `python-app`, for the stage environment. These applications are deployed in the `monitor-review` namespace. One of the applications is requesting more vCPU than is available, so find the application and scale it to zero pods. Thus, logs and events for the application are preserved for developers to fix it.

Use the `admin` user with `redhatocp` as the password. For the resources that you must create, you can use the incomplete YAML files in the `~/D0380/Labs/monitoring-review` directory.

- As the `admin` user, connect to the OpenShift cluster and configure the OpenShift alert manager to send all notifications for alerts with the `NodeCPUOvercommit` name to the `ocp-admins@example.com` email address. Use the email configuration information from the following table:

Field	Value
Host	192.168.50.254:25
Username	smtp_training
Password	Red_H4T@!
Sender	alerts@ocp4.example.com
TLS	False

Use one minute for the `group_interval`, `group_wait`, and `repeat_interval` parameters. Use the preconfigured `cpu-overcommit` receiver. You can find an incomplete example for the alert manager configuration in the `~/D0380/labs/monitoring-review/alertmanager.yaml` file.

1. Open a terminal and connect to the OpenShift cluster as the `admin` user with `redhatocp` as the password.

```
[student@workstation ~]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
Login successful.
...output omitted...
```

2. Change to the `~/D0380/labs/monitoring-review` directory.

```
[student@workstation ~]$ cd ~/D0380/labs/monitoring-review
```

3. Create the alert manager configuration file. You can find an incomplete example in the `~/D0380/labs/monitoring-review/alertmanager.yaml` file.

```
global:
  resolve_timeout: 5m
  smtp_smarthost: 192.168.50.254:25
  smtp_hello: localhost
  smtp_auth_username: smtp_training
  smtp_auth_password: Red_H4T@!
  smtp_from: alerts@ocp4.example.com
  smtp_require_tls: false
  ...output omitted...
receivers:
- name: Default
- name: Watchdog
- name: Critical
- name: 'null'
- name: cpu-overcommit
  email_configs:
    - to: ocp-admins@example.com
  ...output omitted...
route:
  group_by:
    - namespace
```

Chapter 6 | OpenShift Monitoring

```
group_interval: 1m
group_wait: 1m
receiver: Default
repeat_interval: 1m
routes:
- matchers:
  - alertname = Watchdog
  receiver: Watchdog
- matchers:
  - alertname = InfoInhibitor
  receiver: 'null'
- matchers:
  - severity = critical
  receiver: Critical
  continue: true
- matchers:
  - alertname = NodeCPUOvercommit
  receiver: cpu-overcommit
```

- 1.4. Apply the alert manager configuration.

```
[student@workstation monitoring-review]$ oc set data secret/alertmanager-main \
-n openshift-monitoring --from-file alertmanager.yaml
secret/alertmanager-main data updated
```

- 1.5. Follow the `alertmanager` container logs and verify that the alert manager configuration is applied successfully. New log messages can take a few minutes to display. Press `Ctrl+C` to exit the `oc logs` command.

```
[student@workstation monitoring-review]$ oc logs -f alertmanager-main-0 \
-c alertmanager -n openshift-monitoring
...output omitted...
ts=2024-02-12T15:41:38.303Z caller=coordinator.go:113 level=info
  component=configuration msg="Loading configuration file" file=/etc/alertmanager/
config_out/alertmanager.env.yaml
ts=2024-02-12T15:41:38.305Z caller=coordinator.go:126 level=info
  component=configuration msg="Completed loading of configuration file" file=/etc/
alertmanager/config_out/alertmanager.env.yaml
```

2. Verify that you receive email alerts from the OpenShift alert manager. The `lab` user on the `utility.lab.example.com` host receives emails that are sent to the `ocp-admins@example.com` email address.

Use the `mutt` command to access the mail messages. The `mutt` command updates in real-time, so it automatically displays any new mail messages.

It can take a few minutes for OpenShift to fire the alert and to start sending emails.

**Note**

The alert manager sends alerts as emails in HTML format. The lab script configures `mutt` to display HTML messages.

- 2.1. Connect to the `utility.lab.example.com` host as the `lab` user.

```
[student@workstation monitoring-review]$ ssh lab@utility.lab.example.com  
...output omitted...
```

- 2.2. Use the `mutt` command to access the mail messages.

```
[lab@utility ~]$ mutt
```

- 2.3. The existing emails that are sent from `alerts@ocp4.example.com` demonstrate that the alert manager sends email notifications for the node CPU overcommitment.

```
q:Quit d:Del u:Undel s:Save m:Mail r:Reply g:Group ?:Help  
  
1 N Feb 12 alerts@ocp4.exa ( 217) [FIRING:1] (NodeCPUOvercommit platform  
openshift-monitoring/k8s critical)  
...output omitted...
```

- 2.4. Select the first alert message and press `Enter` to open the email.

```
i:Exit -:PrevPg <Space>:NextPg v:View Attach... d:Del r:Reply j:Next ?:Help  
Date: Mon, 12 Feb 2024 11:44:44 +0000  
From: alerts@ocp4.example.com  
To: ocp-admins@example.com  
Subject: [FIRING:1] (NodeCPUOvercommit platform openshift-monitoring/k8s  
critical)  
...output omitted...
```

- 2.5. Press `q` to close the message, and then press `q` again to exit `mutt`.

- 2.6. Exit the SSH session.

```
[lab@utility ~]$ exit
```

3. Use the cluster alerting dashboards to verify the firing alert.

- 3.1. Open a web browser window and navigate to `https://console-openshift-console.apps.ocp4.example.com`. Click **Red Hat Identity Management**. Log in as the `admin` user with `redhatocp` as the password.

- 3.2. Navigate to **Observe > Alerting**.

- 3.3. Verify that the `NodeCPUOvercommit` alert is in the firing state.

4. Use the cluster monitoring dashboards to display cluster CPU and memory resource usage information within the OpenShift web console. The compute nodes have four vCPUs. Identify the namespace that requests more than the available vCPU.

- 4.1. Change to the web browser window, and navigate to **Observe > Dashboards**.

- 4.2. Select the **Kubernetes / Compute Resources / Cluster** dashboard.

- 4.3. Scroll to the **CPU Quota** section. Click the **CPU Requests** column header, if necessary more than once, until it turns blue with a downward arrow. Rows are sorted by

Chapter 6 | OpenShift Monitoring

namespace in descending CPU order. The `monitoring-review` namespace is requesting more than the available vCPU from the stage node.

If you do not see the `monitoring-review` namespace with high resource consumption, then reload the page in your browser.

- 4.4. Scroll to the **Memory Requests** section. Click the **Memory Requests** column header, if necessary more than once, until it turns blue with a downward arrow. Rows are sorted by namespaces in descending memory order. The `monitoring-review` namespace is likely to be listed among the five namespaces that request the most memory resources.
5. Use the cluster monitoring dashboard with the namespace that was identified in the previous step to investigate further and find the application that is responsible for requesting more than the available vCPU.
 - 5.1. Select the **Kubernetes / Compute Resources / Namespace (Workloads)** dashboard.
 - 5.2. Select the `monitoring-review` namespace in the **Namespace** drop-down menu.
 - 5.3. Scroll to the **CPU Quota** section. Verify that the `python-app` deployment requests more vCPU than is available.
 - 5.4. Scroll to the **Memory Requests** section. Verify that the `python-app` deployment requests the most memory resources.
6. Use the information from the previous step to fix the problematic deployment by reducing its number of pods to zero. You scale the number of pods to zero to preserve application logs and events for troubleshooting. Verify that after fixing the deployment, the `NodeCPUOvercommit` alert is not fired.
 - 6.1. Change to the terminal window and scale down the `python-app` deployment to zero replicas.

```
[student@workstation monitoring-review]$ oc scale deployment/python-app \
  -n monitoring-review --replicas 0
deployment.apps/python-app scaled
```

- 6.2. Change to the student HOME directory.

```
[student@workstation monitoring-review]$ cd
```

- 6.3. Change to the web browser window, and navigate to **Observe > Alerting**.
- 6.4. After a few minutes, verify that the `NodeCPUOvercommit` alert is not fired.

Evaluation

As the `student` user on the workstation machine, use the `lab` command to grade your work. Correct any reported failures and rerun the command until successful.

```
[student@workstation ~]$ lab grade monitoring-review
```

Finish

As the student user on the workstation machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish monitoring-review
```

Summary

- Red Hat OpenShift Observability collects logs, traces, events, and system metrics to provide real-time monitoring.
- Red Hat OpenShift Logging aggregates all the logs from the pods and nodes of an OpenShift cluster to a centralized location.
- The OpenShift monitoring stack is based on the Prometheus open source project.
- The default monitoring stack is included with OpenShift Container Platform.
- The components of the default monitoring stack are installed in the `openshift-monitoring` project and provide monitoring features for core platform components.
- In OpenShift Container Platform, you can use the Alerts UI to manage alerts, silences, and alerting rules.
 - Alerting rules contain a set of conditions that outline a particular state within a cluster.
 - An alert is fired when the conditions that are defined in an alerting rule are true.
 - You can apply a silence to an alert to prevent sending notifications when the conditions for an alert are true.

Chapter 7

OpenShift Logging

Goal

Deploy OpenShift Logging and query log entries from workloads and cluster nodes.

Sections

- Log Forwarding (and Guided Exercise)
- Centralized Logging (and Guided Exercise)

Lab

- OpenShift Logging

Log Forwarding

Objectives

- Deploy OpenShift Logging for forwarding logs to an external aggregator.

OpenShift Logging

OpenShift Logging collects and aggregates the log messages from all the pods and nodes in your cluster. Users and administrators can use the OpenShift web console to search and consult log entries.

Depending on the workload and the size of your cluster, processing many logs can require significant disk space and compute resources, which would then not be available for your application workloads. In such a scenario, you might need to deploy more compute nodes and increase your storage capacity to handle the extra load.

You can configure OpenShift Logging to forward the logs to a third-party log aggregator for long-term storage, or to an observability platform for further analysis, and minimize the resource requirement on your cluster.

The following examples of third-party logging solutions can receive logs from OpenShift Logging:

- Elasticsearch
- Grafana Loki
- Splunk
- Amazon CloudWatch
- Google Cloud Logging

OpenShift Logging Components

OpenShift Logging is based on these components: a collector, a log store, and a visualization console. You can deploy them together as a complete logging solution, or you can deploy the collector alone and store the logs in an external solution.

Chapter 7 | OpenShift Logging

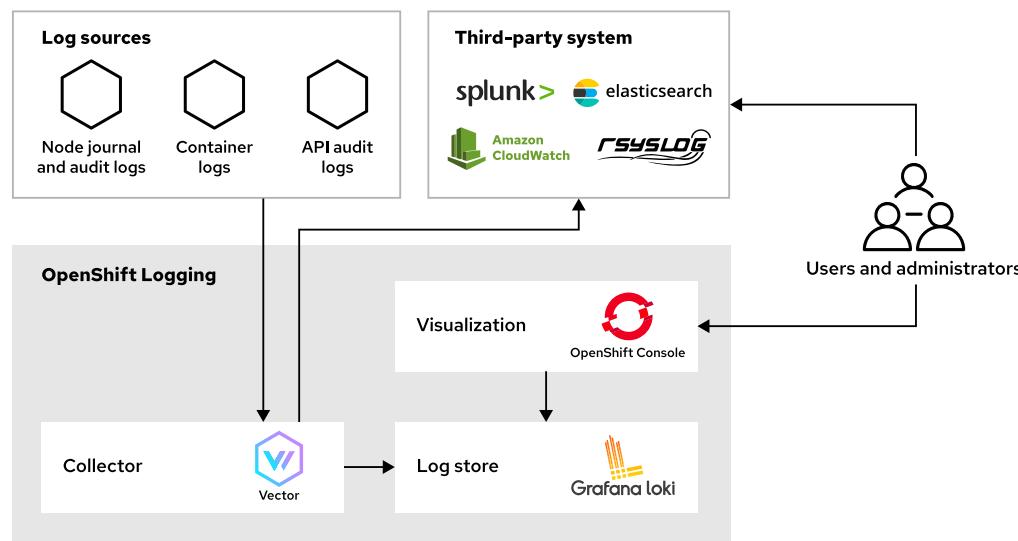


Figure 7.1: OpenShift Logging architecture

Log Collector

The collector is the main component of OpenShift Logging. OpenShift Logging uses Vector to collect logs from all running containers and cluster nodes, and replaces Fluentd, which was the collector in earlier versions of OpenShift Logging.

Vector collects various log types from your cluster that are then grouped into these categories:

Infrastructure

Infrastructure logs include container logs in the `openshift-*`, `kube*`, and `default` namespaces, and system logs from the cluster nodes.

Audit

Audit logs include both Kubernetes API and OpenShift API audit logs, as well as the Linux audit logs from the cluster nodes. These logs might contain sensitive security details, and OpenShift Logging does not store them by default.

Application

Application logs are all container logs from user projects.

To collect those logs, Vector runs as a daemon set in the cluster and therefore runs on all nodes.

In addition to collecting logs, Vector adds metadata to describe where the logs come from, and then forwards the logs to the log store, which is either internal or external to the cluster.

Log Store

The log store uses Grafana Loki to aggregate logs from the entire cluster into a central place and provides access control to logs.

Loki replaces Elasticsearch, which was the log store in earlier versions of the logging subsystem.

The internal log store is an optional component of OpenShift Logging.

Visualization

OpenShift Logging provides a native OpenShift Console plug-in to view and query logs in the internal log store.

The OpenShift Logging UI component replaces Kibana, which was the web interface in earlier versions of OpenShift Logging.

Install and Configure OpenShift Logging

You can deploy OpenShift Logging by installing the OpenShift Logging operator with the Operator Lifecycle Manager (OLM) from OperatorHub on the web console or by using the `oc` command to create the OLM resources. The OpenShift Logging operator manages deploying and configuring the log collector and other resources to support the logging subsystem.

See the references section for instructions to install an operator by using OperatorHub or the `oc` command.

After installation, you configure the logging subsystem by using the `ClusterLogging` custom resource (CR), and configure the log collector by using the `ClusterLogForwarder` CR.

Configure OpenShift Logging Components

The `ClusterLogging` CR configures and manages deploying the OpenShift Logging components.

To configure OpenShift Logging, you must create the cluster logging resource in the `openshift-logging` namespace. The minimum configuration for OpenShift Logging is to enable the log collector, as follows:

```
apiVersion: logging.openshift.io/v1
kind: ClusterLogging
metadata:
  name: instance 1
  namespace: openshift-logging
spec:
  managementState: Managed 2
  collection:
    type: vector 3
```

- 1** The resource name must be `instance`.
- 2** The management state must be `Managed` for the components to receive updates from the OpenShift Logging operator.
- 3** The collector type to deploy. It can be either `vector` or the deprecated `fluentd` collector.

You can use the cluster logging resource to enable or disable logging components, to customize pod placement, and to define pod resource limits. Configuring the log store and logging UI components are discussed in detail in the next section.

In the following example, the log collector component is configured to run on all nodes, including dedicated infrastructure nodes with the `node-role.kubernetes.io/infra=reserved:NoExecute` and `node-role.kubernetes.io/infra=reserved:NoSchedule` taints:

```

apiVersion: logging.openshift.io/v1
kind: ClusterLogging
metadata:
  name: instance
  namespace: openshift-logging
spec:
  managementState: Managed
  collection:
    type: vector
    tolerations: ①
      - effect: NoSchedule
        key: node-role.kubernetes.io/infra
        value: reserved
      - effect: NoExecute
        key: node-role.kubernetes.io/infra
        value: reserved

```

- ① List of tolerations to include in the collector daemon set.



Note

OpenShift Logging automatically adds the `node-role.kubernetes.io/master:NoSchedule` toleration to the collector daemon set so the collector can run on both control planes and compute nodes.

After the cluster logging resource is created, the OpenShift Logging operator starts scheduling pods for the components. The log collector pods are deployed only if the internal log store is enabled, or if log forwarding to an external log aggregator is configured.

Configure Log Collection and Forwarding

The log collector is configured by default to forward infrastructure and application logs to the internal log store that you define in the cluster logging resource. If the internal log store is not deployed, then you must configure the log collector to forward logs to a third-party logging system instead.

Vector can forward logs to various log stores, in addition to the internal Loki deployment, such as Splunk, Amazon CloudWatch, or any logging solution that uses the syslog protocol.

See the references section for a list of external log stores that OpenShift Logging can use.

The `ClusterLogForwarder` custom resource configures the log collector, and defines which logs to collect and where to send them. The following example is a cluster log forwarder resource that configures Vector to forward audit and infrastructure logs to an external Splunk instance:

```

apiVersion: logging.openshift.io/v1
kind: ClusterLogForwarder
metadata:
  name: instance ①
  namespace: openshift-logging
spec:
  outputs:
    - name: splunk-receiver ②

```

```

secret: ③
  name: splunk-auth-token
type: splunk ④
  url: https://mysplunkserver.example.com:8088/services/collector ⑤
pipelines:
  - name: to-splunk ⑥
    inputRefs:
      - audit
      - infrastructure
    outputRefs:
      - splunk-receiver

```

- ①** The resource name must be `instance` and in the `openshift-logging` namespace.
- ②** Name of the log output. You use that name in the log pipeline to refer to that log destination.
- ③** Secret to use for connecting to the log store, if required. In this example, the secret contains the authentication token for the Splunk instance.
- ④** Type of the external log store.
- ⑤** URL of the external log store.
- ⑥** Log pipeline configuration that forwards audit and infrastructure logs to the Splunk instance.

The cluster log forwarder resource is composed of inputs, outputs, and pipelines.

inputs

An input defines the log type to collect. OpenShift Logging provides a predefined input for each log category: `infrastructure`, `audit`, and `application`.

outputs

An output defines a destination for the logs. You can configure multiple log outputs to one or more external log stores.

If the internal log store is configured, then the `default` output becomes available to target the internal Loki instance.

pipelines

A pipeline defines a routing from one or more log inputs to one or more log outputs. You can create several log pipelines and use any combination of log inputs and outputs to forward logs according to your needs.

In the next configuration example, the collector forwards logs to multiple log aggregators:

- The audit logs to a remote syslog server for long-term storage and security compliance
- Both audit logs and infrastructure logs to the internal log store for cluster administrators
- The application logs to an external Splunk instance for developers

```

apiVersion: logging.openshift.io/v1
kind: ClusterLogForwarder
metadata:
  name: instance
  namespace: openshift-logging
spec:
  outputs:

```

```

- name: audit-rsyslog ①
  syslog: ②
    appName: ocp-prod
    facility: local0
    msgID: audit-msg
    procID: audit-proc
    severity: informational
    rfc: RFC5424
  type: syslog
  url: tcp://audit-store.example.com:5514
- name: splunk-receiver ③
  secret:
    name: splunk-auth-token
  type: splunk
  url: https://mysplunkserver.example.com:8088/services/collector
pipelines:
  - name: audit-to-syslog ④
    inputRefs:
      - audit
    outputRefs:
      - audit-rsyslog
  - name: admin-logs ⑤
    inputRefs:
      - audit
      - infrastructure
    outputRefs:
      - default
  - name: app-logs ⑥
    inputRefs:
      - application
    outputRefs:
      - splunk-receiver

```

- ① Log output configuration for the remote syslog
- ② Specific configuration for syslog to format the log messages
- ③ Log output configuration for the Splunk instance
- ④ Log pipeline that forwards audit logs to the remote syslog
- ⑤ Log pipeline that forwards audit and infrastructure logs to the internal log store
- ⑥ Log pipeline that forwards application logs to the Splunk instance

Collect Kubernetes Events

The Event Router is an optional OpenShift Logging component that you can deploy to log Kubernetes events. The Event Router monitors the OpenShift event API and sends the events to the container stdout. The log collector captures the Event Router container logs and forwards them to the log store through the `infrastructure` category.

Because the OpenShift Logging operator does not manage this component, it must be manually deployed and updated. You can deploy the Event Router by using the OpenShift template from the documentation.

Chapter 7 | OpenShift Logging

See the references section for more information about the Event Router and its deployment.

Filtering Log

OpenShift Logging provides filtering capabilities to limit the number of logs to forward to the log store. For example, you can limit the log collection to a specific set of OpenShift projects or application pods by creating a custom log input in the cluster log forwarder resource, as follows:

```
apiVersion: logging.openshift.io/v1
kind: ClusterLogForwarder
metadata:
  name: instance
  namespace: openshift-logging
spec:
  inputs:
    - name: production-apps ①
      application:
        selector:
          matchLabels:
            environment: production
    - name: qa-chain ②
      application:
        selector:
          matchLabels:
            environment: development
      namespaces:
        - qa-testing
        - builders
    ...output omitted...
  pipelines:
    - name: to-splunk ③
      inputRefs:
        - qa-chain
        - production-apps
      outputRefs:
        - splunk-receiver
```

- ① The `production-apps` log input collects application logs from pods with the `environment: production` label.
- ② The `qa-chain` log input collects application logs from pods with the `environment: development` label in the `qa-testing` and `builders` projects.
- ③ The log pipeline forwards both `qa-chain` and `production-apps` log inputs to the Splunk instance.

In addition to application logs, you can filter Kubernetes API audit events.

Each call to the Kubernetes and OpenShift APIs generates an audit event that is forwarded to the log store. Although these events can be valuable for security audits, they represent a high volume of data, which can increase your storage requirement and network bandwidth usage, and therefore increase the cost of your logging solution.

To mitigate this effect, you can define audit filters that remove unwanted or low-value events from the audit logs, and reduce the data that is sent to the log store.

Audit filters are defined in the cluster log forwarder resource and list a set of rules that match the events to remove. You can then apply these filters to the selected log pipeline.

In the following example, an audit filter is created to remove all audit events from infrastructure service accounts, and to remove audit events from updates to leases resources by Kubernetes system users:

```
spec:  
  filters:  
    - name: unwanted-events ①  
      type: kubeAPIAudit  
      kubeAPIAudit:  
        rules:  
          - level: None ②  
            namespaces:  
              - openshift-*  
              - kube*  
            userGroups:  
              - system:serviceaccounts:openshift-*  
              - system:nodes  
            - level: None ③  
              resources:  
                - group: coordination.k8s.io  
                  resources:  
                    - leases  
              users:  
                - system:kube*  
                - system:apiserver  
              verbs:  
                - update  
...output omitted...  
  pipelines:  
    - name: audit-to-syslog  
      inputRefs:  
        - audit  
      filterRefs: ④  
        - unwanted-events  
      outputRefs:  
        - audit-rsyslog
```

- ① Name of the filter. You use that name in the log pipeline to refer to that filter configuration.
- ② The first rule removes all audit events that were created from users in the system:serviceaccounts:openshift-* and system:nodes groups, for resources in the openshift-* and kube* namespaces.
- ③ The second rule removes audit events that were generated from the system:kube* and system:apiserver users, who update leases resources in the coordination.k8s.io API group.
- ④ The filterRefs field lists all filters to apply to the log pipeline.

The KubeAPIAudit field uses the same syntax as the Kubernetes audit policy resources. See the references section for more information about the Kubernetes audit policy and its syntax.

Troubleshoot Log Forwarding

You can verify the configuration of the cluster log forwarder by reviewing the status of the resource. The OpenShift Logging operator validates each log input, output, and pipeline for configuration errors:

```
[user@host ~]$ oc -n openshift-logging describe clusterlogforwarder/instance
...output omitted...
Status:
  Conditions:
    Message:          clusterlogforwarder is not ready
    Reason:           ValidationFailure
    Status:            True
    Type:              Validation
  Inputs:
    Critical - Apps:
      Last Transition Time: 2024-01-18T14:24:20Z
      Status:                True
      Type:                  Ready
  Outputs:
    Audit - Syslog:
      Last Transition Time: 2024-01-18T14:24:20Z
      Status:                True
      Type:                  Ready
    Infra - Syslog:
      Last Transition Time: 2024-01-18T14:24:20Z
      Status:                True
      Type:                  Ready
  Pipelines:
    Audit - Syslog:
      Last Transition Time: 2024-01-18T14:24:20Z
      Status:                True
      Type:                  Ready
    Critical - Apps - Syslog:
      Last Transition Time: 2024-01-18T14:24:20Z
      Message:               invalid: unrecognized outputs: [app-syslog], no valid
  outputs
    Reason:             Invalid
    Status:              False
    Type:                Ready
...output omitted...
```

In this example, one of the configured log pipelines refers to an unknown log output (app-syslog). After the configuration is fixed, the cluster log forwarder resource enters the ready state.

You can also verify the status of the logging configuration by reviewing Kubernetes events about the logging resources:

```
[user@host ~]$ oc get event -n openshift-logging \
--field-selector=involvedObject.name=instance
LAST SEEN    TYPE      REASON          OBJECT
MESSAGE
2m25s        Normal    ReconcilingLoggingCR   clusterlogging/instance
  Reconciling logging resource
78s          Normal    ReconcilingLoggingCR   clusterlogforwarder/instance
  Reconciling logging resource
4m11s        Warning   Invalid           clusterlogforwarder/instance
  clusterlogforwarder is not ready
78s          Normal    Ready             clusterlogforwarder/instance
  ClusterLogForwarder is valid
```

Each time that the log forwarder configuration is updated, the OpenShift Logging operator redeploys the collector pods.

If the cluster log forwarder configuration is correct, but logs are still not forwarded to the external log store, then you can review the collector logs for errors:

```
[user@host ~]$ oc -n openshift-logging logs daemonset/collector
...output omitted...
... ERROR vector::topology::builder: msg="Healthcheck failed." error=Connect
  error: Connection refused (os error 111) component_kind="sink"
  component_type="socket" component_id=infra_syslog component_name=infra_syslog
```

In this example, Vector cannot reach the external log store that is configured in the `infra_syslog` log output. Verify that the remote server is accessible from the OpenShift cluster and that the log forwarder configuration is correct.

If you enable the cluster monitoring on the `openshift-logging` namespace during the operator installation or by adding the `openshift.io/cluster-monitoring: true` label to the namespace, then you can use the monitoring dashboard to inspect the collector status.

From the **Logging / Collection** monitoring dashboard, you can see how many logs are forwarded to each configured log store and the log rate for each application that is running in your cluster.

Chapter 7 | OpenShift Logging

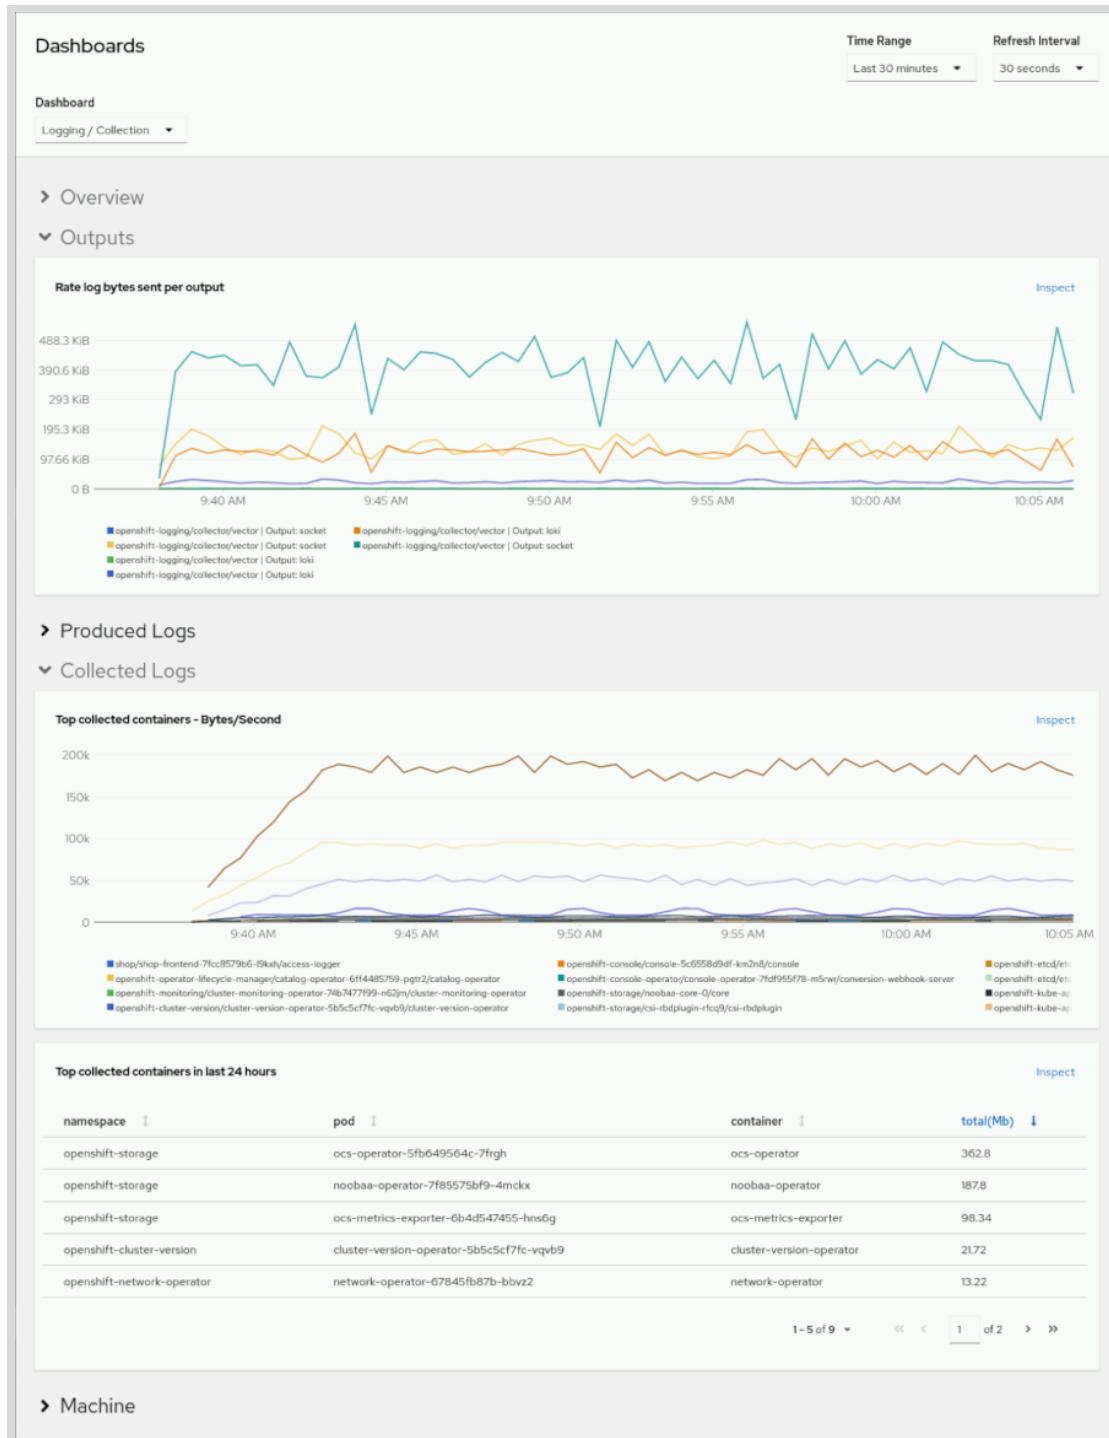


Figure 7.2: OpenShift Logging monitoring dashboard



References

For more information about compatible log destinations, refer to the *Log Output Types* chapter in the Red Hat OpenShift Container Platform 4.14 *Logging* documentation at

https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/logging/index#logging-output-types

For more information about installing the OpenShift Logging operator, refer to the *Installing Logging* chapter in the Red Hat OpenShift Container Platform 4.14 *Logging* documentation at

https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/logging/index#cluster-logging-deploying

For more information about installing the Event Router, refer to the *Collecting and Storing Kubernetes Events* chapter in the Red Hat OpenShift Container Platform 4.14 *Logging* documentation at

https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/logging/index#cluster-logging-eventrouter

Kubernetes Audit Policy

<https://kubernetes.io/docs/tasks/debug/debug-cluster/audit/#audit-policy>

ClusterLogForwarder API Audit Filter

<https://github.com/openshift/cluster-logging-operator/blob/master/docs/features/logforwarding/filters/api-audit-filter.adoc>

► Guided Exercise

Log Forwarding

Deploy OpenShift Logging for forwarding logs to an external syslog server.

Outcomes

- Deploy the OpenShift Logging operator.
- Configure the OpenShift Logging operator to forward logs to a remote syslog.

Before You Begin

As the student user on the workstation machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start logging-forward
```

Instructions

Your company requires you to send OpenShift logs to the company's centralized syslog server for long-term storage.

You must forward the following log types:

- Kubernetes API audit logs
- Kubernetes events
- CoreOS journal logs
- CoreOS audit logs
- Infrastructure container logs
- Critical application container logs



Note

For critical application pods, the `logging: critical` label is assigned. Such applications are available in the `shop` OpenShift project to verify the log forwarding configuration.

Install the OpenShift Logging operator and configure log forwarding to the syslog server. Deploy the OpenShift Event Router to capture Kubernetes events and forward them to the syslog server.

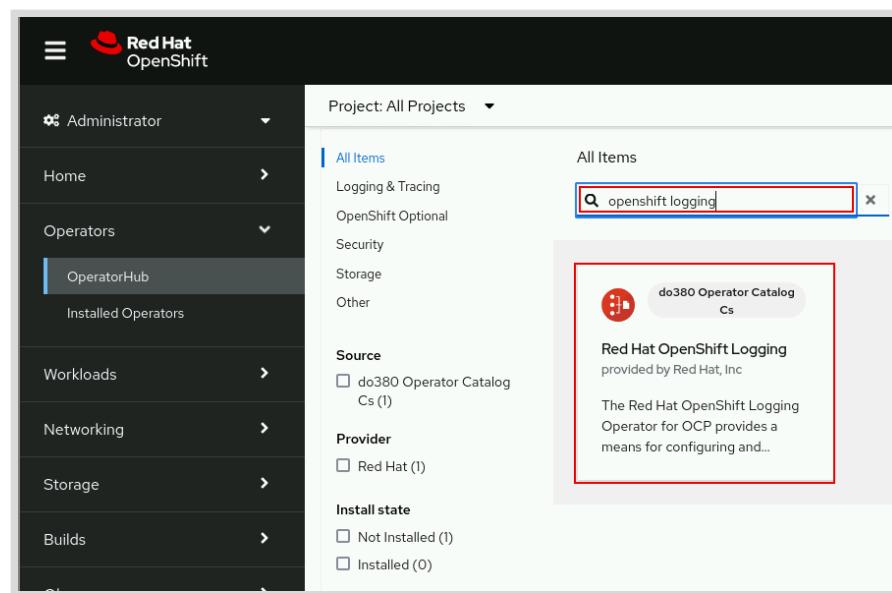
Verify that all required log types are forwarded and available on the syslog server. Forwarded OpenShift logs are stored in the `/var/log/openshift` path on the utility machine.

► 1. As the `admin` user, install the OpenShift Logging operator.

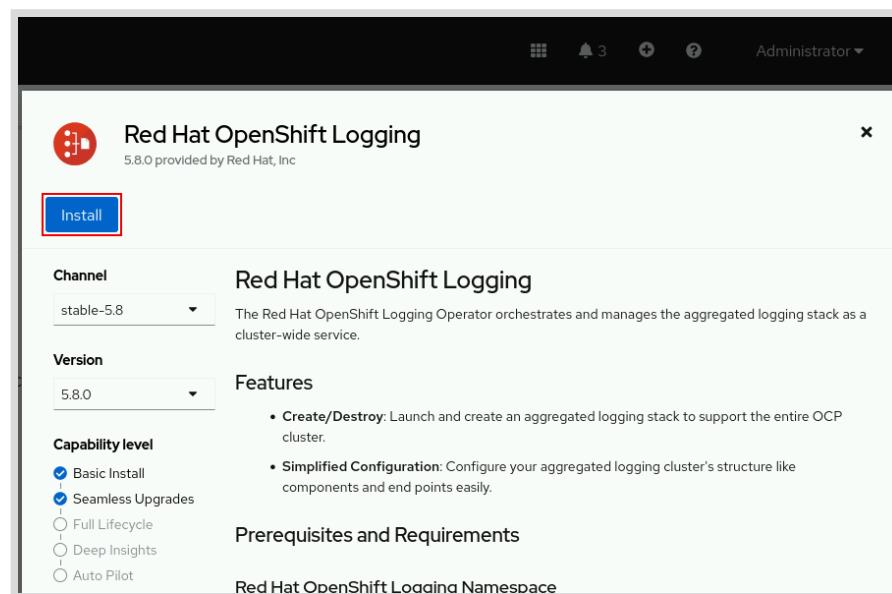
1. Open a web browser and navigate to `https://console-openshift-console.apps.ocp4.example.com`.
2. Click Red Hat Identity Management and log in as the `admin` user with `redhatocp` as the password.

Chapter 7 | OpenShift Logging

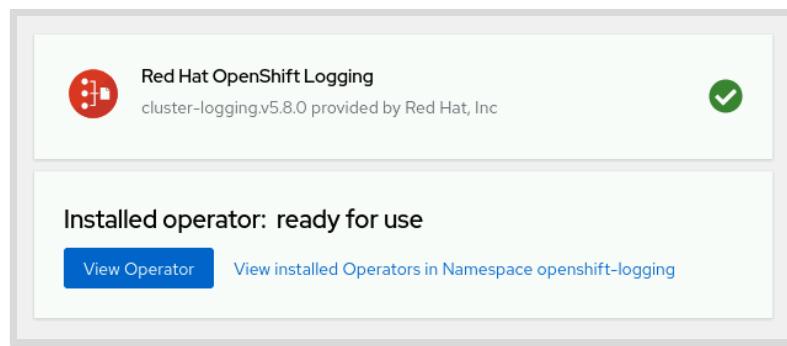
13. Click Operators > OperatorHub. In the Filter by keyword field, type openshift logging to locate the OpenShift Logging operator, and then click Red Hat OpenShift Logging.



14. The web console displays information about the Red Hat OpenShift Logging operator. Click Install to proceed to the Install Operator page.



15. Click Install to install the operator with the default options.
16. Wait until the installation is complete and the web console displays the ready for use message.



▶ 2. Configure OpenShift Logging to enable the Vector log collector.

- 2.1. From the terminal, log in to your OpenShift cluster as the `admin` user.

```
[student@workstation ~]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
Login successful.

...output omitted...
```

- 2.2. Change to the `openshift-logging` project.

```
[student@workstation ~]$ oc project openshift-logging
Now using project "openshift-logging" on server ...
```

- 2.3. Change to the `~/D0380/labs/logging-forward` directory.

```
[student@workstation ~]$ cd ~/D0380/labs/logging-forward
[student@workstation logging-forward]$
```

- 2.4. Edit and complete the `ClusterLogging` resource definition in the `~/D0380/labs/logging-forward/clusterlogging.yaml` file.

```
apiVersion: logging.openshift.io/v1
kind: ClusterLogging
metadata:
  name: instance
  namespace: openshift-logging
spec:
  managementState: Managed
  collection:
    type: vector
```



Note

You can use the solution in the `~/D0380/solutions/logging-forward/clusterlogging.yaml` file.

- 2.5. Apply the cluster logging configuration.

```
[student@workstation logging-forward]$ oc apply -f clusterlogging.yml
clusterlogging.logging.openshift.io/instance created
```

► 3. Configure OpenShift Logging to forward logs to the syslog server.

The syslog server is configured to filter OpenShift audit, infrastructure, and application logs into separate files in the /var/log/openshift path by using the msgID syslog attribute with the audit, infra, and apps values, respectively.

Configure the cluster log forwarder and define three log pipelines to send each log type to the syslog server by using the matching msgID value.

Ensure that only application logs with the `logging: critical` label are forwarded.

The syslog server DNS name is `utility.lab.example.com` and the service listens to the TCP port 514.

3.1. Edit and complete the `ClusterLogForwarder` resource definition in the `~/DO380/labs/logging-forward/clusterlogforwarder.yaml` file.

```
apiVersion: logging.openshift.io/v1
kind: ClusterLogForwarder
metadata:
  name: instance
  namespace: openshift-logging
spec:
  inputs:
    - name: critical-apps
      application:
        selector: ①
        matchLabels:
          logging: critical
  outputs:
    - name: audit-syslog ②
      type: syslog
      url: tcp://utility.lab.example.com:514
      syslog:
        msgID: audit
      ...output omitted...
    - name: apps-syslog ③
      type: syslog
      url: tcp://utility.lab.example.com:514
      syslog:
        msgID: apps
      ...output omitted...
    - name: infra-syslog ④
      type: syslog
      url: tcp://utility.lab.example.com:514
      syslog:
        msgID: infra
      ...output omitted...
  pipelines:
    - name: critical-apps-syslog
      inputRefs:
        - critical-apps
```

Chapter 7 | OpenShift Logging

```

outputRefs:
  - apps-syslog
- name: infra-syslog
  inputRefs:
    - infrastructure
  outputRefs:
    - infra-syslog
- name: audit-syslog
  inputRefs:
    - audit
  outputRefs:
    - audit-syslog

```

- ➊ Label selector to include application pod logs with the **logging: critical** label only
- ➋ Syslog parameters for the audit log type
- ➌ Syslog parameters for the application log type
- ➍ Syslog parameters for the infrastructure log type
- ➎ Log pipelines for each log type

**Note**

You can use the solution in the `~/D0380/solutions/logging-forward/clusterlogforwarder.yaml` file.

3.2. Apply the configuration for the cluster log forwarder.

```
[student@workstation logging-forward]$ oc apply -f clusterlogforwarder.yaml
clusterlogforwarder.logging.openshift.io/instance created
```

3.3. Verify that the OpenShift Logging operator deploys the collector pod on each node.

```
[student@workstation logging-forward]$ oc get daemonset -l component=collector
NAME      DESIRED   CURRENT   READY     UP-TO-DATE   AVAILABLE   ...
collector  6          6          6          6          6          ...
```

▶ 4. Deploy the Event Router component to capture Kubernetes events.

4.1. Use the OpenShift template from the `~/D0380/labs/logging-forward/eventrouter.yaml` file to deploy the Event Router component.

```
[student@workstation logging-forward]$ oc process -f eventrouter.yaml \
| oc apply -f -
serviceaccount/eventrouter created
clusterrole.rbac.authorization.k8s.io/event-reader created
clusterrolebinding.rbac.authorization.k8s.io/event-reader-binding created
configmap/eventrouter created
deployment.apps/eventrouter created
```

Chapter 7 | OpenShift Logging

- 4.2. Verify that the Event Router is running.

```
[student@workstation logging-forward]$ oc get pod -l component=eventrouter
NAME          READY   STATUS    RESTARTS   AGE
eventrouter-59dfb998cc-9v7gb   1/1     Running   0          8m17s
```

- 5. Verify that the syslog service received the logs on the utility machine.

- 5.1. Open a new terminal and connect to the utility machine with SSH as the root user.

```
[student@workstation logging-forward]$ ssh root@utility
```

**Note**

The remainder of this activity refers to this new terminal as the utility terminal, and the first terminal as the workstation terminal.

- 5.2. Check that the OpenShift log files are created in the /var/log/openshift path.

```
[root@utility ~]# ls -l /var/log/openshift
total 298564
-rw----- 1 root root 863869 Jan 11 11:38 apps.log
-rw----- 1 root root 70381151 Jan 11 11:38 audit.log
-rw----- 1 root root 204538257 Jan 11 11:38 infra.log
```

- 6. Verify that the journal logs from the cluster nodes are forwarded to the syslog server.

- 6.1. On the utility terminal, change to the /var/log/openshift directory.

```
[root@utility ~]# cd /var/log/openshift
[root@utility openshift]#
```

- 6.2. Monitor the infrastructure log and filter the output with the jq command to list events from the sshd daemon on the cluster nodes.

```
[root@utility openshift]# tail -f infra.log | egrep -o "\{.*\}" \
| jq '. | select(.systemd.u.SYSLOG_IDENTIFIER=="sshd") \
| .hostname + " " + .message'
```

**Note**

The select() function of the jq command filters the log messages that match the provided expression. In this specific case, only messages with the .systemd.u.SYSLOG_IDENTIFIER field with the value sshd are listed.

- 6.3. On the workstation terminal, connect to the master01 cluster node with the ssh command as the core user.

Chapter 7 | OpenShift Logging

```
[student@workstation logging-forward]$ ssh core@master01.ocp4.example.com
[core@master01 ~]$
```

- 6.4. On the **utility** terminal, verify that the SSH connection is logged. Press **Ctrl+C** to exit the **tail** command.

```
"master01 main: sshd: ssh-rsa algorithm is disabled"
"master01 Accepted publickey for core from 172.25.250.9 port 51158 ssh2: RSA
SHA256:M8i...BT0"
"master01 pam_unix(sshd:session): session opened for user core(uid=1000) by
(uid=0)"
^C
[root@utility openshift]#
```

- 6.5. On the **workstation** terminal, close the SSH connection with the **exit** command or **Ctrl+D**.

```
[core@master01 ~]$ exit
logout
Connection to master01.ocp4.example.com closed.
[student@workstation logging-forward]$
```

- ▶ 7. Verify that the audit logs from the cluster nodes are forwarded to the syslog server.

- 7.1. On the **utility** terminal, review the audit log. Use the **jq** command to list the **USER_LOGIN** Linux audit events that are generated from the previous SSH connection to the **master01** node.

```
[root@utility openshift]# egrep -ho "\{.*\}" audit.log* \
| jq '. | select(.audit.linux.type == "USER_LOGIN") \
| .["@timestamp" + " " + .hostname + " " + .message'
...output omitted...
"2024-01-12T11:08:25.618+00:00 master01 type=USER_LOGIN
msg=audit(1705057705.618:937): pid=106902 uid=0 auid=1000 ses=11
subj=system_u:system_r:sshd_t:s0-s0:c0.c1023 msg='op=login id=1000
exe=\"/usr/sbin/sshd\" hostname=? addr=172.25.250.9 terminal=/dev/pts/0
res=success'\u0001dUID=\"root\" AUID=\"core\" ID=\"core\""
```

**Note**

To use a JSON key that contains special characters, such as **.**, **@** or **-** with the **jq** command, you must surround the key name with double quotes.

- ▶ 8. Verify that the infrastructure container logs are forwarded to the syslog server.

- 8.1. On the **utility** terminal, monitor the infrastructure log. Use the **jq** command to display container logs from the **openshift-storage** namespace. Press **Ctrl+C** to exit the **tail** command.

```
[root@utility openshift]# tail -f infra.log | egrep -o "\{.*\}" \
| jq '. | select(.kubernetes.namespace_name == "openshift-storage") \
| .kubernetes.pod_name + " " + .message'
...output omitted...
"noobaa-core-0 ... time=\"2024-01-12T11:45:54Z\" level=info msg=\"...\""
"noobaa-operator-... time=\"2024-01-12T11:45:54Z\" level=info msg=\"...\""
^C
[root@utility openshift]#
```

▶ 9. Verify that the Kubernetes audit logs are forwarded to the syslog server.

- 9.1. On the utility terminal, monitor the audit log. Use the jq command to display audit events about the developer user.

```
[root@utility openshift]# tail -f audit.log | egrep -o "\{.*\}" \
| jq -c '. | select(.user.username == "developer") \
| [.user.username, .annotations, .verb, .objectRef]'
```

- 9.2. On the workstation terminal, log in to the OpenShift cluster as the developer user with developer as the password.

```
[student@workstation logging-forward]$ oc login -u developer -p developer \
https://api.ocp4.example.com:6443
Login successful.

...output omitted...
```

- 9.3. Create the logger-app project.

```
[student@workstation logging-forward]$ oc new-project logger-app
Now using project "logger-app" on server "https://api.ocp4.example.com:6443".

...output omitted...
```

- 9.4. On the utility terminal, verify that the Kubernetes audit events from the developer user are logged. Press Ctrl+C to exit the tail command.

Chapter 7 | OpenShift Logging

```

["developer", {"authorization.k8s.io/decision": "allow", "authorization.k8s.io/reason": "RBAC: allowed by ClusterRoleBinding \\"basic-users\\" of ClusterRole \\"basic-user\\" to Group \\"system:authenticated\\\"", "get", {"apiGroup": "user.openshift.io", "apiVersion": "v1", "name": "~", "resource": "users"}]

...output omitted...

["developer", {"authorization.k8s.io/decision": "allow", "authorization.k8s.io/reason": "RBAC: allowed by ClusterRoleBinding \\"self-provisioners\\" of ClusterRole \\"self-provisioner\\" to Group \\"system:authenticated:oauth\\\"", "create", {"apiGroup": "project.openshift.io", "apiVersion": "v1", "name": "logger-app", "resource": "projectrequests"}]
^C

[root@utility openshift]#

```

► **10.** Verify that the Kubernetes events are forwarded to the syslog server.

- 10.1. On the **utility** terminal, monitor the infrastructure log. Use the `jq` command to display the Kubernetes events on the `logger-app` namespace.

```

[root@utility openshift]# tail -f infra.log | egrep -o "\{.*\}" \
| jq -c '.
| select(.kubernetes.event.involvedObject.namespace == "logger-app")
| [.kubernetes.event.involvedObject.kind,
.kubernetes.event.involvedObject.name, .message]'

```

- 10.2. On the **workstation** terminal, deploy the `logger` application by using the definition in the `~/DO380/labs/logging-forward/logger.yml` file.

```

[student@workstation logging-forward]$ oc apply -f logger.yml
deployment.apps/logger-app created

```

- 10.3. On the **utility** terminal, verify that the Kubernetes events from the `logger-app` namespace are logged. Press `Ctrl+C` to exit the `tail` command.

```

...output omitted...
["Pod", "logger-app-...", "Add eth0 [10.11.0.18/23] from ovn-kubernetes"]
["Pod", "logger-app-...", "Container image \"registry.ocp4.example.com:8443/redhattraining/access-logger:v0.1\" already present on machine"]
["Pod", "logger-app-...", "Created container access-logger"]
["Pod", "logger-app-...", "Started container access-logger"]
^C

```

► **11.** Verify that only for critical applications, such as the `shop-web` application in the `shop` project, the container logs are forwarded to the syslog server.

- 11.1. On the **utility** terminal, review the application log. Verify that container logs from the `shop-web` application are forwarded.

```
[root@utility openshift]# grep -m1 shop-web apps.log | egrep -o "\{.*\}" | jq
{
  "@timestamp": "2024-01-11T16:23:14.051244304Z",
  ...
  ...output omitted...
  ...
  "container_name": "shop-web",
  "labels": {
    "app": "shop-web",
    "logging": "critical",
    "pod-template-hash": "5cf65ffc4"
  },
  "namespace_name": "shop",
  ...
  ...output omitted...
  ...
  "log_type": "application",
  "message": "2024/01/11 16:23:14 \"GET /products?id=371 HTTP/1.1\" 403
  \"Mozilla/5.0 (X11; Linux x86_64; rv:25.0) Gecko/20100101 Firefox/33.0\"",
}
}
```

- 11.2. Verify that container logs from the `logger-app` application that was deployed in the previous step are not forwarded.

```
[root@utility openshift]# grep logger-app apps.log
no output expected
```

- 11.3. Close the `utility` terminal and return to the `workstation` terminal.

▶ **12.** Clean up the resources.

- 12.1. Change to the student `HOME` directory.

```
[student@workstation logging-forward]$ cd
[student@workstation ~]$
```

- 12.2. Log in to your OpenShift cluster as the `admin` user.

```
[student@workstation ~]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
Login successful.

...output omitted...
```

- 12.3. Delete the `logger-app` project.

```
[student@workstation ~]$ oc delete project logger-app
project.project.openshift.io "logger-app" deleted
```

- 12.4. Delete the `ClusterLogging` and `ClusterLogForwarder` resources.

```
[student@workstation ~]$ oc -n openshift-logging \
    delete clusterlogging,clusterlogforwarder --all
clusterlogging.logging.openshift.io "instance" deleted
clusterlogforwarder.logging.openshift.io "instance" deleted
```

12.5. Remove the Event Router deployment.

```
[student@workstation ~]$ oc -n openshift-logging delete deployment eventrouter
deployment.apps "eventrouter" deleted
```

Finish

On the workstation machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish logging-forward
```

Centralized Logging

Objectives

- Deploy OpenShift Logging for short-term log retention and aggregation.

Log Retention

By default, Kubernetes stores logs from pods in the local disk on the nodes. These logs are ephemeral, because the information is lost when the pod is deleted. You can keep the pods in the cluster instead of deleting them, to keep access to the logs. However, Kubernetes garbage collectors delete the pods when such disk space is needed, and access to the logs is lost. Kubernetes also removes the logs when a node is restarted, for example after applying an update.

Administrators and developers need access to the logs to debug the applications in the event of issues, to analyze application performance, or to review object state changes. A one-week retention period is typically enough for users and administrators to review these logs.

OpenShift Logging uses Grafana Loki as the internal log store to aggregate logs from the entire cluster into a central place and to provide access control to logs. Loki replaces the deprecated Elasticsearch log store from earlier versions of OpenShift Logging. OpenShift Logging can also use other external log stores, such as Splunk, Amazon CloudWatch, or Google Cloud Logging.

Loki Log Store

Loki is a horizontally scalable, highly available, and multitenant log aggregation system. Loki indexes only a few fixed labels during ingestion, and then compresses and stores the log data in chunks in object stores. The Loki operator supports object stores such as AWS S3, Azure, and OpenShift Data Foundation. The Loki operator includes the `LokiStack` CR to manage the Loki deployment.

The steps for using Loki as the log store for OpenShift Logging are as follows:

1. Install the Loki operator.
2. Configure the Loki object storage.
3. Create a `LokiStack` CR with the object storage credentials.
4. Configure the `ClusterLogging` CR from OpenShift Logging to use Loki as the log store.

Install the Loki Operator

OpenShift Logging does not automatically deploy Loki. You must deploy Loki by installing the Loki operator from OperatorHub or by using the CLI. See the references section for instructions.

The Loki operator creates and manages the components of the log store.

Configure Loki Object Storage

Storing logs for your cluster might require significant disk space. Because the required disk space depends on the size of your cluster and application workloads, you might need to plan your storage solution with application developers.

**Note**

This section focuses only on how to use ODF as the Loki object storage.

For more details about ODF, refer to the *DO370: Enterprise Kubernetes Storage with Red Hat OpenShift Data Foundation* training course.

For more information about how to configure other S3-compatible object stores, refer to https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/logging/index#logging-loki-storage_installing-log-storage

Similar to a persistent volume claim, you can use an object bucket claim to request an S3-compatible bucket. ODF creates a secret and a configuration map with the same name as the object bucket claim; between them, they provide the credentials and information to access the bucket.

The following excerpt is an example of an object bucket claim definition in the `openshift-logging` namespace:

```
apiVersion: objectbucket.io/v1alpha1
kind: ObjectBucketClaim
metadata:
  name: loki-bucket-odf
  namespace: openshift-logging
spec:
  generateBucketName: loki-bucket-odf
  storageClassName: openshift-storage.noobaa.io
```

ODF creates a configuration map with the bucket information in the `openshift-logging` namespace. You can retrieve one part of the credentials for the bucket from the configuration map as follows:

```
[user@host ~]$ oc get configmap/loki-bucket-odf -o yaml
apiVersion: v1
data:
  BUCKET_HOST: s3.openshift-storage.svc
  BUCKET_NAME: loki-bucket-odf-69d9...ab91
  BUCKET_PORT: "443"
  BUCKET_REGION: ""
  BUCKET_SUBREGION: ""
kind: ConfigMap
...output omitted...
```

You can retrieve the other part of the bucket credentials from the secret as follows:

```
[user@host ~]$ oc extract --to=-- secret/loki-bucket-odf -n openshift-logging
# AWS_ACCESS_KEY_ID
00s0Y8oaZtPfp18DRjSa
# AWS_SECRET_ACCESS_KEY
z1QtDH37lUvLnkjpe3E4aS8yQI56CLPozGJ0t31
```

After you retrieve the credentials, create a secret that contains those credentials in the `openshift-logging` namespace for Loki to use the credentials:

```
[user@host ~]$ oc create secret generic logging-loki-odf \
-n openshift-logging \
--from-literal=access_key_id=${ACCESS_KEY_ID} \
--from-literal=access_key_secret=${SECRET_ACCESS_KEY} \
--from-literal=bucketnames=${BUCKET_NAME} \
--from-literal=endpoint=https:// ${BUCKET_HOST}:${BUCKET_PORT}
```

Create a LokiStack Instance

After installing the Loki operator and creating the object bucket, you create and configure the `LokiStack` instance with the bucket credentials.

The following example is a `LokiStack` resource that uses the `logging-loki-odf` secret that contains the bucket credentials:

```
apiVersion: loki.grafana.com/v1
kind: LokiStack
metadata:
  name: logging-loki
  namespace: openshift-logging
spec:
  limits: ①
    global:
      retention: ②
        days: 20
        stream: ③
          - days: 4
            priority: 1
            selector: '{kubernetes_namespace_name=~"test.+"}'
  size: 1x.demo ④
  storage:
    secret:
      name: logging-loki-odf ⑤
      type: s3 ⑥
    tls:
      caName: openshift-service-ca.crt ⑦
  storageClassName: ocs-external-storagecluster-ceph-rbd ⑧
  tenants:
    mode: openshift-logging
```

- ① Defines the limits that Loki applies to the log streams. See the references section for more information about the limit types for log streams.
- ② This limit defines 20 days for log retention.
- ③ The retention policy that is based on log streams. Stream-based retention policies are useful if lowering the retention period for a log stream. For example, you might switch to debug mode temporarily in a project to troubleshoot a critical issue. However, debug logs can quickly fill your storage if the retention period is long. Then, you reduce the retention period for that project for a few days. The selector for the stream-based retention policy is specified in the LogQL query language, which is explained later in this section.

Chapter 7 | OpenShift Logging

- ④ The size of the LokiStack instance. You can configure it by using sizes such as `1x.small` and `1x.medium`.
- ⑤ The name for the secret that contains the bucket credentials.
- ⑥ The corresponding storage type.
- ⑦ Loki must configure the Certificate Authority (CA) for the object storage endpoint if the object storage endpoint certificate is not trusted. For ODF, you can use the OpenShift service CA in the `openshift-service-ca.crt` configuration map. For other object stores, create a configuration map that contains the CA certificate.
- ⑧ The name of a storage class for temporary storage. You can list the storage classes for your cluster by using the `oc get storageclasses` command.

**Note**

Red Hat recommends setting the global retention at the object storage level for cost-efficient pruning of logs.

See the references section for more information about the lifecycle bucket configuration and the `LokiStack` CR.

After you create the `LokiStack` instance, the Loki operator starts scheduling pods for components. You can verify the installation by listing the pods in the `openshift-logging` project and verifying that the Loki pods are running.

Configure the OpenShift Logging Log Store

The OpenShift Logging ClusterLogging CR configures and manages deploying the OpenShift Logging components.

To configure OpenShift Logging to use Loki as the log store, you must create a cluster logging resource in the `openshift-logging` namespace and set `lokistack` as the log store type. Then, OpenShift Logging uses Loki to store the infrastructure and application logs.

The following excerpt shows a minimum configuration for OpenShift Logging to enable Loki as the log store:

```
apiVersion: logging.openshift.io/v1
kind: ClusterLogging
metadata:
  name: instance
  namespace: openshift-logging
spec:
  managementState: Managed
  logStore:
    type: lokistack ①
    lokistack:
      name: logging-loki ②
  collection:
    type: vector
```

- ① The log store type.

- ❷ The name for the internal LokiStack instance.

After you create the `ClusterLogging` instance, the OpenShift Logging operator starts scheduling pods for components. You can verify the installation by listing the pods in the `openshift-logging` project and verifying that the collector pods are running.

Audit Logs

By default, the OpenShift Logging operator includes logs for the infrastructure and the application. However, the Logging operator does not include audit logs, because these logs might contain sensitive security details.

To include audit logs in the log store, you must create a `ClusterLogForwarder` CR to configure the log collector to collect them.

The following excerpt shows a minimum configuration to configure the `ClusterLogForwarder` CR to forward the audit logs together with the infrastructure and application logs to the internal Loki instance:

```
apiVersion: logging.openshift.io/v1
kind: ClusterLogForwarder
metadata:
  name: instance
  namespace: openshift-logging
spec:
  pipelines:
    - name: all-to-default
      inputRefs: ❶
        - infrastructure
        - application
        - audit
      outputRefs: ❷
        - default
```

- ❶ The type of logs to collect, which include infrastructure, audit, and application logs.
❷ The destination for the logs, with the internal Loki instance being the default output.

View, Search, and Query Logs

OpenShift Logging provides a native OpenShift Console plug-in to view, search, and query logs in the internal log store.

To use the OpenShift web console for logging visualization, first enable the OpenShift Logging console plug-in. You can use the `Operators > Installed Operators` menu, by selecting the Red Hat OpenShift Logging operator, and ensuring that the console plug-in is enabled as in the following image:

The screenshot shows the 'Operator details' page for the 'Red Hat OpenShift Logging' operator. The top navigation bar includes 'Project: openshift-logging' and 'Actions'. Below the header, there's a summary card for the operator, followed by tabs for 'Details', 'YAML', 'Subscription', 'Events', 'All instances', 'Cluster Log Forwarder', and 'Cluster Logging'. The 'Details' tab is selected. Under 'Provided APIs', there are two sections: 'Cluster Log Forwarder' and 'Cluster Logging'. The 'Cluster Logging' section includes a 'Provider' field (Red Hat, Inc.), a 'Created at' field (Jan 23, 2024, 10:45 AM), and a 'Console plugin' status (Enabled). A red box highlights the 'Console plugin Enabled' status.

Then, create a cluster logging resource in the `openshift-logging` namespace and set `ocp-console` as the visualization type. The following excerpt shows a minimum configuration for OpenShift Logging to enable the OpenShift web console for visualization:

```
apiVersion: logging.openshift.io/v1
kind: ClusterLogging
metadata:
  name: instance
  namespace: openshift-logging
spec:
  ...
  visualization:
    type: ocp-console
```

After you enable the web console for visualization, you can view your logs in the **Observe > Logs** menu. The following image shows the **Logs** menu and the filters that you can apply to your logs by using the web console:

The screenshot shows the OpenShift Logging interface. At the top, there are several filter and search controls: a histogram button (1), a time range dropdown set to 'Last 1 day' (1), a refresh dropdown (1), a 'Content' dropdown (2), a search bar (3), a 'Severity' dropdown (3), a 'Log Type' dropdown set to 'infrastructure' (4), and a 'Show Resources' link (4). Below these is a red box labeled (5) containing a LogQL query: '{ log_type="infrastructure" } | json'. To the right of the query is a 'Run Query' button. The main area displays log entries with columns for 'Date' and 'Message'. The first entry is: 'Jan 22, 2024, 10:09:11.750 I0122 15:09:11.713637 1 log.go:194] reconciling (monitoring.coreos.com/v1, Kind=ServiceMonitor) openshift-multus/monitor-network'. The second entry is: 'Jan 22, 2024, 10:09:11.750 I0122 15:09:11.713489 1 log.go:194] Apply / Create of (apps/v1, Kind=DaemonSet) openshift-multus/network-metrics-daemon was successful'. The third entry is: 'Jan 22, 2024, 10:09:11.618 I0122 15:09:11.510221 1 log.go:194] reconciling (apps/v1, Kind=DaemonSet) openshift-multus/network-metrics-daemon'.

- ➊ Time range for the logs.
- ➋ Filter the logs by content, namespace, pod, or container.
- ➌ Log severity. Possible values are `critical`, `error`, `warning`, `debug`, `info`, `trace`, and `unknown`.
- ➍ Log type. Possible values are `application`, `infrastructure`, or `audit`.
- ➎ The log query in LogQL query language. You can show this field by clicking **Show Query**.

LogQL Queries

Loki organizes logs into streams that are called *chunks*. A *chunk* is a sequence of log entries that use the same set of labels. Labels are key-value pairs that help identify and filter log entries. Labels can include information such as the application name, environment, severity, or any other relevant metadata. Loki uses LogQL as the query language for filtering and searching logs with labels and other criteria.

By using LogQL in the OpenShift web console, you can extract and filter specific information from the logs.

A LogQL query consists of the following parameters:

- The log stream selector, which accepts the following operators: `=` equals, `!=` not equals, `=~` regex pattern matches, and `!~` regex pattern does not match.
- A filter expression, which accepts the following operators: `|=` equals, `!=` not equals, `|~` regex pattern matches, and `!~` regex pattern does not match.

When using LogQL in OpenShift Logging, you must always use at least the `log_type` log stream selector for the log type.

For example, the following LogQL query shows the infrastructure logs:

```
log_type="infrastructure" | json
```

Chapter 7 | OpenShift Logging

Adding the `| json` string to your query extracts all JSON properties as labels for ease of viewing. After you type the LogQL query, click **Run Query** to apply it to your logs.

By using LogQL you can filter your logs with more advanced expressions. For example, the following LogQL query filters the audit logs only from `audit.log` files in any directory:

```
log_type="audit" | json | file=~".*audit.log"
```

You can use multiple log stream selectors and filter expressions to increase filter granularity. The following example shows a LogQL query for the audit logs that contain the `error` string but not the `warning` string, and that return a 403 or 503 status.

```
log_type="audit" | json |= "error" != "warning" |~ "status [45]03"
```



Note

For more information about querying Loki by using the LogQL query language, refer to <https://grafana.com/docs/loki/latest/logql/>

Vector adds OpenShift metadata to pod logs to provide fields for querying and filtering. You can use these fields in your LogQL queries for more precise filtering.

Some fields that you can use in your queries include the following strings:

hostname

The hostname of the OpenShift node for the pod that generates the log.

kubernetes_container_name

The name of the container that generates the log.

kubernetes_namespace_name

The name of the project that contains the pod that generates the log.

kubernetes_pod_name

The name of the pod that generates the log.

level

The log level.

message

The log message.

Log Access Permissions

By default, the OpenShift Logging operator grants only administrator users access to the logs. Thus, regular users cannot access the logs. As an administrator, you can configure access to the logs for users and groups by using RBAC rules.

The OpenShift Logging operator provides the following cluster roles to simplify RBAC rules:

cluster-logging-application-view

Grants read permission to application logs.

cluster-logging-infrastructure-view

Grants read permission to infrastructure logs.

cluster-logging-audit-view
Grants read permission to audit logs.



References

For more information about logging on OpenShift, refer to the *About Logging* section in the Red Hat OpenShift Container Platform 4.14 *Logging* documentation at

https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/logging/index#cluster-logging

For more information about installing the Loki operator, refer to the *Installing Log Storage* section in the Red Hat OpenShift Container Platform 4.14 *Logging* documentation at

https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/logging/index#installing-log-storage

For more information about Loki limits for log streams, refer to the *LimitsSpec* section in the *Loki Operator API* documentation at

<https://loki-operator.dev/docs/api.md/#loki-grafana-com-v1-LimitsSpec>

Setting a Lifecycle Configuration on a Bucket

<https://docs.aws.amazon.com/AmazonS3/latest/userguide/how-to-set-lifecycle-configuration-intro.html>

For more information about lifecycle bucket configuration, refer to the *Lifecycle Bucket Configuration in Multicloud Object Gateway* section in the Red Hat OpenShift Data Foundation 4.14 *Managing Hybrid and Multicloud Resources* documentation at

https://access.redhat.com/documentation/en-us/red_hat_openshift_data.foundation/4.14/html-single/managing_hybrid_and_multicloud_resources/index#con.lifecycle.bucket-configuration-in-multicloud-object-gateway_rhodf

Knowledgebase – How to Configure Lifecycle Policy in MCG

<https://access.redhat.com/solutions/7058004>

LokiStack

<https://loki-operator.dev/docs/api.md/#loki-grafana-com-v1-LokiStack>

For more information about the Loki deployment size, refer to the *Installing Log Storage* section in the Red Hat OpenShift Container Platform 4.14 *Log Storage* documentation at

https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/logging/index#loki-deployment-sizing_installing-log-storage

For more information about the secret type values, refer to the *Loki Object Storage* section in the Red Hat OpenShift Container Platform 4.14 *Log Storage* documentation at

https://access.redhat.com/documentation/en-us/openshift_container_platform/4.14/html-single/logging/index#logging-loki-storage_installing-log-storage

► Guided Exercise

Centralized Logging

Deploy OpenShift Logging for short-term log retention and aggregation.

Outcomes

- Use Loki as the log store for OpenShift Logging.
- Use Vector as the collector and the OpenShift web UI for log visualization.

Before You Begin

As the student user on the workstation machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start logging-central
```

Instructions

Your company requires you to configure the deployed OpenShift Logging operator for short-term log retention and aggregation. You must configure OpenShift Logging by using Vector, Loki, and the OpenShift web console. Moreover, you must apply the `cluster-logging-application-view` cluster role to the `ocpdevs` group, so the `developer` user can retrieve application logs for the `testing-logs` project.

- 1. As the OpenShift `admin` user, install the Loki operator in the `openshift-operators-redhat` namespace.
 - 1.1. Open a web browser and navigate to <https://console-openshift-console.apps.ocp4.example.com>.
 - 1.2. Click **Red Hat Identity Management** and log in as the `admin` user with `redhatocp` as the password.
 - 1.3. Navigate to **Operators > OperatorHub**.
 - 1.4. Click **Loki Operator**, and then click **Install**.
 - 1.5. Select **Enable Operator recommended cluster monitoring on this Namespace** and click **Install**. The Operator Lifecycle Manager (OLM) can take a few minutes to install the operator. Click **View Operator** to navigate to the operator details.
- 2. Create an S3-compatible object storage bucket with OpenShift Data Foundation for the Loki operator. Retrieve the credentials for the bucket, and create a secret with those credentials.
 - 2.1. From the terminal, log in to your OpenShift cluster as the `admin` user.

```
[student@workstation ~]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
Login successful.
...output omitted...
```

2.2. Change to the ~/D0380/labs/logging-central directory.

```
[student@workstation ~]$ cd ~/D0380/labs/logging-central
```

2.3. Create an ObjectBucketClaim resource YAML file for a bucket called loki-bucket-odf in the openshift-logging namespace. The Loki operator uses this bucket. You can find an incomplete example for the resource in the ~/D0380/labs/logging-central/objectbucket.yaml file.

```
apiVersion: objectbucket.io/v1alpha1
kind: ObjectBucketClaim
metadata:
  name: loki-bucket-odf
  namespace: openshift-logging
spec:
  generateBucketName: loki-bucket-odf
  storageClassName: openshift-storage.noobaa.io
```

2.4. Create the ObjectBucketClaim resource.

```
[student@workstation logging-central]$ oc create -f objectbucket.yaml
objectbucketclaim.objectbucket.io/loki-bucket-odf created
```

2.5. Verify that the object bucket claim is created and in the Bound phase.

```
[student@workstation logging-central]$ oc get obc -n openshift-logging
NAME           STORAGE-CLASS          PHASE   AGE
loki-bucket-odf  openshift-storage.noobaa.io  Bound   12m
```

2.6. Retrieve the S3 bucket information and credentials, and store them in environment variables. When an object bucket claim is created, OpenShift Data Foundation creates a secret and a configuration map with the same name as for the bucket information and credentials. The bucket credentials would differ on your system.

```
[student@workstation logging-central]$ BUCKET_HOST=$(oc get -n openshift-logging \
configmap loki-bucket-odf -o jsonpath='{.data.BUCKET_HOST}'); \
BUCKET_NAME=$(oc get -n openshift-logging configmap loki-bucket-odf \
-o jsonpath='{.data.BUCKET_NAME}'); \
BUCKET_PORT=$(oc get -n openshift-logging configmap loki-bucket-odf \
-o jsonpath='{.data.BUCKET_PORT}'); \
ACCESS_KEY_ID=$(oc get -n openshift-logging secret loki-bucket-odf \
-o jsonpath='{.data.AWS_ACCESS_KEY_ID}' | base64 -d); \
SECRET_ACCESS_KEY=$(oc get -n openshift-logging secret loki-bucket-odf \
-o jsonpath='{.data.AWS_SECRET_ACCESS_KEY}' | base64 -d)
```

Chapter 7 | OpenShift Logging

- 2.7. Create a secret called `logging-loki-odf` in the `openshift-logging` namespace with the bucket credentials.

```
[student@workstation logging-central]$ oc create secret generic logging-loki-odf \
-n openshift-logging \
--from-literal=access_key_id=${ACCESS_KEY_ID} \
--from-literal=access_key_secret=${SECRET_ACCESS_KEY} \
--from-literal=bucketnames=${BUCKET_NAME} \
--from-literal=endpoint=https://${BUCKET_HOST}:${BUCKET_PORT}
secret/logging-loki-odf created
```

- 3. Create a LokiStack instance by using the bucket as the storage.

- 3.1. Create a LokiStack resource YAML file for an instance called `logging-loki` in the `openshift-logging` namespace. This instance uses the bucket as the storage. You can find an incomplete example for the resource in the `~/D0380/labs/logging-central/lokistack.yaml` file.

```
apiVersion: loki.grafana.com/v1
kind: LokiStack
metadata:
  name: logging-loki
  namespace: openshift-logging
spec:
  size: 1x.demo
  storage:
    secret:
      name: logging-loki-odf
      type: s3
    tls:
      caName: openshift-service-ca.crt
  storageClassName: ocs-external-storagecluster-ceph-rbd
  tenants:
    mode: openshift-logging
```

- 3.2. Create the LokiStack resource.

```
[student@workstation logging-central]$ oc create -f lokistack.yaml
lokistack.loki.grafana.com/logging-loki created
```

- 3.3. Verify that the LokiStack pods are up and running.

```
[student@workstation logging-central]$ oc get pods -n openshift-logging
NAME                               READY   STATUS    RESTARTS   AGE
cluster-logging-operator-554849f7dd-9tcz2   1/1     Running   0          20m
logging-loki-compactor-0                  1/1     Running   0          98s
logging-loki-distributor-64c798c4c5-6wpxp   1/1     Running   0          99s
logging-loki-gateway-68fb59cdf5-6b8mt      2/2     Running   0          98s
logging-loki-gateway-68fb59cdf5-7qf4s      2/2     Running   0          98s
logging-loki-index-gateway-0                1/1     Running   0          98s
```

Chapter 7 | OpenShift Logging

logging-loki-ingester-0	1/1	Running	0	99s
logging-loki-querier-577b55f8d5-f4cfb	1/1	Running	0	98s
logging-loki-query-frontend-775755684d-f94bd	1/1	Running	0	98s

- ▶ 4. Create a `ClusterLogging` instance by using `loki` as the log store, `vector` as the collector, and the `ocp-console` as the visualization console.
- 4.1. Create a `ClusterLogging` resource YAML file by using `loki` as the log store, `vector` as the collector, and the `ocp-console` as the visualization console. You can find an incomplete example for the resource in the `~/D0380/labs/logging-central/clusterlogging.yaml` file.

```
apiVersion: logging.openshift.io/v1
kind: ClusterLogging
metadata:
  name: instance
  namespace: openshift-logging
spec:
  managementState: Managed
  logStore:
    type: lokistack
    lokistack:
      name: logging-loki
  collection:
    type: vector
  visualization:
    type: ocp-console
```

- 4.2. Create the `ClusterLogging` resource.

```
[student@workstation logging-central]$ oc create -f clusterlogging.yaml
clusterlogging.logging.openshift.io/instance created
```

- 4.3. Verify that the `ClusterLogging` pods are up and running.

```
[student@workstation logging-central]$ oc get pods -n openshift-logging
NAME                               READY   STATUS    RESTARTS   AGE
cluster-logging-operator-554849f7dd-9tcz2   1/1    Running   0          24m
collector-7rb2n                      1/1    Running   0          93s
collector-bkj8x                      1/1    Running   0          93s
collector-cng5z                      1/1    Running   0          93s
collector-hx2gd                      1/1    Running   0          93s
collector-x92zq                      1/1    Running   0          93s
collector-xlqw9                      1/1    Running   0          93s
...output omitted...
logging-view-plugin-5b9b5b7bdc-tvkqk     1/1    Running   0          94s
```

- ▶ 5. Enable the console plug-in for the OpenShift Logging operator. Verify that you have access to the logs.
- 5.1. Change to the web console browser. Click **Operators > Installed Operators**, and select **All Projects** from the drop-down menu.

- 5.2. Click Red Hat OpenShift Logging, click **Console plugin**, select **Enable**, and click **Save**.
- 5.3. Reload the web console, and navigate to **Observe > Logs**. If the **Observe > Logs** menu is not available, then wait until the web console shows the **Web console update is available** message and reload the web console. You have access to the logs for the application and infrastructure resources. By default, the **ClusterLogging** instance includes logs for the application and infrastructure, but not the audit logs. Observe the application logs, which are selected by default.
- 5.4. From the drop-down menu, select the infrastructure logs and observe them.

Date	Message
Jan 19, 2024, 04:34:15.260	run-runc-68862018b107a95a2979522d1e3ff24b493579a14b2b03a4c7f96bd39 cc9b52a-runc.q4gbLP.mount: Deactivated successfully.
Jan 19, 2024, 04:34:14.456	[32mJan-19 9:34:14.310[35m [WebServer/38] [36m [L0][39m core.servr.node_services.nodes_store:: bulk_update: success { nInserted: [3 3m0[39m, nModified: [33m1[39m }
Jan 19, 2024, 04:34:14.456	[32mJan-19 9:34:14.303[35m [WebServer/38] [36m [L0][39m core.servr.node_services.nodes_store:: bulk_update: executing bulk with [33 m1[39m updates and [33m0[39m inserts
Jan 19, 2024, 04:34:14.456	[32mJan-19 9:34:14.302[35m [HostedAgents/35] [36m [L0][39m core.rp c.ice:: TLS CLOSED: tcp://[10.9.2.15]:60100=>tcp://[::ffff:10.9.2.15]:55506
Jan 19, 2024, 04:34:14.456	[32mJan-19 9:34:14.302[35m [HostedAgents/35] [36m [L0][39m core.rp

- 5.5. From the drop-down menu, select the audit logs and observe the **No datapoints found** message. You receive this message because the **ClusterLogging** instance does not forward the audit logs.
6. Include the audit logs by creating a log forwarder for the application, infrastructure, and audit logs to the **LokiStack** resource.

- 6.1. Change to the terminal window, and create an **ClusterLogForwarder** resource YAML file for a log forwarder called **instance** in the **openshift-logging** namespace. The log forwarder must forward the application, infrastructure, and audit logs. You can find an incomplete example for the resource in the **~/DO380/labs/logging-central/forwarder.yaml** file.

```

apiVersion: logging.openshift.io/v1
kind: ClusterLogForwarder
metadata:
  name: instance
  namespace: openshift-logging
spec:

```

```
pipelines:
- name: all-to-default
  inputRefs:
  - infrastructure
  - application
  - audit
  outputRefs:
  - default
```

- 6.2. Create the `ClusterLogForwarder` resource.

```
[student@workstation logging-central]$ oc create -f forwarder.yaml
clusterlogforwarder.logging.openshift.io/instance created
```

- 6.3. Change to the web console browser and reload it. You have access to the audit logs.

- ▶ 7. Verify that the infrastructure and audit logs are sent to Loki, by trying to open an SSH connection to one of the compute nodes.

- 7.1. Change to the terminal window. Try to open an SSH connection to the `worker01.ocp4.example.com` node, which is rejected.

```
[student@workstation logging-central]$ ssh worker01.ocp4.example.com
Warning: Permanently added 'worker01.ocp4.example.com' (ED25519) to the list of
known hosts.
student@worker01.ocp4.example.com: Permission denied (publickey,gssapi-
keyex,gssapi-with-mic).
```

- 7.2. Change to the web console browser, and click the refresh button in the upper-right corner to update the logs.

- 7.3. Ensure that the audit logs are selected in the drop-down menu. Search for SSH connection messages by typing `ssh` in the `Search by Content` field and clicking `Run Query`.

- 7.4. Unfold the information for the first audit log, which shows an attempt to log to the `worker01` node with a failed result.

- 7.5. From the drop-down menu, select the `infrastructure` logs. Click `Clear all filters` and click `Show Query`.

- 7.6. Modify the query to filter logs only from the sshd service. The query should read as follows:

```
{ log_type="infrastructure" } | json | systemd_u_SYSLOG_IDENTIFIER="sshd"
```

- 7.7. Click Run Query and verify that you receive logs from the sshd service.
- 8. Create a project and a pod as the developer user and verify that this user has access to the pod logs.

- 8.1. Change to the terminal window and log in to the OpenShift cluster as the developer user with developer as the password.

```
[student@workstation logging-central]$ oc login -u developer -p developer
Login successful.
...output omitted...
```

- 8.2. Create the testing-logs project.

```
[student@workstation logging-central]$ oc new-project testing-logs
Now using project "testing-logs" on server "https://api.ocp4.example.com:6443".
...output omitted...
```

- 8.3. Create a test-date pod.

```
[student@workstation logging-central]$ oc run test-date --restart 'Never' \
--image registry.ocp4.example.com:8443/ubi9/ubi -- date
pod/test-date created
```

- 8.4. Verify that the pod is in the Completed status.

```
[student@workstation logging-central]$ oc get pods
NAME      READY   STATUS    RESTARTS   AGE
test-date  0/1     Completed  0          56s
```

- 8.5. Review the logs for the pod by using the terminal.

```
[student@workstation logging-central]$ oc logs test-date
Tue Jan 16 11:15:15 UTC 2024
```

- 8.6. Delete the pod.

```
[student@workstation logging-central]$ oc delete pod test-date
pod "test-date" deleted
```

- 8.7. Try to retrieve the logs for the pod, and verify that the developer user no longer has access to the pod logs.

```
[student@workstation logging-central]$ oc logs test-date
Error from server (NotFound): pods "test-date" not found
```

► **9.** Log in as the developer user and verify that the user has no access to the logs.

- 9.1. Change to the browser window, open a private window, and navigate to <https://console-openshift-console.apps.ocp4.example.com>.
- 9.2. Click **Red Hat Identity Management** and log in as the developer user with developer as the password. Click **Skip tour**.
- 9.3. Navigate to **Observe**. Verify that the **testing-logs** project is selected. Change to the **Logs** tab. The developer user has no permissions to retrieve the pod logs.

The screenshot shows the Red Hat OpenShift web console interface. At the top, there's a navigation bar with the Red Hat OpenShift logo and a dropdown menu for 'developer'. Below the header, a 'Project: testing-logs' dropdown is visible. The main area is titled 'Observe' and has tabs for 'Dashboard', 'Metrics', 'Events', and 'Logs', with 'Logs' being the active tab. Above the log table, there are filters for 'Content' (with a search bar), 'Severity' (set to 'Info'), and buttons for 'Show Resources', 'Run Query', and 'Show Query'. Below these are buttons for 'Date' (sorted by 'Date') and 'Message'. A single log entry is shown with a red 'Forbidden' icon and the text 'Missing permissions to get logs'. Below the log entry, a note says 'Make sure you have the required role to get application logs in this namespace.' and 'Ask your administrator to grant you this role.'

- 10. Review the `test-date` pod logs by using the OpenShift logging operator. The `admin` user has access to the stored pod logs.

- 10.1. Change to the web console browser where the `admin` user is logged in. From the drop-down menu where the infrastructure logs are currently selected, select the application logs. Click the refresh button in the upper-right corner to update the logs.
- 10.2. Modify the query to filter the results for the `testing-logs` namespace. The query should read as follows:

```
{ log_type="application", kubernetes_namespace_name="testing-logs" } | json
```

- 10.3. Click **Run Query**. Retrieve the information for the only entry, which shows the logs from the deleted pod.

- 11. Give the `ocpdevs` group permission to view the logs in the `test-logs` project, and verify that the `developer` user has access to the logs. Give the permission to the `ocpdevs` group by assigning it the `cluster-logging-application-view` cluster role.

- 11.1. Change to the terminal window and log in to the OpenShift cluster as the `admin` user with `redhatocp` as the password.

```
[student@workstation logging-central]$ oc login -u admin -p redhatocp
Login successful.
...output omitted...
```

- 11.2. Review the required role to provide access to the application logs to the `ocpdevs` group. You can find an example in the `~/D0380/labs/logging-central/ocpdevs-role.yaml` file.

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: view-application-logs
  namespace: testing-logs
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: cluster-logging-application-view
subjects:
- kind: Group
  name: ocpdevs
  apiGroup: rbac.authorization.k8s.io
```

- 11.3. Apply the role to the ocpdevs group.

```
[student@workstation logging-central]$ oc create -f ocpdevs-role.yaml
rolebinding.rbac.authorization.k8s.io/view-application-logs created
```

- 11.4. Change to the web console in the private browser where the developer user is logged in, and refresh it.
- 11.5. Verify that the developer user has access to the deleted pod logs.
- 11.6. Close both the web browser windows and change to the student HOME directory in the terminal window.

```
[student@workstation logging-central]$ cd
```

Finish

On the workstation machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish logging-central
```

▶ Lab

OpenShift Logging

Configure OpenShift Logging for short-term and long-term log retention and aggregation.

Outcomes

- Configure OpenShift Logging to forward logs to an external aggregator for long-term storage.
- Configure OpenShift Logging with Loki for short-term log retention.
- Configure OpenShift Logging to collect logs from specific applications.

Before You Begin

As the student user on the `workstation` machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start logging-review
```

Instructions

The security policy of your company requires you to send OpenShift audit logs to a dedicated syslog server for long-term storage.

The developer team asks you to provide them access to the production application logs and to the CI job logs that are running in the `build-ci` namespace. To limit the footprint of the log storage, the developers agree to keep only seven days of logs for the `build-ci` namespace.

Configure OpenShift Logging to forward only audit logs to the syslog server. The syslog server DNS name is `utility.lab.example.com` and the service listens on the TCP port 514. Forwarded audit logs are stored in the `/var/log/openshift/audit.log` file on the `utility` machine.

Configure Loki as the internal log store for both infrastructure logs for the cluster administrators and application logs for the developers.

An S3 bucket is available for you, in the lab environment, to configure as log storage for Loki. The bucket information and credentials are available in the `~/D0380/labs/logging-review/s3bucket.env` file on the `workstation` machine.

Ensure that only application logs with the `environment=production` label or in the `build-ci` namespace are collected. Ensure that audit logs are not stored in the internal log store.

Configure the global log retention in Loki to 30 days. Configure the log retention for pods that are running in the `build-ci` namespace to seven days.

Provide access to the application logs to the `ocpdevs` group. You can use the `developer` user, which is in the `ocpdevs` group, to verify that the permissions are correct.

Chapter 7 | OpenShift Logging

Use the following applications in the cluster to verify that the logging configuration is working as expected:

Label	Projects
environment=production	<ul style="list-style-type: none"> shop-prod supportbot-prod
environment=staging	<ul style="list-style-type: none"> shop-stg supportbot-stg
environment=development	<ul style="list-style-type: none"> shop-dev supportbot-dev build-ci

1. Create a secret for Loki with the object storage credentials from the `~/D0380/labs/logging-review/s3bucket.env` environment file.
2. Create and configure a `logging-loki` LokiStack instance to use the S3 bucket, and set the log retention for the `build-ci` namespace.
You can use the partial resource definition in the `~/D0380/labs/logging-review/lokistack.yml` file.
3. Configure OpenShift Logging to forward audit logs to the syslog server, and to forward the infrastructure and application logs to the internal log store.
Ensure that only application logs with the `environment: production` label or in the `build-ci` namespace are collected. Ensure that audit logs are not stored in the Loki instance.
You can use the partial resource definition in the `~/D0380/labs/logging-review/clusterlogforwarder.yml` file.
4. Configure OpenShift Logging to deploy Vector and the web console plug-in, and use the Loki instance that you configure in a previous step as the log store.
You can use the partial resource definition in the `~/D0380/labs/logging-review/clusterlogging.yml` file.
5. Connect to the `utility` server and verify that the syslog service receives the audit logs. Audit logs are stored in the `/var/log/openshift/audit.log` file on the `utility` machine.
6. Enable the web console plug-in for the OpenShift Logging operator and verify that infrastructure and application logs are available.
7. Grant view access to the application logs to the `ocpdevs` group.
8. Verify that the `developer` user can access the application logs from the web console.
Verify that only application logs with the `environment: production` label or in the `build-ci` namespace are collected.

Evaluation

As the student user on the `workstation` machine, use the `lab` command to grade your work. Correct any reported failures and rerun the command until successful.

```
[student@workstation ~]$ lab grade logging-review
```

Finish

As the student user on the workstation machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish logging-review
```

► Solution

OpenShift Logging

Configure OpenShift Logging for short-term and long-term log retention and aggregation.

Outcomes

- Configure OpenShift Logging to forward logs to an external aggregator for long-term storage.
- Configure OpenShift Logging with Loki for short-term log retention.
- Configure OpenShift Logging to collect logs from specific applications.

Before You Begin

As the student user on the `workstation` machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start logging-review
```

Instructions

The security policy of your company requires you to send OpenShift audit logs to a dedicated syslog server for long-term storage.

The developer team asks you to provide them access to the production application logs and to the CI job logs that are running in the `build-ci` namespace. To limit the footprint of the log storage, the developers agree to keep only seven days of logs for the `build-ci` namespace.

Configure OpenShift Logging to forward only audit logs to the syslog server. The syslog server DNS name is `utility.lab.example.com` and the service listens on the TCP port 514. Forwarded audit logs are stored in the `/var/log/openshift/audit.log` file on the `utility` machine.

Configure Loki as the internal log store for both infrastructure logs for the cluster administrators and application logs for the developers.

An S3 bucket is available for you, in the lab environment, to configure as log storage for Loki. The bucket information and credentials are available in the `~/D0380/labs/logging-review/s3bucket.env` file on the `workstation` machine.

Ensure that only application logs with the `environment=production` label or in the `build-ci` namespace are collected. Ensure that audit logs are not stored in the internal log store.

Configure the global log retention in Loki to 30 days. Configure the log retention for pods that are running in the `build-ci` namespace to seven days.

Provide access to the application logs to the `ocpdevs` group. You can use the `developer` user, which is in the `ocpdevs` group, to verify that the permissions are correct.

Chapter 7 | OpenShift Logging

Use the following applications in the cluster to verify that the logging configuration is working as expected:

Label	Projects
environment=production	<ul style="list-style-type: none"> shop-prod supportbot-prod
environment=staging	<ul style="list-style-type: none"> shop-stg supportbot-stg
environment=development	<ul style="list-style-type: none"> shop-dev supportbot-dev build-ci

1. Create a secret for Loki with the object storage credentials from the ~/D0380/labs/logging-review/s3bucket.env environment file.
 - 1.1. Connect to the OpenShift cluster as the `admin` user with `redhatocp` as the password.

```
[student@workstation ~]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
Login successful.
...output omitted...
```

- 1.2. Change to the ~/D0380/labs/logging-review directory.

```
[student@workstation ~]$ cd ~/D0380/labs/logging-review
```

- 1.3. Change to the openshift-logging project.

```
[student@workstation logging-review]$ oc project openshift-logging
Now using project "openshift-logging" on server "https://
api.ocp4.example.com:6443".
```

- 1.4. Use the ~/D0380/labs/logging-review/s3bucket.env environment file to create the `logging-loki-odf` secret in the `openshift-logging` namespace.

```
[student@workstation logging-review]$ oc create secret generic logging-loki-odf \
-n openshift-logging --from-env-file=s3bucket.env
secret/logging-loki-odf created
```

2. Create and configure a `logging-loki` LokiStack instance to use the S3 bucket, and set the log retention for the `build-ci` namespace.

You can use the partial resource definition in the ~/D0380/labs/logging-review/lokistack.yml file.

- 2.1. Modify the partial resource definition in the ~/D0380/labs/logging-review/lokistack.yml file as follows:

```
apiVersion: loki.grafana.com/v1
kind: LokiStack
metadata:
```

```

name: logging-loki
namespace: openshift-logging
spec:
  limits:
    global:
      retention:
        days: 30
      streams:
        - selector: '{kubernetes_namespace_name="build-ci"}'
          priority: 1
          days: 7
  size: 1x.demo
  storage:
    tls:
      caName: openshift-service-ca.crt
    secret:
      name: logging-loki-odf
      type: s3
  storageClassName: ocs-external-storagecluster-ceph-rbd
  tenants:
    mode: openshift-logging

```

2.2. Create the LokiStack resource.

```
[student@workstation logging-review]$ oc create -f lokistack.yml
lokistack.loki.grafana.com/logging-loki created
```

2.3. Verify that the LokiStack pods are up and running.

```
[student@workstation logging-review]$ oc get deployment,statefulset \
-l app.kubernetes.io/name=lokistack
NAME                                         READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/logging-loki-distributor     1/1     1           1           86s
deployment.apps/logging-loki-gateway         2/2     2           2           85s
deployment.apps/logging-loki-querier         1/1     1           1           85s
deployment.apps/logging-loki-query-frontend   1/1     1           1           85s

NAME                                         READY   AGE
statefulset.apps/logging-loki-compactor      1/1     85s
statefulset.apps/logging-loki-index-gateway   1/1     85s
statefulset.apps/logging-loki-ingester        1/1     86s
```

3. Configure OpenShift Logging to forward audit logs to the syslog server, and to forward the infrastructure and application logs to the internal log store.

Ensure that only application logs with the `environment: production` label or in the `build-ci` namespace are collected. Ensure that audit logs are not stored in the Loki instance.

You can use the partial resource definition in the `~/D0380/labs/logging-review/clusterlogforwarder.yml` file.

- 3.1. Modify the partial resource definition in the `~/D0380/labs/logging-review/clusterlogforwarder.yml` file as follows:

```

apiVersion: logging.openshift.io/v1
kind: ClusterLogForwarder
metadata:
  name: instance
  namespace: openshift-logging
spec:
  inputs:
    - name: production-apps
      application:
        selector:
          matchLabels:
            environment: production
    - name: ci
      application:
        namespaces:
          - build-ci
  outputs:
    - name: audit-syslog
      type: syslog
      url: tcp://utility.lab.example.com:514
      syslog:
        msgID: audit
        appName: ocp-lab
        facility: user
        procID: vector
        rfc: RFC5424
        severity: informational
  pipelines:
    - name: to-syslog
      inputRefs:
        - audit
      outputRefs:
        - audit-syslog
    - name: to-loki
      inputRefs:
        - infrastructure
        - ci
        - production-apps
      outputRefs:
        - default

```

3.2. Apply the configuration for the cluster log forwarder.

```
[student@workstation logging-review]$ oc apply -f clusterlogforwarder.yaml
clusterlogforwarder.logging.openshift.io/instance created
```

4. Configure OpenShift Logging to deploy Vector and the web console plug-in, and use the Loki instance that you configure in a previous step as the log store.

You can use the partial resource definition in the ~/D0380/labs/logging-review/clusterlogging.yaml file.

- 4.1. Modify the partial resource definition in the ~/D0380/labs/logging-review/clusterlogging.yaml file as follows:

```

apiVersion: logging.openshift.io/v1
kind: ClusterLogging
metadata:
  name: instance
  namespace: openshift-logging
spec:
  managementState: Managed
  logStore:
    type: lokistack
    lokistack:
      name: logging-loki
  collection:
    type: vector
  visualization:
    type: ocp-console

```

4.2. Apply the cluster logging configuration.

```
[student@workstation logging-review]$ oc apply -f clusterlogging.yml
clusterlogging.logging.openshift.io/instance created
```

4.3. Verify that the OpenShift Logging operator deploys the collector pod on each node.

```
[student@workstation logging-review]$ oc get daemonset -l component=collector
NAME        DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE   ...
collector   6          6          6          6          6          ...
```

5. Connect to the **utility** server and verify that the syslog service receives the audit logs. Audit logs are stored in the `/var/log/openshift/audit.log` file on the **utility** machine.

5.1. Connect to the **utility** machine with SSH as the **root** user.

```
[student@workstation logging-review]$ ssh root@utility
```

5.2. Check that the OpenShift audit log file exists in the `/var/log/openshift` path.

```
[root@utility ~]# ls -l /var/log/openshift
total 978112
-rw----- 1 root root 462107806 Jan 30 10:10 audit.log
```

5.3. Review the content of the audit log.

```
[root@utility ~]# tail -1 /var/log/openshift/audit.log
"2024-01-30T10:10:25.618+00:00 ...output omitted..."
```

5.4. Disconnect from the **utility** machine.

```
[root@utility ~]# exit
logout
Connection to utility closed.
[student@workstation logging-review]$
```

6. Enable the web console plug-in for the OpenShift Logging operator and verify that infrastructure and application logs are available.
 - 6.1. Open a web browser and navigate to <https://console-openshift-console.apps.ocp4.example.com>. Click Red Hat Identity Management and log in as the admin user with redhatocp as the password.
 - 6.2. Click Operators > Installed Operators, and select All Projects from the drop-down menu.
 - 6.3. Click Red Hat OpenShift Logging, click Console plugin, select Enable, and click Save.
 - 6.4. Reload the web console, and navigate to Observe > Logs. If the Observe > Logs menu is not available, then wait until the web console shows the Web console update is available message and reload the web console.
Verify that the application logs are available.
 - 6.5. From the drop-down menu, select infrastructure and verify that the infrastructure logs are available.

Date	Message
Jan 19, 2024, 04:34:15.260	run-runc-68862018b107a95a2979522d1e3ff24b493579a14b2b03a4c7f96bd39 cc9b52a-runc.q4gbLP.mount: Deactivated successfully.
Jan 19, 2024, 04:34:14.456	[32mJan-19 9:34:14.310[35m [WebServer/38] [36m [L0][39m core.serve r.node_services.nodes_store:: bulk_update: success { nInserted: [3 3m0[39m, nModified: [33m1[39m }]
Jan 19, 2024, 04:34:14.456	[32mJan-19 9:34:14.303[35m [WebServer/38] [36m [L0][39m core.serve r.node_services.nodes_store:: bulk_update: executing bulk with [33 m1[39m updates and [33m0[39m inserts
Jan 19, 2024, 04:34:14.456	[32mJan-19 9:34:14.302[35m [HostedAgents/35] [36m [L0][39m core.rp c.ice:: TLS CLOSED: tcp://[10.9.2.15]:60100=>tcp://[::ffff:10.9.2. 15]:55506
Jan 19, 2024, 04:34:14.456	[32mJan-19 9:34:14.302[35m [HostedAgents/35] [36m [L0][39m core.rp

- 6.6. From the drop-down menu, select audit and confirm that no audit logs are stored in the log store.
7. Grant view access to the application logs to the ocpdevs group.
 - 7.1. Add the cluster-logging-application-view role to the ocpdevs group.

```
[student@workstation logging-review]$ oc adm policy add-cluster-role-to-group \
cluster-logging-application-view ocpdevs
clusterrole.rbac.authorization.k8s.io/cluster-logging-application-view added:
"ocpdevs"
```

8. Verify that the developer user can access the application logs from the web console.
Verify that only application logs with the environment: production label or in the build-ci namespace are collected.
 - 8.1. Open a new private browser window, and navigate to <https://console-openshift-console.apps.ocp4.example.com>.
 - 8.2. Click Red Hat Identity Management and log in as the developer user with developer as the password. Click Skip tour.
 - 8.3. Navigate to Observe and select build-ci from the project drop-down menu. Change to the Logs tab. Verify that the application logs are available.
 - 8.4. Change to the shop-prod project and verify that the application logs are available.
 - 8.5. Change to the shop-dev project, and verify that no application logs are available, because the application does not have the required label.
 - 8.6. Close both the web browser windows and change to the student HOME directory. in the terminal window.

```
[student@workstation logging-review]$ cd
```

Evaluation

As the student user on the workstation machine, use the lab command to grade your work. Correct any reported failures and rerun the command until successful.

```
[student@workstation ~]$ lab grade logging-review
```

Finish

As the student user on the workstation machine, use the lab command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish logging-review
```

Summary

- By default, Kubernetes stores logs from pods in the local disk on the nodes, and the information is lost when the pod is deleted.
- The OpenShift Logging operator collects and aggregates the log messages in your cluster, to keep access to the logs.
- OpenShift Logging is based on the following components:
 - A collector, such as Vector, to collect logs from all running containers and cluster nodes.
 - A log store, such as Grafana Loki, to aggregate logs from the entire cluster into a central place and to provide access control to logs.
 - A visualization console, such as the OpenShift Logging UI, to view and query logs in the internal log store.
- You can configure OpenShift Logging to forward logs to a third-party logging system, or use an internal log store for short-term log retention.
- Although by default, the OpenShift Logging operator includes logs for the infrastructure and the application, you must configure a log forwarder to include audit logs in the log store.
- You can deploy the Event Router component in OpenShift Logging to log Kubernetes events.
- Loki uses LogQL as the query language for filtering and searching logs.

Chapter 8

Comprehensive Review

Goal

Review tasks from *Red Hat OpenShift Administration III: Scaling Deployments in the Enterprise*.

Sections

- Comprehensive Review

Lab

- Cluster Administration

Comprehensive Review

Objectives

After completing this section, you should have reviewed and refreshed the knowledge and skills that you learned in *Red Hat OpenShift Administration III: Scaling Deployments in the Enterprise*.

Reviewing Red Hat OpenShift Administration III: Scaling Deployments in the Enterprise

Before beginning the comprehensive review for this course, you should be comfortable with the topics covered in each chapter. Do not hesitate to ask the instructor for extra guidance or clarification on these topics.

Chapter 1, Authentication and Identity Management

Configure OpenShift clusters to authenticate by using LDAP and OIDC enterprise identity systems and to recognize groups that those systems define.

Chapter 2, Backup, Restore, and Migration of Applications with OADP

Back up and restore application settings and data with OpenShift API for Data Protection (OADP).

Chapter 3, Cluster Partitioning

Configure a subset of cluster nodes to be dedicated to a type of workload.

Chapter 4, Pod Scheduling

Configure workloads to run on a dedicated set of cluster nodes and prevent other workloads from using those cluster nodes.

Chapter 5, OpenShift GitOps

Deploy OpenShift GitOps for managing clusters and applications.

Chapter 6, OpenShift Monitoring

Troubleshoot performance and availability issues with applications and clusters.

Chapter 7, OpenShift Logging

Deploy OpenShift Logging and query log entries from workloads and cluster nodes.

▶ Lab

Cluster Administration

Use Red Hat OpenShift GitOps for cluster administration.

Configure an OIDC identity provider.

Configure a one-time backup with OADP and restore from it.

Configure OpenShift Logging for short-term log retention and aggregation.

Configure alert forwarding and inspect alerts.

Outcomes

- Configure Red Hat Single Sign-On (SSO) as an OIDC identity provider (IdP) for OpenShift by using Red Hat OpenShift GitOps.
- Deploy the OpenShift Logging operator and configure it to use Loki as the log store, Vector as the collector, and the OpenShift web UI for log visualization.
- Add permission for a user to read the logs from a project.
- Back up an application and restore the application to a different namespace.
- Configure alert forwarding.
- Use monitoring to identify an application problem.

As the `student` user on the `workstation` machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start comprevew-review1
```

Specifications

- Use GitOps to configure the cluster to use Red Hat SSO for authentication. Red Hat SSO provides the `filipmansur` user, with the `redhat_sso` password, in the `ocp_rhsso` client. This user is part of the `etherpad-devs` group.

Use the following parameters to configure Red Hat SSO in OpenShift:

Attribute	Value
Name	RHSSO_OIDC
Client ID	ocp_rhsso
Client secret	QGEP6zoLo6BUGbib5oCkGwtZ8EAlmMgW

- To use OpenShift GitOps, you can use the `admin` OpenShift user with `redhatocp` as the password. The `admin` user is part of the `ocpadmins`, so grant administrator rights in Argo

Chapter 8 | Comprehensive Review

CD to users the `ocpadmins` group. The operator adds a link to the default instance in the application menu of the OpenShift console.

- The classroom has a GitLab instance that you can use to create any necessary Git repositories. GitLab is available at the `https://git.ocp4.example.com` URL. You can use the `developer` user with `d3v3lop3r` as the password. The lab scripts expect a `comprevew-review1` repository for cleanup.

The lab scripts configure the username, email, and authentication for Git in the `workstation` machine.

- Argo CD accesses only trusted repositories. GitLab uses a certificate that is signed by the classroom CA. This CA is included in the certificates that are trusted by the cluster. You can use the `config.openshift.io/inject-trusted-cabundle` label to inject the cluster trusted certificates into a configuration map, and then configure Argo CD to trust the certificate. The injected certificate is in the `ca-bundle.crt` file in the configuration map, and Argo CD uses the `/etc/pki/ca-trust/extracted/pem/tls-ca-bundle.pem` for trusted certificates in the repository server container.
- The `~/D0380/labs/comprevew-review1/sso_config.yaml` file contains an incomplete authentication configuration.
- The lab scripts deploy Etherpad in the `comprevew-review1` namespace. The Etherpad resources have the `app.kubernetes.io/name` label with `etherpad` as the value, and the supporting database resources have the `app.kubernetes.io/name` label with `mariadb` as the value.
 - Create an `etherpad-backup` backup schedule. The `~/D0380/labs/comprevew-review1/schedule-db-backup.yaml` file contains an example to create the schedule.
 - You can define an alias to access the `velero` binary by using the following command:

```
[user@host ~]$ alias velero='\\
oc -n openshift-adp exec deployment/velero -c velero -it -- ./velero'
```

- Trigger an immediate backup from the schedule, and restore the backup to the `etherpad-test-restore` namespace.
- Configure the deployed OpenShift Logging operator for short-term log retention and aggregation. The Loki operator is deployed in the cluster.
 - An S3 bucket is available for you, in the lab environment, to configure as log storage for Loki. The bucket information and credentials are available in the `~/D0380/labs/comprevew-review1/s3bucket.env` file on the `workstation` machine.
 - Create a LokiStack instance with the `1x.demo` size. Use `logging-loki` as the name for the LokiStack resource. You can find an incomplete example for the resource in the `~/D0380/labs/comprevew-review1/logging/lokistack.yaml` file.
 - Use Loki as the log store, Vector as the collector, and the OpenShift web UI for log visualization. You can find an incomplete example for the resource in the `~/D0380/labs/comprevew-review1/logging/clusterlogging.yaml` file.
 - Configure the Logging operator to include the audit logs, by using the `ClusterLogForwarder` resource. You can find an incomplete example for the resource in the `~/D0380/labs/comprevew-review1/logging/forwarder.yaml` file.

- Users in the `etherpad-devs` group have read and write permissions on the `comprevew-review1` namespace. Apply the necessary permissions so users in that group also have access to the application logs for that namespace. Red Hat SSO provides the `filipmansur` user from the `etherpad-devs` group. You can find an incomplete example for the role binding in the `~/D0380/labs/comprevew-review1/logging/group-role.yaml` file.
- Configure monitoring to send alerts by using a webhook.
 - The lab scripts deploy a webhook debugger service to the `utility` machine. This debugger starts a web server on port 8000 that prints the received payloads in the `/home/student/persistent_alerts` file in the `utility` machine.
 - OpenShift can send webhooks to this debugger at the `http://utility.lab.example.com:8000` URL.
 - This exercise generates alerts with `alertname` labels that start with the `Persistent` text. You can use the `alertname=~Persistent.*` regular expression filter so that the debugger prints only alerts for this exercise.
 - You can reduce the group and repeat intervals of the Alertmanager configuration to receive alerts sooner.
 - The debugger registers successful receipt of alerts for grading.
- Both the original Etherpad deployment and the restore deployment trigger critical monitoring alerts. Review the two critical alerts that the Etherpad deployments fire, and solve them.

Evaluation

As the `student` user on the `workstation` machine, use the `lab` command to grade your work. Correct any reported failures and rerun the command until successful.

```
[student@workstation ~]$ lab grade comprevew-review1
```

Finish

As the `student` user on the `workstation` machine, use the `lab` command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish comprevew-review1
```

► Solution

Cluster Administration

Use Red Hat OpenShift GitOps for cluster administration.

Configure an OIDC identity provider.

Configure a one-time backup with OADP and restore from it.

Configure OpenShift Logging for short-term log retention and aggregation.

Configure alert forwarding and inspect alerts.

Outcomes

- Configure Red Hat Single Sign-On (SSO) as an OIDC identity provider (IdP) for OpenShift by using Red Hat OpenShift GitOps.
- Deploy the OpenShift Logging operator and configure it to use Loki as the log store, Vector as the collector, and the OpenShift web UI for log visualization.
- Add permission for a user to read the logs from a project.
- Back up an application and restore the application to a different namespace.
- Configure alert forwarding.
- Use monitoring to identify an application problem.

As the student user on the workstation machine, use the `lab` command to prepare your environment for this exercise, and to ensure that all required resources are available.

```
[student@workstation ~]$ lab start compreview-review1
```

1. Install the OpenShift GitOps operator.
 - 1.1. Open a web browser and navigate to <https://console-openshift-console.apps.ocp4.example.com>.
 - 1.2. Click **Red Hat Identity Management** and log in as the **admin** user with **redhatocp** as the password.
 - 1.3. Navigate to **Operators > OperatorHub**.
 - 1.4. Click **Red Hat OpenShift GitOps**, and then click **Install**.
 - 1.5. Review the default configuration and click **Install**. The OLM can take a few minutes to install the operator. Click **View Operator** to navigate to the operator details.
2. Configure the default Argo CD instance to trust the classroom certificate to access repositories. Argo CD accesses only trusted repositories. You can use the

Chapter 8 | Comprehensive Review

`config.openshift.io/inject-trusted-cabundle` label to inject the classroom certificate into a configuration map, and then configure Argo CD to trust the certificate. The injected certificate is in the `ca-bundle.crt` file in the configuration map, and Argo CD uses the `/etc/pki/ca-trust/extracted/pem/tls-ca-bundle.pem` path for trusted certificates in the repository server container.

Grant administrator rights to users in the `ocpadmins` group.

- 2.1. Use the terminal to log in to the OpenShift cluster as the `admin` user with `redhatocp` as the password.

```
[student@workstation ~]$ oc login -u admin -p redhatocp \
https://api.ocp4.example.com:6443
...output omitted...
```

- 2.2. Create a `cluster-root-ca-bundle` configuration map in the `openshift-gitops` namespace.

```
[student@workstation ~]$ oc create configmap -n openshift-gitops \
cluster-root-ca-bundle
```

- 2.3. Add the `config.openshift.io/inject-trusted-cabundle` label to the configuration map with the `true` value. OpenShift injects the cluster certificates into a configuration map with this label. This bundle contains the signing certificate for the classroom GitLab instance.

```
[student@workstation ~]$ oc label configmap -n openshift-gitops \
cluster-root-ca-bundle config.openshift.io/inject-trusted-cabundle=true
configmap/cluster-root-ca-bundle labeled
```

- 2.4. Edit the Argo CD default instance to inject the certificates.

You can use the following command to edit the resource:

```
[student@workstation ~]$ oc edit argocd -n openshift-gitops openshift-gitops
```

Edit the resource to mount the `ca-bundle.crt` file from the configuration map in the `/etc/pki/ca-trust/extracted/pem/tls-ca-bundle.pem` path of the repository server container.

```
...output omitted...
spec:
...output omitted...
repo:
  resources:
    limits:
      cpu: "1"
      memory: 1Gi
    requests:
      cpu: 250m
      memory: 256Mi
  volumeMounts:
    - mountPath: /etc/pki/ca-trust/extracted/pem/tls-ca-bundle.pem
      name: cluster-root-ca-bundle
```

Chapter 8 | Comprehensive Review

```
subPath: ca-bundle.crt
volumes:
- configMap:
  name: cluster-root-ca-bundle
  name: cluster-root-ca-bundle
resourceExclusions: |
...output omitted...
```

- 2.5. Edit the Argo CD default instance to grant administrator rights to the `ocpadmins` group.

```
...output omitted...
spec:
...output omitted...
prometheus:
  enabled: false
  ingress:
    enabled: false
  route:
    enabled: false
rbac:
  defaultPolicy: ""
  policy: |
    g, system:cluster-admins, role:admin
    g, cluster-admins, role:admin
    g, ocpadmins, role:admin
  scopes: '[groups]'
redis:
...output omitted...
```

3. Create a `comprevew-review1` public repository for the authentication configuration in the classroom GitLab at <https://git.ocp4.example.com>. Use the developer GitLab user with `d3v3lop3r` as the password.
- 3.1. Open a web browser and navigate to <https://git.ocp4.example.com>. Log in as the developer user with `d3v3lop3r` as the password.
 - 3.2. Click **New project**, and then click **Create blank project**. Use `comprevew-review1` as the project slug (repository name), select the **Public** visibility level, and use the default values for all other fields. Click **Create project**.
4. Populate the repository with the OAuth CR file so OpenShift synchronizes users from the Red Hat SSO OIDC client.
- 4.1. Click **Clone**, and then copy the <https://git.ocp4.example.com/developer/comprevew-review1.git> URL.
 - 4.2. In the terminal window, change to the `~/D0380/labs/comprevew-review1` directory.

```
[student@workstation ~]$ cd ~/D0380/labs/comprevew-review1
```

- 4.3. In a terminal, run the following command to clone the new repository.

Chapter 8 | Comprehensive Review

```
[student@workstation compreview-review1]$ git clone \
https://git.ocp4.example.com/developer/compreview-review1.git
Cloning into 'compreview-review1'...
...output omitted...
```

4.4. Change to the cloned repository directory.

```
[student@workstation compreview-review1]$ cd compreview-review1
```

4.5. Create the `rhsso-oidc-client-secret` OpenShift secret for the Red Hat SSO client secret by using the client secret from Red Hat SSO parameters in the table.

```
[student@workstation compreview-review1]$ oc create secret generic \
rhsso-oidc-client-secret \
--from-literal clientSecret=QGEP6zoLo6BUGbib5oCkGwtZ8EA1mMgW \
-n openshift-config
secret/rhsso-oidc-client-secret created
```

4.6. Create the OAuth CR YAML file. You can find an example for the CR in the `~/DO380/labs/compreview-review1/sso_config.yaml` file. The YAML file includes an LDAP IdP that you must preserve, because it provides the `admin` and `developer` users. Do not remove the LDAP IdP, and add the OIDC IdP for Red Hat SSO.

```
apiVersion: config.openshift.io/v1
kind: OAuth
metadata:
  name: cluster
  annotations:
    argocd.argoproj.io/sync-options: ServerSideApply=true,Validate=false
spec:
  identityProviders:
    - ldap:
      ...output omitted...
    - openID:
        claims:
          email:
            - email
        name:
          - name
        preferredUsername:
          - preferred_username
      groups:
        - groups
    clientID: ocp_rhsso
    clientSecret:
      name: rhsso-oidc-client-secret
    extraScopes: []
    issuer: >-
      https://sso.ocp4.example.com:8080/auth/realms/internal_devs
```

Chapter 8 | Comprehensive Review

```
mappingMethod: claim
name: RHSSO_OIDC
type: OpenID
```

- 4.7. Copy the modified `sso_config.yml` file to the repository.

```
[student@workstation compreview-review1]$ cp ./sso_config.yml .
```

- 4.8. Add the `sso_config.yml` file to the Git index.

```
[student@workstation compreview-review1]$ git add sso_config.yml
```

- 4.9. Commit the changes.

```
[student@workstation compreview-review1]$ git commit -m "Add SSO to OAuth"
[main f3c0ef1] Add SSO to OAuth
...output omitted...
```

- 4.10. Push the changes to the repository. The `developer` user is configured in the lab scripts as the default Git user.

```
[student@workstation compreview-review1]$ git push
...output omitted...
```

- 4.11. Change to the `~/DO380/labs/compreview-review1` directory.

```
[student@workstation compreview-review1]$ cd ..
```

- 5.** Log in to the default Argo CD instance with the `admin` user.

- 5.1. Open a separate tab and open the default Argo CD instance. You can use the application menu to access this URL, by clicking **Cluster Argo CD** in the application menu, or use the `https://openshift-gitops-server-openshift-gitops.apps.ocp4.example.com` URL.

The browser displays a warning because Argo CD uses a self-signed certificate. Trust the certificate. Argo CD might take a few minutes before showing the login page.

- 5.2. Click **LOG IN VIA OPENSHIFT**, and then click **Red Hat Identity Management**. Log in as the `admin` user with `redhatocp` as the password, and then allow the `user:info` permission.

- 6.** Create an Argo CD application with the repository and observe the results.

- 6.1. On the Argo CD browser tab, click **CREATE APPLICATION**.

- 6.2. Create an application with the information in the following table. Then, click **CREATE**.

Field	Value
Application Name	sso-oauth
Project Name	default
Repository URL	https://git.ocp4.example.com/developer/comprevew-review1.git
Path	.
Cluster URL	https://kubernetes.default.svc

- 6.3. Click **sso-oauth** to view the application.
- 6.4. Click **SYNC** to display the synchronization panel, and then click **SYNCHRONIZE**.
Argo CD starts synchronizing the application. After about one minute, the console shows the application as synchronized and healthy.
- 6.5. Change to the terminal window and verify the status for the OAuth pods. Wait for the OAuth pods to be redeployed. It can take a few minutes for OpenShift to redeploy the pods.

```
[student@workstation comprevew-review1]$ watch oc get pods \
-n openshift-authentication
Every 2.0s: oc get pods -n openshift-authentication  workstation: Thu Feb  1
06:11:52 2024

NAME                      READY   STATUS    RESTARTS   AGE
oauth-openshift-69d79b5598-85knt  1/1     Running   0          85s
oauth-openshift-69d79b5598-q2nwk  1/1     Running   0          58s
oauth-openshift-69d79b5598-sj2fj  1/1     Running   0          114s
^C
```

- 6.6. Verify that you can log in to the cluster as the **filipmansur** user with **redhat_sso** as the password.

```
[student@workstation comprevew-review1]$ oc login -u filipmansur -p redhat_sso
Login successful.
...output omitted...
```

- 6.7. Log in to the OpenShift cluster as the **admin** user.

```
[student@workstation comprevew-review1]$ oc login -u admin -p redhatocp
...output omitted...
```

7. Create an etherpad-backup backup schedule of Etherpad. The `~/D0380/labs/comprevew-review1/schedule-db-backup.yml` file contains an example to create the backup definition. The backup must include resources with the `app.kubernetes.io/name` label with `etherpad` as the value, and the `app` label with `mariadb` as the value. The backup must also include the required resources to preserve the UID and GID of the namespace for OADP to restore, and to ensure that the backup hooks are executed during the backup.

Chapter 8 | Comprehensive Review

- 7.1. Edit the example to match the following text.

```
apiVersion: velero.io/v1
kind: Schedule
metadata:
  name: etherpad-backup
  namespace: openshift-adp
spec:
  schedule: "0 7 * * 0"
  paused: true
  template:
    ttl: 360h0m0s
    includedNamespaces:
      - comprevew-review1
    orLabelSelectors:
      - matchLabels:
          app.kubernetes.io/name: etherpad
      - matchLabels:
          app.kubernetes.io/name: mariadb
      - matchLabels:
          kubernetes.io/metadata.name: comprevew-review1
    includedResources:
      - deployment
      - route
      - service
      - pvc
      - persistentvolume
      - secret
      - service
      - namespace
      - pods
    hooks:
...output omitted...
```

- 7.2. Apply the configuration for the schedule resource.

```
[student@workstation comprevew-review1]$ oc apply -f schedule-db-backup.yml
schedule.velero.io/etherpad-backup created
```

- 7.3. Create an alias to access the `velero` binary from the Velero deployment in the `openshift-adp` namespace.

```
[student@workstation comprevew-review1]$ alias velero='\
  oc -n openshift-adp exec deployment/velero -c velero -it -- ./velero'
```

- 7.4. Verify the status of the schedule with the `velero` command.

```
[student@workstation comprevew-review1]$ velero get schedule
NAME           STATUS  ...  PAUSED
etherpad-backup  New    ...  true
```

8. Trigger an immediate backup from the scheduled backup, and restore it to the `etherpad-test-restore` project.

Chapter 8 | Comprehensive Review

The backup and restore should complete without errors or warnings.

- 8.1. Use the `velero` command to start a backup by using the schedule definition from the previous step. Note the name of the backup that the command creates, to use in the next step.

```
[student@workstation comprevie-review1]$ velero backup create \
--from-schedule etherpad-backup
INFO[0000] No Schedule.template.metadata.labels set - using Schedule.labels for
backup object  backup=openshift-adp/etherpad-backup-20240131113134  labels="map[]"
Creating backup from schedule, all other filters are ignored.
Backup request "etherpad-backup-20240131113134" submitted successfully.
Run velero backup describe etherpad-backup-20240131113134 or velero backup logs
etherpad-backup-20240131113134 for more details.
```

**Note**

The S3 object storage that is configured in the lab environment uses a custom certificate that is signed with the OpenShift service CA. You must add the CA certificate to the `velero backup describe --details` and `velero backup logs` commands as follows:

```
[user@host]$ velero backup logs \
--cacert=/run/secrets/kubernetes.io/serviceaccount/service-ca.crt \
etherpad-backup-20231115113447
```

- 8.2. Monitor the status of the backup and verify that the backup ends with the **Completed** status. The backup process takes several minutes.

```
[student@workstation comprevie-review1]$ velero get backup
NAME                  STATUS     ERRORS   WARNINGS   ...
etherpad-backup-20231115113447  Completed   0         0          ...
```

- 8.3. Restore the backup to the `etherpad-test-restore` namespace.

```
[student@workstation comprevie-review1]$ velero restore create etherpad-test \
--from-backup etherpad-backup-20240129143859 \
--namespace-mappings comprevie-review1:etherpad-test-restore
Restore request "etherpad-test" submitted successfully.
Run velero restore describe etherpad-test or velero restore logs etherpad-test for
more details.
```

- 8.4. Use the `velero` command to get the status of the restore. Monitor the output to verify that the restore status is **Completed**. The restore process takes several minutes.

```
[student@workstation comprevie-review1]$ velero get restore
NAME      ...  STATUS      ...
etherpad-test  ...  Completed  ...
```

- 8.5. Review the restored resources in the `etherpad-test-restore` project.

Chapter 8 | Comprehensive Review

The pod for the etherpad deployment requires more time to be ready, because the pod requires the database deployment to be ready first.

```
[student@workstation compreview-review1]$ oc get -n etherpad-test-restore \
pod,deployment,route
NAME                                READY   STATUS    RESTARTS   AGE
pod/etherpad-58f64cfb7d-8qrws      1/1     Running   3 (44s ago)  84s
pod/mariadb-66dc48b5f7-svlst      1/1     Running   0          84s

NAME                  READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/etherpad  1/1       1           1           84s
deployment.apps/mariadb   1/1       1           1           84s

NAME   HOST/PORT
...    etherpad-etherpad-test-restore.apps.ocp4.example.com  ...
```

- 8.6. Visit the <https://etherpad-etherpad-test-restore.apps.ocp4.example.com> URL to verify that the restored Etherpad works.
9. Use the bucket credentials to create the `logging-loki-odf` secret in the `openshift-logging` namespace. Create a `LokiStack` resource YAML file for an instance called `logging-loki` in the `openshift-logging` namespace. Create a `ClusterLogging` resource YAML file by using the `loki` log store, the `vector` collector, and the `ocp-console` visualization type.
- 9.1. Use the `~/D0380/labs/compreview-review1/s3bucket.env` environment file to create the `logging-loki-odf` secret in the `openshift-logging` namespace.

```
[student@workstation compreview-review1]$ oc create secret generic \
logging-loki-odf -n openshift-logging --from-env-file=s3bucket.env
secret/logging-loki-odf created
```

- 9.2. Create a `LokiStack` resource YAML file for an instance called `logging-loki` in the `openshift-logging` namespace. This instance uses the bucket as the storage. You can find an incomplete example for the resource in the `~/D0380/labs/compreview-review1/logging/lokistack.yaml` file.

```
apiVersion: loki.grafana.com/v1
kind: LokiStack
metadata:
  name: logging-loki
  namespace: openshift-logging
spec:
  size: 1x.demo
  storage:
    secret:
      name: logging-loki-odf
      type: s3
    tls:
      caName: openshift-service-ca.crt
  storageClassName: ocs-external-storagecluster-ceph-rbd
  tenants:
    mode: openshift-logging
```

Chapter 8 | Comprehensive Review

9.3. Create the LokiStack resource.

```
[student@workstation compreview-review1]$ oc create -f logging/lokistack.yaml
lokistack.loki.grafana.com/logging-loki created
```

9.4. Create a ClusterLogging resource YAML file by using the loki log store, the vector collector, and the ocp-console visualization type. You can find an incomplete example for the resource in the ~/D0380/labs/compreview-review1/logging/clusterlogging.yaml file.

```
apiVersion: logging.openshift.io/v1
kind: ClusterLogging
metadata:
  name: instance
  namespace: openshift-logging
spec:
  managementState: Managed
  logStore:
    type: lokistack
    lokistack:
      name: logging-loki
  collection:
    type: vector
  visualization:
    type: ocp-console
```

9.5. Create the ClusterLogging resource.

```
[student@workstation compreview-review1]$ oc create -f logging/clusterlogging.yaml
clusterlogging.logging.openshift.io/instance created
```

9.6. Verify that the ClusterLogging and LokiStack pods are up and running.

```
[student@workstation compreview-review1]$ oc get pods -n openshift-logging
NAME                               READY   STATUS    RESTARTS   AGE
cluster-logging-operator-554849f7dd-75hjk   1/1     Running   0          6m7s
collector-2dqsh                      1/1     Running   0          18s
collector-5fp2q                      1/1     Running   0          19s
collector-c6zjv                      1/1     Running   0          19s
collector-gggtc                      1/1     Running   0          13s
collector-k5zs8                      1/1     Running   0          13s
collector-wgz7q                      1/1     Running   0          12s
logging-loki-compactor-0              1/1     Running   0          83s
logging-loki-distributor-c55478c4c-9kt2k  1/1     Running   0          83s
logging-loki-gateway-75d6fccb68-krk49   2/2     Running   0          82s
logging-loki-gateway-75d6fccb68-vghdg   2/2     Running   0          82s
logging-loki-index-gateway-0           1/1     Running   0          82s
logging-loki-ingester-0               1/1     Running   0          83s
logging-loki-querier-6f7d8b7564-4glpp  1/1     Running   0          83s
logging-loki-query-frontend-678dddf864-947hn 1/1     Running   0          83s
logging-view-plugin-5b9b5b7bdc-zrp9t   1/1     Running   0          35s
```

Chapter 8 | Comprehensive Review

10. Enable the console plug-in for the OpenShift Logging operator. Verify that you have access to the logs.
 - 10.1. Change to the web console browser. Click **Operators > Installed Operators**, and select **All Projects** from the drop-down menu.
 - 10.2. Click **Red Hat OpenShift Logging**, click **Console plugin**, select **Enable**, and click **Save**.
 - 10.3. Reload the web console, and navigate to **Observe > Logs**. If the **Observe > Logs** menu is not available, then wait until the web console shows the **Web console update is available** message and reload the web console. You have access to logs for the application and infrastructure resources. By default, the **ClusterLogging** instance includes logs for the application and infrastructure, but not the audit logs. Observe the application logs, which are selected by default.
11. Include the audit logs by creating a log forwarder for the application, infrastructure, and audit logs to the **LokiStack** resource.
 - 11.1. Change to the terminal window, and create an **ClusterLogForwarder** resource YAML file for a log forwarder called **instance** in the **openshift-logging** namespace. The log forwarder must forward the application, infrastructure, and audit logs. You can find an incomplete example for the resource in the **~/D0380/labs/comprevew-review1/logging/forwarder.yaml** file.

```
apiVersion: logging.openshift.io/v1
kind: ClusterLogForwarder
metadata:
  name: instance
  namespace: openshift-logging
spec:
  pipelines:
    - name: all-to-default
      inputRefs:
        - infrastructure
        - application
        - audit
      outputRefs:
        - default
```

- 11.2. Create the **ClusterLogForwarder** resource.

```
[student@workstation compreview-review1]$ oc create -f logging/forwarder.yaml
clusterlogforwarder.logging.openshift.io/instance created
```

- 11.3. Change to the web console browser and reload it. You have access to the audit logs.
12. Give the **filipmansur** user permission to view the logs in the **comprevew-review1** project, and verify that the user has access to the logs. Give the permission to the **filipmansur** user by assigning the **cluster-logging-application-view** cluster role through the **etherpad-devs** group.
 - 12.1. Change to the terminal window.
 - 12.2. Review the required role to provide access to the application logs to the **filipmansur** user. You can find an example in the **~/D0380/labs/comprevew-review1/logging/group-role.yaml** file.

Chapter 8 | Comprehensive Review

```

apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: view-application-logs
  namespace: compreview-review1
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: cluster-logging-application-view
subjects:
- kind: Group
  name: etherpad-devs
  apiGroup: rbac.authorization.k8s.io

```

- 12.3. Apply the role to the filipmansur user.

```
[student@workstation compreview-review1]$ oc create -f \
logging/group-role.yaml
rolebinding.rbac.authorization.k8s.io/view-application-logs created
```

- 12.4. Change to the browser window, open a private window, and navigate to <https://console-openshift-console.apps.ocp4.example.com>.
 - 12.5. Click **RHSSO_OIDC** and log in as the **filipmansur** user with **redhat_sso** as the password. Click **Skip tour**.
 - 12.6. Navigate to **Observe**. Verify that the **compreview-review1** project is selected. Change to the **Logs** tab. The **filipmansur** user has access to the application logs.
- 13.** Configure alerts.
- 13.1. In the OpenShift web console, change to the non-private window. Navigate to **Administration > Cluster Settings**, click the **Configuration** tab, and then click **Alertmanager**.
 - 13.2. In the **Alert routing** tile, click **Edit**. Change the **Group interval** and **Repeat interval** fields to **1m**.
 - 13.3. In the **Receivers** tile, click **Create Receiver**. Create a receiver with the information in the following table:

Field	Value
Receiver name	persistent
Receiver type	Webhook
URL	http://utility.lab.example.com:8000

Specify the **alertname=~Persistent.*** routing label.

Then, click **Create**.

- 14.** Review that monitoring shows that the **mariadb** persistent volume claims are nearly full.

Chapter 8 | Comprehensive Review

- 14.1. In the OpenShift web console, navigate to **Observe > Alerting**. The console should list the **PersistentVolumeUsageCritical** and **PersistentVolumeUsageNearFull** alerts. Besides the original Etherpad deployment, you created a second deployment with a restore. Based on the disk usage in each deployment, the number and type of alerts can vary.
- 14.2. Change to the terminal window and connect to the **utility** machine with SSH as the **student** user.

```
[student@workstation compreview-review1]$ ssh utility  
...output omitted...
```

- 14.3. After approximately one minute after the receiver is created, the webhook debugger shows the alerts in the `/home/student/persistent_alerts` file.

```
[student@utility ~]$ head persistent_alerts  
{'alerts': [ {'annotations': {'description': 'PVC mariadb utilization has '  
           'crossed 75%. Free up some space '  
           'or expand the PVC.',  
           'message': 'PVC mariadb is nearing full. Data '  
           'deletion or PVC expansion is '  
           'required.',  
           'severity_level': 'warning',  
           'storage_type': 'ceph'},  
           'endsAt': '0001-01-01T00:00:00Z',  
           'fingerprint': 'd051e03da5866d5b',  
...output omitted...
```

- 14.4. Disconnect from the utility machine.

```
[student@utility ~]$ exit  
logout  
Connection to utility closed.  
[student@workstation compreview-review1]$
```

- 14.5. Change to the student HOME directory.

```
[student@workstation compreview-review1]$ cd  
[student@workstation ~]$
```

15. Expand the PVCs.

- 15.1. Navigate to **Storage > PersistentVolumeClaims**. Select **All Projects** from the project drop-down menu. Locate the two **mariadb** persistent volume claims.
- 15.2. For each claim, click its name to view the details. Each claim has 190 MiB capacity and about 40 MiB available.
For each claim, select **Expand PVC** from the **Actions** list. Edit the total size of the claim to 1900 MiB, and then click **Expand**.
- 15.3. If you navigate to **Observe > Alerting**, then the alerts disappear after a few minutes.

Evaluation

As the **student** user on the **workstation** machine, use the **lab** command to grade your work. Correct any reported failures and rerun the command until successful.

```
[student@workstation ~]$ lab grade compreview-review1
```

Finish

As the **student** user on the **workstation** machine, use the **lab** command to complete this exercise. This step is important to ensure that resources from previous exercises do not impact upcoming exercises.

```
[student@workstation ~]$ lab finish compreview-review1
```

